

Data Science for Biological, Medical and Health Research: Notes for 432

Thomas E. Love, Ph.D.

Built 2022-02-06 16:43:55

Contents

Introduction	9
R Packages used in these notes	11
Dealing with Conflicts	12
General Theme for <code>ggplot</code> work	13
Data used in these notes	13
1 Building Table 1	15
1.1 Data load	15
1.2 Two examples from the <i>New England Journal of Medicine</i>	16
1.3 The MR CLEAN trial	17
1.4 Simulated <code>fakestroke</code> data	19
1.5 Building Table 1 for <code>fakestroke</code> : Attempt 1	20
1.6 <code>fakestroke</code> Table 1: Attempt 2	22
1.7 Obtaining a more detailed Summary	25
1.8 Exporting the Completed Table 1 from R to Excel or Word	29
1.9 A Controlled Biological Experiment - The Blood-Brain Barrier .	31
1.10 The <code>bloodbrain.csv</code> file	31
1.11 A Table 1 for <code>bloodbrain</code>	32
2 BRFSS SMART Data Building	39
2.1 Key resources	39
2.2 Ingesting the Raw Data	40
2.3 Ingesting from our CSV file	41
2.4 What does the raw data look like?	41
2.5 Cleaning the BRFSS Data	42
2.6 Imputing Age and Income as Quantitative from Thin Air	105
2.7 Clean Data in the State of Ohio	108
2.8 Clean Cleveland-Elyria Data	109
3 Dealing with Missingness: Single Imputation	111
3.1 Selecting Some Variables from the <code>smart_cle</code> data	111
3.2 <code>smart_cle1</code> : Seeing our Missing Data	112
3.3 Missing-data mechanisms	117

3.4	Options for Dealing with Missingness	118
3.5	Complete Case (and Available Case) analyses	118
3.6	Single Imputation	119
3.7	Multiple Imputation	119
3.8	Approach 1: Building a Complete Case Analysis: <code>smart_cle1_cc</code>	119
3.9	Approach 2: Single Imputation to create <code>smart_cle1_sh</code>	120
4	Summarizing the <code>smart_cle1</code> data	133
4.1	General Approaches to Obtaining Numeric Summaries	134
4.2	Counting as exploratory data analysis	137
4.3	Can <code>bmi</code> predict <code>physhealth</code> ?	143
5	Analysis of Variance with SMART	157
5.1	A One-Factor Analysis of Variance	157
5.2	A Two-Factor ANOVA (without Interaction)	165
5.3	A Two-Factor ANOVA (with Interaction)	169
6	Analysis of Variance	177
6.1	The <code>bonding</code> data: A Designed Dental Experiment	177
6.2	A One-Factor Analysis of Variance	178
6.3	A Two-Way ANOVA: Looking at Two Factors	182
6.4	A Means Plot (with standard deviations) to check for interaction	183
6.5	Fitting the Two-Way ANOVA model with Interaction	185
6.6	Comparing Individual Combinations of <code>resin</code> and <code>light</code>	188
6.7	The <code>bonding</code> model without Interaction	189
6.8	<code>cortisol</code> : A Hypothetical Clinical Trial	191
6.9	Creating a factor combining <code>sex</code> and <code>waist</code>	192
6.10	A Means Plot for the <code>cortisol</code> trial (with standard errors)	193
6.11	A Two-Way ANOVA model for <code>cortisol</code> with Interaction	194
6.12	A Two-Way ANOVA model for <code>cortisol</code> without Interaction	198
7	Analysis of Covariance	201
7.1	An Emphysema Study	201
7.2	Does <code>sex</code> affect the mean change in <code>theophylline</code> ?	202
7.3	Is there an association between <code>age</code> and <code>sex</code> in this study?	203
7.4	Adding a quantitative covariate, <code>age</code> , to the model	203
7.5	Rerunning the ANCOVA model after simple imputation	205
7.6	Looking at a factor-covariate interaction	206
7.7	Centering the Covariate to Facilitate ANCOVA Interpretation	207
8	Analysis of Covariance with the SMART data	209
8.1	A New Small Study: Predicting BMI	209
8.2	<code>c8_m1</code> : A simple t-test model	211
8.3	<code>c8_m2</code> : Adding another predictor (two-way ANOVA without interaction)	213

8.4	c8_m3: Adding the interaction term (Two-way ANOVA with interaction)	216
8.5	c8_m4: Using <code>female</code> and <code>physhealth</code> in a model for <code>bmi</code>	218
8.6	Making Predictions with a Linear Regression Model	221
8.7	Centering the model	222
8.8	Rescaling an input by subtracting the mean and dividing by 2 standard deviations	224
8.9	c8_m5: What if we add more variables?	227
8.10	c8_m6: Would adding self-reported health help?	229
8.11	Key Regression Assumptions for Building Effective Prediction Models	230
9	Adding Non-linear Terms to a Linear Regression Model	237
9.1	The <code>pollution</code> data	237
9.2	Fitting a straight line model to predict <code>y</code> from <code>x2</code>	238
9.3	Quadratic polynomial model to predict <code>y</code> using <code>x2</code>	240
9.4	Orthogonal Polynomials	244
9.5	Fit a cubic polynomial to predict <code>y</code> from <code>x3</code>	248
9.6	Fitting a restricted cubic spline in a linear regression	251
9.7	“Spending” Degrees of Freedom	256
9.8	Spending DF on Non-Linearity: The Spearman Plot	259
10	Using <code>ols</code> from the <code>rms</code> package to fit linear models	263
10.1	Fitting a model with <code>ols</code>	263
10.2	ANOVA for an <code>ols</code> model	265
10.3	Effect Estimates	266
10.4	The <code>Predict</code> function for an <code>ols</code> model	268
10.5	Checking Influence via <code>dfbeta</code>	270
10.6	Model Validation and Correcting for Optimism	273
10.7	Building a Nomogram for Our Model	274
11	Logistic Regression: The Foundations	277
11.1	A First Attempt: A Linear Probability Model	277
11.2	Logistic Regression	279
11.3	The Logistic Regression Model	280
11.4	The Link Function	280
11.5	The logit or log odds	281
11.6	Interpreting the Coefficients of a Logistic Regression Model	281
11.7	The Logistic Regression has non-constant variance	282
11.8	Fitting a Logistic Regression Model to our Simulated Data	282
11.9	Plotting the Logistic Regression Model	284
12	Logistic Regression and the <code>resect</code> data	287
12.1	The <code>resect</code> data	287
12.2	Running A Simple Logistic Regression Model	288
12.3	Logistic Regression using <code>glm</code>	289

12.4 Interpreting the Model Summary	295
12.5 Plotting a Simple Logistic Regression Model	298
12.6 How well does Model A classify subjects?	300
12.7 The Confusion Matrix	301
12.8 Using the <code>confusionMatrix</code> tool from the <code>caret</code> package	302
12.9 Receiver Operating Characteristic Curve Analysis	303
12.10 The ROC Plot for <code>res_modA</code>	310
12.11 Assessing Residual Plots from Model A	312
12.12 Model B: A “Kitchen Sink” Logistic Regression Model	314
12.13 Plotting Model B	316
12.14 Logistic Regression using <code>lrm</code>	320
12.15 Model D: An Augmented Kitchen Sink Model	329
12.16 Model E: Fitting a Reduced Model in light of Model D	337
12.17 Concordance: Comparing Model C, D and E’s predictions	343
12.18 Conclusions	346
13 A Study of Prostate Cancer	347
13.1 Data Load and Background	347
13.2 Code Book	348
13.3 Additions for Later Use	349
13.4 Fitting and Evaluating a Two-Predictor Model	350
13.5 Exploring Model <code>c11_prost_A</code>	351
13.6 Plotting Model <code>c11_prost_A</code>	356
13.7 Cross-Validation of Model <code>c11_prost_A</code>	359
14 Multiple Imputation and Linear Regression	367
14.1 Developing a <code>smart_16</code> data set	367
14.2 Obtaining a Simple Imputation with <code>mice</code>	369
14.3 Linear Regression: Considering a Transformation of the Outcome	370
14.4 Linear Regression: Considering Non-Linearity in the Predictors .	372
14.5 “Main Effects” Linear Regression with <code>lm</code> on the Complete Cases	373
14.6 “Augmented” Linear Regression with <code>lm</code> on the Complete Cases .	376
14.7 Using <code>mice</code> to perform Multiple Imputation	382
14.8 Running the Linear Regression in <code>lm</code> with Multiple Imputation .	384
14.9 Fit the Multiple Imputation Model with <code>aregImpute</code>	386
14.10 Fit Linear Regression using <code>ols</code> and <code>fit.mult.impute</code>	390
15 Using <code>tidymodels</code> to fit linear regressions	397
16 Using <code>tidymodels</code> to fit logistic regressions	399
17 Using <code>tidymodels</code> approaches: Next Steps	401
18 Colorectal Cancer Screening and Some Special Cases	403
18.1 Logistic Regression for Aggregated Data	403
18.2 Probit Regression	408

19 Modeling a Count Outcome in Ohio SMART	415
19.1 Preliminaries	415
19.2 A subset of the Ohio SMART data	416
19.3 Exploratory Data Analysis (in the 18-49 group)	419
19.4 Modeling Strategies Explored Here	422
19.5 The OLS Approach	423
19.6 OLS model on $\log(\text{physhealth} + 1)$ days	428
19.7 A Poisson Regression Model	433
19.8 Overdispersion in a Poisson Model	442
19.9 Fitting the Quasi-Poisson Model	443
19.10 Poisson and Quasi-Poisson models using <code>Glm</code> from the <code>rms</code> package	449
19.11 Negative Binomial Model	453
19.12 The Problem: Too Few Zeros	463
19.13 The Zero-Inflated Poisson Regression Model	464
19.14 The Zero-Inflated Negative Binomial Regression Model	471
19.15 A “hurdle” model (with Poisson)	478
19.16 A “hurdle” model (with negative binomial for overdispersion)	487
19.17 A Tobit (Censored) Regression Model	495
20 Modeling an Ordinal Categorical Outcome in Ohio SMART	503
20.1 Preliminaries	503
20.2 A subset of the Ohio SMART data	503
20.3 Building Cross-Tabulations	505
20.4 Graphing Categorical Data	510
20.5 Building a Model for <code>genh</code> using <code>veg_day</code>	512
20.6 Interpreting Model <code>m1</code>	516
20.7 The Assumption of Proportional Odds	520
20.8 Can model <code>m1</code> be fit using <code>rms</code> tools?	523
20.9 Building a Three-Predictor Model	524
20.10 A Larger Model, including income group	533
20.11 References for this Chapter	537
21 Analyzing Literary Styles with Multinomial Logistic Regression	539
21.1 The Authorship Example	539
21.2 Focus on 11 key words	540
21.3 A Multinomial Logistic Regression Model	543
21.4 Model 2	545
21.5 Classification Table	547
21.6 Probability Curves based on a Single Predictor	548
22 Exploring Time To Event / Survival Data	553
22.1 An Outline of Key Topics Discussed in these Notes	554
22.2 Foundations of Survival Analysis	554
22.3 A First Example: Recurrent Lobar Intracerebral Hemorrhage	556
22.4 Comparing Survival Across the Two Genotypes	561

22.5 Testing the difference between two survival curves	563
22.6 A “Fancy” K-M Plot with a number at risk table	564
22.7 The Hazard Function	567
23 Cox Regression Models for Survival Data: Example 1	573
23.1 Sources used in building this material	574
23.2 Fitting a Cox Model in R with <code>coxph</code>	574
23.3 Fitting a Cox Model using <code>cph</code> from the <code>rms</code> package	580
24 Cox Regression Models for Survival Data: Example 2	589
24.1 A Second Example: The <code>leukem</code> data	589
24.2 Model A: <code>coxph</code> Model for Survival Time using <code>age</code> at diagnosis	591
24.3 Building Model A with <code>cph</code> for the <code>leukem</code> data	594
24.4 Model B: Fitting a 5-Predictor Model with <code>coxph</code>	601
24.5 Model B2: A Stepwise Reduction of Model B	605
24.6 Model C: Using a Spearman Plot to pick a model	608

Introduction

These Notes provide a series of examples using R to work through issues that are likely to come up in PQHS/CRSP/MPHP 432.

While these Notes share some of the features of a textbook, they are neither comprehensive nor completely original. The main purpose is to give students in 432 a set of common materials on which to draw during the course. In class, we will sometimes:

- reiterate points made in this document,
- amplify what is here,
- simplify the presentation of things done here,
- use new examples to show some of the same techniques,
- refer to issues not mentioned in this document,

but what we don't (always) do is follow these notes very precisely. We assume instead that you will read the materials and try to learn from them, just as you will attend classes and try to learn from them. We welcome feedback of all kinds on this document or anything else via Piazza.

What you will mostly find are brief explanations of a key idea or summary, accompanied (most of the time) by R code and a demonstration of the results of applying that code.

Everything you see here is available to you as HTML or PDF. You will also have access to the R Markdown files, which contain the code which generates everything in the document, including all of the R results. We will demonstrate the use of R Markdown (this document is generated with the additional help of an R package called bookdown) and R Studio (the “program” which we use to interface with the R language) in class.

To download the data and R code related to these notes, visit the appropriate link at the 432 course website.

R Packages used in these notes

Here, we'll load in some packages used in these notes. The list of R Packages we will use in 432 is more extensive, and is available on our course website.

```
library(here)
library(janitor)
library(magrittr)
library(conflicted)

library(tableone)

library(broom)
library(haven)
library(janitor)
library(patchwork)
library(Hmisc)
library(rms)

library(MASS)
library(visdat)
library(naniar)
library(caret)
library(simputation)
library(car)
library(mice)
library(leaps)
library(lars)
library(Epi)
library(pROC)
library(ROCR)
library(VGAM)
library(ggridges)
```

```
library(pander)
library(arm)
library(survival)
library(survminer)
library(kableExtra)

## and of course, we conclude with...

library(tidymodels)
library(tidyverse)
```

Dealing with Conflicts

I'm loading a lot of packages here, and sometimes individual functions are in conflict. R's default conflict resolution system gives precedence to the most recently loaded package. This can make it hard to detect conflicts, particularly when introduced by an update to an existing package.

Using the code below helps the entire book run properly. You may or may not need to look into the `conflicted` package for your work.

```
conflict_prefer("filter", "dplyr")
```

```
[conflicted] Will prefer dplyr::filter over any other package
conflict_prefer("select", "dplyr")
```

```
[conflicted] Will prefer dplyr::select over any other package
conflict_prefer("Predict", "rms")
```

```
[conflicted] Will prefer rms::Predict over any other package
conflict_prefer("impute_median", "simputation")
```

```
[conflicted] Will prefer simputation::impute_median over any other package
conflict_prefer("summarize", "dplyr")
```

```
[conflicted] Will prefer dplyr::summarize over any other package
specify_decimal <- function(x, k) format(round(x, k), nsmall=k)
```

General Theme for ggplot work

```
theme_set(theme_bw())
```

Data used in these notes

All data sets used in these notes are available on our Data and Code website.

Dr. Love is in the process of moving all of the data loads below to their individual chapters.

```
prost <- read_csv("data/prost.csv")
pollution <- read_csv("data/pollution.csv")

bonding <- read_csv("data/bonding.csv")
cortisol <- read_csv("data/cortisol.csv")
emphysema <- read_csv("data/emphysema.csv")
resect <- read_csv("data/resect.csv")
colscr <- read_csv("data/screening.csv")
colscr2 <- read_csv("data/screening2.csv")
authorship <- read_csv("data/authorship.csv")
hem <- read_csv("data/hem.csv")
leukem <- read_csv("data/leukem.csv")
```


Chapter 1

Building Table 1

Many scientific articles involve direct comparison of results from various exposures, perhaps treatments. In 431, we studied numerous methods, including various sorts of hypothesis tests, confidence intervals, and descriptive summaries, which can help us to understand and compare outcomes in such a setting. One common approach is to present what's often called Table 1. Table 1 provides a summary of the characteristics of a sample, or of groups of samples, which is most commonly used to help understand the nature of the data being compared.

1.1 Data load

Let's load two data sets for this Chapter. All data sets used in these notes are available on our Data and Code website.

```
fakestroke <- read_csv("data/fakestroke.csv")  
  
Rows: 500 Columns: 18  
-- Column specification -----  
Delimiter: ","  
chr (8): studyid, trt, sex, location, hx.isch, mrankin, iv.altep, ia.occlus  
dbl (10): age, nihss, afib, dm, sbp, time.iv, aspects, extra.ica, time.rand,...  
  
i Use `spec()` to retrieve the full column specification for this data.  
i Specify the column types or set `show_col_types = FALSE` to quiet this message.  
  
bloodbrain <- read_csv("data/bloodbrain.csv")  
  
Rows: 34 Columns: 11  
-- Column specification -----  
Delimiter: ","
```

```

chr (2): solution, sex
dbl (9): case, brain, liver, tlratio, sactime, postin, wt.init, wt.loss, wt.....

i Use `spec()` to retrieve the full column specification for this data.
i Specify the column types or set `show_col_types = FALSE` to quiet this message.

```

1.2 Two examples from the *New England Journal of Medicine*

1.2.1 A simple Table 1

Table 1 is especially common in the context of clinical research. Consider the excerpt below, from a January 2015 article in the *New England Journal of Medicine* (Tolaney et al. 2015).

Table 1. Baseline Characteristics of the Patients.*	
Characteristic	Patients (N=406)
	<i>no. (%)</i>
Age group	
<50 yr	132 (32.5)
50–59 yr	137 (33.7)
60–69 yr	96 (23.6)
≥70 yr	41 (10.1)
Sex	
Female	405 (99.8)
Male	1 (0.2)
Race†	
White	351 (86.5)
Black	28 (6.9)
Asian	11 (2.7)
Other	16 (3.9)

This (partial) table reports baseline characteristics on age group, sex and race, describing 406 patients with HER2-positive¹ invasive breast cancer that began the protocol therapy. Age, sex and race (along with severity of illness) are the most commonly identified characteristics in a Table 1.

In addition to the measures shown in this excerpt, the full Table also includes detailed information on the primary tumor for each patient, including its size, nodal status and histologic grade. Footnotes tell us that the percentages shown are subject to rounding, and may not total 100, and that the race information was self-reported.

¹HER2 = human epidermal growth factor receptor type 2. Over-expression of this occurs in 15–20% of invasive breast cancers, and has been associated with poor outcomes.

1.2.2 A group comparison

A more typical Table 1 involves a group comparison, for example in this excerpt from Roy et al. (2008). This Table 1 describes a multi-center randomized clinical trial comparing two different approaches to caring for patients with heart failure and atrial fibrillation².

Variable	Rhythm-Control Group (N = 682)	Rate-Control Group (N = 694)
Male sex (%)	78	85
Age (yr)	66±11	67±11
Body-mass index†	27.8±5.4	28.0±5.1
Nonwhite race (%)‡	16	13
NYHA class III or IV (%)		
At baseline	32	31
During previous 6 mo	76	76
Predominant cardiac diagnosis (%)§		
Coronary artery disease	48	48
Valvular heart disease	5	5
Nonischemic cardiomyopathy	36	39
Congenital heart disease	1	1
Hypertensive heart disease	10	7

The article provides percentages, means and standard deviations across groups, but note that it does not provide p values for the comparison of baseline characteristics. This is a common feature of NEJM reports on randomized clinical trials, where we anticipate that the two groups will be well matched at baseline. Note that the patients in this study were *randomly* assigned to either the rhythm-control group or to the rate-control group, using blocked randomization stratified by study center.

1.3 The MR CLEAN trial

Berkhemer et al. (2015) reported on the MR CLEAN trial, involving 500 patients with acute ischemic stroke caused by a proximal intracranial arterial occlusion. The trial was conducted at 16 medical centers in the Netherlands, where 233 were randomly assigned to the intervention (intraarterial treatment plus usual care) and 267 to control (usual care alone.) The primary outcome was the modified Rankin scale score at 90 days; this categorical scale measures functional outcome, with scores ranging from 0 (no symptoms) to 6 (death). The fundamental conclusion of Berkhemer et al. (2015) was that in patients with acute ischemic stroke caused by a proximal intracranial occlusion of the anterior circulation, intraarterial treatment administered within 6 hours after stroke onset was effective and safe.

²The complete Table 1 appears on pages 2668-2669 of Roy et al. (2008), but I have only reproduced the first page and the footnote in this excerpt.

Here's the Table 1 from Berkhemer et al. (2015).

Table 1. Baseline Characteristics of the 500 Patients.*		
Characteristic	Intervention (N=233)	Control (N=267)
Age — yr		
Median	65.8	65.7
Interquartile range	54.5–76.0	55.5–76.4
Male sex — no. (%)	135 (57.9)	157 (58.8)
NIHSS score†		
Median (interquartile range)	17 (14–21)	18 (14–22)
Range	3–30	4–38
Location of stroke in left hemisphere — no. (%)	116 (49.8)	153 (57.3)
History of ischemic stroke — no. (%)	29 (12.4)	25 (9.4)
Atrial fibrillation — no. (%)	66 (28.3)	69 (25.8)
Diabetes mellitus — no. (%)	34 (14.6)	34 (12.7)
Prestroke modified Rankin scale score — no. (%)‡		
0	190 (81.5)	214 (80.1)
1	21 (9.0)	29 (10.9)
2	12 (5.2)	13 (4.9)
>2	10 (4.3)	11 (4.1)
Systolic blood pressure — mm Hg§	146±26.0	145±24.4
Treatment with IV alteplase — no. (%)	203 (87.1)	242 (90.6)
Time from stroke onset to start of IV alteplase — min		
Median	85	87
Interquartile range	67–110	65–116
ASPECTS — median (interquartile range)¶	9 (7–10)	9 (8–10)
Intracranial arterial occlusion — no./total no. (%)		
Intracranial ICA	1/233 (0.4)	3/266 (1.1)
ICA with involvement of the M1 middle cerebral artery segment	59/233 (25.3)	75/266 (28.2)
M1 middle cerebral artery segment	154/233 (66.1)	165/266 (62.0)
M2 middle cerebral artery segment	18/233 (7.7)	21/266 (7.9)
A1 or A2 anterior cerebral artery segment	1/233 (0.4)	2/266 (0.8)
Extracranial ICA occlusion — no./total no. (%) **	75/233 (32.2)	70/266 (26.3)
Time from stroke onset to randomization — min††		
Median	204	196
Interquartile range	152–251	149–266
Time from stroke onset to groin puncture — min		
Median	260	NA
Interquartile range	210–313	

The Table was accompanied by the following notes.

* The intervention group was assigned to intraarterial treatment plus usual care, and the control group was assigned to usual care alone. Plus-minus values are means ±SD. ICA denotes internal carotid artery, IV intravenous, and NA not applicable.

† Scores on the National Institutes of Health Stroke Scale (NIHSS) range from 0 to 42, with higher scores indicating more severe neurologic deficits. The NIHSS is a 15-item scale, and values for 30 of the 7500 items were missing (0.4%). The highest number of missing items for a single patient was 6.

‡ Scores on the modified Rankin scale of functional disability range from 0 (no symptoms) to 6 (death). A score of 2 or less indicates functional independence.

§ Data on systolic blood pressure at baseline were missing for one patient assigned to the control group.

¶ The Alberta Stroke Program Early Computed Tomography Score (ASPECTS) is a measure of the extent of stroke. Scores ranges from 0 to 10, with higher scores indicating fewer early ischemic changes. Scores were not available for four patients assigned to the control group: noncontrast computed tomography was not performed in one patient, and three patients had strokes in the territory of the anterior cerebral artery.

|| Vessel imaging was not performed in one patient in the control group, so the level of occlusion was not known.

** Extracranial ICA occlusions were reported by local investigators.

†† Data were missing for two patients in the intervention group.

1.4 Simulated fakestroke data

Consider the simulated data, available on our Data and Code website in the `fakestroke.csv` file, which I built to let us mirror the Table 1 for MR CLEAN (Berkhemer et al. 2015). The `fakestroke.csv` file contains the following 18 variables for 500 patients.

Variable	Description
<code>studyid</code>	Study ID # (z001 through z500)
<code>trt</code>	Treatment group (Intervention or Control)
<code>age</code>	Age in years
<code>sex</code>	Male or Female
<code>nihss</code>	NIH Stroke Scale Score (can range from 0-42; higher scores indicate more severe neurological deficits)
<code>location</code>	Stroke Location - Left or Right Hemisphere
<code>hx.isch</code>	History of Ischemic Stroke (Yes/No)
<code>afib</code>	Atrial Fibrillation (1 = Yes, 0 = No)
<code>dm</code>	Diabetes Mellitus (1 = Yes, 0 = No)
<code>mrankin</code>	Pre-stroke modified Rankin scale score (0, 1, 2 or > 2) indicating functional disability - complete range is 0 (no symptoms) to 6 (death)
<code>sbp</code>	Systolic blood pressure, in mm Hg
<code>iv.altep</code>	Treatment with IV alteplase (Yes/No)
<code>time.iv</code>	Time from stroke onset to start of IV alteplase (minutes) if <code>iv.altep</code> =Yes
<code>aspects</code>	Alberta Stroke Program Early Computed Tomography score, which measures extent of stroke from 0 - 10; higher scores indicate fewer early ischemic changes
<code>ia.occlus</code>	Intracranial arterial occlusion, based on vessel imaging - five categories ³
<code>extra.ica</code>	Extracranial ICA occlusion (1 = Yes, 0 = No)
<code>time.rand</code>	Time from stroke onset to study randomization, in minutes
<code>time.punc</code>	Time from stroke onset to groin puncture, in minutes (only if Intervention)

Here's a quick look at the simulated data in `fakestroke`.

```
fakestroke
```

```
# A tibble: 500 x 18
  studyid trt      age sex   nihss location hx.isch afib     dm mrankin    sbp
  <chr>   <chr>    <dbl> <chr> <dbl> <chr>   <chr>   <dbl> <dbl> <chr>   <dbl>
1 z001    Control    53 Male    21 Right   No       0     0 2     127
2 z002    Interve~   51 Male    23 Left    No       1     0 0     137
3 z003    Control    68 Fema~   11 Right   No       0     0 0     138
```

```

4 z004    Control    28 Male     22 Left    No      0   0 0 122
5 z005    Control    91 Male    24 Right   No      0   0 0 162
6 z006    Control    34 Fema~   18 Left    No      0   0 2 166
7 z007    Interve~  75 Male    25 Right   No      0   0 0 140
8 z008    Control    89 Fema~   18 Right   No      0   0 0 157
9 z009    Control    75 Male    25 Left    No      1   0 2 129
10 z010   Interve~  26 Fema~   27 Right   No      0   0 0 143
# ... with 490 more rows, and 7 more variables: iv.altep <chr>, time.iv <dbl>,
#   aspects <dbl>, ia.occlus <chr>, extra.ica <dbl>, time.rand <dbl>,
#   time.punc <dbl>

```

1.5 Building Table 1 for `fakestroke`: Attempt 1

Our goal, then, is to take the data in `fakestroke.csv` and use it to generate a Table 1 for the study that compares the 233 patients in the Intervention group to the 267 patients in the Control group, on all of the other variables (except study ID #) available. I'll use the `tableone` package of functions available in R to help me complete this task. We'll make a first attempt, using the `CreateTableOne` function in the `tableone` package. To use the function, we'll need to specify:

- the `vars` or variables we want to place in the rows of our Table 1 (which will include just about everything in the `fakestroke` data except the `studyid` code and the `trt` variable for which we have other plans, and the `time.punc` which applies only to subjects in the Intervention group.)
 - A useful trick here is to use the `dput` function, specifically something like `dput(names(fakestroke))` can be used to generate a list of all of the variables included in the `fakestroke` tibble, and then this can be copied and pasted into the `vars` specification, saving some typing.
- the `strata` which indicates the levels want to use in the columns of our Table 1 (for us, that's `trt`)

```

fs.vars <- c("age", "sex", "nihss", "location",
           "hx.isch", "afib", "dm", "mrankin", "sbp",
           "iv.altep", "time.iv", "aspects",
           "ia.occlus", "extra.ica", "time.rand")

fs.trt <- c("trt")

att1 <- CreateTableOne(data = fakestroke,
                        vars = fs.vars,
                        strata = fs.trt)

print(att1)

```

Stratified by trt			
Control	Intervention	p	test

n	267	233	
age (mean (SD))	65.38 (16.10)	63.93 (18.09)	0.343
sex = Male (%)	157 (58.8)	135 (57.9)	0.917
nihss (mean (SD))	18.08 (4.32)	17.97 (5.04)	0.787
location = Right (%)	114 (42.7)	117 (50.2)	0.111
hx.isch = Yes (%)	25 (9.4)	29 (12.4)	0.335
afib (mean (SD))	0.26 (0.44)	0.28 (0.45)	0.534
dm (mean (SD))	0.13 (0.33)	0.12 (0.33)	0.923
mrankin (%)			0.922
> 2	11 (4.1)	10 (4.3)	
0	214 (80.1)	190 (81.5)	
1	29 (10.9)	21 (9.0)	
2	13 (4.9)	12 (5.2)	
sbp (mean (SD))	145.00 (24.40)	146.03 (26.00)	0.647
iv.altep = Yes (%)	242 (90.6)	203 (87.1)	0.267
time.iv (mean (SD))	87.96 (26.01)	98.22 (45.48)	0.003
aspects (mean (SD))	8.65 (1.47)	8.35 (1.64)	0.033
ia.occlus (%)			0.795
A1 or A2	2 (0.8)	1 (0.4)	
ICA with M1	75 (28.2)	59 (25.3)	
Intracranial ICA	3 (1.1)	1 (0.4)	
M1	165 (62.0)	154 (66.1)	
M2	21 (7.9)	18 (7.7)	
extra.ica (mean (SD))	0.26 (0.44)	0.32 (0.47)	0.150
time.rand (mean (SD))	213.88 (70.29)	202.51 (57.33)	0.051

1.5.1 Some of this is very useful, and other parts need to be fixed.

1. The 1/0 variables (`afib`, `dm`, `extra.ica`) might be better if they were treated as the factors they are, and reported as the Yes/No variables are reported, with counts and percentages rather than with means and standard deviations.
2. In some cases, we may prefer to re-order the levels of the categorical (factor) variables, particularly the `mrankin` variable, but also the `ia.occlus` variable. It would also be more typical to put the Intervention group to the left and the Control group to the right, so we may need to adjust our `trt` variable's levels accordingly.
3. For each of the quantitative variables (`age`, `nihss`, `sbp`, `time.iv`, `aspects`, `extra.ica`, `time.rand` and `time.punc`) we should make a decision whether a summary with mean and standard deviation is appropriate, or whether we should instead summarize with, say, the median and quartiles. A mean and standard deviation really only yields an appropriate summary when the data are least approximately Normally distributed. This will make the *p* values a bit more reasonable, too. The `test` column in the first attempt

- will soon have something useful to tell us.
4. If we'd left in the `time.punc` variable, we'd get some warnings, having to do with the fact that `time.punc` is only relevant to patients in the Intervention group.

1.5.2 `fakestroke` Cleaning Up Categorical Variables

Let's specify each of the categorical variables as categorical explicitly. This helps the `CreateTableOne` function treat them appropriately, and display them with counts and percentages. This includes all of the 1/0, Yes/No and multi-categorical variables.

```
fs.factorvars <- c("sex", "location", "hx.isch", "afib", "dm",
  "mrankin", "iv.altep", "ia.occlus", "extra.ica")
```

Then we simply add a `factorVars = fs.factorvars` call to the `CreateTableOne` function.

We also want to re-order some of those categorical variables, so that the levels are more useful to us. Specifically, we want to:

- place Intervention before Control in the `trt` variable,
- reorder the `mrankin` scale as 0, 1, 2, > 2, and
- rearrange the `ia.occlus` variable to the order⁴ presented in Berkhemer et al. (2015).

To accomplish this, we'll use the `fct_relevel` function from the `forcats` package (loaded with the rest of the core `tidyverse` packages) to reorder our levels manually.

```
fakestroke <- fakestroke %>%
  mutate(trt = fct_relevel(trt, "Intervention", "Control"),
    mrankin = fct_relevel(mrankin, "0", "1", "2", "> 2"),
    ia.occlus = fct_relevel(ia.occlus, "Intracranial ICA",
      "ICA with M1", "M1", "M2",
      "A1 or A2")
  )
```

1.6 `fakestroke` Table 1: Attempt 2

```
att2 <- CreateTableOne(data = fakestroke,
  vars = fs.vars,
  factorVars = fs.factorvars,
```

⁴We might also have considered reordering the `ia.occlus` factor by its frequency, using the `fct_infreq` function

	Stratified by trt		p	test
n	Intervention	Control		
age (mean (SD))	63.93 (18.09)	65.38 (16.10)	0.343	
sex = Male (%)	135 (57.9)	157 (58.8)	0.917	
nihss (mean (SD))	17.97 (5.04)	18.08 (4.32)	0.787	
location = Right (%)	117 (50.2)	114 (42.7)	0.111	
hx.isch = Yes (%)	29 (12.4)	25 (9.4)	0.335	
afib = 1 (%)	66 (28.3)	69 (25.8)	0.601	
dm = 1 (%)	29 (12.4)	34 (12.7)	1.000	
mrankin (%)				0.922
0	190 (81.5)	214 (80.1)		
1	21 (9.0)	29 (10.9)		
2	12 (5.2)	13 (4.9)		
> 2	10 (4.3)	11 (4.1)		
sbp (mean (SD))	146.03 (26.00)	145.00 (24.40)	0.647	
iv.altep = Yes (%)	203 (87.1)	242 (90.6)	0.267	
time.iv (mean (SD))	98.22 (45.48)	87.96 (26.01)	0.003	
aspects (mean (SD))	8.35 (1.64)	8.65 (1.47)	0.033	
ia.occlus (%)				0.795
Intracranial ICA	1 (0.4)	3 (1.1)		
ICA with M1	59 (25.3)	75 (28.2)		
M1	154 (66.1)	165 (62.0)		
M2	18 (7.7)	21 (7.9)		
A1 or A2	1 (0.4)	2 (0.8)		
extra.ica = 1 (%)	75 (32.2)	70 (26.3)	0.179	
time.rand (mean (SD))	202.51 (57.33)	213.88 (70.29)	0.051	

The categorical data presentation looks much improved.

1.6.1 What summaries should we show?

Now, we'll move on to the issue of making a decision about what type of summary to show for the quantitative variables. Since the `fakestroke` data are just simulated and only match the summary statistics of the original results, not the details, we'll adopt the decisions made by Berkhemer et al. (2015), which were to use medians and interquartile ranges to summarize the distributions of all of the continuous variables **except** systolic blood pressure.

- Specifying certain quantitative variables as *non-normal* causes R to show them with medians and the 25th and 75th percentiles, rather than means and standard deviations, and also causes those variables to be tested using

non-parametric tests, like the Wilcoxon signed rank test, rather than the t test. The `test` column indicates this with the word `nonnorm`.

- In real data situations, what should we do? The answer is to look at the data. I would not make the decision as to which approach to take without first plotting (perhaps in a histogram or a Normal Q-Q plot) the observed distributions in each of the two samples, so that I could make a sound decision about whether Normality was a reasonable assumption. If the means and medians are meaningfully different from each other, this is especially important.
- To be honest, though, if the variable in question is a relatively unimportant covariate and the p values for the two approaches are nearly the same, I'd say that further investigation is rarely important,
- Specifying *exact* tests for certain categorical variables (we'll try this for the `location` and `mrankin` variables) can be done, and these changes will be noted in the `test` column, as well.
 - In real data situations, I would rarely be concerned about this issue, and often choose Pearson (approximate) options across the board. This is reasonable so long as the number of subjects falling in each category is reasonably large, say above 10. If not, then an exact test may be a tiny improvement.
 - Paraphrasing Rosenbaum (2017), having an exact rather than an approximate test result is about as valuable as having a nice crease in your trousers.

To finish our Table 1, then, we need to specify which variables should be treated as non-Normal in the `print` statement - notice that we don't need to redo the `CreateTableOne` for this change.

```
print(att2,
      nonnormal = c("age", "nihss", "time.iv", "aspects", "time.rand"),
      exact = c("location", "mrankin"))
```

	Stratified by trt	
	Intervention	Control
n	233	267
age (median [IQR])	65.80 [54.50, 76.00]	65.70 [55.75, 76.20]
sex = Male (%)	135 (57.9)	157 (58.8)
nihss (median [IQR])	17.00 [14.00, 21.00]	18.00 [14.00, 22.00]
location = Right (%)	117 (50.2)	114 (42.7)
hx.isch = Yes (%)	29 (12.4)	25 (9.4)
afib = 1 (%)	66 (28.3)	69 (25.8)
dm = 1 (%)	29 (12.4)	34 (12.7)
mrankin (%)		
0	190 (81.5)	214 (80.1)
1	21 (9.0)	29 (10.9)
2	12 (5.2)	13 (4.9)
> 2	10 (4.3)	11 (4.1)

```

sbp (mean (SD))      146.03 (26.00)      145.00 (24.40)
iv.altep = Yes (%)    203 (87.1)          242 (90.6)
time.iv (median [IQR]) 85.00 [67.00, 110.00] 87.00 [65.00, 116.00]
aspects (median [IQR]) 9.00 [7.00, 10.00]    9.00 [8.00, 10.00]
ia.occlus (%)

  Intracranial ICA      1 ( 0.4)          3 ( 1.1)
  ICA with M1            59 (25.3)         75 (28.2)
  M1                      154 (66.1)        165 (62.0)
  M2                      18 ( 7.7)          21 ( 7.9)
  A1 or A2                1 ( 0.4)          2 ( 0.8)
extra.ica = 1 (%)       75 (32.2)          70 (26.3)
time.rand (median [IQR]) 204.00 [152.00, 249.50] 196.00 [149.00, 266.00]

Stratified by trt
  p      test

n
age (median [IQR])      0.579 nonnorm
sex = Male (%)           0.917
nihss (median [IQR])     0.453 nonnorm
location = Right (%)     0.106 exact
hx.isch = Yes (%)        0.335
afib = 1 (%)              0.601
dm = 1 (%)                1.000
mrankin (%)               0.917 exact

  0
  1
  2
  > 2

sbp (mean (SD))          0.647
iv.altep = Yes (%)        0.267
time.iv (median [IQR])    0.596 nonnorm
aspects (median [IQR])    0.075 nonnorm
ia.occlus (%)              0.795

  Intracranial ICA
  ICA with M1
  M1
  M2
  A1 or A2
extra.ica = 1 (%)         0.179
time.rand (median [IQR])   0.251 nonnorm

```

1.7 Obtaining a more detailed Summary

If this was a real data set, we'd want to get a more detailed description of the data to make decisions about things like potentially collapsing categories of a

variable, or whether or not a normal distribution was useful for a particular continuous variable, etc. You can do this with the **summary** command applied to a created Table 1, which shows, among other things, the effect of changing from normal to non-normal *p* values for continuous variables, and from approximate to “exact” *p* values for categorical factors.

Again, as noted above, in a real data situation, we’d want to plot the quantitative variables (within each group) to make a smart decision about whether a t test or Wilcoxon approach is more appropriate.

Note in the summary below that we have some missing values here. Often, we’ll present this information within the Table 1, as well.

```
summary(att2)
```

```
### Summary of continuous variables ###

trt: Intervention
    n miss p.miss mean sd median p25 p75 min max skew kurt
age     233   0    0.0   64 18      66  54  76  23  96 -0.34 -0.52
nihss   233   0    0.0   18  5      17  14  21  10  28  0.48 -0.74
sbp     233   0    0.0  146 26     146 129 164  78 214 -0.07 -0.22
time.iv 233   30   12.9  98 45     85  67 110  42 218  1.03  0.08
aspects  233   0    0.0    8  2      9   7  10  5  10 -0.56 -0.98
time.rand 233   2    0.9  203 57    204 152 250 100 300  0.01 -1.16
-----
trt: Control
    n miss p.miss mean sd median p25 p75 min max skew kurt
age     267   0    0.0   65 16      66  56  76  24  94 -0.296 -0.28
nihss   267   0    0.0   18  4      18  14  22  11  25  0.017 -1.24
sbp     267   1    0.4  145 24     145 128 161  82 231  0.156  0.08
time.iv 267   25   9.4   88 26     87  65 116  44 130  0.001 -1.32
aspects  267   4    1.5    9  1      9   8  10  5  10 -1.071  0.36
time.rand 267   0    0.0  214 70    196 149 266 120 360  0.508 -0.93

p-values
          pNormal pNonNormal
age       0.342813660 0.57856976
nihss    0.787487252 0.45311695
sbp      0.647157646 0.51346132
time.iv  0.003073372 0.59641104
aspects  0.032662901 0.07464683
time.rand 0.050803672 0.25134327

Standardize mean differences
      1 vs 2
age      0.08478764
```

```

nihss      0.02405390
sbp        0.04100833
time.iv    0.27691223
aspects    0.19210662
time.rand  0.17720957
=====
```

```
### Summary of categorical variables ###
```

trt: Intervention

	var	n	miss	p.miss	level	freq	percent	cum.percent
	sex	233	0	0.0	Female	98	42.1	42.1
					Male	135	57.9	100.0
	location	233	0	0.0	Left	116	49.8	49.8
					Right	117	50.2	100.0
	hx.isch	233	0	0.0	No	204	87.6	87.6
					Yes	29	12.4	100.0
	afib	233	0	0.0	0	167	71.7	71.7
					1	66	28.3	100.0
	dm	233	0	0.0	0	204	87.6	87.6
					1	29	12.4	100.0
	mrankin	233	0	0.0	0	190	81.5	81.5
					1	21	9.0	90.6
					2	12	5.2	95.7
					> 2	10	4.3	100.0
	iv.altep	233	0	0.0	No	30	12.9	12.9
					Yes	203	87.1	100.0
	ia.occlus	233	0	0.0	Intracranial	ICA	1	0.4
					ICA with M1	59	25.3	25.8
					M1	154	66.1	91.8
					M2	18	7.7	99.6
					A1 or A2	1	0.4	100.0
	extra.ica	233	0	0.0	0	158	67.8	67.8
					1	75	32.2	100.0

trt: Control

var	n	miss	p.miss	level	freq	percent	cum.percent
sex	267	0	0.0	Female	110	41.2	41.2
				Male	157	58.8	100.0
location	267	0	0.0	Left	153	57.3	57.3
				Right	114	42.7	100.0
hx.isch	267	0	0.0	No	242	90.6	90.6
				Yes	25	9.4	100.0
afib	267	0	0.0	0	198	74.2	74.2
				1	69	25.8	100.0
dm	267	0	0.0	0	233	87.3	87.3
				1	34	12.7	100.0
mrankin	267	0	0.0	0	214	80.1	80.1
				1	29	10.9	91.0
				2	13	4.9	95.9
				> 2	11	4.1	100.0
iv.altep	267	0	0.0	No	25	9.4	9.4
				Yes	242	90.6	100.0
ia.occlus	267	1	0.4	Intracranial	ICA	3	1.1
				ICA with M1	75	28.2	29.3
				M1	165	62.0	91.4
				M2	21	7.9	99.2
				A1 or A2	2	0.8	100.0
extra.ica	267	1	0.4	0	196	73.7	73.7
				1	70	26.3	100.0

p-values

	pApprox	pExact
sex	0.9171387	0.8561188
location	0.1113553	0.1056020
hx.isch	0.3352617	0.3124683
afib	0.6009691	0.5460206
dm	1.0000000	1.0000000
mrankin	0.9224798	0.9173657
iv.altep	0.2674968	0.2518374
ia.occlus	0.7945580	0.8189090
extra.ica	0.1793385	0.1667574

```
Standardize mean differences
      1 vs 2
sex      0.017479025
location 0.151168444
hx.isch  0.099032275
afib     0.055906317
dm       0.008673478
mrankin  0.062543164
iv.altep  0.111897009
ia.occlus 0.117394890
extra.ica 0.129370206
```

In this case, I have simulated the data to mirror the results in the published Table 1 for this study. In no way have I captured the full range of the real data, or any of the relationships in that data, so it's more important here to see what's available in the analysis, rather than to interpret it closely in the clinical context.

1.8 Exporting the Completed Table 1 from R to Excel or Word

Once you've built the table and are generally satisfied with it, you'll probably want to be able to drop it into Excel or Word for final cleanup.

1.8.1 Approach A: Save and open in Excel

One option is to **save the Table 1** to a .csv file within our `data` subfolder (note that the `data` folder must already exist), which you can then open directly in Excel. This is the approach I generally use. Note the addition of some `quote`, `noSpaces` and `printToggle` selections here.

```
fs.table1$save <- print(att2,
  nonnormal = c("age", "nihss", "time.iv", "aspects", "time.rand"),
  exact = c("location", "mrankin"),
  quote = FALSE, noSpaces = TRUE, printToggle = FALSE)

write.csv(fs.table1$save, file = "data/fs-table1.csv")
```

When I then open the `fs-table1.csv` file in Excel, it looks like this:

	A	B	C	D	E
1		Intervention	Control	p	test
2	n	233	267		
3	age (median [IQR])	65.80 [54.50, 76.00]	65.70 [55.75, 76.20]	0.579	nonnorm
4	sex = Male (%)	135 (57.9)	157 (58.8)	0.917	
5	nihss (median [IQR])	17.00 [14.00, 21.00]	18.00 [14.00, 22.00]	0.453	nonnorm
6	location = Right (%)	117 (50.2)	114 (42.7)	0.111	
7	hx.isch = Yes (%)	29 (12.4)	25 (9.4)	0.335	
8	afib = 1 (%)	66 (28.3)	69 (25.8)	0.601	
9	dm = 1 (%)	29 (12.4)	34 (12.7)	1	
10	mrankin (%)			0.922	
11		0 190 (81.5)	214 (80.1)		
12		1 21 (9.0)	29 (10.9)		
13		2 12 (5.2)	13 (4.9)		
14	> 2	10 (4.3)	11 (4.1)		
15	sbp (mean (sd))	146.03 (26.00)	145.00 (24.40)	0.647	
16	iv.altep = Yes (%)	203 (87.1)	242 (90.6)	0.267	
17	time.iv (median [IQR])	85.00 [67.00, 110.00]	87.00 [65.00, 116.00]	0.596	nonnorm
18	aspects (median [IQR])	9.00 [7.00, 10.00]	9.00 [8.00, 10.00]	0.075	nonnorm
19	ia.occlus (%)			0.795	
20	Intracranial ICA	1 (0.4)	3 (1.1)		
21	ICA with M1	59 (25.3)	75 (28.2)		
22	M1	154 (66.1)	165 (62.0)		
23	M2	18 (7.7)	21 (7.9)		
24	A1 or A2	1 (0.4)	2 (0.8)		
25	extra.ica = 1 (%)	75 (32.2)	70 (26.3)	0.179	
26	time.rand (median [IQR])	204.00 [152.00, 249.50]	196.00 [149.00, 266.00]	0.251	nonnorm
27	time.punc (median [IQR])	260.00 [212.00, 313.00]	NA [NA, NA]	NA	nonnorm
28					

And from here, I can either drop it directly into Word, or present it as is, or start tweaking it to meet formatting needs.

1.8.2 Approach B: Produce the Table so you can cut and paste it

```
print(att2,
      nonnormal = c("age", "nihss", "time.iv", "aspects", "time.rand"),
      exact = c("location", "mrankin"),
      quote = TRUE, noSpaces = TRUE)
```

This will look like a mess by itself, but if you:

1. copy and paste that mess into Excel
2. select Text to Columns from the Data menu
3. select Delimited, then Space and select Treat consecutive delimiters as one

you should get something usable again.

Or, in Word,

1. insert the text
2. select the text with your mouse
3. select Insert ... Table ... Convert Text to Table
4. place a quotation mark in the “Other” area under Separate text at ...

After dropping blank columns, the result looks pretty good.

1.9 A Controlled Biological Experiment - The Blood-Brain Barrier

My source for the data and the following explanatory paragraph is page 307 from Ramsey and Schafer (2002). The original data come from Barnett et al. (1995).

The human brain (and that of rats, coincidentally) is protected from the bacteria and toxins that course through the bloodstream by something called the blood-brain barrier. After a method of disrupting the barrier was developed, researchers tested this new mechanism, as follows. A series of 34 rats were inoculated with human lung cancer cells to induce brain tumors. After 9-11 days they were infused with either the barrier disruption (BD) solution or, as a control, a normal saline (NS) solution. Fifteen minutes later, the rats received a standard dose of a particular therapeutic antibody (L6-F(ab')2. The key measure of the effectiveness of transmission across the brain-blood barrier is the ratio of the antibody concentration in the brain tumor to the antibody concentration in normal tissue outside the brain. The rats were then sacrificed, and the amounts of antibody in the brain tumor and in normal tissue from the liver were measured. The study's primary objective is to determine whether the antibody concentration in the tumor increased when the blood-barrier disruption infusion was given, and if so, by how much?

1.10 The `bloodbrain.csv` file

Consider the data, available on our Data and Code website in the `bloodbrain.csv` file, which includes the following variables:

Variable	
<code>case</code>	identification number for the rat (1 - 34)
<code>brain</code>	an outcome: Brain tumor antibody count (per gram)

³The five categories are Intracranial ICA, ICA with involvement of the M1 middle cerebral artery segment, M1 middle cerebral artery segment, M2 middle cerebral artery segment, A1 or A2 anterior cerebral artery segment

<code>case</code>	identification number for the rat (1 - 34)
<code>brain</code>	an outcome: Brain tumor antibody count (per gram)

Variable	Description
<code>liver</code>	an outcome: Liver antibody count (per gram)
<code>tlratio</code>	an outcome: tumor / liver concentration ratio
<code>solution</code>	the treatment: BD (barrier disruption) or NS (normal saline)
<code>sactime</code>	a design variable: Sacrifice time (hours; either 0.5, 3, 24 or 72)
<code>postin</code>	covariate: Days post-inoculation of lung cancer cells (9, 10 or 11)
<code>sex</code>	covariate: M or F
<code>wt.init</code>	covariate: Initial weight (grams)
<code>wt.loss</code>	covariate: Weight loss (grams)
<code>wt.tumor</code>	covariate: Tumor weight (10^{-4} grams)

And here's what the data look like in R.

```
bloodbrain
```

```
# A tibble: 34 x 11
  case   brain    liver tlratio solution sactime postin sex   wt.init wt.loss
  <dbl>  <dbl>  <dbl>  <dbl> <chr>     <dbl>  <dbl> <chr>  <dbl>  <dbl>
1     1  41081 1456164  0.0282 BD      0.5    10 F    239   5.9
2     2  44286 1602171  0.0276 BD      0.5    10 F    225    4
3     3 102926 1601936  0.0642 BD      0.5    10 F    224  -4.9
4     4  25927 1776411  0.0146 BD      0.5    10 F    184   9.8
5     5  42643 1351184  0.0316 BD      0.5    10 F    250    6
6     6  31342 1790863  0.0175 NS      0.5    10 F    196   7.7
7     7  22815 1633386  0.0140 NS      0.5    10 F    200   0.5
8     8  16629 1618757  0.0103 NS      0.5    10 F    273    4
9     9  22315 1567602  0.0142 NS      0.5    10 F    216   2.8
10    10  77961 1060057  0.0735 BD      3     10 F    267   2.6
# ... with 24 more rows, and 1 more variable: wt.tumor <dbl>
```

1.11 A Table 1 for `bloodbrain`

Barnett et al. (1995) did not provide a Table 1 for these data, so let's build one to compare the two `solutions` (BD vs. NS) on the covariates and outcomes, plus the natural logarithm of the tumor/liver concentration ratio (`tlratio`). We'll opt to treat the sacrifice time (`sactime`) and the days post-inoculation of lung cancer cells (`postin`) as categorical rather than quantitative variables.

```
bloodbrain <- bloodbrain %>%
  mutate(logTL = log(tlratio))
```

```
dput(names(bloodbrain))

c("case", "brain", "liver", "tlratio", "solution", "sactime",
  "postin", "sex", "wt.init", "wt.loss", "wt.tumor", "logTL")

OK - there's the list of variables we'll need. I'll put the outcomes at the bottom
of the table.

bb.vars <- c("sactime", "postin", "sex", "wt.init", "wt.loss",
           "wt.tumor", "brain", "liver", "tlratio", "logTL")

bb.factors <- c("sactime", "sex", "postin")

bb.att1 <- CreateTableOne(data = bloodbrain,
                           vars = bb.vars,
                           factorVars = bb.factors,
                           strata = c("solution"))

summary(bb.att1)

### Summary of continuous variables ###

solution: BD
      n miss p.miss   mean     sd median    p25    p75    min    max skew
wt.init 17    0      0 243 3e+01 2e+02 2e+02 3e+02 2e+02 3e+02 -0.39
wt.loss 17    0      0 3 5e+00 4e+00 1e+00 6e+00 -5e+00 1e+01 -0.10
wt.tumor 17    0      0 157 8e+01 2e+02 1e+02 2e+02 2e+01 4e+02  0.53
brain    17    0      0 56043 3e+04 5e+04 4e+04 8e+04 6e+03 1e+05  0.29
liver    17    0      0 672577 7e+05 6e+05 2e+04 1e+06 2e+03 2e+06  0.35
tlratio  17    0      0 2 3e+00 1e-01 6e-02 3e+00 1e-02 9e+00  1.58
logTL   17    0      0 -1 2e+00 -2e+00 -3e+00 1e+00 -4e+00 2e+00  0.08
      kurt
wt.init  0.7
wt.loss  0.2
wt.tumor 1.0
brain   -0.6
liver   -1.7
tlratio  1.7
logTL   -1.7
-----
solution: NS
      n miss p.miss   mean     sd median    p25    p75    min    max skew
wt.init 17    0      0 240 3e+01 2e+02 2e+02 3e+02 2e+02 3e+02  0.33
wt.loss 17    0      0 4 4e+00 3e+00 2e+00 7e+00 -4e+00 1e+01 -0.09
wt.tumor 17    0      0 209 1e+02 2e+02 2e+02 3e+02 3e+01 5e+02  0.63
brain    17    0      0 23887 1e+04 2e+04 1e+04 3e+04 1e+03 5e+04  0.30
liver    17    0      0 664975 7e+05 7e+05 2e+04 1e+06 9e+02 2e+06  0.40
```

```

tlratio 17    0      0      1 2e+00 5e-02 3e-02 9e-01 1e-02 7e+00 2.27
logTL   17    0      0     -2 2e+00 -3e+00 -3e+00 -7e-02 -5e+00 2e+00 0.27
      kurt
wt.init -0.48
wt.loss  0.08
wt.tumor 0.77
brain    -0.35
liver    -1.56
tlratio  4.84
logTL   -1.61

p-values
      pNormal  pNonNormal
wt.init 0.807308940 0.641940278
wt.loss 0.683756156 0.876749808
wt.tumor 0.151510151 0.190482094
brain   0.001027678 0.002579901
liver   0.974853609 0.904045603
tlratio 0.320501715 0.221425879
logTL   0.351633525 0.221425879

Standardize mean differences
      1 vs 2
wt.init 0.08435244
wt.loss 0.14099823
wt.tumor 0.50397184
brain   1.23884159
liver   0.01089667
tlratio 0.34611465
logTL   0.32420504

=====
### Summary of categorical variables ###

solution: BD
      var n miss p.miss level freq percent cum.percent
sactime 17    0    0.0    0.5    5    29.4      29.4
          3    4    23.5      52.9
          24   4    23.5      76.5
          72   4    23.5     100.0
postin 17    0    0.0     9    1    5.9      5.9
          10   14   82.4     88.2
          11   2   11.8     100.0

```

```

sex 17    0    0.0      F 13    76.5    76.5
                  M  4    23.5
                                         100.0

-----
solution: NS
      var   n miss p.miss level freq percent cum.percent
sactime 17    0    0.0    0.5     4    23.5    23.5
              3     5    29.4    52.9
              24    4    23.5    76.5
              72    4    23.5   100.0

postin 17    0    0.0      9    2    11.8    11.8
              10   13    76.5    88.2
              11    2    11.8
                                         100.0

sex 17    0    0.0      F 13    76.5    76.5
                  M  4    23.5
                                         100.0

p-values
      pApprox pExact
sactime 0.9739246    1
postin  0.8309504    1
sex      1.0000000    1

Standardize mean differences
      1 vs 2
sactime 0.1622214
postin  0.2098877
sex      0.0000000

```

Note that, in this particular case, the decisions we make about normality vs. non-normality (for quantitative variables) and the decisions we make about approximate vs. exact testing (for categorical variables) won't actually change the implications of the p values. Each approach gives similar results for each variable. Of course, that's not always true.

1.11.1 Generate final Table 1 for bloodbrain

I'll choose to treat `tlratio` and its logarithm as non-Normal, but otherwise, use t tests, but admittedly, that's an arbitrary decision, really.

```
print(bb.att1, nonnormal = c("tlratio", "logTL"))
```

Stratified by solution	NS
BD	NS

n	17	17
sactime (%)		
0.5	5 (29.4)	4 (23.5)
3	4 (23.5)	5 (29.4)
24	4 (23.5)	4 (23.5)
72	4 (23.5)	4 (23.5)
postin (%)		
9	1 (5.9)	2 (11.8)
10	14 (82.4)	13 (76.5)
11	2 (11.8)	2 (11.8)
sex = M (%)	4 (23.5)	4 (23.5)
wt.init (mean (SD))	242.82 (27.23)	240.47 (28.54)
wt.loss (mean (SD))	3.34 (4.68)	3.94 (3.88)
wt.tumor (mean (SD))	157.29 (84.00)	208.53 (116.68)
brain (mean (SD))	56043.41 (33675.40)	23887.18 (14610.53)
liver (mean (SD))	672577.35 (694479.58)	664975.47 (700773.13)
tlratio (median [IQR])	0.12 [0.06, 2.84]	0.05 [0.03, 0.94]
logTL (median [IQR])	-2.10 [-2.74, 1.04]	-2.95 [-3.41, -0.07]
Stratified by solution		
p	test	
n		
sactime (%)	0.974	
0.5		
3		
24		
72		
postin (%)	0.831	
9		
10		
11		
sex = M (%)	1.000	
wt.init (mean (SD))	0.807	
wt.loss (mean (SD))	0.684	
wt.tumor (mean (SD))	0.152	
brain (mean (SD))	0.001	
liver (mean (SD))	0.975	
tlratio (median [IQR])	0.221 nonnorm	
logTL (median [IQR])	0.221 nonnorm	

Or, we can get an Excel-readable version placed in a `data` subfolder, using

```
bb.t1 <- print(bb.att1, nonnormal = c("tlratio", "logTL"), quote = FALSE,
                noSpaces = TRUE, printToggle = FALSE)

write.csv(bb.t1, file = "data/bb-table1.csv")
```

which, when dropped into Excel, will look like this:

	A	B	C	D	E	
1		BD	NS	p		test
2 n		17		17		
3 sex = M (%)	4 (23.5)	4 (23.5)		1		
4 sactime (%)				0.974		
5	0.5 5 (29.4)	4 (23.5)				
6	3 4 (23.5)	5 (29.4)				
7	24 4 (23.5)	4 (23.5)				
8	72 4 (23.5)	4 (23.5)				
9 postin (%)				0.831		
10	9 1 (5.9)	2 (11.8)				
11	10 14 (82.4)	13 (76.5)				
12	11 2 (11.8)	2 (11.8)				
13 wt.init (mean (sd))	242.82 (27.23)	240.47 (28.54)		0.807		
14 wt.loss (mean (sd))	3.34 (4.68)	3.94 (3.88)		0.684		
15 wt.tumor (mean (sd))	157.29 (84.00)	208.53 (116.68)		0.152		
16 brain (mean (sd))	56043.41 (33675.40)	23887.18 (14610.53)		0.001		
17 liver (mean (sd))	672577.35 (694479.58)	664975.47 (700773.13)		0.975		
18 tlratio (median [IQR])	0.12 [0.06, 2.84]	0.05 [0.03, 0.94]		0.221	nonnorm	
19 logTL (median [IQR])	-2.10 [-2.74, 1.04]	-2.95 [-3.41, -0.07]		0.221	nonnorm	
20						

One thing I would definitely clean up here, in practice, is to change the presentation of the *p* value for `sex` from 1 to > 0.99, or just omit it altogether. I'd also drop the computer-eze where possible, add units for the measures, round a lot, identify the outcomes carefully, and use notes to indicate deviations from the main approach.

1.11.2 A More Finished Version (after Cleanup in Word)

Table 1. Comparing Rats Receiving BD to those Receiving NS on Available Covariates and Design Variables, and Key Outcomes

	Barrier Disruption (BD: treatment)	Normal Saline (NS: control)	p
# of Rats	17	17	
Sex = Male	4 (23.5)	4 (23.5)	-
Sacrifice Time (hours)			0.97
0.5	5 (29.4)	4 (23.5)	
3	4 (23.5)	5 (29.4)	
24	4 (23.5)	4 (23.5)	
72	4 (23.5)	4 (23.5)	
Days post-inoculation of lung cancer cells			0.83
9	1 (5.9)	2 (11.8)	
10	14 (82.4)	13 (76.5)	
11	2 (11.8)	2 (11.8)	
Initial Weight (g)	243 (27)	240 (29)	0.81
Weight Loss (g)	3.3 (4.7)	3.9 (3.9)	0.68
Tumor Weight (10^{-4} g)	157.3 (84.0)	208.5 (116.7)	0.15
Key Outcomes: mean (sd) unless otherwise indicated			
Brain Tumor Antibody Count (per g)	56,043 (33,675)	23,887 (14,611)	0.001
Liver Antibody Count (per g)	672,577 (694,480)	664,975 (700,773)	0.98
Tumor/Liver Ratio (median [Q25, Q75])	0.12 [0.06, 2.84]	0.05 [0.03, 0.94]	0.22
Natural Log of Tumor/Liver Ratio (median [Q25, Q75])	-2.10 [-2.74, 1.04]	-2.95 [-3.41, -0.07]	0.22

Table 1 Notes:

- Categorical variables are summarized with counts, percentages and p values based on approximate chi-square tests.
- Continuous variables, unless otherwise indicated, are summarized with means, standard deviations and p values based on t tests.
- The Tumor / Liver ratio and its natural logarithm are summarized with the median and quartiles and a p value from a non-parametric (Wilcoxon signed rank) test.

Chapter 2

BRFSS SMART Data Building

The Centers for Disease Control analyzes Behavioral Risk Factor Surveillance System (BRFSS) survey data for specific metropolitan and micropolitan statistical areas (MMSAs) in a program called the Selected Metropolitan/Micropolitan Area Risk Trends of BRFSS (SMART BRFSS.)

In this work, we will focus on data from the 2017 SMART, and in particular on data from the state of Ohio, and from the Cleveland-Elyria, OH, Metropolitan Statistical Area. The purpose of this survey is to provide localized health information that can help public health practitioners identify local emerging health problems, plan and evaluate local responses, and efficiently allocate resources to specific needs.

In this chapter, I describe some cleaning of the BRFSS SMART data, and break it out into national, statewide, and local samples.

The data files produced by this chapter include:

- `smart_ohio.Rds` which includes data on approximately 100 variables for over 7000 subjects in six MMSAs that are at least partially located in the state of Ohio.
- `smart_cle.Rds` which includes data on those same variables for a little over 1000 subjects in the Cleveland-Elyria-Lorain OH MMSA.

2.1 Key resources

- the “raw” data, in the form of the 2017 SMART BRFSS MMSA Data, found in a zipped SAS Transport Format file. The data were released in

October 2018.

- the MMSA Variable Layout which simply lists the variables included in the data file
- the Calculated Variables PDF which describes the risk factors by data variable names - there is also an online summary matrix of these calculated variables.
- the lengthy 2017 Survey Questions PDF which lists all questions asked as part of the BRFSS in 2017
- the enormous Codebook for the 2017 BRFSS Survey PDF which identifies the variables by name for us.

Also, for each subject, we are also provided with a sampling weight, in `_MMSAWT`, which will help us incorporate the sampling design later. These weights are at the MMSA level, and are used for generating MMSA-level estimates for variables in the data set. Details on the weighting methodology are available at this PDF.

2.2 Ingesting the Raw Data

To create the data files we'll use, I used the `read_xpt` function from the `haven` package to bring in the SAS XPT data file that is provided by CDC. The codes I used (but won't use in these Notes) were:

```
smart_raw <- read_xpt("MMSA2017/MMSA2017.xpt")
```

This gives the nationwide data, which has 230,875 rows and 177 columns.

But for the purposes of putting these Notes online, I needed to crank down the sample size enormously. To that end, I created a new data file, which I developed by

- importing the `MMSA2017.xpt` file as above
- filtering away all observations except those from MMSAs which include Ohio in their name, and
- saving the result, which now has 7,412 rows and 177 columns.

The code (again, not run here) that I used to filter to the OH-based MMSAs was:

```
smart_ohio_raw <- smart_raw %>%
  filter(str_detect(MMSANAME, "OH"))

write_csv(smart_ohio_raw, "data/smart_ohio_raw.csv")
```

So, for purposes of these notes, our complete data set is actually coming from `smart_ohio_raw.csv` and consists only of the 7,412 observations associated with the six MMSAs that include Ohio in their names.

2.3 Ingesting from our CSV file

Note that the `smart_ohio_raw.csv` and other data files we're developing in this Chapter are available on our Data and Code website

```
smart_ohio_raw <- read_csv("data/smart_ohio_raw.csv")
```

```
Warning: One or more parsing issues, see `problems()` for details
dim(smart_ohio_raw)
```

```
[1] 7412 177
```

2.4 What does the raw data look like?

Here is a list of all variable names included in this file. We're not going to use all of those variables, but this will give you a sense of what is available.

```
names(smart_ohio_raw)
```

```
[1] "DISPCODE" "STATERE1" "SAFETIME" "HHADULT" "GENHLTH" "PHYSHLTH"
[7] "MENTHLTH" "POORHLTH" "HLTHPLN1" "PERSDOC2" "MEDCOST" "CHECKUP1"
[13] "BPHIGH4" "BPMEDS" "CHOLCHK1" "TOLDHI2" "CHOLMED1" "CVDINFR4"
[19] "CVDCRHD4" "CVDSTRK3" "ASTHMA3" "ASTHNOW" "CHCSCNCR" "CHCOCNCR"
[25] "CHCCOPD1" "HAVARTH3" "ADDEPEV2" "CHCKIDNY" "DIABETE3" "DIABAGE2"
[31] "LMTJOIN3" "ARTHDIS2" "ARTHSOCL" "JOINPAI1" "SEX" "MARITAL"
[37] "EDUCA" "RENTHOM1" "NUMHHOL2" "NUMPHON2" "CPDEMO1A" "VETERAN3"
[43] "EMPLOY1" "CHILDREN" "INCOME2" "INTERNET" "WEIGHT2" "HEIGHT3"
[49] "PREGNANT" "DEAF" "BLIND" "DECIDE" "DIFFWALK" "DIFFDRES"
[55] "DIFFALON" "SMOKE100" "SMOKDAY2" "STOPSMK2" "LASTSMK2" "USENOW3"
[61] "ECIGARET" "ECIGNOW" "ALCDAY5" "AVEDRNK2" "DRNK3GE5" "MAXDRNKS"
[67] "FRUIT2" "FRUITJU2" "FVGREEN1" "FRENCHF1" "POTATOE1" "VEGETAB2"
[73] "EXERANY2" "EXTRACT11" "EXEROFT1" "EXERHMM1" "EXRACT21" "EXEROFT2"
[79] "EXERHMM2" "STRENGTH" "SEATBELT" "FLUSHOT6" "FLSHTMY2" "PNEUVAC3"
[85] "SHINGLE2" "HIVTSTD6" "HIVTSTD3" "HIVRISK5" "CASTHDX2" "CASTHNO2"
[91] "CALLBCKZ" "WDUSENOW" "WDINFTRK" "WDHOWFT" "WDSHARE" "NAMTRIBE"
[97] "NAMOTHR" "URBNRRL" "STSTR" "IMPSEX" "RFHLTH" "PHYS14D"
[103] "MENT14D" "HCVU651" "RFHYPE5" "CHOLCH1" "RFCHOL1" "MICHD"
[109] "LTASTH1" "CASTHM1" "ASTHMS1" "DRDXAR1" "LMTACT1" "LMTWRK1"
[115] "LMTSCL1" "PRACE1" "MRACE1" "HISPANC" "RACE" "RACEG21"
[121] "RACEGR3" "AGEG5YR" "AGE65YR" "AGE80" "AGE_G" "WTKG3"
[127] "BMI5" "BMI5CAT" "RFBMI5" "EDUCAG" "INCOMG" "SMOKER3"
[133] "RFSMOK3" "ECIGSTS" "CURECIG" "DRNKANY5" "RFBING5" "DRNKWEK"
[139] "RFDRHV5" "FTJUDA2_" "FRUTDA2_" "GREND1_" "FRNCHDA_" "POTADA1_"
[145] "VEGEDA2_" "MISFRT1" "MISVEG1" "FRTRRES1" "VEGRES1" "FRUTSU1"
[151] "VEGESU1" "FRTLTT1A" "VEGLT1A" "FRT16A" "VEG23A" "FRUIT1"
```

```
[157] "_VEGETE1"  "_TOTINDA"  "_MINAC11"  "_MINAC21"  "_PACAT1"   "_PAINDEX1"
[163] "_PA150R2"  "_PA300R2"  "_PA30021"  "_PASTRNG"  "_PAREC1"   "_PASTAE1"
[169] "_RFSEAT2"  "_RFSEAT3"  "_FLSHOT6"  "_PNEUMO2"  "_AIDTST3"  "_MMSA"
[175] "_MMSAWT"    "SEQNO"     "MMSANAME"
```

2.5 Cleaning the BRFSS Data

2.5.1 Identifying Information

The identifying variables for each subject are gathered in `SEQNO`, which I'll leave alone.

- Each statistical (geographic) area is identified by a `_MMSA` variable, which I'll rename `mmsa_code`, and by an `MMSANAME` which I'll rename as `mmsa_name`
- For each subject, we are also provided with a sampling weight, in `_MMSAWT`, which will help us incorporate the sampling design later in the semester. We'll rename this as `mmsa_wt`. Details on the weighting methodology are available at https://www.cdc.gov/brfss/annual_data/2017/pdf/2017_SMART_BRFSS_MMSA_Methodology-508.pdf

```
smart_ohio_raw <- smart_ohio_raw %>%
  mutate(mmsa_code = `_MMSA`,
        mmsa_name = `MMSANAME`,
        mmsa_wt = `_MMSAWT`)

smart_ohio_raw %>% count(mmsa_code, mmsa_name)

# A tibble: 6 x 3
  mmsa_code mmsa_name                               n
  <dbl> <chr>                                         <int>
1 17140 Cincinnati, OH-KY-IN, Metropolitan Statistical Area 1737
2 17460 Cleveland-Elyria, OH, Metropolitan Statistical Area 1133
3 18140 Columbus, OH, Metropolitan Statistical Area 2033
4 19380 Dayton, OH, Metropolitan Statistical Area 587
5 26580 Huntington-Ashland, WV-KY-OH, Metropolitan Statistical Area 1156
6 45780 Toledo, OH, Metropolitan Statistical Area 766
```

Those names are very long. I'll build some shorter ones, by dropping everything after the comma.

```
smart_ohio_raw <- smart_ohio_raw %>%
  mutate(mmsa = str_replace_all(string = mmsa_name, pattern="\\", replacement=""))

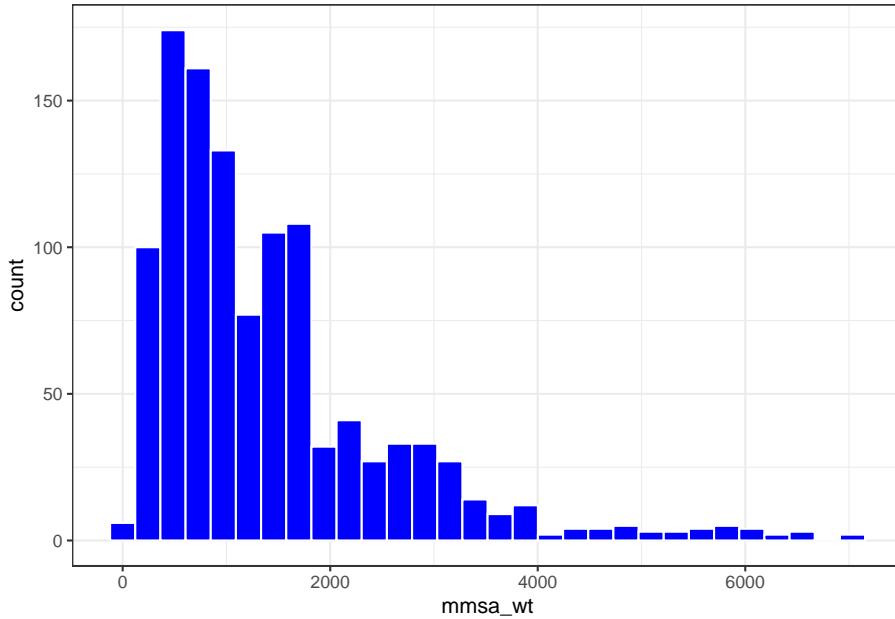
smart_ohio_raw %>% count(mmsa, mmsa_name)

# A tibble: 6 x 3
```

mmsa	mmsa_name	n
<chr>	<chr>	<int>
1 "Cincinnati "	Cincinnati, OH-KY-IN, Metropolitan Statistical Ar~	1737
2 "Cleveland-Elyria "	Cleveland-Elyria, OH, Metropolitan Statistical Ar~	1133
3 "Columbus "	Columbus, OH, Metropolitan Statistical Area	2033
4 "Dayton "	Dayton, OH, Metropolitan Statistical Area	587
5 "Huntington-Ashland "	Huntington-Ashland, WV-KY-OH, Metropolitan Statis~	1156
6 "Toledo "	Toledo, OH, Metropolitan Statistical Area	766

And here are the sampling weights for the subjects in the Cleveland-Elyria MSA.

```
smart_ohio_raw %>%
  filter(mmsa_code == 17460) %>%
  ggplot(., aes(x = mmsa_wt)) +
  geom_histogram(bins = 30, fill = "blue", col = "white")
```



2.5.2 Survey Method

2.5.2.1 DISPCODE and its cleanup to completed

DISPCODE which is 1100 if the subject completed the interview, and 1200 if they partially completed the interview. We'll create a variable called `completed` that indicates (1 = complete, 0 = not) whether the subject completed the interview.

```
smart_ohio_raw <- smart_ohio_raw %>%
  mutate(completed = 12 - (DISPCODE/100))

smart_ohio_raw %>% count(DISPCODE, completed)

# A tibble: 2 x 3
  DISPCODE completed     n
  <dbl>      <dbl> <int>
1     1100        1   6277
2     1200        0   1135
```

2.5.2.2 STATERE1 and SAFETIME and their reduction to landline

BRFSSS is conducted by telephone. The next two variables help us understand whether the subject was contacted via land line or via cellular phone.

- **STATERE1** is 1 if the subject is a resident of the state (only asked of people in the land line version of the survey).
- **SAFTIME** is 1 if this is a safe time to talk (only asked of people in the cell phone version of the survey).
- We'll use **STATERE1** and **SAFTIME** to create an indicator variable **landline** that specifies how the respondent was surveyed (1 = land line, 0 = cell phone), as follows...

```
smart_ohio_raw <- smart_ohio_raw %>%
  mutate(landline = replace_na(STATERE1, 0))

smart_ohio_raw %>% count(STATERE1, SAFETIME, landline)

# A tibble: 2 x 4
  STATERE1 SAFETIME landline     n
  <dbl>    <dbl>    <dbl> <int>
1       1       1       1   3649
2      NA       1       0   3763
```

2.5.2.3 HHADULT and its cleanup to hhadults

- **HHADULT** is the response to “How many members of your household, including yourself, are 18 years of age or older?”
 - The permitted responses range from 1-76, with special values 77 for Don’t Know/Not Sure and 99 for refused, with BLANK for missing or not asked.
 - So we should change all numerical values above 76 to NA for our analyses (the blanks are already regarded as NAs by R in the ingestion process.)

```

smart_ohio_raw %>% tabyl(HHADULT)

  HHADULT     n    percent valid_percent
  1   274 0.0369670804  0.236206897
  2   603 0.0813545602  0.519827586
  3   170 0.0229357798  0.146551724
  4    73 0.0098488937  0.062931034
  5    28 0.0037776579  0.024137931
  6     4 0.0005396654  0.003448276
  7     3 0.0004047491  0.002586207
  8     1 0.0001349164  0.000862069
  10    1 0.0001349164  0.000862069
  11    1 0.0001349164  0.000862069
  99    2 0.0002698327  0.001724138
NA 6252 0.8434970318          NA

smart_ohio_raw <- smart_ohio_raw %>%
  mutate(hhadults = HHADULT,
        hhadults = replace(hhadults, hhadults > 76, NA))

smart_ohio_raw %>% count(HHADULT, hhadults) %>% tail()

# A tibble: 6 x 3
  HHADULT hhadults     n
  <dbl>    <dbl> <int>
1      7       7     3
2      8       8     1
3     10      10     1
4     11      11     1
5     99      NA     2
6     NA      NA  6252

```

2.5.3 Health Status (1 item)

The next variable describes relate to the subject's health status.

2.5.3.1 GENHLTH and its cleanup to genhealth

- GENHLTH, the General Health variable, which is the response to “Would you say that in general your health is ...”
 - 1 = Excellent
 - 2 = Very good
 - 3 = Good
 - 4 = Fair
 - 5 = Poor

- 7 = Don't know/Not sure
- 9 = Refused
- BLANK = Not asked or missing

To clean up the GENHLTH data into a new variable called genhealth we'll need to - convince R that the 7 and 9 values are in fact best interpreted as NA, - and perhaps change the variable to a factor and incorporate the names into the levels.

```
smart_ohio_raw <- smart_ohio_raw %>%
  mutate(genhealth = fct_recode(factor(GENHLTH),
                                "1_Excellent" = "1",
                                "2_VeryGood" = "2",
                                "3_Good" = "3",
                                "4_Fair" = "4",
                                "5_Poor" = "5",
                                NULL = "7",
                                NULL = "9"))

smart_ohio_raw %>% count(GENHLTH, genhealth)

# A tibble: 7 x 3
  GENHLTH genhealth      n
  <dbl> <fct>     <int>
1       1 1_Excellent  1057
2       2 2_VeryGood   2406
3       3 3_Good       2367
4       4 4_Fair        1139
5       5 5_Poor        428
6       7 <NA>           10
7       9 <NA>            5
```

2.5.4 Healthy Days - Health-Related Quality of Life (3 items)

The next three variables describe the subject's health-related quality of life.

2.5.4.1 PHYSHLTH and its cleanup to physhealth

PHYSHLTH‘, the Number of Days Physical Health Not Good variable, which is the response to “Now thinking about your physical health, which includes physical illness and injury, for how many days during the past 30 days was your physical health not good?”

- Values of 1-30 are numeric and reasonable.
- A value of 88 indicates “none” and should be recoded to 0.

- 77 is the code for Don't know/Not sure
- 99 is the code for Refused
- BLANK indicates Not asked or missing, and R recognizes this as NA properly.

To clean up PHYSHLTH to a new variable called physhealth, we'll need: - to convince R that the 77 and 99 values are in fact best interpreted as NA, and - to convince R that the 88 should be interpreted as 0.

```
smart_ohio_raw <- smart_ohio_raw %>%
  mutate(physhealth = PHYSHLTH,
        physhealth = replace(physhealth, physhealth %in% c(77, 99), NA),
        physhealth = replace(physhealth, physhealth == 88, 0))

smart_ohio_raw %>% count(PHYSHLTH, physhealth) %>% tail()

# A tibble: 6 x 3
  PHYSHLTH physhealth     n
  <dbl>      <dbl> <int>
1     28       28     12
2     29       29     14
3     30       30    677
4     77       NA    123
5     88       0    4380
6     99       NA     15
```

Note that we present the `tail` of the counts in this case so we can see what happens to the key values (77, 88, 99) of our original variable PHYSHLTH.

2.5.4.2 MENTHLTH and its cleanup to menthealth

MENTHLTH', the Number of Days Mental Health Not Good variable, which is the response to "Now thinking about your mental health, which includes stress, depression, and problems with emotions, for how many days during the past 30 days was your mental health not good?"

- This is coded just like the PHYSHLTH variable, so we need to do the same cleaning we did there.

To clean up MENTHLTH to a new variable called menthealth, we'll need: - to convince R that the 77 and 99 values are in fact best interpreted as NA, and - to convince R that the 88 should be interpreted as 0.

```
smart_ohio_raw <- smart_ohio_raw %>%
  mutate(menthealth = MENTHLTH,
        menthealth = replace(menthealth, menthealth %in% c(77, 99), NA),
        menthealth = replace(menthealth, menthealth == 88, 0))
```

```
smart_ohio_raw %>% count(MENTHLTH, menthealth) %>% tail()

# A tibble: 6 x 3
  MENTHLTH menthealth     n
    <dbl>      <dbl> <int>
1     28        28     7
2     29        29    10
3     30        30   475
4     77        NA    86
5     88         0  4823
6     99        NA    28
```

2.5.4.3 POORHLTH and its cleanup to poorhealth

POORHLTH, the Poor Physical or Mental Health variable, which is the response to “During the past 30 days, for about how many days did poor physical or mental health keep you from doing your usual activities, such as self-care, work, or recreation?”

- Again, we recode just like the PHYSHLTH variable.

```
smart_ohio_raw <- smart_ohio_raw %>%
  mutate(poorhealth = POORHLTH,
        poorhealth = replace(poorhealth, poorhealth %in% c(77, 99), NA),
        poorhealth = replace(poorhealth, poorhealth == 88, 0))

smart_ohio_raw %>% count(POORHLTH, poorhealth) %>% tail()

# A tibble: 6 x 3
  POORHLTH poorhealth     n
    <dbl>      <dbl> <int>
1     29        29     4
2     30        30   382
3     77        NA    64
4     88         0  2194
5     99        NA    11
6     NA        NA  3337
```

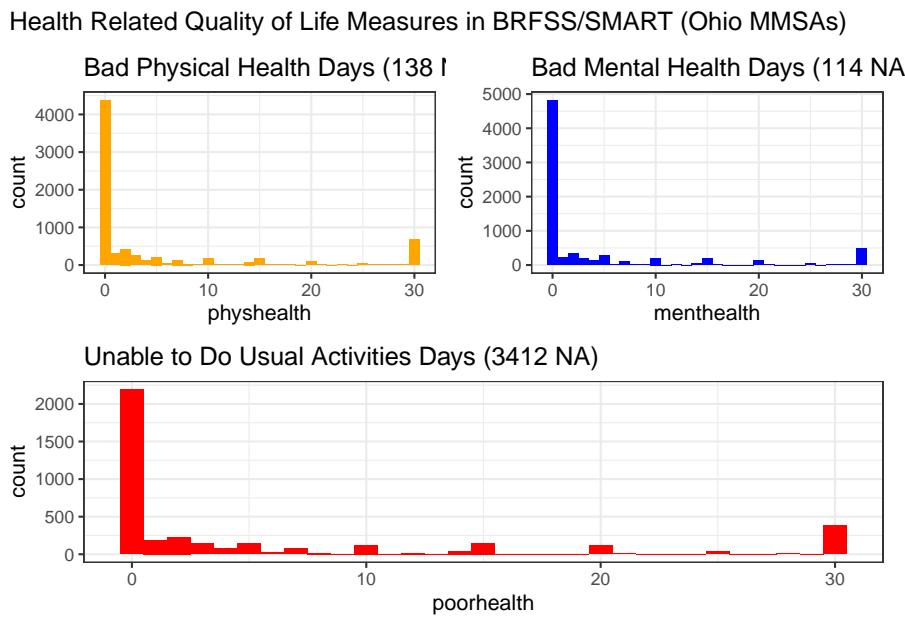
There's a lot more missingness in the poorhealth counts than in the other health-related quality of life measures. There's also a strong mode at 0, and a smaller mode at 30 in each variable.

```
p1 <- ggplot(smart_ohio_raw, aes(x = physhealth)) +
  geom_histogram(binwidth = 1, fill = "orange") +
  labs(title = paste0("Bad Physical Health Days (",
                     sum(is.na(smart_ohio_raw$physhealth))),
```

```

    " NA"))
p2 <- ggplot(smart_ohio_raw, aes(x = menthealth)) +
  geom_histogram(binwidth = 1, fill = "blue") +
  labs(title = paste0("Bad Mental Health Days (",
                     sum(is.na(smart_ohio_raw$menthealth)),
                     " NA")))
p3 <- ggplot(smart_ohio_raw, aes(x = poorhealth)) +
  geom_histogram(binwidth = 1, fill = "red") +
  labs(title = paste0("Unable to Do Usual Activities Days (",
                     sum(is.na(smart_ohio_raw$poorhealth)),
                     " NA")))
(p1 + p2) / p3 +
  plot_annotation(title = "Health Related Quality of Life Measures in BRFSS/SMART (Ohio MMSAs)")

```



2.5.5 Health Care Access (4 items)

The next four variables relate to the subject's health care access.

2.5.5.1 HLTHPLN1 and its cleanup to healthplan

HLTHPLN1, the Have any health care coverage variable, is the response to “Do you have any kind of health care coverage, including health insurance, prepaid plans such as HMOs, or government plans such as Medicare, or Indian Health Service?”

- 1 = Yes
- 2 = No
- 7 = Don’t know/Not sure
- 9 = Refused

To clean up the HLTHPLN1 data into a new variable called `healthplan` we’ll - convince R that the 7 and 9 values are in fact best interpreted as `NA`, - and turn it into an indicator variable, e.g., we will leave the variable as numeric, but change the values to 1 = Yes and 0 = No.

```
smart_ohio_raw <- smart_ohio_raw %>%
  mutate(healthplan = HLTHPLN1,
        healthplan = replace(healthplan, healthplan %in% c(7, 9), NA),
        healthplan = replace(healthplan, healthplan == 2, 0))

smart_ohio_raw %>% count(HLTHPLN1, healthplan)

# A tibble: 4 x 3
  HLTHPLN1 healthplan     n
    <dbl>      <dbl> <int>
1       1          1   6994
2       2          0    398
3       7         NA     10
4       9         NA     10
```

2.5.5.2 PERSDOC2 and its cleanup to hasdoc and to numdocs2

PERSDOC2, the Multiple Health Care Professionals variable, is the response to “Do you have one person you think of as your personal doctor or health care provider?” where if the response is “No,” the survey then asks “Is there more than one or is there no person who you think of as your personal doctor or health care provider?”

- 1 = Yes, only one
- 2 = More than one
- 3 = No
- 7 = Don’t know/Not sure
- 9 = Refused
- BLANK = Not asked or missing

To clean up the PERSDOC2 data into a new variable called `hasdoc` we'll - convince R that the 7 and 9 values are in fact best interpreted as NA, - and turn it into an indicator variable, e.g., we will leave the variable as numeric, but change the values to 1 = Yes and 0 = No, so that the original 1 and 2 become 1, and the original 3 becomes 0.

```
smart_ohio_raw <- smart_ohio_raw %>%
  mutate(hasdoc = PERSDOC2,
        hasdoc = replace(hasdoc, hasdoc %in% c(7, 9), NA),
        hasdoc = replace(hasdoc, hasdoc %in% c(1, 2), 1),
        hasdoc = replace(hasdoc, hasdoc == 3, 0))

smart_ohio_raw %>% count(PERSDOC2, hasdoc)

# A tibble: 5 x 3
  PERSDOC2 hasdoc     n
  <dbl>    <dbl> <int>
1       1      1   5784
2       2      1    623
3       3      0    990
4       7     NA     14
5       9     NA      1
```

2.5.5.3 MEDCOST and its cleanup to costprob

MEDCOST, the Could Not See Doctor Because of Cost variable, is the response to "Was there a time in the past 12 months when you needed to see a doctor but could not because of cost?"

- 1 = Yes
- 2 = No
- 7 = Don't know/Not sure
- 9 = Refused
- BLANK = Not asked or missing

This is just like HLTHPLAN.

```
smart_ohio_raw <- smart_ohio_raw %>%
  mutate(costprob = MEDCOST,
        costprob = replace(costprob, costprob %in% c(7, 9), NA),
        costprob = replace(costprob, costprob == 2, 0))

smart_ohio_raw %>% count(MEDCOST, costprob)

# A tibble: 4 x 3
  MEDCOST costprob     n
  <dbl>    <dbl> <int>
1       1      1   5784
2       2      0    623
3       7     NA     14
4       9     NA      1
```

	1	2	0	NA	14
1	1	2	0	NA	714
2					6680
3					14
4					4

2.5.5.4 CHECKUP1 and its cleanup to t_checkup

CHECKUP1, the Length of time since last routine checkup variable, is the response to “About how long has it been since you last visited a doctor for a routine checkup? [A routine checkup is a general physical exam, not an exam for a specific injury, illness, or condition.]”

- 1 = Within past year (anytime less than 12 months ago)
- 2 = Within past 2 years (1 year but less than 2 years ago)
- 3 = Within past 5 years (2 years but less than 5 years ago)
- 4 = 5 or more years ago
- 7 = Don’t know/Not sure
- 8 = Never
- 9 = Refused
- BLANK = Not asked or missing

To clean up the CHECKUP1 data into a new variable called t_checkup we’ll - convince R that the 7 and 9 values are in fact best interpreted as NA, - relabel options 1, 2, 3, 4 and 8 while turning the variable into a factor.

```
smart_ohio_raw <- smart_ohio_raw %>%
  mutate(t_checkup = fct_recode(factor(CHECKUP1),
    "1_In-past-year" = "1",
    "2_1-to-2-years" = "2",
    "3_2-to-5-years" = "3",
    "4_5_plus_years" = "4",
    "8_Never" = "8",
    NULL = "7",
    NULL = "9"))

smart_ohio_raw %>% count(CHECKUP1, t_checkup)

# A tibble: 7 x 3
  CHECKUP1 t_checkup      n
  <dbl> <fct>     <int>
1       1 1_In-past-year  5803
2       2 2_1-to-2-years   714
3       3 3_2-to-5-years   413
4       4 4_5_plus_years   376
5       7 <NA>            68
6       8 8_Never           32
7       9 <NA>              6
```

2.5.6 Blood Pressure (2 measures)

2.5.6.1 BPHIGH4 and its cleanup to bp_high

BPHIGH4 is asking about awareness of a hypertension diagnosis. It's the response to the question: "Have you EVER been told by a doctor, nurse or other health professional that you have high blood pressure?" In addition, if the answer was "Yes" and the respondent is female, they were then asked "Was this only when you were pregnant?"

The available codes are:

- 1 = Yes
- 2 = Yes, but female told only during pregnancy
- 3 = No
- 4 = Told borderline high or pre-hypertensive
- 7 = Don't know/Not sure
- 9 = Refused
- BLANK = Not asked or missing

To clean up the BPHIGH4 data into a new variable called `bp_high` we'll - convince R that the 7 and 9 values are in fact best interpreted as NA, - relabel (and re-order) options 1, 2, 3, 4 while turning the variable into a factor.

```
smart_ohio_raw <- smart_ohio_raw %>%
  mutate(bp_high = fct_recode(factor(BPHIGH4),
    "0_No" = "3",
    "1_Yes" = "1",
    "2_Only_while_pregnant" = "2",
    "4_Borderline" = "4",
    NULL = "7",
    NULL = "9"),
  bp_high = fct_relevel(bp_high,
    "0_No", "1_Yes",
    "2_Only_while_pregnant",
    "4_Borderline"))

smart_ohio_raw %>% count(BPHIGH4, bp_high)

# A tibble: 6 x 3
  BPHIGH4 bp_high          n
  <dbl> <fct>        <int>
1     1 1_Yes        3161
2     2 2_Only_while_pregnant  67
3     3 0_No         4114
4     4 4_Borderline   49
5     7 <NA>          19
6     9 <NA>           2
```

2.5.6.2 BPMEDS and its cleanup to bp_meds

BPMEDS is the response to the question “Are you currently taking medicine for your high blood pressure?”

- 1 = Yes
- 2 = No
- 7 = Don’t know/Not sure
- 9 = Refused
- BLANK = Not asked or missing

To clean up the BPMEDS data into a new variable called `bp_meds` we’ll treat it just as we did with `HLTHPLN1` and - convince R that the 7 and 9 values are in fact best interpreted as `NA`, - and turn it into an indicator variable, e.g., we will leave the variable as numeric, but change the values to 1 = Yes and 0 = No.

```
smart_ohio_raw <- smart_ohio_raw %>%
  mutate(bp_meds = BPMEDS,
        bp_meds = replace(bp_meds, bp_meds %in% c(7, 9), NA),
        bp_meds = replace(bp_meds, bp_meds == 2, 0))

smart_ohio_raw %>% count(BPMEDS, bp_meds)
```

```
# A tibble: 5 x 3
  BPMEDS bp_meds     n
  <dbl>   <dbl> <int>
1      1       1  2675
2      2       0   481
3      7      NA    4
4      9      NA    1
5     NA      NA 4251
```

What is the relationship between our two blood pressure variables? Only the people with `bp_meds = “1_Yes”` were asked the `bp_meds` question.

```
smart_ohio_raw %>% tabyl(bp_high, bp_meds)
```

	bp_high	0	1	NA_
0_No	0	0	4114	
1_Yes	481	2675	5	
2_Only_while_pregnant	0	0	67	
4_Borderline	0	0	49	
<NA>	0	0	21	

2.5.7 Cholesterol (3 items)

2.5.7.1 CHOLCHK1 and its cleanup to t_chol

CHOLCHK1, the Length of time since cholesterol was checked, is the response to “Blood cholesterol is a fatty substance found in the blood. About how long has it been since you last had your blood cholesterol checked?”

- 1 = Never
- 2 = Within past year (anytime less than 12 months ago)
- 3 = Within past 2 years (1 year but less than 2 years ago)
- 4 = Within past 5 years (2 years but less than 5 years ago)
- 5 = 5 or more years ago
- 7 = Don’t know/Not sure
- 9 = Refused
- BLANK = Not asked or missing

To clean up the CHOLCHK1 data into a new variable called t_chol we’ll - convince R that the 7 and 9 values are in fact best interpreted as NA, - relabel options 1, 2, 3, 4 and 8 while turning the variable into a factor.

```
smart_ohio_raw <- smart_ohio_raw %>%
  mutate(t_chol = fct_recode(factor(CHOLCHK1),
    "1_Never" = "1",
    "2_In-past-year" = "2",
    "3_1-to-2-years" = "3",
    "4_2-to-5-years" = "4",
    "5_5_plus_years" = "5",
    NULL = "7",
    NULL = "9"))

smart_ohio_raw %>% count(CHOLCHK1, t_chol)

# A tibble: 8 x 3
  CHOLCHK1 t_chol      n
  <dbl> <fct>     <int>
1       1 1_Never     424
2       2 2_In-past-year 5483
3       3 3_1-to-2-years  559
4       4 4_2-to-5-years   289
5       5 5_5_plus_years   272
6       7 <NA>          376
7       9 <NA>            8
8      NA <NA>           1
```

The next two measures are not gathered from the people who answered “Never” to this question.

2.5.7.2 TOLDHI2 and its cleanup to chol_high

TOLDHI2 is asking about awareness of a diagnosis of high cholesterol. It's the response to the question: "Have you EVER been told by a doctor, nurse or other health professional that your blood cholesterol is high?"

The available codes are:

- 1 = Yes
- 2 = No
- 7 = Don't know/Not sure
- 9 = Refused
- BLANK = Not asked or missing

To clean up the TOLDHI2 data into a new variable called `chol_high` we'll treat it like BPMEDS and HLTHPLN1 - convince R that the 7 and 9 values are in fact best interpreted as NA, - and turn it into an indicator variable, e.g., we will leave the variable as numeric, but change the values to 1 = Yes and 0 = No.

```
smart_ohio_raw <- smart_ohio_raw %>%
  mutate(chol_high = TOLDHI2,
        chol_high = replace(chol_high, chol_high %in% c(7, 9), NA),
        chol_high = replace(chol_high, chol_high == 2, 0))

smart_ohio_raw %>% count(TOLDHI2, chol_high)

# A tibble: 5 x 3
  TOLDHI2 chol_high     n
  <dbl>      <dbl> <int>
1       1          1  2612
2       2          0  4286
3       7         NA    70
4       9         NA     4
5      NA         NA  440
```

2.5.7.3 CHOLMED1 and its cleanup to chol_meds

CHOLMED1 is the response to the question "Are you currently taking medicine prescribed by a doctor or other health professional for your blood cholesterol?"

- 1 = Yes
- 2 = No
- 7 = Don't know/Not sure
- 9 = Refused
- BLANK = Not asked or missing

To clean up the CHOLMED1 data into a new variable called `chol_meds` we'll treat it just as we did with HLTHPLN1 and - convince R that the 7 and 9 values are in

fact best interpreted as NA, - and turn it into an indicator variable, e.g., we will leave the variable as numeric, but change the values to 1 = Yes and 0 = No.

```
smart_ohio_raw <- smart_ohio_raw %>%
  mutate(chol_meds = CHOLMED1,
        chol_meds = replace(chol_meds, chol_meds %in% c(7, 9), NA),
        chol_meds = replace(chol_meds, chol_meds == 2, 0))

smart_ohio_raw %>% count(CHOLMED1, chol_meds)

# A tibble: 4 x 3
  CHOLMED1 chol_meds     n
  <dbl>      <dbl> <int>
1       1          1   1781
2       2          0    826
3       7         NA     5
4      NA         NA  4800
```

2.5.8 Chronic Health Conditions (14 items)

2.5.8.1 Self-reported diagnosis history (11 items)

The next few variables describe whether or not the subject meets a particular standard, and are all coded in the raw data the same way:

- 1 = Yes
- 2 = No
- 7 = Don't know/Not sure
- 9 = Refused
- BLANK = Not asked or missing

and we'll recode them all to 1 = Yes, 0 = No, otherwise NA, as we've done previously.

The questions are all started with “Has a doctor, nurse, or other health professional ever told you that you had any of the following? For each, tell me Yes, No, or you're Not sure.”

Original	Revised	Details
CVDINFR4	hx_mi	(Ever told) you had a heart attack, also called a myocardial infarction?
CVDCRHD4	hx_chd	(Ever told) you had angina or coronary heart disease?
CVDSTRK3	hx_stroke	(Ever told) you had a stroke?
ASTHMA3	hx_asthma	(Ever told) you had asthma?
ASTHNOW	now_asthma	Do you still have asthma? (only asked of those with Yes in ASTHMA3)

Original	Revised	Details
CHCSCNCR	hx_skinc	(Ever told) you had skin cancer?
CHCOCNCR	hx_otherc	(Ever told) you had any other types of cancer?
CHCCOPD1	hx_copd	(Ever told) you have Chronic Obstructive Pulmonary Disease or COPD, emphysema or chronic bronchitis?
HAVARTH3	hx_arthr	(Ever told) you have some form of arthritis, rheumatoid arthritis, gout, lupus, or fibromyalgia? (Arthritis diagnoses include: rheumatism, polymyalgia rheumatica; osteoarthritis (not osteoporosis); tendonitis, bursitis, bunion, tennis elbow; carpal tunnel syndrome, tarsal tunnel syndrome; joint infection, etc.)
ADDEPEV2	hx_depress	(Ever told) you that you have a depressive disorder, including depression, major depression, dysthymia, or minor depression?
CHCKIDNY	hx_kidney	(Ever told) you have kidney disease? Do NOT include kidney stones, bladder infection or incontinence.

```
smart_ohio_raw <- smart_ohio_raw %>%
  mutate(hx_mi = CVDINFR4,
        hx_mi = replace(hx_mi, hx_mi %in% c(7, 9), NA),
        hx_mi = replace(hx_mi, hx_mi == 2, 0),
        hx_chd = CVDCRHD4,
        hx_chd = replace(hx_chd, hx_chd %in% c(7, 9), NA),
        hx_chd = replace(hx_chd, hx_chd == 2, 0),
        hx_stroke = CVDSTRK3,
        hx_stroke = replace(hx_stroke, hx_stroke %in% c(7, 9), NA),
        hx_stroke = replace(hx_stroke, hx_stroke == 2, 0),
        hx_asthma = ASTHMA3,
        hx_asthma = replace(hx_asthma, hx_asthma %in% c(7, 9), NA),
        hx_asthma = replace(hx_asthma, hx_asthma == 2, 0),
        now_asthma = ASTHNOW,
        now_asthma = replace(now_asthma, now_asthma %in% c(7, 9), NA),
        now_asthma = replace(now_asthma, now_asthma == 2, 0),
        hx_skinc = CHCSCNCR,
        hx_skinc = replace(hx_skinc, hx_skinc %in% c(7, 9), NA),
        hx_skinc = replace(hx_skinc, hx_skinc == 2, 0),
        hx_otherc = CHCOCNCR,
        hx_otherc = replace(hx_otherc, hx_otherc %in% c(7, 9), NA),
        hx_otherc = replace(hx_otherc, hx_otherc == 2, 0),
        hx_copd = CHCCOPD1,
        hx_copd = replace(hx_copd, hx_copd %in% c(7, 9), NA),
        hx_copd = replace(hx_copd, hx_copd == 2, 0),
        hx_arthr = HAVARTH3,
```

```

hx_arthr = replace(hx_arthr, hx_arthr %in% c(7, 9), NA),
hx_arthr = replace(hx_arthr, hx_arthr == 2, 0),
hx_depress = ADDEPEV2,
hx_depress = replace(hx_depress, hx_depress %in% c(7, 9), NA),
hx_depress = replace(hx_depress, hx_depress == 2, 0),
hx_kidney = CHCKIDNY,
hx_kidney = replace(hx_kidney, hx_kidney %in% c(7, 9), NA),
hx_kidney = replace(hx_kidney, hx_kidney == 2, 0))

```

We definitely should have written a function to do that, of course.

2.5.8.2 `_ASTHMS1` and its cleanup to `asthma`

`_ASTHMS1` categorizes subjects by asthma status as:

- 1 = Current
- 2 = Former
- 3 = Never
- 9 = Don't Know / Not Sure / Refused / Missing

We'll turn this into a factor with appropriate levels and NA information.

```

smart_ohio_raw <- smart_ohio_raw %>%
  mutate(asthma = fct_recode(
    factor(`_ASTHMS1`),
    "Current" = "1",
    "Former" = "2",
    "Never" = "3",
    NULL = "9"))

smart_ohio_raw %>% count(`_ASTHMS1`, asthma)

# A tibble: 4 x 3
`_ASTHMS1` asthma     n
<dbl> <fct>   <int>
1      1 Current    734
2      2 Former     248
3      3 Never     6376
4      9 <NA>       54

```

2.5.8.3 `DIABETE3` and its cleanup to `hx_diabetes` and `dm_status`

`DIABETE3`, the (Ever told) you have diabetes variable, is the response to “(Ever told) you have diabetes (If Yes and respondent is female, ask Was this only when you were pregnant?. If Respondent says pre-diabetes or borderline diabetes, use response code 4.)”

- 1 = Yes
- 2 = Yes, but female told only during pregnancy
- 3 = No
- 4 = No, pre-diabetes or borderline diabetes
- 7 = Don't know/Not sure
- 9 = Refused
- BLANK = Not asked or missing

I'll create one variable called `hx_diabetes` which is 1 if `DIABETE3` = 1, and 0 otherwise, with appropriate NAs, like our other variables. Then I'll create `dm_status` to include all of this information in a factor, but again recode the missing values properly.

```
smart_ohio_raw <- smart_ohio_raw %>%
  mutate(hx_diabetes = DIABETE3,
        hx_diabetes = replace(hx_diabetes, hx_diabetes %in% c(7, 9), NA),
        hx_diabetes = replace(hx_diabetes, hx_diabetes %in% 2:4, 0),
        dm_status = fct_recode(factor(DIABETE3),
                               "Diabetes" = "1",
                               "Pregnancy-Induced" = "2",
                               "No-Diabetes" = "3",
                               "Pre-Diabetes" = "4",
                               NULL = "7",
                               NULL = "9"),
        dm_status = fct_relevel(dm_status,
                               "No-Diabetes",
                               "Pre-Diabetes",
                               "Pregnancy-Induced",
                               "Diabetes"))

smart_ohio_raw %>% count(DIABETE3, hx_diabetes, dm_status)

# A tibble: 6 x 4
  DIABETE3 hx_diabetes dm_status      n
  <dbl>     <dbl> <fct>       <int>
1       1         1 Diabetes      1098
2       2         0 Pregnancy-Induced    67
3       3         0 No-Diabetes    6100
4       4         0 Pre-Diabetes   133
5       7        NA <NA>          12
6       9        NA <NA>           2
```

2.5.8.4 DIABAGE2 and its cleanup to dm_age

`DIABAGE2`, the Age When Told Diabetic variable, is the response to “How old were you when you were told you have diabetes?” It is asked only of people with

`DIABETE3 = 1` (Yes).

- The response is 1-97, with special values 98 for Don't Know/Not Sure and 99 for refused, with BLANK for missing or not asked. People 97 years of age and above were listed as 97.

```
smart_ohio_raw <- smart_ohio_raw %>%
  mutate(dm_age = DIABAGE2,
        dm_age = replace(dm_age, dm_age > 97, NA))

smart_ohio_raw %>% count(DIABAGE2, dm_age) %>% tail()

# A tibble: 6 x 3
  DIABAGE2 dm_age     n
  <dbl>    <dbl> <int>
1       84      84     1
2       85      85     2
3       90      90     1
4       98      NA    61
5       99      NA     4
6       NA      NA   6314
```

2.5.9 Arthritis Burden (4 items)

The first two measures are only asked of people with `hx_arthr = 1`, and are coded as:

- 1 = Yes
- 2 = No
- 7 = Don't know/Not sure
- 9 = Refused
- BLANK = Not asked or missing

and we'll recode them to 1 = Yes, 0 = No, otherwise NA, as we've done previously.

2.5.9.1 LMTJOIN3 (Limited because of joint symptoms), and its cleanup to `arth_lims`

This is the response to “Are you now limited in any way in any of your usual activities because of arthritis or joint symptoms?”

```
smart_ohio_raw <- smart_ohio_raw %>%
  mutate(arth_lims = LMTJOIN3,
        arth_lims = replace(arth_lims, arth_lims %in% c(7, 9), NA),
        arth_lims = replace(arth_lims, arth_lims == 2, 0))

smart_ohio_raw %>% count(hx_arthr, LMTJOIN3, arth_lims)
```

```
# A tibble: 6 x 4
  hx_arthr LMTJOIN3 arth_lims     n
  <dbl>      <dbl>    <dbl> <int>
1 0          NA        NA  4587
2 1          1         1   1378
3 1          2         0   1388
4 1          7         NA   17
5 1          9         NA   2
6 NA         NA        NA  40
```

2.5.9.2 ARTHDIS2 (Does Arthritis Affect Whether You Work), and its cleanup to arth_work

This is the response to “Do arthritis or joint symptoms now affect whether you work, the type of work you do or the amount of work you do?”

```
smart_ohio_raw <- smart_ohio_raw %>%
  mutate(arth_work = ARTHDIS2,
        arth_work = replace(arth_work, arth_work %in% c(7, 9), NA),
        arth_work = replace(arth_work, arth_work == 2, 0))

smart_ohio_raw %>% count(ARTHDIS2, arth_work)

# A tibble: 5 x 3
  ARTHDIS2 arth_work     n
  <dbl>      <dbl> <int>
1 1          1     925
2 2          0    1808
3 7          NA    42
4 9          NA    10
5 NA         NA  4627
```

2.5.9.3 ARTHSOCL (Social Activities Limited Because of Joint Symptoms) and its cleanup to arth_soc

This is the response to “During the past 30 days, to what extent has your arthritis or joint symptoms interfered with your normal social activities, such as going shopping, to the movies, or to religious or social gatherings?”

The responses are:

- 1 = A lot
- 2 = A little
- 3 = Not at all
- 7 = Don't know/Not sure
- 9 = Refused

- BLANK = Not asked or missing

```
smart_ohio_raw <- smart_ohio_raw %>%
  mutate(arth_soc = fct_recode(factor(ARTHSOCL),
                               "A lot" = "1",
                               "A little" = "2",
                               "Not at all" = "3",
                               NULL = "7",
                               NULL = "9"))

smart_ohio_raw %>% count(ARTHSOCL, arth_soc)

# A tibble: 6 x 3
  ARTHSOCL arth_soc     n
  <dbl> <fct>     <int>
1       1 A lot      606
2       2 A little    734
3       3 Not at all  1427
4       7 <NA>        15
5       9 <NA>         3
6      NA <NA>      4627
```

2.5.9.4 JOINPAI1 (How Bad Was Joint Pain - scale of 0-10) and its cleanup to joint_pain

This is the response to the following question: “Please think about the past 30 days, keeping in mind all of your joint pain or aching and whether or not you have taken medication. On a scale of 0 to 10 where 0 is no pain or aching and 10 is pain or aching as bad as it can be, DURING THE PAST 30 DAYS, how bad was your joint pain ON AVERAGE?”

The available values are 0-10, plus codes 77 (Don’t Know / Not Sure), 99 (Refused) and BLANK.

To clean up JOINPAI1 to a new variable called `joint_pain`, we’ll need to convince R that the 77 and 99 values are, like BLANK, in fact best interpreted as `NA`.

```
smart_ohio_raw <- smart_ohio_raw %>%
  mutate(joint_pain = JOINPAI1,
        joint_pain = replace(joint_pain, joint_pain %in% c(77, 99), NA))

smart_ohio_raw %>% count(JOINPAI1, joint_pain) %>% tail()

# A tibble: 6 x 3
  JOINPAI1 joint_pain     n
  <dbl>     <dbl> <int>
1       8         8   277
```

2	9	9	72
3	10	10	158
4	77	NA	28
5	99	NA	5
6	NA	NA	4627

2.5.10 Demographics (25 items)

2.5.10.1 `_AGEG5YR`, which we'll edit into `agegroup`

The `_AGEG5YR` variable is a calculated variable (by CDC) obtained from the subject's age. Since the `age` data are not available, we instead get these groupings, which we'll rearrange into the `agegroup` factor.

<code>_AGEG5YR</code>	Age range	<code>agegroup</code>
1	18 <= AGE <= 24	18-24
2	25 <= AGE <= 29	25-29
3	30 <= AGE <= 34	30-34
4	35 <= AGE <= 39	35-39
5	40 <= AGE <= 44	40-44
6	45 <= AGE <= 49	45-49
7	50 <= AGE <= 54	50-54
8	55 <= AGE <= 59	55-59
9	60 <= AGE <= 64	60-64
10	65 <= AGE <= 69	65-69
11	70 <= AGE <= 74	70-74
12	75 <= AGE <= 79	75-79
13	AGE >= 80	80plus
14	Don't Know, Refused or Missing	NA

```
smart_ohio_raw <- smart_ohio_raw %>%
  mutate(agegroup = fct_recode(factor(`_AGEG5YR`),
    "18-24" = "1",
    "25-29" = "2",
    "30-34" = "3",
    "35-39" = "4",
    "40-44" = "5",
    "45-49" = "6",
    "50-54" = "7",
    "55-59" = "8",
    "60-64" = "9",
    "65-69" = "10",
    "70-74" = "11",
```

```

    "75-79" = "12",
    "80-96" = "13",
    NULL = "14"))

smart_ohio_raw %>% count(`_AGEG5YR`, agegroup)

# A tibble: 14 x 3
`_AGEG5YR` agegroup     n
<dbl> <fct>     <int>
1     1 18-24     448
2     2 25-29     327
3     3 30-34     375
4     4 35-39     446
5     5 40-44     426
6     6 45-49     509
7     7 50-54     604
8     8 55-59     786
9     9 60-64     837
10    10 65-69    810
11    11 70-74    685
12    12 75-79    499
13    13 80-96    592
14    14 <NA>      68

```

2.5.10.2 `_MRACE1` recoded to `race`

We'll create three variables describing race/ethnicity. The first comes from the `_MRACE1` variable categorized by CDC, and the available responses are:

- 1 = White only
- 2 = Black or African-American only
- 3 = American Indian or Alaskan Native only
- 4 = Asian only
- 5 = Native Hawaiian or Pacific Islander only
- 6 = Other race only
- 7 = Multiracial
- 77 = Don't know / Not Sure
- 99 = Refused
- BLANK = Missing

We'll create a factor out of this information, with appropriate level names.

```

smart_ohio_raw <- smart_ohio_raw %>%
  mutate(race = fct_recode(factor(`_MRACE1`),
    "White" = "1",
    "Black or African A" = "2",
    "American Indian or Alaskan Native" = "3",
    "Asian" = "4",
    "Native Hawaiian or Pacific Islander" = "5",
    "Other race" = "6",
    "Multiracial" = "7",
    "Don't know / Not Sure" = "77",
    "Refused" = "99",
    "Missing" = ""))

```

```

    "Amer Indian or Alaskan" = "3",
    "Asian" = "4",
    "Hawaiian or Pac Island" = "5",
    "Other Race" = "6",
    "Multiracial" = "7",
    NULL = "77",
    NULL = "99"))

smart_ohio_raw %>% count(`_MRACE1`, race)

# A tibble: 9 x 3
`_MRACE1` race          n
<dbl> <fct>      <int>
1       1 White        6177
2       2 Black or African A   739
3       3 Amer Indian or Alaskan  66
4       4 Asian         115
5       5 Hawaiian or Pac Island  5
6       6 Other Race     43
7       7 Multiracial    153
8       77 <NA>          14
9       99 <NA>         100

```

2.5.10.3 `_HISPANC` recoded to `hispanic`

The `_HISPANC` variable specifies whether or not the respondent is of Hispanic or Latinx origin. The available responses are:

- 1 = Hispanic, Latinx or Spanish origin
- 2 = Not of Hispanic, Latinx or Spanish origin
- 9 = Don't Know, Refused, or Missing

We'll turn the 9s into NA, and create an indicator variable (1 = Hispanic or Latinx, 0 = not)

```

smart_ohio_raw <- smart_ohio_raw %>%
  mutate(hispanic = 2 - `_HISPANC`,
        hispanic = replace(hispanic, hispanic < 0, NA))

smart_ohio_raw %>% count(`_HISPANC`, hispanic)

# A tibble: 3 x 3
`_HISPANC` hispanic     n
<dbl>      <dbl> <int>
1           1          1   146
2           2          0  7217

```

3	9	NA	49
---	---	----	----

2.5.10.4 `_RACEGR3` recoded to `race_eth`

The `_RACEGR3` variable is a five-level combination of race and ethnicity. The responses are:

- 1 = White non-Hispanic
- 2 = Black non-Hispanic
- 3 = Other race non-Hispanic
- 4 = Multiracial non-Hispanic
- 5 = Hispanic
- 9 = Don't Know / Not Sure / Refused

We'll create a factor out of this information, with appropriate level names.

```
smart_ohio_raw <- smart_ohio_raw %>%
  mutate(race_eth = fct_recode(
    factor(`_RACEGR3`),
    "White non-Hispanic" = "1",
    "Black non-Hispanic" = "2",
    "Other race non-Hispanic" = "3",
    "Multiracial non-Hispanic" = "4",
    "Hispanic" = "5",
    NULL = "9"))

smart_ohio_raw %>% count(`_RACEGR3`, race_eth)

# A tibble: 6 x 3
  `_RACEGR3` race_eth          n
  <dbl> <fct>        <int>
1 1 White non-Hispanic 6086
2 2 Black non-Hispanic 725
3 3 Other race non-Hispanic 193
4 4 Multiracial non-Hispanic 143
5 5 Hispanic           146
6 9 <NA>                119
```

2.5.10.5 `SEX` recoded to `female`

The available levels of `SEX` are:

- 1 = Male
- 2 = Female
- 9 = Refused

We'll recode that to `female = 1` for Female, 0 Male, otherwise NA. Note the trick here is to subtract one from the coded `SEX` to get the desired `female`, but this requires that we move 8 to NA, rather than 9.

```
smart_ohio_raw <- smart_ohio_raw %>%
  mutate(female = SEX - 1,
        female = replace(female, female == 8, NA))

smart_ohio_raw %>% count(SEX, female)

# A tibble: 2 x 3
  SEX   female     n
  <dbl> <dbl> <int>
1     1       0  3136
2     2       1  4276
```

2.5.10.6 MARITAL status, revised to marital

The available levels of MARITAL are:

- 1 = Married
- 2 = Divorced
- 3 = Widowed
- 4 = Separated
- 5 = Never married
- 6 = A member of an unmarried couple
- 9 = Refused
- BLANK = Not asked or missing

We'll just turn this into a factor, and move 9 to NA.

```
smart_ohio_raw <- smart_ohio_raw %>%
  mutate(marital = fct_recode(factor(MARITAL),
    "Married" = "1",
    "Divorced" = "2",
    "Widowed" = "3",
    "Separated" = "4",
    "Never_Married" = "5",
    "Unmarried_Couple" = "6",
    NULL = "9"))

smart_ohio_raw %>% count(MARITAL, marital)

# A tibble: 7 x 3
  MARITAL   marital      n
  <dbl> <fct> <int>
1       1 Married    3668
```

2	2 Divorced	1110
3	3 Widowed	978
4	4 Separated	142
5	5 Never_Married	1248
6	6 Unmarried_Couple	208
7	9 <NA>	58

2.5.10.7 EDUCA recoded to educgroup

The available levels of EDUCA (Education Level) are responses to: “What is the highest grade or year of school you completed?”

- 1 = Never attended school or only kindergarten
- 2 = Grades 1 through 8 (Elementary)
- 3 = Grades 9 through 11 (Some high school)
- 4 = Grade 12 or GED (High school graduate)
- 5 = College 1 year to 3 years (Some college or technical school)
- 6 = College 4 years or more (College graduate)
- 9 = Refused
- BLANK = Not asked or missing

We'll just turn this into a factor, and move 9 to NA.

```
smart_ohio_raw <- smart_ohio_raw %>%
  mutate(educgroup = fct_recode(factor(EDUCA),
    "Kindergarten" = "1",
    "Elementary" = "2",
    "Some_HS" = "3",
    "HS_Grad" = "4",
    "Some_College" = "5",
    "College_Grad" = "6",
    NULL = "9"))

smart_ohio_raw %>% count(EDUCA, educgroup)
```

```
# A tibble: 7 x 3
  EDUCA educgroup      n
  <dbl> <fct>     <int>
1     1 Kindergarten     3
2     2 Elementary      117
3     3 Some_HS        332
4     4 HS_Grad       2209
5     5 Some_College   2079
6     6 College_Grad   2646
7     9 <NA>            26
```

2.5.10.8 RENTHOM1 recoded to home_own

The available levels of RENTHOM1 (Own or Rent Home) are responses to: “Do you own or rent your home? (Home is defined as the place where you live most of the time/the majority of the year.)”

- 1 = Own
- 2 = Rent
- 3 = Other Arrangement
- 7 = Don’t know/Not Sure
- 9 = Refused
- BLANK = Not asked or missing

We’ll recode as `home_own` = 1 if they own their home, and 0 otherwise, and dealing with missingness properly.

```
smart_ohio_raw <- smart_ohio_raw %>%
  mutate(home_own = RENTHOM1,
        home_own = replace(home_own, home_own %in% c(7,9), NA),
        home_own = replace(home_own, home_own %in% c(2,3), 0))

smart_ohio_raw %>% count(RENTHOM1, home_own)

# A tibble: 5 x 3
RENTHOM1 home_own     n
<dbl>    <dbl> <int>
1       1      1  5216
2       2      0  1793
3       3      0   348
4       7     NA    28
5       9     NA    27
```

2.5.10.9 CPDEMO1A and its cleanup to cell_own

CPDEMO1A is the response to “Including phones for business and personal use, do you have a cell phone for personal use?”

Available responses are:

- 1 = Yes
- 2 = No
- 7 = Don’t know/Not sure
- 9 = Refused
- BLANK = Not asked or missing

and we’ll recode them to 1 = Yes, 0 = No, otherwise NA, as we’ve done previously.

```
smart_ohio_raw <- smart_ohio_raw %>%
  mutate(cell_own = 2 - CPDEMO1A,
        cell_own = replace(cell_own, cell_own < 0, NA))

smart_ohio_raw %>% count(CPDEMO1A, cell_own)

# A tibble: 5 x 3
  CPDEMO1A cell_own     n
  <dbl>      <dbl> <int>
1       1         1   2930
2       2         0    698
3       7        NA     2
4       9        NA    19
5      NA        NA   3763
```

2.5.10.10 VETERAN3 and its cleanup to veteran

VETERAN3, the Are You A Veteran variable, is the response to “Have you ever served on active duty in the United States Armed Forces, either in the regular military or in a National Guard or military reserve unit? (Active duty does not include training for the Reserves or National Guard, but DOES include activation, for example, for the Persian Gulf War.)”

- 1 = Yes
- 2 = No
- 7 = Don’t know/Not sure
- 9 = Refused
- BLANK = Not asked or missing

```
smart_ohio_raw <- smart_ohio_raw %>%
  mutate(veteran = VETERAN3,
        veteran = replace(veteran, veteran %in% c(7, 9), NA),
        veteran = replace(veteran, veteran == 2, 0))

smart_ohio_raw %>% count(VETERAN3, veteran)

# A tibble: 3 x 3
  VETERAN3 veteran     n
  <dbl>      <dbl> <int>
1       1         1   927
2       2         0  6479
3       9        NA     6
```

2.5.10.11 EMPLOY1 and its cleanup to employment

EMPLOY1, the Employment Status variable, is the response to “Are you currently ... ?”

- 1 = Employed for wages
- 2 = Self-employed
- 3 = Out of work for 1 year or more
- 4 = Out of work for less than 1 year
- 5 = A homemaker
- 6 = A student
- 7 = Retired
- 8 = Unable to work
- 9 = Refused
- BLANK = Not asked or missing

We'll just turn this into a factor, and move 9 to NA.

```
smart_ohio_raw <- smart_ohio_raw %>%
  mutate(employment = fct_recode(factor(EMPLOY1),
    "Employed_for_wages" = "1",
    "Self-employed" = "2",
    "Outofwork_1yearormore" = "3",
    "Outofwork_lt1year" = "4",
    "Homemaker" = "5",
    "Student" = "6",
    "Retired" = "7",
    "Unable_to_work" = "8",
    NULL = "9"))

smart_ohio_raw %>% count(EMPLOY1, employment)

# A tibble: 9 x 3
  EMPLOY1   employment       n
  <dbl> <fct>        <int>
1     1 Employed_for_wages  3119
2     2 Self-employed      466
3     3 Outofwork_1yearormore 254
4     4 Outofwork_lt1year   134
5     5 Homemaker          411
6     6 Student             190
7     7 Retired             2202
8     8 Unable_to_work     603
9     9 <NA>                 33
```

2.5.10.12 CHILDREN and its cleanup to kids

CHILDREN, the Number of Children in Household variable, is the response to “How many children less than 18 years of age live in your household?”

- 1-87 = legitimate responses
- 88 = None
- 99 = Refused
- BLANK = Not asked or missing

```
smart_ohio_raw <- smart_ohio_raw %>%
  mutate(kids = CHILDREN,
        kids = replace(kids, kids == 99, NA),
        kids = replace(kids, kids == 88, 0))

smart_ohio_raw %>% count(CHILDREN, kids) %>% tail()

# A tibble: 6 x 3
  CHILDREN   kids     n
  <dbl> <dbl> <int>
1       6      6     7
2       7      7     5
3       8      8     2
4      12     12     1
5      88      0  5449
6      99     NA    43
```

2.5.10.13 INCOME2 to incomegroup

The available levels of INCOME2 (Income Level) are responses to: “Is your annual household income from all sources ...”

- 1 = Less than \$10,000
- 2 = \$10,000 to less than \$15,000
- 3 = \$15,000 to less than \$20,000
- 4 = \$20,000 to less than \$25,000
- 5 = \$25,000 to less than \$35,000
- 6 = \$35,000 to less than \$50,000
- 7 = \$50,000 to less than \$75,000
- 8 = \$75,000 or more
- 77 = Don’t know/Not sure
- 99 = Refused
- BLANK = Not asked or missing

We’ll just turn this into a factor, and move 77 and 99 to NA.

```
smart_ohio_raw <- smart_ohio_raw %>%
  mutate(incomegroup = fct_recode(factor(`INCOME2`),
    "0-9K" = "1",
    "10-14K" = "2",
    "15-19K" = "3",
    "20-24K" = "4",
    "25-34K" = "5",
    "35-49K" = "6",
    "50-74K" = "7",
    "75K+" = "8",
    NULL = "77",
    NULL = "99"))
```

```
smart_ohio_raw %>% count(`INCOME2`, incomegroup)
```

```
# A tibble: 11 x 3
  INCOME2 incomegroup     n
  <dbl> <fct>      <int>
1     1 0-9K        285
2     2 10-14K       306
3     3 15-19K       477
4     4 20-24K       589
5     5 25-34K       685
6     6 35-49K       922
7     7 50-74K       928
8     8 75K+        1910
9     77 <NA>         610
10    99 <NA>         678
11    NA <NA>          22
```

2.5.10.14 INTERNET and its cleanup to internet30

INTERNET, the Internet use in the past 30 days variable, is the response to “Have you used the internet in the past 30 days?”

- 1 = Yes
- 2 = No
- 7 = Don’t know/Not sure
- 9 = Refused
- BLANK = Not asked or missing

```
smart_ohio_raw <- smart_ohio_raw %>%
  mutate(internet30 = INTERNET,
    internet30 = replace(internet30, internet30 %in% c(7, 9), NA),
    internet30 = replace(internet30, internet30 == 2, 0))
```

```
smart_ohio_raw %>% count(INTERNET, internet30)

# A tibble: 5 x 3
  INTERNET internet30     n
  <dbl>        <dbl> <int>
1      1            1   6020
2      2            0   1335
3      7            NA    10
4      9            NA    10
5     NA            NA    37
```

2.5.10.15 WTKG3 is weight_kg

WTKG3 is computed by CDC, as the respondent's weight in kilograms with two implied decimal places. We calculate the actual weight in kg, with the following:

```
smart_ohio_raw <- smart_ohio_raw %>%
  mutate(weight_kg = WTKG3/100)

smart_ohio_raw %>% count(WTKG3, weight_kg) %>% tail()

# A tibble: 6 x 3
  WTKG3 weight_kg     n
  <dbl>     <dbl> <int>
1 19051     191.     1
2 19278     193.     1
3 19504     195.     1
4 20412     204.     2
5 20865     209.     1
6     NA       NA    462
```

2.5.10.16 HEIGHT3 is replaced with height_m

HEIGHT3 is strangely gathered to allow people to specify their height in either feet and inches or in meters and centimeters.

- 200-711 indicates height in feet (first digit) and inches (second two digits)
- 9000 - 9998 indicates height in meters (second digit) and centimeters (last two digits)
- 7777 = Don't know/Not sure
- 9999 = Refused

Note that there is one impossible value of 575 in the data set. We'll make that an NA, and we'll also make NA any heights below 3 feet, or above 2.24 meters. Specifically, we calculate the actual height in meters, with the following:

```
smart_ohio_raw <- smart_ohio_raw %>%
  mutate(height_m = case_when(
    HEIGHT3 >= 300 & HEIGHT3 <= 511 ~ round((12*floor(HEIGHT3/100) + (HEIGHT3 - 100)/100)),
    HEIGHT3 >= 600 & HEIGHT3 <= 711 ~ round((12*floor(HEIGHT3/100) + (HEIGHT3 - 100)/100)),
    HEIGHT3 >= 9000 & HEIGHT3 <= 9224 ~ ((HEIGHT3 - 9000)/100)))

smart_ohio_raw %>% count(HEIGHT3, height_m) %>% tail()

# A tibble: 6 x 3
  HEIGHT3 height_m     n
  <dbl>     <dbl> <int>
1     607     2.01     2
2     608     2.03     6
3     609     2.06     1
4    7777     NA     27
5    9999     NA     86
6       NA     NA     67
```

2.5.10.17 bmi is calculated from height_m and weight_kg

We'll calculate body-mass index from height and weight.

```
smart_ohio_raw <- smart_ohio_raw %>%
  mutate(bmi = round(weight_kg/(height_m)^2, 2))

smart_ohio_raw %>% count(height_m, weight_kg, bmi) # %>% tail()

# A tibble: 1,806 x 4
  height_m weight_kg     bmi     n
  <dbl>     <dbl> <dbl> <int>
1     1.35     39.0  21.4     1
2     1.35     52.2  28.6     1
3     1.4      89.8  45.8     1
4     1.42     31.8  15.8     1
5     1.42     45.4  22.5     1
6     1.42     55.8  27.7     1
7     1.42     58.5  29.0     1
8     1.42     59.9  29.7     1
9     1.42     60.8  30.1     1
10    1.42     71.2  35.3     1
# ... with 1,796 more rows
```

2.5.10.18 bmigroup is calculated from bmi

We'll then divide the respondents into adult BMI categories, in the usual way.

- BMI < 18.5 indicates underweight
- BMI from 18.5 up to 25 indicates normal weight
- BMI from 25 up to 30 indicates overweight
- BMI of 30 and higher indicates obesity

```
smart_ohio_raw <- smart_ohio_raw %>%
  mutate(bmigroup = factor(cut2(as.numeric(bmi),
                                cuts = c(18.5, 25.0, 30.0))))
```

```
smart_ohio_raw %>% count(bmigroup)
```

```
# A tibble: 5 x 2
  bmigroup      n
  <fct>     <int>
1 [13.3,18.5)   119
2 [18.5,25.0)  2017
3 [25.0,30.0)  2445
4 [30.0,75.5]  2338
5 <NA>         493
```

2.5.10.19 PREGNANT and its cleanup to pregnant

PREGNANT, the Pregnancy Status variable, is the response to “To your knowledge, are you now pregnant?”

- 1 = Yes
- 2 = No
- 7 = Don’t know/Not sure
- 9 = Refused
- BLANK = Not asked or missing (includes SEX = male)

```
smart_ohio_raw <- smart_ohio_raw %>%
  mutate(pregnant = PREGNANT,
        pregnant = replace(pregnant, pregnant %in% c(7, 9), NA),
        pregnant = replace(pregnant, pregnant == 2, 0))
```

```
smart_ohio_raw %>% count(PREGNANT, pregnant)
```

```
# A tibble: 5 x 3
  PREGNANT pregnant      n
  <dbl>     <dbl> <int>
1       1         1    41
2       2         0  1329
3       7         NA     3
4       9         NA     3
5      NA        NA  6036
```

2.5.10.20 DEAF and its cleanup to deaf

DEAF, the Are you deaf or do you have serious difficulty hearing variable, is the response to “Are you deaf or do you have serious difficulty hearing?”

- 1 = Yes
- 2 = No
- 7 = Don’t know/Not sure
- 9 = Refused
- BLANK = Not asked or missing

```
smart_ohio_raw <- smart_ohio_raw %>%
  mutate(deaf = DEAF,
         deaf = replace(deaf, deaf %in% c(7, 9), NA),
         deaf = replace(deaf, deaf == 2, 0))

smart_ohio_raw %>% count(DEAF, deaf)

# A tibble: 5 x 3
  DEAF   deaf     n
  <dbl> <dbl> <int>
1     1     1    708
2     2     0   6551
3     7    NA     15
4     9    NA      4
5    NA    NA    134
```

2.5.10.21 BLIND and its cleanup to blind

BLIND, the Blind or Difficulty seeing variable, is the response to “Are you blind or do you have serious difficulty seeing, even when wearing glasses?”

- 1 = Yes
- 2 = No
- 7 = Don’t know/Not sure
- 9 = Refused
- BLANK = Not asked or missing

```
smart_ohio_raw <- smart_ohio_raw %>%
  mutate(blind = BLIND,
        blind = replace(blind, blind %in% c(7, 9), NA),
        blind = replace(blind, blind == 2, 0))

smart_ohio_raw %>% count(BLIND, blind)

# A tibble: 5 x 3
  BLIND blind     n
  <dbl> <dbl> <int>
1     1     1    708
2     2     0   6551
3     7    NA     15
4     9    NA      4
5    NA    NA    134
```

```
<dbl> <dbl> <int>
1     1     1   415
2     2     0  6834
3     7    NA    14
4     9    NA     1
5    NA    NA   148
```

2.5.10.22 DECIDE and its cleanup to decide

DECIDE, the Difficulty Concentrating or Remembering variable, is the response to “Because of a physical, mental, or emotional condition, do you have serious difficulty concentrating, remembering, or making decisions?”

- 1 = Yes
- 2 = No
- 7 = Don’t know/Not sure
- 9 = Refused
- BLANK = Not asked or missing

```
smart_ohio_raw <- smart_ohio_raw %>%
  mutate(decide = DECIDE,
        decide = replace(decide, decide %in% c(7, 9), NA),
        decide = replace(decide, decide == 2, 0))

smart_ohio_raw %>% count(DECIDE, decide)

# A tibble: 5 x 3
  DECIDE decide      n
  <dbl>   <dbl> <int>
1     1     1   870
2     2     0  6348
3     7    NA    30
4     9    NA     2
5    NA    NA   162
```

2.5.10.23 DIFFWALK and its cleanup to diffwalk

DIFFWALK, the Difficulty Walking or Climbing Stairs variable, is the response to “Do you have serious difficulty walking or climbing stairs?”

- 1 = Yes
- 2 = No
- 7 = Don’t know/Not sure
- 9 = Refused
- BLANK = Not asked or missing

```
smart_ohio_raw <- smart_ohio_raw %>%
  mutate(diffwalk = DIFFWALK,
         diffwalk = replace(diffwalk, diffwalk %in% c(7, 9), NA),
         diffwalk = replace(diffwalk, diffwalk == 2, 0))

smart_ohio_raw %>% count(DIFFWALK, diffwalk)

# A tibble: 5 x 3
  DIFFWALK diffwalk     n
  <dbl>      <dbl> <int>
1       1        1  1482
2       2        0  5738
3       7       NA    19
4       9       NA     2
5      NA       NA   171
```

2.5.10.24 DIFFDRES and its cleanup to diffdress

DIFFDRES, the Difficulty Dressing or Bathing variable, is the response to “Do you have difficulty dressing or bathing?”

- 1 = Yes
- 2 = No
- 7 = Don’t know/Not sure
- 9 = Refused
- BLANK = Not asked or missing

```
smart_ohio_raw <- smart_ohio_raw %>%
  mutate(diffdress = DIFFDRES,
         diffdress = replace(diffdress, diffdress %in% c(7, 9), NA),
         diffdress = replace(diffdress, diffdress == 2, 0))

smart_ohio_raw %>% count(DIFFDRES, diffdress)

# A tibble: 5 x 3
  DIFFDRES diffdress     n
  <dbl>      <dbl> <int>
1       1        1   352
2       2        0  6868
3       7       NA    12
4       9       NA     1
5      NA       NA   179
```

2.5.10.25 DIFFALON and its cleanup to diffalone

DIFFALON, the Difficulty Doing Errands Alone variable, is the response to “Because of a physical, mental, or emotional condition, do you have difficulty doing errands alone such as visiting a doctor’s office or shopping?”

- 1 = Yes
- 2 = No
- 7 = Don’t know/Not sure
- 9 = Refused
- BLANK = Not asked or missing

```
smart_ohio_raw <- smart_ohio_raw %>%
  mutate(diffalone = DIFFALON,
        diffalone = replace(diffalone, diffalone %in% c(7, 9), NA),
        diffalone = replace(diffalone, diffalone == 2, 0))

smart_ohio_raw %>% count(DIFFALON, diffalone)

# A tibble: 5 x 3
  DIFFALON diffalone     n
  <dbl>      <dbl> <int>
1       1          1    636
2       2          0   6560
3       7         NA     15
4       9         NA      4
5      NA         NA   197
```

2.5.11 Tobacco Use (2 items)

2.5.11.1 SMOKE100 and its cleanup to smoke100

SMOKE100, the Smoked at Least 100 Cigarettes variable, is the response to “Have you smoked at least 100 cigarettes in your entire life? [Note: 5 packs = 100 cigarettes]”

- 1 = Yes
- 2 = No
- 7 = Don’t know/Not sure
- 9 = Refused

```
smart_ohio_raw <- smart_ohio_raw %>%
  mutate(smoke100 = SMOKE100,
        smoke100 = replace(smoke100, smoke100 %in% c(7, 9), NA),
        smoke100 = replace(smoke100, smoke100 == 2, 0))

smart_ohio_raw %>% count(SMOKE100, smoke100)
```

```
# A tibble: 5 x 3
  SMOKE100 smoke100     n
  <dbl>    <dbl> <int>
1      1        1  3294
2      2        0 3881
3      7       NA   31
4      9       NA    4
5     NA       NA  202
```

2.5.11.2 `_SMOKER3` and its cleanup to `smoker`

`_SMOKER3`, is a calculated variable which categorizes subjects by their smoking status:

- 1 = Current smoker who smokes daily
- 2 = Current smoker but not every day
- 3 = Former smoker
- 4 = Never smoked
- 9 = Don't Know / Refused / Missing

We'll reclassify this as a factor with appropriate labels and NAs.

```
smart_ohio_raw <- smart_ohio_raw %>%
  mutate(smoker = fct_recode(factor(`_SMOKER3`),
    "Current_daily" = "1",
    "Current_not_daily" = "2",
    "Former" = "3",
    "Never" = "4",
    NULL = "9"))

smart_ohio_raw %>% count(`_SMOKER3`, smoker)
```

```
# A tibble: 5 x 3
  `_SMOKER3` smoker             n
  <dbl> <fct>           <int>
1      1 Current_daily     990
2      2 Current_not_daily 300
3      3 Former            1999
4      4 Never             3881
5      9 <NA>              242
```

2.5.12 E-Cigarettes (2 items)

2.5.12.1 ECIGARET and its cleanup to ecig_ever

ECIGARET, the Ever used an e-cigarette variable, is the response to “Have you ever used an e-cigarette or other electronic vaping product, even just one time, in your entire life?”

- 1 = Yes
- 2 = No
- 7 = Don’t know/Not sure
- 9 = Refused

```
smart_ohio_raw <- smart_ohio_raw %>%
  mutate(ecig_ever = ECIGARET,
        ecig_ever = replace(ecig_ever, ecig_ever %in% c(7, 9), NA),
        ecig_ever = replace(ecig_ever, ecig_ever == 2, 0))

smart_ohio_raw %>% count(ECIGARET, ecig_ever)

# A tibble: 5 x 3
  ECIGARET ecig_ever     n
  <dbl>      <dbl> <int>
1       1         1   1354
2       2         0   5799
3       7         NA     9
4       9         NA     3
5      NA         NA   247
```

2.5.12.2 _ECIGSTS and its cleanup to ecigs

_ECIGSTS, is a calculated variable which categorizes subjects by their smoking status:

- 1 = Current and uses daily
- 2 = Current user but not every day
- 3 = Former user
- 4 = Never used e-cigarettes
- 9 = Don’t Know / Refused / Missing

We’ll reclassify this as a factor with appropriate labels and NAs.

```
smart_ohio_raw <- smart_ohio_raw %>%
  mutate(ecigs = fct_recode(factor(`_ECIGSTS`),
                            "Current_daily" = "1",
                            "Current_not_daily" = "2",
                            "Former" = "3",
```

```

    "Never" = "4",
    NULL = "9"))

smart_ohio_raw %>% count(`_ECIGSTS`, ecigs)

# A tibble: 5 x 3
`_ECIGSTS` ecigs      n
<dbl> <fct>     <int>
1          1 Current_daily    102
2          2 Current_not_daily 165
3          3 Former            1085
4          4 Never             5799
5          9 <NA>              261

```

2.5.13 Alcohol Consumption (6 items)

2.5.13.1 ALCDAY5 and its cleanup to alcdays

ALCDAY5, the Days in past 30 had alcoholic beverage variable, is the response to “During the past 30 days, how many days per week or per month did you have at least one drink of any alcoholic beverage such as beer, wine, a malt beverage or liquor?”

- 101-107 = # of days per week (101 = 1 day per week, 107 = 7 days per week)
- 201-230 = # of days in past 30 days (201 = 1 day in last 30, 230 = 30 days in last 30)
- 777 = Don’t know/Not sure
- 888 = No drinks in past 30 days
- 999 = Refused
- BLANK = Not asked or Missing

We’re going to convert this to a single numeric value. Answers in days per week (in the past 7 days) will be converted (after rounding) to days in the past 30. This is a little bit of a mess, really, but we can do it.

```

smart_ohio_raw <- smart_ohio_raw %>%
  mutate(alcdays = as.numeric(ALCDAY5)) %>%
  mutate(alcdays = replace(alcdays, alcdays == 888, 0),
        alcdays = replace(alcdays, alcdays %in% c(777, 999), NA)) %>%
  mutate(alcdays = case_when(ALCDAY5 > 199 & ALCDAY5 < 231 ~ ALCDAY5 - 200,
                             ALCDAY5 > 100 & ALCDAY5 < 108 ~ round((ALCDAY5 - 100)*30)
                             TRUE ~ alcdays))

smart_ohio_raw %>% count(ALCDAY5, alcdays)

```

```
# A tibble: 39 x 3
  ALCDAY5 alcdays     n
  <dbl>    <dbl> <int>
1    101      4    263
2    102      9    197
3    103     13   142
4    104     17    76
5    105     21    53
6    106     26    18
7    107     30   114
8    201      1   621
9    202      2   448
10   203      3   233
# ... with 29 more rows
```

2.5.13.2 AVEDRNK2 and its cleanup to avgdrinks

AVEDRNK2, the Avg alcoholic drinks per day in past 30 variable, is the response to “One drink is equivalent to a 12-ounce beer, a 5-ounce glass of wine, or a drink with one shot of liquor. During the past 30 days, on the days when you drank, about how many drinks did you drink on the average? (A 40 ounce beer would count as 3 drinks, or a cocktail drink with 2 shots would count as 2 drinks.)”

- 1-76 = # of drinks per day
- 77 = Don’t know/Not sure
- 99 = Refused
- BLANK = Not asked or Missing (always happens when ALCDAY5 = 777, 888 or 999)

```
smart_ohio_raw <- smart_ohio_raw %>%
  mutate(avgdrinks = AVEDRNK2,
        avgdrinks = replace(avgdrinks, avgdrinks > 76, NA))
```

```
smart_ohio_raw %>% count(AVEDRNK2, avgdrinks) %>% tail()
```

```
# A tibble: 6 x 3
  AVEDRNK2 avgdrinks     n
  <dbl>    <dbl> <int>
1      42      42     1
2      60      60     2
3      76      76     1
4      77      NA    46
5      99      NA     5
6      NA      NA  3876
```

2.5.13.3 MAXDRNKS and its cleanup to maxdrinks

MAXDRINKS, the most drinks on a single occasion in the past 30 days variable, is the response to “During the past 30 days, what is the largest number of drinks you had on any occasion?”

- 1-76 = # of drinks
- 77 = Don’t know/Not sure
- 99 = Refused
- BLANK = Not asked or Missing (always happens when ALCDAY5 = 777, 888 or 999)

```
smart_ohio_raw <- smart_ohio_raw %>%
  mutate(maxdrinks = MAXDRNKS,
        maxdrinks = replace(maxdrinks, maxdrinks > 76, NA))
```

```
smart_ohio_raw %>% count(MAXDRNKS, maxdrinks) %>% tail()
```

	MAXDRNKS	maxdrinks	n
	<dbl>	<dbl>	<int>
1	42	42	1
2	48	48	1
3	76	76	2
4	77	NA	94
5	99	NA	11
6	NA	NA	3899

2.5.13.4 _RFBING5 and its cleanup to binge

_RFBING5 identifies binge drinkers (males having five or more drinks on one occasion, females having four or more drinks on one occasion in the past 30 days)

The values are

- 1 = No
- 2 = Yes
- 9 = Don’t Know / Refused / Missing

People who reported no alcdays are reported here as “No,” so we’ll adjust this into an indicator variable, and create the necessary NAs.

```
smart_ohio_raw <- smart_ohio_raw %>%
  mutate(binge = `_RFBING5` - 1,
        binge = replace(binge, binge > 1, NA))

smart_ohio_raw %>% count(`_RFBING5`, binge)
```

```
# A tibble: 3 x 3
`_RFBING5` binge     n
<dbl> <dbl> <int>
1      1     0  6035
2      2     1 1000
3      9    NA  377
```

2.5.13.5 _DRNKWEK and its cleanup to drinks_wk

_DRNKWEK provides the computed number of alcoholic drinks per week, with two implied decimal places. The code 99900 is used for “Don’t know / Not sure / Refused / Missing” so we’ll fix that, and also divide by 100 to get an average with a decimal point.

Note: We’re also going to treat all results of 100 or more drinks per week as incorrect, and thus indicate them as missing data here.

```
smart_ohio_raw <- smart_ohio_raw %>%
  mutate(drinks_wk = `_DRNKWEK` / 100,
        drinks_wk = replace(drinks_wk, drinks_wk > 99, NA))

smart_ohio_raw %>% count(`_DRNKWEK`, drinks_wk) %>% tail(12)

# A tibble: 12 x 3
`_DRNKWEK` drinks_wk     n
<dbl>       <dbl> <int>
1      9333     93.3     2
2     10000     NA       1
3     10500     NA       2
4     11667     NA       1
5     14000     NA       2
6     16800     NA       2
7     17500     NA       1
8     18200     NA       1
9     28000     NA       1
10    29400     NA       1
11    53200     NA       1
12    99900     NA     379
```

2.5.13.6 _RFDRHV5 and its cleanup to drink_heavy

_RFDRHV5 identifies heavy drinkers (males having 14 or more drinks per week, females having 7 or more drinks per week)

The values are

- 1 = No
- 2 = Yes
- 9 = Don't Know / Refused / Missing

People who reported no `alcdays` are reported here as “No,” so we’ll adjust this into an indicator variable, and create the necessary NAs.

```
smart_ohio_raw <- smart_ohio_raw %>%
  mutate(drink_heavy = `_RFDRHV5` - 1,
        drink_heavy = replace(drink_heavy, drink_heavy > 1, NA))

smart_ohio_raw %>% count(`_RFDRHV5`, drink_heavy)

# A tibble: 3 x 3
  `_RFDRHV5` drink_heavy     n
  <dbl>       <dbl> <int>
1 1             0    6607
2 2             1    426
3 9            NA    379
```

2.5.14 Fruits and Vegetables (8 items)

2.5.14.1 _FRUTSU1 and its cleanup to fruit_day

`_FRUTSU1` provides the computed number of fruit servings consumed per day, with two implied decimal places. We’ll divide by 100 to insert the decimal point.

Note: We’re also going to treat all results exceeding 16 servings per day as implausible, and thus indicate them as missing data here, following some CDC procedures.

```
smart_ohio_raw <- smart_ohio_raw %>%
  mutate(fruit_day = `_FRUTSU1` / 100,
        fruit_day = replace(fruit_day, fruit_day > 16, NA))

smart_ohio_raw %>% count(`_FRUTSU1`, fruit_day) %>% tail()

# A tibble: 6 x 3
  `_FRUTSU1` fruit_day     n
  <dbl>       <dbl> <int>
1 913        9.13      1
2 1000       10         4
3 1400       14         1
4 3000       NA         1
5 7600       NA         1
6 NA          NA        555
```

2.5.14.2 _VEGESU1 and its cleanup to veg_day

_VEGESU1 provides the computed number of vegetable servings consumed per day, with two implied decimal places. We'll divide by 100 to insert the decimal point.

Note: We're also going to treat all results exceeding 23 servings per day as implausible, and thus indicate them as missing data here, following some CDC procedures.

```
smart_ohio_raw <- smart_ohio_raw %>%
  mutate(veg_day = `_VEGESU1` / 100,
        veg_day = replace(veg_day, veg_day > 23, NA))

smart_ohio_raw %>% count(`_VEGESU1`, veg_day) %>% tail()

# A tibble: 6 x 3
  `_VEGESU1` veg_day     n
  <dbl>     <dbl> <int>
1     1414    14.1      1
2     1603    16.0      1
3     1891    18.9      1
4     2167    21.7      1
5     3150     NA       1
6       NA     NA      666
```

2.5.14.3 FTJUDA2_ and its cleanup to eat_juice

FTJUDA2_ provides the servings of fruit juice consumed per day, with two implied decimal places. We'll divide by 100 to insert the decimal point.

Note: We're also going to treat all results exceeding 16 servings per day as implausible, and thus indicate them as missing data here.

```
smart_ohio_raw <- smart_ohio_raw %>%
  mutate(eat_juice = `FTJUDA2_` / 100,
        eat_juice = replace(eat_juice, eat_juice > 16, NA))

smart_ohio_raw %>% count(`FTJUDA2_`, eat_juice) %>% tail()

# A tibble: 6 x 3
  FTJUDA2_ eat_juice     n
  <dbl>     <dbl> <int>
1     500      5       6
2     600      6       1
3     700      7       1
4    1200     12       1
```

5	7500	NA	1
6	NA	NA	469

2.5.14.4 FRUTDA2_ and its cleanup to eat_fruit

FRUTDA2_ provides the servings of fruit consumed per day, with two implied decimal places. We'll divide by 100 to insert the decimal point.

Note: We're also going to treat all results exceeding 16 servings per day as implausible, and thus indicate them as missing data here.

```
smart_ohio_raw <- smart_ohio_raw %>%
  mutate(eat_fruit = `FRUTDA2_` / 100,
        eat_fruit = replace(eat_fruit, eat_fruit > 16, NA))

smart_ohio_raw %>% count(`FRUTDA2_`, eat_fruit) %>% tail()

# A tibble: 6 x 3
  FRUTDA2_ eat_fruit     n
  <dbl>      <dbl> <int>
1    700       7     5
2    800       8     3
3    900       9     1
4   1000      10     1
5   3000      NA     1
6     NA      NA    456
```

2.5.14.5 GRENDAA1_ and its cleanup to eat_greenveg

GRENDA1_ provides the servings of dark green vegetables consumed per day, with two implied decimal places. We'll divide by 100 to insert the decimal point.

Note: We're also going to treat all results exceeding 16 servings per day as implausible, and thus indicate them as missing data here.

```
smart_ohio_raw <- smart_ohio_raw %>%
  mutate(eat_greenveg = `GRENDA1_` / 100,
        eat_greenveg = replace(eat_greenveg, eat_greenveg > 16, NA))

smart_ohio_raw %>% count(`GRENDA1_`, eat_greenveg) %>% tail()

# A tibble: 6 x 3
  GRENDA1_ eat_greenveg     n
  <dbl>      <dbl> <int>
1    700       7     4
2    786      7.86    1
3    800       8     2
```

```

4      2000      NA      1
5      3000      NA      1
6       NA      NA    447

```

2.5.14.6 FRNCHDA_ and its cleanup to eat_fries

FRNCHDA_ provides the servings of french fries consumed per day, with two implied decimal places. We'll divide by 100 to insert the decimal point.

Note: We're also going to treat all results exceeding 16 servings per day as implausible, and thus indicate them as missing data here.

```

smart_ohio_raw <- smart_ohio_raw %>%
  mutate(eat_fries = `FRNCHDA_` / 100,
        eat_fries = replace(eat_fries, eat_fries > 16, NA))

smart_ohio_raw %>% count(`FRNCHDA_`, eat_fries) %>% tail()

# A tibble: 6 x 3
  FRNCHDA_ eat_fries     n
  <dbl>      <dbl> <int>
1     300      3     9
2     314     3.14     1
3     400      4     3
4     500      5     1
5     700      7     1
6      NA      NA    453

```

2.5.14.7 POTADA1_ and its cleanup to eat_potato

POTADA1_ provides the servings of potatoes consumed per day, with two implied decimal places. We'll divide by 100 to insert the decimal point.

Note: We're also going to treat all results exceeding 16 servings per day as implausible, and thus indicate them as missing data here.

```

smart_ohio_raw <- smart_ohio_raw %>%
  mutate(eat_potato = `POTADA1_` / 100,
        eat_potato = replace(eat_potato, eat_potato > 16, NA))

smart_ohio_raw %>% count(`POTADA1_`, eat_potato) %>% tail()

# A tibble: 6 x 3
  POTADA1_ eat_potato     n
  <dbl>      <dbl> <int>
1     314     3.14     1
2     329     3.29     1

```

3	400	4	3
4	471	4.71	1
5	700	7	1
6	NA	NA	501

2.5.14.8 VEGEDA2_ and its cleanup to eat_otherveg

VEGEDA2_ provides the servings of other vegetables consumed per day, with two implied decimal places. We'll divide by 100 to insert the decimal point.

Note: We're also going to treat all results exceeding 16 servings per day as implausible, and thus indicate them as missing data here.

```
smart_ohio_raw <- smart_ohio_raw %>%
  mutate(eat_otherveg = `VEGEDA2_` / 100,
        eat_otherveg = replace(eat_otherveg, eat_otherveg > 16, NA))

smart_ohio_raw %>% count(`VEGEDA2_`, eat_otherveg) %>% tail()

# A tibble: 6 x 3
  VEGEDA2_ eat_otherveg     n
  <dbl>      <dbl> <int>
1     600       6     3
2     700       7    11
3     800       8     1
4    1000      10     2
5    1100      11     1
6       NA      NA   509
```

2.5.15 Exercise and Physical Activity (8 items)

2.5.15.1 _TOTINDA and its cleanup to exerany

_TOTINDA, the Exercise in Past 30 Days variable, is the response to “During the past month, other than your regular job, did you participate in any physical activities or exercises such as running, calisthenics, golf, gardening, or walking for exercise?”

- 1 = Yes
- 2 = No
- 7 = Don't know/Not sure
- 9 = Refused
- BLANK = Not asked or missing

This is just like HLTHPLAN.

```
smart_ohio_raw <- smart_ohio_raw %>%
  mutate(exerany = `TOTINDA`,
        exerany = replace(exerany, exerany %in% c(7, 9), NA),
        exerany = replace(exerany, exerany == 2, 0))

smart_ohio_raw %>% count(`TOTINDA`, exerany)

# A tibble: 3 x 3
`TOTINDA` exerany     n
<dbl>    <dbl> <int>
1         1      1  4828
2         2      0  2137
3         9     NA   447
```

2.5.15.2 _PACAT1 and its cleanup to activity

_PACAT1 contains physical activity categories, estimated from responses to the BRFSS. The categories are:

- 1 = Highly Active
- 2 = Active
- 3 = Insufficiently Active
- 4 = Inactive
- 9 = Don't Know / Not Sure / Refused / Missing

So we'll create a factor.

```
smart_ohio_raw <- smart_ohio_raw %>%
  mutate(activity = factor(`_PACAT1`),
         activity = fct_recode(activity,
                               "Highly_Active" = "1",
                               "Active" = "2",
                               "Insufficiently_Active" = "3",
                               "Inactive" = "4",
                               NULL = "9"))

smart_ohio_raw %>% count(`_PACAT1`, activity)

# A tibble: 5 x 3
`_PACAT1` activity           n
<dbl>    <fct>          <int>
1         1 Highly_Active  2053
2         2 Active          1132
3         3 Insufficiently_Active 1293
4         4 Inactive        2211
5         9 <NA>            723
```

2.5.15.3 _PAINDX1 and its cleanup to rec_aerobic

_PAINDX1 indicates whether the respondent's stated levels of physical activity meet recommendations for aerobic activity. The responses are:

- 1 = Yes
- 2 = No
- 9 = Don't know/Not sure/Refused/Missing

```
smart_ohio_raw <- smart_ohio_raw %>%
  mutate(rec_aerobic = 2 - `_PAINDX1`,
         rec_aerobic = replace(rec_aerobic, rec_aerobic < 0, NA))

smart_ohio_raw %>% count(`_PAINDX1`, rec_aerobic)

# A tibble: 3 x 3
`_PAINDX1` rec_aerobic     n
<dbl>        <dbl> <int>
1            1            1  3228
2            2            0  3504
3            9           NA   680
```

2.5.15.4 _PASTRNG and its cleanup to rec_strength

_PASTRNG indicates whether the respondent's stated levels of physical activity meet recommendations for strength-building activity. The responses are:

- 1 = Yes
- 2 = No
- 9 = Don't know/Not sure/Refused/Missing

```
smart_ohio_raw <- smart_ohio_raw %>%
  mutate(rec_strength = 2 - `_PASTRNG`,
         rec_strength = replace(rec_strength, rec_strength < 0, NA))

smart_ohio_raw %>% count(`_PASTRNG`, rec_strength)

# A tibble: 3 x 3
`_PASTRNG` rec_strength     n
<dbl>        <dbl> <int>
1            1            1  1852
2            2            0  5004
3            9           NA   556
```

2.5.15.5 EXTRACT11 and its cleanup to exer1_type

Respondents are asked “What type of physical activity or exercise did you spend the most time doing during the past month?” and these responses are gathered into a set of 76 named categories, including an “other” category. Codes 77 (Don’t Know / Not Sure) and 99 (Refused) are dropped into NA in my code below, and Code 98 (“Other type of activity”) remains. Then I went through the tedious work of converting the factor levels from numbers to names, following the value labels provided by BRFSS.

```
smart_ohio_raw <- smart_ohio_raw %>%
  mutate(exer1_type = factor(EXTRACT11),
        exer1_type = fct_recode(
          exer1_type,
          "Active Gaming Devices" = "1",
          "Aerobics video or class" = "2",
          "Backpacking" = "3",
          "Badminton" = "4",
          "Basketball" = "5",
          "Bicycling machine" = "6",
          "Bicycling" = "7",
          "Boating" = "8",
          "Bowling" = "9",
          "Boxing" = "10",
          "Calisthenics" = "11",
          "Canoeing" = "12",
          "Carpentry" = "13",
          "Dancing" = "14",
          "Elliptical machine" = "15",
          "Fishing" = "16",
          "Frisbee" = "17",
          "Gardening" = "18",
          "Golf with cart" = "19",
          "Golf without cart" = "20",
          "Handball" = "21",
          "Hiking" = "22",
          "Hockey" = "23",
          "Horseback riding" = "24",
          "Hunting large game" = "25",
          "Hunting small game" = "26",
          "Inline skating" = "27",
          "Jogging" = "28",
          "Lacrosse" = "29",
          "Mountain climbing" = "30",
          "Mowing lawn" = "31",
          "Paddleball" = "32",
```

```
"Painting house" = "33",
"Pilates" = "34",
"Racquetball" = "35",
"Raking lawn" = "36",
"Running" = "37",
"Rock climbing" = "38",
"Rope skipping" = "39",
"Rowing machine" = "40",
"Rugby" = "41",
"Scuba diving" = "42",
"Skateboarding" = "43",
"Skating" = "44",
"Sledding" = "45",
"Snorkeling" = "46",
"Snow blowing" = "47",
"Snow shoveling" = "48",
"Snow skiing" = "49",
"Snowshoeing" = "50",
"Soccer" = "51",
"Softball/Baseball" = "52",
"Squash" = "53",
"Stair Climbing" = "54",
"Stream fishing" = "55",
"Surfing" = "56",
"Swimming" = "57",
"Swimming in laps" = "58",
"Table tennis" = "59",
"Tai Chi" = "60",
"Tennis" = "61",
"Touch football" = "62",
"Volleyball" = "63",
"Walking" = "64",
"Waterskiing" = "66",
"Weight lifting" = "67",
"Wrestling" = "68",
"Yoga" = "69",
"Child Care" = "71",
"Farm Work" = "72",
"Household Activities" = "73",
"Martial Arts" = "74",
"Upper Body Cycle" = "75",
"Yard Work" = "76",
"Other Activities" = "98",
NULL = "77",
NULL = "99")
```

```
)
```

```
Warning: Unknown levels in `f`: 3, 17, 21, 32, 36, 41, 42, 45, 47, 53, 55, 56,
59
```

The warning generated here is caused by the fact that some of the available types of exercise were not mentioned by people in our sample. Looking at the last few results, we can see how many people fell into several categories.

```
smart_ohio_raw %>% count(EXTRACT11, exer1_type) %>% tail()
```

```
# A tibble: 6 x 3
  EXTRACT11 exer1_type      n
  <dbl> <fct>        <int>
1      75 Upper Body Cycle    6
2      76 Yard Work         78
3      77 <NA>              10
4      98 Other Activities   276
5      99 <NA>              4
6      NA <NA>            2588
```

The most common activities are:

```
smart_ohio_raw %>% count(exer1_type, sort = TRUE) %>% head(10)
```

```
# A tibble: 10 x 2
  exer1_type      n
  <fct>        <int>
1 Walking       2605
2 <NA>          2602
3 Running        324
4 Other Activities  276
5 Gardening       242
6 Weight lifting   189
7 Aerobics video or class 103
8 Bicycling machine  103
9 Bicycling        96
10 Golf with cart    90
```

2.5.15.6 EXTRACT21 and its cleanup to exer2_type

As a follow-up, respondents are asked “What other type of physical activity gave you the next most exercise during the past month?” and these responses are also gathered into the same set of 76 named categories, including an “other” category, but now also adding a “No Other Activity” category (code 88). Codes 77 (Don’t Know / Not Sure) and 99 (Refused) are dropped into NA in my code below, and Code 98 (“Other type of activity”) remains. Then I went through the tedious

work of converting the factor levels from numbers to names, following the value labels provided by BRFSS. I'm sure there's a better way to do this.

```
smart_ohio_raw <- smart_ohio_raw %>%
  mutate(exer2_type = factor(EXTRACT21),
        exer2_type = fct_recode(
          exer2_type,
          "Active Gaming Devices" = "1",
          "Aerobics video or class" = "2",
          "Backpacking" = "3",
          "Badminton" = "4",
          "Basketball" = "5",
          "Bicycling machine" = "6",
          "Bicycling" = "7",
          "Boating" = "8",
          "Bowling" = "9",
          "Boxing" = "10",
          "Calisthenics" = "11",
          "Canoeing" = "12",
          "Carpentry" = "13",
          "Dancing" = "14",
          "Elliptical machine" = "15",
          "Fishing" = "16",
          "Frisbee" = "17",
          "Gardening" = "18",
          "Golf with cart" = "19",
          "Golf without cart" = "20",
          "Handball" = "21",
          "Hiking" = "22",
          "Hockey" = "23",
          "Horseback riding" = "24",
          "Hunting large game" = "25",
          "Hunting small game" = "26",
          "Inline skating" = "27",
          "Jogging" = "28",
          "Lacrosse" = "29",
          "Mountain climbing" = "30",
          "Mowing lawn" = "31",
          "Paddleball" = "32",
          "Painting house" = "33",
          "Pilates" = "34",
          "Racquetball" = "35",
          "Raking lawn" = "36",
          "Running" = "37",
          "Rock climbing" = "38",
          "Rope skipping" = "39",
```

```

"Rowing machine" = "40",
"Rugby" = "41",
"Scuba diving" = "42",
"Skateboarding" = "43",
"Skating" = "44",
"Sledding" = "45",
"Snorkeling" = "46",
"Snow blowing" = "47",
"Snow shoveling" = "48",
"Snow skiing" = "49",
"Snowshoeing" = "50",
"Soccer" = "51",
"Softball/Baseball" = "52",
"Squash" = "53",
"Stair Climbing" = "54",
"Stream fishing" = "55",
"Surfing" = "56",
"Swimming" = "57",
"Swimming in laps" = "58",
"Table tennis" = "59",
"Tai Chi" = "60",
"Tennis" = "61",
"Touch football" = "62",
"Volleyball" = "63",
"Walking" = "64",
"Waterskiing" = "66",
"Weight lifting" = "67",
"Wrestling" = "68",
"Yoga" = "69",
"Child Care" = "71",
"Farm Work" = "72",
"Household Activities" = "73",
"Martial Arts" = "74",
"Upper Body Cycle" = "75",
"Yard Work" = "76",
"No Other Activity" = "88",
"Other Activities" = "98",
NULL = "77",
NULL = "99")
)

```

```

Warning: Unknown levels in `f`: 3, 21, 30, 39, 41, 46, 50, 62
smart_ohio_raw %>% count(EXTRACT21, exer2_type) %>% tail()

```

```
# A tibble: 6 x 3
```

```
EXTRACT21 exer2_type          n
           <dbl> <fct>        <int>
1       76 Yard Work      153
2       77 <NA>          26
3       88 No Other Activity 1854
4       98 Other Activities 246
5       99 <NA>          19
6       NA <NA>         2627
```

The most common activity types in this group are:

```
smart_ohio_raw %>% count(exer2_type, sort = TRUE) %>% head(10)
```

```
# A tibble: 10 x 2
  exer2_type          n
  <fct>        <int>
1 <NA>          2672
2 No Other Activity 1854
3 Walking        629
4 Weight lifting 272
5 Other Activities 246
6 Gardening       202
7 Household Activities 169
8 Yard Work       153
9 Running         148
10 Bicycling      118
```

2.5.15.7 _MINAC11 and its cleanup to exer1_min

_MINAC11 is minutes of physical activity per week for the first activity (listed as `exer1_type` above.) Since there are only about 10,080 minutes in a typical week, we'll treat as implausible any values larger than 4200 minutes (which would indicate 70 hours per week.)

```
smart_ohio_raw <- smart_ohio_raw %>%
  mutate(exer1_min = `_MINAC11`,
        exer1_min = replace(exer1_min, exer1_min > 4200, NA))
```

```
smart_ohio_raw %>% count(`_MINAC11`, exer1_min) %>% tail()
```

```
# A tibble: 6 x 3
`_MINAC11` exer1_min     n
<dbl>      <dbl> <int>
1      3780      3780     8
2      3959      3959     1
3      3960      3960     1
4      4193      4193     6
```

```
5      27000      NA      1
6       NA      NA  2760
```

2.5.15.8 `_MINAC21` and its cleanup to `exer2_min`

`_MINAC21` is minutes of physical activity per week for the second activity (listed as `exer2_type` above.) Again, we'll treat as implausible any values larger than 4200 minutes (which would indicate 70 hours per week.)

```
smart_ohio_raw <- smart_ohio_raw %>%
  mutate(exer2_min = `_MINAC21`,
        exer2_min = replace(exer2_min, exer2_min > 4200, NA))

smart_ohio_raw %>% count(`_MINAC21`, exer2_min) %>% tail()

# A tibble: 6 x 3
`_MINAC21` exer2_min     n
<dbl>      <dbl> <int>
1      3360      3360     3
2      3780      3780     7
3      4193      4193     3
4      6120       NA     1
5      8400       NA     1
6       NA       NA  2770
```

2.5.16 Seatbelt Use (1 item)

2.5.16.1 SEATBELT and its cleanup to `seatbelt`

This question asks “How often do you use seat belts when you drive or ride in a car?” Possible responses are:

- 1 = Always
- 2 = Nearly always
- 3 = Sometimes
- 4 = Seldom
- 5 = Never
- 7 = Don’t know / Not sure
- 8 = Never drive or ride in a car
- 9 = Refused

We'll treat codes 7, 8 and 9 as NA, and turn this into a factor.

```
smart_ohio_raw <- smart_ohio_raw %>%
  mutate(seatbelt = fct_recode(factor(SEATBELT),
                               "Always" = "1",
```

```

    "Nearly_always" = "2",
    "Sometimes" = "3",
    "Seldom" = "4",
    "Never" = "5",
    NULL = "7",
    NULL = "8",
    NULL = "9"))

smart_ohio_raw %>% count(SEATBELT, seatbelt)

# A tibble: 9 x 3
  SEATBELT seatbelt      n
  <dbl>   <fct>     <int>
1       1 Always     6047
2       2 Nearly_always 409
3       3 Sometimes   191
4       4 Seldom      81
5       5 Never       148
6       7 <NA>          7
7       8 <NA>         21
8       9 <NA>          2
9      NA <NA>        506

```

2.5.17 Immunization (3 items)

2.5.17.1 FLUSHOT6 and its cleanup to vax_flu

FLUSHOT6 gives the response to “During the past 12 months, have you had either a flu shot or a flu vaccine that was sprayed in your nose?” The responses are:

- 1 = Yes
- 2 = No
- 7 = Don’t know/Not sure
- 9 = Refused

```

smart_ohio_raw <- smart_ohio_raw %>%
  mutate(vax_flu = 2 - FLUSHOT6,
        vax_flu = replace(vax_flu, vax_flu < 0, NA))

smart_ohio_raw %>% count(FLUSHOT6, vax_flu)

# A tibble: 5 x 3
  FLUSHOT6 vax_flu      n
  <dbl>    <dbl>     <int>
1       1      1  3453
2       2      0  3410

```

```

3      7      NA    26
4      9      NA     3
5     NA      NA   520

```

2.5.17.2 PNEUVAC3 and its cleanup to vax_pneumo

PNEUVAC3 gives the response to “A pneumonia shot or pneumococcal vaccine is usually given only once or twice in a person’s lifetime and is different from the flu shot. Have you ever had a pneumonia shot?” The responses are:

- 1 = Yes
- 2 = No
- 7 = Don’t know/Not sure
- 9 = Refused

```

smart_ohio_raw <- smart_ohio_raw %>%
  mutate(vax_pneumo = 2 - PNEUVAC3,
         vax_pneumo = replace(vax_pneumo, vax_pneumo < 0, NA))

smart_ohio_raw %>% count(PNEUVAC3, vax_pneumo)

# A tibble: 5 x 3
  PNEUVAC3 vax_pneumo     n
  <dbl>      <dbl> <int>
1       1          1  3112
2       2          0  3262
3       7          NA   509
4       9          NA     3
5      NA          NA   526

```

2.5.17.3 SHINGLE2 and its cleanup to vax_shingles

SHINGLE2 gives the response to “Have you ever had the shingles or zoster vaccine?” The responses are:

- 1 = Yes
- 2 = No
- 7 = Don’t know/Not sure
- 9 = Refused

```

smart_ohio_raw <- smart_ohio_raw %>%
  mutate(vax_shingles = 2 - SHINGLE2,
         vax_shingles = replace(vax_shingles, vax_shingles < 0, NA))

smart_ohio_raw %>% count(SHINGLE2, vax_shingles)

```

```
# A tibble: 4 x 3
  SHINGLE2 vax_shingles     n
  <dbl>          <dbl> <int>
1      1            1  1503
2      2            0  2979
3      7            NA   78
4     NA           NA 2852
```

2.5.18 HIV/AIDS (2 items)

2.5.18.1 HIVTST6 and its cleanup to hiv_test

HIVTST6 gives the response to “Have you ever been tested for HIV? Do not count tests you may have had as part of a blood donation. Include testing fluid from your mouth.” The responses are:

- 1 = Yes
- 2 = No
- 7 = Don’t know/Not sure
- 9 = Refused

```
smart_ohio_raw <- smart_ohio_raw %>%
  mutate(hiv_test = 2 - HIVTST6,
        hiv_test = replace(hiv_test, hiv_test < 0, NA))

smart_ohio_raw %>% count(HIVTST6, hiv_test)
```

```
# A tibble: 5 x 3
  HIVTST6 hiv_test     n
  <dbl>    <dbl> <int>
1      1      1  2017
2      2      0  4565
3      7     NA   260
4      9     NA    14
5     NA     NA   556
```

2.5.18.2 HIVRISK5 and its cleanup to hiv_risk

HIVRISK5 gives the response to “I am going to read you a list. When I am done, please tell me if any of the situations apply to you. You do not need to tell me which one. You have injected any drug other than those prescribed for you in the past year. You have been treated for a sexually transmitted disease or STD in the past year. You have given or received money or drugs in exchange for sex in the past year.” The responses are:

- 1 = Yes

2.6. IMPUTING AGE AND INCOME AS QUANTITATIVE FROM THIN AIR105

- 2 = No
- 7 = Don't know/Not sure
- 9 = Refused

```
smart_ohio_raw <- smart_ohio_raw %>%
  mutate(hiv_risk = 2 - HIVRISK5,
        hiv_risk = replace(hiv_risk, hiv_risk < 0, NA))

smart_ohio_raw %>% count(HIVRISK5, hiv_risk)
```

```
# A tibble: 5 x 3
  HIVRISK5 hiv_risk     n
  <dbl>      <dbl> <int>
1       1         1   277
2       2         0  6537
3       7        NA     2
4       9        NA    17
5      NA        NA   579
```

2.6 Imputing Age and Income as Quantitative from Thin Air

This section is purely for teaching purposes. I would never use the variables created in this section for research work.

2.6.1 age_imp: Imputing Age Data

I want a quantitative age variable, so I'm going to create an imputed `age_imp` value for each subject based on their `agegroup`. For each age group, I will assume that each of the ages represented by a value in that age group will be equally likely, and will draw from the relevant uniform distribution to impute age.

```
set.seed(2020432002)

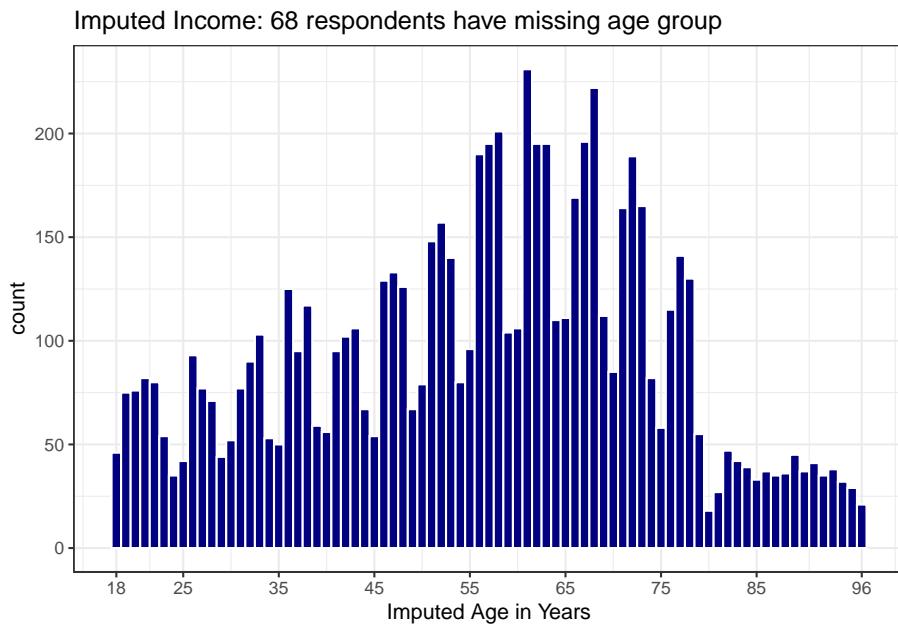
smart_ohio_raw <- smart_ohio_raw %>%
  mutate(age_low = as.numeric(str_sub(as.character(agegroup), 1, 2))) %>%
  mutate(age_high = as.numeric(str_sub(as.character(agegroup), 4, 5))) %>%
  rowwise() %>%
  mutate(age_imp = ifelse(!is.na(agegroup),
                        round(runif(1, min = age_low, max = age_high), 0),
                        NA))

smart_ohio_raw %>% count(agegroup, age_imp) #>% tail()
```

```
# A tibble: 80 x 3
# Rowwise:
#> #> #>   agegroup age_imp     n
#> #> #>   <fct>      <dbl> <int>
#> 1 18-24        18     46
#> 2 18-24        19     75
#> 3 18-24        20     76
#> 4 18-24        21     82
#> 5 18-24        22     80
#> 6 18-24        23     54
#> 7 18-24        24     35
#> 8 25-29        25     42
#> 9 25-29        26     93
#> 10 25-29       27     77
#> # ... with 70 more rows
```

Here is a histogram of the `age_imp` variable.

```
ggplot(smart_ohio_raw, aes(x = age_imp)) +
  geom_histogram(fill = "navy", col = "white",
                 binwidth = 1) +
  scale_x_continuous(breaks = c(18, 25, 35, 45, 55, 65, 75, 85, 96)) +
  labs(x = "Imputed Age in Years",
       title = paste0("Imputed Income: ",
                     sum(is.na(smart_ohio_raw$age_imp)),
                     " respondents have missing age group"))
```



2.6. IMPUTING AGE AND INCOME AS QUANTITATIVE FROM THIN AIR107

2.6.2 inc_imp: Imputing Income Data

I want a quantitative income variable, so I'm going to create an imputed `inc_imp` value for each subject based on their `incomegroup`. For most income groups, I will assume that each of the incomes represented by a value in that income group will be equally likely, and will draw from the relevant uniform distribution to impute income. The exception is the highest income group, where I will impute a value drawn from a distribution that places all values at \$75,000 or more, but has a substantial right skew and long tail.

```
set.seed(2020432001)

smart_ohio_raw <- smart_ohio_raw %>%
  mutate(inc_imp = case_when(
    incomegroup == "0-9K" ~ round(runif(1, min = 100, max = 9999)),
    incomegroup == "10-14K" ~ round(runif(1, min = 10000, max = 14999)),
    incomegroup == "15-19K" ~ round(runif(1, min = 15000, max = 19999)),
    incomegroup == "20-24K" ~ round(runif(1, min = 20000, max = 24999)),
    incomegroup == "25-34K" ~ round(runif(1, min = 25000, max = 34999)),
    incomegroup == "35-49K" ~ round(runif(1, min = 35000, max = 49999)),
    incomegroup == "50-74K" ~ round(runif(1, min = 50000, max = 74999)),
    incomegroup == "75K+" ~ round((rnorm(n = 1, mean = 0, sd = 300)^2) + 74999)))
  )

smart_ohio_raw %>% count(incomegroup, inc_imp) %>% tail()

# A tibble: 6 x 3
# Rowwise:
  incomegroup inc_imp     n
  <fct>        <dbl> <int>
1 75K+         774009     1
2 75K+         798174     1
3 75K+         806161     1
4 75K+         847758     1
5 75K+         1085111    1
6 <NA>          NA      1310
```

Here are density plots of the `inc_imp` variable. The top picture shows the results on a linear scale, and the bottom shows them on a log (base 10) scale.

```
p1 <- ggplot(smart_ohio_raw, aes(x = inc_imp/1000)) +
  geom_density(fill = "darkgreen", col = "white") +
  labs(x = "Imputed Income in Thousands of Dollars",
       title = "Imputed Income on the Linear scale") +
  scale_x_continuous(breaks = c(25, 75, 250, 1000))

p2 <- ggplot(smart_ohio_raw, aes(x = inc_imp/1000)) +
  geom_density(fill = "darkgreen", col = "white") +
```

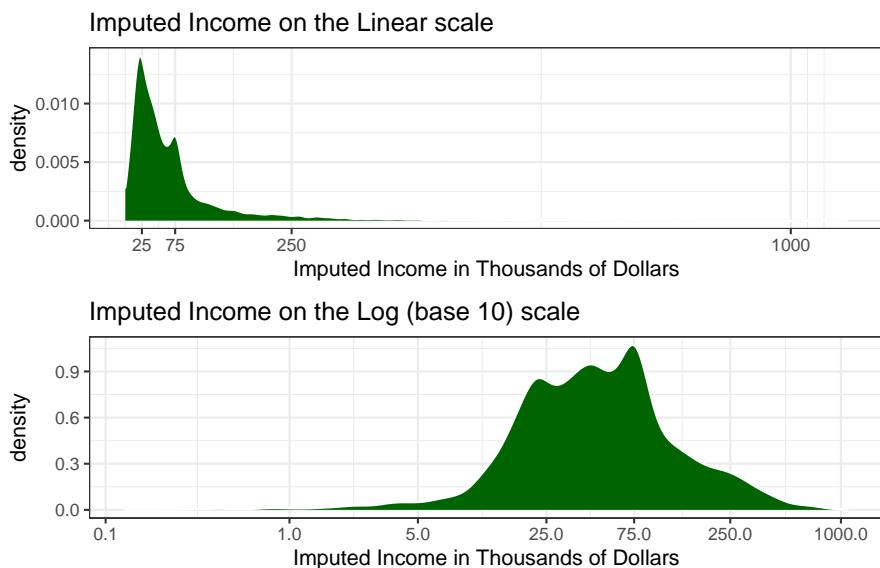
```

  labs(x = "Imputed Income in Thousands of Dollars",
       title = "Imputed Income on the Log (base 10) scale") +
  scale_x_log10(breaks = c(0.1, 1, 5, 25, 75, 250, 1000))

p1 / p2 +
  plot_annotation(title =
    paste0("Imputed Income: ", sum(is.na(smart_ohio_raw$inc_imp)))
  )

```

Imputed Income: 1310 respondents have missing income group



2.7 Clean Data in the State of Ohio

There are six MMSAs associated with the state of Ohio. We're going to create a `smart_ohio` that includes each of them. First, I'll ungroup the data that I created earlier, so I get a clean tibble.

```
smart_ohio_raw <- smart_ohio_raw %>% ungroup()
```

Next, I'll select the variables I want to retain (they are the ones I created, plus `SEQNO`.)

```
smart_ohio <- smart_ohio_raw %>%
  select(SEQNO, mmsa, mmsa_code, mmsa_name, mmsa_wt, completed,
         landline, hhadults,
         genhealth, physhealth, menhealth, poorhealth,
         agegroup, age_imp, race, hispanic, race_eth,
```

```

female, marital, kids, educgroup, home_own,
veteran, employment, incomegroup, inc_imp,
cell_own, internet30,
weight_kg, height_m, bmi, bmigroup,
pregnant, deaf, blind, decide,
diffwalk, diffdress, diffalone,
smoke100, smoker, ecig_ever, ecigs,
healthplan, hasdoc, costprob, t_checkup,
bp_high, bp_meds,
t_chol, chol_high, chol_meds,
asthma, hx_asthma, now_asthma,
hx_mi, hx_chd, hx_stroke, hx_skinc, hx_otherc,
hx_copd, hx_depress, hx_kidney,
hx_diabetes, dm_status, dm_age,
hx_arthr, arth_lims, arth_work, arth_soc,
joint_pain, alcdays, avgdrinks, maxdrinks,
binge, drinks_wk, drink_heavy,
fruit_day, veg_day, eat_juice, eat_fruit,
eat_greenveg, eat_fries, eat_potato,
eat_otherveg, exerany, activity, rec_aerobic,
rec_strength, exer1_type, exer2_type,
exer1_min, exer2_min, seatbelt,
vax_flu, vax_pneumo, vax_shingles,
hiv_test, hiv_risk)

saveRDS(smart_ohio, "data/smart_ohio.Rds")

write_csv(smart_ohio, "data/smart_ohio.csv")

```

The `smart_ohio` file should contain 99 variables, describing 7412 respondents.

2.8 Clean Cleveland-Elyria Data

2.8.1 Cleveland - Elyria Data

The `mmsa_name` variable is probably the simplest way for us to filter our data down to the MMSA we are interested in. Here, I'm using the `str_detect` function to identify the values of `mmsa_name` that contain the text "Cleveland."

```

smart_cle <- smart_ohio %>%
  filter(str_detect(mmsa_name, 'Cleveland'))

saveRDS(smart_cle, "data/smart_cle.Rds")

```

In the Cleveland-Elyria MSA, we have 1133 observations on the same 99 variables. We'll build a variety of smaller subsets from these data, eventually.

Chapter 3

Dealing with Missingness: Single Imputation

3.1 Selecting Some Variables from the `smart_cle` data

```
smart_cle <- readRDS("data/smart_cle.Rds")  
  
smart_cle1 <- smart_cle %>%  
  select(SEQNO, physhealth, genhealth, bmi,  
         age_imp, female, race_eth, internet30,  
         smoke100, activity, drinks_wk, veg_day)
```

The `smart_cle.Rds` data file available on the Data and Code page of our website describes information on 99 variables for 1133 respondents to the BRFSS 2017, who live in the Cleveland-Elyria, OH, Metropolitan Statistical Area. The variables in the `smart_cle1.csv` file are listed below, along with the items that generate these responses.

Variable	Description
SEQNO	respondent identification number (all begin with 2016)
physhealth	Now thinking about your physical health, which includes physical illness and injury, for how many days during the past 30 days was your physical health not good?
genhealth	Would you say that in general, your health is ... (five categories: Excellent, Very Good, Good, Fair or Poor)
bmi	Body mass index, in kg/m ²
age_imp	Age, imputed, in years

Variable	Description
<code>female</code>	Sex, 1 = female, 0 = male
<code>race_eth</code>	Race and Ethnicity, in five categories
<code>internet30</code>	Have you used the internet in the past 30 days? (1 = yes, 0 = no)
<code>smoke100</code>	Have you smoked at least 100 cigarettes in your life? (1 = yes, 0 = no)
<code>activity</code>	Physical activity (Highly Active, Active, Insufficiently Active, Inactive)
<code>drinks_wk</code>	On average, how many drinks of alcohol do you consume in a week?
<code>veg_day</code>	How many servings of vegetables do you consume per day, on average?

```
str(smart_cle1)
```

```
tibble [1,133 x 12] (S3: tbl_df/tbl/data.frame)
$ SEQNO      : num [1:1133] 2.02e+09 2.02e+09 2.02e+09 2.02e+09 2.02e+09 ...
$ physhealth: num [1:1133] 4 0 0 0 0 2 2 0 0 0 ...
$ genhealth : Factor w/ 5 levels "1_Excellent",...: 1 1 3 3 3 2 3 2 4 1 ...
$ bmi        : num [1:1133] NA 23.1 26.9 26.5 24.2 ...
$ age_imp    : num [1:1133] 51 28 37 36 88 43 23 34 58 54 ...
$ female     : num [1:1133] 1 1 1 1 0 0 0 0 0 1 ...
$ race_eth   : Factor w/ 5 levels "White non-Hispanic",...: 1 1 3 1 1 1 1 3 2 1 ...
$ internet30: num [1:1133] 1 1 0 1 1 1 1 1 1 ...
$ smoke100   : num [1:1133] 1 0 0 1 1 1 0 0 0 1 ...
$ activity   : Factor w/ 4 levels "Highly_Active",...: 4 4 3 1 1 NA 1 1 1 1 ...
$ drinks_wk  : num [1:1133] 0.7 0 0 4.67 0.93 0 2 0 0 0.47 ...
$ veg_day    : num [1:1133] NA 3 4.06 2.07 1.31 NA 1.57 0.83 0.49 1.72 ...
```

3.2 smart_cle1: Seeing our Missing Data

The `naniar` package provides several useful functions for summarizing missingness in our data set. Like all tidy data sets, our `smart_cle1` tibble contains rows which describe observations, sometimes called *cases*, and also contains columns which describe variables.

Overall, there are 1133 cases, and 1133 observations in our `smart_cle1` tibble.

- We can obtain a count of the number of missing cells in the entire tibble.

```
smart_cle1 %>% n_miss()
```

```
[1] 479
```

- We can use the `miss_var_summary` function to get a sorted table of each variable by number missing.

```
miss_var_summary(smart_cle1) %>% knitr::kable()
```

variable	n_miss	pct_miss
activity	109	9.6204766
veg_day	101	8.9143866
bmi	91	8.0317741
drinks_wk	66	5.8252427
smoke100	40	3.5304501
race_eth	26	2.2947926
physhealth	24	2.1182701
age_imp	11	0.9708738
internet30	7	0.6178288
genhealth	4	0.3530450
SEQNO	0	0.0000000
female	0	0.0000000

- Or we can use the `miss_var_table` function to tabulate the number of variables that have each observed level of missingness.

```
miss_var_table(smart_cle1)
```

	# A tibble: 11 x 3	n_miss_in_var	n_vars	pct_vars
		<int>	<int>	<dbl>
1		0	2	16.7
2		4	1	8.33
3		7	1	8.33
4		11	1	8.33
5		24	1	8.33
6		26	1	8.33
7		40	1	8.33
8		66	1	8.33
9		91	1	8.33
10		101	1	8.33
11		109	1	8.33

- Or we can get a count for a specific variable, like `activity`:

```
smart_cle1 %>% select(activity) %>% n_miss()
```

```
[1] 109
```

- We can also use `prop_miss_case` or `pct_miss_case` to specify the proportion (or percentage) of missing observations across an entire data set, or within a specific variable.

```
prop_miss_case(smart_cle1)

[1] 0.2127096
```

```
smart_cle1 %>% select(activity) %>% pct_miss_case(.)
```

```
[1] 9.620477
```

- We can also use `prop_miss_var` or `pct_miss_var` to specify the proportion (or percentage) of variables with missing observations across an entire data set.

```
prop_miss_var(smart_cle1)
```

```
[1] 0.8333333
```

```
pct_miss_var(smart_cle1)
```

```
[1] 83.33333
```

- We use `miss_case_table` to identify the number of missing values for each of the cases (rows) in our tibble.

```
miss_case_table(smart_cle1)
```

```
# A tibble: 7 x 3
  n_miss_in_case n_cases pct_cases
  <int>     <int>     <dbl>
1 0          892      78.7
2 1          129      11.4
3 2           51      4.50
4 3           22      1.94
5 4           21      1.85
6 5           10      0.883
7 6            8      0.706
```

- Use `miss_case_summary` to specify individual observations and count their missing values.

```
miss_case_summary(smart_cle1)
```

```
# A tibble: 1,133 x 3
  case n_miss pct_miss
  <int>   <int>     <dbl>
1 17       6      50
2 42       6      50
3 254      6      50
4 425      6      50
5 521      6      50
6 729      6      50
```

```

7   757      6      50
8  1051      6      50
9    89      5     41.7
10   94      5     41.7
# ... with 1,123 more rows

```

The case numbers identified here are row numbers. Extract the data for case 17, for instance, with the `slice` function.

```
smart_cle1 %>% slice(17)
```

```

# A tibble: 1 x 12
  SEQNO physhealth genhealth   bmi age_imp female race_eth internet30 smoke100
  <dbl>      <dbl> <fct>     <dbl> <dbl> <fct>        <dbl>      <dbl>
1 2.02e9       0 1_Excell~    NA     50  0 White n~       NA      NA
# ... with 3 more variables: activity <fct>, drinks_wk <dbl>, veg_day <dbl>

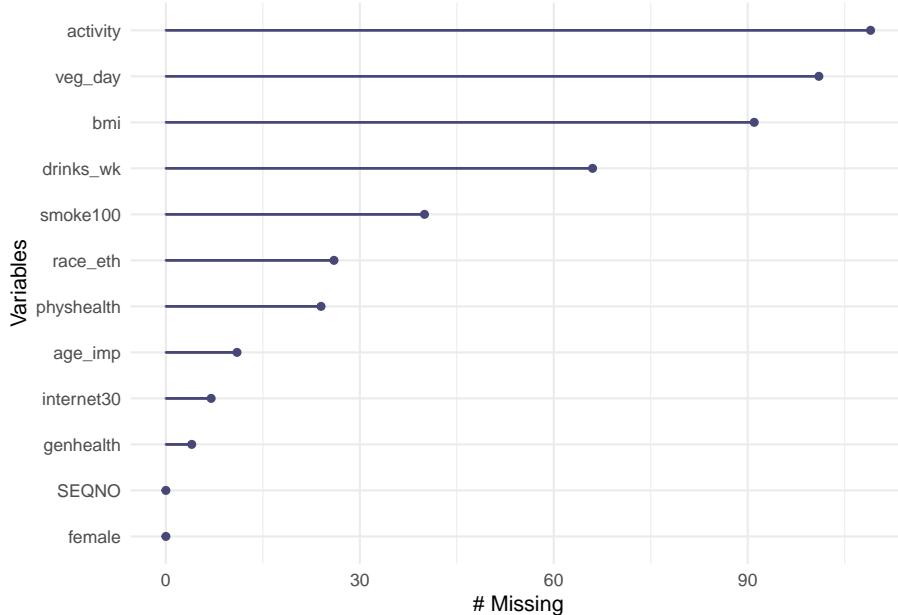
```

3.2.1 Plotting Missingness

The `gg_miss_var` function plots the number of missing observations in each variable in our data set.

```
gg_miss_var(smart_cle1)
```

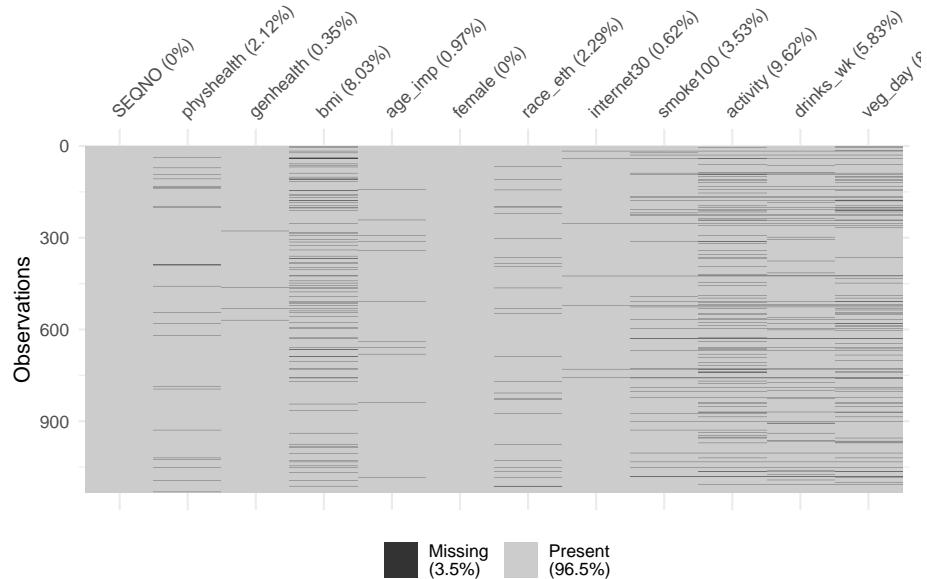
Warning: It is deprecated to specify `guide = FALSE` to remove a guide. Please use `guide = "none"` instead.



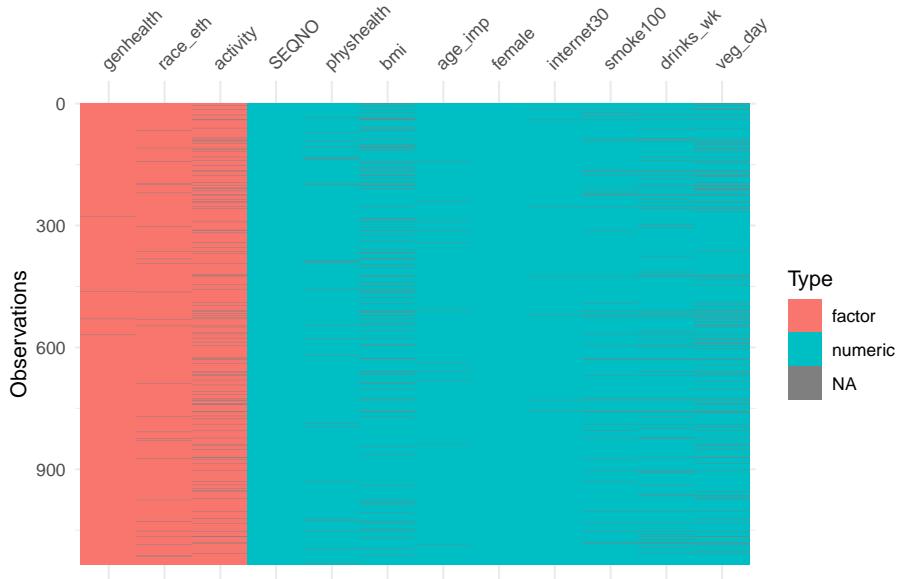
So the most commonly missing variable is `activity` which, as we've seen, has 109 missing values.

To get a general sense of the missingness in our data, we might use either the `vis_dat` or the `vis_miss` function from the `visdat` package.

```
vis_miss(smart_cle1)
```



```
vis_dat(smart_cle1)
```



3.3 Missing-data mechanisms

My source for this description of mechanisms is Chapter 25 of Gelman and Hill (2007), and that chapter is available at this link.

1. **MCAR = Missingness completely at random.** A variable is missing completely at random if the probability of missingness is the same for all units, for example, if for each subject, we decide whether to collect the `diabetes` status by rolling a die and refusing to answer if a “6” shows up. If data are missing completely at random, then throwing out cases with missing data does not bias your inferences.
2. **Missingness that depends only on observed predictors.** A more general assumption, called **missing at random** or **MAR**, is that the probability a variable is missing depends only on available information. Here, we would have to be willing to assume that the probability of nonresponse to `diabetes` depends only on the other, fully recorded variables in the data. It is often reasonable to model this process as a logistic regression, where the outcome variable equals 1 for observed cases and 0 for missing. When an outcome variable is missing at random, it is acceptable to exclude the missing cases (that is, to treat them as NA), as long as the regression controls for all the variables that affect the probability of missingness.
3. **Missingness that depends on unobserved predictors.** Missingness is no longer “at random” if it depends on information that has not been

recorded and this information also predicts the missing values. If a particular treatment causes discomfort, a patient is more likely to drop out of the study. This missingness is not at random (unless “discomfort” is measured and observed for all patients). If missingness is not at random, it must be explicitly modeled, or else you must accept some bias in your inferences.

4. **Missingness that depends on the missing value itself.** Finally, a particularly difficult situation arises when the probability of missingness depends on the (potentially missing) variable itself. For example, suppose that people with higher earnings are less likely to reveal them.

Essentially, situations 3 and 4 are referred to collectively as **non-random missingness**, and cause more trouble for us than 1 and 2.

3.4 Options for Dealing with Missingness

There are several available methods for dealing with missing data that are MCAR or MAR, but they basically boil down to:

- Complete Case (or Available Case) analyses
- Single Imputation
- Multiple Imputation

3.5 Complete Case (and Available Case) analyses

In **Complete Case** analyses, rows containing NA values are omitted from the data before analyses commence. This is the default approach for many statistical software packages, and may introduce unpredictable bias and fail to include some useful, often hard-won information.

- A complete case analysis can be appropriate when the number of missing observations is not large, and the missing pattern is either MCAR (missing completely at random) or MAR (missing at random.)
- Two problems arise with complete-case analysis:
 1. If the units with missing values differ systematically from the completely observed cases, this could bias the complete-case analysis.
 2. If many variables are included in a model, there may be very few complete cases, so that most of the data would be discarded for the sake of a straightforward analysis.
- A related approach is *available-case* analysis where different aspects of a problem are studied with different subsets of the data, perhaps identified on the basis of what is missing in them.

3.6 Single Imputation

In **single imputation** analyses, NA values are estimated/replaced *one time* with *one particular data value* for the purpose of obtaining more complete samples, at the expense of creating some potential bias in the eventual conclusions or obtaining slightly *less* accurate estimates than would be available if there were no missing values in the data.

- A single imputation can be just a replacement with the mean or median (for a quantity) or the mode (for a categorical variable.) However, such an approach, though easy to understand, underestimates variance and ignores the relationship of missing values to other variables.
- Single imputation can also be done using a variety of models to try to capture information about the NA values that are available in other variables within the data set.
- The **simputation** package can help us execute single imputations using a wide variety of techniques, within the pipe approach used by the **tidyverse**. Another approach I have used in the past is the **mice** package, which can also perform single imputations.

3.7 Multiple Imputation

Multiple imputation, where NA values are repeatedly estimated/replaced with multiple data values, for the purpose of obtaining more complete samples *and* capturing details of the variation inherent in the fact that the data have missingness, so as to obtain *more* accurate estimates than are possible with single imputation.

- We'll postpone the discussion of multiple imputation for a while.

3.8 Approach 1: Building a Complete Case Analysis: **smart_cle1_cc**

In the 431 course, we usually dealt with missing data by restricting our analyses to respondents with complete data on all variables. Let's start by doing that here. We'll create a new tibble called **smart_cle1_cc** which includes all respondents with complete data on all of these variables.

```
smart_cle1_cc <- smart_cle1 %>%
  drop_na()

dim(smart_cle1_cc)
```

```
[1] 892 12
```

Our `smart_cle1_cc` tibble now has many fewer observations than its predecessors, but all of the variables in this complete cases tibble have no missing observations.

	Data Set	Rows	Columns	Missingness?
	<code>smart_cle</code>	1133	99	Quite a bit.
	<code>smart_cle1</code>	1133	12	Quite a bit.
	<code>smart_cle1_cc</code>	892	12	None.

3.9 Approach 2: Single Imputation to create `smart_cle1_sh`

Next, we'll create a data set which has all of the rows in the original `smart_cle1` tibble, but deals with missingness by imputing (estimating / filling in) new values for each of the missing values. To do this, we'll make heavy use of the `imputation` package in R.

The `imputation` package is designed for single imputation work. Note that we'll eventually adopt a **multiple imputation** strategy in some of our modeling work, and we'll use some specialized tools to facilitate that later.

To begin, we'll create a “shadow” in our tibble to track what we'll need to impute.

```
smart_cle1_sh <- bind_shadow(smart_cle1)

names(smart_cle1_sh)

[1] "SEQNO"          "physhealth"      "genhealth"      "bmi"
[5] "age_imp"        "female"         "race_eth"       "internet30"
[9] "smoke100"       "activity"        "drinks_wk"     "veg_day"
[13] "SEQNO_NA"       "physhealth_NA"  "genhealth_NA"   "bmi_NA"
[17] "age_imp_NA"    "female_NA"      "race_eth_NA"   "internet30_NA"
[21] "smoke100_NA"   "activity_NA"    "drinks_wk_NA"  "veg_day_NA"
```

Note that the `bind_shadow()` function doubles the number of variables in our tibble, specifically by creating a new variable for each that takes the value `!NA` or `NA`. For example, consider

```
smart_cle1_sh %>% count(activity, activity_NA)

# A tibble: 5 x 3
  activity      activity_NA     n
  <fct>        <fct>      <int>
1 active        active        5
2 inactive     inactive       1
3 missing       missing       1
4 not missing  not missing   3
5 not active   not active   1
```

3.9. APPROACH 2: SINGLE IMPUTATION TO CREATE SMART_CLE1_SH121

1 Highly_Active	!NA	338
2 Active	!NA	173
3 Insufficiently_Active	!NA	201
4 Inactive	!NA	312
5 <NA>	NA	109

The `activity_NA` variable takes the value `!NA` (meaning not missing) when the value of the `activity` variable is known, and takes the value `NA` for observations where the `activity` variable is missing. This background tracking will be helpful to us when we try to assess the impact of imputation on some of our summaries.

3.9.1 What Type of Missingness Do We Have?

There are three types of missingness that we might assume in any given setting: missing completely at random (MCAR), missing at random (MAR) and missing not at random (MNAR). Together, MCAR and MAR are sometimes called *ignorable* non-response, which essentially means that imputation provides a way to useful estimates. MNAR or missing NOT at random is sometimes called non-ignorable missingness, implying that even high-quality imputation may not be sufficient to provide useful information to us.

Missing Completely at Random means that the missing data points are a random subset of the data. Essentially, there is nothing that makes some data more likely to be missing than others. If the data truly match the standard for MCAR, then a complete-case analysis will be about as good as an analysis after single or multiple imputation.

Missing at Random means that there is a systematic relationship between the observed data and the missingness mechanism. Another way to say this is that the missing value is not related to the reason why it is missing, but is related to the other variables collected in the study. The implication is that the missingness can be accounted for by studying the variables with complete information. Imputation strategies can be very helpful here, incorporating what we know (or think we know) about the relationships between the results that are missing and the results that we see.

- Wikipedia provides a nice example. If men are less likely to fill in a depression survey, but this has nothing to do with their level of depression after accounting for the fact that they are male, then the missingness can be assumed MAR.
- Determining whether missingness is MAR or MNAR can be tricky. We'll spend more time discussing this later.

Missing NOT at Random means that the missing value is related to the reason why it is missing.

- Continuing the Wikipedia example, if men failed to fill in a depression survey because of their level of depression, then this would be MNAR.

- Single imputation is most helpful in the MAR situation, although it is also appropriate when we assume MCAR.
- Multiple imputation will, similarly, be more helpful in MCAR and MAR situations than when data are missing NOT at random.

It's worth noting that many people are unwilling to impute values for outcomes or key predictors in a modeling setting, but are happy to impute for less important covariates. For now, we'll assume MCAR or MAR for all of the missingness in our `smart_cle1` data, which will allow us to adopt a single imputation strategy.

3.9.2 Single imputation into `smart_cle1_sh`

Which variables in `smart_cle1_sh` contain missing data?

```
miss_var_summary(smart_cle1_sh)
```

```
# A tibble: 24 x 3
  variable   n_miss pct_miss
  <chr>      <int>    <dbl>
1 activity     109    9.62
2 veg_day      101    8.91
3 bmi          91     8.03
4 drinks_wk    66     5.83
5 smoke100     40     3.53
6 race_eth     26     2.29
7 physhealth   24     2.12
8 age_imp       11    0.971
9 internet30    7    0.618
10 genhealth    4    0.353
# ... with 14 more rows
```

We will impute these variables using several different strategies, all supported nicely by the `simputation` package.

These include imputation methods based solely on the distribution of the complete cases of the variable being imputed.

- `impute_median`: impute the median value of all non-missing observations into the missing values for the variable
- `impute_rhd`: random “hot deck” imputation involves drawing at random from the complete cases for that variable

Also available are imputation strategies that impute predicted values from models using other variables in the data set besides the one being imputed.

- `impute_pmm`: imputation using predictive mean matching
- `impute_rlm`: imputation using robust linear models
- `impute_cart`: imputation using classification and regression trees

3.9. APPROACH 2: SINGLE IMPUTATION TO CREATE SMART_CLE1_SH123

- `impute_knn`: imputation using k-nearest neighbors methods

3.9.3 Imputing Binary Categorical Variables

Here, we'll arbitrarily impute our 1/0 variables as follows:

- For `internet30` we'll use the `impute_rhd` approach to draw a random observation from the existing set of 1s and 0s in the complete `internet30` data.
- For `smoke100` we'll use a method called predictive mean matching (`impute_pmm`) which takes the result from a model based on the (imputed) `internet30` value and whether or not the subject is `female`, and converts it to the nearest value in the observed `smoke100` data. This is a good approach for imputing discrete variables.

These are completely arbitrary choices, for demonstration purposes.

```
set.seed(2020001)
smart_cle1_sh <- smart_cle1_sh %>%
  data.frame() %>%
  impute_rhd(.,
              internet30 ~ 1) %>%
  impute_pmm(., smoke100 ~ internet30 + female) %>%
 tbl_df()

Warning: `tbl_df()` was deprecated in dplyr 1.0.0.
Please use `tibble::as_tibble()` instead.
This warning is displayed once every 8 hours.
Call `lifecycle::last_lifecycle_warnings()` to see where this warning was generated.
smart_cle1_sh %>% count(smoke100, smoke100_NA)

# A tibble: 4 x 3
  smoke100 smoke100_NA     n
  <dbl> <fct>      <int>
1       0 !NA        579
2       0 NA         21
3       1 !NA        514
4       1 NA         19

smart_cle1_sh %>% count(internet30, internet30_NA)

# A tibble: 4 x 3
  internet30 internet30_NA     n
  <dbl> <fct>      <int>
1       0 !NA        207
2       0 NA          1
3       1 !NA        919
```

4	1 NA	6
---	------	---

Other approaches that may be used with 1/0 variables include `impute_knn` and `impute_pmm`.

3.9.4 Imputing Quantitative Variables

We'll demonstrate a different approach for imputing each of the quantitative variables with missing observations. Again, we're making purely arbitrary decisions here about what to include in each imputation. In practical work, we'd want to be a bit more thoughtful about this.

Note that I'm choosing to use `impute_pmm` with the `physhealth` and `age_imp` variables. This is (in part) because I want my imputations to be integers, as the other observations are for those variables. `impute_rhd` would also accomplish this.

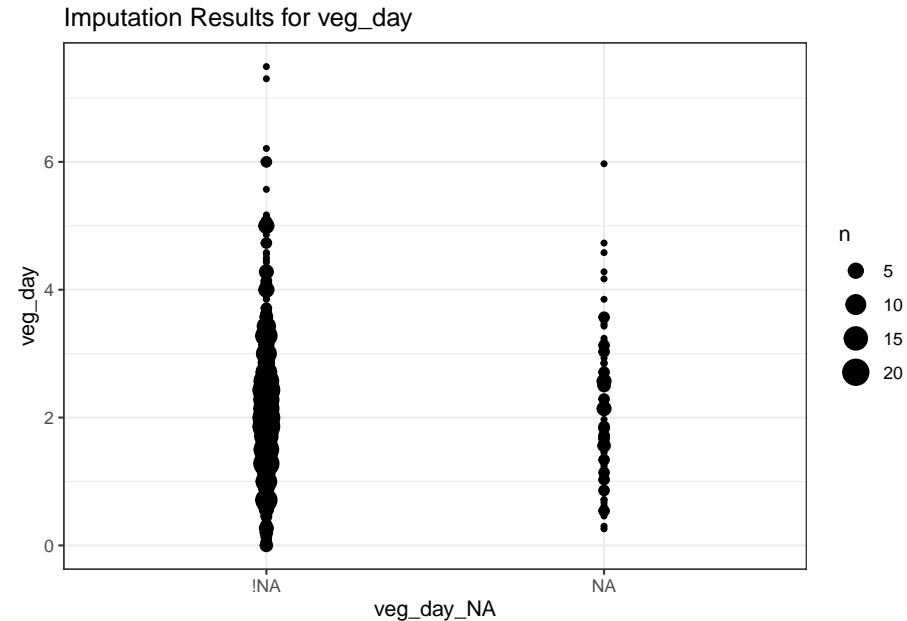
```
set.seed(2020001)
smart_cle1_sh <- smart_cle1_sh %>%
  data.frame() %>%
  impute_rhd(., veg_day ~ 1) %>%
  impute_median(., drinks_wk ~ 1) %>%
  impute_pmm(., physhealth ~
              drinks_wk + female + smoke100) %>%
  impute_pmm(., age_imp ~ drinks_wk + physhealth) %>%
  impute_rlm(., bmi ~ physhealth + smoke100) %>%
tbl_df()
```

3.9.5 Imputation Results

Let's plot a few of these results, so we can see what imputation has done to the distribution of these quantities.

```
1. veg_day
ggplot(smart_cle1_sh, aes(x = veg_day_NA, y = veg_day)) +
  geom_count() +
  labs(title = "Imputation Results for veg_day")
```

3.9. APPROACH 2: SINGLE IMPUTATION TO CREATE SMART_CLE1_SH125



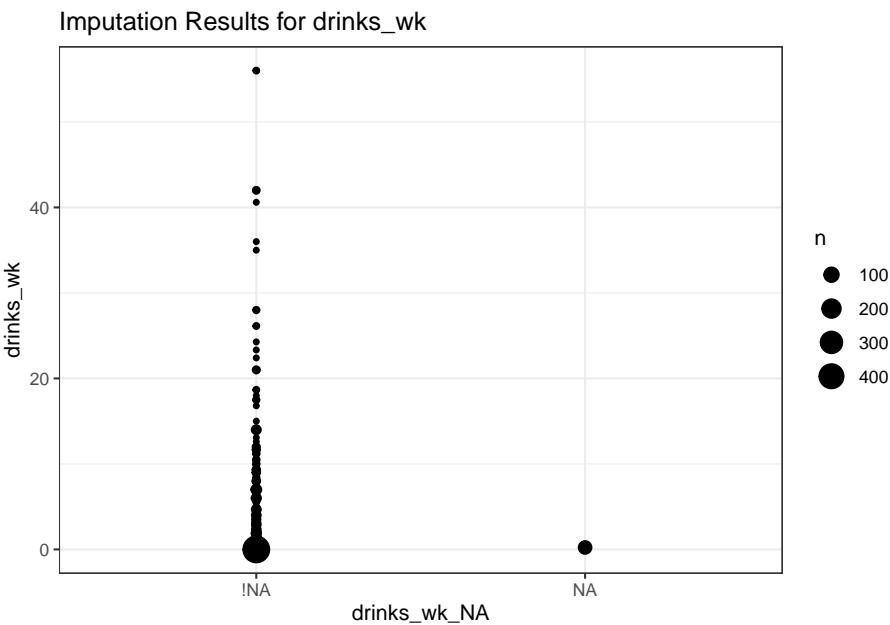
```
smart_cle1_sh %%%
mosaic::favstats(veg_day ~ veg_day_NA)
```

```
Registered S3 method overwritten by 'mosaic':
method                  from
fortify.SpatialPolygonsDataFrame ggplot2
```

	veg_day_NA	min	Q1	median	Q3	max	mean	sd	n	missing
1	!NA	0.00	1.2675	1.72	2.42	7.49	1.912548	1.038403	1032	0
2	NA	0.26	1.3400	1.86	2.72	5.97	2.085050	1.062316	101	0

2. drinks_wk for which we imputed the median value...

```
ggplot(smart_cle1_sh, aes(x = drinks_wk_NA, y = drinks_wk)) +
  geom_count() +
  labs(title = "Imputation Results for drinks_wk")
```



```
smart_cle1_sh %>% filter(drinks_wk_NA == "NA") %>%
  tabyl(drinks_wk)
```

drinks_wk	n	percent
0.23	66	1

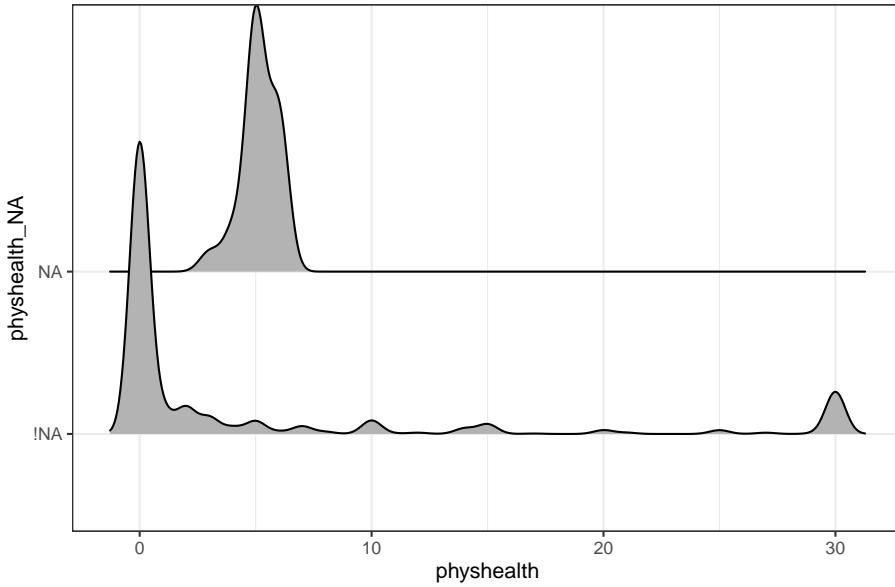
3. physhealth, a count between 0 and 30...

```
ggplot(smart_cle1_sh,
       aes(x = physhealth, y = physhealth_NA)) +
  geom_density_ridges() +
  labs(title = "Imputation Results for physhealth")
```

Picking joint bandwidth of 0.426

3.9. APPROACH 2: SINGLE IMPUTATION TO CREATE SMART_CLE1_SH127

Imputation Results for physhealth

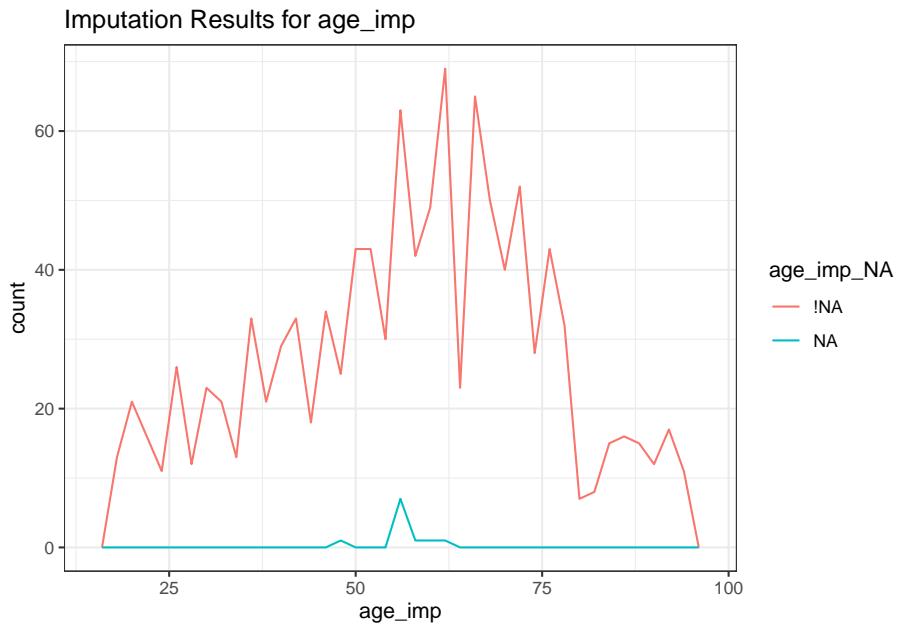


```
smart_cle1_sh %>% filter(physhealth_NA == "NA") %>%
  tabyl(physhealth)
```

physhealth	n	percent
3	1	0.04166667
4	2	0.08333333
5	13	0.54166667
6	8	0.33333333

4. age_imp, in (integer) years

```
ggplot(smart_cle1_sh,
       aes(x = age_imp, color = age_imp_NA)) +
  geom_freqpoly(binwidth = 2) +
  labs(title = "Imputation Results for age_imp")
```



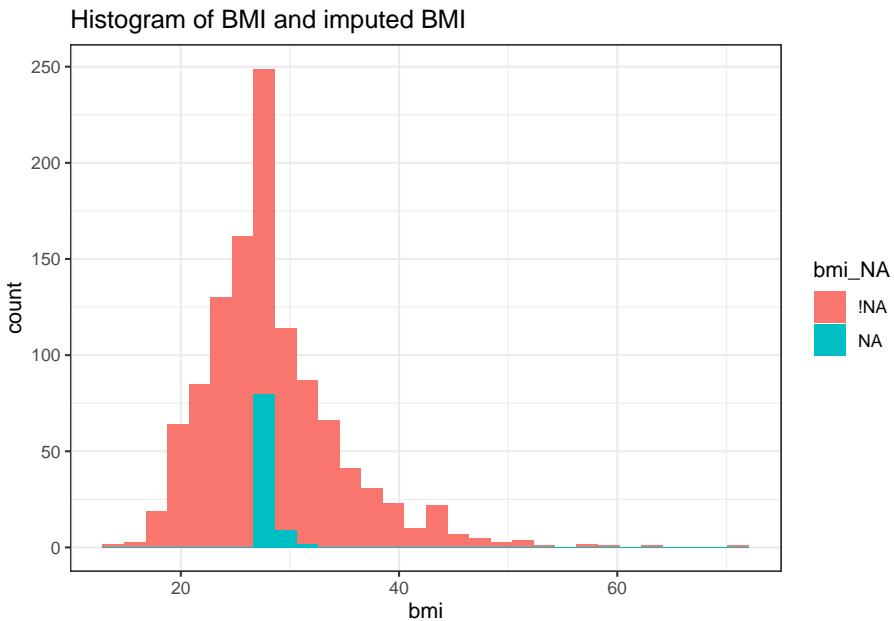
```
smart_cle1_sh %>% filter(age_imp_NA == "NA") %>%
  tabyl(age_imp)
```

age_imp	n	percent
48	1	0.09090909
57	7	0.63636364
58	1	0.09090909
61	1	0.09090909
63	1	0.09090909

5. bmi or body mass index

```
ggplot(smart_cle1_sh, aes(x = bmi, fill = bmi_NA)) +
  geom_histogram(bins = 30) +
  labs(title = "Histogram of BMI and imputed BMI")
```

3.9. APPROACH 2: SINGLE IMPUTATION TO CREATE SMART_CLE1_SH129



```
smart_cle1_sh %$% mosaic::favstats(bmi ~ bmi_NA)
```

bmi_NA	min	Q1	median	Q3	max	mean	sd	n
1 !NA	13.3000	24.1100	27.30000	31.68000	70.56000	28.40947	6.6289286	1042
2 NA	27.0693	27.0693	27.50229	27.66574	30.75898	27.66057	0.8964101	91
missing								
1	0							
2	0							

3.9.6 Imputing Multi-Categorical Variables

The three multi-categorical variables we have left to impute are `activity`, `race_eth` and `genhealth`, and each is presented as a factor in R, rather than as a character variable.

We'll arbitrarily decide to impute

- `activity` and `genhealth` with a classification tree using `physhealth`, `bmi` and `smoke100`,
- and then impute `race_eth` with a random draw from the distribution of complete cases.

```
set.seed(2020001)
smart_cle1_sh <- smart_cle1_sh %>%
  data.frame() %>%
  impute_cart(., activity + genhealth ~
```

```
physhealth + bmi + smoke100) %>%
  impute_rhd(., race_eth ~ 1) %>%
 tbl_df()
```

Let's check our results.

```
smart_cle1_sh %>% count(activity_NA, activity)
```

```
# A tibble: 6 x 3
  activity_NA activity             n
  <fct>      <fct>            <int>
  1 !NA        Highly_Active     338
  2 !NA        Active           173
  3 !NA        Insufficiently_Active 201
  4 !NA        Inactive          312
  5 NA         Highly_Active     90
  6 NA         Inactive          19
```

```
smart_cle1_sh %>% count(race_eth_NA, race_eth)
```

```
# A tibble: 9 x 3
  race_eth_NA race_eth             n
  <fct>      <fct>            <int>
  1 !NA        White non-Hispanic 805
  2 !NA        Black non-Hispanic 222
  3 !NA        Other race non-Hispanic 24
  4 !NA        Multiracial non-Hispanic 22
  5 !NA        Hispanic          34
  6 NA         White non-Hispanic 19
  7 NA         Black non-Hispanic 4
  8 NA         Multiracial non-Hispanic 2
  9 NA         Hispanic          1
```

```
smart_cle1_sh %>% count(genhealth_NA, genhealth)
```

```
# A tibble: 7 x 3
  genhealth_NA genhealth             n
  <fct>      <fct>            <int>
  1 !NA        1_Excellent     164
  2 !NA        2_VeryGood      383
  3 !NA        3_Good          364
  4 !NA        4_Fair           158
  5 !NA        5_Poor           60
  6 NA         2_VeryGood       3
  7 NA         3_Good            1
```

And now, we should have no missing values in the data, at all.

3.9. APPROACH 2: SINGLE IMPUTATION TO CREATE SMART_CLE1_SH131

```
miss_case_table(smart_cle1_sh)

# A tibble: 1 x 3
  n_miss_in_case n_cases pct_cases
  <int>     <int>      <dbl>
1          0     1133       100
```

3.9.7 Saving the new tibbles

```
saveRDS(smart_cle1_cc, here("data", "smart_cle1_cc.Rds"))
saveRDS(smart_cle1_sh, here("data", "smart_cle1_sh.Rds"))
```


Chapter 4

Summarizing the `smart_cle1` data

In this chapter, we'll work with the two data files we built in the previous chapter.

```
smart_cle1_sh <- readRDS(here("data", "smart_cle1_sh.Rds"))
smart_cle1_cc <- readRDS(here("data", "smart_cle1_cc.Rds"))
```

Those files (_`sh` contains single imputations, and a shadow set of variables which have _`NA` at the end of their names, while _`cc` contains only the complete cases) describe information on the following variables from the BRFSS 2017, who live in the Cleveland-Elyria, OH, Metropolitan Statistical Area.

Variable	Description
<code>SEQNO</code>	respondent identification number (all begin with 2016)
<code>physhealth</code>	Now thinking about your physical health, which includes physical illness and injury, for how many days during the past 30 days was your physical health not good?
<code>genhealth</code>	Would you say that in general, your health is ... (five categories: Excellent, Very Good, Good, Fair or Poor)
<code>bmi</code>	Body mass index, in kg/m ²
<code>age_imp</code>	Age, imputed, in years
<code>female</code>	Sex, 1 = female, 0 = male
<code>race_eth</code>	Race and Ethnicity, in five categories
<code>internet30</code>	Have you used the internet in the past 30 days? (1 = yes, 0 = no)
<code>smoke100</code>	Have you smoked at least 100 cigarettes in your life? (1 = yes, 0 = no)
<code>activity</code>	Physical activity (Highly Active, Active, Insufficiently Active, Inactive)

Variable	Description
drinks_wk	On average, how many drinks of alcohol do you consume in a week?
veg_day	How many servings of vegetables do you consume per day, on average?

4.1 General Approaches to Obtaining Numeric Summaries

4.1.1 summary for a data frame

Of course, we can use the usual `summary` to get some basic information about the data.

```
summary(smart_cle1_cc)

      SEQNO          physhealth        genhealth         bmi
Min. :2.017e+09  Min. : 0.000  1_Excellent:164  Min. :13.30
1st Qu.:2.017e+09 1st Qu.: 0.000  2_VeryGood :386   1st Qu.:24.38
Median :2.017e+09 Median : 0.000  3_Good     :365   Median :27.31
Mean   :2.017e+09 Mean  : 4.681  4_Fair     :158   Mean  :28.35
3rd Qu.:2.017e+09 3rd Qu.: 4.000  5_Poor    : 60   3rd Qu.:31.08
Max.   :2.017e+09 Max.  :30.000                    Max.  :70.56

      age_imp       female           race_eth
Min.   :18.00  Min.   :0.0000  White non-Hispanic   :824
1st Qu.:45.00  1st Qu.:0.0000  Black non-Hispanic   :226
Median :58.00  Median :1.0000  Other race non-Hispanic : 24
Mean   :57.33  Mean   :0.5931  Multiracial non-Hispanic: 24
3rd Qu.:70.00  3rd Qu.:1.0000  Hispanic            : 35
Max.   :95.00  Max.   :1.0000

      internet30      smoke100           activity       drinks_wk
Min.   :0.0000  Min.   :0.0000  Highly_Active   :428  Min.   : 0.000
1st Qu.:1.0000  1st Qu.:0.0000  Active          :173  1st Qu.: 0.000
Median :1.0000  Median :0.0000  Insufficiently_Active:201 Median : 0.230
Mean   :0.8164  Mean   :0.4704  Inactive        :331  Mean   : 2.474
3rd Qu.:1.0000  3rd Qu.:1.0000
Max.   :1.0000  Max.   :1.0000

      veg_day      SEQNO_NA  physhealth_NA  genhealth_NA  bmi_NA  age_imp_NA
Min.   :0.000  !NA:1133  !NA:1109    !NA:1129    !NA:1042  !NA:1122
1st Qu.:1.270 NA :  0   NA : 24      NA :  4     NA :  91   NA :  11
Median :1.730
Mean   :1.928
3rd Qu.:2.430
```

4.1. GENERAL APPROACHES TO OBTAINING NUMERIC SUMMARIES 135

```
Max. : 7.490
female_NA  race_eth_NA internet30_NA smoke100_NA activity_NA drinks_wk_NA
!NA:1133  !NA:1107  !NA:1126  !NA:1093  !NA:1024  !NA:1067
NA : 0  NA : 26  NA : 7  NA : 40  NA : 109  NA : 66
```

```
veg_day_NA
!NA:1032
NA : 101
```

4.1.2 The `inspect` function from the `mosaic` package

```
smart_cle1_cc %>% mosaic::inspect()

categorical variables:
  name   class levels   n missing
1 genhealth factor      5 1133      0
2 race_eth factor      5 1133      0
3 activity factor       4 1133      0
                                         distribution
1 2_VeryGood (34.1%) , 3_Good (32.2%) ...
2 White non-Hispanic (72.7%) ...
3 Highly_Active (37.8%) ...

quantitative variables:
  name   class      min        Q1       median        Q3
...1    SEQNO numeric 2.017e+09 2.017e+09 2.017001e+09 2.017001e+09
...2  physhealth numeric 0.000e+00 0.000e+00 0.000000e+00 4.000000e+00
...3      bmi numeric 1.330e+01 2.438e+01 2.731000e+01 3.108000e+01
...4    age_imp numeric 1.800e+01 4.500e+01 5.800000e+01 7.000000e+01
...5     female numeric 0.000e+00 0.000e+00 1.000000e+00 1.000000e+00
...6  internet30 numeric 0.000e+00 1.000e+00 1.000000e+00 1.000000e+00
...7    smoke100 numeric 0.000e+00 0.000e+00 0.000000e+00 1.000000e+00
...8   drinks_wk numeric 0.000e+00 0.000e+00 2.300000e-01 2.100000e+00
...9    veg_day numeric 0.000e+00 1.270e+00 1.730000e+00 2.430000e+00
                                         max        mean        sd   n missing
...1 2.017001e+09 2.017001e+09 327.2132332 1133      0
...2 3.000000e+01 4.681377e+00  9.1208987 1133      0
```

```

...3 7.056000e+01 2.834932e+01   6.3651826 1133      0
...4 9.500000e+01 5.732568e+01   18.0803278 1133      0
...5 1.000000e+00 5.931156e-01   0.4914699 1133      0
...6 1.000000e+00 8.164166e-01   0.3873150 1133      0
...7 1.000000e+00 4.704325e-01   0.4993454 1133      0
...8 5.600000e+01 2.473689e+00   5.6900315 1133      0
...9 7.490000e+00 1.927926e+00   1.0412415 1133      0

shade variables:
  name class levels    n missing
1     SEQNO_NA shade      2 1133      0
2 physhealth_NA shade      2 1133      0
3 genhealth_NA shade      2 1133      0
4     bmi_NA shade      2 1133      0
5 age_imp_NA shade      2 1133      0
6 female_NA shade      2 1133      0
7 race_eth_NA shade      2 1133      0
8 internet30_NA shade      2 1133      0
9 smoke100_NA shade      2 1133      0
10 activity_NA shade      2 1133      0
11 drinks_wk_NA shade      2 1133      0
12 veg_day_NA shade      2 1133      0

distribution
1 !NA (100%), NA (0%)
2 !NA (97.9%), NA (2.1%)
3 !NA (99.6%), NA (0.4%)
4 !NA (92%), NA (8%)
5 !NA (99%), NA (1%)
6 !NA (100%), NA (0%)
7 !NA (97.7%), NA (2.3%)
8 !NA (99.4%), NA (0.6%)
9 !NA (96.5%), NA (3.5%)
10 !NA (90.4%), NA (9.6%)
11 !NA (94.2%), NA (5.8%)
12 !NA (91.1%), NA (8.9%)

```

4.1.3 The `describe` function in Hmisc

This provides some useful additional summaries, including a list of the lowest and highest values (which is very helpful when checking data.)

```

smart_cle1_cc %>%
  select(bmi, genhealth, female) %>%
  Hmisc::describe()

```

```

3 Variables      1133 Observations
-----
bmi
  n   missing  distinct      Info      Mean      Gmd      .05      .10
  1133       0       558        1     28.35     6.681    20.09    21.37
  .25       .50       .75        .90      .95
  24.38     27.31     31.08      36.37     40.44

lowest : 13.30 13.64 15.59 15.71 15.75, highest: 56.31 57.12 58.98 63.00 70.56
-----
genhealth
  n   missing  distinct
  1133       0       5

lowest : 1_Excellent 2_VeryGood 3_Good      4_Fair      5_Poor
highest: 1_Excellent 2_VeryGood 3_Good      4_Fair      5_Poor

Value      1_Excellent 2_VeryGood      3_Good      4_Fair      5_Poor
Frequency      164          386          365          158          60
Proportion      0.145        0.341        0.322        0.139        0.053
-----
female
  n   missing  distinct      Info      Sum      Mean      Gmd
  1133       0       2       0.724      672     0.5931     0.4831
-----
```

- The **Info** measure is used for quantitative and binary variables. It is a relative information measure that increases towards 1 for variables with no ties, and is smaller for variables with many ties.
- The **Gmd** is the Gini mean difference. It is a measure of spread (or dispersion), where larger values indicate greater spread in the distribution, like the standard deviation or the interquartile range. It is defined as the mean absolute difference between any pairs of observations.

See the Help file for `describe` in the `Hmisc` package for more details on these measures, and on the settings for `describe`.

4.2 Counting as exploratory data analysis

Counting and/or tabulating things can be amazingly useful. Suppose we want to understand the `genhealth` values, after our single imputation.

```
smart_cle1_sh %>% count(genhealth) %>%
  mutate(percent = 100*n / sum(n))
```

```
# A tibble: 5 x 3
  genhealth      n percent
  <fct>     <int>   <dbl>
1 1_Excellent    164    14.5
2 2_VeryGood     386    34.1
3 3_Good         365    32.2
4 4_Fair          158    13.9
5 5_Poor          60     5.30
```

We might use `tabyl` to do this job...

```
smart_cle1_sh %>%
  tabyl(genhealth) %>%
  adorn_pct_formatting(digits = 1) %>%
  knitr::kable()
```

genhealth	n	percent
1_Excellent	164	14.5%
2_VeryGood	386	34.1%
3_Good	365	32.2%
4_Fair	158	13.9%
5_Poor	60	5.3%

4.2.1 Did `genhealth` vary by smoking status?

```
smart_cle1_sh %>%
  count(genhealth, smoke100) %>%
  mutate(percent = 100*n / sum(n))
```

```
# A tibble: 10 x 4
  genhealth  smoke100      n percent
  <fct>     <dbl> <int>   <dbl>
1 1_Excellent    0    105    9.27
2 1_Excellent    1     59    5.21
3 2_VeryGood     0    220   19.4
4 2_VeryGood     1    166   14.7
5 3_Good          0    184   16.2
6 3_Good          1    181   16.0
7 4_Fair          0     67    5.91
8 4_Fair          1     91    8.03
9 5_Poor          0     24    2.12
10 5_Poor         1     36    3.18
```

Suppose we want to find the percentage within each smoking status group. Here's one approach...

```
smart_cle1_sh %>%
  count(smoke100, genhealth) %>%
  group_by(smoke100) %>%
  mutate(prob = 100*n / sum(n))

# A tibble: 10 x 4
# Groups:   smoke100 [2]
  smoke100  genhealth      n    prob
  <dbl>     <fct>     <int>  <dbl>
1       0 1_Excellent    105 17.5
2       0 2_VeryGood     220 36.7
3       0 3_Good         184 30.7
4       0 4_Fair          67 11.2
5       0 5_Poor          24  4
6       1 1_Excellent     59 11.1
7       1 2_VeryGood     166 31.1
8       1 3_Good          181 34.0
9       1 4_Fair          91 17.1
10      1 5_Poor          36  6.75
```

And here's another ...

```
smart_cle1_sh %>%
  tabyl(smoke100, genhealth) %>%
  adorn_totals(where = c("row", "col")) %>%
  adorn_percentages(denominator = "row") %>%
  adorn_pct_formatting(digits = 1) %>%
  adorn_ns(position = "front")

smoke100 1_Excellent 2_VeryGood 3_Good 4_Fair 5_Poor
0 105 (17.5%) 220 (36.7%) 184 (30.7%) 67 (11.2%) 24 (4.0%)
1 59 (11.1%) 166 (31.1%) 181 (34.0%) 91 (17.1%) 36 (6.8%)
Total 164 (14.5%) 386 (34.1%) 365 (32.2%) 158 (13.9%) 60 (5.3%)
Total
600 (100.0%)
533 (100.0%)
1133 (100.0%)
```

4.2.2 What's the distribution of physhealth?

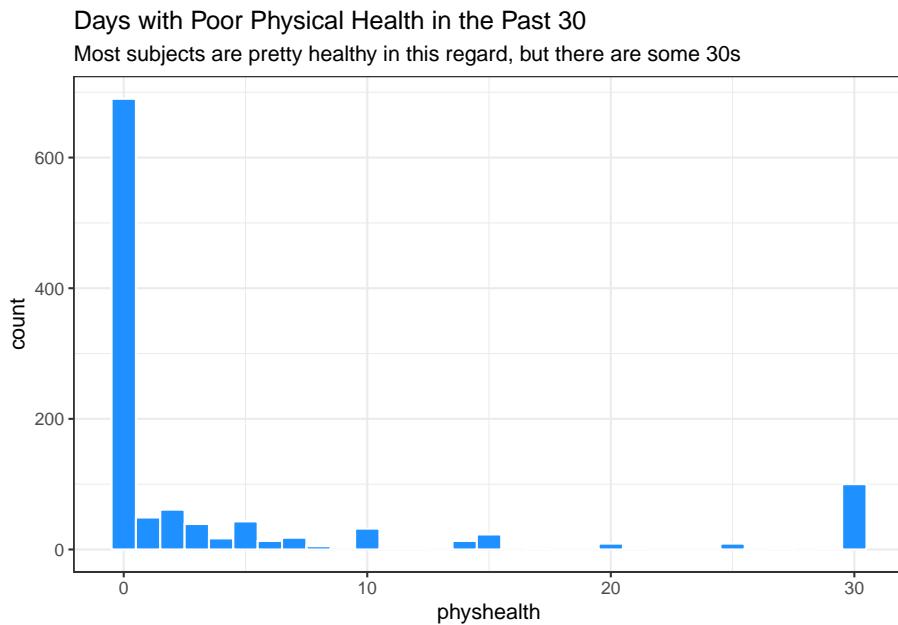
We can count quantitative variables with discrete sets of possible values, like `physhealth`, which is captured as an integer (that must fall between 0 and 30.)

```
smart_cle1_sh %>% count(physhealth)
```

```
# A tibble: 21 x 2
  physhealth     n
  <dbl> <int>
1      0    690
2      1     49
3      2     61
4      3     39
5      4     17
6      5     43
7      6     13
8      7     18
9      8      5
10     10    32
# ... with 11 more rows
```

Of course, a natural summary of a quantitative variable like this would be graphical.

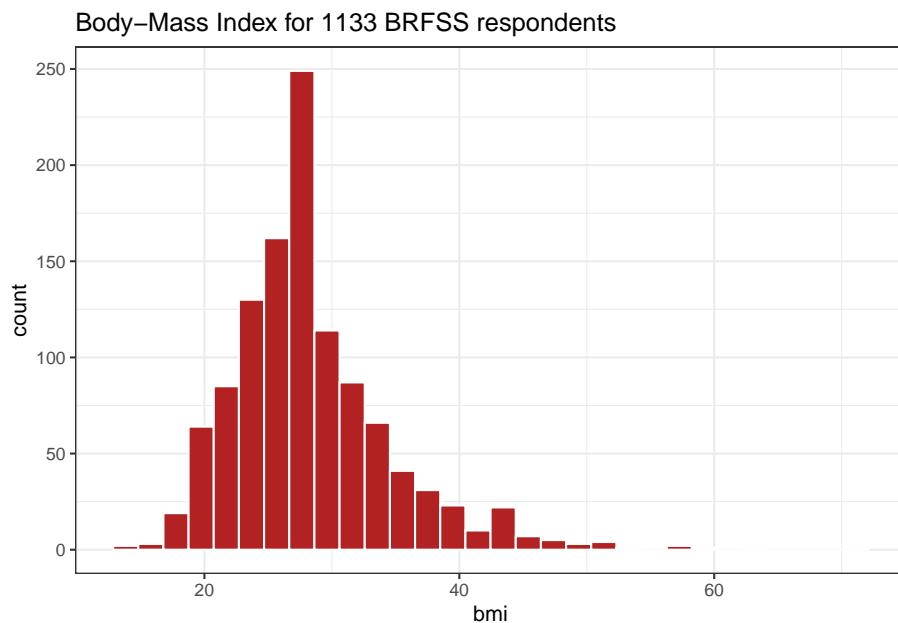
```
ggplot(smart_cle1_sh, aes(physhealth)) +
  geom_histogram(binwidth = 1,
                 fill = "dodgerblue", col = "white") +
  labs(title = "Days with Poor Physical Health in the Past 30",
       subtitle = "Most subjects are pretty healthy in this regard, but there are some outliers")
```



4.2.3 What's the distribution of `bmi`?

`bmi` is the body-mass index, an indicator of size (thickness, really.)

```
ggplot(smart_cle1_sh, aes(bmi)) +
  geom_histogram(bins = 30,
                 fill = "firebrick", col = "white") +
  labs(title = paste0("Body-Mass Index for ",
                      nrow(smart_cle1_sh),
                      " BRFSS respondents"))
```



4.2.4 How many of the respondents have a BMI below 30?

```
smart_cle1_sh %>% count(bmi < 30) %>%
  mutate(proportion = n / sum(n))
```

```
# A tibble: 2 x 3
`bmi < 30`    n proportion
<lgl>      <int>     <dbl>
1 FALSE        330     0.291
2 TRUE         803     0.709
```

4.2.5 How many of the respondents with a BMI < 30 are highly active?

```
smart_cle1_sh %>%
  filter(bmi < 30) %>%
  tabyl(activity) %>%
  adorn_pct_formatting()

      activity    n percent
Highly_Active 343   42.7%
        Active 133   16.6%
Insufficiently_Active 129   16.1%
      Inactive 198   24.7%
```

4.2.6 Is obesity associated with smoking history?

```
smart_cle1_sh %>% count(smoke100, bmi < 30) %>%
  group_by(smoke100) %>%
  mutate(percent = 100*n/sum(n))

# A tibble: 4 x 4
# Groups:   smoke100 [2]
  smoke100 `bmi < 30`    n percent
  <dbl> <lgl>     <int>   <dbl>
1 0 FALSE       163    27.2
2 0 TRUE        437    72.8
3 1 FALSE       167    31.3
4 1 TRUE        366    68.7
```

4.2.7 Comparing drinks_wk summaries by obesity status

Can we compare the `drinks_wk` means, medians and 75th percentiles for respondents whose BMI is below 30 to the respondents whose BMI is not?

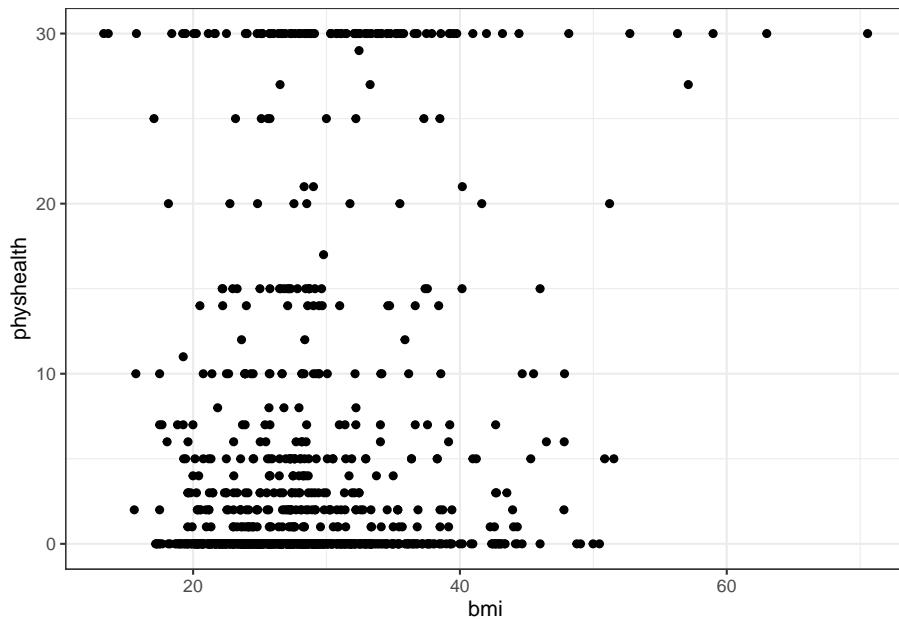
```
smart_cle1_sh %>%
  group_by(bmi < 30) %>%
  summarize(mean(drinks_wk), median(drinks_wk),
            q75 = quantile(drinks_wk, 0.75))

# A tibble: 2 x 4
`bmi < 30` `mean(drinks_wk)` `median(drinks_wk)` `q75
<lgl>           <dbl>           <dbl> <dbl>
1 FALSE          1.67          0.23  1.17
2 TRUE           2.80          0.23  2.8
```

4.3 Can bmi predict physhealth?

We'll start with an effort to predict physhealth using bmi. A natural graph would be a scatterplot.

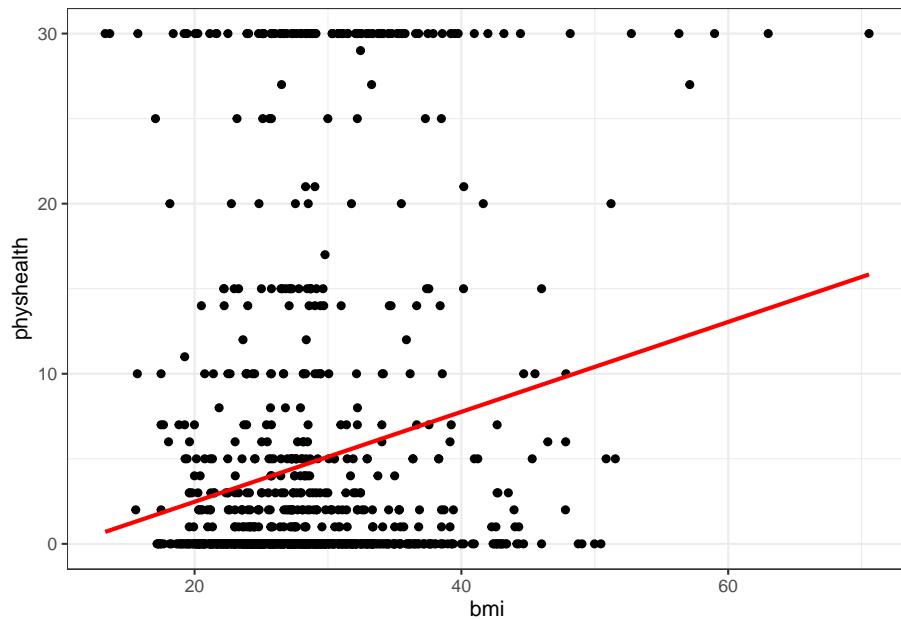
```
ggplot(data = smart_cle1_sh, aes(x = bmi, y = physhealth)) +
  geom_point()
```



A good question to ask ourselves here might be: “In what BMI range can we make a reasonable prediction of physhealth?”

Now, we might take the plot above and add a simple linear model ...

```
ggplot(data = smart_cle1_sh, aes(x = bmi, y = physhealth)) +
  geom_point() +
  geom_smooth(method = "lm", se = FALSE, col = "red")  
  
`geom_smooth()` using formula 'y ~ x'
```



which shows the same least squares regression model that we can fit with the `lm` command.

4.3.1 Fitting a Simple Regression Model

```
model_A <- lm(physhealth ~ bmi, data = smart_cle1_sh)
```

```
model_A
```

```
Call:
```

```
lm(formula = physhealth ~ bmi, data = smart_cle1_sh)
```

```
Coefficients:
```

(Intercept)	bmi
-2.8121	0.2643

```
summary(model_A)
```

```
Call:
```

```
lm(formula = physhealth ~ bmi, data = smart_cle1_sh)
```

```
Residuals:
```

Min	1Q	Median	3Q	Max
-10.5258	-4.5943	-3.5608	-0.5106	29.2965

```
Coefficients:
```

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	-2.81208	1.21672	-2.311	0.021 *
bmi	0.26433	0.04188	6.312	3.95e-10 ***

Signif. codes:	0 ***	0.001 **	0.01 *	0.05 .
	'	'	'	'
	'	'	'	'

Residual standard error: 8.968 on 1131 degrees of freedom
 Multiple R-squared: 0.03403, Adjusted R-squared: 0.03317
 F-statistic: 39.84 on 1 and 1131 DF, p-value: 3.95e-10

```
confint(model_A, level = 0.95)
```

	2.5 %	97.5 %
(Intercept)	-5.1993624	-0.4247909
bmi	0.1821599	0.3464915

The model coefficients can be obtained by printing the model object, and the `summary` function provides several useful descriptions of the model's residuals, its statistical significance, and quality of fit.

4.3.2 Model Summary for a Simple (One-Predictor) Regression

The fitted model predicts `physhealth` using a prediction equation we can read off from the model coefficient estimates. Specifically, we have:

```
coef(model_A)
```

	bmi
(Intercept)	-2.8120766
bmi	0.2643257

so the equation is `physhealth` = -2.82 + 0.265 `bmi`.

Each of the 1133 respondents included in the `smart_cle1_sh` data makes a contribution to this model.

4.3.2.1 Residuals

Suppose Harry is one of the people in that group, and Harry's data is `bmi` = 20, and `physhealth` = 3.

- Harry's *observed* value of `physhealth` is just the value we have in the data for them, in this case, observed `physhealth` = 3 for Harry.
- Harry's *fitted* or *predicted* `physhealth` value is the result of calculating $-2.82 + 0.265 \text{ bmi}$ for Harry. So, if Harry's BMI was 20, then Harry's predicted `physhealth` value is $-2.82 + 0.265 (20) = 2.48$.

- The *residual* for Harry is then his *observed* outcome minus his *fitted* outcome, so Harry has a residual of $3 - 2.48 = 0.52$.
- Graphically, a residual represents vertical distance between the observed point and the fitted regression line.
- Points above the regression line will have positive residuals, and points below the regression line will have negative residuals. Points on the line have zero residuals.

The residuals are summarized at the top of the `summary` output for linear model.

```
summary(model_A)
```

```
Call:
lm(formula = physhealth ~ bmi, data = smart_cle1_sh)

Residuals:
    Min      1Q  Median      3Q     Max 
-10.5258 -4.5943 -3.5608 -0.5106 29.2965 

Coefficients:
            Estimate Std. Error t value Pr(>|t|)    
(Intercept) -2.81208   1.21672  -2.311   0.021 *  
bmi         0.26433   0.04188   6.312 3.95e-10 *** 
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 8.968 on 1131 degrees of freedom
Multiple R-squared:  0.03403, Adjusted R-squared:  0.03317 
F-statistic: 39.84 on 1 and 1131 DF,  p-value: 3.95e-10
```

- The mean residual will always be zero in an ordinary least squares model, but a five number summary of the residuals is provided by the summary, as is an estimated standard deviation of the residuals (called here the Residual standard error.)
- In the `smart_cle1_sh` data, the minimum residual was -10.53, so for one subject, the observed value was 10.53 days smaller than the predicted value. This means that the prediction was 10.53 days too large for that subject.
- Similarly, the maximum residual was 29.30 days, so for one subject the prediction was 29.30 days too small. Not a strong performance.
- In a least squares model, the residuals are assumed to follow a Normal distribution, with mean zero, and standard deviation (for the `smart_cle1_sh` data) of about 9.0 days. We know this because the residual standard error is specified as 8.968 later in the linear model output. Thus, by the definition of a Normal distribution, we'd expect
 - about 68% of the residuals to be between -9 and +9 days,
 - about 95% of the residuals to be between -18 and +18 days,
 - about all (99.7%) of the residuals to be between -27 and +27 days.

4.3.2.2 Coefficients section

The `summary` for a linear model shows Estimates, Standard Errors, t values and *p* values for each coefficient fit.

```
summary(model_A)
```

Call:

```
lm(formula = physhealth ~ bmi, data = smart_cle1_sh)
```

Residuals:

Min	1Q	Median	3Q	Max
-10.5258	-4.5943	-3.5608	-0.5106	29.2965

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	-2.81208	1.21672	-2.311	0.021 *
bmi	0.26433	0.04188	6.312	3.95e-10 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 8.968 on 1131 degrees of freedom

Multiple R-squared: 0.03403, Adjusted R-squared: 0.03317

F-statistic: 39.84 on 1 and 1131 DF, p-value: 3.95e-10

- The Estimates are the point estimates of the intercept and slope of `bmi` in our model.
- In this case, our estimated slope is 0.265, which implies that if Harry's BMI is 20 and Sally's BMI is 21, we predict that Sally's `physhealth` will be 0.265 days larger than Harry's.
- The Standard Errors are also provided for each estimate. We can create rough 95% uncertainty intervals for these estimated coefficients by adding and subtracting two standard errors from each coefficient, or we can get a slightly more accurate answer with the `confint` function.
- Here, the 95% uncertainty interval for the slope of `bmi` is estimated to be (0.18, 0.35). This is a good measure of the uncertainty in the slope that is captured by our model. We are 95% confident in the process of building this interval, but this doesn't mean we're 95% sure that the true slope is actually in that interval.

Also available are a *t* value (just the Estimate divided by the Standard Error) and the appropriate *p* value for testing the null hypothesis that the true value of the coefficient is 0 against a two-tailed alternative.

- If a slope coefficient is statistically detectably different from 0, this implies that 0 will not be part of the uncertainty interval obtained through `confint`.
- If the slope was zero, it would suggest that `bmi` would add no predictive value to the model. But that's unlikely here.

If the `bmi` slope coefficient is associated with a small p value, as in the case of our `model_A`, it suggests that the model including `bmi` is statistically detectably better at predicting `physhealth` than the model without `bmi`.

- Without `bmi` our `model_A` would become an *intercept-only* model, in this case, which would predict the mean `physhealth` for everyone, regardless of any other information.

4.3.2.3 Model Fit Summaries

```
summary(model_A)

Call:
lm(formula = physhealth ~ bmi, data = smart_cle1_sh)

Residuals:
    Min      1Q  Median      3Q     Max 
-10.5258 -4.5943 -3.5608 -0.5106 29.2965 

Coefficients:
            Estimate Std. Error t value Pr(>|t|)    
(Intercept) -2.81208   1.21672  -2.311   0.021 *  
bmi         0.26433   0.04188   6.312 3.95e-10 *** 
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 8.968 on 1131 degrees of freedom
Multiple R-squared:  0.03403, Adjusted R-squared:  0.03317 
F-statistic: 39.84 on 1 and 1131 DF,  p-value: 3.95e-10
```

The `summary` of a linear model also displays:

- The residual standard error and associated degrees of freedom for the residuals.
- For a simple (one-predictor) least regression like this, the residual degrees of freedom will be the sample size minus 2.
- The multiple R-squared (or coefficient of determination)
- This is interpreted as the proportion of variation in the outcome (`physhealth`) accounted for by the model, and will always fall between 0 and 1 as a result.
- Our model_A accounts for a mere 3.4% of the variation in `physhealth`.
- The Adjusted R-squared value “adjusts” for the size of our model in terms of the number of coefficients included in the model.
- The adjusted R-squared will always be smaller than the Multiple R-squared.
- We still hope to find models with relatively large adjusted R^2 values.

- In particular, we hope to find models where the adjusted R^2 isn't substantially less than the Multiple R-squared.
- The adjusted R-squared is usually a better estimate of likely performance of our model in new data than is the Multiple R-squared.
- The adjusted R-squared result is no longer interpretable as a proportion of anything - in fact, it can fall below 0.
- We can obtain the adjusted R^2 from the raw R^2 , the number of observations N and the number of predictors p included in the model, as follows:

$$R_{adj}^2 = 1 - \frac{(1 - R^2)(N - 1)}{N - p - 1},$$

- The F statistic and p value from a global ANOVA test of the model.
 - Obtaining a statistically significant result here is usually pretty straightforward, since the comparison is between our model, and a model which simply predicts the mean value of the outcome for everyone.
 - In a simple (one-predictor) linear regression like this, the t statistic for the slope is just the square root of the F statistic, and the resulting p values for the slope's t test and for the global F test will be identical.
- To see the complete ANOVA F test for this model, we can run `anova(model_A)`.

```
anova(model_A)
```

Analysis of Variance Table

```
Response: physhealth
          Df Sum Sq Mean Sq F value    Pr(>F)
bmi         1   3204   3204.4  39.84 3.95e-10 ***
Residuals 1131  90968     80.4
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

4.3.3 Using the `broom` package

The `broom` package has three functions of particular use in a linear regression model:

4.3.3.1 The `tidy` function

`tidy` builds a data frame/tibble containing information about the coefficients in the model, their standard errors, t statistics and p values.

```
tidy(model_A)

# A tibble: 2 x 5
  term      estimate std.error statistic p.value
  <chr>     <dbl>     <dbl>     <dbl>    <dbl>
1 (Intercept) -2.81      1.22     -2.31 2.10e- 2
2 bmi         0.264     0.0419     6.31 3.95e-10

It's often useful to include other summaries in this tidying, for instance:

tidy(model_A, conf.int = TRUE, conf.level = 0.9) %>%
  select(term, estimate, conf.low, conf.high)

# A tibble: 2 x 4
  term      estimate conf.low conf.high
  <chr>     <dbl>     <dbl>     <dbl>
1 (Intercept) -2.81     -4.82     -0.809
2 bmi         0.264      0.195     0.333
```

4.3.3.2 The `glance` function

`glance`¹ builds a data frame/tibble containing summary statistics about the model, including

- the (raw) multiple R^2 and adjusted R^2
- `sigma` which is the residual standard error
- the `F statistic`, `p.value` model `df` and `df.residual` associated with the global ANOVA test, plus
- several statistics that will be useful in comparing models down the line:
 - the model's log likelihood function value, `logLik`
 - the model's Akaike's Information Criterion value, `AIC`
 - the model's Bayesian Information Criterion value, `BIC`
 - and the model's deviance statistic

```
glance(model_A)

# A tibble: 1 x 12
  r.squared adj.r.squared sigma statistic p.value    df logLik    AIC    BIC
  <dbl>        <dbl>     <dbl>     <dbl>    <dbl> <dbl> <dbl> <dbl> <dbl>
1 0.0340       0.0332    8.97     39.8 3.95e-10     1 -4092. 8190. 8205.
# ... with 3 more variables: deviance <dbl>, df.residual <int>, nobs <int>
```

4.3.3.3 The `augment` function

`augment` builds a data frame/tibble which adds fitted values, residuals and other diagnostic summaries that describe each observation to the original data used to

fit the model, and this includes

- `.fitted` and `.resid`, the fitted and residual values, in addition to
- `.hat`, the leverage value for this observation
- `.cooksdi`, the Cook's distance measure of *influence* for this observation
- `.stdresid`, the standardized residual (think of this as a z-score - a measure of the residual divided by its associated standard deviation `.sigma`)
- and `se.fit` which will help us generate prediction intervals for the model downstream

Note that each of the new columns begins with `.` to avoid overwriting any data.

```
head(augment(model_A))
```

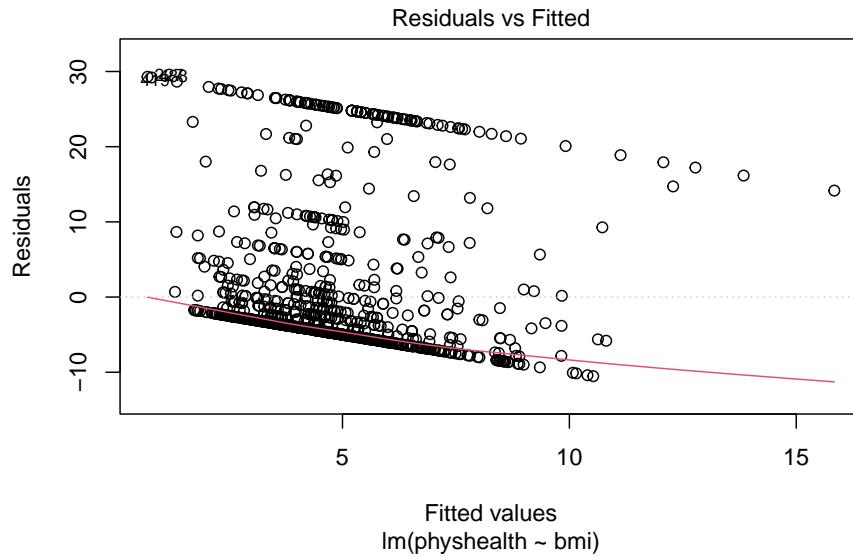
```
# A tibble: 6 x 8
  physhealth bmi .fitted .resid     .hat .sigma     .cooksdi .std.resid
  <dbl> <dbl>   <dbl>  <dbl>    <dbl>  <dbl>      <dbl>       <dbl>
1      4 27.9    4.57 -0.572 0.000886  8.97 0.00000181 -0.0638
2      0 23.0    3.28 -3.28 0.00149   8.97 0.000100  -0.366 
3      0 26.9    4.31 -4.31 0.000927  8.97 0.000107  -0.480 
4      0 26.5    4.20 -4.20 0.000956  8.97 0.000105  -0.468 
5      0 24.2    3.60 -3.60 0.00125   8.97 0.000101  -0.401 
6      2 27.7    4.51 -2.51 0.000891  8.97 0.0000351 -0.281
```

For more on the `broom` package, you may want to look at this vignette.

4.3.4 How does the model do? (Residuals vs. Fitted Values)

- Remember that the R^2 value was about 3.4%.

```
plot(model_A, which = 1)
```

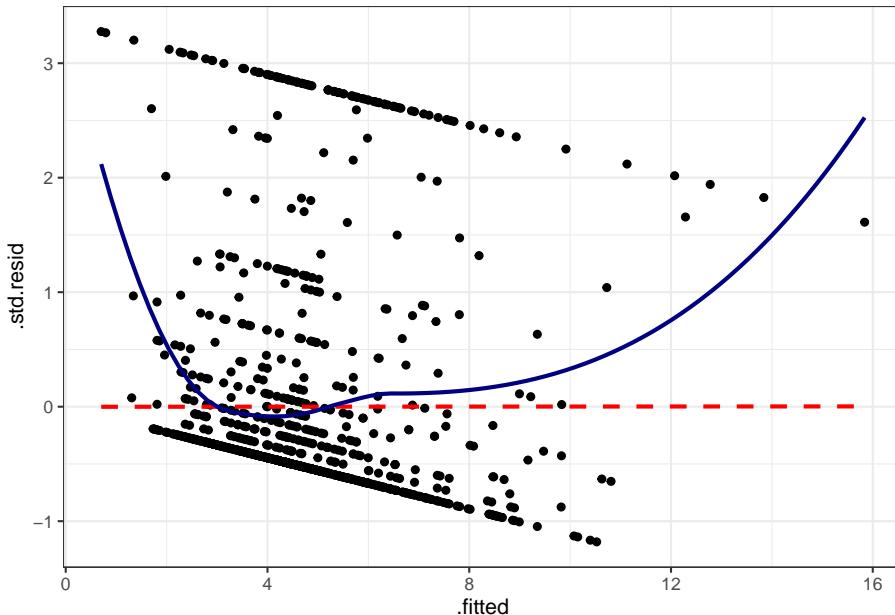


This is a plot of residuals vs. fitted values. The goal here is for this plot to look like a random scatter of points, perhaps like a “fuzzy football,” and that’s **not** what we have. Why?

If you prefer, here’s a `ggplot2` version of a similar plot, now looking at standardized residuals instead of raw residuals, and adding a loess smooth and a linear fit to the result.

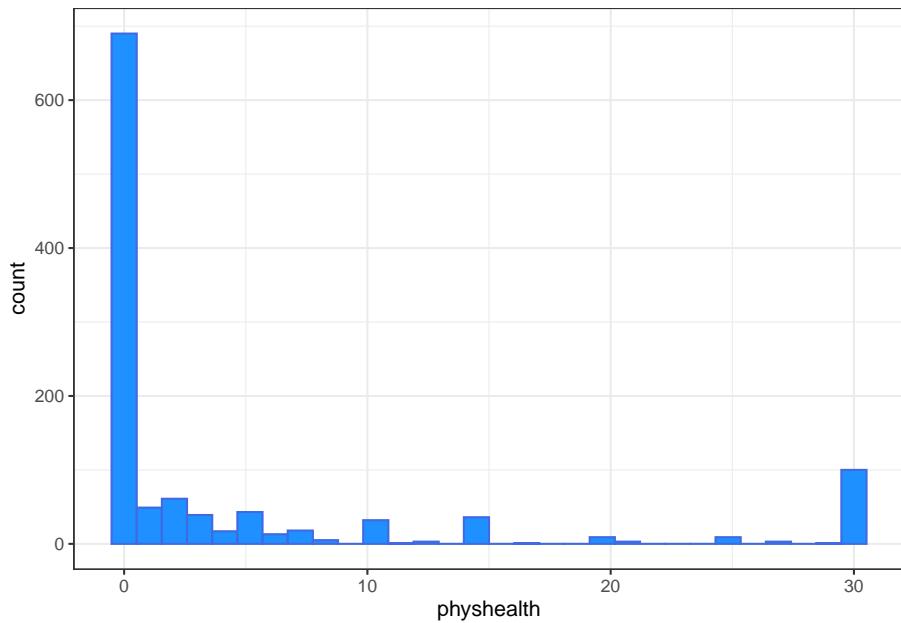
```
ggplot(augment(model_A), aes(x = .fitted, y = .std.resid)) +
  geom_point() +
  geom_smooth(method = "lm", se = FALSE, col = "red", linetype = "dashed") +
  geom_smooth(method = "loess", se = FALSE, col = "navy") +
  theme_bw()

`geom_smooth()` using formula 'y ~ x'
`geom_smooth()` using formula 'y ~ x'
```



The problem we're having here becomes, I think, a little more obvious if we look at what we're predicting. Does `physhealth` look like a good candidate for a linear model?

```
ggplot(smart_cle1_sh, aes(x = physhealth)) +  
  geom_histogram(bins = 30, fill = "dodgerblue",  
                 color = "royalblue")
```



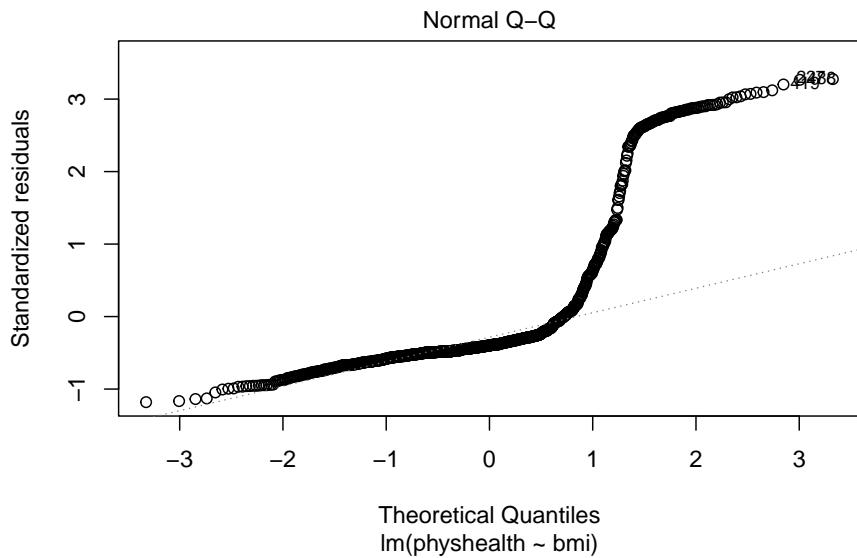
```
smart_cle1_sh %>% count(physhealth == 0, physhealth == 30)
```

physhealth == 0	physhealth == 30	n
1 FALSE	FALSE	343
2 FALSE	TRUE	100
3 TRUE	FALSE	690

No matter what model we fit, if we are predicting `physhealth`, and most of the data are values of 0 and 30, we have limited variation in our outcome, and so our linear model will be somewhat questionable just on that basis.

A normal Q-Q plot of the standardized residuals for our `model_A` shows this problem, too.

```
plot(model_A, which = 2)
```



We're going to need a method to deal with this sort of outcome, that has both a floor and a ceiling. We'll get there eventually, but linear regression alone doesn't look promising.

All right, so that didn't go anywhere great. We'll try again, with a new outcome, in the next chapter.

Chapter 5

Analysis of Variance with SMART

In this chapter, we'll work with the `smart_cle1_sh` data file again.

```
smart_cle1_sh <- readRDS(here("data", "smart_cle1_sh.Rds"))
```

The variables we'll look at in this chapter are as follows.

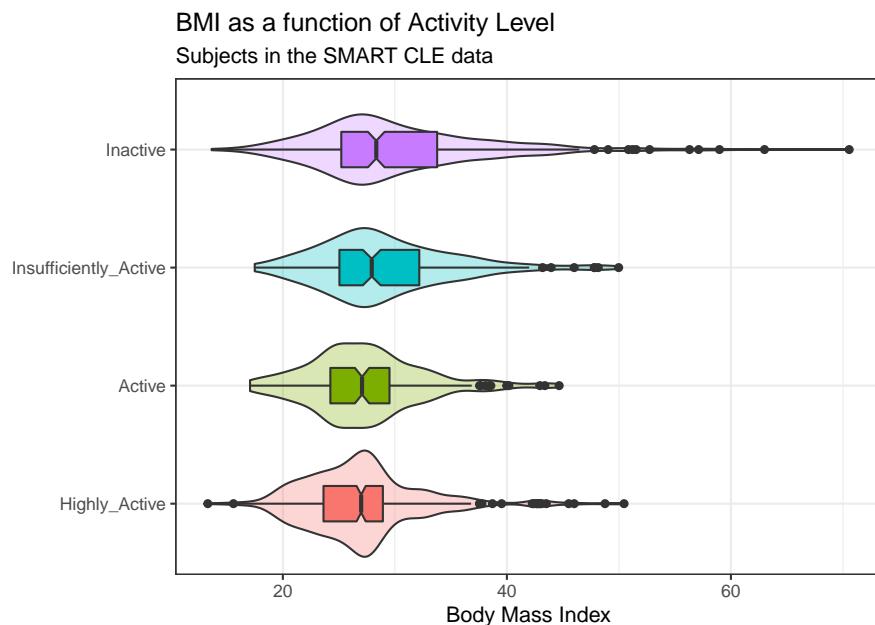
Variable	Description
<code>SEQNO</code>	respondent identification number (all begin with 2016)
<code>bmi</code>	Body mass index, in kg/m ²
<code>female</code>	Sex, 1 = female, 0 = male
<code>smoke100</code>	Have you smoked at least 100 cigarettes in your life? (1 = yes, 0 = no)
<code>activity</code>	Physical activity (Highly Active, Active, Insufficiently Active, Inactive)
<code>drinks_wk</code>	On average, how many drinks of alcohol do you consume in a week?
<code>physhealth</code>	Now thinking about your physical health, which includes physical illness and injury, for how many days during the past 30 days was your physical health not good?

5.1 A One-Factor Analysis of Variance

We'll be predicting body mass index, at first using a single factor as a predictor: the `activity` level.

5.1.1 Can activity be used to predict bmi?

```
ggplot(smart_cle1_sh, aes(x = activity, y = bmi,
                           fill = activity)) +
  geom_violin(alpha = 0.3) +
  geom_boxplot(width = 0.3, notch = TRUE) +
  guides(fill = "none") +
  coord_flip() +
  labs(title = "BMI as a function of Activity Level",
       subtitle = "Subjects in the SMART CLE data",
       x = "", y = "Body Mass Index")
```



Here's a numerical summary of the distributions of `bmi` within each `activity` group.

```
smart_cle1_sh %$% mosaic::favstats(bmi ~ activity)
```

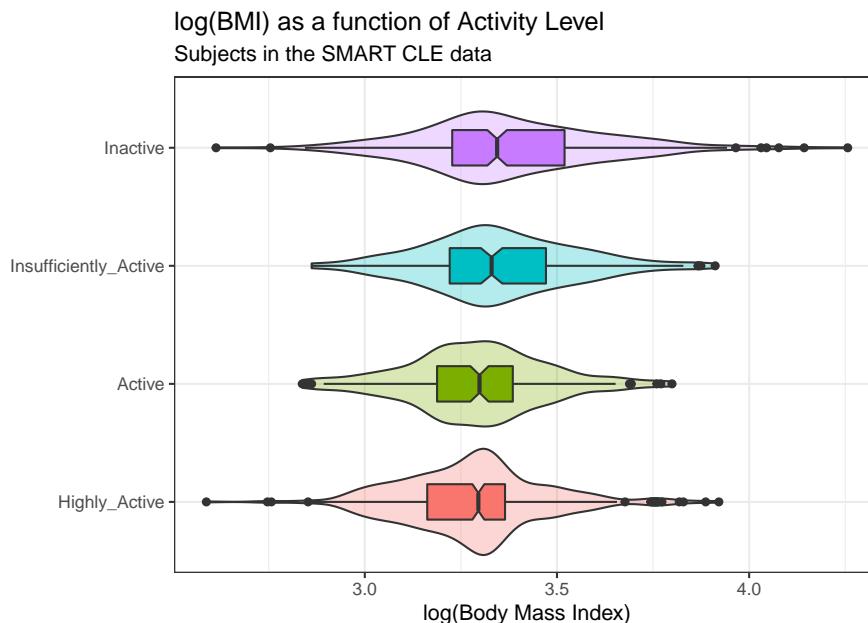
	activity	min	Q1	median	Q3	max	mean	sd
1	Highly_Active	13.30	23.6275	26.9900	28.930	50.46	27.02253	5.217496
2	Active	17.07	24.2400	27.06930	29.520	44.67	27.36157	5.151796
3	Insufficiently_Active	17.49	25.0500	27.93776	32.180	49.98	29.04328	6.051823
4	Inactive	13.64	25.2150	28.34000	33.775	70.56	30.15978	7.832675
	n missing							
1	428	0						
2	173	0						

```
3 201      0
4 331      0
```

5.1.2 Should we transform `bmi`?

The analysis of variance is something of a misnomer. What we're doing is using the variance to say something about population means. In light of the apparent right skew of the `bmi` results in each `activity` group, might it be a better choice to use a logarithmic transformation? We'll use the natural logarithm here, which in R, is symbolized by `log`.

```
ggplot(smart_cle1_sh, aes(x = activity, y = log(bmi),
                           fill = activity)) +
  geom_violin(alpha = 0.3) +
  geom_boxplot(width = 0.3, notch = TRUE) +
  guides(fill = "none") +
  coord_flip() +
  labs(title = "log(BMI) as a function of Activity Level",
       subtitle = "Subjects in the SMART CLE data",
       x = "", y = "log(Body Mass Index)")
```



The logarithmic transformation yields distributions that look much more symmetric in each `activity` group, so we'll proceed to build our regression model predicting `log(bmi)` using `activity`. Here's the numerical summary of these logged results:

```
smart_cle1_sh %$% mosaic::favstats(log(bmi) ~ activity)

      activity      min      Q1   median      Q3      max     mean
1 Highly_Active 2.587764 3.162411 3.295466 3.364879 3.921181 3.279246
2           Active 2.837323 3.188004 3.298400 3.385068 3.799302 3.292032
3 Insufficiently_Active 2.861629 3.220874 3.329979 3.471345 3.911623 3.348383
4           Inactive 2.613007 3.227439 3.344274 3.519721 4.256463 3.376468
      sd    n missing
1 0.1851478 428      0
2 0.1850568 173      0
3 0.2007241 201      0
4 0.2411196 331      0
```

5.1.3 Building the ANOVA model

```
model_5a <- smart_cle1_sh %$% lm(log(bmi) ~ activity)
```

```
model_5a
```

Call:
`lm(formula = log(bmi) ~ activity)`

Coefficients:

	(Intercept)	activityActive
	3.27925	0.01279
activityInsufficiently_Active	0.06914	activityInactive
		0.09722

The `activity` data is categorical and there are four levels. The model equation is:

$$\begin{aligned} \log(\text{bmi}) = & 3.279 + 0.013 (\text{activity} = \text{Active}) \\ & + 0.069 (\text{activity} = \text{Insufficiently Active}) \\ & + 0.097 (\text{activity} = \text{Inactive}) \end{aligned}$$

where, for example, (`activity = Active`) is 1 if `activity` is Active, and 0 otherwise. The fourth level (Highly Active) is not shown here and is used as a baseline. Thus the model above can be interpreted as follows.

activity	Predicted log(bmi)	Predicted bmi
Highly Active	3.279	$\exp(3.279) = 26.55$
Active	$3.279 + 0.013 = 3.292$	$\exp(3.292) = 26.90$
Insufficiently Active	$3.279 + 0.069 = 3.348$	$\exp(3.348) = 28.45$
Inactive	$3.279 + 0.097 = 3.376$	$\exp(3.376) = 29.25$

Those predicted `log(bmi)` values should look familiar. They are just the means of `log(bmi)` in each group, but I'm sure you'll also notice that the predicted `bmi` values are not exact matches for the observed means of `bmi`.

```
smart_cle1_sh %>% group_by(activity) %>%
  summarise(mean(log(bmi)), mean(bmi))
```

	<code>activity</code>	<code>mean(log(bmi))</code>	<code>mean(bmi)</code>
	<code><fct></code>	<code><dbl></code>	<code><dbl></code>
1	Highly_Active	3.28	27.0
2	Active	3.29	27.4
3	Insufficiently_Active	3.35	29.0
4	Inactive	3.38	30.2

5.1.4 The ANOVA table

Now, let's press on to look at the ANOVA results for this model.

```
anova(model_5a)
```

Analysis of Variance Table

```
Response: log(bmi)
          Df Sum Sq Mean Sq F value    Pr(>F)
activity     3  2.060  0.68652  16.225 2.496e-10 ***
Residuals 1129 47.772  0.04231
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

- The total variation in `log(bmi)`, our outcome, is captured by the sums of squares here. $SS(\text{Total}) = 2.058 + 47.770 = 49.828$
- Here, the `activity` variable (with 4 levels, so $4-1 = 3$ degrees of freedom) accounts for 4.13% ($2.058 / 49.828$) of the variation in `log(bmi)`. Another way of saying this is that the model R^2 or η^2 is 0.0413.
- The variation accounted for by the `activity` categories meets the standard for a statistically detectable result, according to the ANOVA F test, although that's not really important.
- The square root of the Mean Square(Residuals) is the residual standard error, σ , we've seen in the past. $MS(\text{Residual})$ estimates the variance (0.0423), so the residual standard error is $\sqrt{0.0423} \approx 0.206$.

5.1.5 The Model Coefficients

To address the question of effect size for the various levels of `activity` on `log(bmi)`, we could look directly at the regression model coefficients. For that,

we might look at the model `summary`.

```
summary(model_5a)

Call:
lm(formula = log(bmi) ~ activity)

Residuals:
    Min      1Q  Median      3Q     Max 
-0.76346 -0.12609 -0.00286  0.11055  0.88000 

Coefficients:
            Estimate Std. Error t value Pr(>|t|)    
(Intercept) 3.279246  0.009943 329.806 < 2e-16 ***
activityActive 0.012785  0.018532   0.690    0.49    
activityInsufficiently_Active 0.069137  0.017589   3.931 8.99e-05 ***
activityInactive 0.097221  0.015056   6.457 1.58e-10 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.2057 on 1129 degrees of freedom
Multiple R-squared:  0.04133, Adjusted R-squared:  0.03878 
F-statistic: 16.22 on 3 and 1129 DF,  p-value: 2.496e-10
```

If we want to see the confidence intervals around these estimates, we could use

```
confint(model_5a, conf.level = 0.95)
```

	2.5 %	97.5 %
(Intercept)	3.25973769	3.29875522
activityActive	-0.02357630	0.04914707
activityInsufficiently_Active	0.03462572	0.10364764
activityInactive	0.06767944	0.12676300

The model suggests, based on these 1133 subjects, that (remember that the baseline category is Highly Active)

- a 95% confidence (uncertainty) interval for the difference between Active and Highly Active subjects in log(BMI) ranges from -0.024 to 0.049
- a 95% confidence (uncertainty) interval for the difference between Insufficiently Active and Highly Active subjects in log(BMI) ranges from 0.035 to 0.104
- a 95% confidence (uncertainty) interval for the difference between Inactive and Highly Active subjects in log(BMI) ranges from 0.068 to 0.127
- the model accounts for 4.13% of the variation in log(BMI), so that knowing the respondent's activity level somewhat reduces the size of the prediction errors as compared to an intercept only model that would predict the overall mean log(BMI), regardless of activity level, for all subjects.

- from the summary of residuals, we see that one subject had a residual of 0.88 - that means they were predicted to have a log(BMI) 0.88 lower than their actual log(BMI) and one subject had a log(BMI) that is 0.76 larger than their actual log(BMI), at the extremes.

5.1.6 Using `broom::tidy` to explore the coefficients

A better strategy for displaying the coefficients in any regression model is to use the `tidy` function from the `broom` package.

```
tidy(model_5a, conf.int = TRUE, conf.level = 0.95) %>%
  knitr::kable(digits = 3)
```

term	estimate	std.error	statistic	p.value	conf.low	conf.high
(Intercept)	3.279	0.010	329.806	0.00	3.260	3.299
activityActive	0.013	0.019	0.690	0.49	-0.024	0.049
activityInsufficiently_Active	0.069	0.018	3.931	0.00	0.035	0.104
activityInactive	0.097	0.015	6.457	0.00	0.068	0.127

5.1.7 Using `broom::glance` to summarize the model's fit

```
glance(model_5a) %>% select(1:3) %>%
  knitr::kable(digits = c(4, 4, 3))
```

r.squared	adj.r.squared	sigma
0.0413	0.0388	0.206

- The `r.squared` or R^2 value is interpreted for a linear model as the percentage of variation in the outcome (here, `log(bmi)`) that is accounted for by the model.
- The `adj.r.squared` or adjusted R^2 value incorporates a small penalty for the number of predictors included in the model. Adjusted R^2 is useful for models with more than one predictor, not simple regression models like this one. Like R^2 and most of these other summaries, its primary value comes when making comparisons between models for the same outcome.
- The `sigma` or σ is the residual standard error. Doubling this value gives us a good idea of the range of errors made by the model (approximately 95% of the time if the normal distribution assumption for the residuals holds perfectly.)

```
glance(model_5a) %>% select(4:7) %>%
  knitr::kable(digits = c(2, 3, 0, 2))
```

statistic	p.value	df	logLik
16.22	0	3	185.99

- The `statistic` and `p.value` shown here refer to the ANOVA F test and p value. They test the null hypothesis that the `activity` information is of no use in separating out the `bmi` data, or, equivalently, that the true R^2 is 0.
- The `df` indicates the model degrees of freedom, and in this case simply specifies the number of parameters fitted attributed to the model. Models that require more `df` for estimation require larger sample sizes.
- The `logLik` is the log likelihood for the model. This is a function of the sample size, but we can compare the fit of multiple models by comparing this value across different models for the same outcome. You want to maximize the log-likelihood.

```
glance(model_5a) %>% select(8:9) %>%
  knitr::kable(digits = 2)
```

AIC	BIC
-361.98	-336.82

- The `AIC` (or Akaike information criterion) and `BIC` (Bayes information criterion) are also used only to compare models. You want to minimize AIC and BIC in selecting a model. AIC and BIC are unique only up to a constant, so different packages or routines in R may give differing values, but in comparing two models - the difference in AIC (or BIC) should be consistent.

5.1.8 Using `broom::augment` to make predictions

We can obtain residuals and predicted (fitted) values for the points used to fit the model with `augment` from the `broom` package.

```
augment(model_5a, se_fit = TRUE) %>%
  select(1:5) %>% slice(1:4) %>%
  knitr::kable(digits = 3)
```

log(bmi)	activity	.fitted	.se.fit	.resid
3.330	Inactive	3.376	0.011	-0.047
3.138	Inactive	3.376	0.011	-0.239
3.293	Insufficiently_Active	3.348	0.015	-0.055
3.278	Highly_Active	3.279	0.010	-0.002

- The `.fitted` value is the predicted value of `log(bmi)` for this subject.
- The `.se.fit` value shows the standard error associated with the fitted value.
- The `.resid` is the residual value (observed - fitted `log(bmi)`)

```
augment(model_5a, se_fit = TRUE) %>%
  select(1:2, 6:9) %>% slice(1:4) %>%
  knitr::kable(digits = 3)
```

log(bmi)	activity	.hat	.sigma	.cooksdi	.std.resid
3.330	Inactive	0.003	0.206	0.000	-0.227
3.138	Inactive	0.003	0.206	0.001	-1.163
3.293	Insufficiently_Active	0.005	0.206	0.000	-0.269
3.278	Highly_Active	0.002	0.206	0.000	-0.008

- The `.hat` value shows the leverage index associated with the observation (this is a function of the predictors - higher leveraged points have more unusual predictor values)
- The `.sigma` value shows the estimate of the residual standard deviation if this observation were to be dropped from the model, and thus indexes how much of an outlier this observation's residual is.
- The `.cooksdi` or Cook's distance value shows the influence that the observation has on the model - it is one of a class of leave-one-out diagnostic measures. Larger values of Cook's distance indicate more influential points.
- The `.std.resid` shows the standardized residual (which is designed to have mean 0 and standard deviation 1, facilitating comparisons across models for differing outcomes)

5.2 A Two-Factor ANOVA (without Interaction)

Let's add `race_eth` to the predictor set for `log(BMI)`.

```
model_5b <- smart_cle1_sh %%
  lm(log(bmi) ~ activity + race_eth)

anova(model_5b)
```

Analysis of Variance Table

```
Response: log(bmi)
          Df Sum Sq Mean Sq F value    Pr(>F)
activity      3  2.060  0.68652 16.5090 1.676e-10 ***
race_eth      4  0.989  0.24716  5.9435 9.843e-05 ***
Residuals 1125 46.783  0.04158
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

Notice that the ANOVA model assesses these variables sequentially, so the $SS(\text{activity}) = 2.058$ is accounted for before we consider the $SS(\text{race_eth}) = 0.990$. Thus, in total, the model accounts for $2.058 + 0.990 = 3.048$ of the sums of squares in `log(bmi)` in these data.

If we flip the order in the model, like this:

```
smart_cle1_sh %%%
lm(log(bmi) ~ race_eth + activity) %>%
  anova()
```

Analysis of Variance Table

```
Response: log(bmi)
          Df Sum Sq Mean Sq F value    Pr(>F)
race_eth     4  1.119  0.27981  6.7287 2.371e-05 ***
activity      3  1.929  0.64299 15.4620 7.332e-10 ***
Residuals 1125 46.783  0.04158
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

- After flipping the order of the predictors, `race_eth` accounts for a larger Sum of Squares than it did previously, but `activity` accounts for a smaller amount, and the total between `race_eth` and `activity` remains the same, as $1.121 + 1.927$ is still 3.048.

5.2.1 Model Coefficients

The model coefficients are unchanged regardless of the order of the variables in our two-factor ANOVA model.

```
tidy(model_5b, conf.int = TRUE, conf.level = 0.95) %>%
  select(term, estimate, std.error, conf.low, conf.high) %>%
  knitr::kable(digits = 3)
```

term	estimate	std.error	conf.low	conf.high
(Intercept)	3.268	0.010	3.247	3.288
activityActive	0.012	0.018	-0.024	0.048
activityInsufficiently_Active	0.073	0.018	0.039	0.108
activityInactive	0.092	0.015	0.063	0.122
race_ethBlack non-Hispanic	0.066	0.015	0.036	0.096
race_ethOther race non-Hispanic	-0.086	0.042	-0.169	-0.002
race_ethMultiracial non-Hispanic	0.020	0.042	-0.063	0.103
race_ethHispanic	0.012	0.035	-0.057	0.082

The `model_5b` equation is:

```
log(BMI) = 3.268
+ 0.012 (activity = Active)
+ 0.073 (activity = Insufficiently Active)
+ 0.092 (activity = Inactive)
+ 0.066 (race_eth = Black non-Hispanic)
- 0.086 (race_eth = Other race non-Hispanic)
+ 0.020 (race_eth = Multiracial non-Hispanic)
```

```
+ 0.012 (race_eth = Hispanic)
```

and we can make predictions by filling in appropriate 1s and 0s for the indicator variables in parentheses.

For example, the predicted $\log(\text{BMI})$ for a White Highly Active person is 3.268, as White and Highly Active are the baseline categories in our two factors.

For all other combinations, we can make predictions as follows:

```
new_dat = tibble(
  race_eth = rep(c("White non-Hispanic",
                  "Black non-Hispanic",
                  "Other race non-Hispanic",
                  "Multiracial non-Hispanic",
                  "Hispanic"), 4),
  activity = c(rep("Highly_Active", 5),
              rep("Active", 5),
              rep("Insufficiently_Active", 5),
              rep("Inactive", 5))
)

augment(model_5b, newdata = new_dat)
```

Warning: 'newdata' had 20 rows but variables found have 1133 rows

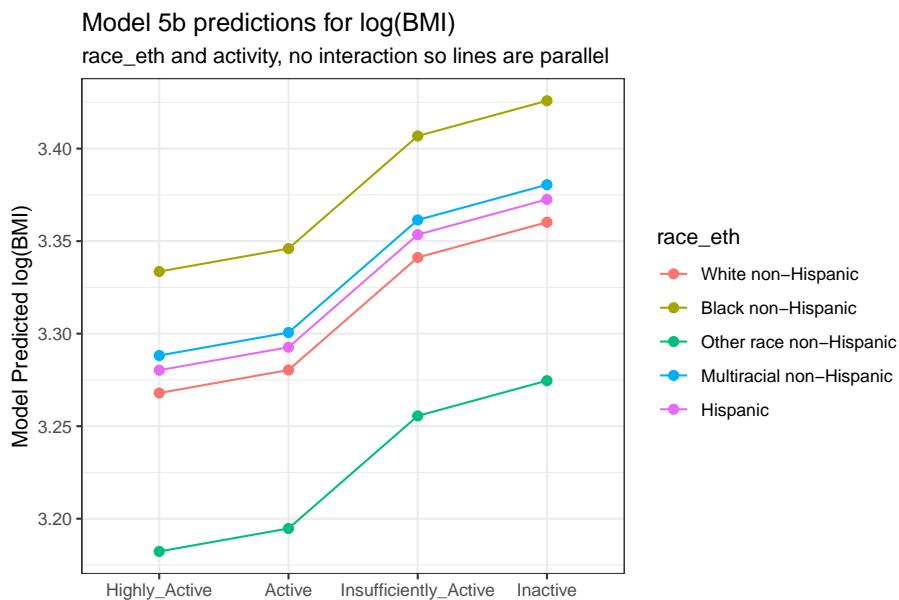
	race_eth	activity	.fitted
	<chr>	<chr>	<dbl>
1	White non-Hispanic	Highly_Active	3.27
2	Black non-Hispanic	Highly_Active	3.33
3	Other race non-Hispanic	Highly_Active	3.18
4	Multiracial non-Hispanic	Highly_Active	3.29
5	Hispanic	Highly_Active	3.28
6	White non-Hispanic	Active	3.28
7	Black non-Hispanic	Active	3.35
8	Other race non-Hispanic	Active	3.19
9	Multiracial non-Hispanic	Active	3.30
10	Hispanic	Active	3.29
11	White non-Hispanic	Insufficiently_Active	3.34
12	Black non-Hispanic	Insufficiently_Active	3.41
13	Other race non-Hispanic	Insufficiently_Active	3.26
14	Multiracial non-Hispanic	Insufficiently_Active	3.36
15	Hispanic	Insufficiently_Active	3.35
16	White non-Hispanic	Inactive	3.36
17	Black non-Hispanic	Inactive	3.43
18	Other race non-Hispanic	Inactive	3.27
19	Multiracial non-Hispanic	Inactive	3.38

```

20 Hispanic           Inactive      3.37
augment(model_5b, newdata = new_dat) %>%
  mutate(race_eth = fct_relevel(factor(race_eth),
                                "White non-Hispanic",
                                "Black non-Hispanic",
                                "Other race non-Hispanic",
                                "Multiracial non-Hispanic",
                                "Hispanic"),
         activity = fct_relevel(factor(activity),
                                "Highly_Active",
                                "Active",
                                "Insufficiently_Active",
                                "Inactive")) %>%
  ggplot(., aes(x = activity, y = .fitted,
                col = race_eth, group = race_eth)) +
  geom_point(size = 2) +
  geom_line() +
  labs(title = "Model 5b predictions for log(BMI)",
       subtitle = "race_eth and activity, no interaction so lines are parallel",
       y = "Model Predicted log(BMI)",
       x = "")

```

Warning: 'newdata' had 20 rows but variables found have 1133 rows



The lines joining the points for each `race_eth` category are parallel to each other.

The groups always hold the same position relative to each other, regardless of their activity levels, and vice versa. There is no interaction in this model allowing the predicted effects of, say, `activity` on `log(BMI)` values to differ for the various `race_eth` groups. To do that, we'd have to fit the two-factor ANOVA model incorporating an interaction term.

5.3 A Two-Factor ANOVA (with Interaction)

Let's add the interaction of `activity` and `race_eth` (symbolized in R by `activity * race_eth`) to the model for `log(BMI)`.

```
model_5c <- smart_cle1_sh %>%
  lm(log(bmi) ~ activity * race_eth)

anova(model_5c)
```

Analysis of Variance Table

	Df	Sum Sq	Mean Sq	F value	Pr(>F)						
activity	3	2.060	0.68652	16.4468	1.839e-10 ***						
race_eth	4	0.989	0.24716	5.9211	0.0001026 ***						
activity:race_eth	12	0.324	0.02700	0.6469	0.8028368						
Residuals	1113	46.459	0.04174								

Signif. codes:	0	'***'	0.001	'**'	0.01	'*'	0.05	'.'	0.1	' '	1

The ANOVA model shows that the $SS(\text{interaction}) = SS(\text{activity:race_eth})$ is 0.324, and uses 12 degrees of freedom. The model including the interaction term now accounts for $2.058 + 0.990 + 0.324 = 3.372$, which is 6.8% of the variation in `log(BMI)` overall (which is calculated as $SS(\text{Total}) = 2.058 + 0.990 + 0.324 + 46.456 = 49.828$.)

5.3.1 Model Coefficients

The model coefficients now include additional product terms that incorporate indicator variables for both `activity` and `race_eth`. For each of the product terms to take effect, both their `activity` and `race_eth` status must yield a 1 in the indicator variables.

```
tidy(model_5c, conf.int = TRUE, conf.level = 0.95) %>%
  select(term, estimate, std.error, conf.low, conf.high) %>%
  knitr::kable(digits = 3)
```

term	estimate	std.error	conf.low	conf.high
(Intercept)	3.264	0.011	3.242	3.286
activityActive	0.021	0.021	-0.021	0.063
activityInsufficiently_Active	0.079	0.020	0.039	0.119
activityInactive	0.097	0.018	0.063	0.131
race_ethBlack non-Hispanic	0.062	0.026	0.011	0.113
race_ethOther race non-Hispanic	-0.070	0.078	-0.223	0.088
race_ethMultiracial non-Hispanic	0.067	0.060	-0.051	0.188
race_ethHispanic	0.110	0.060	-0.008	0.239
activityActive:race_ethBlack non-Hispanic	-0.001	0.048	-0.096	0.094
activityInsufficiently_Active:race_ethBlack non-Hispanic	0.005	0.046	-0.086	0.096
activityInactive:race_ethBlack non-Hispanic	0.008	0.037	-0.065	0.081
activityActive:race_ethOther race non-Hispanic	-0.065	0.165	-0.389	0.059
activityInsufficiently_Active:race_ethOther race non-Hispanic	-0.035	0.101	-0.233	0.031
activityInactive:race_ethOther race non-Hispanic	0.033	0.129	-0.221	0.291
activityActive:race_ethMultiracial non-Hispanic	-0.208	0.134	-0.470	0.058
activityInsufficiently_Active:race_ethMultiracial non-Hispanic	-0.050	0.120	-0.285	0.025
activityInactive:race_ethMultiracial non-Hispanic	-0.056	0.110	-0.272	0.004
activityActive:race_ethHispanic	-0.104	0.096	-0.291	0.082
activityInsufficiently_Active:race_ethHispanic	-0.240	0.214	-0.660	0.074
activityInactive:race_ethHispanic	-0.169	0.082	-0.331	0.042

The model_5c equation is:

```

log(BMI) = 3.264
+ 0.021 (activity = Active)
+ 0.079 (activity = Insufficiently Active)
+ 0.097 (activity = Inactive)
+ 0.062 (race_eth = Black non-Hispanic)
- 0.070 (race_eth = Other race non-Hispanic)
+ 0.067 (race_eth = Multiracial non-Hispanic)
+ 0.110 (race_eth = Hispanic)
- 0.002 (activity = Active)(race_eth = Black non-Hispanic)
+ 0.005 (Insufficiently Active)(Black non-Hispanic)
+ 0.008 (Inactive)(Black non-Hispanic)
- 0.065 (Active)(Other race non-Hispanic)
- 0.035 (Insufficiently Active)(Other race non-Hispanic)
+ 0.033 (Inactive)(Other race non-Hispanic)
- 0.208 (Active)(Multiracial non-Hispanic)
- 0.050 (Insufficiently Active)(Multiracial non-Hispanic)
- 0.056 (Inactive)(Multiracial non-Hispanic)
- 0.104 (Active)(Hispanic)
- 0.240 (Insufficiently Active)(Hispanic)
- 0.169 (Inactive)(Hispanic)

```

and again, we can make predictions by filling in appropriate 1s and 0s for the indicator variables in parentheses.

For example, the predicted $\log(\text{BMI})$ for a White Highly Active person is 3.264, as White and Highly Active are the baseline categories in our two factors.

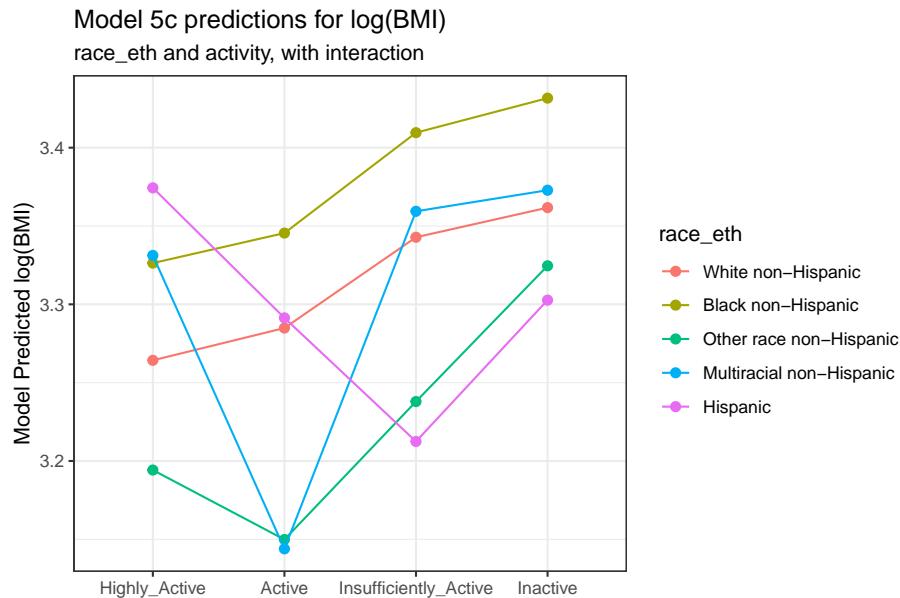
But the predicted $\log(\text{BMI})$ for a Hispanic Inactive person would be $3.264 + 0.097 + 0.110 - 0.169 = 3.302$.

Again, we'll plot the predicted $\log(\text{BMI})$ predictions for each possible combination.

```
new_dat = tibble(
  race_eth = rep(c("White non-Hispanic",
                  "Black non-Hispanic",
                  "Other race non-Hispanic",
                  "Multiracial non-Hispanic",
                  "Hispanic"), 4),
  activity = c(rep("Highly_Active", 5),
              rep("Active", 5),
              rep("Insufficiently_Active", 5),
              rep("Inactive", 5))
)

augment(model_5c, newdata = new_dat) %>%
  mutate(race_eth = fct_relevel(factor(race_eth),
                                "White non-Hispanic",
                                "Black non-Hispanic",
                                "Other race non-Hispanic",
                                "Multiracial non-Hispanic",
                                "Hispanic"),
        activity = fct_relevel(factor(activity),
                               "Highly_Active",
                               "Active",
                               "Insufficiently_Active",
                               "Inactive")) %>%
  ggplot(., aes(x = activity, y = .fitted,
                col = race_eth, group = race_eth)) +
  geom_point(size = 2) +
  geom_line() +
  labs(title = "Model 5c predictions for log(BMI)",
       subtitle = "race_eth and activity, with interaction",
       y = "Model Predicted log(BMI)",
       x = "")
```

Warning: 'newdata' had 20 rows but variables found have 1133 rows



Note that the lines joining the points for each `race_eth` category are no longer parallel to each other. The race-ethnicity group relative positions on $\log(\text{BMI})$ is now changing depending on the `activity` status.

5.3.2 Is the interaction term necessary?

We can assess this in three ways, in order of importance:

1. With an interaction plot
2. By assessing the fraction of the variation in the outcome accounted for by the interaction
3. By assessing whether the interaction accounts for statistically detectable outcome variation

5.3.2.1 The Interaction Plot

A simple interaction plot is just a plot of the unadjusted outcome means, stratified by the two factors. For example, consider this plot for our two-factor ANOVA model. To obtain this plot, we first summarize the means within each group.

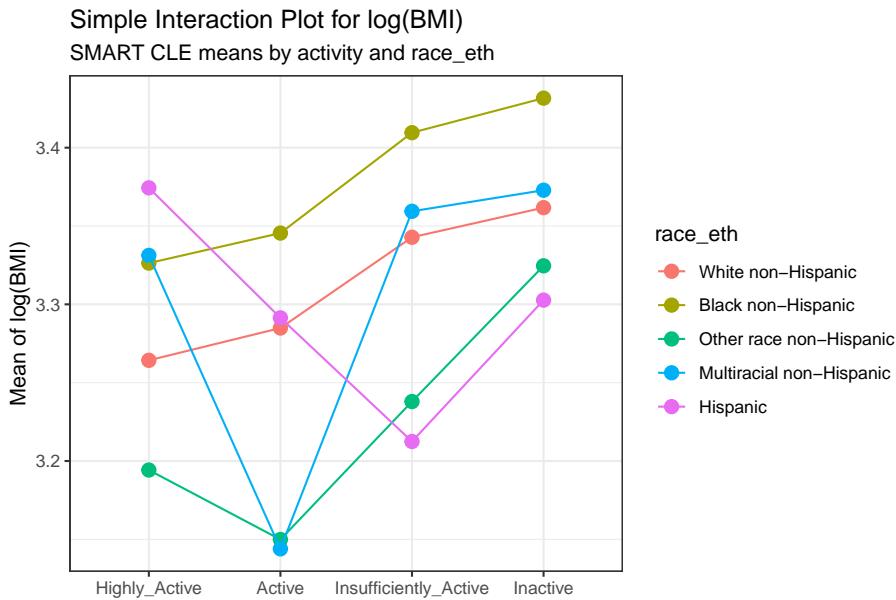
```
summaries_5 <- smart_cle1_sh %>%
  group_by(activity, race_eth) %>%
  summarize(n = n(), mean = mean(log(bmi)),
            sd = sd(log(bmi)))
```

`summarise()` has grouped output by 'activity'. You can override using the `groups` argument.

```
summaries_5
```

```
# A tibble: 20 x 5
# Groups:   activity [4]
  activity       race_eth      n  mean     sd
  <fct>        <fct>      <int> <dbl>  <dbl>
1 Highly_Active White non-Hispanic    320  3.26  0.176
2 Highly_Active Black non-Hispanic     77  3.33  0.190
3 Highly_Active Other race non-Hispanic 7  3.19  0.198
4 Highly_Active Multiracial non-Hispanic 12  3.33  0.187
5 Highly_Active Hispanic                12  3.37  0.296
6 Active         White non-Hispanic    129  3.28  0.173
7 Active         Black non-Hispanic     31  3.35  0.224
8 Active         Other race non-Hispanic 2  3.15  0.0845
9 Active         Multiracial non-Hispanic 3  3.14  0.121
10 Active        Hispanic                 8  3.29  0.213
11 Insufficiently_Active White non-Hispanic 150  3.34  0.194
12 Insufficiently_Active Black non-Hispanic 35  3.41  0.213
13 Insufficiently_Active Other race non-Hispanic 11  3.24  0.137
14 Insufficiently_Active Multiracial non-Hispanic 4  3.36  0.374
15 Insufficiently_Active Hispanic                1  3.21  NA
16 Inactive       White non-Hispanic    225  3.36  0.238
17 Inactive       Black non-Hispanic     83  3.43  0.247
18 Inactive       Other race non-Hispanic 4  3.32  0.238
19 Inactive       Multiracial non-Hispanic 5  3.37  0.129
20 Inactive       Hispanic                 14  3.30  0.264

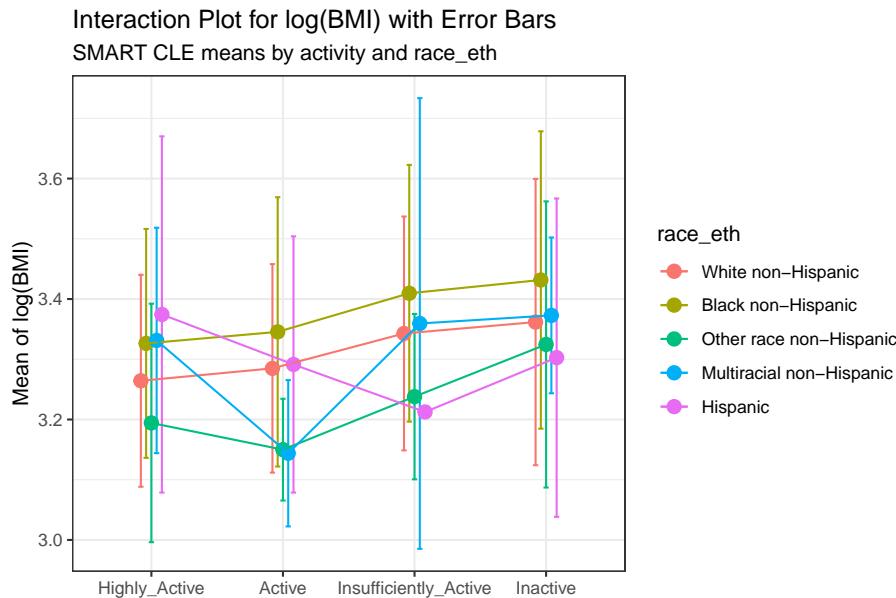
ggplot(summaries_5, aes(x = activity, y = mean,
                        color = race_eth,
                        group = race_eth)) +
  geom_point(size = 3) +
  geom_line() +
  labs(title = "Simple Interaction Plot for log(BMI)",
       subtitle = "SMART CLE means by activity and race_eth",
       x = "", y = "Mean of log(BMI)")
```



The interaction plot suggests that there is a modest interaction here. The White non-Hispanic and Black non-Hispanic groups appear pretty parallel (and they are the two largest groups) and Other race non-Hispanic has a fairly similar pattern, but the other two groups (Hispanic and Multiracial non-Hispanic) bounce around quite a bit based on activity level.

An alternative would be to include a small “dodge” for each point and include error bars (means \pm standard deviation) for each combination.

```
pd = position_dodge(0.2)
ggplot(summaries_5, aes(x = activity, y = mean,
                        color = race_eth,
                        group = race_eth)) +
  geom_errorbar(aes(ymin = mean - sd,
                     ymax = mean + sd),
                width = 0.2, position = pd) +
  geom_point(size = 3, position = pd) +
  geom_line(position = pd) +
  labs(title = "Interaction Plot for log(BMI) with Error Bars",
       subtitle = "SMART CLE means by activity and race_eth",
       x = "", y = "Mean of log(BMI)")
```



Here, we see a warning flag because we have one combination (which turns out to be Insufficiently Active and Hispanic) with only one observation in it, so a standard deviation cannot be calculated. In general, I'll stick with the simpler means plot most of the time.

5.3.2.2 Does the interaction account for substantial variation?

In this case, we can look at the fraction of the overall sums of squares accounted for by the interaction.

```
anova(model_5c)
```

Analysis of Variance Table

```
Response: log(bmi)
          Df Sum Sq Mean Sq F value    Pr(>F)
activity      3  2.060  0.68652 16.4468 1.839e-10 ***
race_eth      4  0.989  0.24716  5.9211 0.0001026 ***
activity:race_eth 12  0.324  0.02700  0.6469 0.8028368
Residuals   1113 46.459  0.04174
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

Here we have

$$\eta^2(\text{Interaction}) = \frac{0.324}{2.058 + 0.990 + 0.324 + 46.456} = 0.0065$$

so the interaction accounts for 0.65% of the variation in `bmi`. That looks pretty modest.

5.3.2.3 Does the interaction account for statistically detectable variation?

We can test this directly with the p value from the ANOVA table, which shows $p = 0.803$, which is far above any of our usual standards for a statistically detectable effect.

On the whole, I don't think the interaction term is especially helpful in improving this model.

In the next chapter, we'll look at two different examples of ANOVA models, now in more designed experiments. We'll also add some additional details on how the analyses might proceed.

We'll return to the SMART CLE data later in these Notes.

Chapter 6

Analysis of Variance

6.1 The bonding data: A Designed Dental Experiment

The bonding data describe a designed experiment into the properties of four different resin types (`resin` = A, B, C, D) and two different curing light sources (`light` = Halogen, LED) as they relate to the resulting bonding strength (measured in MPa¹) on the surface of teeth. The source is Kim (2014).

The experiment involved making measurements of bonding strength under a total of 80 experimental setups, or runs, with 10 runs completed at each of the eight combinations of a light source and a resin type. The data are gathered in the `bonding.csv` file.

```
bonding
```

```
# A tibble: 80 x 4
  run_ID   light   resin strength
  <chr>    <chr>   <chr>     <dbl>
1 R101     LED     B          12.8
2 R102     Halogen B          22.2
3 R103     Halogen B          24.6
4 R104     LED     A          17
5 R105     LED     C          32.2
6 R106     Halogen B          27.1
7 R107     LED     A          23.4
8 R108     Halogen A          23.5
9 R109     Halogen D          37.3
```

¹The MPa is defined as the failure load (in Newtons) divided by the entire bonded area, in mm².

```
10 R110    Halogen A      19.7  
# ... with 70 more rows
```

6.2 A One-Factor Analysis of Variance

Suppose we are interested in the distribution of the `strength` values for the four different types of `resin`.

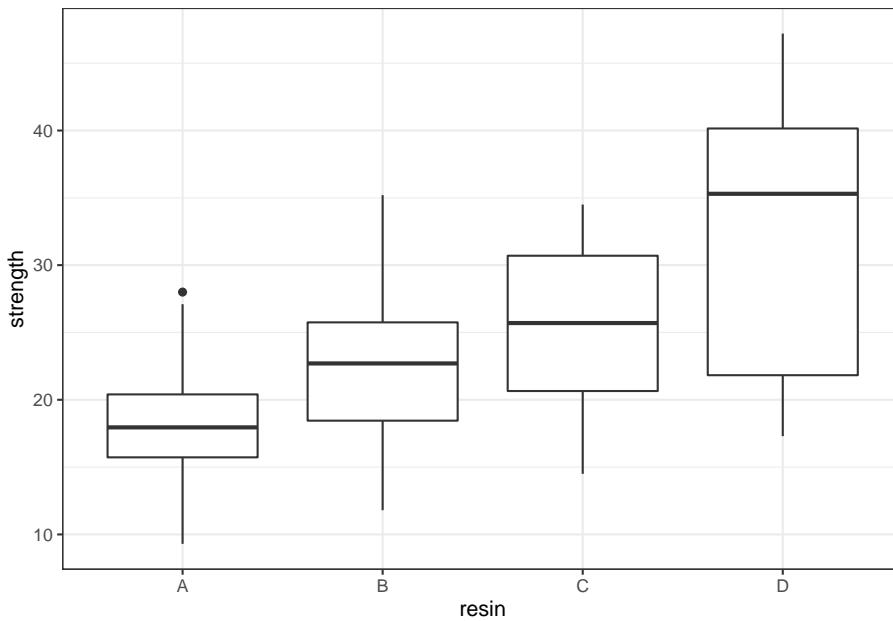
```
bonding %>% group_by(resin) %>% summarize(n = n(), mean(strength), median(strength))
```

```
# A tibble: 4 x 4  
  resin     n `mean(strength)` `median(strength)`  
  <chr> <int>          <dbl>            <dbl>  
1 A         20            18.4             18.0  
2 B         20            22.2             22.7  
3 C         20            25.2             25.7  
4 D         20            32.1             35.3
```

I'd begin serious work with a plot.

6.2.1 Look at the Data!

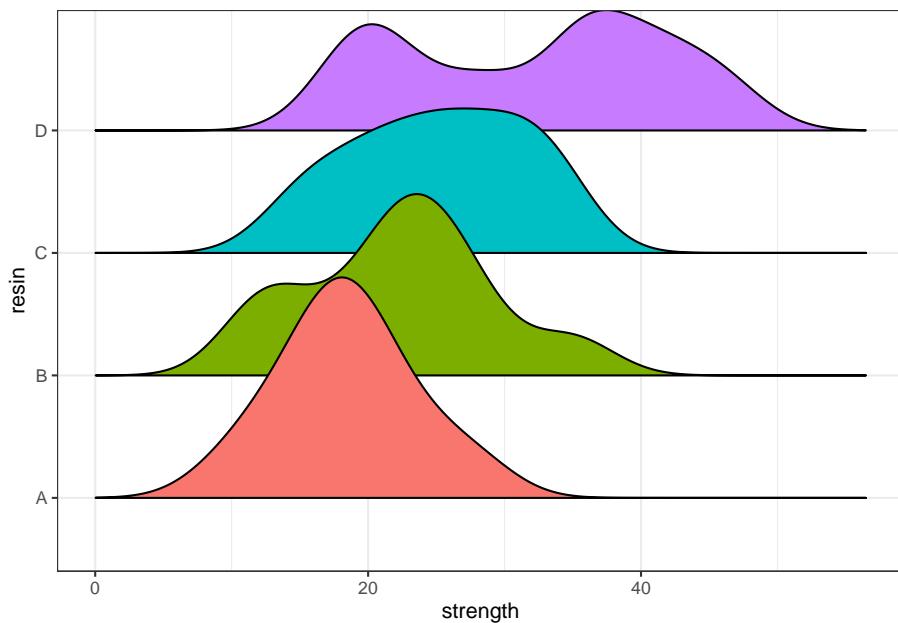
```
ggplot(bonding, aes(x = resin, y = strength)) +  
  geom_boxplot()
```



Another good plot for this purpose is a ridgeline plot.

```
ggplot(bonding, aes(x = strength, y = resin, fill = resin)) +  
  geom_density_ridges2() +  
  guides(fill = "none")
```

Picking joint bandwidth of 3.09



6.2.2 Table of Summary Statistics

With the small size of this experiment ($n = 20$ for each `resin` type), graphical summaries may not perform as well as they often do. We'll also produce a quick table of summary statistics for `strength` within each `resin` type.

```
bonding %$% mosaic::favstats(strength ~ resin)
```

	resin	min	Q1	median	Q3	max	mean	sd	n	missing
1	A	9.3	15.725	17.95	20.40	28.0	18.415	4.805948	20	0
2	B	11.8	18.450	22.70	25.75	35.2	22.230	6.748263	20	0
3	C	14.5	20.650	25.70	30.70	34.5	25.155	6.326425	20	0
4	D	17.3	21.825	35.30	40.15	47.2	32.075	9.735063	20	0

Since the means and medians within each group are fairly close, and the distributions (with the possible exception of `resin` D) are reasonably well approximated by the Normal, I'll fit an ANOVA model².

```
anova(lm(strength ~ resin, data = bonding))
```

Analysis of Variance Table

Response: strength

²If the data weren't approximately Normally distributed, we might instead consider a rank-based alternative to ANOVA, like the Kruskal-Wallis test.

```
Df Sum Sq Mean Sq F value    Pr(>F)
resin      3 1999.7  666.57  13.107 5.52e-07 ***
Residuals 76 3865.2   50.86
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

It appears that the `resin` types have a significant association with mean `strength` of the bonds. Can we identify which `resin` types have generally higher or lower `strength`?

```
TukeyHSD(aov(lm(strength ~ resin, data = bonding)))
```

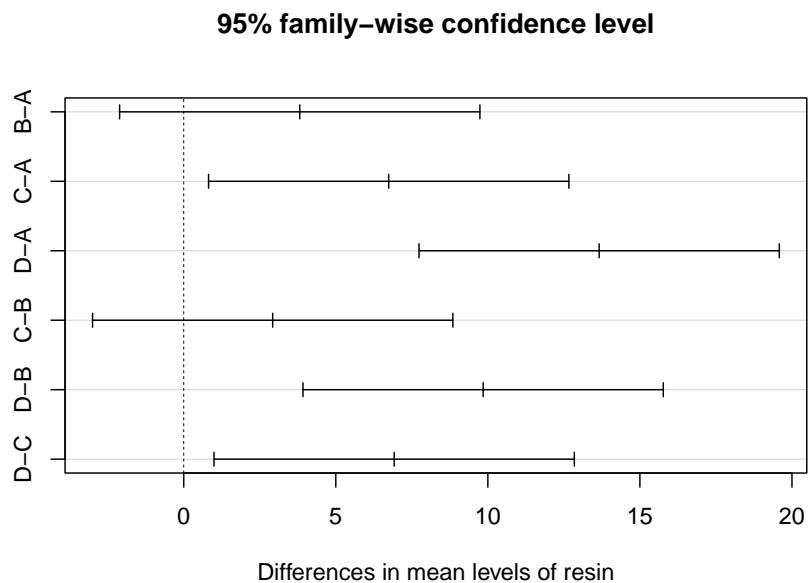
```
Tukey multiple comparisons of means
95% family-wise confidence level

Fit: aov(formula = lm(strength ~ resin, data = bonding))

$resin
   diff      lwr      upr     p adj
B-A  3.815 -2.1088676  9.738868 0.3351635
C-A  6.740  0.8161324 12.663868 0.0193344
D-A 13.660  7.7361324 19.583868 0.0000003
C-B  2.925 -2.9988676  8.848868 0.5676635
D-B  9.845  3.9211324 15.768868 0.0002276
D-C  6.920  0.9961324 12.843868 0.0154615
```

Based on these confidence intervals (which have a family-wise 95% confidence level), we see that D is associated with significantly larger mean `strength` than A or B or C, and that C is also associated with significantly larger mean `strength` than A. This may be easier to see in a plot of these confidence intervals.

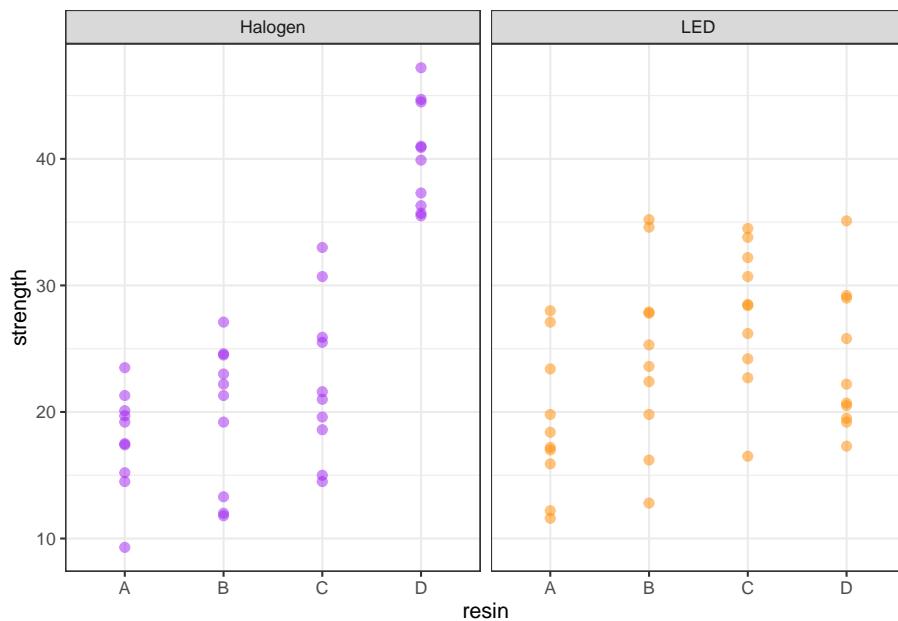
```
plot(TukeyHSD(aov(lm(strength ~ resin, data = bonding))))
```



6.3 A Two-Way ANOVA: Looking at Two Factors

Now, we'll now add consideration of the `light` source into our study. We can look at the distribution of the `strength` values at the combinations of both `light` and `resin`, with a plot like this one.

```
ggplot(bonding, aes(x = resin, y = strength, color = light)) +
  geom_point(size = 2, alpha = 0.5) +
  facet_wrap(~ light) +
  guides(color = "none") +
  scale_color_manual(values = c("purple", "darkorange")) +
  theme_bw()
```



6.4 A Means Plot (with standard deviations) to check for interaction

Sometimes, we'll instead look at a plot simply of the means (and, often, the standard deviations) of strength at each combination of `light` and `resin`. We'll start by building up a data set with the summaries we want to plot.

```
bond.sum <- bonding %>%
  group_by(resin, light) %>%
  summarise(mean.str = mean(strength), sd.str = sd(strength))

`summarise()` has grouped output by 'resin'. You can override using the
`.groups` argument.

bond.sum

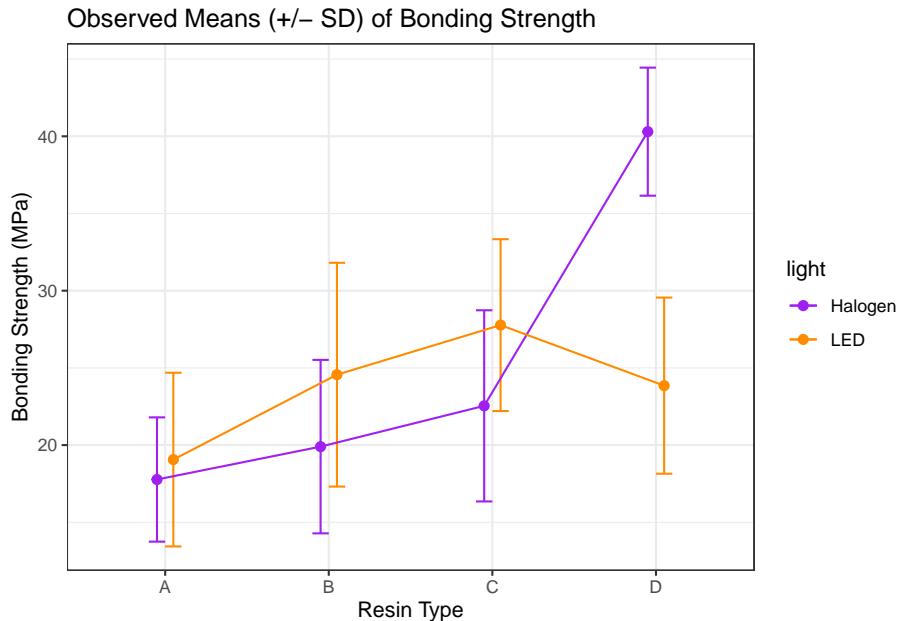
# A tibble: 8 x 4
# Groups:   resin [4]
  resin light   mean.str sd.str
  <chr> <chr>     <dbl>  <dbl>
1 A     Halogen    17.8   4.02
2 A     LED        19.1   5.63
3 B     Halogen    19.9   5.62
4 B     LED        24.6   7.25
```

5	C	Halogen	22.5	6.19
6	C	LED	27.8	5.56
7	D	Halogen	40.3	4.15
8	D	LED	23.8	5.70

Now, we'll use this new data set to plot the means and standard deviations of strength at each combination of resin and light.

```
## The error bars will overlap unless we adjust the position.
pd <- position_dodge(0.2) # move them .1 to the left and right

ggplot(bond.sum, aes(x = resin, y = mean.str, col = light)) +
  geom_errorbar(aes(ymin = mean.str - sd.str,
                     ymax = mean.str + sd.str),
                width = 0.2, position = pd) +
  geom_point(size = 2, position = pd) +
  geom_line(aes(group = light), position = pd) +
  scale_color_manual(values = c("purple", "darkorange")) +
  theme_bw() +
  labs(y = "Bonding Strength (MPa)", x = "Resin Type",
       title = "Observed Means (+/- SD) of Bonding Strength")
```



Is there evidence of a meaningful interaction between the resin type and the light source on the bonding strength in this plot?

- Sure. A meaningful interaction just means that the strength associated with different resin types depends on the light source.

- With LED light, it appears that `resin` C leads to the strongest bonding strength.
- With Halogen light, though, it seems that `resin` D is substantially stronger.
- Note that the lines we see here connecting the `light` sources aren't in parallel (as they would be if we had zero interaction between `resin` and `light`), but rather, they cross.

6.4.1 Summarizing the data after grouping by `resin` and `light`

We might want to look at a numerical summary of the strengths within these groups, too.

```
bonding %>% mosaic::favstats(strength ~ resin + light) %>%
  select(resin.light, median, mean, sd, n, missing)
```

	resin.light	median	mean	sd	n	missing
1	A.Halogen	18.35	17.77	4.024108	10	0
2	B.Halogen	21.75	19.90	5.617631	10	0
3	C.Halogen	21.30	22.54	6.191069	10	0
4	D.Halogen	40.40	40.30	4.147556	10	0
5	A.LED	17.80	19.06	5.625181	10	0
6	B.LED	24.45	24.56	7.246792	10	0
7	C.LED	28.45	27.77	5.564980	10	0
8	D.LED	21.45	23.85	5.704043	10	0

6.5 Fitting the Two-Way ANOVA model with Interaction

```
c3_m1 <- lm(strength ~ resin * light, data = bonding)
```

```
summary(c3_m1)
```

Call:

```
lm(formula = strength ~ resin * light, data = bonding)
```

Residuals:

Min	1Q	Median	3Q	Max
-11.760	-3.663	-0.320	3.697	11.250

Coefficients:

Estimate	Std. Error	t value	Pr(> t)
----------	------------	---------	----------

```
(Intercept) 17.770   1.771 10.033 2.57e-15 ***
resinB      2.130   2.505  0.850  0.3979
resinC      4.770   2.505  1.904  0.0609 .
resinD     22.530   2.505  8.995 2.13e-13 ***
lightLED    1.290   2.505  0.515  0.6081
resinB:lightLED 3.370   3.542  0.951  0.3446
resinC:lightLED 3.940   3.542  1.112  0.2697
resinD:lightLED -17.740  3.542 -5.008 3.78e-06 ***
---
Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 5.601 on 72 degrees of freedom
Multiple R-squared: 0.6149, Adjusted R-squared: 0.5775
F-statistic: 16.42 on 7 and 72 DF, p-value: 9.801e-13
```

6.5.1 The ANOVA table for our model

In a two-way ANOVA model, we begin by assessing the interaction term. If it's important, then our best model is the model including the interaction. If it's not important, we will often move on to consider a new model, fit without an interaction.

The ANOVA table is especially helpful in this case, because it lets us look specifically at the interaction effect.

```
anova(c3_m1)
```

Analysis of Variance Table

```
Response: strength
          Df Sum Sq Mean Sq F value    Pr(>F)
resin      3 1999.72 666.57 21.2499 5.792e-10 ***
light      1   34.72   34.72  1.1067   0.2963
resin:light 3 1571.96 523.99 16.7043 2.457e-08 ***
Residuals 72 2258.52   31.37
---
Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

6.5.2 Is the interaction important?

In this case, the interaction:

- is evident in the means plot, and
- is highly statistically significant, and
- accounts for a sizable fraction (27%) of the overall variation

$$\eta^2_{interaction} = \frac{SS(\text{resin:light})}{SS(\text{Total})} = \frac{1571.96}{1999.72 + 34.72 + 1571.96 + 2258.52} = 0.268$$

If the interaction were *either* large or significant we would be inclined to keep it in the model. In this case, it's both, so there's no real reason to remove it.

6.5.3 Interpreting the Interaction

Recall the model equation, which is:

```
c3_m1
```

Call:

```
lm(formula = strength ~ resin * light, data = bonding)
```

Coefficients:

	(Intercept)	resinB	resinC	resinD
lightLED	17.77	2.13	4.77	22.53
	1.29	3.37	3.94	-17.74

so we have:

$$strength = 17.77 + 2.13resinB + 4.77resinC + 22.53resinD + 1.29lightLED + 3.37resinB*lightLED + 3.94resinC*lightLED$$

So, if `light` = Halogen, our equation is:

$$strength = 17.77 + 2.13resinB + 4.77resinC + 22.53resinD$$

And if `light` = LED, our equation is:

$$strength = 19.06 + 5.50resinB + 8.71resinC + 4.79resinD$$

Note that both the intercept and the slopes change as a result of the interaction. The model yields a different prediction for every possible combination of a `resin` type and a `light` source.

6.6 Comparing Individual Combinations of resin and light

To make comparisons between individual combinations of a `resin` type and a `light` source, using something like Tukey's HSD approach for multiple comparisons, we first refit the model using the `aov` structure, rather than `lm`.

```
c3m1_aov <- aov(strength ~ resin * light, data = bonding)
```

```
summary(c3m1_aov)
```

	Df	Sum Sq	Mean Sq	F value	Pr(>F)						
resin	3	1999.7	666.6	21.250	5.79e-10 ***						
light	1	34.7	34.7	1.107	0.296						
resin:light	3	1572.0	524.0	16.704	2.46e-08 ***						
Residuals	72	2258.5	31.4								

Signif. codes:	0	'***'	0.001	'**'	0.01	'*'	0.05	'..'	0.1	' '	1

And now, we can obtain Tukey HSD comparisons (which will maintain an overall 95% family-wise confidence level) across the `resin` types, the `light` sources, and the combinations, with the `TukeyHSD` command. This approach is only completely appropriate if these comparisons are pre-planned, and if the design is balanced (as this is, with the same sample size for each combination of a `light` source and `resin` type.)

```
TukeyHSD(c3m1_aov)
```

```
Tukey multiple comparisons of means
95% family-wise confidence level
```

```
Fit: aov(formula = strength ~ resin * light, data = bonding)
```

```
$resin
  diff      lwr      upr      p adj
B-A  3.815 -0.843129  8.473129 0.1461960
C-A  6.740  2.081871 11.398129 0.0016436
D-A 13.660  9.001871 18.318129 0.0000000
C-B  2.925 -1.733129  7.583129 0.3568373
D-B  9.845  5.186871 14.503129 0.0000026
D-C  6.920  2.261871 11.578129 0.0011731
```

```
$light
  diff      lwr      upr      p adj
LED-Halogen -1.3175 -3.814042 1.179042 0.2963128
```

```
$`resin:light`
```

	diff	lwr	upr	p adj
B:Halogen-A:Halogen	2.13	-5.68928258	9.949283	0.9893515
C:Halogen-A:Halogen	4.77	-3.04928258	12.589283	0.5525230
D:Halogen-A:Halogen	22.53	14.71071742	30.349283	0.0000000
A:LED-A:Halogen	1.29	-6.52928258	9.109283	0.9995485
B:LED-A:Halogen	6.79	-1.02928258	14.609283	0.1361092
C:LED-A:Halogen	10.00	2.18071742	17.819283	0.0037074
D:LED-A:Halogen	6.08	-1.73928258	13.899283	0.2443200
C:Halogen-B:Halogen	2.64	-5.17928258	10.459283	0.9640100
D:Halogen-B:Halogen	20.40	12.58071742	28.219283	0.0000000
A:LED-B:Halogen	-0.84	-8.65928258	6.979283	0.9999747
B:LED-B:Halogen	4.66	-3.15928258	12.479283	0.5818695
C:LED-B:Halogen	7.87	0.05071742	15.689283	0.0473914
D:LED-B:Halogen	3.95	-3.86928258	11.769283	0.7621860
D:Halogen-C:Halogen	17.76	9.94071742	25.579283	0.0000000
A:LED-C:Halogen	-3.48	-11.29928258	4.339283	0.8591455
B:LED-C:Halogen	2.02	-5.79928258	9.839283	0.9922412
C:LED-C:Halogen	5.23	-2.58928258	13.049283	0.4323859
D:LED-C:Halogen	1.31	-6.50928258	9.129283	0.9995004
A:LED-D:Halogen	-21.24	-29.05928258	-13.420717	0.0000000
B:LED-D:Halogen	-15.74	-23.55928258	-7.920717	0.0000006
C:LED-D:Halogen	-12.53	-20.34928258	-4.710717	0.0001014
D:LED-D:Halogen	-16.45	-24.26928258	-8.630717	0.0000002
B:LED-A:LED	5.50	-2.31928258	13.319283	0.3665620
C:LED-A:LED	8.71	0.89071742	16.529283	0.0185285
D:LED-A:LED	4.79	-3.02928258	12.609283	0.5471915
C:LED-B:LED	3.21	-4.60928258	11.029283	0.9027236
D:LED-B:LED	-0.71	-8.52928258	7.109283	0.9999920
D:LED-C:LED	-3.92	-11.73928258	3.899283	0.7690762

One conclusion from this is that the combination of D and Halogen is significantly stronger than each of the other seven combinations.

6.7 The bonding model without Interaction

It seems incorrect in this situation to fit a model without the interaction term, but we'll do so just so you can see what's involved.

```
c3_m2 <- lm(strength ~ resin + light, data = bonding)
```

```
summary(c3_m2)
```

Call:

```
lm(formula = strength ~ resin + light, data = bonding)
```

Residuals:

	Min	1Q	Median	3Q	Max
	-14.1163	-4.9531	0.1187	4.4613	14.4663

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	19.074	1.787	10.676	< 2e-16 ***
resinB	3.815	2.260	1.688	0.09555 .
resinC	6.740	2.260	2.982	0.00386 **
resinD	13.660	2.260	6.044	5.39e-08 ***
lightLED	-1.317	1.598	-0.824	0.41229

Signif. codes:	0 '***'	0.001 '**'	0.01 '*'	0.05 '.'
	0.1 ' '	1		

Residual standard error: 7.147 on 75 degrees of freedom
 Multiple R-squared: 0.3469, Adjusted R-squared: 0.312
 F-statistic: 9.958 on 4 and 75 DF, p-value: 1.616e-06

In the no-interaction model, if `light` = Halogen, our equation is:

$$\text{strength} = 19.07 + 3.82\text{resinB} + 6.74\text{resinC} + 13.66\text{resinD}$$

And if `light` = LED, our equation is:

$$\text{strength} = 17.75 + 3.82\text{resinB} + 6.74\text{resinC} + 13.66\text{resinD}$$

So, in the no-interaction model, only the intercept changes.

`anova(c3_m2)`

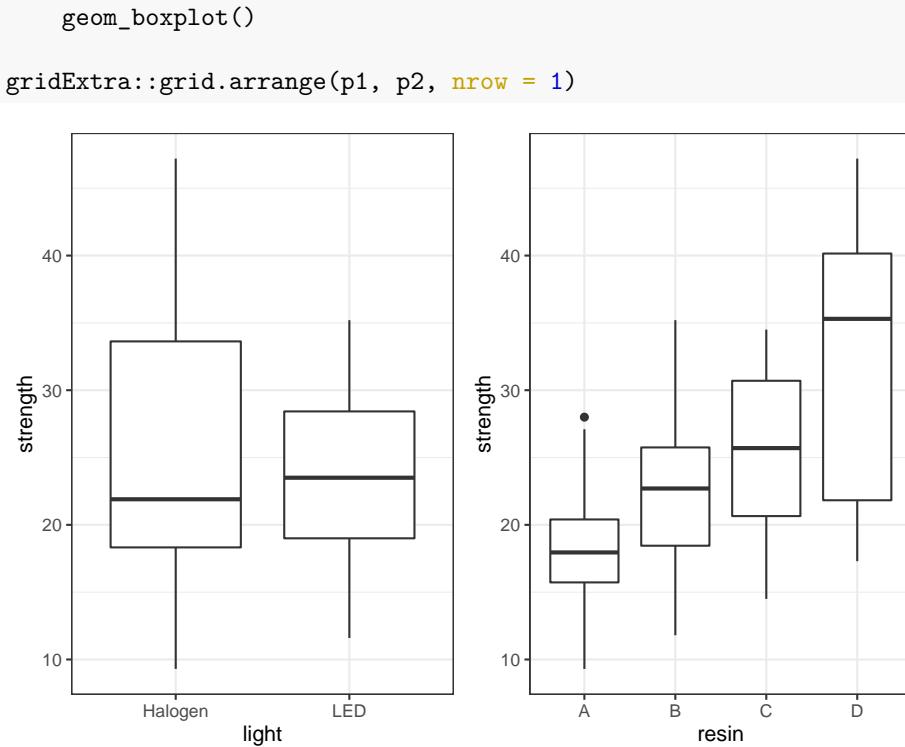
Analysis of Variance Table

Response: strength	Df	Sum Sq	Mean Sq	F value	Pr(>F)
resin	3	1999.7	666.57	13.0514	6.036e-07 ***
light	1	34.7	34.72	0.6797	0.4123
Residuals	75	3830.5	51.07		

Signif. codes:	0 '***'	0.001 '**'	0.01 '*'	0.05 '.'	0.1 ' '
	1				

And, it appears, if we ignore the interaction, then `resin` type has a significant impact on `strength` but `light` source doesn't. This is clearer when we look at boxplots of the separated `light` and `resin` groups.

```
p1 <- ggplot(bonding, aes(x = light, y = strength)) +
  geom_boxplot()
p2 <- ggplot(bonding, aes(x = resin, y = strength)) +
```



6.8 cortisol: A Hypothetical Clinical Trial

156 adults who complained of problems with a high-stress lifestyle were enrolled in a hypothetical clinical trial of the effectiveness of a behavioral intervention designed to help reduce stress levels, as measured by salivary cortisol.

The subjects were randomly assigned to one of three intervention groups (usual care, low dose, and high dose.) The “low dose” subjects received a one-week intervention with a follow-up at week 5. The “high dose” subjects received a more intensive three-week intervention, with follow up at week 5.

Since cortisol levels rise and fall with circadian rhythms, the cortisol measurements were taken just after rising for all subjects. These measurements were taken at baseline, and again at five weeks. The difference (baseline - week 5) in cortisol level (in micrograms / l) serves as the primary outcome.

6.8.1 Codebook and Raw Data for `cortisol`

The data are gathered in the `cortisol` data set. Included are:

Variable	Description
subject	subject identification code
interv	intervention group (UC = usual care, Low, High)
waist	waist circumference at baseline (in inches)
sex	male or female
cort.1	salivary cortisol level (microg/l) week 1
cort.5	salivary cortisol level (microg/l) week 5

`cortisol`

```
# A tibble: 156 x 6
  subject interv  waist sex    cort.1   cort.5
  <dbl> <chr>    <dbl> <chr>   <dbl>   <dbl>
1     1001 UC      48.3 M     13.4   13.3
2     1002 Low     58.3 M     17.8   16.6
3     1003 High    43.0 M     14.4   12.7
4     1004 Low     44.9 M      9.0    9.8
5     1005 High    46.1 M     14.2   14.2
6     1006 UC      41.3 M     14.8   15.1
7     1007 Low     51.0 F     13.7   16.0
8     1008 UC      42.0 F     17.3   18.7
9     1009 Low     24.7 F     15.3   15.8
10    1010 Low     59.4 M     12.4   11.7
# ... with 146 more rows
```

6.9 Creating a factor combining sex and waist

Next, we'll put the `waist` and `sex` data in the `cortisol` example together. We want to build a second categorical variable (called `fat_est`) combining this information, to indicate "healthy" vs. "unhealthy" levels of fat around the waist.

- Male subjects whose waist circumference is 40 inches or more, and
- Female subjects whose waist circumference is 35 inches or more, will fall in the "unhealthy" group.

```
cortisol <- cortisol %>%
  mutate(
    fat_est = factor(case_when(
      sex == "M" & waist >= 40 ~ "unhealthy",
      sex == "F" & waist >= 35 ~ "unhealthy",
      TRUE                  ~ "healthy")),
    cort_diff = cort.1 - cort.5)

summary(cortisol)
```

subject	interv	waist	sex	
Min.	:1001	Length:156	Min. :20.80	Length:156
1st Qu.	:1040	Class :character	1st Qu.:33.27	Class :character
Median	:1078	Mode :character	Median :40.35	Mode :character
Mean	:1078		Mean :40.42	
3rd Qu.	:1117		3rd Qu.:47.77	
Max.	:1156		Max. :59.90	
cort.1	cort.5	fat_est	cort_diff	
Min.	: 6.000	Min. : 4.2	healthy : 56	Min. : -2.3000
1st Qu.	: 9.675	1st Qu.: 9.6	unhealthy:100	1st Qu.: -0.5000
Median	:12.400	Median :12.6		Median : 0.2000
Mean	:12.686	Mean :12.4		Mean : 0.2821
3rd Qu.	:16.025	3rd Qu.:15.7		3rd Qu.: 1.2000
Max.	:19.000	Max. :19.7		Max. : 2.0000

6.10 A Means Plot for the cortisol trial (with standard errors)

Again, we'll start by building up a data set with the summaries we want to plot.

```
cort.sum <- cortisol %>%
  group_by(interv, fat_est) %>%
  summarise(mean.cort = mean(cort_diff),
            se.cort = sd(cort_diff)/sqrt(n()))

`summarise()` has grouped output by 'interv'. You can override using the
`.groups` argument.

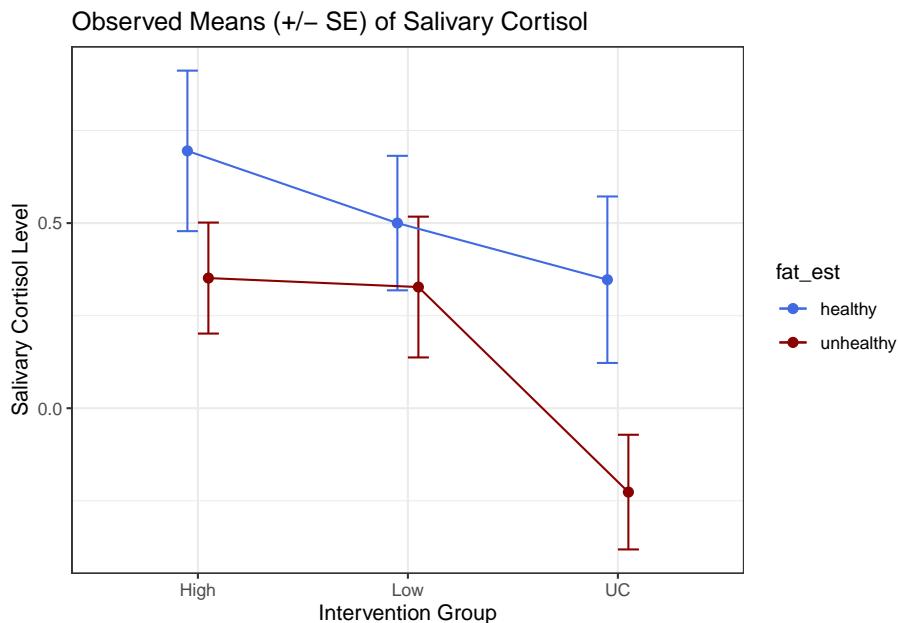
cort.sum

# A tibble: 6 x 4
# Groups:   interv [3]
  interv fat_est  mean.cort se.cort
  <chr>   <fct>     <dbl>    <dbl>
1 High    healthy    0.695    0.217
2 High    unhealthy   0.352    0.150
3 Low     healthy    0.5      0.182
4 Low     unhealthy   0.327    0.190
5 UC      healthy    0.347    0.225
6 UC      unhealthy   -0.226   0.155
```

Now, we'll use this new data set to plot the means and standard errors.

```
## The error bars will overlap unless we adjust the position.
pd <- position_dodge(0.2) # move them .1 to the left and right
```

```
ggplot(cort.sum, aes(x = interv, y = mean.cort, col = fat_est)) +
  geom_errorbar(aes(ymin = mean.cort - se.cort,
                     ymax = mean.cort + se.cort),
                width = 0.2, position = pd) +
  geom_point(size = 2, position = pd) +
  geom_line(aes(group = fat_est), position = pd) +
  scale_color_manual(values = c("royalblue", "darkred")) +
  theme_bw() +
  labs(y = "Salivary Cortisol Level", x = "Intervention Group",
       title = "Observed Means (+/- SE) of Salivary Cortisol")
```



6.11 A Two-Way ANOVA model for cortisol with Interaction

```
c3_m3 <- lm(cort_diff ~ interv * fat_est, data = cortisol)
anova(c3_m3)
```

Analysis of Variance Table

```
Response: cort_diff
          Df  Sum Sq Mean Sq F value Pr(>F)
```

6.11. A TWO-WAY ANOVA MODEL FOR CORTISOL WITH INTERACTION195

```
interv          2    7.847   3.9235   4.4698  0.01301 *
fat_est         1    4.614   4.6139   5.2564  0.02326 *
interv:fat_est  2    0.943   0.4715   0.5371  0.58554
Residuals      150  131.666   0.8778
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

Does it seem like we need the interaction term in this case?

```
summary(c3_m3)
```

```
Call:
lm(formula = cort_diff ~ interv * fat_est, data = cortisol)
```

```
Residuals:
    Min      1Q      Median      3Q      Max
-2.62727 -0.75702  0.08636  0.84848  2.12647
```

```
Coefficients:
```

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	0.6950	0.2095	3.317	0.00114 **
intervLow	-0.1950	0.3001	-0.650	0.51689
intervUC	-0.3479	0.3091	-1.126	0.26206
fat_estunhealthy	-0.3435	0.2655	-1.294	0.19774
intervLow:fat_estunhealthy	0.1708	0.3785	0.451	0.65256
intervUC:fat_estunhealthy	-0.2300	0.3846	-0.598	0.55068

```
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
Residual standard error: 0.9369 on 150 degrees of freedom
Multiple R-squared:  0.0924,    Adjusted R-squared:  0.06214
F-statistic: 3.054 on 5 and 150 DF,  p-value: 0.01179
```

How do you reconcile the apparent difference in significance levels between this regression summary and the ANOVA table above?

6.11.1 Notes on this Question

When we're evaluating a two-factor ANOVA model with an interaction, we are choosing between models with either:

1. just one factor
2. both factors, but only as main effects
3. both factors and an interaction between them

But we don't get to pick models that include any other combination of terms. For this two-way ANOVA, then, our choices are:

- a model with ‘interv only’
- a model with ‘fat_est only’
- a model with both `interv` and `fat_est` but not their interaction
- a model with `interv` and `fat_est` and their interaction

Those are the only modeling options available to us.

First, consider the ANOVA table, repeated below...

```
anova(c3_m3)
```

Analysis of Variance Table

```
Response: cort_diff
          Df  Sum Sq Mean Sq F value Pr(>F)
interv      2   7.847  3.9235  4.4698 0.01301 *
fat_est     1   4.614  4.6139  5.2564 0.02326 *
interv:fat_est  2   0.943  0.4715  0.5371 0.58554
Residuals   150 131.666  0.8778
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

The conclusions here are as follows:

1. The interaction effect (`interv:fat_est`) has a large p value (0.58554) and assesses whether the two interaction terms (product terms) included in the model add detectable predictive value to the main effects model that includes only `interv` and `fat_est` alone. You are right to say that this ANOVA is sequential, which means that the p value for the interaction effect is looking at the additional effect of the interaction once we already have the main effects `interv` and `fat_est` included.
2. The `interv` and `fat_est` terms aren't usually evaluated with significance tests or interpreted in the ANOVA for this setting, since if we intend to include the interaction term (as this model does) we also need these main effects. If we wanted to look at those terms individually in a model without the interaction, then we'd want to fit that model (without the interaction term) to do so.

Next, let's look at the summary of the `c3_m3` model, specifically the coefficients...

```
summary(c3_m3)
```

Call:

```
lm(formula = cort_diff ~ interv * fat_est, data = cortisol)
```

Residuals:

Min	1Q	Median	3Q	Max
-2.62727	-0.75702	0.08636	0.84848	2.12647

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	0.6950	0.2095	3.317	0.00114 **
intervLow	-0.1950	0.3001	-0.650	0.51689
intervUC	-0.3479	0.3091	-1.126	0.26206
fat_estunhealthy	-0.3435	0.2655	-1.294	0.19774
intervLow:fat_estunhealthy	0.1708	0.3785	0.451	0.65256
intervUC:fat_estunhealthy	-0.2300	0.3846	-0.598	0.55068

Signif. codes:	0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1			

Residual standard error: 0.9369 on 150 degrees of freedom
 Multiple R-squared: 0.0924, Adjusted R-squared: 0.06214
 F-statistic: 3.054 on 5 and 150 DF, p-value: 0.01179

So here, we see two p values associated with the interaction terms (the two product terms at the bottom of the regression) but these aren't especially helpful, because we're either going to include the interaction (in which case both of these terms will be in the model) or we're not going to include the interaction (in which case neither of these terms will be in the model.)

So the p values provided here aren't very helpful - like all such p values for t tests, they are looking at the value of the term in their row as the last predictor in to the model, essentially comparing the full model to the model without that specific component, but none of those tests enable us to decide which of the 4 available model choices is our best fit.

Now, let's consider the reason why, for example, the p value for `fat_est` in the `summary()` which is looking at comparing the following models ...

- a model including `interv` (which has 2 coefficients to account for its 3 categories), `fat_est` (which has 1 coefficient to account for its 2 categories), and the `interv*fat_est` interaction terms (which are 2 terms)
- a model including `interv` and the `interv*fat_est` interaction (but somehow not the main effect of `fat_est`, which actually makes no sense: if we include the interaction we always include the main effect)

to the p value for `fat_est` in the ANOVA which is looking at comparing

- the model with `interv` alone to
- the model with `interv` and `fat_est` as main effects, but no interaction

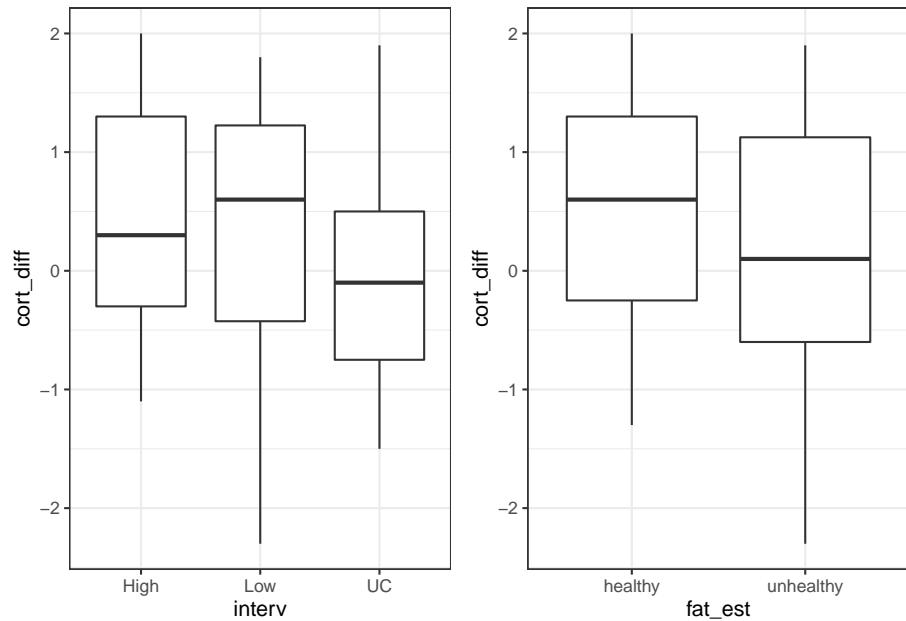
Only the ANOVA p value is therefore in any way useful, and it suggests that once you have the main effect of `interv`, adding `fat_est`'s main effect adds statistically detectable value (p = 0.023)

6.12 A Two-Way ANOVA model for cortisol without Interaction

6.12.1 The Graph

```
p1 <- ggplot(cortisol, aes(x = interv, y = cort_diff)) +
  geom_boxplot()
p2 <- ggplot(cortisol, aes(x = fat_est, y = cort_diff)) +
  geom_boxplot()

gridExtra::grid.arrange(p1, p2, nrow = 1)
```



6.12.2 The ANOVA Model

```
c3_m4 <- lm(cort_diff ~ interv + fat_est, data = cortisol)
anova(c3_m4)
```

Analysis of Variance Table

Response: cort_diff	Df	Sum Sq	Mean Sq	F value	Pr(>F)
---------------------	----	--------	---------	---------	--------

```

interv      2    7.847  3.9235  4.4972  0.01266 *
fat_est     1    4.614  4.6139  5.2886  0.02283 *
Residuals  152 132.609  0.8724
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

```

How do these results compare to those we saw in the model with interaction?

6.12.3 The Regression Summary

```

summary(c3_m4)

Call:
lm(formula = cort_diff ~ interv + fat_est, data = cortisol)

Residuals:
    Min      1Q   Median      3Q      Max 
-2.55929 -0.74527  0.05457  0.86456  2.05489 

Coefficients:
            Estimate Std. Error t value Pr(>|t|)    
(Intercept)  0.70452   0.16093   4.378 2.22e-05 ***
intervLow   -0.08645   0.18232  -0.474  0.63606    
intervUC    -0.50063   0.18334  -2.731  0.00707 **  
fat_estunhealthy -0.35878   0.15601  -2.300  0.02283 *  
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.934 on 152 degrees of freedom
Multiple R-squared:  0.0859,    Adjusted R-squared:  0.06785 
F-statistic: 4.761 on 3 and 152 DF,  p-value: 0.00335

```

6.12.4 Tukey HSD Comparisons

Without the interaction term, we can make direct comparisons between levels of the intervention, and between levels of the `fat_est` variable. This is probably best done here in a Tukey HSD comparison.

```

TukeyHSD(aov(cort_diff ~ interv + fat_est, data = cortisol))

Tukey multiple comparisons of means
  95% family-wise confidence level

Fit: aov(formula = cort_diff ~ interv + fat_est, data = cortisol)

```

```
$interv
```

	diff	lwr	upr	p adj
Low-High	-0.09074746	-0.5222655	0.34077063	0.8724916
UC-High	-0.51642619	-0.9500745	-0.08277793	0.0150150
UC-Low	-0.42567873	-0.8613670	0.01000948	0.0570728

```
$fat_est
```

	diff	lwr	upr	p adj
unhealthy-healthy	-0.3582443	-0.6662455	-0.05024305	0.0229266

What conclusions can we draw, at a 5% significance level?

Chapter 7

Analysis of Covariance

7.1 An Emphysema Study

My source for this example is Riffenburgh (2006), section 18.4. Serum theophylline levels (in mg/dl) were measured in 16 patients with emphysema at baseline, then 5 days later (at the end of a course of antibiotics) and then at 10 days after baseline. Clinicians anticipate that the antibiotic will increase the theophylline level. The data are stored in the `emphysema.csv` data file, and note that the age for patient 5 is not available.

7.1.1 Codebook

Variable	Description
<code>patient</code>	ID code
<code>age</code>	patient's age in years
<code>sex</code>	patient's sex (F or M)
<code>st_base</code>	patient's serum theophylline at baseline (mg/dl)
<code>st_day5</code>	patient's serum theophylline at day 5 (mg/dl)
<code>st_day10</code>	patient's serum theophylline at day 10 (mg/dl)

We're going to look at the change from baseline to day 5 as our outcome of interest, since the clinical expectation is that the antibiotic (azithromycin) will increase theophylline levels.

```
emphysema <- emphysema %>%
  mutate(st_delta = st_day5 - st_base)
```

```
emphysema
```

```
# A tibble: 16 x 7
  patient age sex st_base st_day5 st_day10 st_delta
  <dbl> <dbl> <chr>   <dbl>    <dbl>    <dbl>    <dbl>
1      1    61 F     14.1     2.3     10.3    -11.8
2      2    70 F     7.2      5.4      7.3     -1.8
3      3    65 M     14.2     11.9     11.3     -2.3
4      4    65 M     10.3     10.7     13.8     0.400
5      5    NA M     9.9      10.7     11.7     0.800
6      6    76 M     5.2      6.8      4.2      1.6
7      7    72 M     10.4     14.6     14.1      4.2
8      8    69 F     10.5     7.2      5.4     -3.3
9      9    66 M     5         5       5.1      0
10     10   62 M     8.6      8.1      7.4     -0.5
11     11   65 F     16.6     14.9     13        -1.70
12     12   71 M     16.4     18.6     17.1      2.20
13     13   51 F     12.2     11       12.3     -1.2
14     14   71 M     6.6      3.7      4.5     -2.9
15     15   64 F     15.4     15.2     13.6     -0.200
16     16   50 M     10.2     10.8     11.2     0.600
```

7.2 Does `sex` affect the mean change in theophylline?

```
emphysema %$% mosaic::favstats(st_delta)
```

	min	Q1	median	Q3	max	mean	sd	n	missing
	-11.8	-1.925	-0.35	0.65	4.2	-0.99375	3.484149	16	0

```
emphysema %$% mosaic::favstats(st_delta ~ sex)
```

sex	min	Q1	median	Q3	max	mean	sd	n	missing
1 F	-11.8	-2.925	-1.75	-1.325	-0.2	-3.333333	4.267864	6	0
2 M	-2.9	-0.375	0.50	1.400	4.2	0.410000	2.067446	10	0

Overall, the mean change in theophylline during the course of the antibiotic is -0.99, but this is -3.33 for female patients and 0.41 for male patients.

A one-way ANOVA model looks like this:

```
anova(lm(st_delta ~ sex, data = emphysema))
```

Analysis of Variance Table

7.3. IS THERE AN ASSOCIATION BETWEEN AGE AND SEX IN THIS STUDY?203

```
Response: st_delta
          Df  Sum Sq Mean Sq F value Pr(>F)
sex        1  52.547  52.547  5.6789 0.03189 *
Residuals 14 129.542   9.253
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

The ANOVA F test finds a statistically significant difference between the mean `st_delta` among males and the mean `st_delta` among females. But is there more to the story?

7.3 Is there an association between age and sex in this study?

```
emphysema %$% mosaic::favstats(age ~ sex)
```

	sex	min	Q1	median	Q3	max	mean	sd	n	missing
1	F	51	61.75	64.5	68	70	63.33333	6.889606	6	0
2	M	50	65.00	66.0	71	76	66.44444	7.568208	9	1

But we note that the male patients are also older than the female patients, on average (mean age for males is 66.4, for females 63.3)

- Does the fact that male patients are older affect change in theophylline level?
- And how should we deal with the one missing `age` value (in a male patient)?

7.4 Adding a quantitative covariate, age, to the model

We could fit an ANOVA model to predict `st_delta` using `sex` and `age` directly, but only if we categorized `age` into two or more groups. Because `age` is not categorical, we cannot include it in an ANOVA. But if `age` is an influence, and we don't adjust for it, it may well bias the outcome of our initial ANOVA. With a quantitative variable like `age`, we will need a method called ANCOVA, for analysis of covariance.

7.4.1 The ANCOVA model

ANCOVA in this case is just an ANOVA model with our outcome (`st_delta`) adjusted for a continuous covariate, called `age`. For the moment, we'll ignore

the one subject with missing `age` and simply fit the regression model with `sex` and `age`.

```
summary(lm(st_delta ~ sex + age, data = emphysema))

Call:
lm(formula = st_delta ~ sex + age, data = emphysema)

Residuals:
    Min      1Q  Median      3Q     Max 
-8.3352 -0.4789  0.6948  1.5580  3.5202 

Coefficients:
            Estimate Std. Error t value Pr(>|t|)    
(Intercept) -6.90266   7.92948 -0.871   0.4011    
sexM         3.52466   1.75815  2.005   0.0681 .  
age          0.05636   0.12343  0.457   0.6561    
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 3.255 on 12 degrees of freedom
(1 observation deleted due to missingness)
Multiple R-squared:  0.2882,    Adjusted R-squared:  0.1696 
F-statistic:  2.43 on 2 and 12 DF,  p-value: 0.13
```

This model assumes that the slope of the regression line between `st_delta` and `age` is the same for both sexes.

Note that the model yields `st_delta = -6.9 + 3.52 (sex = male) + 0.056 age`, or

- `st_delta = -6.9 + 0.056 age` for female patients, and
- `st_delta = (-6.9 + 3.52) + 0.056 age = -3.38 + 0.056 age` for male patients.

Note that we can test this assumption of equal slopes by fitting an alternative model (with a product term between `sex` and `age`) that doesn't require the assumption, and we'll do that later.

7.4.2 The ANCOVA Table

First, though, we'll look at the ANCOVA table.

```
anova(lm(st_delta ~ sex + age, data = emphysema))
```

Analysis of Variance Table

Response: `st_delta`

```
Df  Sum Sq Mean Sq F value    Pr(>F)
sex      1  49.284  49.284  4.6507  0.05203 .
age      1   2.209   2.209  0.2085  0.65612
Residuals 12 127.164 10.597
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

When we tested `sex` without accounting for `age`, we found a p value of 0.032, which is less than our usual cutpoint of 0.05. But when we adjusted for `age`, we find that `sex` loses significance, even though `age` is not a significant influence on `st_delta` by itself, according to the ANCOVA table.

7.5 Rerunning the ANCOVA model after simple imputation

We could have *imputed* the missing `age` value for patient 5, rather than just deleting that patient. Suppose we do the simplest potentially reasonable thing to do: insert the mean `age` in where the NA value currently exists.

```
emph_imp <- replace_na(emphysema, list(age = mean(emphysema$age, na.rm = TRUE)))

emph_imp
```

	patient	age	sex	st_base	st_day5	st_day10	st_delta
1	1	61	F	14.1	2.3	10.3	-11.8
2	2	70	F	7.2	5.4	7.3	-1.8
3	3	65	M	14.2	11.9	11.3	-2.3
4	4	65	M	10.3	10.7	13.8	0.400
5	5	65.2	M	9.9	10.7	11.7	0.800
6	6	76	M	5.2	6.8	4.2	1.6
7	7	72	M	10.4	14.6	14.1	4.2
8	8	69	F	10.5	7.2	5.4	-3.3
9	9	66	M	5	5	5.1	0
10	10	62	M	8.6	8.1	7.4	-0.5
11	11	65	F	16.6	14.9	13	-1.70
12	12	71	M	16.4	18.6	17.1	2.20
13	13	51	F	12.2	11	12.3	-1.2
14	14	71	M	6.6	3.7	4.5	-2.9
15	15	64	F	15.4	15.2	13.6	-0.200
16	16	50	M	10.2	10.8	11.2	0.600

More on simple imputation and missing data is coming soon.

For now, we can rerun the ANCOVA model on this new data set, after imputation...

```
anova(lm(st_delta ~ sex + age, data = emph_imp))
```

Analysis of Variance Table

```
Response: st_delta
  Df Sum Sq Mean Sq F value Pr(>F)
sex      1 52.547 52.547  5.3623 0.03755 *
age      1   2.151   2.151  0.2195 0.64721
Residuals 13 127.392   9.799
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

When we do this, we see that now the `sex` variable returns to a *p* value below 0.05. Our complete case analysis (which omitted patient 5) gives us a different result than the ANCOVA based on the data after mean imputation.

7.6 Looking at a factor-covariate interaction

Let's run a model including the interaction (product) term between `age` and `sex`, which implies that the slope of `age` on our outcome (`st_delta`) depends on the patient's sex. We'll use the imputed data again. Here is the new ANCOVA table, which suggests that the interaction of `age` and `sex` is small (because it accounts for only a small amount of the total Sum of Squares) and not significant (*p* = 0.91).

```
anova(lm(st_delta ~ sex * age, data = emph_imp))
```

Analysis of Variance Table

```
Response: st_delta
  Df Sum Sq Mean Sq F value Pr(>F)
sex      1 52.547 52.547  4.9549 0.04594 *
age      1   2.151   2.151  0.2028 0.66051
sex:age   1   0.130   0.130  0.0123 0.91355
Residuals 12 127.261 10.605
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

Since the interaction term is neither substantial nor significant, we probably don't need it here. But let's look at its interpretation anyway, just to fix ideas. To do that, we'll need the coefficients from the underlying regression model.

```
tidy(lm(st_delta ~ sex * age, data = emph_imp))
```

```
# A tibble: 4 x 5
  term      estimate std.error statistic p.value
  <chr>     <dbl>     <dbl>     <dbl>     <dbl>
1 (Intercept) -5.65     13.5    -0.420    0.682
2 sexM         1.72      16.8     0.102    0.920
3 age          0.0365    0.211    0.173    0.866
4 sexM:age     0.0289    0.260    0.111    0.914
```

Our ANCOVA model for `st_delta` incorporating the `age` x `sex` product term is $-5.65 + 1.72$ (sex = M) + 0.037 age + 0.029 (sex = M)(age). So that means:

- our model for females is `st_delta` = $-5.65 + 0.037$ age
- our model for males is `st_delta` = $(-5.65 + 1.72) + (0.037 + 0.029)$ age, or $-3.93 + 0.066$ age

but, again, our conclusion from the ANCOVA table is that this increase in complexity (letting both the slope and intercept vary by `sex`) doesn't add much in the way of predictive value for our `st_delta` outcome.

7.7 Centering the Covariate to Facilitate ANCOVA Interpretation

When developing an ANCOVA model, we will often **center** or even **center and rescale** the covariate to facilitate interpretation of the product term. In this case, let's center `age` and rescale it by dividing by two standard deviations.

```
emph_imp %$% mosaic::favstats(age)

min   Q1 median   Q3 max mean      sd n missing
50 63.5 65.1 70.25 76 65.2 6.978061 16      0
```

Note that in our imputed data, the mean `age` is 65.2 and the standard deviation of `age` is 7 years.

So we build the rescaled `age` variable that I'll call `age_z`, and then use it to refit our model.

```
emph_imp <- emph_imp %>%
  mutate(age_z = (age - mean(age)) / (2 * sd(age)))

anova(lm(st_delta ~ sex * age_z, data = emph_imp))
```

Analysis of Variance Table

```
Response: st_delta
           Df Sum Sq Mean Sq F value Pr(>F)
sex          1  52.547  52.547  4.9549 0.04594 *
```

```

age_z      1   2.151   2.151  0.2028  0.66051
sex:age_z  1   0.130   0.130  0.0123  0.91355
Residuals 12 127.261 10.605
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
tidy(lm(st_delta ~ sex * age_z, data = emph_imp))

# A tibble: 4 x 5
  term      estimate std.error statistic p.value
  <chr>     <dbl>     <dbl>     <dbl>     <dbl>
1 (Intercept) -3.27      1.39     -2.35     0.0364
2 sexM         3.60      1.74      2.08     0.0601
3 age_z        0.510     2.95      0.173    0.866
4 sexM:age_z   0.403     3.63      0.111    0.914

```

Comparing the two models, we have:

- (unscaled): $st_delta = -5.65 + 1.72 (sex = M) + 0.037 age + 0.029 (sex = M) \times (age)$
- (rescaled): $st_delta = -3.27 + 3.60 (sex = M) + 0.510 \text{ rescaled } age_z + 0.402 (sex = M) \times (\text{rescaled } age_z)$

In essence, the rescaled model on `age_z` is:

- $st_delta = -3.27 + 0.510 age_z$ for female subjects, and
- $st_delta = (-3.27 + 3.60) + (0.510 + 0.402) age_z = 0.33 + 0.912 age_z$ for male subjects

Interpreting the centered, rescaled model, we have:

- no change in the ANOVA results or R-squared or residual standard deviation compared to the uncentered, unscaled model, but
- the intercept (-3.27) now represents the `st_delta` for a female of average age,
- the `sex` slope (3.60) represents the (male - female) difference in predicted `st_delta` for a person of average age,
- the `age_z` slope (0.510) represents the difference in predicted `st_delta` for a female one standard deviation older than the mean age as compared to a female one standard deviation younger than the mean age, and
- the product term's slope (0.402) represents the male - female difference in the slope of `age_z`, so that if you add the `age_z` slope (0.510) and the interaction slope (0.402) you see the difference in predicted `st_delta` for a male one standard deviation older than the mean age as compared to a male one standard deviation younger than the mean age.

Chapter 8

Analysis of Covariance with the SMART data

In this chapter, we'll work with the `smart_cle1_sh` data file again.

```
smart_cle1_sh <- readRDS(here("data", "smart_cle1_sh.Rds"))
```

8.1 A New Small Study: Predicting BMI

We'll begin by investigating the problem of predicting `bmi`, at first with just three regression inputs: `smoke100`, `female` and `physhealth`, in our `smart_cle1_sh` data set.

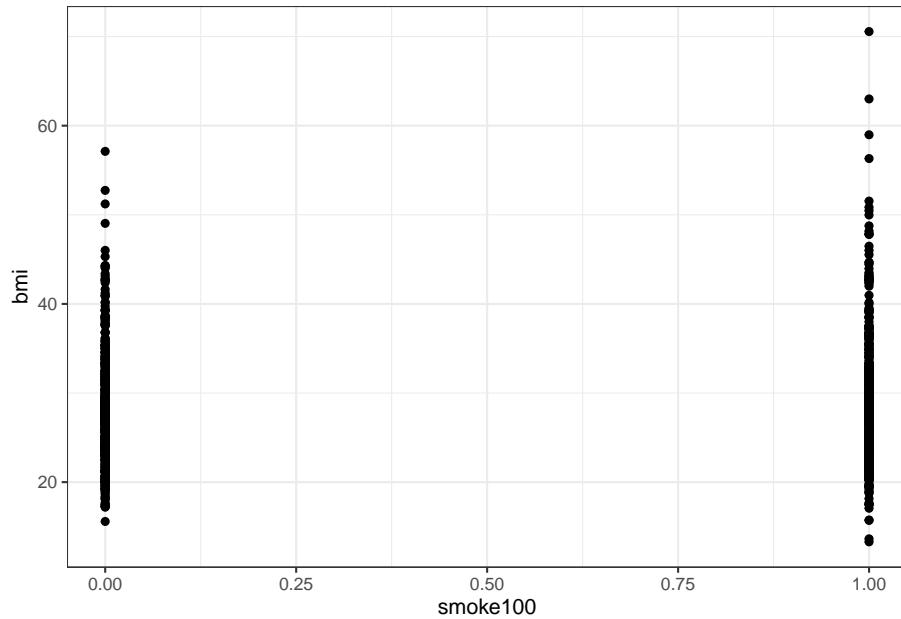
- The outcome of interest is `bmi`.
- Inputs to the regression model are:
 - `smoke100` = 1 if the subject has smoked 100 cigarettes in their lifetime
 - `female` = 1 if the subject is female, and 0 if they are male
 - `physhealth` = number of poor physical health days in past 30 (treated as quantitative)

8.1.1 Does `smoke100` predict `bmi` well?

8.1.1.1 Graphical Assessment

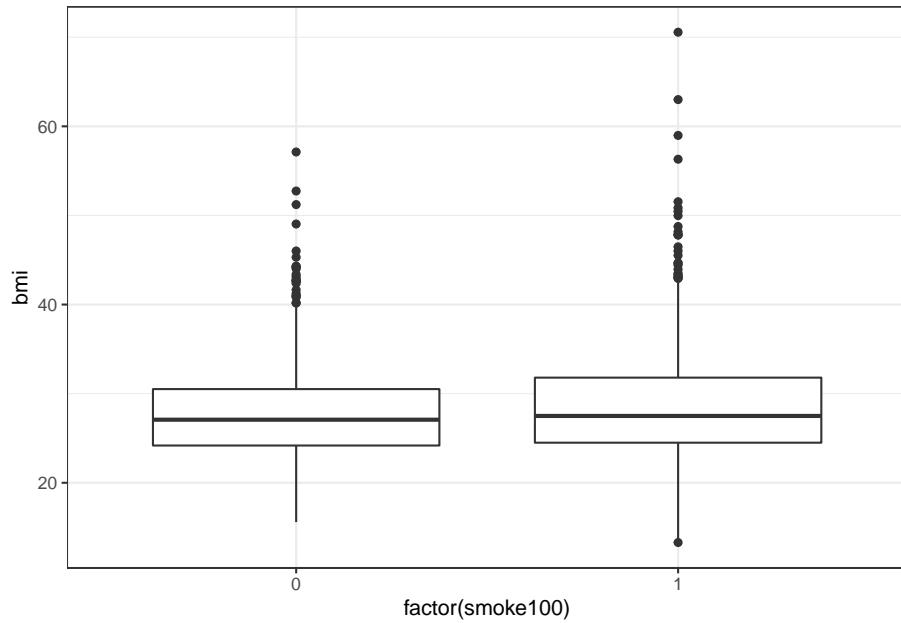
```
ggplot(smart_cle1_sh, aes(x = smoke100, y = bmi)) +  
  geom_point()
```

210 CHAPTER 8. ANALYSIS OF COVARIANCE WITH THE SMART DATA



Not so helpful. We should probably specify that `smoke100` is a factor, and try another plotting approach.

```
ggplot(smart_cle1_sh, aes(x = factor(smoke100), y = bmi)) +  
  geom_boxplot()
```



The median BMI looks a little higher for those who have smoked 100 cigarettes. Let's see if a model reflects that.

8.2 c8_m1: A simple t-test model

```
c8_m1 <- lm(bmi ~ smoke100, data = smart_cle1_sh)
c8_m1

Call:
lm(formula = bmi ~ smoke100, data = smart_cle1_sh)

Coefficients:
(Intercept)      smoke100
              27.9390     0.8722

summary(c8_m1)

Call:
lm(formula = bmi ~ smoke100, data = smart_cle1_sh)

Residuals:
    Min      1Q  Median      3Q      Max 
-15.511 -4.019 -0.870   2.841  41.749 

Coefficients:
            Estimate Std. Error t value Pr(>|t|)    
(Intercept) 27.9390     0.2594 107.722 <2e-16 ***
smoke100    0.8722     0.3781   2.307   0.0213 *  
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 6.353 on 1131 degrees of freedom
Multiple R-squared:  0.004682, Adjusted R-squared:  0.003802 
F-statistic: 5.321 on 1 and 1131 DF,  p-value: 0.02125

confint(c8_m1)

          2.5 %    97.5 %
(Intercept) 27.4301016 28.447875
smoke100    0.1303021  1.614195

tidy(c8_m1, conf.int = TRUE, conf.level = 0.95)

# A tibble: 2 x 7
  term       estimate std.error statistic p.value conf.low conf.high
  <chr>        <dbl>     <dbl>      <dbl>     <dbl>     <dbl>      <dbl>
```

212CHAPTER 8. ANALYSIS OF COVARIANCE WITH THE SMART DATA

```

1 (Intercept) 27.9      0.259     108.      0      27.4      28.4
2 smoke100    0.872      0.378     2.31   0.0213    0.130     1.61
glance(c8_m1)

# A tibble: 1 x 12
  r.squared adj.r.squared sigma statistic p.value    df logLik    AIC    BIC
      <dbl>        <dbl>    <dbl>     <dbl>    <dbl>    <dbl>    <dbl>    <dbl>
1 0.00468      0.00380  6.35     5.32  0.0213     1 -3702. 7409. 7424.
# ... with 3 more variables: deviance <dbl>, df.residual <int>, nobs <int>

```

The model suggests, based on these 1133 subjects, that

- our best prediction for non-smokers is $BMI = 27.94 \text{ kg/m}^2$, and
- our best prediction for those who have smoked 100 cigarettes is $BMI = 27.94 + 0.87 = 28.81 \text{ kg/m}^2$.
- the mean difference between smokers and non-smokers is $+0.87 \text{ kg/m}^2$ in BMI
- a 95% confidence (uncertainty) interval for that mean smoker - non-smoker difference in BMI ranges from 0.13 to 1.61
- the model accounts for 0.47% of the variation in BMI, so that knowing the respondent's tobacco status does very little to reduce the size of the prediction errors as compared to an intercept only model that would predict the overall mean BMI for each of our subjects.
- the model makes some enormous errors, with one subject being predicted to have a BMI 42 points lower than his/her actual BMI.

Note that this simple regression model just gives us the t-test.

```
t.test(bmi ~ smoke100, var.equal = TRUE, data = smart_cle1_sh)
```

Two Sample t-test

```

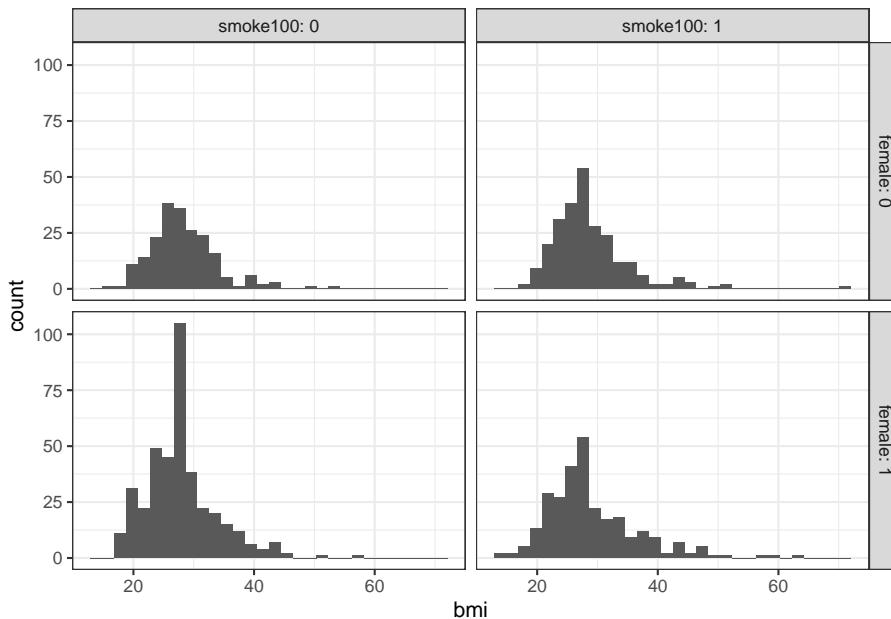
data: bmi by smoke100
t = -2.3066, df = 1131, p-value = 0.02125
alternative hypothesis: true difference in means between group 0 and group 1 is not eq
95 percent confidence interval:
-1.6141946 -0.1303021
sample estimates:
mean in group 0 mean in group 1
27.93899      28.81124

```

8.3 c8_m2: Adding another predictor (two-way ANOVA without interaction)

When we add in the information about `female` (sex) to our original model, we might first picture the data. We could look at separate histograms,

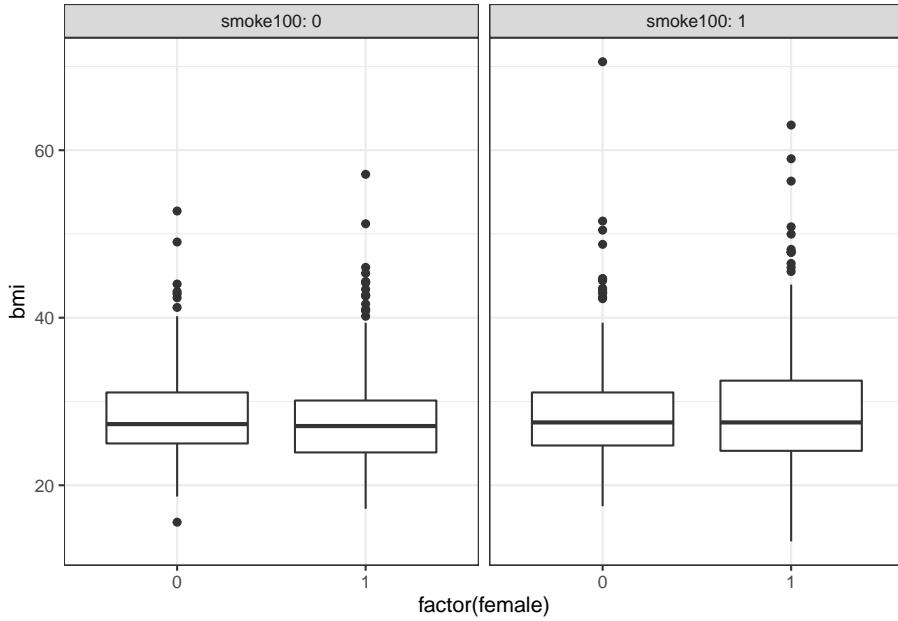
```
ggplot(smart_cle1_sh, aes(x = bmi)) +
  geom_histogram(bins = 30) +
  facet_grid(female ~ smoke100, labeller = label_both)
```



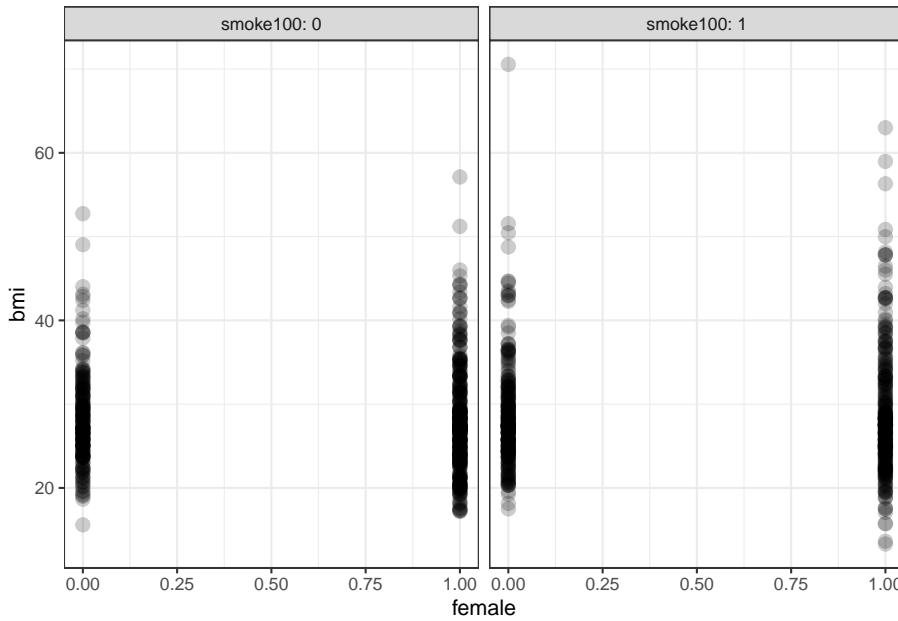
or maybe boxplots?

```
ggplot(smart_cle1_sh, aes(x = factor(female), y = bmi)) +
  geom_boxplot() +
  facet_wrap(~ smoke100, labeller = label_both)
```

214 CHAPTER 8. ANALYSIS OF COVARIANCE WITH THE SMART DATA



```
ggplot(smart_cle1_sh, aes(x = female, y = bmi))+
  geom_point(size = 3, alpha = 0.2) +
  theme_bw() +
  facet_wrap(~ smoke100, labeller = label_both)
```



OK. Let's try fitting a model.

```
c8_m2 <- lm(bmi ~ female + smoke100, data = smart_cle1_sh)
c8_m2
```

Call:

```
lm(formula = bmi ~ female + smoke100, data = smart_cle1_sh)
```

Coefficients:

(Intercept)	female	smoke100
28.0265	-0.1342	0.8555

This new model predicts only four predicted values:

- $bmi = 28.0265$ if the subject is male and has not smoked 100 cigarettes (so `female` = 0 and `smoke100` = 0)
- $bmi = 28.0265 - 0.1342 = 27.8923$ if the subject is female and has not smoked 100 cigarettes (`female` = 1 and `smoke100` = 0)
- $bmi = 28.0265 + 0.8555 = 28.8820$ if the subject is male and has smoked 100 cigarettes (so `female` = 0 and `smoke100` = 1), and, finally
- $bmi = 28.0265 - 0.1342 + 0.8555 = 28.7478$ if the subject is female and has smoked 100 cigarettes (so both `female` and `smoke100` = 1).

Another way to put this is that for those who have not smoked 100 cigarettes, the model is:

- $bmi = 28.0265 - 0.1342 \text{ female}$

and for those who have smoked 100 cigarettes, the model is:

- $bmi = 28.8820 - 0.1342 \text{ female}$

Only the intercept of the `bmi-female` model changes depending on `smoke100`.

```
summary(c8_m2)
```

Call:

```
lm(formula = bmi ~ female + smoke100, data = smart_cle1_sh)
```

Residuals:

Min	1Q	Median	3Q	Max
-15.448	-3.972	-0.823	2.774	41.678

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	28.0265	0.3620	77.411	<2e-16 ***
female	-0.1342	0.3875	-0.346	0.7291
smoke100	0.8555	0.3814	2.243	0.0251 *

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

216CHAPTER 8. ANALYSIS OF COVARIANCE WITH THE SMART DATA

```
Residual standard error: 6.356 on 1130 degrees of freedom
Multiple R-squared:  0.004788, Adjusted R-squared:  0.003027
F-statistic: 2.718 on 2 and 1130 DF,  p-value: 0.06642
confint(c8_m2)
```

```
          2.5 %    97.5 %
(Intercept) 27.3161140 28.7368281
female      -0.8944773  0.6259881
smoke100    0.1072974  1.6037825
```

The slopes of both `female` and `smoke100` have confidence intervals that are completely below zero, indicating that both `female` sex and `smoke100` appear to be associated with reductions in `bmi`.

The R^2 value suggests that 0.4788% of the variation in `bmi` is accounted for by this ANOVA model.

In fact, this regression (on two binary indicator variables) is simply a two-way ANOVA model without an interaction term.

```
anova(c8_m2)
```

Analysis of Variance Table

```
Response: bmi
           Df Sum Sq Mean Sq F value Pr(>F)
female       1   16    16.301  0.4036 0.52538
smoke100    1   203   203.296  5.0330 0.02506 *
Residuals 1130  45644   40.393
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

8.4 c8_m3: Adding the interaction term (Two-way ANOVA with interaction)

Suppose we want to let the effect of `female` vary depending on the `smoke100` status. Then we need to incorporate an interaction term in our model.

```
c8_m3 <- lm(bmi ~ female * smoke100, data = smart_cle1_sh)
c8_m3
```

```
Call:
lm(formula = bmi ~ female * smoke100, data = smart_cle1_sh)
```

```
Coefficients:
(Intercept)        female        smoke100  female:smoke100
```

28.2690	-0.5064	0.4119	0.7536
---------	---------	--------	--------

So, for example, for a male who has smoked 100 cigarettes, this model predicts

- $bmi = 28.269 - 0.506(0) + 0.412(1) + 0.754(0)(1) = 28.269 + 0.412 = 28.681$

And for a female who has smoked 100 cigarettes, the model predicts

- $bmi = 28.269 - 0.506(1) + 0.412(1) + 0.754(1)(1) = 28.269 - 0.506 + 0.412 + 0.754 = 28.929$

For those who have not smoked 100 cigarettes, the model is:

- $bmi = 28.269 - 0.506 \text{female}$

But for those who have smoked 100 cigarettes, the model is:

- $bmi = (28.269 + 0.412) + (-0.506 + 0.754) \text{female}$, or „,
- $bmi = 28.681 - 0.248 \text{female}$

Now, both the slope and the intercept of the `bmi-female` model change depending on `smoke100`.

```
summary(c8_m3)
```

Call:

```
lm(formula = bmi ~ female * smoke100, data = smart_cle1_sh)
```

Residuals:

Min	1Q	Median	3Q	Max
-15.628	-3.938	-0.829	2.759	41.879

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	28.2690	0.4396	64.301	<2e-16 ***
female	-0.5064	0.5446	-0.930	0.353
smoke100	0.4119	0.5946	0.693	0.489
female:smoke100	0.7536	0.7750	0.972	0.331

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 6.356 on 1129 degrees of freedom

Multiple R-squared: 0.005621, Adjusted R-squared: 0.002979

F-statistic: 2.127 on 3 and 1129 DF, p-value: 0.09507

```
confint(c8_m3)
```

	2.5 %	97.5 %
(Intercept)	27.4063783	29.1315563
female	-1.5749026	0.5621793
smoke100	-0.7547605	1.5786121

```
female:smoke100 -0.7670239 2.2742178
```

In fact, this regression (on two binary indicator variables and a product term) is simply a two-way ANOVA model with an interaction term.

```
anova(c8_m3)
```

Analysis of Variance Table

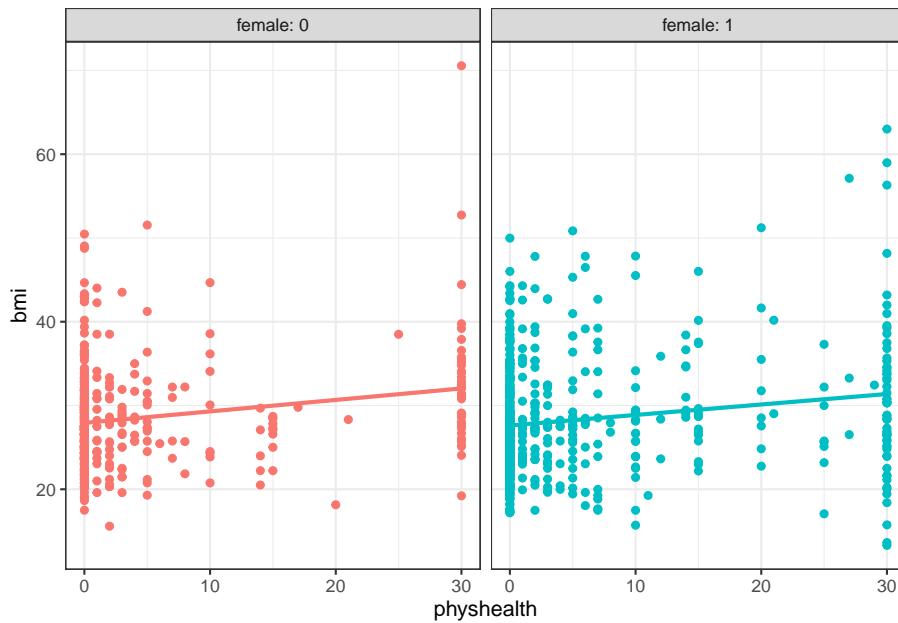
```
Response: bmi
          Df Sum Sq Mean Sq F value Pr(>F)
female        1    16   16.301  0.4035 0.52539
smoke100      1   203  203.296  5.0327 0.02507 *
female:smoke100  1     38   38.194  0.9455 0.33107
Residuals    1129 45606  40.395
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

The interaction term doesn't change very much here. Its uncertainty interval includes zero, and the overall model still accounts for less than 1% of the variation in `bmi`.

8.5 c8_m4: Using `female` and `physhealth` in a model for `bmi`

```
ggplot(smart_cle1_sh, aes(x = physhealth, y = bmi, color = factor(female))) +
  geom_point() +
  guides(col = "none") +
  geom_smooth(method = "lm", se = FALSE) +
  facet_wrap(~ female, labeller = label_both)

`geom_smooth()` using formula 'y ~ x'
```



Does the difference in slopes of bmi and physhealth for males and females appear to be substantial and important?

```
c8_m4 <- lm(bmi ~ female * physhealth, data = smart_cle1_sh)
```

```
summary(c8_m4)
```

Call:

```
lm(formula = bmi ~ female * physhealth, data = smart_cle1_sh)
```

Residuals:

Min	1Q	Median	3Q	Max
-18.069	-3.825	-0.624	2.516	38.526

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	27.92386	0.32196	86.731	< 2e-16 ***
female	-0.30335	0.42346	-0.716	0.474
physhealth	0.13700	0.03277	4.180	3.14e-05 ***
female:physhealth	-0.01203	0.04191	-0.287	0.774

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 6.262 on 1129 degrees of freedom

Multiple R-squared: 0.03486, Adjusted R-squared: 0.03229

F-statistic: 13.59 on 3 and 1129 DF, p-value: 1.027e-08

220CHAPTER 8. ANALYSIS OF COVARIANCE WITH THE SMART DATA

Does it seem as though the addition of `physhealth` has improved our model substantially over a model with `female` alone (which, you recall, was `c8_m1`)?

Since the `c8_m4` model contains the `c8_m1` model's predictors as a subset and the outcome is the same for each model, we consider the models *nested* and have some extra tools available to compare them.

- I might start by looking at the basic summaries for each model.

```
glance(c8_m4)
```

```
# A tibble: 1 x 12
  r.squared adj.r.squared sigma statistic      p.value      df logLik     AIC     BIC
  <dbl>         <dbl> <dbl>      <dbl>        <dbl> <dbl> <dbl> <dbl> <dbl>
1 0.0349       0.0323  6.26      13.6 0.0000000103     3 -3684. 7378. 7403.
# ... with 3 more variables: deviance <dbl>, df.residual <int>, nobs <int>
```

```
glance(c8_m1)
```

```
# A tibble: 1 x 12
  r.squared adj.r.squared sigma statistic p.value      df logLik     AIC     BIC
  <dbl>         <dbl> <dbl>      <dbl>        <dbl> <dbl> <dbl> <dbl>
1 0.00468      0.00380  6.35      5.32 0.0213      1 -3702. 7409. 7424.
# ... with 3 more variables: deviance <dbl>, df.residual <int>, nobs <int>
```

- The R^2 is much larger for the model with `physhealth`, but still very tiny.
- Smaller AIC and smaller BIC statistics are more desirable. Here, there's little to choose from, so `c8_m4` looks better, too.
- We might also consider a significance test by looking at an ANOVA model comparison. This is only appropriate because `c8_m1` is nested in `c8_m4`.

```
anova(c8_m4, c8_m1)
```

Analysis of Variance Table

```
Model 1: bmi ~ female * physhealth
Model 2: bmi ~ smoke100
  Res.Df   RSS Df Sum of Sq    F    Pr(>F)
1   1129 44265
2   1131 45649 -2     -1384 17.649 2.834e-08 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

The addition of the `physhealth` term appears to be a statistically detectable improvement, not that that means very much.

8.6 Making Predictions with a Linear Regression Model

Recall model 4, which yields predictions for body mass index on the basis of the main effects of sex (`female`) and days of poor physical health (`physhealth`) and their interaction.

```
c8_m4
```

```
Call:
lm(formula = bmi ~ female * physhealth, data = smart_cle1_sh)

Coefficients:
            (Intercept)          female      physhealth female:physhealth
              27.92386           -0.30335            0.13700            -0.01203
```

8.6.1 Fitting an Individual Prediction and 95% Prediction Interval

What do we predict for the `bmi` of a subject who is `female` and had 8 poor physical health days in the past 30?

```
c8_new1 <- tibble(female = 1, physhealth = 8)
predict(c8_m4, newdata = c8_new1, interval = "prediction", level = 0.95)

    fit      lwr      upr
1 28.62022 16.32454 40.9159
```

The predicted `bmi` for this new subject is shown above. The prediction interval shows the bounds of a 95% uncertainty interval for a predicted `bmi` for an individual female subject who has 8 days of poor physical health out of the past 30. From the `predict` function applied to a linear model, we can get the prediction intervals for any new data points in this manner.

8.6.2 Confidence Interval for an Average Prediction

- What do we predict for the **average body mass index of a population of subjects** who are female and have `physhealth = 8`?

```
predict(c8_m4, newdata = c8_new1, interval = "confidence", level = 0.95)

    fit      lwr      upr
1 28.62022 28.12256 29.11788
```

- How does this result compare to the prediction interval?

8.6.3 Fitting Multiple Individual Predictions to New Data

- How does our prediction change for a respondent if they instead have 7, or 9 poor physical health days? What if they are male, instead of female?

```
c8_new2 <- tibble(subjectid = 1001:1006, female = c(1, 1, 1, 0, 0, 0), physhealth = c(7, 8, 9, 7, 8, 9))
pred2 <- predict(c8_m4, newdata = c8_new2, interval = "prediction", level = 0.95) %>%
  bind_cols(c8_new2, pred2)
result2
```

	subjectid	female	physhealth	fit	lwr	upr
	<int>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>
1	1001	1	7	28.5	16.2	40.8
2	1002	1	8	28.6	16.3	40.9
3	1003	1	9	28.7	16.4	41.0
4	1004	0	7	28.9	16.6	41.2
5	1005	0	8	29.0	16.7	41.3
6	1006	0	9	29.2	16.9	41.5

The `result2` tibble contains predictions for each scenario.

- Which has a bigger impact on these predictions and prediction intervals? A one category change in `female` or a one hour change in `physhealth`?

8.7 Centering the model

Our model `c8_m4` has four predictors (the constant, `physhealth`, `female` and their interaction) but just two inputs (`female` and `physhealth`.) If we center the quantitative input `physhealth` before building the model, we get a more interpretable interaction term.

```
smart_cle1_sh_c <- smart_cle1_sh %>%
  mutate(physhealth_c = physhealth - mean(physhealth))

c8_m4_c <- lm(bmi ~ female * physhealth_c, data = smart_cle1_sh_c)

summary(c8_m4_c)
```

Call:

```
lm(formula = bmi ~ female * physhealth_c, data = smart_cle1_sh_c)
```

Residuals:

Min	1Q	Median	3Q	Max
-18.069	-3.825	-0.624	2.516	38.526

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	28.56520	0.29213	97.784	< 2e-16 ***
female	-0.35969	0.37917	-0.949	0.343
physhealth_c	0.13700	0.03277	4.180	3.14e-05 ***
female:physhealth_c	-0.01203	0.04191	-0.287	0.774

Signif. codes:	0 ***	0.001 **	0.01 *	0.05 .
	'	'	'	'

Residual standard error: 6.262 on 1129 degrees of freedom

Multiple R-squared: 0.03486, Adjusted R-squared: 0.03229

F-statistic: 13.59 on 3 and 1129 DF, p-value: 1.027e-08

What has changed as compared to the original c8_m4?

- Our original model was $bmi = 27.92 - 0.30 \text{ female} + 0.14 \text{ physhealth} - 0.01 \text{ female} \times \text{physhealth}$
- Our new model is $bmi = 28.57 - 0.36 \text{ female} + 0.14 \text{ centered physhealth} - 0.01 \text{ female} \times \text{centered physhealth}$.

So our new model on centered data is:

- $28.57 + 0.14 \text{ centered physhealth_c}$ for male subjects, and
- $(28.57 - 0.30) + (0.14 - 0.01) \text{ centered physhealth_c}$, or $28.27 - 0.13 \text{ centered physhealth_c}$ for female subjects.

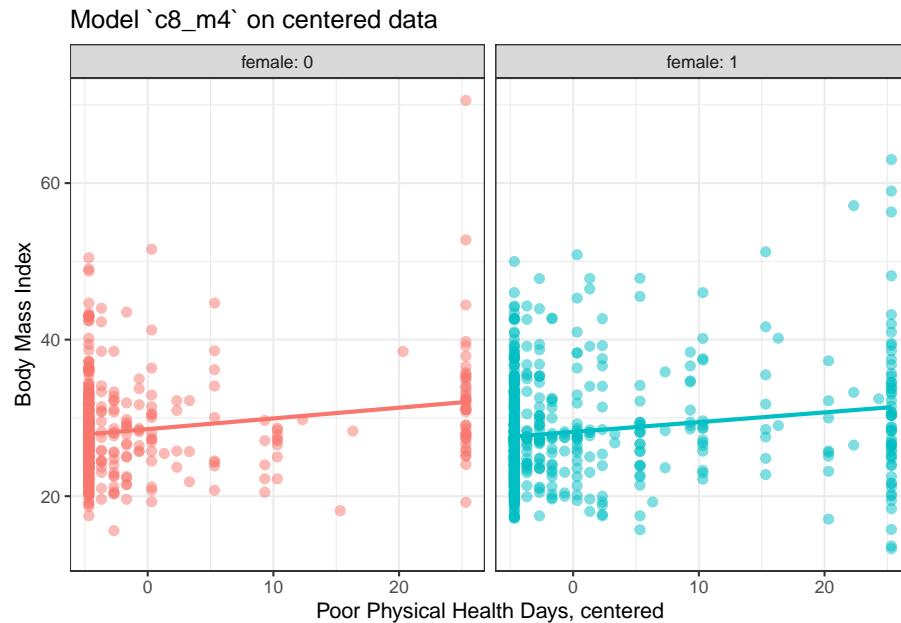
In our new (centered physhealth_c) model,

- the main effect of **female** now corresponds to a predictive difference (female - male) in **bmi** with **physhealth** at its mean value, 4.68 days,
- the intercept term is now the predicted **bmi** for a male respondent with an average **physhealth**, and
- the product term corresponds to the change in the slope of centered **physhealth_c** on **bmi** for a female rather than a male subject, while
- the residual standard deviation and the R-squared values remain unchanged from the model before centering.

8.7.1 Plot of Model 4 on Centered physhealth: c8_m4_c

```
ggplot(smart_cle1_sh_c, aes(x = physhealth_c, y = bmi, group = female, col = factor(female))) +
  geom_point(alpha = 0.5, size = 2) +
  geom_smooth(method = "lm", se = FALSE) +
  guides(color = "none") +
  labs(x = "Poor Physical Health Days, centered", y = "Body Mass Index",
       title = "Model `c8_m4` on centered data") +
  facet_wrap(~ female, labeller = label_both)
```

```
`geom_smooth()` using formula 'y ~ x'
```



8.8 Rescaling an input by subtracting the mean and dividing by 2 standard deviations

Centering helped us interpret the main effects in the regression, but it still leaves a scaling problem.

- The `female` coefficient estimate is much larger than that of `physhealth`, but this is misleading, considering that we are comparing the complete change in one variable (sex = female or not) to a 1-day change in `physhealth`.
- Gelman and Hill (2007) recommend all continuous predictors be scaled by dividing by 2 standard deviations, so that:
 - a 1-unit change in the rescaled predictor corresponds to a change from 1 standard deviation below the mean, to 1 standard deviation above.
 - an unscaled binary (1/0) predictor with 50% probability of occurring will be exactly comparable to a rescaled continuous predictor done in this way.

```
smart_cle1_sh_rescale <- smart_cle1_sh %>%
  mutate(physhealth_z = (physhealth - mean(physhealth))/(2*sd(physhealth)))
```

8.8. RESCALING AN INPUT BY SUBTRACTING THE MEAN AND DIVIDING BY 2 STANDARD DEVIATIONS

8.8.1 Refitting model c8_m4 to the rescaled data

```
c8_m4_z <- lm(bmi ~ female * physhealth_z, data = smart_cle1_sh_rescale)

summary(c8_m4_z)

Call:
lm(formula = bmi ~ female * physhealth_z, data = smart_cle1_sh_rescale)

Residuals:
    Min      1Q  Median      3Q     Max 
-18.069 -3.825 -0.624  2.516 38.526 

Coefficients:
            Estimate Std. Error t value Pr(>|t|)    
(Intercept) 28.5652    0.2921 97.784 < 2e-16 ***
female      -0.3597    0.3792 -0.949   0.343    
physhealth_z  2.4991    0.5978  4.180 3.14e-05 ***
female:physhealth_z -0.2195    0.7645 -0.287   0.774    
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 6.262 on 1129 degrees of freedom
Multiple R-squared:  0.03486, Adjusted R-squared:  0.03229 
F-statistic: 13.59 on 3 and 1129 DF,  p-value: 1.027e-08
```

8.8.2 Interpreting the model on rescaled data

What has changed as compared to the original c8_m4?

- Our original model was $bmi = 27.92 - 0.30 \text{female} + 0.14 \text{physhealth} - 0.01 \text{female} \times \text{physhealth}$
- Our model on centered physhealth was $bmi = 28.57 - 0.36 \text{female} + 0.14 \text{centered physhealth} - 0.01 \text{female} \times \text{centered physhealth}$.
- Our new model on rescaled physhealth is $bmi = 28.58 - 0.37 \text{female} + 2.50 \text{rescaled physhealth}_z - 0.22 \text{female} \times \text{rescaled physhealth}_z$.

So our rescaled model is:

- $28.58 + 2.50 \text{rescaled physhealth}_z$ for male subjects, and
- $(28.58 - 0.37) + (2.50 - 0.22) \text{rescaled physhealth}_z$, or $28.21 + 2.28 \text{rescaled physhealth}_z$ for female subjects.

In this new rescaled (physhealth_z) model, then,

- the main effect of `female`, -0.37, still corresponds to a predictive difference (`female` - male) in `bmi` with `physhealth` at its mean value, 4.68 days,

- the intercept term is still the predicted `bmi` for a male respondent with an average `physhealth` count, and
- the residual standard deviation and the R-squared values remain unchanged,

as before, but now we also have that:

- the coefficient of `physhealth_z` indicates the predictive difference in `bmi` associated with a change in `physhealth` of 2 standard deviations (from one standard deviation below the mean of 4.68 to one standard deviation above 4.68.)
 - Since the standard deviation of `physhealth` is 9.12 (see below), this covers a massive range of potential values of `physhealth` from 0 all the way up to $4.68 + 2(9.12) = 22.92$ days.

```
mosaic::favstats(~ physhealth, data = smart_cle1_sh)
```

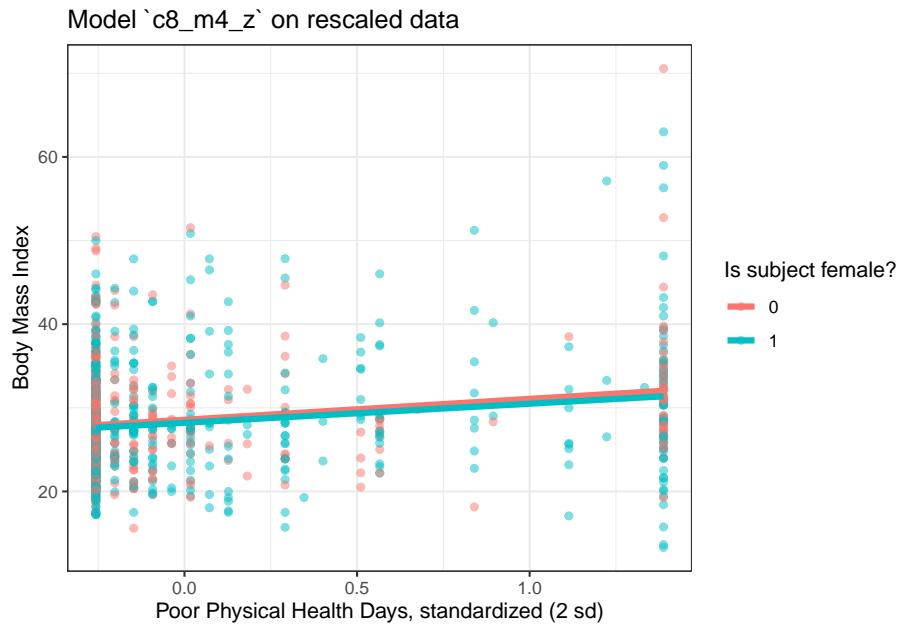
	min	Q1	median	Q3	max	mean	sd	n	missing
	0	0	0	4	30	4.681377	9.120899	1133	0

- the coefficient of the product term (-0.22) corresponds to the change in the coefficient of `physhealth_z` for females as compared to males.

8.8.3 Plot of model on rescaled data

```
ggplot(smart_cle1_sh_rescale, aes(x = physhealth_z, y = bmi,
                                    group = female, col = factor(female))) +
  geom_point(alpha = 0.5) +
  geom_smooth(method = "lm", se = FALSE, size = 1.5) +
  scale_color_discrete(name = "Is subject female?")
  labs(x = "Poor Physical Health Days, standardized (2 sd)", y = "Body Mass Index",
       title = "Model `c8_m4_z` on rescaled data")
```

`geom_smooth()` using formula 'y ~ x'



There's very little difference here.

8.9 c8_m5: What if we add more variables?

We can boost our R^2 a bit, to nearly 5%, by adding in two new variables, related to whether or not the subject (in the past 30 days) used the internet, and the average number of alcoholic drinks per week consumed by the subject.

```
c8_m5 <- lm(bmi ~ female + smoke100 + physhealth + internet30 + drinks_wk,
               data = smart_cle1_sh)
summary(c8_m5)
```

Call:

```
lm(formula = bmi ~ female + smoke100 + physhealth + internet30 +
    drinks_wk, data = smart_cle1_sh)
```

Residuals:

Min	1Q	Median	3Q	Max
-18.358	-3.846	-0.657	2.534	38.049

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	27.52400	0.56076	49.083	< 2e-16 ***
female	-0.43272	0.38510	-1.124	0.26140

228CHAPTER 8. ANALYSIS OF COVARIANCE WITH THE SMART DATA

```

smoke100      0.82654    0.37739    2.190  0.02872 *
physhealth    0.12469    0.02074    6.012 2.47e-09 ***
internet30   0.44287    0.48830    0.907  0.36462
drinks_wk     -0.10193   0.03352   -3.041  0.00241 **
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

```

Residual standard error: 6.231 on 1127 degrees of freedom
Multiple R-squared: 0.04582, Adjusted R-squared: 0.04159
F-statistic: 10.82 on 5 and 1127 DF, p-value: 3.48e-10

1. Here's the ANOVA for this model. What can we study with this?

```
anova(c8_m5)
```

Analysis of Variance Table

```

Response: bmi
          Df Sum Sq Mean Sq F value    Pr(>F)
female      1    16   16.30  0.4198  0.517171
smoke100    1   203  203.30  5.2354  0.022316 *
physhealth   1  1508 1508.08 38.8372 6.497e-10 ***
internet30  1    15   14.69  0.3783  0.538650
drinks_wk    1   359  359.05  9.2466  0.002414 **
Residuals   1127 43762   38.83
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

```

2. Consider the revised output below. Now what can we study?

```
anova(lm(bmi ~ smoke100 + internet30 + drinks_wk + female + physhealth,
        data = smart_cle1_sh))
```

Analysis of Variance Table

```

Response: bmi
          Df Sum Sq Mean Sq F value    Pr(>F)
smoke100     1    215  214.75  5.5304  0.0188606 *
internet30   1      8    7.81  0.2010  0.6539723
drinks_wk     1    444  443.79 11.4288  0.0007479 ***
female       1     32   31.58  0.8132  0.3673566
physhealth    1   1403 1403.49 36.1438 2.472e-09 ***
Residuals   1127 43762   38.83
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

```

3. What does the output below let us conclude?

```
anova(lm(bmi ~ smoke100 + internet30 + drinks_wk + female + physhealth,
         data = smart_cle1_sh),
      lm(bmi ~ smoke100 + female + drinks_wk,
         data = smart_cle1_sh))
```

Analysis of Variance Table

```
Model 1: bmi ~ smoke100 + internet30 + drinks_wk + female + physhealth
Model 2: bmi ~ smoke100 + female + drinks_wk
Res.Df   RSS Df Sum of Sq    F    Pr(>F)
1     1127 43762
2     1129 45166 -2    -1403.7 18.075 1.877e-08 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

4. What does it mean for the models to be “nested”?

8.10 c8_m6: Would adding self-reported health help?

And we can do even a bit better than that by adding in a multi-categorical measure: self-reported general health.

```
c8_m6 <- lm(bmi ~ female + smoke100 + physhealth + internet30 + drinks_wk + genhealth,
             data = smart_cle1_sh)
summary(c8_m6)

Call:
lm(formula = bmi ~ female + smoke100 + physhealth + internet30 +
drinks_wk + genhealth, data = smart_cle1_sh)

Residuals:
    Min      1Q  Median      3Q      Max 
-19.216  -3.659  -0.736   2.669  36.810 

Coefficients:
            Estimate Std. Error t value Pr(>|t|)    
(Intercept) 25.20736   0.71106  35.450 < 2e-16 ***
female      -0.31949   0.37667  -0.848  0.3965  
smoke100     0.45866   0.37214   1.232  0.2180  
physhealth    0.04353   0.02506   1.737  0.0827  
internet30   0.93270   0.48273   1.932  0.0536  
drinks_wk     -0.07712   0.03294  -2.341  0.0194 *  
genhealth2_VeryGood 1.21169   0.56838   2.132  0.0332 * 

```

230CHAPTER 8. ANALYSIS OF COVARIANCE WITH THE SMART DATA

```

genhealth3_Good      3.22783    0.58009    5.564 3.29e-08 ***
genhealth4_Fair       4.14497    0.73284    5.656 1.96e-08 ***
genhealth5_Poor       5.86335    1.09253    5.367 9.73e-08 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

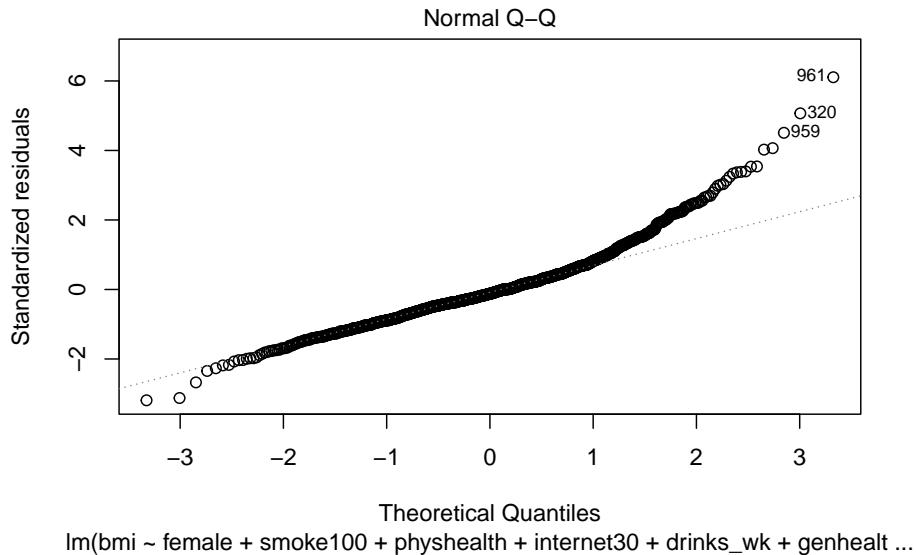
```

Residual standard error: 6.089 on 1123 degrees of freedom
Multiple R-squared: 0.09206, Adjusted R-squared: 0.08478
F-statistic: 12.65 on 9 and 1123 DF, p-value: < 2.2e-16

1. If Harry and Marty have the same values of `female`, `smoke100`, `physhealth`, `internet30` and `drinks_wk`, but Harry rates his health as Good, and Marty rates his as Fair, then what is the difference in the predictions? Who is predicted to have a larger BMI, and by how much?

2. What does this normal probability plot of the residuals suggest?

```
plot(c8_m6, which = 2)
```



8.11 Key Regression Assumptions for Building Effective Prediction Models

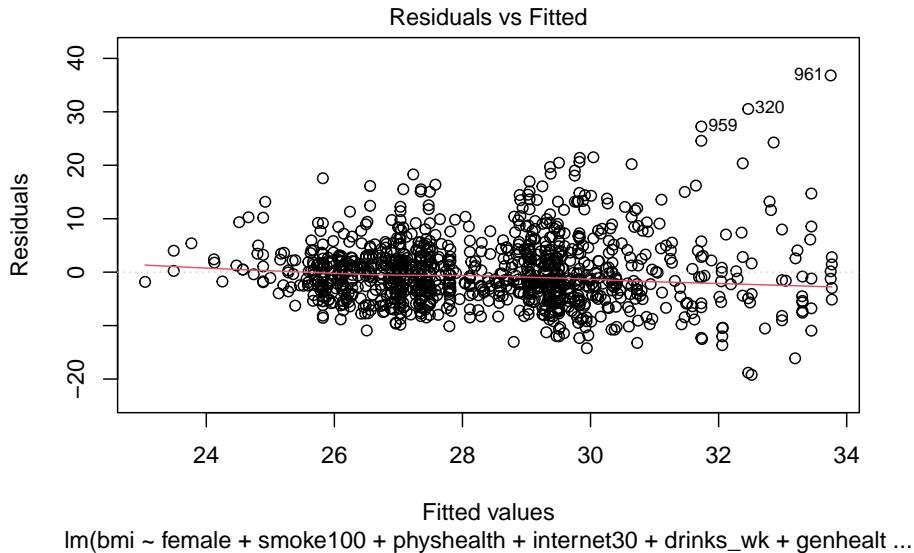
1. Validity - the data you are analyzing should map to the research question you are trying to answer.

- The outcome should accurately reflect the phenomenon of interest.
 - The model should include all relevant predictors. (It can be difficult to decide which predictors are necessary, and what to do with predictors that have large standard errors.)
 - The model should generalize to all of the cases to which it will be applied.
 - Can the available data answer our question reliably?
2. Additivity and linearity - most important assumption of a regression model is that its deterministic component is a linear function of the predictors. We often think about transformations in this setting.
 3. Independence of errors - errors from the prediction line are independent of each other
 4. Equal variance of errors - if this is violated, we can more efficiently estimate parameters using *weighted least squares* approaches, where each point is weighted inversely proportional to its variance, but this doesn't affect the coefficients much, if at all.
 5. Normality of errors - not generally important for estimating the regression line

8.11.1 Checking Assumptions in model c8_m6

1. How does the assumption of linearity behind this model look?

```
plot(c8_m6, which = 1)
```

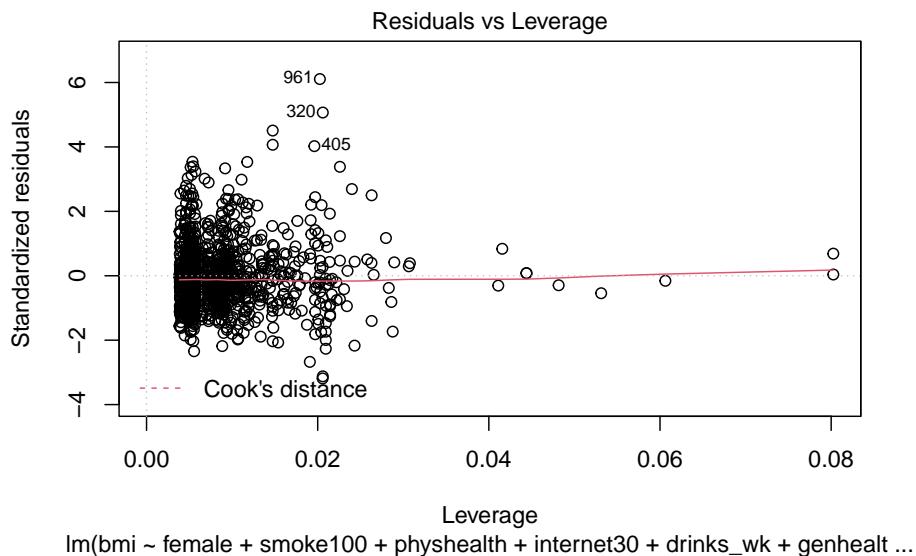


232CHAPTER 8. ANALYSIS OF COVARIANCE WITH THE SMART DATA

We see no strong signs of serious non-linearity here. There's no obvious curve in the plot, for example. We may have a problem with increasing variance as we move to the right.

- What can we conclude from the plot below?

```
plot(c8_m6, which = 5)
```



This plot can help us identify points with large standardized residuals, large leverage values, and large influence on the model (as indicated by large values of Cook's distance.) In this case, I see no signs of any points used in the model with especially large influence, although there are some poorly fitted points (with especially large standardized residuals.)

We might want to identify the point listed here as 961, which appears to have an enormous standardized residual. To do so, we can use the `slice` function from `dplyr`.

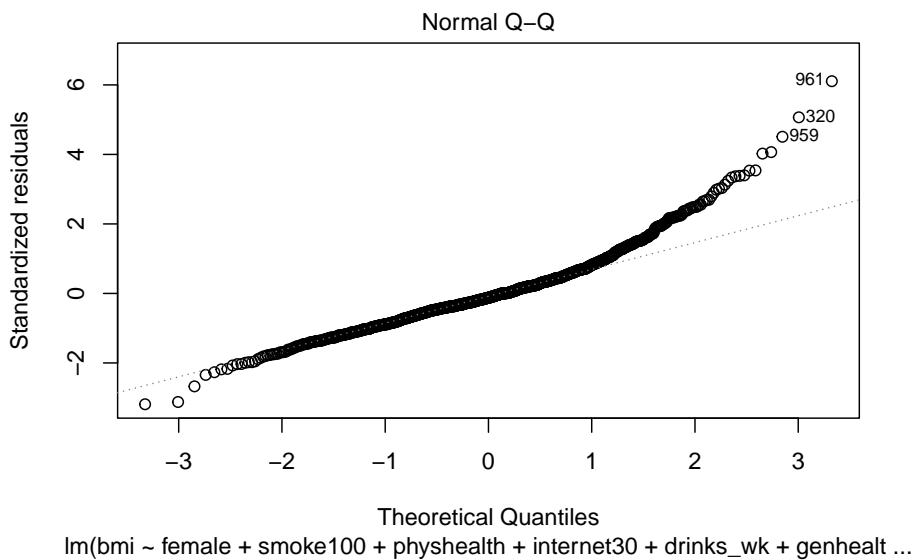
```
smart_cle1_sh %>% slice(961) %>% select(SEQNO)
```

```
# A tibble: 1 x 1
  SEQNO
  <dbl>
1 2017000961
```

Now we know exactly which subject we're talking about.

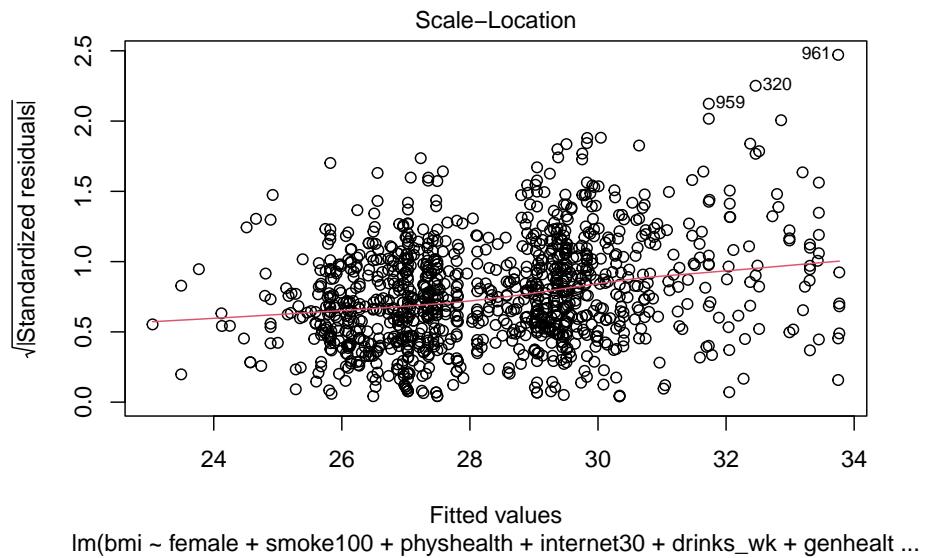
- What other residual plots are available with `plot` and how do we interpret them?

```
plot(c8_m6, which = 2)
```



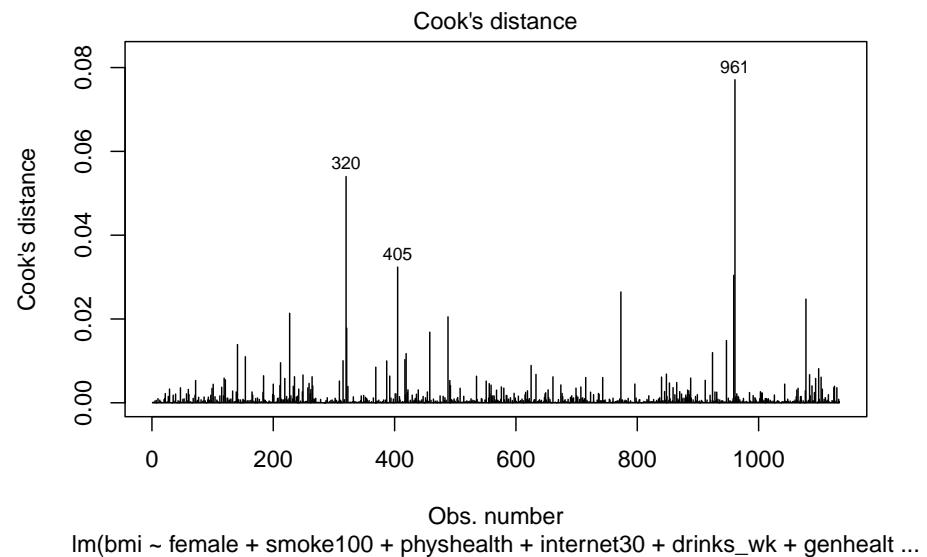
This plot is simply a Normal Q-Q plot of the standardized residuals from our model. We're looking here for serious problems with the assumption of Normality.

```
plot(c8_m6, which = 3)
```



This is a scale-location plot, designed to help us see non-constant variance in the residuals as we move across the fitted values as a linear trend, rather than as a fan shape, by plotting the square root of the residuals on the vertical axis.

```
plot(c8_m6, which = 4)
```



Finally, this is an index plot of the Cook's distance values, allowing us to identify points that are particularly large. Remember that a value of 0.5 (or perhaps even 1.0) is a reasonable boundary for a substantially influential point.

Chapter 9

Adding Non-linear Terms to a Linear Regression Model

9.1 The pollution data

Consider the `pollution` data set, which contain 15 independent variables and a measure of mortality, describing 60 US metropolitan areas in 1959-1961. The data come from McDonald and Schwing (1973), and are available at <http://www4.stat.ncsu.edu/~boos/var.select/pollution.html> and our web site.

```
pollution
```

```
# A tibble: 60 x 16
  x1    x2    x3    x4    x5    x6    x7    x8    x9    x10   x11   x12   x13
  <dbl> <dbl>
1 36    27    71    8.1   3.34  11.4   81.5  3243   8.8   42.6  11.7   21    15
2 35    23    72   11.1   3.14   11    78.8  4281   3.5   50.7  14.4    8    10
3 44    29    74   10.4   3.21   9.8   81.6  4260   0.8   39.4  12.4    6     6
4 47    45    79    6.5   3.41  11.1   77.5  3125   27.1  50.2  20.6   18     8
5 43    35    77    7.6   3.44   9.6   84.6  6441   24.4  43.7  14.3   43    38
6 53    45    80    7.7   3.45  10.2   66.8  3325   38.5  43.1  25.5   30    32
7 43    30    74   10.9   3.23  12.1   83.9  4679   3.5   49.2  11.3   21    32
8 45    30    73    9.3   3.29  10.6   86    2140   5.3   40.4  10.5    6     4
9 36    24    70    9     3.31  10.5   83.2  6582   8.1   42.5  12.6   18    12
10   36   27    72   9.5   3.36  10.7   79.3  4213   6.7   41    13.2   12     7
# ... with 50 more rows, and 3 more variables: x14 <dbl>, x15 <dbl>, y <dbl>
```

Here's a codebook:

Variable	Description
Variable	Description
y	Total Age Adjusted Mortality Rate
x1	Mean annual precipitation in inches
x2	Mean January temperature in degrees Fahrenheit
x3	Mean July temperature in degrees Fahrenheit
x4	Percent of 1960 SMSA population that is 65 years of age or over
x5	Population per household, 1960 SMSA
x6	Median school years completed for those over 25 in 1960 SMSA
x7	Percent of housing units that are found with facilities
x8	Population per square mile in urbanized area in 1960
x9	Percent of 1960 urbanized area population that is non-white
x10	Percent employment in white-collar occupations in 1960 urbanized area
x11	Percent of families with income under 3; 000 in 1960 urbanized area
x12	Relative population potential of hydrocarbons, HC
x13	Relative pollution potential of oxides of nitrogen, NOx
x14	Relative pollution potential of sulfur dioxide, SO2
x15	Percent relative humidity, annual average at 1 p.m.

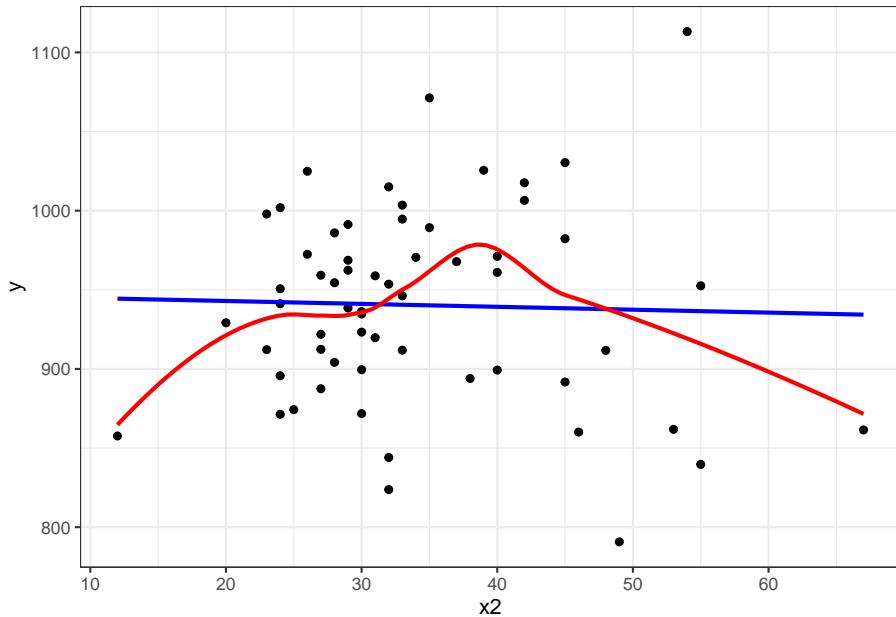
9.2 Fitting a straight line model to predict y from x2

Consider the relationship between y, the age-adjusted mortality rate, and x2, the mean January temperature, across these 60 areas. I'll include both a linear model (in blue) and a loess smooth (in red.) Does the relationship appear to be linear?

```
ggplot(pollution, aes(x = x2, y = y)) +
  geom_point() +
  geom_smooth(method = "lm", col = "blue", se = F) +
  geom_smooth(method = "loess", col = "red", se = F)

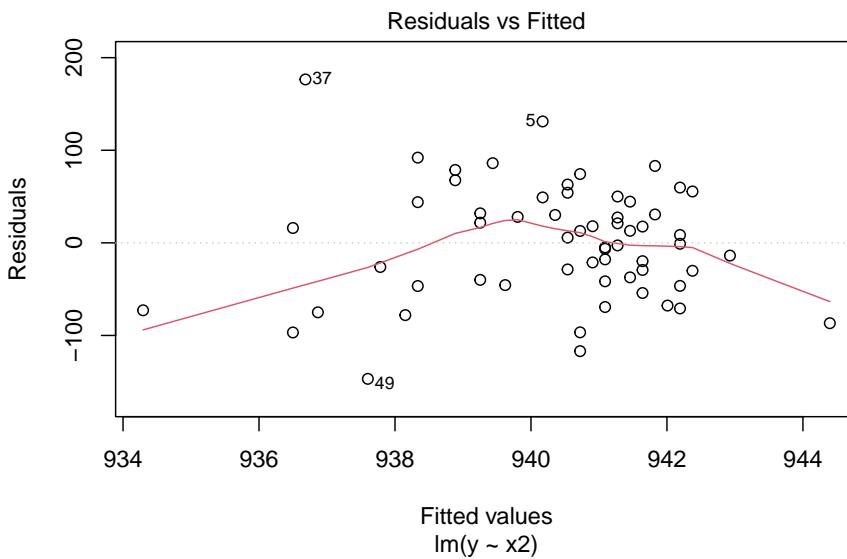
`geom_smooth()` using formula 'y ~ x'
`geom_smooth()` using formula 'y ~ x'
```

9.2. FITTING A STRAIGHT LINE MODEL TO PREDICT Y FROM X2 239



Suppose we plot the residuals that emerge from the linear model shown in blue, above. Do we see a curve in a plot of residuals against fitted values?

```
plot(lm(y ~ x2, data = pollution), which = 1)
```



9.3 Quadratic polynomial model to predict y using x2

A polynomial in the variable x of degree D is a linear combination of the powers of x up to D.

For example:

- Linear: $y = \beta_0 + \beta_1 x$
- Quadratic: $y = \beta_0 + \beta_1 x + \beta_2 x^2$
- Cubic: $y = \beta_0 + \beta_1 x + \beta_2 x^2 + \beta_3 x^3$
- Quartic: $y = \beta_0 + \beta_1 x + \beta_2 x^2 + \beta_3 x^3 + \beta_4 x^4$
- Quintic: $y = \beta_0 + \beta_1 x + \beta_2 x^2 + \beta_3 x^3 + \beta_4 x^4 + \beta_5 x^5$

Fitting such a model creates a **polynomial regression**.

9.3.1 The raw quadratic model

Let's look at a **quadratic model** which predicts y using x2 and the square of x2, so that our model is of the form:

$$y = \beta_0 + \beta_1 x_2 + \beta_2 x_2^2 + \text{error}$$

There are several ways to fit this exact model.

- One approach is to calculate the square of x2 within our pollution data set, and then feed both x2 and x2squared to lm.
- Another approach uses the I function within our lm to specify the use of both x2 and its square.
- Yet another approach uses the poly function within our lm, which can be used to specify raw models including x2 and x2squared.

```
pollution <- pollution %>%
  mutate(x2squared = x2^2)

mod2a <- lm(y ~ x2 + x2squared, data = pollution)
mod2b <- lm(y ~ x2 + I(x2^2), data = pollution)
mod2c <- lm(y ~ poly(x2, degree = 2, raw = TRUE), data = pollution)
```

Each of these approaches produces the same model, as they are just different ways of expressing the same idea.

```
summary(mod2a)
```

Call:

```
lm(formula = y ~ x2 + x2squared, data = pollution)
```

9.3. QUADRATIC POLYNOMIAL MODEL TO PREDICT Y USING X2 241

Residuals:

	Min	1Q	Median	3Q	Max
	-148.977	-38.651	6.889	35.312	189.346

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)							
(Intercept)	785.77449	79.54086	9.879	5.87e-14 ***							
x2	8.87640	4.27394	2.077	0.0423 *							
x2squared	-0.11704	0.05429	-2.156	0.0353 *							

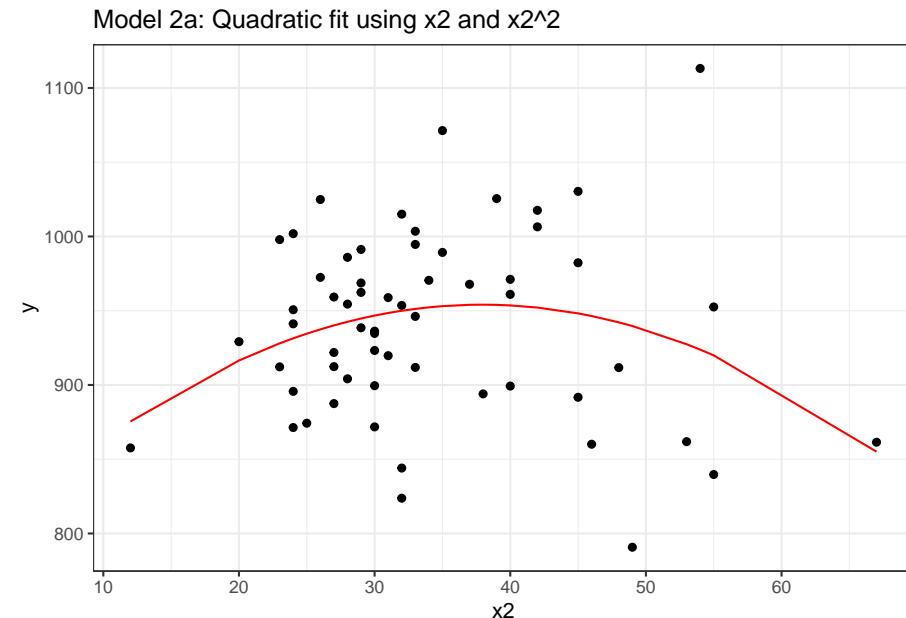
Signif. codes:	0	'***'	0.001	'**'	0.01	'*'	0.05	'. '	0.1	' '	1

Residual standard error: 60.83 on 57 degrees of freedom
 Multiple R-squared: 0.07623, Adjusted R-squared: 0.04382
 F-statistic: 2.352 on 2 and 57 DF, p-value: 0.1044

And if we plot the fitted values for this mod2 using whatever approach you like, we get exactly the same result.

```
mod2a.aug <- augment(mod2a, pollution)

ggplot(mod2a.aug, aes(x = x2, y = y)) +
  geom_point() +
  geom_line(aes(x = x2, y = .fitted), col = "red") +
  labs(title = "Model 2a: Quadratic fit using x2 and x2^2")
```



```

mod2b.aug <- augment(mod2b, pollution)

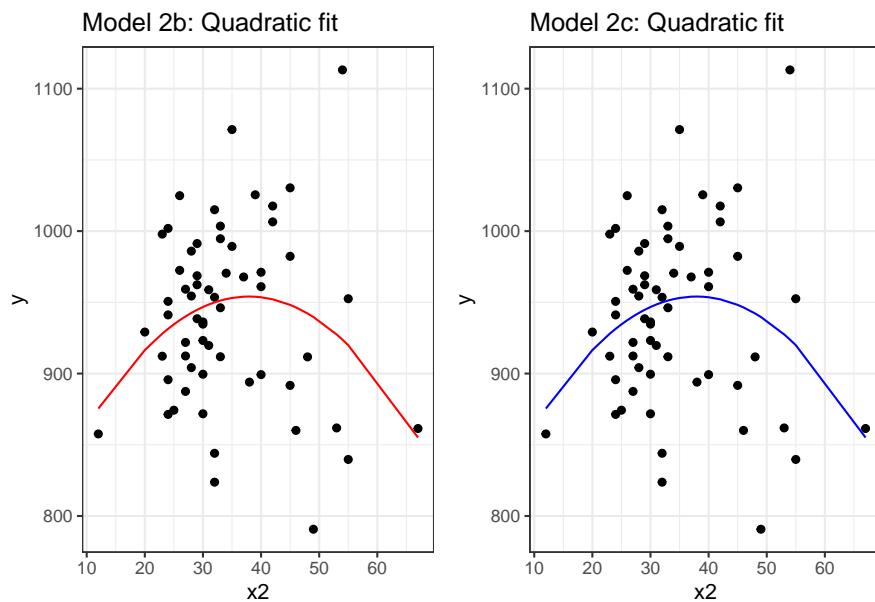
mod2c.aug <- augment(mod2c, pollution)

p1 <- ggplot(mod2b.aug, aes(x = x2, y = y)) +
  geom_point() +
  geom_line(aes(x = x2, y = .fitted), col = "red") +
  labs(title = "Model 2b: Quadratic fit")

p2 <- ggplot(mod2c.aug, aes(x = x2, y = y)) +
  geom_point() +
  geom_line(aes(x = x2, y = .fitted), col = "blue") +
  labs(title = "Model 2c: Quadratic fit")

p1 + p2

```



9.3.2 Raw quadratic fit after centering x2

Sometimes, we'll center (and perhaps rescale, too) the $x2$ variable before including it in a quadratic fit like this.

```

pollution <- pollution %>%
  mutate(x2_c = x2 - mean(x2))

```

9.3. QUADRATIC POLYNOMIAL MODEL TO PREDICT Y USING X2 243

```
mod2d <- lm(y ~ x2_c + I(x2_c^2), data = pollution)
summary(mod2d)
```

Call:
`lm(formula = y ~ x2_c + I(x2_c^2), data = pollution)`

Residuals:

Min	1Q	Median	3Q	Max
-148.977	-38.651	6.889	35.312	189.346

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	952.25941	9.59896	99.204	<2e-16 ***
x2_c	0.92163	0.93237	0.988	0.3271
I(x2_c^2)	-0.11704	0.05429	-2.156	0.0353 *

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

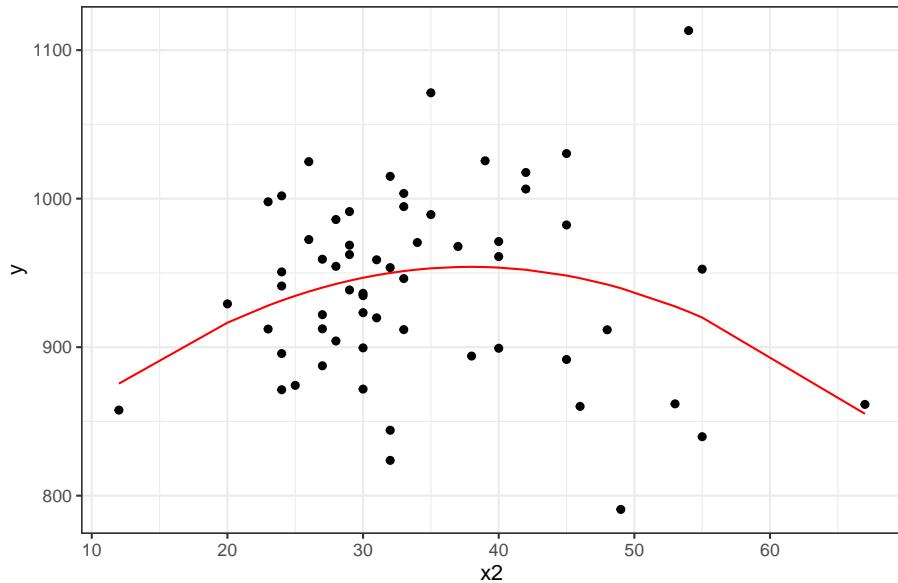
Residual standard error: 60.83 on 57 degrees of freedom
Multiple R-squared: 0.07623, Adjusted R-squared: 0.04382
F-statistic: 2.352 on 2 and 57 DF, p-value: 0.1044

Note that this model looks very different, with the exception of the second order quadratic term. But, it produces the same fitted values as the models we fit previously.

```
mod2d.aug <- augment(mod2d, pollution)

ggplot(mod2d.aug, aes(x = x2, y = y)) +
  geom_point() +
  geom_line(aes(x = x2, y = .fitted), col = "red") +
  labs(title = "Model 2d: Quadratic fit using centered x2 and x2^2")
```

Model 2d: Quadratic fit using centered x2 and x2^2



Or, if you don't believe me yet, look at the four sets of fitted values another way.

```
mosaic::favstats(~ .fitted, data = mod2a.aug)
```

	min	Q1	median	Q3	max	mean	sd	n	missing
	855.1041	936.7155	945.597	950.2883	954.073	940.3585	17.17507	60	0

```
mosaic::favstats(~ .fitted, data = mod2b.aug)
```

	min	Q1	median	Q3	max	mean	sd	n	missing
	855.1041	936.7155	945.597	950.2883	954.073	940.3585	17.17507	60	0

```
mosaic::favstats(~ .fitted, data = mod2c.aug)
```

	min	Q1	median	Q3	max	mean	sd	n	missing
	855.1041	936.7155	945.597	950.2883	954.073	940.3585	17.17507	60	0

```
mosaic::favstats(~ .fitted, data = mod2d.aug)
```

	min	Q1	median	Q3	max	mean	sd	n	missing
	855.1041	936.7155	945.597	950.2883	954.073	940.3585	17.17507	60	0

9.4 Orthogonal Polynomials

Now, let's fit an orthogonal polynomial of degree 2 to predict y using x2.

```

mod2_orth <- lm(y ~ poly(x2, 2), data = pollution)
summary(mod2_orth)

Call:
lm(formula = y ~ poly(x2, 2), data = pollution)

Residuals:
    Min      1Q  Median      3Q     Max 
-148.977 -38.651   6.889  35.312 189.346 

Coefficients:
            Estimate Std. Error t value Pr(>|t|)    
(Intercept)  940.358     7.853 119.746  <2e-16 ***
poly(x2, 2)1 -14.345     60.829  -0.236   0.8144  
poly(x2, 2)2 -131.142    60.829   -2.156   0.0353 *  
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 60.83 on 57 degrees of freedom
Multiple R-squared:  0.07623, Adjusted R-squared:  0.04382 
F-statistic: 2.352 on 2 and 57 DF,  p-value: 0.1044

```

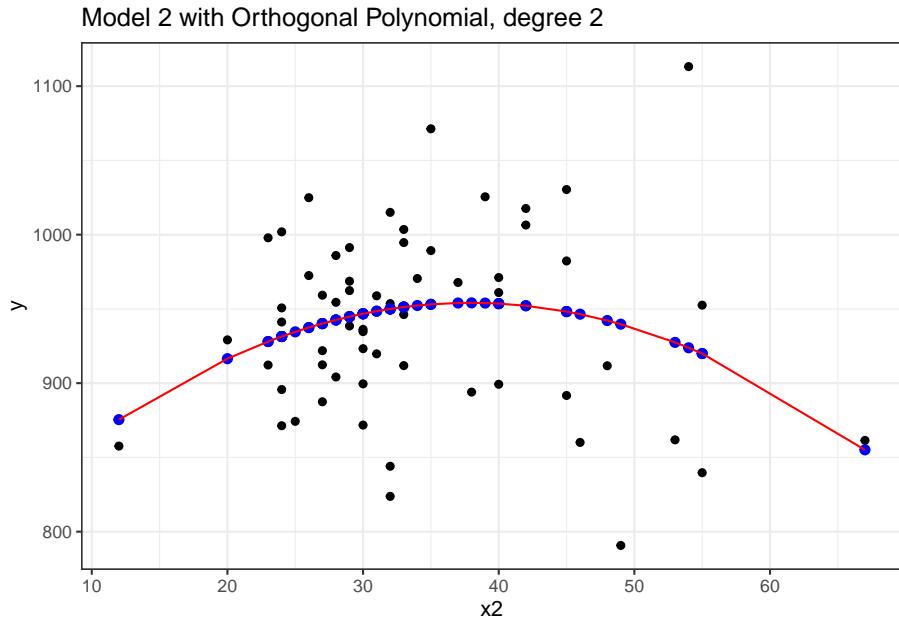
Now this looks very different in the equation, but, again, we can see that this produces exactly the same fitted values as our previous models, and the same model fit summaries. Is it, in fact, the same model? Here, we'll plot the fitted Model 2a in a red line, and this new Model 2 with Orthogonal Polynomials as blue points.

```

mod2orth.aug <- augment(mod2_orth, pollution)

ggplot(mod2orth.aug, aes(x = x2, y = y)) +
  geom_point() +
  geom_point(aes(x = x2, y = .fitted),
             col = "blue", size = 2) +
  geom_line(data = mod2a.aug, aes(x = x2, y = .fitted),
            col = "red") +
  labs(title = "Model 2 with Orthogonal Polynomial, degree 2")

```



Yes, it is again the same model in terms of the predictions it makes for y .

By default, with `raw = FALSE`, the `poly()` function within a linear model computes what is called an **orthogonal polynomial**. An orthogonal polynomial sets up a model design matrix using the coding we've seen previously: `x2` and `x2^2` in our case, and then scales those columns so that each column is **orthogonal** to the previous ones. This eliminates the collinearity (correlation between predictors) and lets our t tests tell us whether the addition of any particular polynomial term improves the fit of the model over the lower orders.

Would the addition of a cubic term help us much in predicting y from $x2$?

```
mod3 <- lm(y ~ poly(x2, 3), data = pollution)
summary(mod3)
```

Call:

```
lm(formula = y ~ poly(x2, 3), data = pollution)
```

Residuals:

Min	1Q	Median	3Q	Max
-146.262	-39.679	5.569	35.984	191.536

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	940.358	7.917	118.772	<2e-16 ***
poly(x2, 3)1	-14.345	61.328	-0.234	0.8159
poly(x2, 3)2	-131.142	61.328	-2.138	0.0369 *

```

poly(x2, 3)3   16.918      61.328    0.276    0.7837
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 61.33 on 56 degrees of freedom
Multiple R-squared:  0.07748, Adjusted R-squared:  0.02806
F-statistic: 1.568 on 3 and 56 DF,  p-value: 0.2073

```

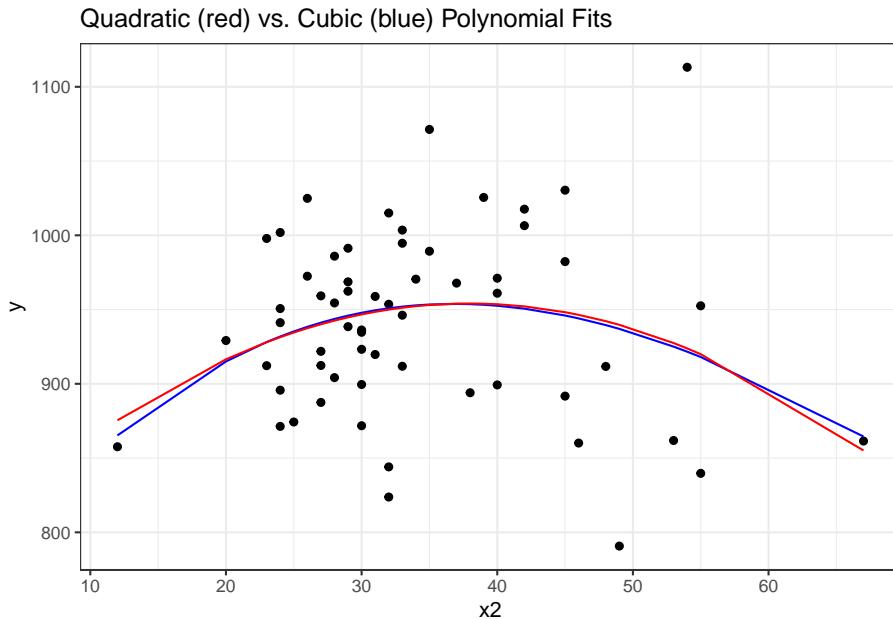
It doesn't appear that the cubic term adds much here, if anything. The p value is not significant for the third degree polynomial, the summaries of fit quality aren't much improved, and as we can see from the plot below, the predictions don't actually change all that much.

```

mod3.aug <- augment(mod3, pollution)

ggplot(mod3.aug, aes(x = x2, y = y)) +
  geom_point() +
  geom_line(aes(x = x2, y = .fitted),
            col = "blue") +
  geom_line(data = mod2orth.aug, aes(x = x2, y = .fitted),
            col = "red") +
  labs(title = "Quadratic (red) vs. Cubic (blue) Polynomial Fits")

```

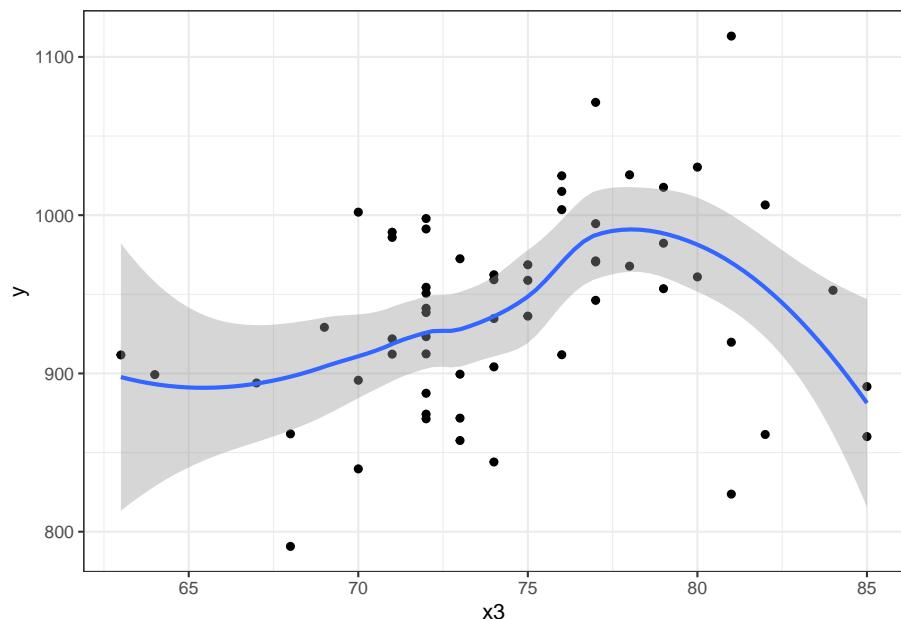


9.5 Fit a cubic polynomial to predict y from x3

What if we consider another predictor instead? Let's look at `x3`, the Mean July temperature in degrees Fahrenheit. Here is the `loess` smooth.

```
ggplot(pollution, aes(x = x3, y = y)) +
  geom_point() +
  geom_smooth(method = "loess")
```

```
`geom_smooth()` using formula 'y ~ x'
```



That looks pretty curvy - perhaps we need a more complex polynomial. We'll consider a linear model (`mod4_L`), a quadratic fit (`mod4_Q`) and a polynomial of degree 3: a **cubic** fit (`mod_4C`)

```
mod4_L <- lm(y ~ x3, data = pollution)
summary(mod4_L)
```

```
Call:
lm(formula = y ~ x3, data = pollution)
```

```
Residuals:
```

Min	1Q	Median	3Q	Max
-139.813	-34.341	4.271	38.197	149.587

```
Coefficients:
```

```

      Estimate Std. Error t value Pr(>|t|)
(Intercept) 670.529     123.140   5.445  1.1e-06 ***
x3          3.618      1.648    2.196   0.0321 *
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

```

Residual standard error: 60.29 on 58 degrees of freedom
 Multiple R-squared: 0.07674, Adjusted R-squared: 0.06082
 F-statistic: 4.821 on 1 and 58 DF, p-value: 0.03213

```

mod4_Q <- lm(y ~ poly(x3, 2), data = pollution)
summary(mod4_Q)

```

Call:
`lm(formula = y ~ poly(x3, 2), data = pollution)`

Residuals:

Min	1Q	Median	3Q	Max
-132.004	-42.184	4.069	47.126	157.396

Coefficients:

```

      Estimate Std. Error t value Pr(>|t|)
(Intercept) 940.358     7.553 124.503  <2e-16 ***
poly(x3, 2)1 132.364    58.504  2.262   0.0275 *
poly(x3, 2)2 -125.270   58.504 -2.141   0.0365 *
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

```

Residual standard error: 58.5 on 57 degrees of freedom
 Multiple R-squared: 0.1455, Adjusted R-squared: 0.1155
 F-statistic: 4.852 on 2 and 57 DF, p-value: 0.01133

```

mod4_C <- lm(y ~ poly(x3, 3), data = pollution)
summary(mod4_C)

```

Call:
`lm(formula = y ~ poly(x3, 3), data = pollution)`

Residuals:

Min	1Q	Median	3Q	Max
-148.004	-29.998	1.441	34.579	141.396

Coefficients:

```

      Estimate Std. Error t value Pr(>|t|)
(Intercept) 940.358     7.065 133.095  < 2e-16 ***
poly(x3, 3)1 132.364    54.728  2.419   0.01886 *
poly(x3, 3)2 -125.270   54.728 -2.289   0.02588 *

```

250 CHAPTER 9. ADDING NON-LINEAR TERMS TO A LINEAR REGRESSION MODEL

```

poly(x3, 3)3 -165.439      54.728   -3.023  0.00377 **
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 54.73 on 56 degrees of freedom
Multiple R-squared:  0.2654,    Adjusted R-squared:  0.226
F-statistic: 6.742 on 3 and 56 DF,  p-value: 0.0005799

```

It looks like the cubic polynomial term is of some real importance here. Do the linear, quadratic and cubic model fitted values look different?

```

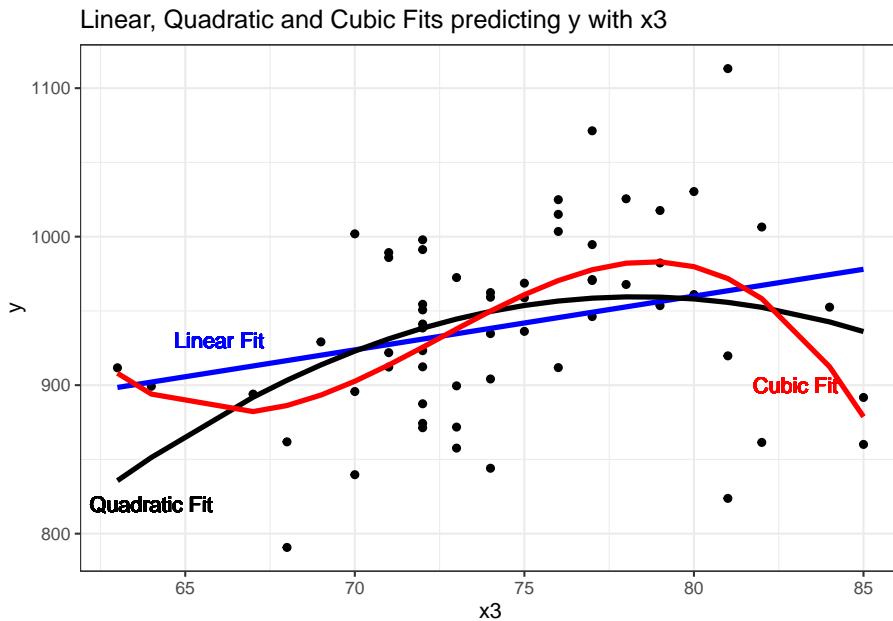
mod4_L.aug <- augment(mod4_L, pollution)

mod4_Q.aug <- augment(mod4_Q, pollution)

mod4_C.aug <- augment(mod4_C, pollution)

ggplot(pollution, aes(x = x3, y = y)) +
  geom_point() +
  geom_line(data = mod4_L.aug, aes(x = x3, y = .fitted),
            col = "blue", size = 1.25) +
  geom_line(data = mod4_Q.aug, aes(x = x3, y = .fitted),
            col = "black", size = 1.25) +
  geom_line(data = mod4_C.aug, aes(x = x3, y = .fitted),
            col = "red", size = 1.25) +
  geom_text(x = 66, y = 930, label = "Linear Fit", col = "blue") +
  geom_text(x = 64, y = 820, label = "Quadratic Fit", col = "black") +
  geom_text(x = 83, y = 900, label = "Cubic Fit", col = "red") +
  labs(title = "Linear, Quadratic and Cubic Fits predicting y with x3") +
  theme_bw()

```



9.6 Fitting a restricted cubic spline in a linear regression

- A **linear spline** is a continuous function formed by connecting points (called **knots** of the spline) by line segments.
- A **restricted cubic spline** is a way to build highly complicated curves into a regression equation in a fairly easily structured way.
- A restricted cubic spline is a series of polynomial functions joined together at the knots.
 - Such a spline gives us a way to flexibly account for non-linearity without over-fitting the model.
 - Restricted cubic splines can fit many different types of non-linearities.
 - Specifying the number of knots is all you need to do in R to get a reasonable result from a restricted cubic spline.

The most common choices are 3, 4, or 5 knots. Each additional knot adds to the non-linearity, and spends an additional degree of freedom:

- 3 Knots, 2 degrees of freedom, allows the curve to “bend” once.
- 4 Knots, 3 degrees of freedom, lets the curve “bend” twice.
- 5 Knots, 4 degrees of freedom, lets the curve “bend” three times.

For most applications, three to five knots strike a nice balance between complicating the model needlessly and fitting data pleasingly. Let’s consider a restricted

cubic spline model for our y based on $x3$ again, but now with:

- in `mod5a`, 3 knots,
- in `mod5b`, 4 knots, and
- in `mod5c`, 5 knots

```
mod5a_rcs <- lm(y ~ rcs(x3, 3), data = pollution)
mod5b_rcs <- lm(y ~ rcs(x3, 4), data = pollution)
mod5c_rcs <- lm(y ~ rcs(x3, 5), data = pollution)
```

Here, for instance, is the summary of the 5-knot model:

```
summary(mod5c_rcs)
```

Call:

```
lm(formula = y ~ rcs(x3, 5), data = pollution)
```

Residuals:

Min	1Q	Median	3Q	Max
-141.522	-32.009	1.674	31.971	147.878

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	468.113	396.319	1.181	0.243
$\text{rcs}(x3, 5)x3$	6.447	5.749	1.121	0.267
$\text{rcs}(x3, 5)x3'$	-25.633	46.810	-0.548	0.586
$\text{rcs}(x3, 5)x3''$	323.137	293.065	1.103	0.275
$\text{rcs}(x3, 5)x3'''$	-612.578	396.270	-1.546	0.128

Residual standard error: 54.35 on 55 degrees of freedom

Multiple R-squared: 0.2883, Adjusted R-squared: 0.2366

F-statistic: 5.571 on 4 and 55 DF, p-value: 0.0007734

We'll begin by storing the fitted values from these three models and other summaries, for plotting.

```
mod5a.aug <- augment(mod5a_rcs, pollution)

mod5b.aug <- augment(mod5b_rcs, pollution)

mod5c.aug <- augment(mod5c_rcs, pollution)

p2 <- ggplot(pollution, aes(x = x3, y = y)) +
  geom_point() +
  geom_smooth(method = "loess", col = "purple", se = F) +
  labs(title = "Loess Smooth") +
  theme_bw()

p3 <- ggplot(mod5a.aug, aes(x = x3, y = y)) +
```

```

geom_point() +
  geom_line(aes(x = x3, y = .fitted),
            col = "blue", size = 1.25) +
  labs(title = "RCS, 3 knots") +
  theme_bw()

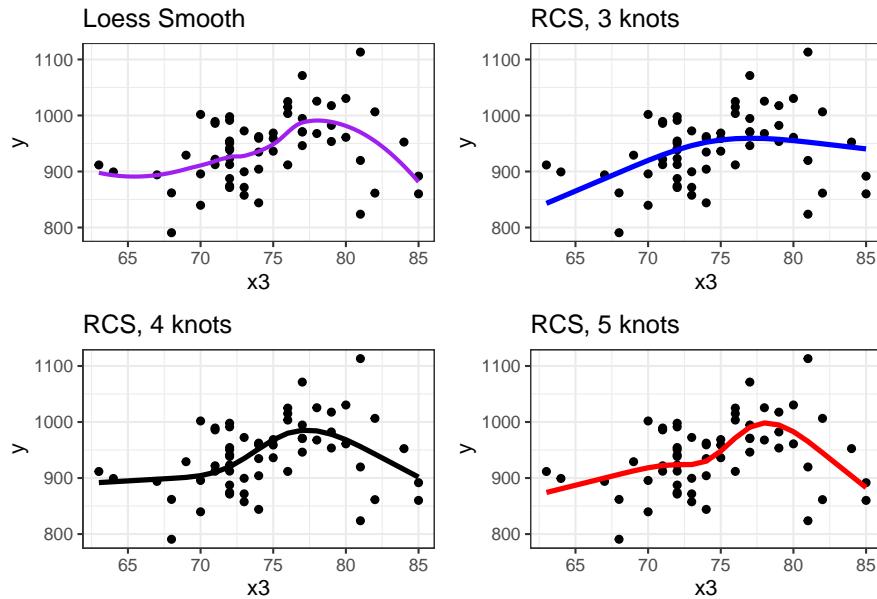
p4 <- ggplot(mod5b.aug, aes(x = x3, y = y)) +
  geom_point() +
  geom_line(aes(x = x3, y = .fitted),
            col = "black", size = 1.25) +
  labs(title = "RCS, 4 knots") +
  theme_bw()

p5 <- ggplot(mod5c.aug, aes(x = x3, y = y)) +
  geom_point() +
  geom_line(aes(x = x3, y = .fitted),
            col = "red", size = 1.25) +
  labs(title = "RCS, 5 knots") +
  theme_bw()

(p2 + p3) / (p4 + p5)

```

`geom_smooth()` using formula 'y ~ x'



Does it seem like the fit improves markedly (perhaps approaching the loess smooth result) as we increase the number of knots?

```
anova(mod5a_rcs, mod5b_rcs, mod5c_rcs)
```

Analysis of Variance Table

```
Model 1: y ~ rcs(x3, 3)
Model 2: y ~ rcs(x3, 4)
Model 3: y ~ rcs(x3, 5)

  Res.Df   RSS Df Sum of Sq    F    Pr(>F)
1     57 194935
2     56 171448  1   23486.9 7.9503 0.006672 **
3     55 162481  1     8967.2 3.0354 0.087057 .
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

Based on an ANOVA comparison, the fourth knot adds significant predictive value ($p = 0.0067$), but the fifth knot is borderline ($p = 0.0871$). From the `glance` function in the `broom` package, we can also look at some key summaries.

```
glance(mod5a_rcs)
```

```
# A tibble: 1 x 12
  r.squared adj.r.squared sigma statistic p.value    df logLik    AIC    BIC
      <dbl>         <dbl> <dbl>     <dbl> <dbl> <dbl> <dbl> <dbl>
1     0.146        0.116  58.5     4.88  0.0111     2 -328.  663.  672.
# ... with 3 more variables: deviance <dbl>, df.residual <int>, nobs <int>
```

```
glance(mod5b_rcs)
```

```
# A tibble: 1 x 12
  r.squared adj.r.squared sigma statistic p.value    df logLik    AIC    BIC
      <dbl>         <dbl> <dbl>     <dbl> <dbl> <dbl> <dbl> <dbl>
1     0.249        0.209  55.3     6.19  0.00104    3 -324.  658.  668.
# ... with 3 more variables: deviance <dbl>, df.residual <int>, nobs <int>
```

```
glance(mod5c_rcs)
```

```
# A tibble: 1 x 12
  r.squared adj.r.squared sigma statistic p.value    df logLik    AIC    BIC
      <dbl>         <dbl> <dbl>     <dbl> <dbl> <dbl> <dbl> <dbl>
1     0.288        0.237  54.4     5.57  0.000773    4 -322.  657.  669.
# ... with 3 more variables: deviance <dbl>, df.residual <int>, nobs <int>
```

Model	Knots	R^2	Adj. R^2	AIC	BIC
5a	3	0.146	0.116	663.4	671.8
5b	4	0.249	0.209	657.7	668.2
5c	5	0.288	0.237	656.5	669.1

Model	Knots	R^2	Adj. R^2	AIC	BIC
-------	-------	-------	------------	-----	-----

Within our sample, the five-knot RCS outperforms the 3- and 4-knot versions on adjusted R^2 and AIC (barely) and does a little worse than the 4-knot RCS on BIC.

Of course, we could also use the cross-validation methods we've developed for other linear regressions to assess predictive capacity of these models. I'll skip that for now.

To see the values of `x3` where the splines place their knots, we can use the `attributes` function.

```
attributes(rcs(pollution$x3, 5))

$dim
[1] 60  4

$dimnames
$dimnames[[1]]
NULL

$dimnames[[2]]
[1] "pollution"     "pollution''"    "pollution'''"   "pollution''''"

$class
[1] "rms"

$name
[1] "pollution"

$label
[1] "pollution"

$assume
[1] "rcspline"

$assume.code
[1] 4

$parms
[1] 68 72 74 77 82

$nonlinear
[1] FALSE  TRUE  TRUE  TRUE
```

```
$colnames
[1] "pollution"      "pollution''"    "pollution'''"   "pollution'''''
```

The knots in this particular 5-knot spline are placed by the computer at 68, 72, 74, 77 and 82, it seems.

There are two kinds of Multivariate Regression Models

1. [Prediction] Those that are built so that we can make accurate predictions.
2. [Explanatory] Those that are built to help understand underlying phenomena.

While those two notions overlap considerably, they do imply different things about how we strategize about model-building and model assessment. Harrell's primary concern is effective use of the available data for **prediction** - this implies some things that will be different from what we've seen in the past.

Harrell refers to multivariable regression modeling strategy as the process of **spending degrees of freedom**. The main job in strategizing about multivariate modeling is to

1. Decide the number of degrees of freedom that can be spent
2. Decide where to spend them
3. Spend them, wisely.

What this means is essentially linked to making decisions about predictor complexity, both in terms of how many predictors will be included in the regression model, and about how we'll include those predictors.

9.7 “Spending” Degrees of Freedom

- “Spending” df includes
 - fitting parameter estimates in models, or
 - examining figures built using the outcome variable Y that tell you how to model the predictors.

If you use a scatterplot of Y vs. X or the residuals of the Y-X regression model vs. X to decide whether a linear model is appropriate, then how many degrees of freedom have you actually spent?

Grambsch and O'Brien conclude that if you wish to preserve the key statistical properties of the various estimation and fitting procedures used in building a model, you can't retrieve these degrees of freedom once they have been spent.

9.7.1 Overfitting and Limits on the # of Predictors

Suppose you have a total sample size of n observations, then you really shouldn't be thinking about estimating more than $n/15$ regression coefficients, at the most.

- If k is the number of parameters in a full model containing all candidate predictors for a stepwise analysis, then k should be no greater than $n/15$.
- k should include all variables screened for association with the response, including interaction terms.
- Sometimes I hold myself to a tougher standard, or $n/50$ predictors, at maximum.

So if you have 97 observations in your data, then you can probably just barely justify the use of a stepwise analysis using the main effects alone of 5 candidate variables (with one additional DF for the intercept term) using the $n/15$ limit.

Harrell (2001) also mentions that if you have a **narrowly distributed** predictor, without a lot of variation to work with, then an even larger sample size n should be required. See Vittinghoff et al. (2012), Section 10.3 for more details.

9.7.2 The Importance of Collinearity

Collinearity denotes correlation between predictors high enough to degrade the precision of the regression coefficient estimates substantially for some or all of the correlated predictors

- Vittinghoff et al. (2012), section 10.4.1
- Can one predictor in a model be predicted well using the other predictors in the model?
 - Strong correlations (for instance, $r \geq 0.8$) are especially troublesome.
- Effects of collinearity
 - decreases precision, in the sense of increasing the standard errors of the parameter estimates
 - decreases power
 - increases the difficulty of interpreting individual predictor effects
 - overall F test is significant, but individual t tests may not be

Suppose we want to assess whether variable X_j is collinear with the other predictors in a model. We run a regression predicting X_j using the other predictors, and obtain the R^2 . The VIF is defined as $1 / (1 - \text{this } R^2)$, and we usually interpret VIFs above 5 as indicating a serious multicollinearity problem (i.e. R^2 values for this predictor of 0.8 and above would thus concern us.)

```
car::vif(lm(y ~ x1 + x2 + x3 + x4 + x5 + x6, data = pollution))
```

x1	x2	x3	x4	x5	x6
2.238862	2.058731	2.153044	4.174448	3.447399	1.792996

Occasionally, you'll see the inverse of VIF reported, and this is called *tolerance*.

- $\text{tolerance} = 1 / \text{VIF}$

9.7.3 Collinearity in an Explanatory Model

- When we are attempting to **identify multiple independent predictors** (the explanatory model approach), then we will need to choose between collinear variables
 - options suggested by Vittinghoff et al. (2012), p. 422, include choosing on the basis of plausibility as a causal factor,
 - choosing the variable that has higher data quality (is measured more accurately or has fewer missing values.)
 - Often, we choose to include a variable that is statistically significant as a predictor, and drop others, should we be so lucky.
- Larger effects, especially if they are associated with predictors that have minimal correlation with the other predictors under study, cause less trouble in terms of potential violation of the $n/15$ rule for what constitutes a reasonable number of predictors.

9.7.4 Collinearity in a Prediction Model

- If we are primarily building a **prediction model** for which inference on the individual predictors is not of interest, then it is totally reasonable to use both predictors in the model, if doing so reduces prediction error.
 - Collinearity doesn't affect predictions in our model development sample.
 - Collinearity doesn't affect predictions on new data so long as the new data have similar relationships between predictors.
 - If our key predictor is correlated strongly with a confounder, then if the predictor remains significant after adjustment for the confounder, then this suggests a meaningful independent effect.
 - * If the effects of the predictor are clearly confounded by the adjustment variable, we again have a clear result.
 - * If neither is statistically significant after adjustment, the data may be inadequate.
 - If the collinearity is between adjustment variables, but doesn't involve the key predictor, then inclusion of the collinear variables is unlikely to cause substantial problems.

9.8 Spending DF on Non-Linearity: The Spearman Plot

We need a flexible approach to assessing non-linearity and fitting models with non-linear predictors. This will lead us to a measure of what Harrell (2001) calls **potential predictive punch** which hides the true form of the regression from the analyst so as to preserve statistical properties, but that lets us make sensible decisions about whether a predictor should be included in a model, and the number of parameters (degrees of freedom, essentially) we are willing to devote to it.

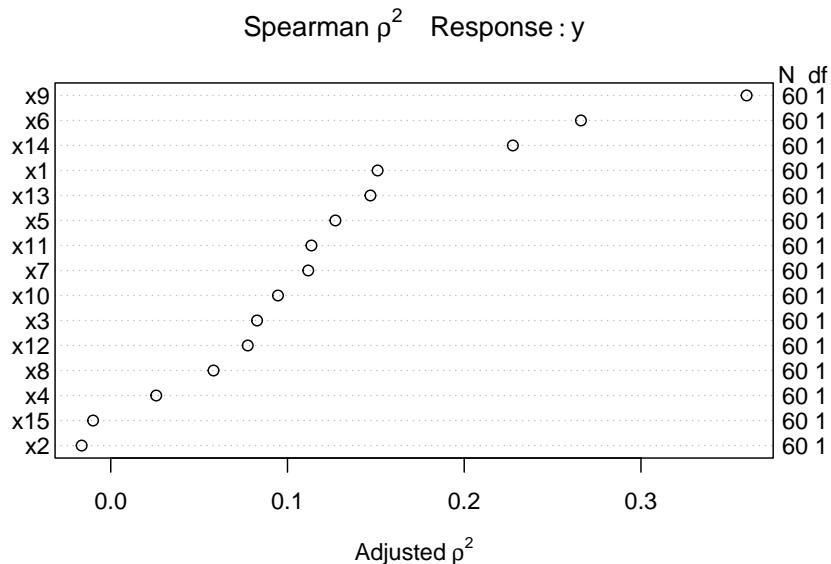
What if we want to consider where best to spend our degrees of freedom on non-linear predictor terms, like interactions, polynomial functions or curved splines to represent our input data? The approach we'll find useful in the largest variety of settings is a combination of

1. a rank correlation assessment of potential predictive punch (using a Spearman ρ^2 plot, available in the `Hmisc` package), followed by
2. the application of restricted cubic splines to fit and assess models.

Let's try such a plot for our fifteen predictors:

```
sp2 <- Hmisc:::spearman2(y ~ x1 + x2 + x3 + x4 + x5 + x6 + x7 +
                           x8 + x9 + x10 + x11 + x12 + x13 +
                           x14 + x15, data = pollution)

plot(sp2)
```



The variable with the largest adjusted squared Spearman ρ statistic in this setting is $x9$, followed by $x6$ and $x14$. With only 60 observations, we might well want to restrict ourselves to a very small model. What the Spearman plot suggests is that we focus any non-linear terms on $x9$ first, and then perhaps $x6$ and $x14$ as they have some potential predictive power. It may or may not work out that the non-linear terms are productive.

9.8.1 Fitting a Big Model to the pollution data

So, one possible model built in reaction this plot might be to fit:

- a restricted cubic spline with 5 knots on $x9$,
- a restricted cubic spline with 3 knots on $x6$,
- a quadratic polynomial on $x14$, and
- a linear fit to $x1$ and $x13$

That's way more degrees of freedom (4 for $x9$, 2 for $x6$, 2 for $x14$ and 1 each for $x1$ and $x13$ makes a total of 10 without the intercept term) than we can really justify with a sample of 60 observations. But let's see what happens.

```
mod_big <- lm(y ~ rcs(x9, 5) + rcs(x6, 3) + poly(x14, 2) +
                 x1 + x13, data = pollution)

anova(mod_big)
```

Analysis of Variance Table

```

Response: y
          Df Sum Sq Mean Sq F value    Pr(>F)
rcs(x9, 5)   4 100164 25040.9 17.8482 4.229e-09 ***
rcs(x6, 3)   2  38306 19152.8 13.6513 1.939e-05 ***
poly(x14, 2) 2  15595  7797.7  5.5579  0.006677 **
x1           1   4787  4787.3  3.4122  0.070759 .
x13          1    712   711.9  0.5074  0.479635
Residuals    49  68747  1403.0
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

```

This `anova` suggests that we have at least some predictive value in each spline (`x9` and `x6`) and some additional value in `x14`, although it's not as clear that the linear terms (`x1` and `x13`) did much good.

9.8.2 Limitations of `lm` for fitting complex linear regression models

We can certainly assess this big, complex model using `lm` in comparison to other models:

- with in-sample summary statistics like adjusted R^2 , AIC and BIC,
- we can assess its assumptions with residual plots, and
- we can also compare out-of-sample predictive quality through cross-validation,

But to really delve into the details of how well this complex model works, and to help plot what is actually being fit, we'll probably want to fit the model using an alternative method for fitting linear models, called `ols`, from the `rms` package developed by Frank Harrell and colleagues. That will be the focus of our next chapter.

Chapter 10

Using `ols` from the `rms` package to fit linear models

At the end of the previous chapter, we had fit a model to the `pollution` data that predicted our outcome $y = \text{Age-Adjusted Mortality Rate}$, using:

- a restricted cubic spline with 5 knots on `x9`
- a restricted cubic spline with 3 knots on `x6`
- a polynomial in 2 degrees on `x14`
- linear terms for `x1` and `x13`

but this model was hard to evaluate in some ways. Now, instead of using `lm` to fit this model, we'll use a new function called `ols` from the `rms` package developed by Frank Harrell and colleagues, in part to support ideas developed in Harrell (2001) for clinical prediction models.

10.1 Fitting a model with `ols`

We will use the `datadist` approach when fitting a linear model with `ols` from the `rms` package, so as to store additional important elements of the model fit.

```
library(rms)

d <- datadist(pollution)
options(datadist = "d")
```

Next, we'll fit the model using `ols` and place its results in `newmod`.

```
newmod <- ols(y ~ rcs(x9, 5) + rcs(x6, 3) + pol(x14, 2) +
  x1 + x13,
```

```

        data = pollution, x = TRUE, y = TRUE)
newmod

Linear Regression Model

ols(formula = y ~ rcs(x9, 5) + rcs(x6, 3) + pol(x14, 2) + x1 +
    x13, data = pollution, x = TRUE, y = TRUE)

      Model Likelihood   Discrimination
      Ratio Test           Indexes
Obs       60     LR chi2     72.02     R2       0.699
sigma37.4566 d.f.          10     R2 adj   0.637
d.f.       49     Pr(> chi2) 0.0000     g       58.961

Residuals

      Min      1Q Median      3Q      Max
-86.189 -18.554 -1.799  18.645 104.307

      Coef      S.E.      t      Pr(>|t|)
Intercept 796.2658 162.3269  4.91 <0.0001
x9         -2.6328  6.3504 -0.41  0.6803
x9'        121.4651 124.4827  0.98 0.3340
x9''       -219.8025 227.6775 -0.97 0.3391
x9'''      151.5700 171.3867  0.88 0.3808
x6         7.6817  15.5230  0.49 0.6229
x6'        -29.4388 18.0531 -1.63 0.1094
x14        0.5652  0.2547  2.22 0.0311
x14^2     -0.0010  0.0010 -0.96 0.3407
x1         1.0717  0.7317  1.46 0.1494
x13        -0.1028  0.1443 -0.71 0.4796

```

Some of the advantages and disadvantages of fitting linear regression models with `ols` or `lm` will reveal themselves over time. For now, one advantage for `ols` is that the entire variance-covariance matrix is saved. Most of the time, there will be some value to considering both `ols` and `lm` approaches.

Most of this output should be familiar, but a few pieces are different.

10.1.1 The Model Likelihood Ratio Test

The **Model Likelihood Ratio Test** compares `newmod` to the null model with only an intercept term. It is a goodness-of-fit test that we'll use in several types of model settings this semester.

- In many settings, the logarithm of the likelihood ratio, multiplied by -2, yields a value which can be compared to a χ^2 distribution. So here, the value 72.02 is -2(log likelihood), and is compared to a χ^2 distribution with 10 degrees of freedom. We reject the null hypothesis that `newmod` is no better than the null model, and conclude instead that at least one of these predictors adds statistically significant value.
 - For `ols`, interpret the model likelihood ratio test like the global (ANOVA) F test in `lm`.
 - The likelihood function is the probability of observing our data under the specified model.
 - We can compare two nested models by evaluating the difference in their likelihood ratios and degrees of freedom, then comparing the result to a χ^2 distribution.

10.1.2 The g statistic

The **g statistic** is new and is referred to as the g-index. it's based on Gini's mean difference and is purported to be a robust and highly efficient measure of variation.

- Here, $g = 58.9$, which implies that if you randomly select two of the 60 areas included in the model, the average difference in predicted `y` (Age-Adjusted Mortality Rate) using this model will be 58.9.
 - Technically, g is Gini's mean difference of the predicted values.

10.2 ANOVA for an ols model

One advantage of the `ols` approach is that when you apply an `anova` to it, it separates out the linear and non-linear components of restricted cubic splines and polynomial terms (as well as product terms, if your model includes them.)

```
anova(newmod)
```

Factor	d.f.	Analysis of Variance		Response: y	
		Partial SS	MS	F	P
x9	4	35219.7647	8804.9412	6.28	0.0004
Nonlinear	3	1339.3081	446.4360	0.32	0.8121
x6	2	9367.6008	4683.8004	3.34	0.0437
Nonlinear	1	3730.7388	3730.7388	2.66	0.1094
x14	2	18679.6957	9339.8478	6.66	0.0028
Nonlinear	1	1298.7625	1298.7625	0.93	0.3407
x1	1	3009.1829	3009.1829	2.14	0.1494
x13	1	711.9108	711.9108	0.51	0.4796

```
TOTAL NONLINEAR  5      6656.1824  1331.2365  0.95 0.4582
REGRESSION       10    159563.8285 15956.3829 11.37 <.0001
ERROR           49    68746.8004  1402.9959
```

Unlike the `anova` approach in `lm`, in `ols` ANOVA, *partial* F tests are presented - each predictor is assessed as “last predictor in” much like the usual *t* tests in `lm`. In essence, the partial sums of squares and F tests here describe the marginal impact of removing each covariate from `newmod`.

We conclude that the non-linear parts of `x9` and `x6` and `x14` combined don’t seem to add much value, but that overall, `x9`, `x6` and `x14` seem to be valuable. So it must be the linear parts of those variables within our model that are doing the lion’s share of the work.

10.3 Effect Estimates

A particularly useful thing to get out of the `ols` approach that is not as easily available in `lm` (without recoding or standardizing our predictors) is a summary of the effects of each predictor in an interesting scale.

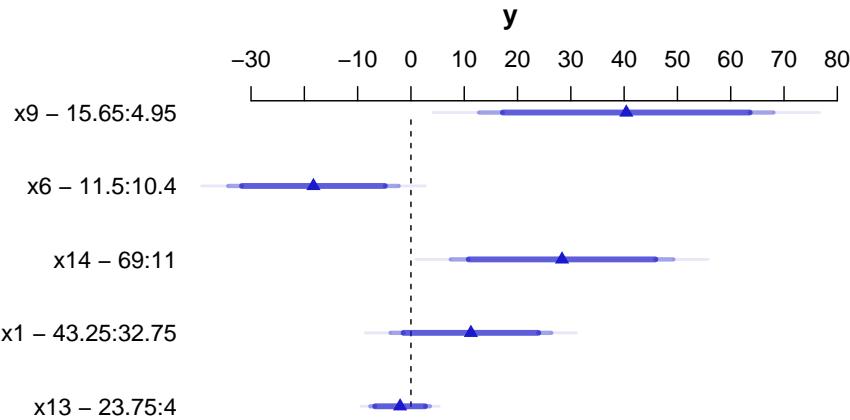
```
summary(newmod)
```

	Effects			Response : y			
Factor	Low	High	Diff. Effect	S.E.	Lower	0.95	Upper
<code>x9</code>	4.95	15.65	10.70	40.4060	14.0790	12.1120	68.6990
<code>x6</code>	10.40	11.50	1.10	-18.2930	8.1499	-34.6710	-1.9153
<code>x14</code>	11.00	69.00	58.00	28.3480	10.6480	6.9503	49.7460
<code>x1</code>	32.75	43.25	10.50	11.2520	7.6833	-4.1878	26.6930
<code>x13</code>	4.00	23.75	19.75	-2.0303	2.8502	-7.7579	3.6973

This “effects summary” shows the effect on `y` of moving from the 25th to the 75th percentile of each variable (along with a standard error and 95% confidence interval) while holding the other variable at the level specified at the bottom of the output.

The most useful way to look at this sort of analysis is often a plot.

```
plot(summary(newmod))
```



For $x9$ note from the `summary` above that the 25th percentile is 4.95 and the 75th is 15.65. Our conclusion is that the estimated effect of moving $x9$ from 4.95 to 15.65 is an increase of 40.4 on y , with a 95% CI of (12.1, 68.7).

For a categorical variable, the low level is shown first and then the high level.

The plot shows the point estimate (arrow head) and then the 90% (narrowest bar), 95% (middle bar) and 99% (widest bar in lightest color) confidence intervals for each predictor's effect.

- It's easier to distinguish this in the $x9$ plot than the one for $x13$.
- Remember that what is being compared is the first value to the second value's impact on the outcome, with other predictors held constant.

10.3.1 Simultaneous Confidence Intervals

These confidence intervals make no effort to deal with the multiple comparisons problem, but just fit individual 95% (or whatever level you choose) confidence intervals for each predictor. The natural alternative is to make an adjustment for multiple comparisons in fitting the confidence intervals, so that the set of (in this case, five - one for each predictor) confidence intervals for effect sizes has a family-wise 95% confidence level. You'll note that the effect estimates and standard errors are unchanged from those shown above, but the confidence limits are a bit wider.

```
summary(newmod, conf.type=c('simultaneous'))
```

	Effects				Response : y			
Factor	Low	High	Diff.	Effect	S.E.	Lower	0.95	Upper
x9	4.95	15.65	10.70	40.4060	14.0790	3.09830	77.7130	
x6	10.40	11.50	1.10	-18.2930	8.1499	-39.88900	3.3024	
x14	11.00	69.00	58.00	28.3480	10.6480	0.13337	56.5630	
x1	32.75	43.25	10.50	11.2520	7.6833	-9.10670	31.6120	
x13	4.00	23.75	19.75	-2.0303	2.8502	-9.58260	5.5221	

Remember that if you're looking for the usual `lm` summary for an `ols` object, use `summary.lm`, and that the `display` function from `arm` does not recognize `ols` objects.

10.4 The Predict function for an ols model

The `Predict` function is very flexible, and can be used to produce individual or simultaneous confidence limits.

```
Predict(newmod, x9 = 12, x6 = 12, x14 = 40, x1 = 40, x13 = 20) # individual limits
```

x9	x6	x14	x1	x13	yhat	lower	upper	
1	12	12	40	40	20	923.0982	893.0984	953.098

Response variable (y): y

Limits are 0.95 confidence limits

```
Predict(newmod, x9 = 5:15) # individual limits
```

x9	x6	x14	x1	x13	yhat	lower	upper	
1	5	11.05	30	38	9	913.7392	889.4802	937.9983
2	6	11.05	30	38	9	916.3490	892.0082	940.6897
3	7	11.05	30	38	9	921.3093	898.9657	943.6529
4	8	11.05	30	38	9	927.6464	907.0355	948.2574
5	9	11.05	30	38	9	934.3853	913.3761	955.3946
6	10	11.05	30	38	9	940.5510	917.8371	963.2648
7	11	11.05	30	38	9	945.2225	921.9971	968.4479
8	12	11.05	30	38	9	948.2885	926.4576	970.1194
9	13	11.05	30	38	9	950.2608	930.3003	970.2213
10	14	11.05	30	38	9	951.6671	932.2370	971.0971
11	15	11.05	30	38	9	953.0342	932.1662	973.9021

Response variable (y): y

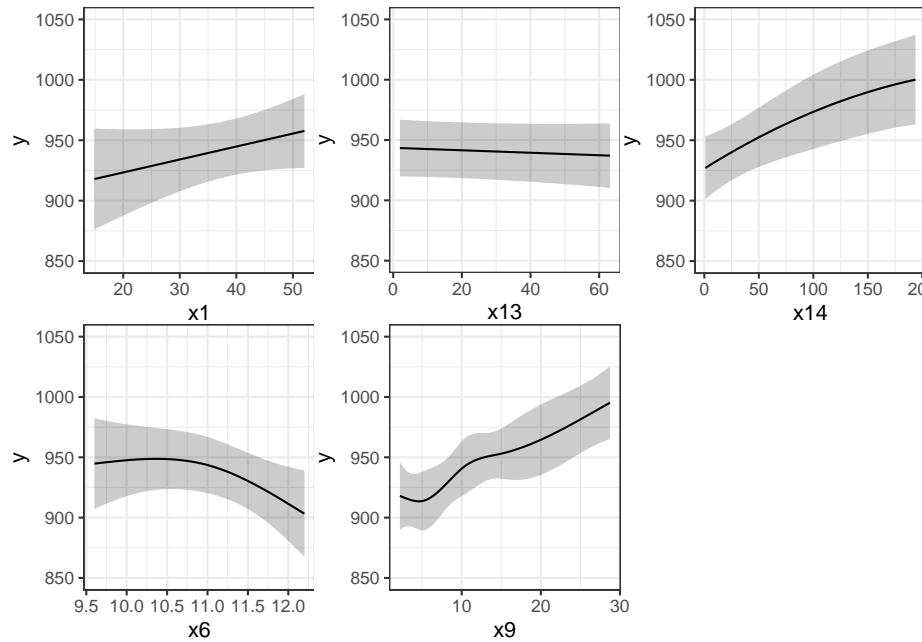
```
Adjust to: x6=11.05 x14=30 x1=38 x13=9
Limits are 0.95 confidence limits
Predict(newmod, x9 = 5:15, conf.type = 'simult')
```

	x9	x6	x14	x1	x13	yhat	lower	upper
1	5	11.05	30	38	9	913.7392	882.4115	945.0669
2	6	11.05	30	38	9	916.3490	884.9158	947.7822
3	7	11.05	30	38	9	921.3093	892.4552	950.1635
4	8	11.05	30	38	9	927.6464	901.0299	954.2630
5	9	11.05	30	38	9	934.3853	907.2544	961.5163
6	10	11.05	30	38	9	940.5510	911.2187	969.8832
7	11	11.05	30	38	9	945.2225	915.2296	975.2154
8	12	11.05	30	38	9	948.2885	920.0965	976.4805
9	13	11.05	30	38	9	950.2608	924.4842	976.0374
10	14	11.05	30	38	9	951.6671	926.5755	976.7587
11	15	11.05	30	38	9	953.0342	926.0856	979.9827

```
Response variable (y): y
Adjust to: x6=11.05 x14=30 x1=38 x13=9
Limits are 0.95 confidence limits
```

The plot below shows the individual effects in `newmod` in five subpanels, using the default approach of displaying the same range of values as are seen in the data. Note that each panel shows point and interval estimates of the effects, and spot the straight lines in `x1` and `x13`, the single bends in `x14` and `x6` and the wiggles in `x9`, corresponding to the amount of non-linearity specified in the model.

```
ggplot(Predict(newmod))
```



10.5 Checking Influence via `dfbeta`

For an `ols` object, we have several tools for looking at residuals. The most interesting to me is `which.influence` which is reliant on the notion of `dfbeta`.

- DFBETA is estimated for each observation in the data, and each coefficient in the model.
- The DFBETA is the difference in the estimated coefficient caused by deleting the observation, scaled by the coefficient's standard error estimated with the observation deleted.
- The `which.influence` command applied to an `ols` model produces a list of all of the predictors estimated by the model, including the intercept.
 - For each predictor, the command lists all observations (by row number) that, if removed from the model, would cause the estimated coefficient (the “beta”) for that predictor to change by at least some particular cutoff.
 - The default is that the DFBETA for that predictor is 0.2 or more.

```
which.influence(newmod)
```

```
$Intercept
[1] 2 11 28 32 37 49 59
```

```
$x9
```

```
[1] 2 3 6 9 31 35 49 57 58

$x6
[1] 2 11 15 28 32 37 50 56 59

$x14
[1] 2 6 7 12 13 16 32 37

$x1
[1] 7 18 32 37 49 57

$x13
[1] 29 32 37
```

The implication here, for instance, is that if we drop row 3 from our data frame, and refit the model, this will have a meaningful impact on the estimate of `x9` but not on the other coefficients. But if we drop, say, row 37, we will affect the estimates of the intercept, `x6`, `x14`, `x1`, and `x13`.

10.5.1 Using the `residuals` command for `dfbetas`

To see the `dfbeta` values, standardized according to the approach I used above, you can use the following code (I'll use `head` to just show the first few rows of results) to get a matrix of the results.

```
head(residuals(newmod, type = "dfbetas"))

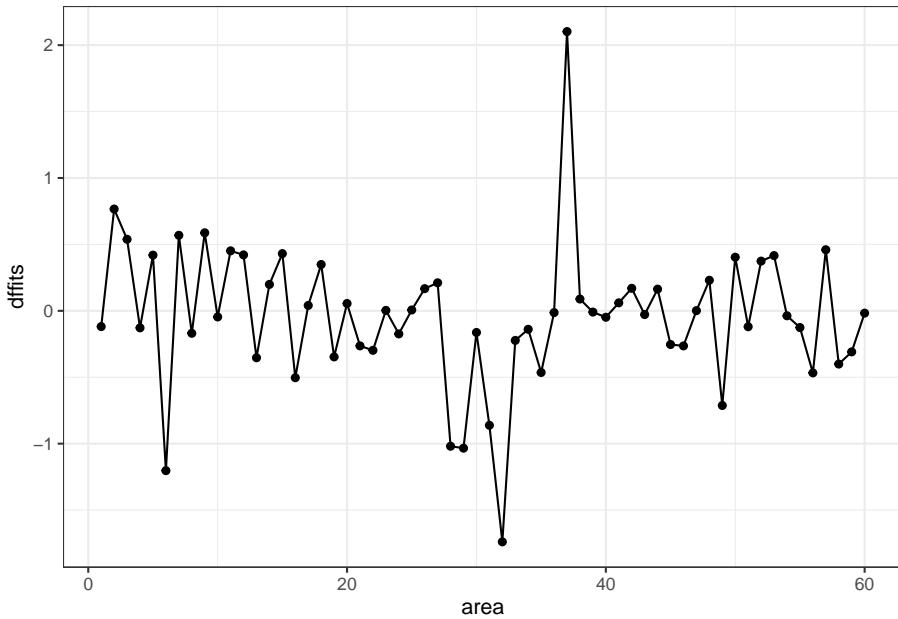
[,1]      [,2]      [,3]      [,4]      [,5]      [,6]
[1,]  0.03071160 -0.023775487 -0.004055111  0.01205425 -0.03260003 -0.02392315
[2,] -0.38276573 -0.048404993 -0.142293606  0.17009666 -0.22350621  0.44737372
[3,]  0.17226780 -0.426153536  0.350913139 -0.32949129  0.25777913 -0.10263448
[4,]  0.06175110 -0.006460916  0.024828272 -0.03009337  0.04154812 -0.06254145
[5,]  0.16875200  0.039839994 -0.058178534  0.06449504 -0.07772208 -0.18058630
[6,]  0.03322073  0.112699877 -0.203543632  0.23987378 -0.35201736 -0.04075617
[,7]      [,8]      [,9]      [,10]     [,11]
[1,]  0.01175375 -0.06494414  0.060929683 -0.011042644  0.03425156
[2,] -0.48562818  0.19372285 -0.212186731 -0.107830147 -0.01503250
[3,]  0.05005284 -0.02049877  0.014059330  0.010793169  0.04924166
[4,]  0.05498432  0.01135031 -0.001877983 -0.005490454 -0.01254111
[5,]  0.16151742  0.02723710  0.065483158  0.003326357 -0.05570035
[6,]  0.02900006 -0.21508009  0.171627718  0.019241676  0.05775536
```

10.5.2 Using the `residuals` command for other summaries

The `residuals` command will also let you get ordinary residuals, leverage values and `dffits` values, which are the normalized differences in predicted values when observations are omitted. See `?residuals.ols` for more details.

```
temp <- data.frame(area = 1:60)
temp$residual <- residuals(newmod, type = "ordinary")
temp$leverage <- residuals(newmod, type = "hat")
temp$dffits <- residuals(newmod, type = "dffits")
tbl_df(temp)
```

```
# A tibble: 60 x 4
  area residual leverage dffits
  <int>    <dbl>     <dbl>    <dbl>
1     1   -13.3    0.0929 -0.119
2     2    81.0    0.0941  0.766
3     3    28.8    0.266   0.539
4     4   -12.5    0.117   -0.128
5     5    27.8    0.204   0.419
6     6   -40.4    0.416   -1.20
7     7    37.0    0.207   0.568
8     8   -14.3    0.145   -0.169
9     9    66.6    0.0863  0.587
10    10   -4.96   0.0997 -0.0460
# ... with 50 more rows
ggplot(temp, aes(x = area, y = dffits)) +
  geom_point() +
  geom_line()
```



It appears that point 37 has the largest (positive) `dffits` value. Recall that point 37 seemed influential on several predictors and the intercept term. Point 32 has the smallest (or largest negative) `dffits`, and also appears to have been influential on several predictors and the intercept.

```
which.max(temp$dffits)
```

```
[1] 37
```

```
which.min(temp$dffits)
```

```
[1] 32
```

10.6 Model Validation and Correcting for Optimism

In 431, we learned about splitting our regression models into **training** samples and **test** samples, performing variable selection work on the training sample to identify two or three candidate models (perhaps via a stepwise approach), and then comparing the predictions made by those models in a test sample.

At the final project presentations, I mentioned (to many folks) that there was a way to automate this process a bit in 432, that would provide some ways to get the machine to split the data for you multiple times, and then average over the results, using a bootstrap approach. This is it.

The `validate` function allows us to perform cross-validation of our models for some summary statistics (and then correct those statistics for optimism in describing likely predictive accuracy) in an easy way.

`validate` develops:

- Resampling validation with or without backward elimination of variables
- Estimates of the *optimism* in measures of predictive accuracy
- Estimates of the intercept and slope of a calibration model

$$(\text{observed } y) = \text{Intercept} + \text{Slope} (\text{predicted } y)$$

with the following code...

```
set.seed(432002); validate(newmod, method = "boot", B = 40)
```

	index.orig	training	test	optimism	index.corrected	n
R-square	0.6989	0.7426	0.5749	0.1676	0.5312	40
MSE	1145.7800	963.9565	1617.4042	-653.4478	1799.2278	40
g	58.9614	59.7891	54.6444	5.1447	53.8168	40
Intercept	0.0000	0.0000	96.6990	-96.6990	96.6990	40
Slope	1.0000	1.0000	0.8961	0.1039	0.8961	40

So, for `R-square` we see that our original estimate was 0.6989

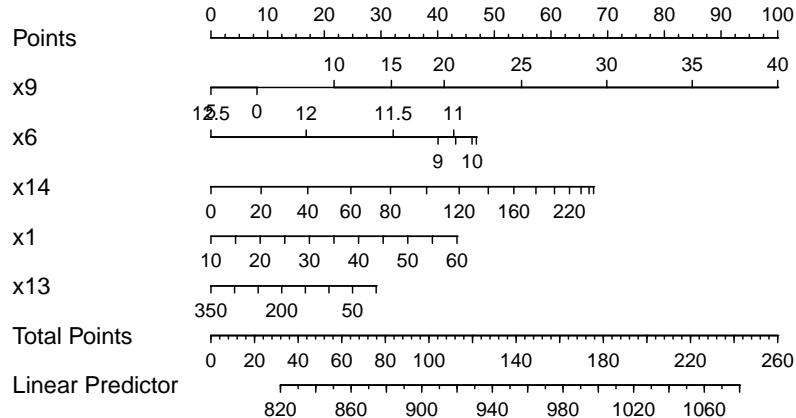
- Our estimated `R-square` across `n = 40` training samples was 0.7426, but in the resulting tests, the average `R-square` was only 0.5749
- This suggests an optimism of $0.7426 - 0.5749 = 0.1676$ (after rounding).
- We then apply that optimism to obtain a new estimate of R^2 corrected for overfitting, at 0.5312, which is probably a better estimate of what our results might look like in new data that were similar to (but not the same as) the data we used in building `newmod` than our initial estimate of 0.6989

We also obtain optimism-corrected estimates of the mean squared error (square of the residual standard deviation), the `g` index, and the intercept and slope of the calibration model. The “corrected” slope is a shrinkage factor that takes overfitting into account.

10.7 Building a Nomogram for Our Model

Another nice feature of an `ols` model object is that we can picture the model with a `nomogram` easily. Here is model `newmod`.

```
plot(nomogram(newmod))
```



For this model, we can use this plot to predict y as follows:

1. find our values of x_9 on the appropriate line
2. draw a vertical line up to the points line to count the points associated with our subject
3. repeat the process to obtain the points associated with x_6 , x_{14} , x_1 , and x_{13} . Sum the points.
4. draw a vertical line down from that number in the Total Points line to estimate y (the Linear Predictor) = Age-Adjusted Mortality Rate.

The impact of the non-linearity is seen in the x_6 results, for example, which turn around from 9-10 to 11-12. We also see non-linearity's effects in the scales of the non-linear terms in terms of points awarded.

An area with a combination of predictor values leading to a total of 100 points, for instance, would lead to a prediction of a Mortality Rate near 905. An area with a total of 140 points would have a predicted Mortality Rate of 955, roughly.

Chapter 11

Logistic Regression: The Foundations

Sources for this material include Harrell (2001), Harrell (2018), Ramsey and Schafer (2002) (chapters 20-21), Vittinghoff et al. (2012) (chapter 5) and Faraway (2006) (chapter 2).

11.1 A First Attempt: A Linear Probability Model

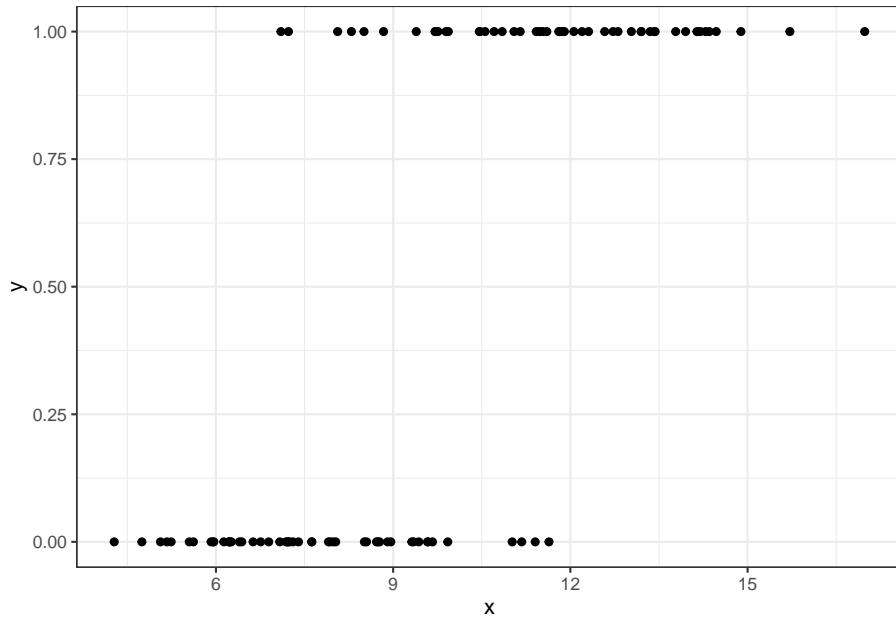
Suppose we want to predict a binary outcome which takes on the value 1 or 0, based on a single quantitative predictor. Let y be a 1/0 outcome, and x be a quantitative predictor in the following simulation.

```
set.seed(432)
sim12 <- data_frame(x = rnorm(100, 10, 3),
                      err = rnorm(100, 0, 2),
                      y = ifelse(x + err > 10, 1, 0))

Warning: `data_frame()` was deprecated in tibble 1.1.0.
Please use `tibble()` instead.
This warning is displayed once every 8 hours.
Call `lifecycle::last_lifecycle_warnings()` to see where this warning was generated.

sim12 <- select(sim12, x, y)

ggplot(sim12, aes(x = x, y = y)) + geom_point()
```



Now, we want to use our variable x here to predict our variable y (which takes on the values 0 and 1).

One approach to doing this would be a linear probability model, as follows:

```
mod12a <- lm(y ~ x, data = sim12)
summary(mod12a)
```

```
Call:
lm(formula = y ~ x, data = sim12)

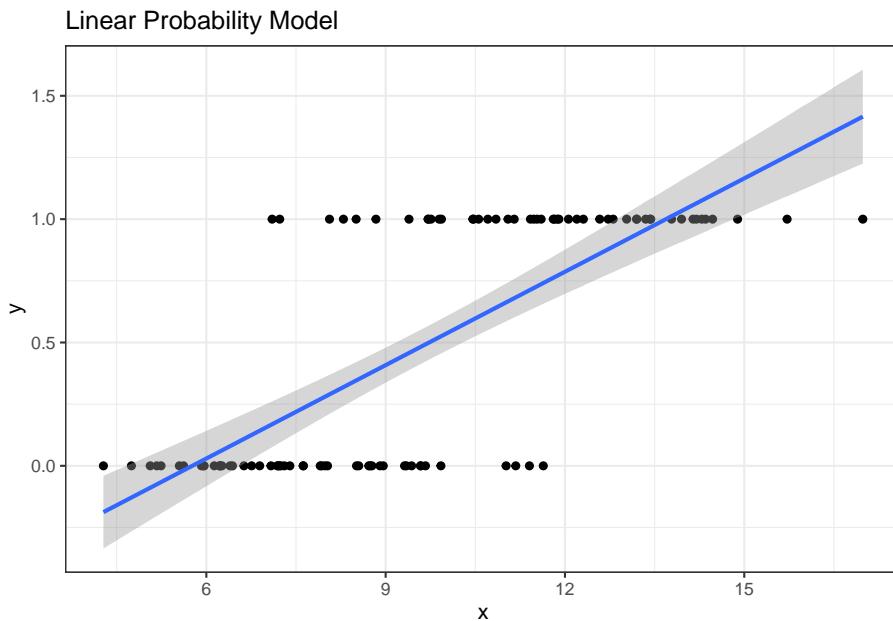
Residuals:
    Min      1Q  Median      3Q     Max 
-0.74104 -0.23411 -0.02894  0.23117  0.83153 

Coefficients:
            Estimate Std. Error t value Pr(>|t|)    
(Intercept) -0.72761   0.12272 -5.929 4.57e-08 ***
x            0.12620   0.01219 10.349 < 2e-16 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.3491 on 98 degrees of freedom
Multiple R-squared:  0.5222,    Adjusted R-squared:  0.5173 
F-statistic: 107.1 on 1 and 98 DF,  p-value: < 2.2e-16
```

Here's a picture of this model. What's wrong here?

```
ggplot(sim12, aes(x = x, y = y)) +
  geom_point() +
  geom_smooth(method = "lm") +
  labs(title = "Linear Probability Model")  
  
`geom_smooth()` using formula 'y ~ x'
```



If y can only take the values 0 and 1 (or, more precisely, if we're trying to predict the value $\pi = \Pr(y = 1)$) then what do we do with the predictions that are outside the range of $(0, 1)$?

11.2 Logistic Regression

Logistic regression is the most common model used when the outcome is binary. Our response variable is assumed to take on two values, zero or one, and we then describe the probability of a “one” response, given a linear function of explanatory predictors. We use logistic regression rather than linear regression for predicting binary outcomes. Linear regression approaches to the problem of predicting probabilities are problematic for several reasons - not least of which being that they predict probabilities greater than one and less than zero. There are several available alternatives, including probit regression and binomial regression, for the problem of predicting a binary outcome.

Logistic regression is part of a class called **generalized linear models** which extend the linear regression model in a variety of ways. There are also several extensions to the logistic regression model, including multinomial logistic regression (which is used for nominal categorical outcomes with more than two levels) and ordered logistic regression (used for ordered multi-categorical outcomes.) The methods involved in binary logistic regression may also be extended to the case where the outcomes are proportions based on counts, often through grouped binary responses (the proportion of cells with chromosomal aberrations, or the proportion of subjects who develop a particular condition.)

Although the models are different in some crucial ways, the practical use of logistic regression tracks well with much of what we've learned about linear regression.

11.3 The Logistic Regression Model

A generalized linear model (or GLM) is a probability model in which the mean of an outcome is related to predictors through a regression equation. A link function g is used to relate the mean, μ , to a linear regression of the predictors X_1, X_2, \dots, X_k .

$$g(\mu) = \beta_0 + \beta_1 X_1 + \beta_2 X_2 + \dots + \beta_k X_k$$

In the case of a logistic regression model,

- the mean μ of our 0/1 outcome is represented by π which describes the probability of a “1” outcome.
- the linking function we use in logistic regression makes use of the logit function, which is built on the natural logarithm.

11.4 The Link Function

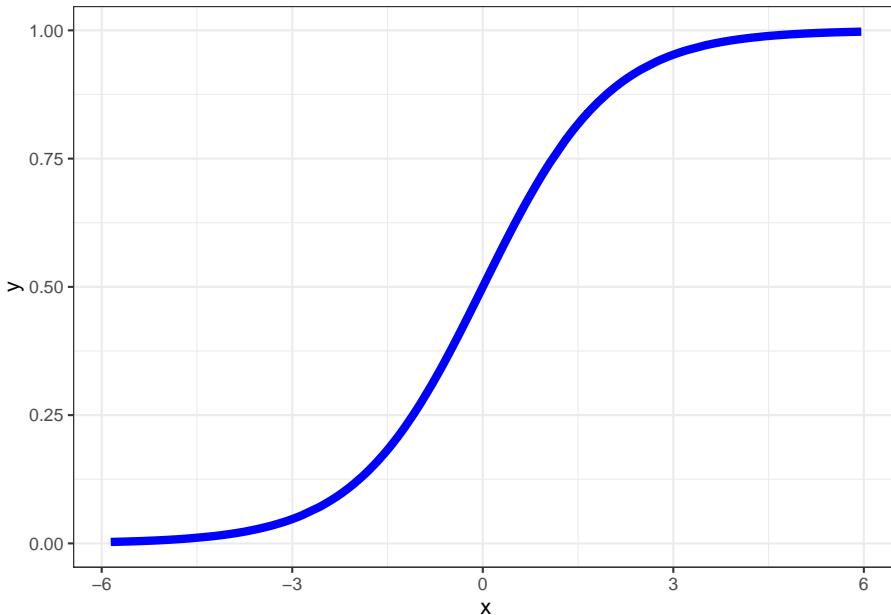
Logistic regression is a non-linear regression approach, since the equation for the mean of the 0/1 Y values conditioned on the values of our predictors X_1, X_2, \dots, X_k turns out to be non-linear in the β coefficients. Its nonlinearity, however, is solely found in its link function, hence the term *generalized* linear model.

The particular link function we use in logistic regression is called the **logit link**.

$$\text{logit}(\pi) = \log\left(\frac{\pi}{1-\pi}\right) = \beta_0 + \beta_1 X_1 + \beta_2 X_2 + \dots + \beta_k X_k$$

The inverse of the logit function is called the **logistic function**. If $\text{logit}(\pi) = \eta$, then $\pi = \frac{\exp(\eta)}{1+\exp(\eta)}$

The plot below displays the logistic function $y = \frac{e^x}{1+e^x}$



As you can see in the figure above, the logistic function $\frac{e^x}{1+e^x}$ takes any value x in the real numbers and returns a value between 0 and 1.

11.5 The logit or log odds

We usually focus on the **logit** in statistical work, which is the inverse of the logistic function.

- If we have a probability $\pi < 0.5$, then $\text{logit}(\pi) < 0$.
- If our probability $\pi > 0.5$, then $\text{logit}(\pi) > 0$.
- Finally, if $\pi = 0.5$, then $\text{logit}(\pi) = 0$.

11.6 Interpreting the Coefficients of a Logistic Regression Model

The critical thing to remember in interpreting a logistic regression model is that the logit is the log odds function. Exponentiating the logit yields the odds.

So, suppose we have a yes/no outcome variable, where yes = 1, and no = 0, and $\pi = \Pr(y = 1)$. Our model holds that:

$$\text{logit}(\pi) = \log\left(\frac{\pi}{1 - \pi}\right) = \beta_0 + \beta_1 X_1 + \beta_2 X_2 + \dots + \beta_k X_k$$

The odds of a yes response (the odds that $Y = 1$) at the level X_1, X_2, \dots, X_k are:

$$\text{Odds}(Y = 1) = \exp(\beta_0 + \beta_1 X_1 + \beta_2 X_2 + \dots + \beta_k X_k)$$

The **probability** of a yes response ($\Pr(y = 1)$, or π) is just

$$\pi = \Pr(Y = 1) = \frac{\text{Odds}(Y = 1)}{1 + \text{Odds}(Y = 1)} = \frac{\exp(\beta_0 + \beta_1 X_1 + \beta_2 X_2 + \dots + \beta_k X_k)}{1 + \exp(\beta_0 + \beta_1 X_1 + \beta_2 X_2 + \dots + \beta_k X_k)}$$

11.7 The Logistic Regression has non-constant variance

In ordinary least squares regression, the variance $\text{Var}(Y|X_1, X_2, \dots, X_k) = \sigma^2$ is a constant that does not depend on the predictor values. This is not the case in logistic regression. The mean and variance specifications of the logistic regression model are quite different.

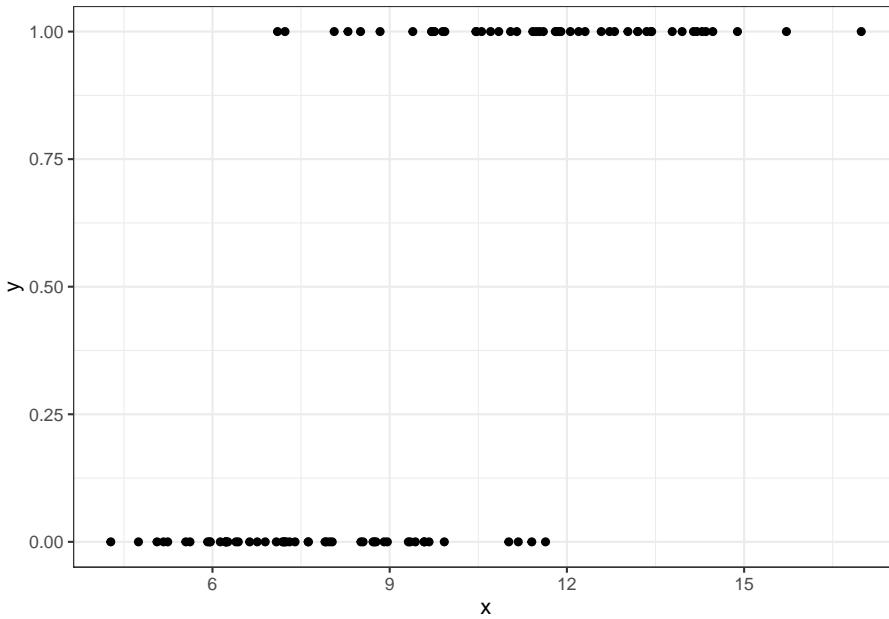
$$\text{logit}(\pi) = \log\left(\frac{\pi}{1 - \pi}\right) = \beta_0 + \beta_1 X_1 + \beta_2 X_2 + \dots + \beta_k X_k \mu[Y|X_1, \dots, X_k] = \pi, \text{Var}[Y|X_1, \dots, X_k] = \pi(1 - \pi)$$

The variance is now a function of the mean, and contains no additional parameter for us to estimate.

11.8 Fitting a Logistic Regression Model to our Simulated Data

Recall the `sim12` data we built earlier.

```
ggplot(sim12, aes(x = x, y = y)) + geom_point()
```



Here is the fitted logistic regression model.

```
model12b <- glm(y ~ x, data = sim12, family = binomial)

model12b

Call: glm(formula = y ~ x, family = binomial, data = sim12)

Coefficients:
(Intercept)          x
-9.1955        0.9566

Degrees of Freedom: 99 Total (i.e. Null);  98 Residual
Null Deviance:    138.6
Residual Deviance: 70.03    AIC: 74.03
```

The logistic regression equation is:

$$\text{logit}(\Pr(y = 1)) = \log \left(\frac{\Pr(y = 1)}{1 - \Pr(y = 1)} \right) = -9.1955 + 0.9566x$$

We can exponentiate the results of this model to get to an equation about odds, and eventually, a prediction about probabilities. Suppose, for instance, that we are interested in the prediction when $x = 12$.

$$\text{logit}(\Pr(y = 1)|X = 12) = \log\left(\frac{\Pr(y = 1)}{1 - \Pr(y = 1)}\right) = -9.1955 + 0.9566 \cdot 12 = 2.2837$$

And we can also get this from the `predict` function applied to our model, although the `predict` approach retains a few more decimal places internally:

```
predict(model12b, newdata = data.frame(x = 12))
```

```
1  
2.284069
```

$$\text{Odds}(Y = 1|X = 12) = \exp(-9.20 + 0.96 \cdot 12) = \exp(2.2837) = 9.812921$$

```
exp(predict(model12b, newdata = data.frame(x = 12)))
```

```
1  
9.81654
```

The estimated **probability** of a yes response ($\Pr(y = 1)$, or π) if $x = 12$ is just

$$\pi = \Pr(Y = 1|X = 12) = \frac{\text{Odds}(Y = 1|X = 12)}{1 + \text{Odds}(Y = 1|X = 12)} = \frac{\exp(-9.20 + 0.96x)}{1 + \exp(-9.20 + 0.96x)} = \frac{9.812921}{1 + 9.812921} = 0.908$$

Does this work out?

```
exp(predict(model12b, newdata = data.frame(x = 12))) /  
(1 + exp(predict(model12b, newdata = data.frame(x = 12))))
```

```
1  
0.907549
```

which is also directly available by running `predict` with `type = "response"`.

```
predict(model12b, newdata = data.frame(x = 12), type = "response")
```

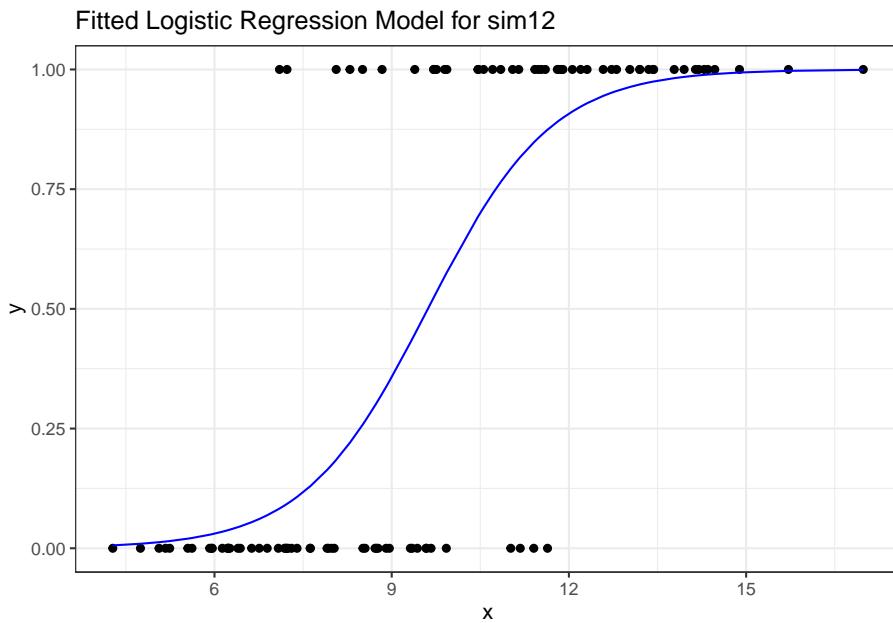
```
1  
0.907549
```

11.9 Plotting the Logistic Regression Model

We can use the `augment` function from the `broom` package to get our fitted probabilities included in the data.

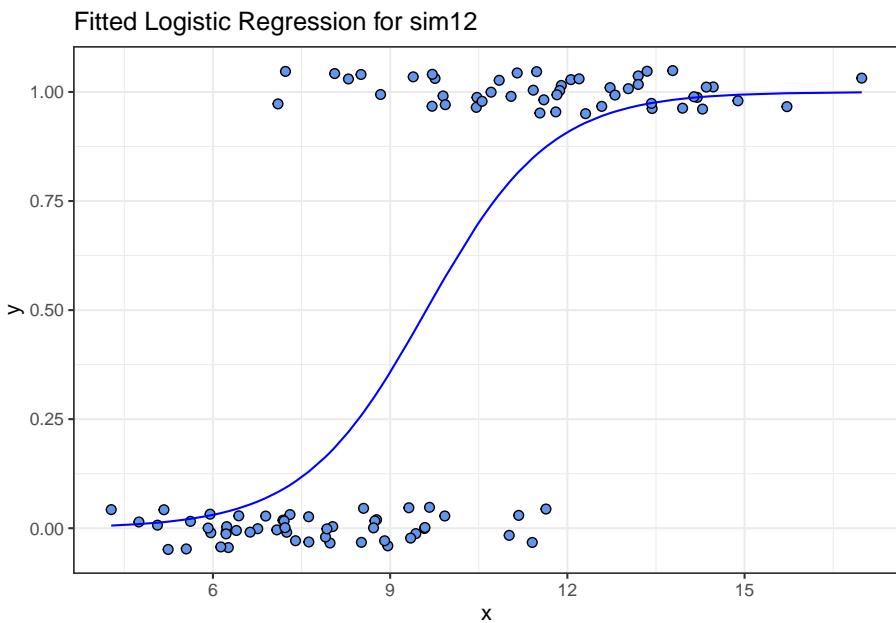
```
m12b.aug <- augment(model12b, sim12, type.predict = "response")

ggplot(m12b.aug, aes(x = x, y = y)) +
  geom_point() +
  geom_line(aes(x = x, y = .fitted), col = "blue") +
  labs(title = "Fitted Logistic Regression Model for sim12")
```



I'll add a little jitter on the vertical scale to the points, so we can avoid overlap, and also make the points a little bigger.

```
ggplot(m12b.aug, aes(x = x, y = y)) +
  geom_jitter(height = 0.05, size = 2, pch = 21,
              fill = "cornflowerblue") +
  geom_line(aes(x = x, y = .fitted), col = "blue") +
  labs(title = "Fitted Logistic Regression for sim12") +
  theme_bw()
```



All right, it's time to move on to fitting models. We'll do that in the next Chapter.

Chapter 12

Logistic Regression and the resect data

12.1 The resect data

My source for these data was Riffenburgh (2006). The data describe 134 patients who had undergone resection of the tracheal carina (most often this is done to address tumors in the trachea), and the `resect.csv` data file contains the following variables:

- `id` = a patient ID #,
- `age` = the patient's age at surgery,
- `prior` = prior tracheal surgery (1 = yes, 0 = no),
- `resection` = extent of the resection (in cm),
- `intubated` = whether intubation was required at the end of surgery (1 = yes, 0 = no), and
- `died` = the patient's death status (1 = dead, 0 = alive).

```
miss_var_summary(resect)
```

```
# A tibble: 6 x 3
  variable n_miss pct_miss
  <chr>     <int>    <dbl>
1 id         0        0
2 age        0        0
3 prior      0        0
4 resection   0        0
5 intubated   0        0
6 died        0        0
```

```

resect %>% count(died, prior)

# A tibble: 4 x 3
  died prior   n
  <dbl> <dbl> <int>
1     0     0    89
2     0     1    28
3     1     0    11
4     1     1     6

resect %>% mosaic::inspect()

quantitative variables:
      name   class min   Q1 median   Q3 max       mean        sd   n
...1      id numeric  1 34.25  67.5 100.75 134 67.5000000 38.8265373 134
...2      age numeric  8 36.00  51.0  61.00  80 47.8432836 15.7775202 134
...3    prior numeric  0  0.00   0.0  0.75   1  0.2537313  0.4367785 134
...4 resection numeric  1  2.00   2.5  4.00   6  2.9634328  1.2402123 134
...5 intubated numeric  0  0.00   0.0  0.00   1  0.1417910  0.3501447 134
...6    died numeric  0  0.00   0.0  0.00   1  0.1268657  0.3340713 134
      missing
...1      0
...2      0
...3      0
...4      0
...5      0
...6      0

```

We have no missing data, and 17 of the 134 patients died. Our goal will be to understand the characteristics of the patients, and how they relate to the binary outcome of interest, death.

12.2 Running A Simple Logistic Regression Model

In the most common scenario, a logistic regression model is used to predict a binary outcome (which can take on the values 0 or 1.) We will eventually fit a logistic regression model in two ways.

1. Through the `glm` function in the base package of R (similar to `lm` for linear regression)
2. Through the `lrm` function available in the `rms` package (similar to `ols` for linear regression)

We'll focus on the `glm` approach first, and save the `lrm` ideas for later in this Chapter.

12.2.1 Logistic Regression Can Be Harder than Linear Regression

- Logistic regression models are fitted using the method of maximum likelihood in `glm`, which requires multiple iterations until convergence is reached.
- Logistic regression models are harder to interpret (for most people) than linear regressions.
- Logistic regression models don't have the same set of assumptions as linear models.
- Logistic regression outcomes (yes/no) carry much less information than quantitative outcomes. As a result, fitting a reasonable logistic regression requires more data than a linear model of similar size.
 - The rule I learned in graduate school was that a logistic regression requires 100 observations to fit an intercept plus another 15 observations for each candidate predictor. That's not terrible, but it's a very large sample size.
 - Frank Harrell recommends that 96 observations + a function of the number of candidate predictors (which depends on the amount of variation in the predictors, but $15 \times$ the number of such predictors isn't too bad if the signal to noise ratio is pretty good) are required just to get reasonable confidence intervals around your predictions.
 - * In a twitter note, Frank suggests that $96 + 8$ times the number of candidate parameters might be reasonable so long as the smallest cell of interest (combination of an outcome and a split of the covariates) is 96 or more observations.
 - Peduzzi et al. (1996) suggest that if we let π be the smaller of the proportions of "yes" or "no" cases in the population of interest, and k be the number of inputs under consideration, then $N = 10k/\pi$ is the minimum number of cases to include, except that if $N < 100$ by this standard, you should increase it to 100, according to Long (1997).
 - * That suggests that if you have an outcome that happens 10% of the time, and you are running a model with 3 predictors, then you could get away with $(10 \times 3)/(0.10) = 300$ observations, but if your outcome happened 40% of the time, you could get away with only $(10 \times 3)/(0.40) = 75$ observations, which you'd round up to 100.

12.3 Logistic Regression using `glm`

We'll begin by attempting to predict death based on the extent of the resection.

```

res_modA <- glm(died ~ resection, data=resect,
                  family="binomial"(link="logit"))

res_modA

Call: glm(formula = died ~ resection, family = binomial(link = "logit"),
          data = resect)

Coefficients:
(Intercept)      resection
-4.4337        0.7417

Degrees of Freedom: 133 Total (i.e. Null); 132 Residual
Null Deviance: 101.9
Residual Deviance: 89.49    AIC: 93.49

```

Note that the `logit` link is the default approach with the `binomial` family, so we could also have used:

```

res_modA <- glm(died ~ resection, data = resect,
                  family = "binomial")

```

which yields the same model.

12.3.1 Interpreting the Coefficients of a Logistic Regression Model

Our model is:

$$\text{logit}(\text{died} = 1) = \log \left(\frac{\Pr(\text{died} = 1)}{1 - \Pr(\text{died} = 1)} \right) = \beta_0 + \beta_1 x = -4.4337 + 0.7417 \times \text{resection}$$

The predicted log odds of death for a subject with a resection of 4 cm is:

$$\log \left(\frac{\Pr(\text{died} = 1)}{1 - \Pr(\text{died} = 1)} \right) = -4.4337 + 0.7417 \times 4 = -1.467$$

The predicted odds of death for a subject with a resection of 4 cm is thus:

$$\frac{\Pr(\text{died} = 1)}{1 - \Pr(\text{died} = 1)} = e^{-4.4337+0.7417\times 4} = e^{-1.467} = 0.2306$$

Since the odds are less than 1, we should find that the probability of death is less than 1/2. With a little algebra, we see that the predicted probability of death for a subject with a resection of 4 cm is:

$$Pr(died = 1) = \frac{e^{-4.4337+0.7417 \times 4}}{1 + e^{-4.4337+0.7417 \times 4}} = \frac{e^{-1.467}}{1 + e^{-1.467}} = \frac{0.2306}{1.2306} = 0.187$$

In general, we can frame the model in terms of a statement about probabilities, like this:

$$Pr(died = 1) = \frac{e^{\beta_0 + \beta_1 x}}{1 + e^{\beta_0 + \beta_1 x}} = \frac{e^{-4.4337 + 0.7417 \times resection}}{1 + e^{-4.4337 + 0.7417 \times resection}}$$

and so by substituting in values for `resection`, we can estimate the model's fitted probabilities of death.

12.3.2 Using `predict` to describe the model's fits

To obtain these fitted odds and probabilities in R, we can use the `predict` function.

- The default predictions are on the scale of the log odds. These predictions are also available through the `type = "link"` command within the `predict` function for a generalized linear model like logistic regression.
- Here are the predicted log odds of death for a subject (Sally) with a 4 cm resection and a subject (Harry) who had a 5 cm resection.

```
predict(res_modA, newdata = tibble(resection = c(4,5)))
```

```
1           2
-1.4669912 -0.7253027
```

- We can also obtain predictions for each subject on the original response (here, probability) scale, backing out of the logit link.

```
predict(res_modA, newdata = tibble(resection = c(4, 5)),
        type = "response")
```

```
1           2
0.1874004 0.3262264
```

So the predicted probability of death is 0.187 for Sally, the subject with a 4 cm resection, and 0.326 for Harry, the subject with a 5 cm resection.

12.3.3 Odds Ratio interpretation of Coefficients

Often, we will exponentiate the estimated slope coefficients of a logistic regression model to help us understand the impact of changing a predictor on the odds of our outcome.

```
exp(coef(res_modA))

(Intercept)    resection
0.01186995  2.09947754
```

To interpret this finding, suppose we have two subjects, Harry and Sally. Harry had a resection that was 1 cm larger than Sally. This estimated coefficient suggests that the estimated odds for death associated with Harry is 2.099 times larger than the odds for death associated with Sally. In general, the odds ratio comparing two subjects who differ by 1 cm on the resection length is 2.099.

To illustrate, again let's assume that Harry's resection was 5 cm, and Sally's was 4 cm. Then we have:

$$\log \left(\frac{Pr(Harrydied)}{1 - Pr(Harrydied)} \right) = -4.4337 + 0.7417 \times 5 = -0.7253, \log \left(\frac{Pr(Sallydied)}{1 - Pr(Sallydied)} \right) = -4.4337 + 0.7417 \times 4 = -0.2307$$

which implies that our estimated odds of death for Harry and Sally are:

$$Odds(Harrydied) = \frac{Pr(Harrydied)}{1 - Pr(Harrydied)} = e^{-4.4337 + 0.7417 \times 5} = e^{-0.7253} = 0.4842 Odds(Sallydied) = \frac{Pr(Sallydied)}{1 - Pr(Sallydied)} = e^{-4.4337 + 0.7417 \times 4} = e^{-0.2307} = 2.099$$

and so the odds ratio is:

$$OR = \frac{Odds(Harrydied)}{Odds(Sallydied)} = \frac{0.4842}{0.2307} = 2.099$$

- If the odds ratio was 1, that would mean that Harry and Sally had the same estimated odds of death, and thus the same estimated probability of death, despite having different sizes of resections.
- Since the odds ratio is greater than 1, it means that Harry has a higher estimated odds of death than Sally, and thus that Harry has a higher estimated probability of death than Sally.
- If the odds ratio was less than 1, it would mean that Harry had a lower estimated odds of death than Sally, and thus that Harry had a lower estimated probability of death than Sally.

Remember that the odds ratio is a fraction describing two positive numbers (odds can only be non-negative) so that the smallest possible odds ratio is 0.

12.3.4 Interpreting the rest of the model output from `glm`

```
res_modA

Call: glm(formula = died ~ resection, family = "binomial", data = resect)

Coefficients:
(Intercept)    resection
-4.4337       0.7417

Degrees of Freedom: 133 Total (i.e. Null); 132 Residual
Null Deviance: 101.9
Residual Deviance: 89.49    AIC: 93.49
```

In addition to specifying the logistic regression coefficients, we are also presented with information on degrees of freedom, deviance (null and residual) and AIC.

- The degrees of freedom indicate the sample size.
 - Recall that we had $n = 134$ subjects in the data. The “Null” model includes only an intercept term (which uses 1 df) and we thus have $n - 1$ (here 133) degrees of freedom available for estimation.
 - In our `res_modA` model, a logistic regression is fit including a single slope (`resection`) and an intercept term. Each uses up one degree of freedom to build an estimate, so we have $n - 2 = 134 - 2 = 132$ residual df remaining.
- The AIC or Akaike Information Criterion (lower values are better) is also provided. This is helpful if we’re comparing multiple models for the same outcome.

12.3.5 Deviance and Comparing Our Model to the Null Model

- The deviance (a measure of the model’s *lack of fit*) is available for both the null model (the model with only an intercept) and for our model (`res_modA`) predicting our outcome, mortality.
- The deviance test, though available in R (see below) isn’t really a test of whether the model works well. Instead, it assumes the model is true, and then tests to see if the coefficients are detectably different from zero. So it isn’t of much practical use.
 - To compare the `deviance` statistics, we can subtract the residual deviance from the null deviance to describe the impact of our model on fit.
 - Null Deviance - Residual Deviance can be compared to a χ^2 distribution with Null DF - Residual DF degrees of freedom to obtain a global test of the in-sample predictive power of our model.
 - We can see this comparison more directly by running `anova` on our model:

```
anova(res_modA, test = "LRT")

Analysis of Deviance Table

Model: binomial, link: logit

Response: died

Terms added sequentially (first to last)

          Df Deviance Resid. Df Resid. Dev  Pr(>Chi)
NULL           133    101.943
resection     1     12.45      132    89.493 0.0004179 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

The `test = "LRT"` section completes a deviance test and provides a p value, which just estimates the probability that a chi-square distribution with a single degree of freedom would exhibit an improvement in deviance as large as 12.45.

The p value for the deviance test here is about 0.0004. But, again, this isn't a test of whether the model is any good - it assumes the model is true, and then tests some consequences.

- Specifically, it tests whether (if the model is TRUE) some of the model's coefficients are non-zero.
- That's not so practically useful, so I discourage you from performing global tests of a logistic regression model with a deviance test.

12.3.6 Using `glance` with a logistic regression model

We can use the `glance` function from the `broom` package to obtain the null and residual deviance and degrees of freedom. Note that the deviance for our model is related to the log likelihood by $-2\log\text{Lik}$.

```
glance(res_modA)

# A tibble: 1 x 8
  null.deviance df.null logLik   AIC   BIC deviance df.residual nobs
        <dbl>     <int>  <dbl> <dbl> <dbl>     <dbl>       <int> <int>
1       102.      133  -44.7  93.5  99.3    89.5       132     134
```

The `glance` result also provides the AIC, and the BIC (Bayes Information Criterion), each of which is helpful in understanding comparisons between multiple models for the same outcome (with smaller values of either criterion indicating better models.) The AIC is based on the deviance, but penalizes you

for making the model more complicated. The BIC does the same sort of thing but with a different penalty.

Again we see that we have a null deviance of 101.94 on 133 degrees of freedom. Including the `resection` information in the model decreased the deviance to 89.49 points on 132 degrees of freedom, so that's a decrease of 12.45 points while using one degree of freedom, a statistically significant reduction in deviance.

12.4 Interpreting the Model Summary

Let's get a more detailed summary of our `res_modA` model, including 95% confidence intervals for the coefficients:

```
summary(res_modA)

Call:
glm(formula = died ~ resection, family = "binomial", data = resect)

Deviance Residuals:
    Min      1Q  Median      3Q     Max 
-1.1844 -0.5435 -0.3823 -0.2663  2.4501 

Coefficients:
            Estimate Std. Error z value Pr(>|z|)    
(Intercept) -4.4337    0.8799 -5.039 4.67e-07 ***
resection     0.7417    0.2230  3.327 0.000879 ***  
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

Null deviance: 101.943  on 133  degrees of freedom
Residual deviance: 89.493  on 132  degrees of freedom
AIC: 93.493

Number of Fisher Scoring iterations: 5
confint(res_modA, level = 0.95)

Waiting for profiling to be done...

              2.5 %    97.5 %    
(Intercept) -6.344472 -2.855856  
resection     0.322898  1.208311
```

Some elements of this summary are very familiar from our work with linear models.

- We still have a five-number summary of residuals, although these are called *deviance* residuals.
- We have a table of coefficients with standard errors, and hypothesis tests, although these are Wald z-tests, rather than the t tests we saw in linear modeling.
- We have a summary of global fit in the comparison of null deviance and residual deviance, but without a formal p value. And we have the AIC, as discussed above.
- We also have some new items related to a *dispersion* parameter and to the number of Fisher Scoring Iterations.

Let's walk through each of these elements.

12.4.1 Wald Z tests for Coefficients in a Logistic Regression

The coefficients output provides the estimated coefficients, and their standard errors, plus a Wald Z statistic, which is just the estimated coefficient divided by its standard error. This is compared to a standard Normal distribution to obtain the two-tailed p values summarized in the `Pr(>|z|)` column.

- The interesting result is `resection`, which has a Wald Z = 3.327, yielding a p value of 0.00088.
- The p value assesses whether the estimated coefficient of `resection`, 0.7417, is statistically detectably different from 0. If the coefficient (on the logit scale) for `resection` was truly 0, this would mean that:
 - the log odds of death did not change based on the `resection` size,
 - the odds of death were unchanged based on the `resection` size (the odds ratio would be 1), and
 - the probability of death was unchanged based on the `resection` size.

In our case, we have a statistically detectable change in the log odds of `died` associated with changes in `resection`, according to this p value. We conclude that `resection` size is associated with a positive impact on death rates (death rates are generally higher for people with larger resections.)

12.4.2 Confidence Intervals for the Coefficients

As in linear regression, we can obtain 95% confidence intervals (to get other levels, change the `level` parameter in `confint`) for the intercept and slope coefficients.

Here, we see, for example, that the coefficient of `resection` has a point estimate of 0.7417, and a confidence interval of (0.3229, 1.208). Since this is on the logit scale, it's not that interpretable, but we will often exponentiate the model and

its confidence interval to obtain a more interpretable result on the odds ratio scale.

```
tidy(res_modA, exponentiate = TRUE, conf.int = TRUE) %>%
  select(term, estimate, conf.low, conf.high)

# A tibble: 2 x 4
  term      estimate conf.low conf.high
  <chr>     <dbl>    <dbl>    <dbl>
1 (Intercept) 0.0119  0.00176  0.0575
2 resection    2.10     1.38     3.35
```

From this output, we can estimate the odds ratio for death associated with a 1 cm increase in resection size is 2.099, with a 95% CI of (1.38, 3.35). - If the odds ratio was 1, it would indicate that the odds of death did not change based on the change in resection size. - Here, it's clear that the estimated odds of death will be larger (odds > 1) for subjects with larger resection sizes. Larger odds of death also indicate larger probabilities of death. This confidence interval indicates that with 95% confidence, we conclude that increases in resection size are associated with statistically detectable increases in the odds of death. - If the odds ratio was less than 1 (remember that it cannot be less than 0) that would mean that subjects with larger resection sizes were associated with smaller estimated odds of death.

12.4.3 Deviance Residuals

In logistic regression, it's certainly a good idea to check to see how well the model fits the data. However, there are a few different types of residuals. The residuals presented here by default are called deviance residuals. Other types of residuals are available for generalized linear models, such as Pearson residuals, working residuals, and response residuals. Logistic regression model diagnostics often make use of multiple types of residuals.

The deviance residuals for each individual subject sum up to the deviance statistic for the model, and describe the contribution of each point to the model likelihood function.

The deviance residual, d_i , for the i^{th} observation in a model predicting y_i (a binary variable), with the estimate being $\hat{\pi}_i$ is:

$$d_i = s_i \sqrt{-2[y_i \log \hat{\pi}_i + (1 - y_i) \log(1 - \hat{\pi}_i)]},$$

where s_i is 1 if $y_i = 1$ and $s_i = -1$ if $y_i = 0$.

Again, the sum of the deviance residuals is the deviance.

12.4.4 Dispersion Parameter

The dispersion parameter is taken to be 1 for `glm` fit using either the `binomial` or `Poisson` families. For other sorts of generalized linear models, the dispersion parameter will be of some importance in estimating standard errors sensibly.

12.4.5 Fisher Scoring iterations

The solution of a logistic regression model involves maximizing a likelihood function. Fisher's scoring algorithm in our `res_modA` needed five iterations to perform the logistic regression fit. All that this tells you is that the model converged, and didn't require a lot of time to do so.

12.5 Plotting a Simple Logistic Regression Model

Let's plot the logistic regression model `res_modA` for `died` using the extent of the resection in terms of probabilities. We can use either of two different approaches:

- we can plot the fitted values from our specific model against the original data, using the `augment` function from the `broom` package, or
- we can create a smooth function called `binomial_smooth` that plots a simple logistic model in an analogous way to `geom_smooth(method = "lm")` for a simple linear regression.

12.5.1 Using `augment` to capture the fitted probabilities

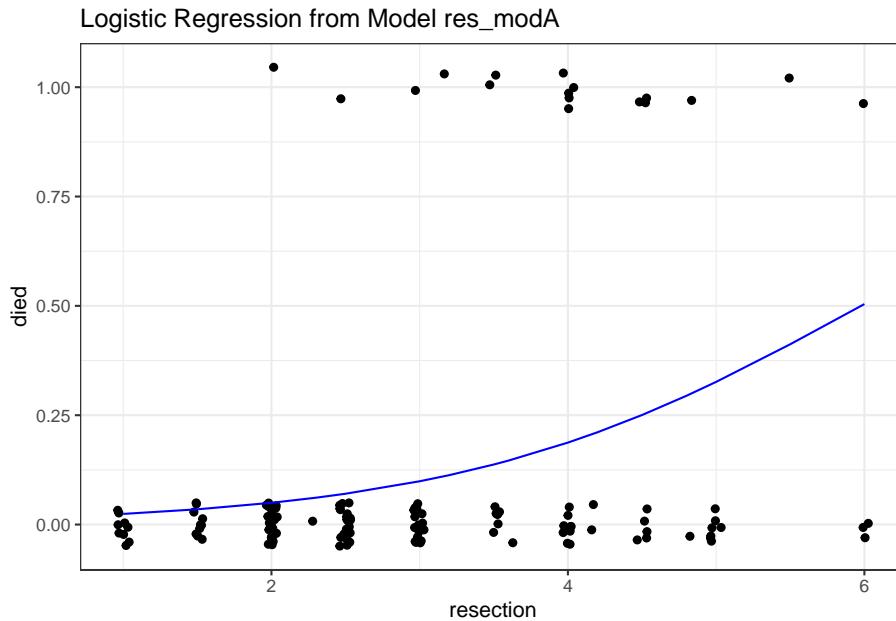
```
res_A_aug <- augment(res_modA, resect,
                      type.predict = "response")
head(res_A_aug)

# A tibble: 6 x 12
  id    age prior resection intubated   died .fitted .resid .std.resid     .hat
  <dbl> <dbl> <dbl>      <dbl>     <dbl> <dbl>   <dbl>   <dbl>      <dbl> <dbl>
1     1     34     1        2.5       0     0  0.0705 -0.382    -0.384 0.0100
2     2     57     0        5       0     0  0.326  -0.889    -0.904 0.0337
3     3     60     1        4       1     1  0.187   1.83     1.84  0.0120
4     4     62     1       4.2       0     0  0.211  -0.689    -0.693 0.0143
5     5     28     0        6       1     1  0.504   1.17     1.22  0.0818
6     6     52     0        3       0     0  0.0990 -0.457    -0.459 0.00922
# ... with 2 more variables: .sigma <dbl>, .cooks.d <dbl>
```

This approach augments the `resect` data set with fitted, residual and other summaries of each observation's impact on the fit, using the "response" type of prediction, which yields the fitted probabilities in the `.fitted` column.

12.5.2 Plotting a Logistic Regression Model's Fitted Values

```
ggplot(res_A_aug, aes(x = resection, y = died)) +
  geom_jitter(height = 0.05) +
  geom_line(aes(x = resection, y = .fitted),
            col = "blue") +
  labs(title = "Logistic Regression from Model res_modA")
```



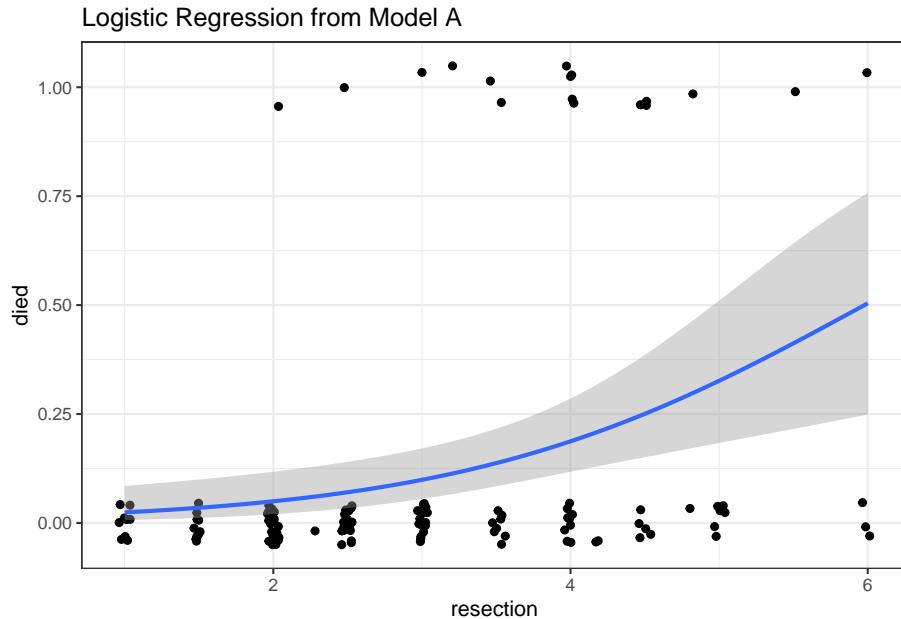
12.5.3 Plotting a Simple Logistic Model using `binomial_smooth`

```
binomial_smooth <- function(...) {
  geom_smooth(method = "glm",
              method.args = list(family = "binomial"), ...)
}

ggplot(resect, aes(x = resection, y = died)) +
  geom_jitter(height = 0.05) +
```

```
binomial_smooth() + ## ...smooth(se=FALSE) to leave out interval
labs(title = "Logistic Regression from Model A") +
theme_bw()

`geom_smooth()` using formula 'y ~ x'
```



As expected, we see an increase in the model probability of death as the extent of the resection grows larger.

12.6 How well does Model A classify subjects?

A natural question to ask is how well does our model classify patients in terms of likelihood of death.

We could specify a particular rule, for example: if the predicted probability of death is 0.5 or greater, then predict “Died.”

```
res_A_aug$rule.5 <- ifelse(res_A_aug$.fitted >= 0.5,
                           "Predict Died", "Predict Alive")

table(res_A_aug$rule.5, res_A_aug$died)
```

	0	1
Predict Alive	114	16

```
Predict Died 3 1
```

And perhaps build the linked table of row probabilities which tells us, for example, that 87.69% of the patients predicted by the model to be alive actually did survive.

```
round(100*prop.table(
  table(res_A_aug$rule.5, res_A_aug$died), 1), 2)
```

```
0      1
Predict Alive 87.69 12.31
Predict Died 75.00 25.00
```

Or the table of column probabilities which tell us, for example, that 97.44% of those who actually survived were predicted by the model to be alive.

```
round(100*prop.table(
  table(res_A_aug$rule.5, res_A_aug$died), 2), 2)
```

```
0      1
Predict Alive 97.44 94.12
Predict Died 2.56 5.88
```

We'll discuss various measures of concordance derived from this sort of classification later.

12.7 The Confusion Matrix

Let's build this misclassification table in standard epidemiological format.

```
confuseA_small <-
  res_A_aug %>%
    mutate(death_predicted = factor(.fitted >= 0.5),
           death_actual = factor(died == "1"),
           death_predicted = fct_relevel(death_predicted, "TRUE"),
           death_actual = fct_relevel(death_actual, "TRUE")) %$%
  table(death_predicted, death_actual)

confuseA_small
```

	death_actual	
death_predicted	TRUE	FALSE
TRUE	1	3
FALSE	16	114

In total, we have 134 observations.

- 115 correct predictions, or 85.8% accuracy
- 17 subjects who died, or 12.6% prevalence of death
- 4 subjects who were predicted to die, or 3.0% detection prevalence.

The sensitivity (also called recall) here is $1 / (1 + 16) = 5.9\%$.

- 5.9% of the subjects who actually died were predicted to die by the model.

The specificity here is $114 / (114 + 3) = 97.4\%$.

- 97.4% of the subjects who actually survived were predicted to survive by the model.

The positive predictive value (PPV: also called precision) is $1 / (1 + 3) = 25\%$

- Our predictions of death were correct 25% of the time.

The negative predictive value (NPV) is $114 / (114 + 16) = 87.7\%$

- Our predictions of survival were correct 87.7% of the time.

12.8 Using the `confusionMatrix` tool from the `caret` package

This provides a more detailed summary of the classification results from our logistic regression model.

```
res_A_aug %$%
  confusionMatrix(
    data = factor(.fitted >= 0.5),
    reference = factor(died == 1),
    positive = "TRUE"
  )

Confusion Matrix and Statistics

             Reference
Prediction   FALSE  TRUE
  FALSE      114   16
  TRUE        3     1

Accuracy : 0.8582
95% CI  : (0.7875, 0.9124)
No Information Rate : 0.8731
P-Value [Acc > NIR] : 0.747802

Kappa : 0.0493
```

12.9. RECEIVER OPERATING CHARACTERISTIC CURVE ANALYSIS303

McNemar's Test P-Value : 0.005905

```
Sensitivity : 0.058824
Specificity : 0.974359
Pos Pred Value : 0.250000
Neg Pred Value : 0.876923
Prevalence : 0.126866
Detection Rate : 0.007463
Detection Prevalence : 0.029851
Balanced Accuracy : 0.516591
```

'Positive' Class : TRUE

- The No Information Rate or NIR is just the percentage of correct predictions we'd get if we just predicted the more common classification (not dead) for every subject.
- Kappa is a correlation statistic ranging from -1 to +1. It measures the inter-rater reliability of our predictions and the true classifications, in this context. Complete agreement would be +1, and complete disagreement would be -1.

12.9 Receiver Operating Characteristic Curve Analysis

One way to assess the predictive accuracy within the model development sample in a logistic regression is to consider an analyses based on the receiver operating characteristic (ROC) curve. ROC curves are commonly used in assessing diagnoses in medical settings, and in signal detection applications.

The accuracy of a “test” can be evaluated by considering two types of errors: false positives and false negatives.

In our `res_modA` model, we use `resection` size to predict whether the patient `died`. Suppose we established a value `R`, so that if the resection size was larger than `R` cm, we would predict that the patient `died`, and otherwise we would predict that the patient did not die.

A good outcome of our model’s “test,” then, would be when the resection size is larger than `R` for a patient who actually died. Another good outcome would be when the resection size is smaller than `R` for a patient who survived.

But we can make errors, too.

- A false positive error in this setting would occur when the resection size is larger than `R` (so we predict the patient dies) but in fact the patient does

not die.

- A false negative error in this case would occur when the resection size is smaller than R (so we predict the patient survives) but in fact the patient dies.

Formally, the true positive fraction (TPF) for a specific resection cutoff R , is the probability of a positive test (a prediction that the patient will die) among the people who have the outcome died = 1 (those who actually die).

$$TPF(R) = Pr(resection > R | subjectdied)$$

Similarly, the false positive fraction (FPF) for a specific cutoff R is the probability of a positive test (prediction that the patient will die) among the people with died = 0 (those who don't actually die)

$$FPF(R) = Pr(resection > R | subjectdidnotdie)$$

The True Positive Rate is referred to as the sensitivity of a diagnostic test, and the True Negative rate (1 - the False Positive rate) is referred to as the specificity of a diagnostic test.

Since the cutoff R is not fixed in advanced, we can plot the value of TPF (on the y axis) against FPF (on the x axis) for all possible values of R , and this is what the ROC curve is. Others refer to the Sensitivity on the Y axis, and 1-Specificity on the X axis, and this is the same idea.

Before we get too far into the weeds, let me show you some simple situations so you can understand what you might learn from the ROC curve. The web page <http://blog.yhat.com/posts/roc-curves.html> provides source materials.

12.9.1 Interpreting the Area under the ROC curve

The AUC or Area under the ROC curve is the amount of space underneath the ROC curve. Often referred to as the c statistic, the AUC represents the quality of your TPR and FPR overall in a single number. The C statistic ranges from 0 to 1, with $C = 0.5$ for a prediction that is no better than random guessing, and $C = 1$ for a perfect prediction model.

Next, I'll build a simulation to demonstrate several possible ROC curves in the sections that follow.

```
set.seed(432999)
sim.temp <- data_frame(x = rnorm(n = 200),
                        prob = exp(x)/(1 + exp(x)),
                        y = as.numeric(1 * runif(200) < prob))
```

```
sim.temp <- sim.temp %>%
  mutate(p_guess = 1,
        p_perfect = y,
        p_bad = exp(-2*x) / (1 + exp(-2*x)),
        p_ok = prob + (1-y)*runif(1, 0, 0.05),
        p_good = prob + y*runif(1, 0, 0.27))
```

12.9.1.1 What if we are guessing?

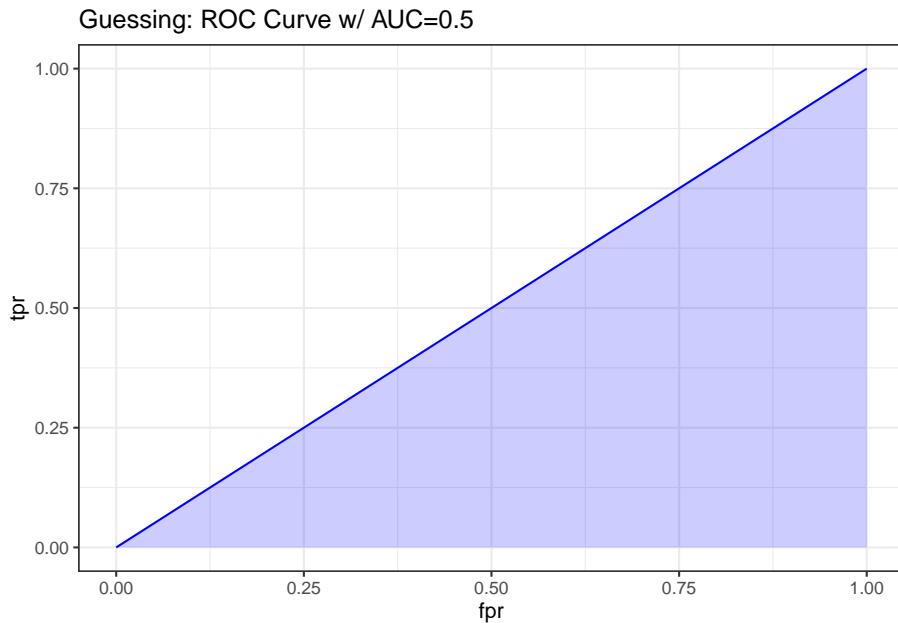
If we're guessing completely at random, then the model should correctly classify a subject (as died or not died) about 50% of the time, so the TPR and FPR will be equal. This yields a diagonal line in the ROC curve, and an area under the curve (C statistic) of 0.5.

There are several ways to do this on the web, but I'll show this one, which has some bizarre code, but that's a function of using a package called `ROCR` to do the work. It comes from this link

```
pred_guess <- prediction(sim.temp$p_guess, sim.temp$y)
perf_guess <- performance(pred_guess, measure = "tpr", x.measure = "fpr")
auc_guess <- performance(pred_guess, measure="auc")

auc_guess <- round(auc_guess@y.values[[1]],3)
roc_guess <- data.frame(fpr=unlist(perf_guess@x.values),
                         tpr=unlist(perf_guess@y.values),
                         model="GLM")

ggplot(roc_guess, aes(x=fpr, ymin=0, ymax=tpr)) +
  geom_ribbon(alpha=0.2, fill = "blue") +
  geom_line(aes(y=tpr), col = "blue") +
  labs(title = paste0("Guessing: ROC Curve w/ AUC=", auc_guess)) +
  theme_bw()
```



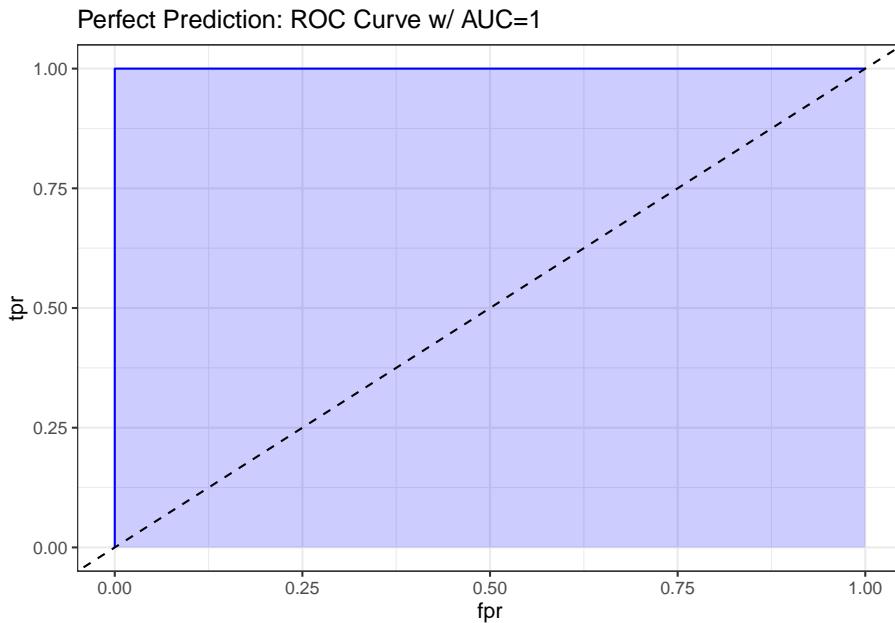
12.9.1.2 What if we classify things perfectly?

If we're classifying subjects perfectly, then we have a TPR of 1 and an FPR of 0. That yields an ROC curve that looks like the upper and left edges of a box. If our model correctly classifies a subject (as died or not died) 100% of the time, the area under the curve (c statistic) will be 1.0. We'll add in the diagonal line here (in a dashed black line) to show how this model compares to random guessing.

```
pred_perf <- prediction(sim.temp$p_perfect, sim.temp$y)
perf_perf <- performance(pred_perf, measure = "tpr", x.measure = "fpr")
auc_perf <- performance(pred_perf, measure="auc")

auc_perf <- round(auc_perf@y.values[[1]],3)
roc_perf <- data.frame(fpr=unlist(perf_perf@x.values),
                        tpr=unlist(perf_perf@y.values),
                        model="GLM")

ggplot(roc_perf, aes(x=fpr, ymin=0, ymax=tpr)) +
  geom_ribbon(alpha=0.2, fill = "blue") +
  geom_line(aes(y=tpr), col = "blue") +
  geom_abline(intercept = 0, slope = 1, lty = "dashed") +
  labs(title = paste0("Perfect Prediction: ROC Curve w/ AUC=", auc_perf)) +
  theme_bw()
```



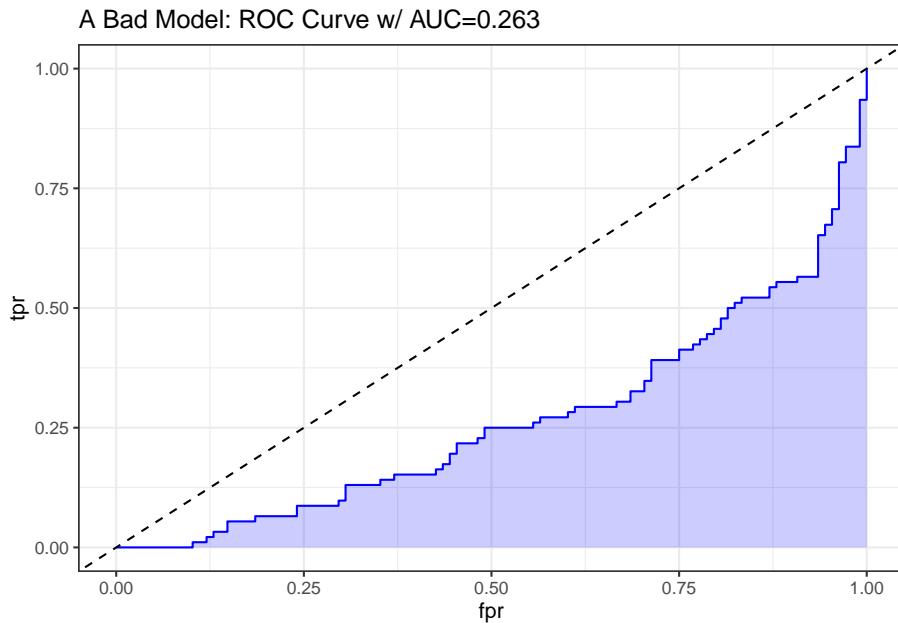
12.9.1.3 What does “worse than guessing” look like?

A bad classifier will appear below and to the right of the diagonal line we'd see if we were completely guessing. Such a model will have a c statistic below 0.5, and will be valueless.

```
pred_bad <- prediction(sim.temp$p_bad, sim.temp$y)
perf_bad <- performance(pred_bad, measure = "tpr", x.measure = "fpr")
auc_bad <- performance(pred_bad, measure="auc")

auc_bad <- round(auc_bad@y.values[[1]],3)
roc_bad <- data.frame(fpr=unlist(perf_bad@x.values),
                      tpr=unlist(perf_bad@y.values),
                      model="GLM")

ggplot(roc_bad, aes(x=fpr, ymin=0, ymax=tpr)) +
  geom_ribbon(alpha=0.2, fill = "blue") +
  geom_line(aes(y=tpr), col = "blue") +
  geom_abline(intercept = 0, slope = 1, lty = "dashed") +
  labs(title = paste0("A Bad Model: ROC Curve w/ AUC=", auc_bad)) +
  theme_bw()
```



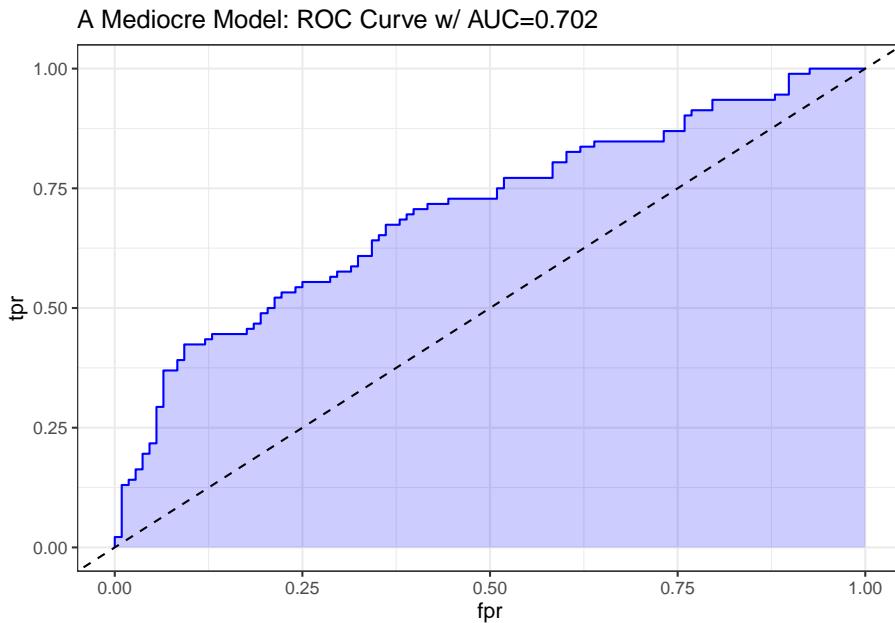
12.9.1.4 What does “better than guessing” look like?

An “OK” classifier will appear above and to the left of the diagonal line we’d see if we were completely guessing. Such a model will have a c statistic above 0.5, and might have some value. The plot below shows a very fairly poor model, but at least it’s better than guessing.

```
pred_ok <- prediction(sim.temp$p_ok, sim.temp$y)
perf_ok <- performance(pred_ok, measure = "tpr", x.measure = "fpr")
auc_ok <- performance(pred_ok, measure="auc")

auc_ok <- round(auc_ok@y.values[[1]],3)
roc_ok <- data.frame(fpr=unlist(perf_ok@x.values),
                      tpr=unlist(perf_ok@y.values),
                      model="GLM")

ggplot(roc_ok, aes(x=fpr, ymin=0, ymax=tpr)) +
  geom_ribbon(alpha=0.2, fill = "blue") +
  geom_line(aes(y=tpr), col = "blue") +
  geom_abline(intercept = 0, slope = 1, lty = "dashed") +
  labs(title = paste0("A Mediocre Model: ROC Curve w/ AUC=", auc_ok)) +
  theme_bw()
```



Sometimes people grasp for a rough guide as to the accuracy of a model's predictions based on the area under the ROC curve. A common thought is to assess the C statistic much like you would a class grade.

C statistic	Interpretation
0.90 to 1.00	model does an excellent job at discriminating "yes" from "no" (A)
0.80 to 0.90	model does a good job (B)
0.70 to 0.80	model does a fair job (C)
0.60 to 0.70	model does a poor job (D)
0.50 to 0.60	model fails (F)
below 0.50	model is worse than random guessing

12.9.1.5 What does “pretty good” look like?

A strong and good classifier will appear above and to the left of the diagonal line we'd see if we were completely guessing, often with a nice curve that is continually increasing and appears to be pulled up towards the top left. Such a model will have a c statistic well above 0.5, but not as large as 1. The plot below shows a stronger model, which appears substantially better than guessing.

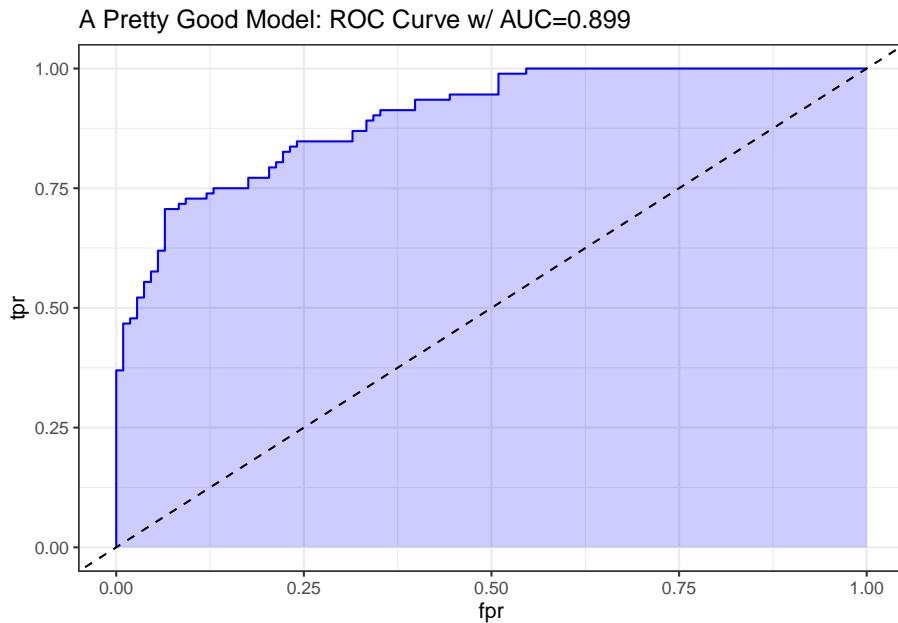
```
pred_good <- prediction(sim.temp$p_good, sim.temp$y)
perf_good <- performance(pred_good, measure = "tpr", x.measure = "fpr")
auc_good <- performance(pred_good, measure="auc")
```

```

auc_good <- round(auc_good@y.values[[1]],3)
roc_good <- data.frame(fpr=unlist(perf_good@x.values),
                        tpr=unlist(perf_good@y.values),
                        model="GLM")

ggplot(roc_good, aes(x=fpr, ymin=0, ymax=tpr)) +
  geom_ribbon(alpha=0.2, fill = "blue") +
  geom_line(aes(y=tpr), col = "blue") +
  geom_abline(intercept = 0, slope = 1, lty = "dashed") +
  labs(title = paste0("A Pretty Good Model: ROC Curve w/ AUC=", auc_good)) +
  theme_bw()

```



12.10 The ROC Plot for res_modA

Let me show you the ROC curve for our `res_modA` model.

```

## requires ROCR package
prob <- predict(res_modA, resect, type="response")
pred <- prediction(prob, resect$died)
perf <- performance(pred, measure = "tpr", x.measure = "fpr")
auc <- performance(pred, measure="auc")

auc <- round(auc@y.values[[1]],3)

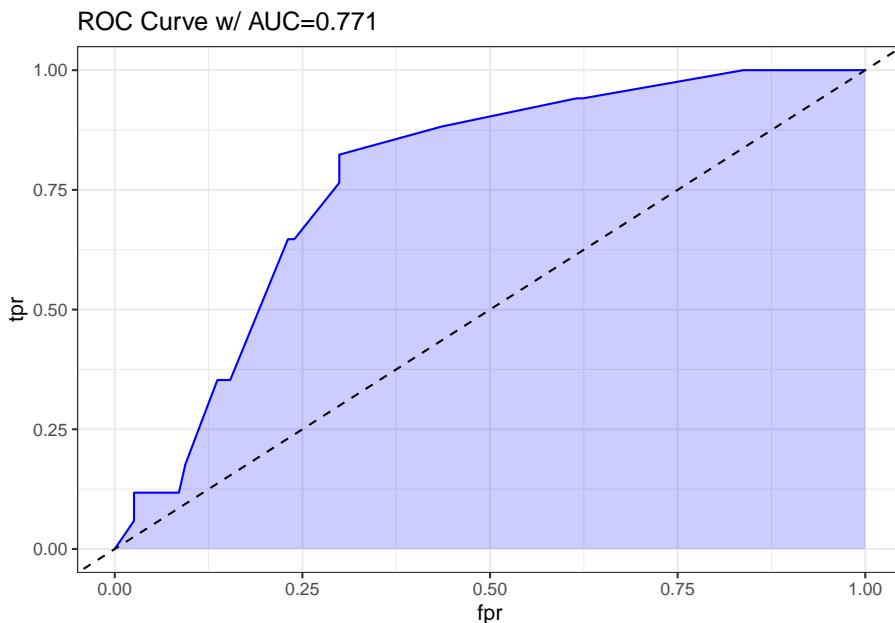
```

```

roc.data <- data.frame(fpr=unlist(perf@x.values),
                        tpr=unlist(perf@y.values),
                        model="GLM")

ggplot(roc.data, aes(x=fpr, ymin=0, ymax=tpr)) +
  geom_ribbon(alpha=0.2, fill = "blue") +
  geom_line(aes(y=tpr), col = "blue") +
  geom_abline(intercept = 0, slope = 1, lty = "dashed") +
  labs(title = paste0("ROC Curve w/ AUC=", auc)) +
  theme_bw()

```



Based on the C statistic (AUC = 0.771) this would rank somewhere near the high end of a “fair” predictive model by this standard, not quite to the level of a “good” model.

12.10.1 Another way to plot the ROC Curve

If we’ve loaded the `pROC` package, we can also use the following (admittedly simpler) approach to plot the ROC curve, without `ggplot2`, and to obtain the C statistic, and a 95% confidence interval around that C statistic.

```

## requires pROC package
roc.modA <-
  roc(resect$died ~ predict(res_modA, type="response"),

```

```

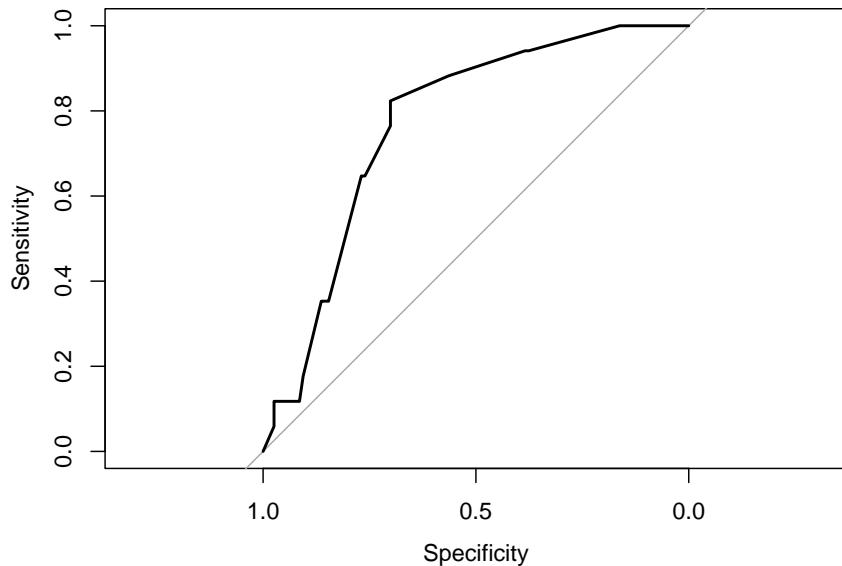
  ci = TRUE)

roc.modA

Call:
roc.formula(formula = resect$died ~ predict(res_modA, type = "response"),
             ci = TRUE)

Data: predict(res_modA, type = "response") in 117 controls (resect$died 0) < 17 cases
Area under the curve: 0.7707
95% CI: 0.67-0.8715 (DeLong)
plot(roc.modA)

```



12.11 Assessing Residual Plots from Model A

Residuals are certainly less informative for logistic regression than they are for linear regression: not only do yes/no outcomes inherently contain less information than continuous ones, but the fact that the adjusted response depends on the fit hampers our ability to use residuals as external checks on the model.

This is mitigated to some extent, however, by the fact that we are also making fewer distributional assumptions in logistic regression,

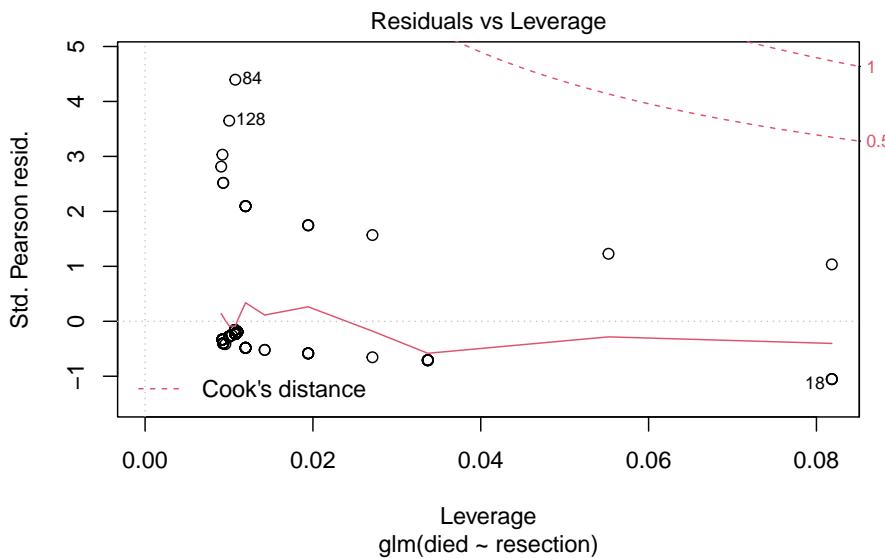
so there is no need to inspect residuals for, say, skewness or heteroskedasticity.

- Patrick Breheny, University of Kentucky, Slides on GLM Residuals and Diagnostics

The usual residual plots are available in R for a logistic regression model, but most of them are irrelevant in the logistic regression setting. The residuals shouldn't follow a standard Normal distribution, and they will not show constant variance over the range of the predictor variables, so plots looking into those issues aren't helpful.

The only plot from the standard set that we'll look at in many settings is plot 5, which helps us assess influence (via Cook's distance contours), and a measure related to leverage (how unusual an observation is in terms of the predictors) and standardized Pearson residuals.

```
plot(res_modA, which = 5)
```



In this case, I don't see any highly influential points, as no points fall outside of the Cook's distance (0.5 or 1) contours.

12.12 Model B: A “Kitchen Sink” Logistic Regression Model

```
res_modB <- glm(died ~ resection + age + prior + intubated,
                  data = resect, family = binomial)

res_modB

Call: glm(formula = died ~ resection + age + prior + intubated, family = binomial,
         data = resect)

Coefficients:
(Intercept)    resection          age        prior      intubated
-5.152886       0.612211      0.001173      0.814691     2.810797

Degrees of Freedom: 133 Total (i.e. Null); 129 Residual
Null Deviance: 101.9
Residual Deviance: 67.36    AIC: 77.36
```

12.12.1 Comparing Model A to Model B

```
anova(res_modA, res_modB)
```

Analysis of Deviance Table

	Model 1: died ~ resection	Model 2: died ~ resection + age + prior + intubated		
	Resid. Df	Resid. Dev	Df	Deviance
1	132	89.493		
2	129	67.359	3	22.134

The addition of `age`, `prior` and `intubated` reduces the lack of fit by 22.134 points, at a cost of 3 degrees of freedom.

```
glance(res_modA)
```

	# A tibble: 1 x 8	null.deviance	df.null	logLik	AIC	BIC	deviance	df.residual	nobs
		<dbl>	<int>	<dbl>	<dbl>	<dbl>	<dbl>	<int>	<int>
1		102.	133	-44.7	93.5	99.3	89.5	132	134

```
glance(res_modB)
```

	# A tibble: 1 x 8	null.deviance	df.null	logLik	AIC	BIC	deviance	df.residual	nobs
		<dbl>	<int>	<dbl>	<dbl>	<dbl>	<dbl>	<int>	<int>
1		102.	133	-44.7	93.5	99.3	89.5	132	134

	<dbl>	<int>	<dbl>	<dbl>	<dbl>	<dbl>	<int>	<int>
1	102.	133	-33.7	77.4	91.8	67.4	129	134

By either AIC or BIC, the larger model (`res_modB`) looks more effective.

12.12.2 Interpreting Model B

```
summary(res_modB)

Call:
glm(formula = died ~ resection + age + prior + intubated, family = binomial,
     data = resect)

Deviance Residuals:
    Min      1Q  Median      3Q      Max
-1.7831 -0.3741 -0.2386 -0.2014  2.5228

Coefficients:
            Estimate Std. Error z value Pr(>|z|)
(Intercept) -5.152886   1.469453 -3.507 0.000454 ***
resection     0.612211   0.282807  2.165 0.030406 *
age          0.001173   0.020646  0.057 0.954700
prior         0.814691   0.704785  1.156 0.247705
intubated     2.810797   0.658395  4.269 1.96e-05 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

Null deviance: 101.943 on 133 degrees of freedom
Residual deviance: 67.359 on 129 degrees of freedom
AIC: 77.359
```

Number of Fisher Scoring iterations: 6

It appears that the `intubated` predictor adds significant value to the model, by the Wald test.

Let's focus on the impact of these variables through odds ratios.

```
tidy(res_modB, exponentiate = TRUE, conf.int = TRUE) %>%
  select(term, estimate, conf.low, conf.high)

# A tibble: 5 x 4
  term       estimate conf.low conf.high
  <chr>        <dbl>    <dbl>    <dbl>
1 (Intercept)  0.00578  0.000241   0.0837
```

2	resection	1.84	1.08	3.35
3	age	1.00	0.962	1.04
4	prior	2.26	0.549	9.17
5	intubated	16.6	4.75	64.6

At a 5% significance level, we might conclude that:

- larger sized `resections` are associated with a meaningful rise (est OR: 1.84, 95% CI 1.08, 3.35) in the odds of death, holding all other predictors constant,
- the need for `intubation` at the end of surgery is associated with a substantial rise (est OR: 16.6, 95% CI 4.7, 64.7) in the odds of death, holding all other predictors constant, but that
- older `age` as well as having a `prior` tracheal surgery appears to be associated with an increase in death risk, but not to an extent that we can declare statistically significant.

12.13 Plotting Model B

Let's think about plotting the fitted values from our model, in terms of probabilities.

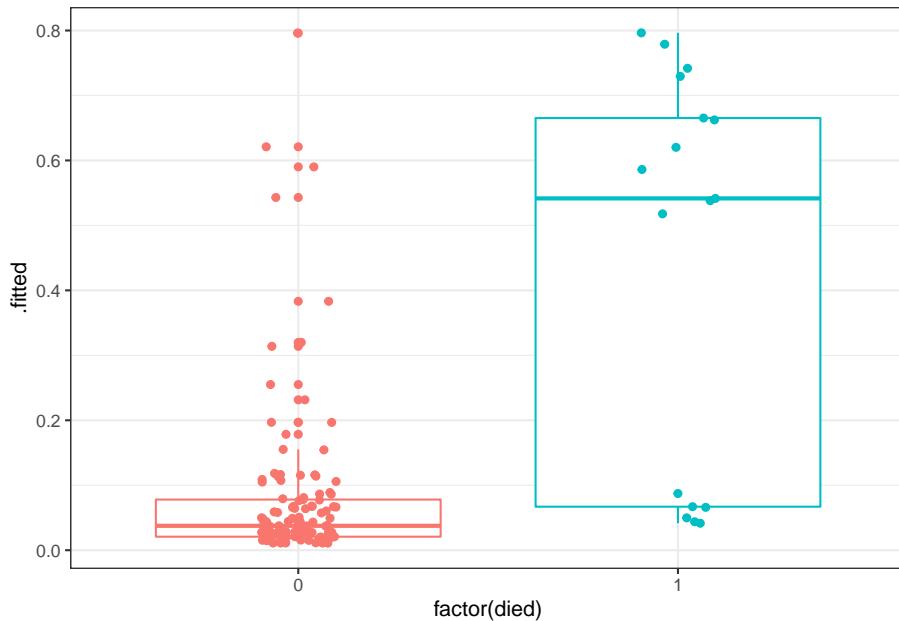
12.13.1 Using `augment` to capture the fitted probabilities

```
res_B_aug <- augment(res_modB, resect,
                      type.predict = "response")
head(res_B_aug)

# A tibble: 6 x 12
  id   age prior resection intubated died .fitted .resid .std.resid .hat
  <dbl> <dbl> <dbl>     <dbl>     <dbl> <dbl> <dbl>     <dbl>     <dbl>
1    1    34    1        2.5      0     0  0.0591 -0.349  -0.354  0.0267
2    2    57    0        5       0     0  0.117  -0.498  -0.508  0.0380
3    3    60    1        4       1     1  0.729  0.794   0.844  0.114
4    4    62    1        4.2      0     0  0.155  -0.581  -0.602  0.0704
5    5    28    0        6       1     1  0.796  0.675   0.724  0.131
6    6    52    0        3       0     0  0.0371 -0.275  -0.277  0.0105
# ... with 2 more variables: .sigma <dbl>, .cooks <dbl>
```

12.13.2 Plotting Model B Fits by Observed Mortality

```
ggplot(res_B_aug, aes(x = factor(died), y = .fitted, col = factor(died))) +
  geom_boxplot() +
  geom_jitter(width = 0.1) +
  guides(col = "none")
```



Certainly it appears as though most of our predicted probabilities (of death) for the subjects who actually survived are quite small, but not all of them. We also have at least 6 big “misses” among the 17 subjects who actually died.

12.13.3 Confusion Matrix for Model B

```
res_B_aug %$%
  confusionMatrix(
    data = factor(.fitted >= 0.5),
    reference = factor(died == 1),
    positive = "TRUE"
  )
```

Confusion Matrix and Statistics

		Reference	
		FALSE	TRUE
Prediction	FALSE	113	6
	TRUE	4	11

```

    Accuracy : 0.9254
    95% CI  : (0.867, 0.9636)
No Information Rate : 0.8731
P-Value [Acc > NIR] : 0.03897

    Kappa : 0.6453

McNemar's Test P-Value : 0.75183

    Sensitivity : 0.64706
    Specificity  : 0.96581
    Pos Pred Value : 0.73333
    Neg Pred Value : 0.94958
    Prevalence   : 0.12687
    Detection Rate : 0.08209
Detection Prevalence : 0.11194
    Balanced Accuracy : 0.80644

'Positive' Class : TRUE

```

12.13.4 The ROC curve for Model B

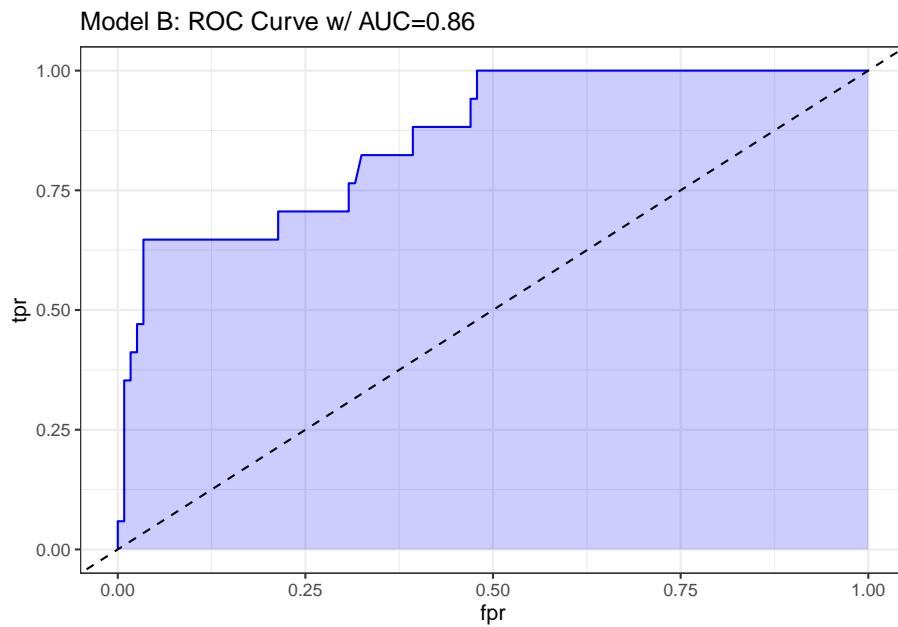
```

## requires ROCR package
prob <- predict(res_modB, resect, type="response")
pred <- prediction(prob, resect$died)
perf <- performance(pred, measure = "tpr", x.measure = "fpr")
auc <- performance(pred, measure="auc")

auc <- round(auc@y.values[[1]],3)
roc.data <- data.frame(fpr=unlist(perf@x.values),
                        tpr=unlist(perf@y.values),
                        model="GLM")

ggplot(roc.data, aes(x=fpr, ymin=0, ymax=tpr)) +
  geom_ribbon(alpha=0.2, fill = "blue") +
  geom_line(aes(y=tpr), col = "blue") +
  geom_abline(intercept = 0, slope = 1, lty = "dashed") +
  labs(title = paste0("Model B: ROC Curve w/ AUC=", auc)) +
  theme_bw()

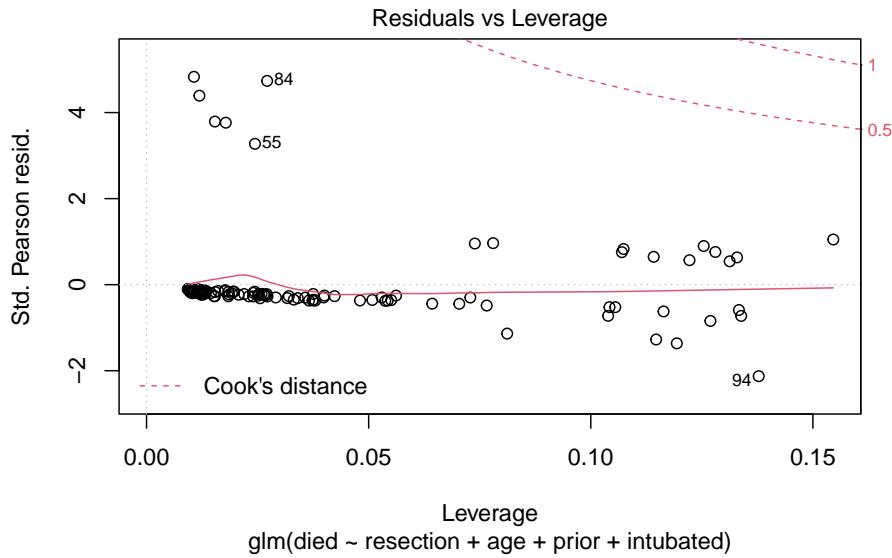
```



The area under the curve (C-statistic) is 0.86, which certainly looks like a more discriminating fit than model A with resection alone.

12.13.5 Residuals, Leverage and Influence

```
plot(res_modB, which = 5)
```



Again, we see no signs of deeply influential points in this model.

12.14 Logistic Regression using `lrm`

To obtain the Nagelkerke R^2 and the C statistic, as well as some other summaries, I'll now demonstrate the use of `lrm` from the `rms` package to fit a logistic regression model.

We'll return to the original model, predicting death using resection size alone.

```
dd <- datadist(resect)
options(datadist="dd")

res_modC <- lrm(died ~ resection, data=resect, x=TRUE, y=TRUE)
res_modC
```

Logistic Regression Model

```
lrm(formula = died ~ resection, data = resect, x = TRUE, y = TRUE)
```

	Model	Likelihood Ratio Test	Discrimination Indexes	Rank Discrim. Indexes
Obs	134	LR chi2 12.45	R2 0.167	C 0.771
0	117	d.f. 1	g 1.037	Dxy 0.541
1	17	Pr(> chi2) 0.0004	gr 2.820	gamma 0.582

```

max |deriv| 2e-06
gp      0.110   tau-a  0.121
Brier   0.103

Coef    S.E.   Wald Z Pr(>|Z|)
Intercept -4.4337 0.8799 -5.04 <0.0001
resection  0.7417 0.2230  3.33  0.0009

```

This output specifies the following:

- **Obs** = The number of observations used to fit the model, with 0 = the number of zeros and 1 = the number of ones in our outcome, `died`. Also specified is the maximum absolute value of the derivative at the point where the maximum likelihood function was estimated. I wouldn't worry about that practically, as all you will care about is whether the iterative function-fitting process converged, and R will warn you in other ways if it doesn't.
- A likelihood ratio test (drop in deviance test) subtracting the residual deviance from the null deviance obtain the Likelihood Ratio χ^2 statistic, subtracting residual df from null df to obtain degrees of freedom, and comparing the resulting test statistic to a χ^2 distribution with the appropriate degrees of freedom to determine a p value.
- A series of discrimination indexes, including the Nagelkerke R^2 , symbolized `R2`, and several others we'll discuss shortly.
- A series of rank discrimination indexes, including the C statistic (area under the ROC curve) and Somers' D (`Dxy`), and several others.
- A table of coefficients, standard errors, Wald Z statistics and p values based on those Wald statistics.

The C statistic is estimated to be 0.771, with an associated (Nagelkerke) $R^2 = 0.167$, both indicating at best mediocre performance for this model, as it turns out.

12.14.1 Interpreting Nagelkerke R^2

There are many ways to calculate R^2 for logistic regression.

- At the URL linked here there is a nice summary of the key issue, which is that there are at least three different ways to think about R^2 in linear regression that are equivalent in that context, but when you move to a categorical outcome, which interpretation you use leads you down a different path for extension to the new type of outcome. In linear regression...
 - You might think of R^2 as a measure of the proportion of variability explained.
 - You might think of R^2 as measuring the improvement from a null model to a fitted model.
 - You might think of R^2 as the square of the correlation coefficient.

- Paul Allison, for instance, describes several at this link in a post entitled “What’s the Best R-Squared for Logistic Regression?”
- Jonathan Bartlett looks at McFadden’s pseudo R^2 in some detail (including some R code) at this link, in a post entitled “R squared in logistic regression”

The Nagelkerke approach that is presented as `R2` in the `lrm` output is as good as most of the available approaches, and has the positive feature that it does reach 1 if the fitted model shows as much improvement as possible over the null model (which predicts the mean response for all subjects, and has $R^2 = 0$). The greater the improvement, the higher the Nagelkerke R^2 .

For model A, our Nagelkerke $R^2 = 0.167$, which is pretty poor. It doesn’t technically mean that 16.7% of any sort of variation has been explained, though.

12.14.2 Interpreting the C statistic and Plotting the ROC Curve

The C statistic is a measure of the area under the receiver operating characteristic curve. This link has some nice material that provides some insight into the C statistic and ROC curve.

- Recall that C ranges from 0 to 1. 0 = BAD, 1 = GOOD.
 - values of C less than 0.5 indicate that your prediction model is not even as good as simple random guessing of “yes” or “no” for your response.
 - C = 0.5 for random guessing
 - C = 1 indicates a perfect classification scheme - one that correctly guesses “yes” for all “yes” patients, and for none of the “no” patients.
- The closer C is to 1, the happier we’ll be, most of the time.
 - Often we’ll call models with $0.5 < C < 0.8$ poor or weak in terms of predictive ability by this measure
 - $0.8 \leq C < 0.9$ are moderately strong in terms of predictive power (indicate good discrimination)
 - $C \geq 0.9$ usually indicates a very strong model in this regard (indicate excellent discrimination)

We’ve seen the ROC curve for this model before, when we looked at model `res_modA` fitted using `glm` in the previous chapter. But, just for completeness, I’ll include it.

Note. I change the initial `predict` call from `type = "response"` for a `glm` fit to `type = "fitted"` in a `lrm` fit. Otherwise, this is the same approach.

```
## requires ROCR package
prob <- predict(res_modC, resect, type="fitted")
pred <- prediction(prob, resect$died)
perf <- performance(pred, measure = "tpr", x.measure = "fpr")
```

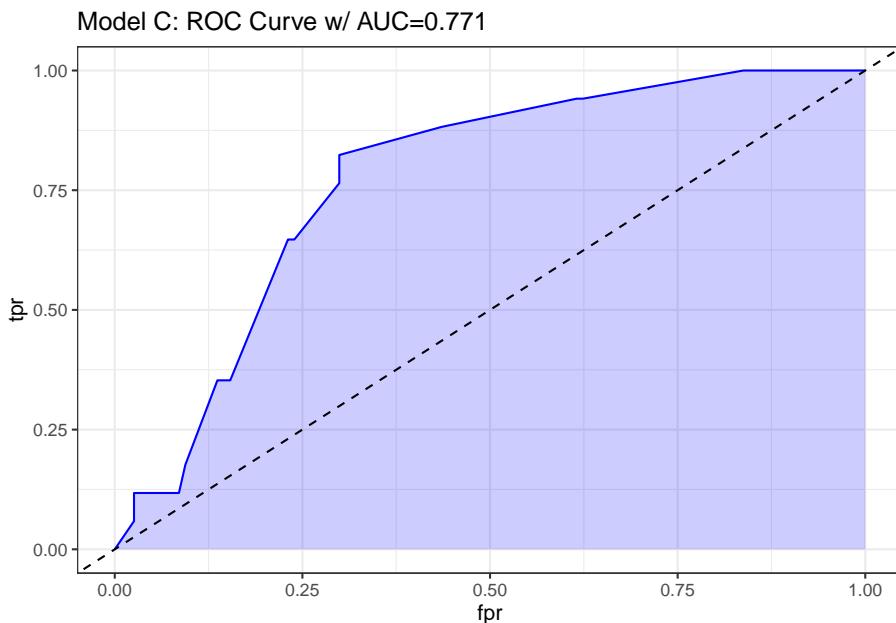
```

auc <- performance(pred, measure="auc")

auc <- round(auc@y.values[[1]],3)
roc.data <- data.frame(fpr=unlist(perf@x.values),
                        tpr=unlist(perf@y.values),
                        model="GLM")

ggplot(roc.data, aes(x=fpr, ymin=0, ymax=tpr)) +
  geom_ribbon(alpha=0.2, fill = "blue") +
  geom_line(aes(y=tpr), col = "blue") +
  geom_abline(intercept = 0, slope = 1, lty = "dashed") +
  labs(title = paste0("Model C: ROC Curve w/ AUC=", auc)) +
  theme_bw()

```



12.14.3 The C statistic and Somers' D

- The C statistic is directly related to **Somers' D statistic**, abbreviated D_{xy} , by the equation $C = 0.5 + (D/2)$.
 - Somers' D and the ROC area only measure how well predicted values from the model can rank-order the responses. For example, predicted probabilities of 0.01 and 0.99 for a pair of subjects are no better than probabilities of 0.2 and 0.8 using rank measures, if the first subject had a lower response value than the second.

- Thus, the C statistic (or D_{xy}) may not be very sensitive ways to choose between models, even though they provide reasonable summaries of the models individually.
- This is especially true when the models are strong. The Nagelkerke R^2 may be more sensitive.
- But as it turns out, we sometimes have to look at the ROC shapes, as the summary statistic alone isn't enough.

In our case, Somers D (D_{xy}) = .541, so the C statistic is 0.771.

12.14.4 Validating the Logistic Regression Model Summary Statistics

Like other regression-fitting tools in **rms**, the **lrm** function has a special **validate** tool to help perform resampling validation of a model, with or without backwards step-wise variable selection. Here, we'll validate our model's summary statistics using 100 bootstrap replications.

```
set.seed(432001)
validate(res_modC, B = 100)
```

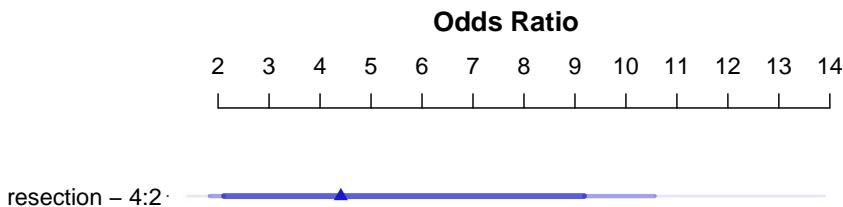
	index.orig	training	test	optimism	index.corrected	n
Dxy	0.5415	0.5422	0.5415	0.0007	0.5408	100
R2	0.1666	0.1748	0.1666	0.0083	0.1583	100
Intercept	0.0000	0.0000	0.1631	-0.1631	0.1631	100
Slope	1.0000	1.0000	1.0463	-0.0463	1.0463	100
Emax	0.0000	0.0000	0.0428	0.0428	0.0428	100
D	0.0854	0.0909	0.0854	0.0055	0.0800	100
U	-0.0149	-0.0149	0.0017	-0.0167	0.0017	100
Q	0.1004	0.1058	0.0837	0.0221	0.0783	100
B	0.1025	0.0986	0.1051	-0.0065	0.1090	100
g	1.0369	1.0677	1.0369	0.0308	1.0061	100
gp	0.1101	0.1080	0.1101	-0.0021	0.1122	100

Recall that our area under the curve (C statistic) = $0.5 + (D_{xy}/2)$, so that we can also use the first row of statistics to validate the C statistic. Accounting for optimism in this manner, our validation-corrected estimates are $D_{xy} = 0.5408$, so $C = 0.7704$, and, from the second row of statistics, we can read off the validated Nagelkerke R^2 , which is 0.1583.

12.14.5 Plotting the Summary of the **lrm** approach

The **summary** function applied to an **lrm** fit shows the effect size comparing the 25th to the 75th percentile of resection.

```
plot(summary(res_modC))
```



```
summary(res_modC)
```

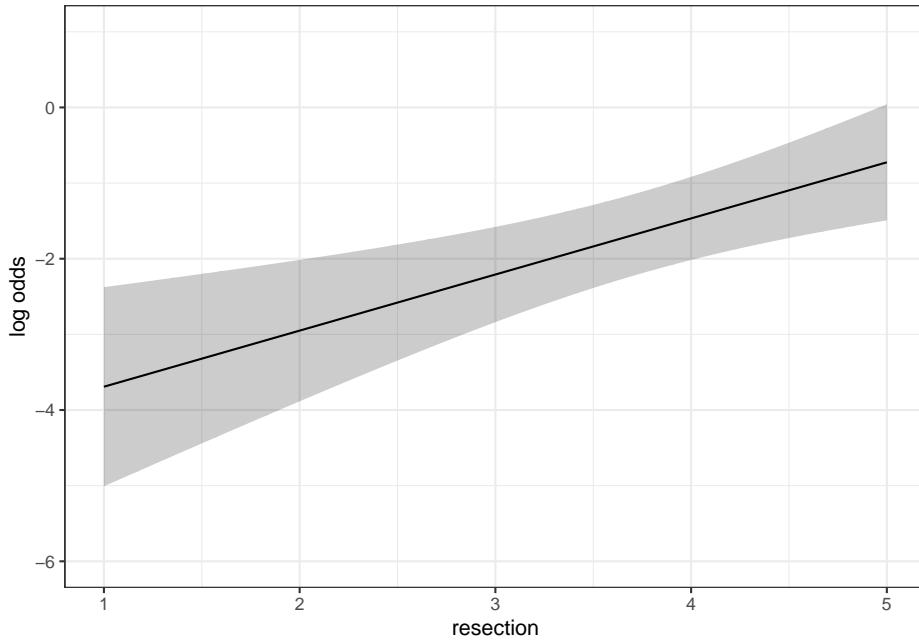
	Effects			Response : died			
Factor	Low	High	Diff.	Effect	S.E.	Lower	0.95 Upper
resection	2	4	2	1.4834	0.44591	0.6094	2.3574
Odds Ratio	2	4	2	4.4078	NA	1.8393	10.5630

So, a move from a resection of 2 cm to a resection of 4 cm is associated with an estimated effect on the log odds of death of 1.48 (with standard error 0.45), or with an estimated effect on the odds ratio for death of 4.41, with 95% CI (1.84, 10.56).

12.14.6 Plot In-Sample Predictions for Model C

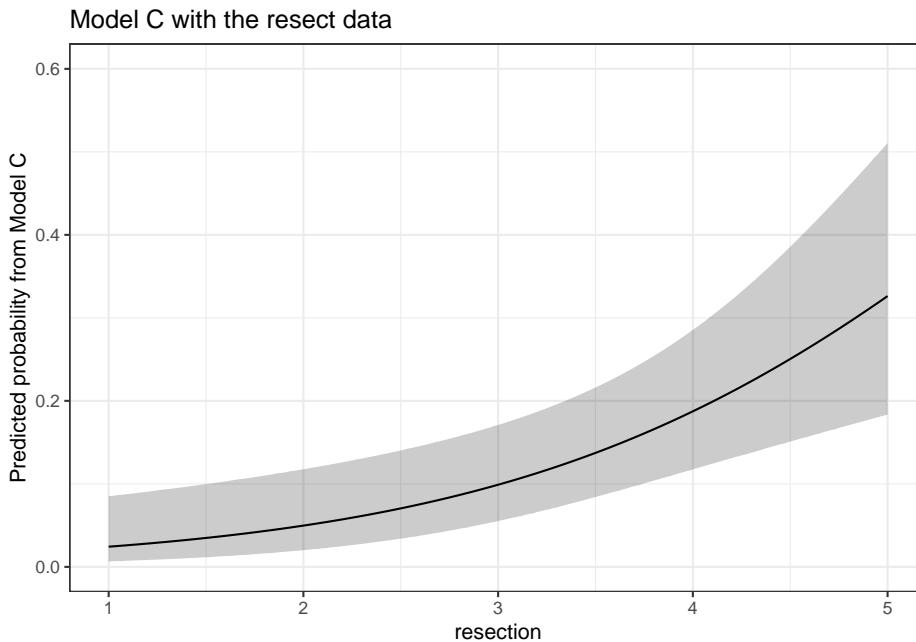
Here we plot the effect of `resection` (and 95% confidence intervals) across the range of observed values of `resection` on the log odds of death. Note the linear effect of `resection` size on the log odds scale.

```
ggplot(Predict(res_modC))
```



By applying the `plogis` function within the `Predict` command, we can plot the effect of `resection` on the estimated probability of death. Note the non-linear effect on this probability in this logistic regression model.

```
ggplot(Predict(res_modC, fun = plogis)) +
  labs(y = "Predicted probability from Model C",
       title = "Model C with the resect data")
```



The `Predict` function itself provides the raw material being captured in this plot.

```
head(Predict(res_modC, fun = plogis))
```

	resection	yhat	lower	upper	.predictor.
resection.1	1.000000	0.02431476	0.006636502	0.08505223	resection
resection.2	1.020101	0.02467096	0.006789313	0.08559056	resection
resection.3	1.040201	0.02503224	0.006945549	0.08613277	resection
resection.4	1.060302	0.02539867	0.007105283	0.08667889	resection
resection.5	1.080402	0.02577033	0.007268589	0.08722896	resection
resection.6	1.100503	0.02614728	0.007435542	0.08778304	resection

Response variable (y):

Limits are 0.95 confidence limits

12.14.7 ANOVA from the `lrm` approach

```
anova(res_modC)
```

Wald Statistics			Response: died	
Factor	Chi-Square	d.f.	P	
resection	11.07	1	9e-04	

TOTAL	11.07	1	9e-04
-------	-------	---	-------

The ANOVA approach applied to a `lrm` fit provides a Wald test for the model as a whole. Here, the use of `resection` is a significant improvement over a null (intercept-only) model. The p value is 9×10^{-4} .

12.14.8 Are any points particularly influential?

I'll use a cutoff for `dfbeta` here of 0.3, instead of the default 0.2, because I want to focus on truly influential points. Note that we have to use the data frame version of `resect` as `show.influence` isn't tibble-friendly.

```
inf.C <- which.influence(res_modC, cutoff=0.3)
inf.C

$Intercept
[1] 84 128

$resection
[1] 84

show.influence(object = inf.C, dframe = data.frame(resect))
```

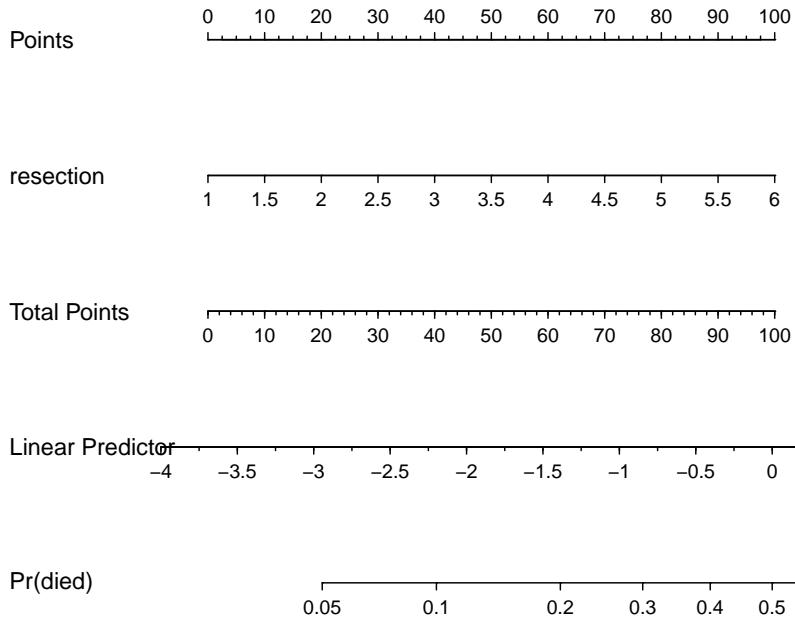
Count resection		
84	2	*2.0
128	1	2.5

It appears that observation 84 may have a meaningful effect on both the intercept and the coefficient for `resection`.

12.14.9 A Nomogram for Model C

We use the `plogis` function within a nomogram call to get R to produce fitted probabilities (of our outcome, `died`) in this case.

```
plot(nomogram(res_modC, fun=plogis,
              fun.at=c(0.05, seq(0.1, 0.9, by = 0.1), 0.95),
              funlabel="Pr(died)"))
```



Since there's no non-linearity in the right hand side of our simple logistic regression model, the nomogram is straightforward. We calculate the points based on the resection by traveling up, and then travel down in a straight vertical line from total points through the linear (log odds) predictor straight to a fitted probability. Note that fitted probabilities above 0.5 are not possible within the range of observed `resection` values in this case.

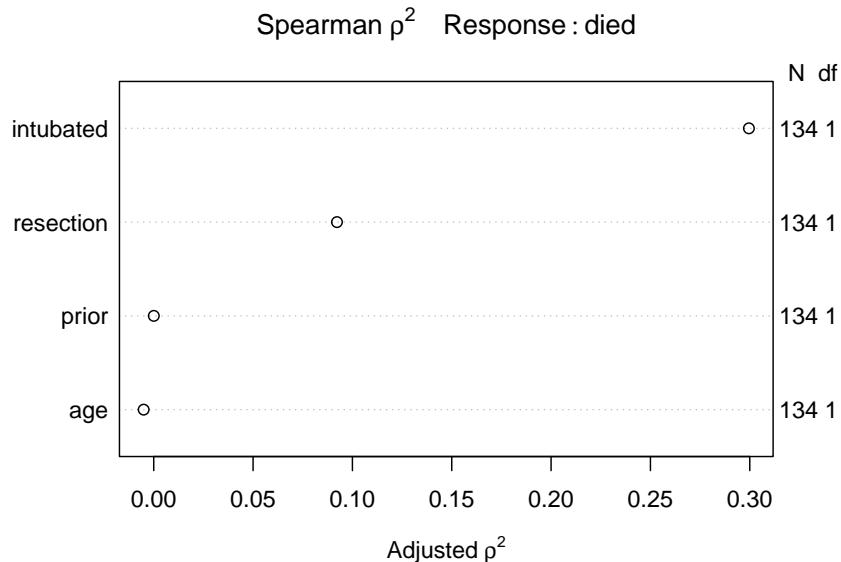
12.15 Model D: An Augmented Kitchen Sink Model

Can we predict survival from the patient's age, whether the patient had prior tracheal surgery or not, the extent of the resection, and whether intubation was required at the end of surgery?

12.15.1 Spearman ρ^2 Plot

Let's start by considering the limited use of non-linear terms for predictors that look important in a Spearman ρ^2 plot.

```
plot(spearman2(died ~ age + prior + resection + intubated, data=resect))
```



The most important variable appears to be whether intubation was required, so I'll include `intubated`'s interaction with the linear effect of the next most (apparently) important variable, `resection`, and also a cubic spline for `resection`, with three knots. Since `prior` and `age` look less important, I'll simply add them as linear terms.

12.15.2 Fitting Model D using `lrm`

Note the use of `%ia%` here. This insures that only the linear part of the `resection` term will be used in the interaction with `intubated`.

```
dd <- datadist(resect)
options(datadist="dd")

res_modD <- lrm(died ~ age + prior + rcs(resection, 3) +
                  intubated + intubated %ia% resection,
                  data=resect, x=TRUE, y=TRUE)
```

12.15.3 Assessing Model D using `lrm`'s tools

```
res_modD
```

Logistic Regression Model

```
lrm(formula = died ~ age + prior + rcs(resection, 3) + intubated +
    intubated %ia% resection, data = resect, x = TRUE, y = TRUE)
```

	Model	Likelihood Ratio Test	Discrimination Indexes		Rank	Discrim.	
Obs	134	LR chi2	38.08	R2	0.464	C	0.880
0	117	d.f.	6	g	2.382	Dxy	0.759
1	17	Pr(> chi2)	<0.0001	gr	10.825	gamma	0.770
max deriv	9e-08			gp	0.172	tau-a	0.169
				Brier	0.067		
		Coef	S.E.	Wald Z	Pr(> Z)		
Intercept		-11.3636	4.9099	-2.31	0.0206		
age		0.0000	0.0210	0.00	0.9988		
prior		0.6269	0.7367	0.85	0.3947		
resection		3.3799	1.9700	1.72	0.0862		
resection'		-4.2104	2.7035	-1.56	0.1194		
intubated		0.4576	2.7848	0.16	0.8695		
intubated * resection		0.6188	0.7306	0.85	0.3970		

- The model likelihood ratio test suggests that at least some of these predictors are helpful.
- The Nagelkerke R^2 of 0.46, and the C statistic of 0.88 indicate a meaningful improvement in discrimination over our model with `resection` alone.
- The Wald Z tests see some potential need to prune the model, as none of the elements reaches statistical significance without the others. The product term between `intubated` and `resection`, in particular, doesn't appear to have helped much, once we already had the main effects.

12.15.4 ANOVA and Wald Tests for Model D

```
anova(res_modD)
```

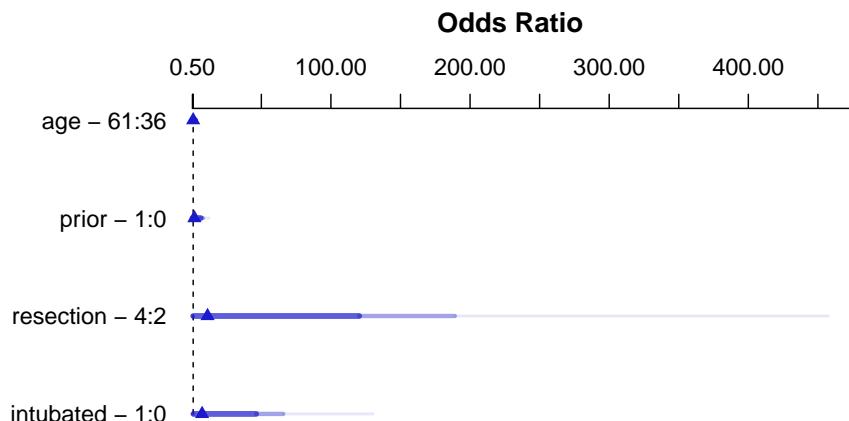
	Wald Statistics	Response: died		
Factor		Chi-Square	d.f.	P
age		0.00	1	0.9988
prior		0.72	1	0.3947

resection (Factor+Higher Order Factors)	4.95	3	0.1753
All Interactions	0.72	1	0.3970
Nonlinear	2.43	1	0.1194
intubated (Factor+Higher Order Factors)	16.45	2	0.0003
All Interactions	0.72	1	0.3970
intubated * resection (Factor+Higher Order Factors)	0.72	1	0.3970
TOTAL NONLINEAR + INTERACTION	2.56	2	0.2783
TOTAL	23.90	6	0.0005

Neither the interaction term nor the non-linearity from the cubic spline appears to be statistically significant, based on the Wald tests via ANOVA. However it is clear that `intubated` has a significant impact as a main effect.

12.15.5 Effect Sizes in Model D

```
plot(summary(res_modD))
```



Adjusted to:resection=2.5 intubated=0

```
summary(res_modD)
```

	Effects			Response : died				
Factor	Low	High	Diff.	Effect	S.E.	Lower	0.95	Upper
age	36	61	25	-0.00080933	0.52409	-1.02800		1.0264
Odds Ratio	36	61	25	0.99919000	NA	0.35772		2.7910
prior	0	1	1	0.62693000	0.73665	-0.81688		2.0707

Odds Ratio	0	1	1	1.87190000	NA	0.44181	7.9307
resection	2	4	2	2.42930000	1.43510	-0.38342	5.2419
Odds Ratio	2	4	2	11.35000000	NA	0.68153	189.0400
intubated	0	1	1	2.00470000	1.11220	-0.17513	4.1845
Odds Ratio	0	1	1	7.42380000	NA	0.83934	65.6610

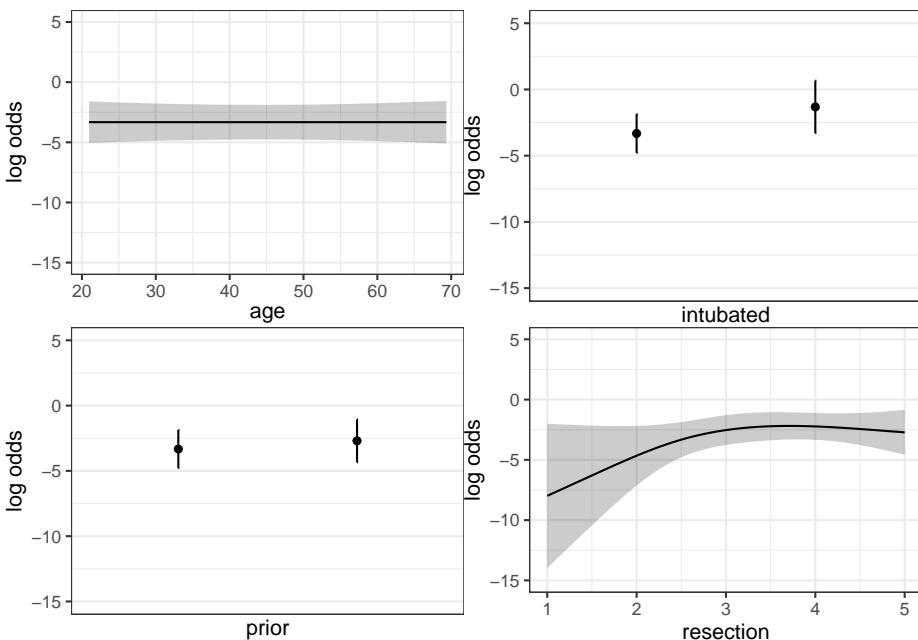
Adjusted to: resection=2.5 intubated=0

The effect sizes are perhaps best described in terms of odds ratios. The odds ratio for death isn't significantly different from 1 for any variable, but the impact of `resection` and `intubated`, though not strong enough to be significant, look more substantial (if poorly estimated) than the effects of `age` and `prior`.

12.15.6 Plot In-Sample Predictions for Model D

Here are plots of the effects across the range of each predictor (holding the others constant) on the log odds scale. Note the non-linear effect of `resection` implied by the use of a spline there.

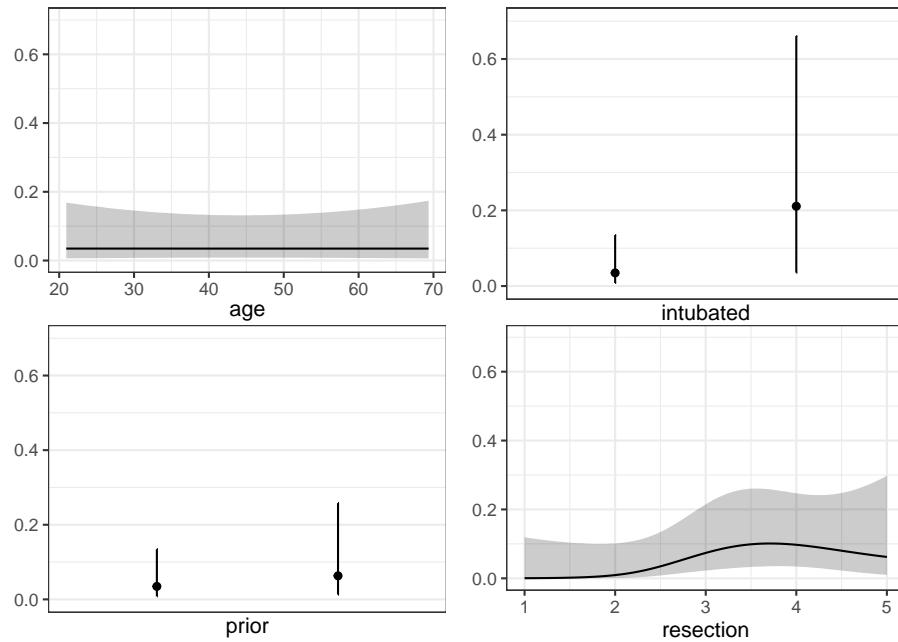
```
ggplot(Predict(res_modD))
```



We can also capture and plot these results on the probability scale, as follows¹.

¹Although I've yet to figure out how to get the y axis relabeled properly without simply dumping the `Predict` results into a new tibble and starting over with creating the plots.

```
ggsignif(Predict(res_modD, fun = plogis))
```



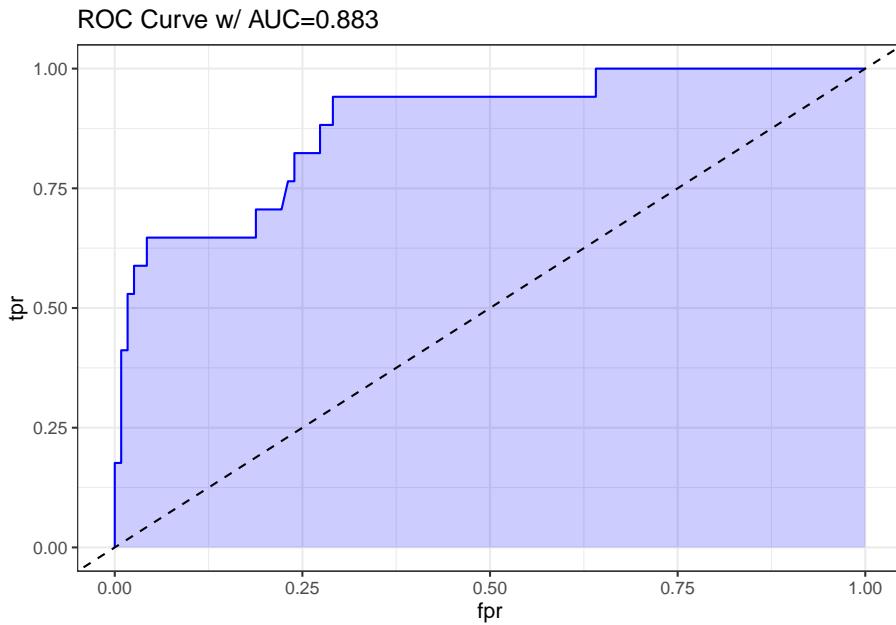
12.15.7 Plotting the ROC curve for Model D

Again, remember to use `type = "fitted"` with a `lrm` fit.

```
## requires ROCR package
prob <- predict(res_modD, resect, type="fitted")
pred <- prediction(prob, resect$died)
perf <- performance(pred, measure = "tpr", x.measure = "fpr")
auc <- performance(pred, measure="auc")

auc <- round(auc@y.values[[1]],3)
roc.data <- data.frame(fpr=unlist(perf@x.values),
                        tpr=unlist(perf@y.values),
                        model="GLM")

ggplot(roc.data, aes(x=fpr, ymin=0, ymax=tpr)) +
  geom_ribbon(alpha=0.2, fill = "blue") +
  geom_line(aes(y=tpr), col = "blue") +
  geom_abline(intercept = 0, slope = 1, lty = "dashed") +
  labs(title = paste0("ROC Curve w/ AUC=", auc)) +
  theme_bw()
```



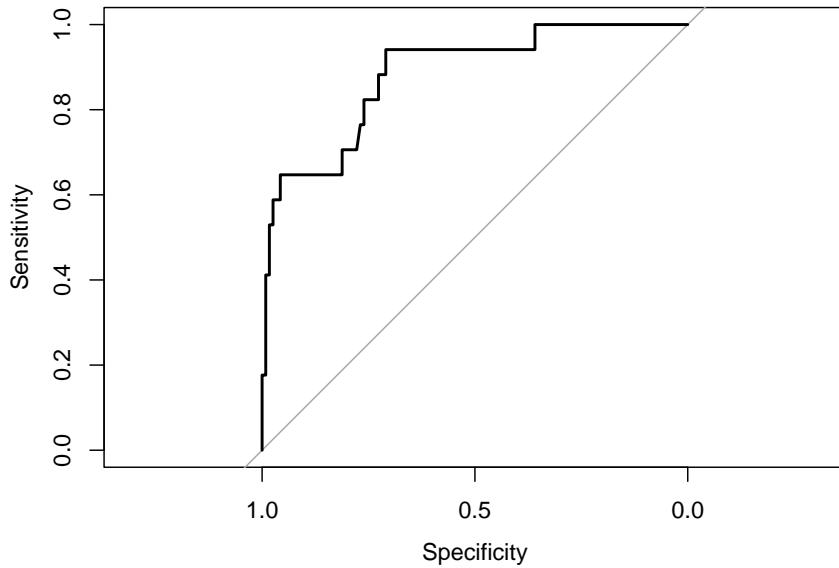
The AUC fitted with `ROCR` (0.883) is slightly different than what `lrm` has told us (0.880), and this also happens if we use the `pROC` approach, demonstrated below.

```
## requires pROC package
roc.modD <-
  roc(resect$died ~ predict(res_modD, type="fitted"),
      ci = TRUE)

roc.modD
```

```
Call:
roc.formula(formula = resect$died ~ predict(res_modD, type = "fitted"),      ci = TRUE)

Data: predict(res_modD, type = "fitted") in 117 controls (resect$died 0) < 17 cases (resect$died
Area under the curve: 0.8826
95% CI: 0.7952-0.97 (DeLong)
plot(roc.modD)
```



12.15.8 Validation of Model D summaries

```
set.seed(432002)
validate(res_modD, B = 100)
```

Divergence or singularity in 5 samples

	index.orig	training	test	optimism	index.corrected	n
Dxy	0.7652	0.8080	0.7352	0.0727	0.6925	95
R2	0.4643	0.5347	0.4119	0.1228	0.3416	95
Intercept	0.0000	0.0000	-0.3533	0.3533	-0.3533	95
Slope	1.0000	1.0000	0.7658	0.2342	0.7658	95
Emax	0.0000	0.0000	0.1308	0.1308	0.1308	95
D	0.2767	0.3415	0.2407	0.1008	0.1759	95
U	-0.0149	-0.0149	0.0883	-0.1032	0.0883	95
Q	0.2916	0.3564	0.1524	0.2040	0.0876	95
B	0.0673	0.0640	0.0736	-0.0096	0.0769	95
g	2.3819	4.0387	2.4635	1.5751	0.8068	95
gp	0.1720	0.1910	0.1632	0.0278	0.1442	95

The C statistic indicates fairly strong discrimination, at $C = 0.88$, although after validation, this looks much weaker (based on $Dxy = 0.6925$, we would have $C = 0.5 + 0.6925/2 = 0.85$) and the Nagelkerke R^2 is also reasonably good, at 0.46, although this, too, is overly optimistic, and we bias-correct through our

12.16. MODEL E: FITTING A REDUCED MODEL IN LIGHT OF MODEL D337

validation study to 0.34.

12.16 Model E: Fitting a Reduced Model in light of Model D

Can you suggest a reduced model (using a subset of the independent variables in model D) that adequately predicts survival?

Based on the anova for model D and the Spearman rho-squared plot, it appears that a two-predictor model using intubation and resection may be sufficient. Neither of the other potential predictors shows a statistically detectable effect in its Wald test.

```
res_modE <- lrm(died ~ intubated + resection, data=resect,
                  x=TRUE, y=TRUE)
res_modE
```

Logistic Regression Model

```
lrm(formula = died ~ intubated + resection, data = resect, x = TRUE,
     y = TRUE)
```

Obs	134	Model Likelihood		Discrimination		Rank	Discrim.
		LR	chi2	Ratio Test	Indexes		
0	117	d.f.	2	R2	0.413	C	0.867
1	17	Pr(> chi2)	<0.0001	g	1.397	Dxy	0.734
max deriv	5e-10			gr	4.043	gamma	0.757
				gp	0.160	tau-a	0.164
				Brier	0.073		
Coef S.E. Wald Z Pr(> Z)							
Intercept	-4.6370	1.0430	-4.45	<0.0001			
intubated	2.8640	0.6479	4.42	<0.0001			
resection	0.5475	0.2689	2.04	0.0418			

The model equation is that the log odds of death is $-4.637 + 2.864 \text{ intubated} + 0.548 \text{ resection}$.

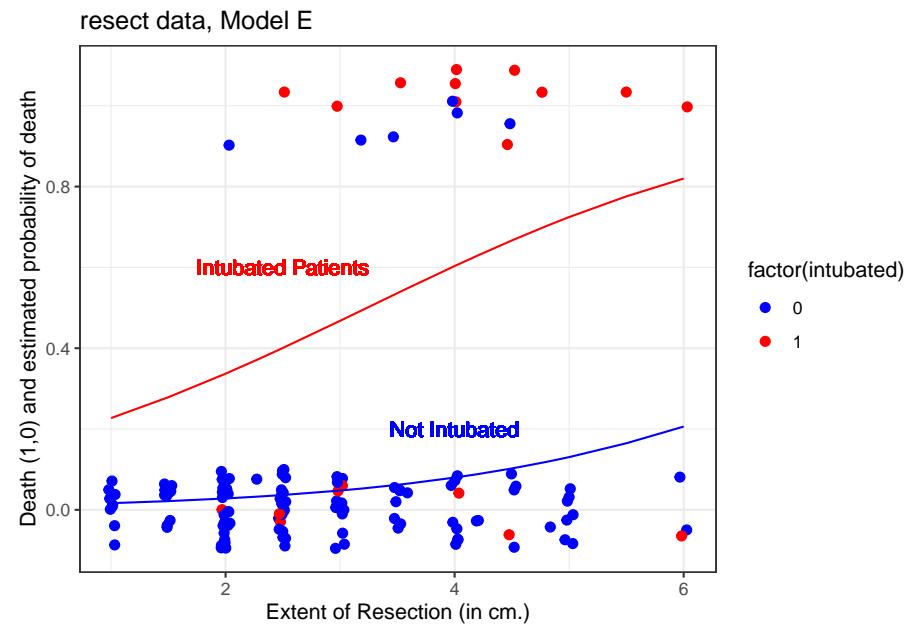
This implies that:

- for intubated patients, the equation is $-1.773 + 0.548 \text{ resection}$, while
- for non-intubated patients, the equation is $-4.637 + 0.548 \text{ resection}$.

We can use the `ilogit` function within the `faraway` package to help plot this.

12.16.1 A Plot comparing the two intubation groups

```
ggplot(resect, aes(x = resection, y = died,
                    col = factor(intubated))) +
  scale_color_manual(values = c("blue", "red")) +
  geom_jitter(size = 2, height = 0.1) +
  geom_line(aes(x = resection,
                 y = faraway::ilogit(-4.637 + 0.548*resection)),
            col = "blue") +
  geom_line(aes(x = resection,
                 y = faraway::ilogit(-1.773 + 0.548*resection)),
            col = "red") +
  geom_text(x = 4, y = 0.2, label = "Not Intubated",
            col="blue") +
  geom_text(x = 2.5, y = 0.6, label = "Intubated Patients",
            col="red") +
  labs(x = "Extent of Resection (in cm.)",
       y = "Death (1,0) and estimated probability of death",
       title = "resect data, Model E")
```



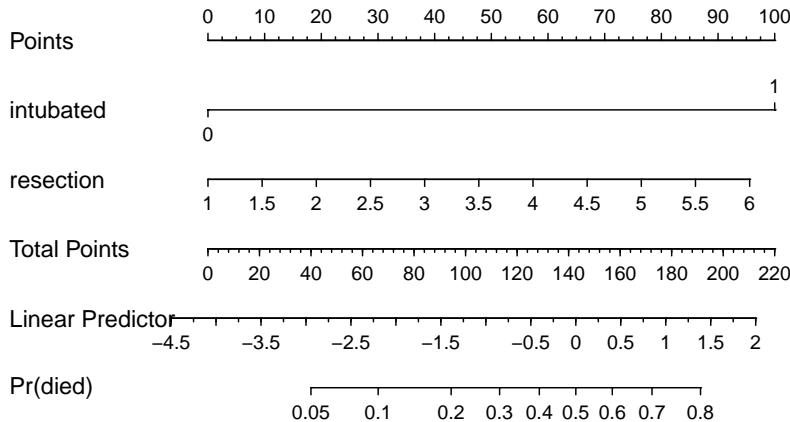
The effect of intubation appears to be very large, compared to the resection size effect.

12.16. MODEL E: FITTING A REDUCED MODEL IN LIGHT OF MODEL D339

12.16.2 Nomogram for Model E

A nomogram of the model would help, too.

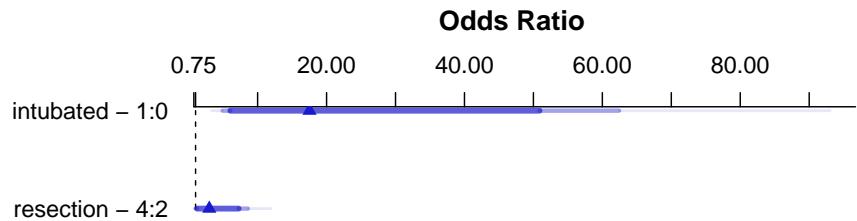
```
plot(nomogram(res_modE, fun=plogis,
               fun.at=c(0.05, seq(0.1, 0.9, by=0.1), 0.95),
               funlabel="Pr(died)"))
```



Again, we see that the effect of intubation is enormous, compared to the effect of resection. Another way to see this is to plot the effect sizes directly.

12.16.3 Effect Sizes from Model E

```
plot(summary(res_modE))
```



```
c effect plot-1.pdf
```

```
summary(res_modE)
```

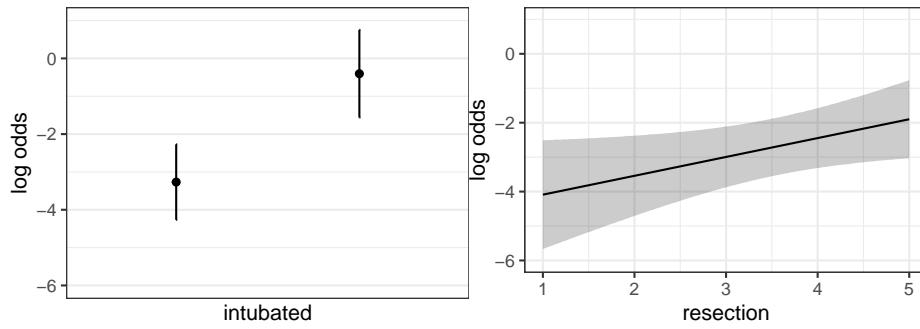
	Effects			Response : died				
Factor	Low	High	Diff.	Effect	S.E.	Lower	0.95	Upper
intubated	0	1	1	2.8640	0.64790	1.59410	4.1338	
Odds Ratio	0	1	1	17.5310	NA	4.92390	62.4160	
resection	2	4	2	1.0949	0.53783	0.04082	2.1491	
Odds Ratio	2	4	2	2.9890	NA	1.04170	8.5769	

12.16.4 Plot In-Sample Predictions for Model E

Here are plots of the effects across the range of each predictor (holding the other constant) on the log odds scale.

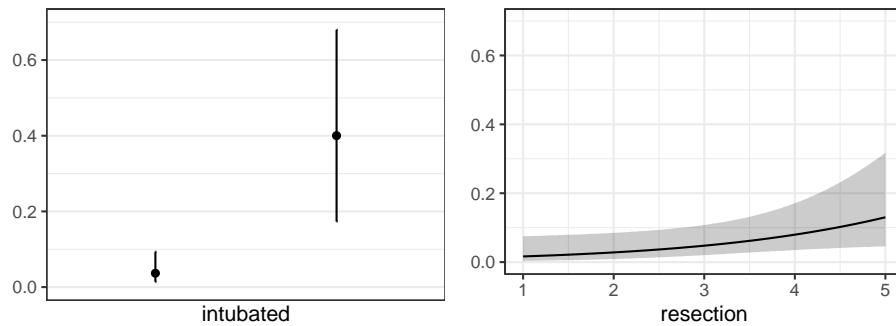
```
ggplot(Predict(res_modE))
```

12.16. MODEL E: FITTING A REDUCED MODEL IN LIGHT OF MODEL D341



We can also capture and plot these results on the probability scale, as follows.

```
ggplot(Predict(res_modE, fun = plogis))
```



12.16.5 ANOVA for Model E

```
anova(res_modE)
```

Factor	Chi-Square	d.f.	P	
intubated	19.54	1	<.0001	
resection	4.14	1	0.0418	
TOTAL	25.47	2	<.0001	

12.16.6 Validation of Model E

```
validate(res_modE, method="boot", B=40)
```

	index.orig	training	test	optimism	index.corrected	n
Dxy	0.7340	0.6896	0.7326	-0.0430	0.7771	40
R2	0.4128	0.3814	0.4025	-0.0211	0.4339	40
Intercept	0.0000	0.0000	0.1367	-0.1367	0.1367	40
Slope	1.0000	1.0000	1.0472	-0.0472	1.0472	40
Emax	0.0000	0.0000	0.0369	0.0369	0.0369	40
D	0.2408	0.2183	0.2339	-0.0157	0.2565	40
U	-0.0149	-0.0149	-0.0001	-0.0148	-0.0001	40
Q	0.2558	0.2332	0.2340	-0.0009	0.2566	40
B	0.0727	0.0727	0.0759	-0.0032	0.0759	40
g	1.3970	1.3391	1.3577	-0.0186	1.4156	40
gp	0.1597	0.1446	0.1563	-0.0117	0.1714	40

Our bootstrap validated assessments of discrimination and goodness of fit look somewhat more reasonable now.

12.16.7 Do any points seem particularly influential?

As a last step, I'll look at influence, and residuals, associated with model E.

```
inf.E <- which.influence(res_modE, cutoff=0.3)
```

```
inf.E
```

```
$Intercept
[1] 84 94

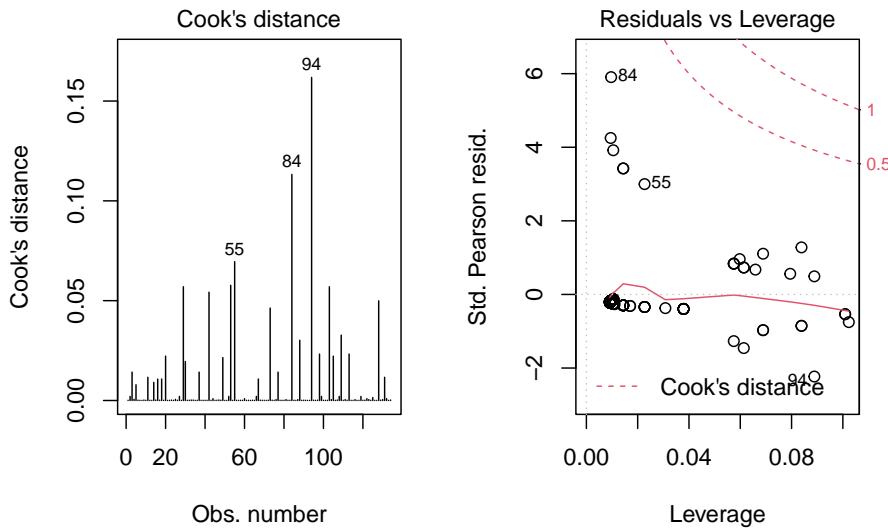
$resection
[1] 84 94
```

```
show.influence(inf.E, dframe = data.frame(resect))
```

Count	resection
84	2 *2
94	2 *6

12.16.8 Fitting Model E using `glm` to get plots about influence

```
res_modEglm <- glm(died ~ intubated + resection,
                     data=resect, family="binomial")
par(mfrow=c(1,2))
plot(res_modEglm, which=c(4:5))
```



Using this `glm` residuals approach, we again see that points 84 and 94 have the largest influence on our model E.

12.17 Concordance: Comparing Model C, D and E's predictions

To start, we'll gather the predictions made by each model (C, D and E) on the probability scale, in one place. Sadly, `augment` from `broom` doesn't work well

with `lrm` fits, so we have to do this on our own.

```
resect_preds <- resect %>%
  mutate(C = predict(res_modC, type = "fitted"),
         D = predict(res_modD, type = "fitted"),
         E = predict(res_modE, type = "fitted"))

head(resect_preds)

# A tibble: 6 x 9
  id    age prior resection intubated died      C      D      E
  <dbl> <dbl> <dbl>     <dbl>     <dbl> <dbl> <dbl> <dbl>
1     1     34     1       2.5      0   0.0705 0.0632 0.0367
2     2     57     0       5       0   0.326   0.0620 0.130
3     3     60     1       4       1   0.187   0.791  0.603
4     4     62     1       4.2      0   0.211   0.158  0.0881
5     5     28     0       6       1   0.504   0.711  0.819
6     6     52     0       3       0   0.0990 0.0737 0.0477
```

And now, we'll use the `gather` command to arrange the models and predicted probabilities in a more useful manner for plotting.

```
res_p <- resect_preds %>%
  gather("model", "prediction", 7:9) %>%
  select(id, died, model, prediction)

head(res_p)

# A tibble: 6 x 4
  id    died model prediction
  <dbl> <dbl> <chr>     <dbl>
1     1     0 C          0.0705
2     2     0 C          0.326
3     3     1 C          0.187
4     4     0 C          0.211
5     5     1 C          0.504
6     6     0 C          0.0990
```

Here's the resulting plot.

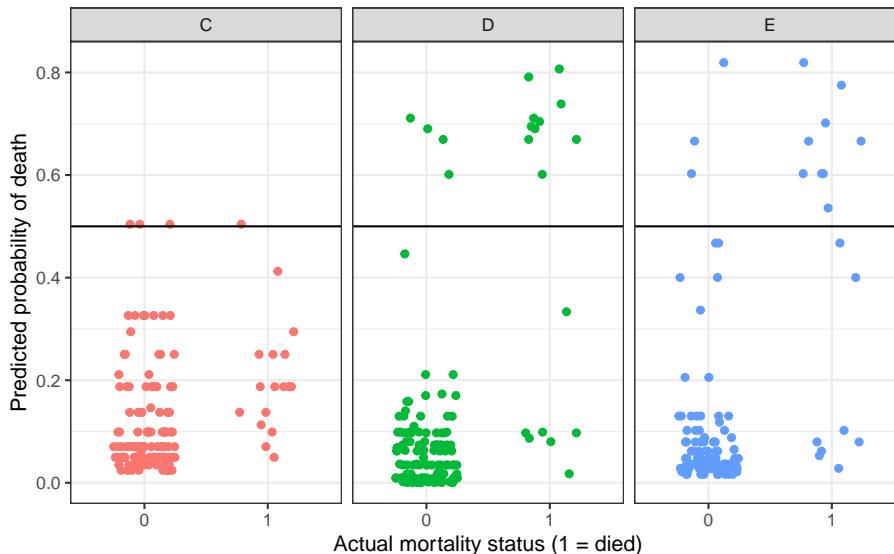
```
ggplot(res_p, aes(x = factor(died), y = prediction,
                   group = model, col = model)) +
  geom_jitter(width = 0.25) +
  geom_hline(yintercept = 0.5) +
  facet_wrap(~ model) +
  guides(color = "none") +
  labs(title = "Comparing Predictions for our Three Models",
       subtitle = "A graphical view of concordance",
```

12.17. CONCORDANCE: COMPARING MODEL C, D AND E'S PREDICTIONS 345

```
x = "Actual mortality status (1 = died)",
y = "Predicted probability of death")
```

Comparing Predictions for our Three Models

A graphical view of concordance



We could specify a particular rule, for example: if the predicted probability of death is 0.5 or greater, then predict “Died.”

```
res_p$rule.5 <- ifelse(res_p$prediction >= 0.5,
                        "Predict Died", "Predict Alive")
```

```
ftable(table(res_p$model, res_p$rule.5, res_p$died))
```

	0	1
C	Predict Alive	114 16
	Predict Died	3 1
D	Predict Alive	113 7
	Predict Died	4 10
E	Predict Alive	114 8
	Predict Died	3 9

And perhaps build the linked table of row probabilities...

```
round(100*prop.table(
  ftable(table(res_p$model, res_p$rule.5, res_p$died)))
,1),2)
```

0	1
0.55	0.45

C	Predict Alive	87.69	12.31
	Predict Died	75.00	25.00
D	Predict Alive	94.17	5.83
	Predict Died	28.57	71.43
E	Predict Alive	93.44	6.56
	Predict Died	25.00	75.00

For example, in model E, 93.44% of those predicted to be alive actually survived, and 75% of those predicted to die actually died.

- Model D does a little better in one direction (94.17% of those predicted by Model D to be alive actually survived) but worse in the other (71.43% of those predicted by Model D to die actually died.)
- Model C does worse than each of the others in both predicting those who survive and those who die.

Note that the approaches discussed here would be useful if we had a new sample to predict on, as well. We could then compare the errors for that new data made by this sort of classification scheme either graphically or in a table.

12.18 Conclusions

It appears that `intubated` status and, to a lesser degree, the extent of the `resection` both play a meaningful role in predicting death associated with tracheal carina resection surgery. Patients who are intubated are associated with worse outcomes (greater risk of death) and more extensive resections are also associated with worse outcomes.

Chapter 13

A Study of Prostate Cancer

13.1 Data Load and Background

The data in `prost.csv` is derived from Stamey et al. (1989) who examined the relationship between the level of prostate-specific antigen and a number of clinical measures in 97 men who were about to receive a radical prostatectomy. The `prost` data, as I'll name it in R, contains 97 rows and 11 columns.

```
prost
```

```
# A tibble: 97 x 10
  subject    lpsa lcavol lweight    age bph      svi    lcp gleason pgg45
  <dbl>   <dbl>   <dbl>   <dbl> <dbl> <chr>   <dbl>   <dbl> <chr>   <dbl>
1       1 -0.431 -0.580    2.77    50 Low      0 -1.39  6        0
2       2 -0.163 -0.994    3.32    58 Low      0 -1.39  6        0
3       3 -0.163 -0.511    2.69    74 Low      0 -1.39  7        20
4       4 -0.163 -1.20     3.28    58 Low      0 -1.39  6        0
5       5  0.372  0.751    3.43    62 Low      0 -1.39  6        0
6       6  0.765 -1.05     3.23    50 Low      0 -1.39  6        0
7       7  0.765  0.737    3.47    64 Medium  0 -1.39  6        0
8       8  0.854  0.693    3.54    58 High    0 -1.39  6        0
9       9  1.05   -0.777   3.54    47 Low      0 -1.39  6        0
10      10  1.05   0.223    3.24    63 Low      0 -1.39  6        0
# ... with 87 more rows
```

Note that a related `prost` data frame is also available as part of several R packages, including the `faraway` package, but there is an error in the `lweight` data for subject 32 in those presentations. The value of `lweight` for subject 32 should not be 6.1, corresponding to a prostate that is 449 grams in size, but instead the `lweight` value should be 3.804438, corresponding to a 44.9 gram

prostate¹.

I've also changed the **gleason** and **bph** variables from their presentation in other settings, to let me teach some additional details.

13.2 Code Book

Variable	Description
subject	subject number (1 to 97)
lpsa	log(prostate specific antigen in ng/ml), our outcome
lcavol	log(cancer volume in cm ³)
lweight	log(prostate weight, in g)
age	age
bph	benign prostatic hyperplasia amount (Low, Medium, or High)
svi	seminal vesicle invasion (1 = yes, 0 = no)
lcp	log(capsular penetration, in cm)
gleason	combined Gleason score (6, 7, or > 7 here)
pgg45	percentage Gleason scores 4 or 5

Notes:

- in general, higher levels of PSA are stronger indicators of prostate cancer. An old standard (established almost exclusively with testing in white males, and definitely flawed) suggested that values below 4 were normal, and above 4 needed further testing. A PSA of 4 corresponds to an **lpsa** of 1.39.
- all logarithms are natural (base *e*) logarithms, obtained in R with the function **log()**
- all variables other than **subject** and **lpsa** are candidate predictors
- the **gleason** variable captures the highest combined Gleason score[[^]Scores range (in these data) from 6 (a well-differentiated, or low-grade cancer) to 9 (a high-grade cancer), although the maximum possible score is 10. 6 is the lowest score used for cancerous prostates. As this combination value increases, the rate at which the cancer grows and spreads should increase. This score refers to the combined Gleason grade, which is based on the sum of two areas (each scored 1-5) that make up most of the cancer.] in a biopsy, and higher scores indicate more aggressive cancer cells. It's stored here as 6, 7, or > 7.
- the **pgg45** variable captures the percentage of individual Gleason scores[[^]The 1-5 scale for individual biopsies are defined so that 1 indicates something that looks like normal prostate tissue, and 5 indicates that the

¹<https://statweb.stanford.edu/~tibs/ElemStatLearn/> attributes the correction to Professor Stephen W. Link.

cells and their growth patterns look very abnormal. In this study, the percentage of 4s and 5s shown in the data appears to be based on 5-20 individual scores in most subjects.] that are 4 or 5, on a 1-5 scale, where higher scores indicate more abnormal cells.

13.3 Additions for Later Use

The code below adds to the `prost` tibble:

- a factor version of the `svi` variable, called `svi_f`, with levels No and Yes,
- a factor version of `gleason` called `gleason_f`, with the levels ordered > 7, 7, and finally 6,
- a factor version of `bph` called `bph_f`, with levels ordered Low, Medium, High,
- a centered version of `lcavol` called `lcavol_c`,
- exponentiated `cavol` and `psa` results derived from the natural logarithms `lcavol` and `lpsa`.

```
prost <- prost %>%
  mutate(svi_f = fct_recode(factor(svi), "No" = "0", "Yes" = "1"),
        gleason_f = fct_relevel(gleason, c("> 7", "7", "6")),
        bph_f = fct_relevel(bph, c("Low", "Medium", "High")),
        lcavol_c = lcavol - mean(lcavol),
        cavol = exp(lcavol),
        psa = exp(lpsa))

glimpse(prost)

Rows: 97
Columns: 16
$ subject   <dbl> 1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13, 14, 15, 16, 17, 1~
$ lpsa       <dbl> -0.4307829, -0.1625189, -0.1625189, -0.1625189, 0.3715636, 0~
$ lcavol     <dbl> -0.5798185, -0.9942523, -0.5108256, -1.2039728, 0.7514161, --
$ lweight    <dbl> 2.769459, 3.319626, 2.691243, 3.282789, 3.432373, 3.228826, ~
$ age        <dbl> 50, 58, 74, 58, 62, 50, 64, 58, 47, 63, 65, 63, 63, 67, 57, ~
$ bph        <chr> "Low", "Low", "Low", "Low", "Low", "Medium", "High", ~
$ svi        <dbl> 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, ~
$ lcp        <dbl> -1.3862944, -1.3862944, -1.3862944, -1.3862944, -1.3862944, ~
$ gleason    <chr> "6", "6", "7", "6", "6", "6", "6", "6", "6", "6", "6", "6", ~
$ pgg45      <dbl> 0, 0, 20, 0, 0, 0, 0, 0, 0, 30, 5, 5, 0, 30, 0, 0, ~
$ svi_f      <fct> No, ~
$ gleason_f  <fct> 6, 6, 7, 6, 6, 6, 6, 6, 6, 7, 7, 7, 6, 7, 6, 6, ~
$ bph_f      <fct> Low, Low, Low, Low, Low, Medium, High, Low, Low, Low, ~
$ lcavol_c   <dbl> -1.9298281, -2.3442619, -1.8608352, -2.5539824, -0.5985935, ~
$ cavol      <dbl> 0.56, 0.37, 0.60, 0.30, 2.12, 0.35, 2.09, 2.00, 0.46, 1.25, ~
```

```
$ psa      <dbl> 0.65, 0.85, 0.85, 0.85, 1.45, 2.15, 2.15, 2.35, 2.85, 2.85, ~
```

13.4 Fitting and Evaluating a Two-Predictor Model

To begin, let's use two predictors (`lcavol` and `svi`) and their interaction in a linear regression model that predicts `lpsa`. I'll call this model `c11_prost_A`

Earlier, we centered the `lcavol` values to facilitate interpretation of the terms. I'll use that centered version (called `lcavol_c`) of the quantitative predictor, and the 1/0 version of the `svi` variable[^We could certainly use the factor version of `svi` here, but it won't change the model in any meaningful way. There's no distinction in model *fitting* via `lm` between a 0/1 numeric variable and a No/Yes factor variable. The factor version of this information will be useful elsewhere, for instance in plotting the model.].

```
c11_prost_A <- lm(lpsa ~ lcavol_c * svi, data = prost)
summary(c11_prost_A)
```

Call:

```
lm(formula = lpsa ~ lcavol_c * svi, data = prost)
```

Residuals:

Min	1Q	Median	3Q	Max
-1.6305	-0.5007	0.1266	0.4886	1.6847

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	2.33134	0.09128	25.540	< 2e-16 ***
lcavol_c	0.58640	0.08207	7.145	1.98e-10 ***
svi	0.60132	0.35833	1.678	0.0967 .
lcavol_c:svi	0.06479	0.26614	0.243	0.8082

Signif. codes:	0 ***	0.001 **	0.01 *	0.05 .
	'	'	'	'

```
Residual standard error: 0.7595 on 93 degrees of freedom
Multiple R-squared:  0.5806,    Adjusted R-squared:  0.5671
F-statistic: 42.92 on 3 and 93 DF,  p-value: < 2.2e-16
```

13.4.1 Using `tidy`

It can be very useful to build a data frame of the model's results. We can use the `tidy` function in the `broom` package to do so.

```
tidy(c11_prost_A)

# A tibble: 4 x 5
  term      estimate std.error statistic p.value
  <chr>     <dbl>     <dbl>     <dbl>    <dbl>
1 (Intercept) 2.33     0.0913    25.5   8.25e-44
2 lcavol_c     0.586    0.0821    7.15   1.98e-10
3 svi         0.601     0.358     1.68   9.67e- 2
4 lcavol_c:svi 0.0648   0.266     0.243  8.08e- 1
```

This makes it much easier to pull out individual elements of the model fit.

For example, to specify the coefficient for **svi**, rounded to three decimal places, I could use

```
tidy(c11_prost_A) %>% filter(term == "svi") %>% select(estimate) %>% round(., 3)
```

- The result is 0.601.
- If you look at the Markdown file, you'll see that the number shown in the bullet point above this one was generated using inline R code, and the function specified above.

13.4.2 Interpretation

1. The intercept, 2.33, for the model is the predicted value of **lpsa** when **lcavol** is at its average and there is no seminal vesicle invasion (e.g. **svi** = 0).
2. The coefficient for **lcavol_c**, 0.59, is the predicted change in **lpsa** associated with a one unit increase in **lcavol** (or **lcavol_c**) when there is no seminal vesicle invasion.
3. The coefficient for **svi**, 0.6, is the predicted change in **lpsa** associated with having no **svi** to having an **svi** while the **lcavol** remains at its average.
4. The coefficient for **lcavol_c:svi**, the product term, which is 0.06, is the difference in the slope of **lcavol_c** for a subject with **svi** as compared to one with no **svi**.

13.5 Exploring Model c11_prost_A

The **glance** function from the **broom** package builds a nice one-row summary for the model.

```
glance(c11_prost_A)
```

```
# A tibble: 1 x 12
  r.squared adj.r.squared sigma statistic p.value    df logLik    AIC    BIC
  <dbl>        <dbl>     <dbl>     <dbl>    <dbl>    <dbl> <dbl>    <dbl> <dbl>
```

```
1      0.581          0.567 0.759      42.9 1.68e-17    3   -109.  228.  241.
# ... with 3 more variables: deviance <dbl>, df.residual <int>, nobs <int>
```

This summary includes, in order,

- the model R^2 , adjusted R^2 and $\hat{\sigma}$, the residual standard deviation,
- the ANOVA F statistic and associated p value,
- the number of degrees of freedom used by the model, and its log-likelihood ratio
- the model's AIC (Akaike Information Criterion) and BIC (Bayesian Information Criterion)
- the model's deviance statistic and residual degrees of freedom

13.5.1 summary for Model c11_prost_A

If necessary, we can also run `summary` on this `c11_prost_A` object to pick up some additional summaries. Since the `svi` variable is binary, the interaction term is, too, so the t test here and the F test in the ANOVA yield the same result.

```
summary(c11_prost_A)

Call:
lm(formula = lpsa ~ lcavol_c * svi, data = prost)

Residuals:
    Min      1Q  Median      3Q      Max 
-1.6305 -0.5007  0.1266  0.4886  1.6847 

Coefficients:
            Estimate Std. Error t value Pr(>|t|)    
(Intercept) 2.33134   0.09128 25.540 < 2e-16 ***
lcavol_c    0.58640   0.08207  7.145 1.98e-10 ***
svi         0.60132   0.35833  1.678  0.0967 .  
lcavol_c:svi 0.06479   0.26614  0.243  0.8082    
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.7595 on 93 degrees of freedom
Multiple R-squared:  0.5806,    Adjusted R-squared:  0.5671 
F-statistic: 42.92 on 3 and 93 DF,  p-value: < 2.2e-16
```

If you've forgotten the details of the pieces of this summary, review the Part C Notes from 431.

13.5.2 Adjusted R^2

R^2 is greedy.

- R^2 will always suggest that we make our models as big as possible, often including variables of dubious predictive value.
- As a result, there are various methods for penalizing R^2 so that we wind up with smaller models.
- The **adjusted** R^2 is often a useful way to compare multiple models for the same response.
 - $R_{adj}^2 = 1 - \frac{(1-R^2)(n-1)}{n-k}$, where n = the number of observations and k is the number of coefficients estimated by the regression (including the intercept and any slopes).
 - So, in this case, $R_{adj}^2 = 1 - \frac{(1-0.5806)(97-1)}{97-4} = 0.5671$
 - The adjusted R^2 value is not, technically, a proportion of anything, but it is comparable across models for the same outcome.
 - The adjusted R^2 will always be less than the (unadjusted) R^2 .

13.5.3 Coefficient Confidence Intervals

Here are the 90% confidence intervals for the coefficients in Model A. Adjust the `level` to get different intervals.

```
confint(c11_prost_A, level = 0.90)

            5 %      95 %
(Intercept) 2.17968697 2.4830012
lcavol_c     0.45004577 0.7227462
svi          0.00599401 1.1966454
lcavol_c:svi -0.37737623 0.5069622
```

What can we conclude from this about the utility of the interaction term?

13.5.4 ANOVA for Model c11_prost_A

The interaction term appears unnecessary. We might wind up fitting the model without it. A complete ANOVA test is available, including a p value, if you want it.

```
anova(c11_prost_A)

Analysis of Variance Table

Response: lpsa
           Df Sum Sq Mean Sq F value    Pr(>F)
lcavol_c     1 69.003 69.003 119.6289 < 2.2e-16 ***
```

```

svi           1  5.237   5.237   9.0801  0.003329 **
lcavol_c:svi 1  0.034   0.034   0.0593  0.808191
Residuals     93 53.643   0.577
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

```

Note that the `anova` approach for a `lm` object is sequential. The first row shows the impact of `lcavol_c` as compared to a model with no predictors (just an intercept). The second row shows the impact of adding `svi` to a model that already contains `lcavol_c`. The third row shows the impact of adding the interaction (product) term to the model with the two main effects. So the order in which the variables are added to the regression model matters for this ANOVA. The F tests here describe the incremental impact of each covariate in turn.

13.5.5 Residuals, Fitted Values and Standard Errors with `augment`

The `augment` function in the `broom` package builds a data frame including the data used in the model, along with predictions (fitted values), residuals and other useful information.

```
c11_prost_A_frame <- augment(c11_prost_A) %>%tbl_df
summary(c11_prost_A_frame)
```

lpsa	lcavol_c	svi	.fitted
Min. :-0.4308	Min. :-2.69708	Min. :0.0000	Min. :0.7498
1st Qu.: 1.7317	1st Qu.:-0.83719	1st Qu.:0.0000	1st Qu.:1.8404
Median : 2.5915	Median : 0.09691	Median :0.0000	Median :2.3950
Mean : 2.4784	Mean : 0.00000	Mean :0.2165	Mean :2.4784
3rd Qu.: 3.0564	3rd Qu.: 0.77703	3rd Qu.:0.0000	3rd Qu.:3.0709
Max. : 5.5829	Max. : 2.47099	Max. :1.0000	Max. :4.5417
.resid	.hat	.sigma	.cooksdf
Min. :-1.6305	Min. :0.01316	Min. :0.7423	Min. :0.0000069
1st Qu.:-0.5007	1st Qu.:0.01562	1st Qu.:0.7569	1st Qu.:0.0007837
Median : 0.1266	Median :0.02498	Median :0.7617	Median :0.0034699
Mean : 0.0000	Mean :0.04124	Mean :0.7595	Mean :0.0111314
3rd Qu.: 0.4886	3rd Qu.:0.04939	3rd Qu.:0.7631	3rd Qu.:0.0103533
Max. : 1.6847	Max. : 0.24627	Max. :0.7636	Max. :0.1341093
.std.resid			
Min. :-2.194508			
1st Qu.:-0.687945			
Median : 0.168917			
Mean : 0.001249			
3rd Qu.: 0.653612			
Max. : 2.261830			

Elements shown here include:

- **.fitted** Fitted values of model (or predicted values)
- **.se.fit** Standard errors of fitted values
- **.resid** Residuals (observed - fitted values)
- **.hat** Diagonal of the hat matrix (these indicate *leverage* - points with high leverage indicate unusual combinations of predictors - values more than 2-3 times the mean leverage are worth some study - leverage is always between 0 and 1, and measures the amount by which the predicted value would change if the observation's y value was increased by one unit - a point with leverage 1 would cause the line to follow that point perfectly)
- **.sigma** Estimate of residual standard deviation when corresponding observation is dropped from model
- **.cooksdi** Cook's distance, which helps identify influential points (values of Cook's d > 0.5 may be influential, values > 1.0 almost certainly are - an influential point changes the fit substantially when it is removed from the data)
- **.std.resid** Standardized residuals (values above 2 in absolute value are worth some study - treat these as normal deviates [Z scores], essentially)

See `?augment.lm` in R for more details.

13.5.6 Making Predictions with c11_prost_A

Suppose we want to predict the `lpsa` for a patient with cancer volume equal to this group's mean, for both a patient with and without seminal vesicle invasion, and in each case, we want to use a 90% prediction interval?

```
newdata <- data.frame(lcavol_c = c(0,0), svi = c(0,1))
predict(c11_prost_A, newdata, interval = "prediction", level = 0.90)
```

	fit	lwr	upr
1	2.331344	1.060462	3.602226
2	2.932664	1.545742	4.319586

Since the predicted value in `fit` refers to the natural logarithm of PSA, to make the predictions in terms of PSA, we would need to exponentiate. The code below will accomplish that task.

```
pred <- predict(c11_prost_A, newdata, interval = "prediction", level = 0.90)
exp(pred)
```

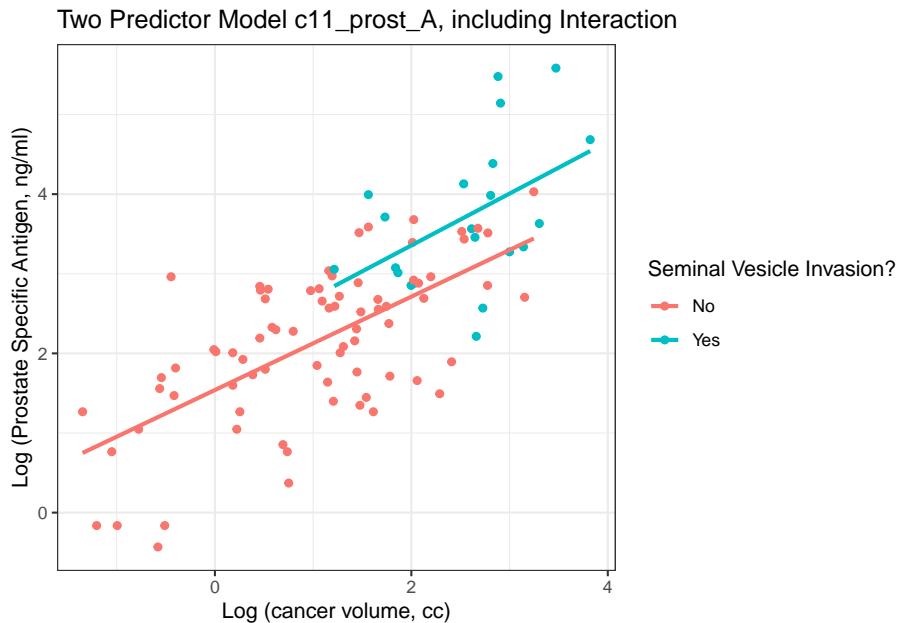
	fit	lwr	upr
1	10.29177	2.887706	36.67978
2	18.77758	4.691450	75.15750

13.6 Plotting Model c11_prost_A

13.6.0.1 Plot logs conventionally

Here, we'll use `ggplot2` to plot the logarithms of the variables as they came to us, on a conventional coordinate scale. Note that the lines are nearly parallel. What does this suggest about our Model A?

```
ggplot(prost, aes(x = lcavol, y = lpsa, group = svi_f, color = svi_f)) +
  geom_point() +
  geom_smooth(method = "lm", se = FALSE) +
  scale_color_discrete(name = "Seminal Vesicle Invasion?") +
  theme_bw() +
  labs(x = "Log (cancer volume, cc)",
       y = "Log (Prostate Specific Antigen, ng/ml)",
       title = "Two Predictor Model c11_prost_A, including Interaction")`geom_smooth()` using formula 'y ~ x'
```



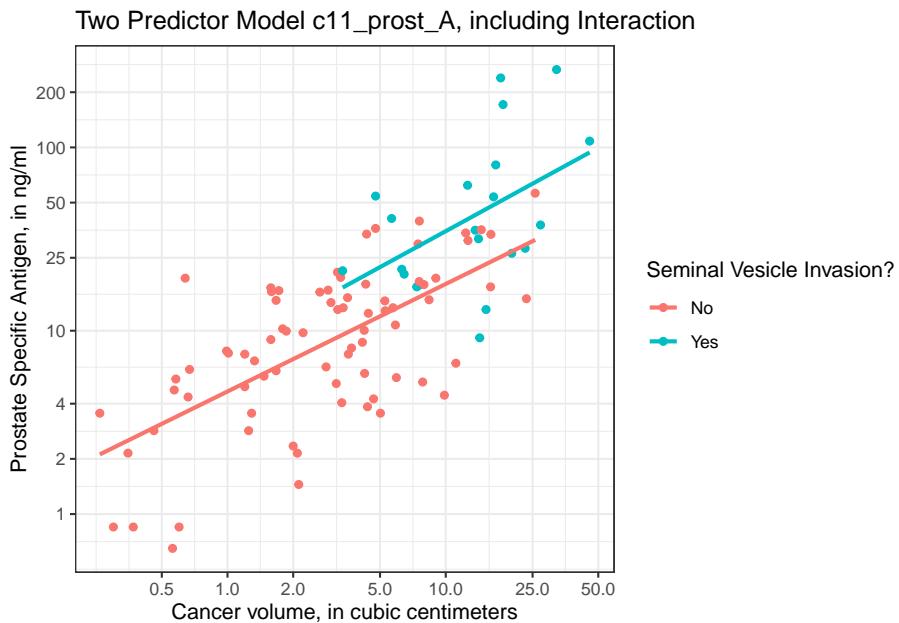
13.6.0.2 Plot on log-log scale

Another approach (which might be easier in some settings) would be to plot the raw values of Cancer Volume and PSA, but use logarithmic axes, again using the natural (base e) logarithm, as follows. If we use the default choice with ‘trans =

"log," we'll find a need to select some useful break points for the grid, as I've done in what follows.

```
ggplot(prost, aes(x = cavol, y = psa, group = svi_f, color = svi_f)) +
  geom_point() +
  geom_smooth(method = "lm", se = FALSE) +
  scale_color_discrete(name = "Seminal Vesicle Invasion?")
  scale_x_continuous(trans = "log",
                      breaks = c(0.5, 1, 2, 5, 10, 25, 50)) +
  scale_y_continuous(trans = "log",
                      breaks = c(1, 2, 4, 10, 25, 50, 100, 200)) +
  theme_bw() +
  labs(x = "Cancer volume, in cubic centimeters",
       y = "Prostate Specific Antigen, in ng/ml",
       title = "Two Predictor Model c11_prost_A, including Interaction")
```

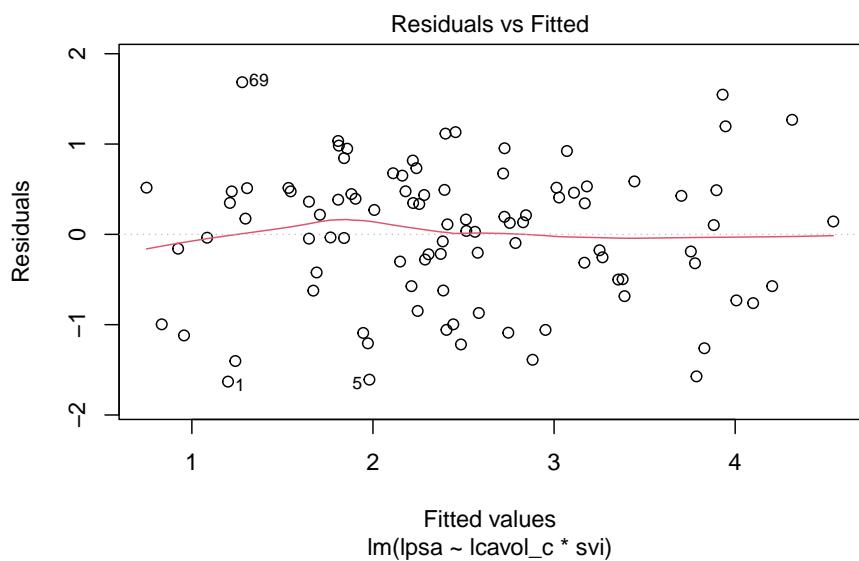
`geom_smooth()` using formula 'y ~ x'



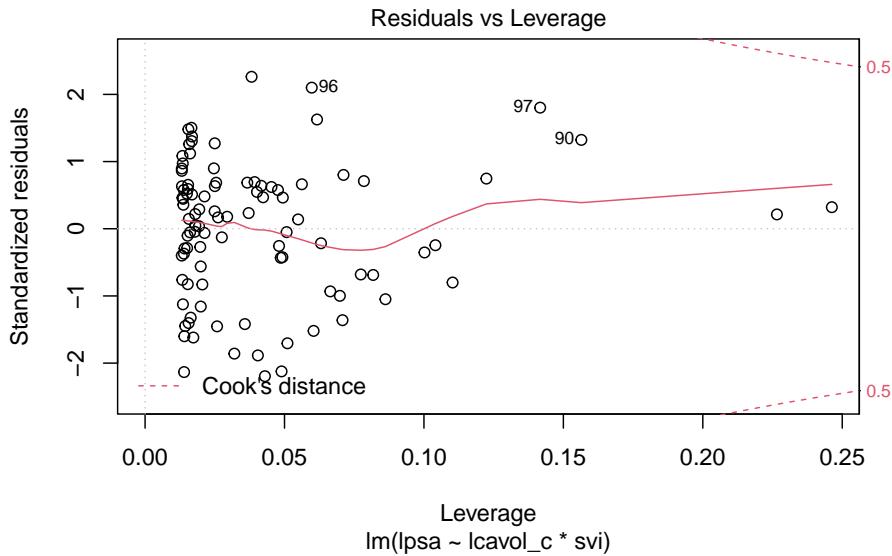
I've used the break point of 4 on the Y axis because of the old rule suggesting further testing for asymptomatic men with PSA of 4 or higher, but the other break points are arbitrary - they seemed to work for me, and used round numbers.

13.6.1 Residual Plots of c11_prost_A

```
plot(c11_prost_A, which = 1)
```



```
plot(c11_prost_A, which = 5)
```



13.7 Cross-Validation of Model `c11_prost_A`

Suppose we want to evaluate whether our model `c11_prost_A` predicts effectively in new data.

13.7.1 A Validation Split Approach

We'll first demonstrate a validation split approach (used, for instance, in 431) which splits our sample into a separate training (perhaps 70% of the data) and test (perhaps 30% of the data) samples, and then:

- fit the model in the training sample,
- use the resulting model to make predictions for `lpsa` in the test sample, and
- evaluate the quality of those predictions, perhaps by comparing the results to what we'd get using a different model.

Our goal will be to cross-validate model `c11_prost_A`, which, you'll recall, uses `lcavol_c`, `svi` and their interaction, to predict `lpsa` in the `prost` data.

We'll start by identifying a random sample of 70% of our `prost` data in a training sample (which we'll call `prost_train`, and leave the rest as our test sample, called `prost_test`. To do this, we'll use the `createDataPartition` function

from the `caret` package. We need only specify the data set and outcome variable, like so.

```
set.seed(4322020)
split_samples <- prost$lpsa %>%
  createDataPartition(p = 0.7, list = FALSE)

prost_train <- prost[split_samples,]
prost_test <- prost[-split_samples,]
```

- Note the need for a comma after `split_samples` in the isolation of the training and test samples.
- Don't forget to pre-specify the random seed, for replicability, as I've done here.

Let's verify that we now have the samples we expect...

```
dim(prost_train)
```

```
[1] 70 16
```

```
dim(prost_test)
```

```
[1] 27 16
```

OK. Next, we'll run the `c11_prost_A` model in the training sample.

```
c11_prost_A_train <- prost_train %$%
  lm(lpsa ~ lcavol_c * svi)
```

```
c11_prost_A_train
```

Call:

```
lm(formula = lpsa ~ lcavol_c * svi)
```

Coefficients:

(Intercept)	lcavol_c	svi	lcavol_c:svi
2.2956	0.5919	0.6282	0.1597

Then we'll use the coefficients from this model to obtain predicted `lpsa` values in the test sample.

```
c11_prost_A_preds <-
  c11_prost_A_train %>% predict(prost_test)
```

```
c11_prost_A_preds[1:3]
```

1	2	3
1.647217	2.210338	2.408746

Now, we can use the `postResample` function from the `caret` package to obtain several key summaries of fit quality for our model. Here, we specify the estimates (or predictions), and then the observed values to the `postResample` function.

```
postResample(c11_prost_A_preds, prost_test$lpsa)
```

```
RMSE    Rsquared      MAE
0.6651311 0.5503311 0.5524172
```

These summary statistics are:

- the RMSE or root mean squared error, which measures the average difference (i.e. prediction error) between the observed known outcome values and the values predicted by the model by first squaring all of the errors, averaging them, and then taking the square root of the result. The lower the RMSE, the better the model.
- the Rsquared or R^2 , which is just the square of the Pearson correlation coefficient relating the predicted and observed values, so we'd like this to be as large as possible, and
- the MAE or mean absolute error, which is a bit less sensitive to outliers than the RMSE, because it measures the average prediction error by taking the absolute value of each error, and then grabbing the average of those values. The lower the MAE, the better the model.

These statistics are more helpful, generally, for comparing multiple models to each other, than for making final decisions on their own. The `caret` package also provides individual functions to gather the elements of `postResample` as follows.

```
prost_A_summaries <- tibble(
  RMSE = RMSE(c11_prost_A_preds, prost_test$lpsa),
  R2 = R2(c11_prost_A_preds, prost_test$lpsa),
  MAE = MAE(c11_prost_A_preds, prost_test$lpsa)
)

prost_A_summaries

# A tibble: 1 x 3
  RMSE    R2     MAE
  <dbl> <dbl> <dbl>
1 0.665 0.550 0.552
```

13.7.2 K-Fold Cross-Validation

One problem with the validation split approach is that with a small data set like `prost`, we may be reluctant to cut our sample size for the training or the testing down because we're afraid that our model building and testing will

be hampered by a small sample size. A potential solution is the idea of **K-fold cross-validation**, which involves partitioning our data into a series of K training-test subsets, and then combining the results. Specifically, we'll try a *5-fold cross validation* here. (K is usually taken to be either 5 or 10.)

The approach includes the following steps.

1. Randomly split the `prost` data into 5 subsets (for 5-fold validation).
2. Reserve one subset and train the model on all other subsets.
3. Test the model on the reserved subset and record the prediction error.
4. Repeat this process until each of the k subsets has served as the test set.
5. Compute the average of the k recorded errors. This is called the cross-validation error and serves as the primary performance metric for the model.

Again using tools from the `caret` packages, we'll first define our `trainControl` approach.

```
set.seed(43220201)
train.control <- trainControl(method = "cv", number = 5)
```

Then we train the model, and obtain the usual summaries of model fit quality.

```
c11_modelA_cv <- train(lpsa ~ lcavol_c * svi,
                        data = prost, method = "lm",
                        trControl = train.control)

c11_modelA_cv
```

Linear Regression

```
97 samples
 2 predictor
```

```
No pre-processing
Resampling: Cross-Validated (5 fold)
Summary of sample sizes: 77, 77, 77, 79, 78
Resampling results:
```

RMSE	R squared	MAE
0.7583897	0.5657784	0.6213951

Tuning parameter 'intercept' was held constant at a value of TRUE

We can then look at the model fit by this cross-validation approach.

```
summary(c11_modelA_cv)
```

```
Call:
lm(formula = .outcome ~ ., data = dat)
```

Residuals:

Min	1Q	Median	3Q	Max
-1.6305	-0.5007	0.1266	0.4886	1.6847

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)							
(Intercept)	2.33134	0.09128	25.540	< 2e-16 ***							
lcavol_c	0.58640	0.08207	7.145	1.98e-10 ***							
svi	0.60132	0.35833	1.678	0.0967 .							
`lcavol_c:svi`	0.06479	0.26614	0.243	0.8082							

Signif. codes:	0	'***'	0.001	'**'	0.01	'*'	0.05	'.'	0.1	' '	1

Residual standard error: 0.7595 on 93 degrees of freedom

Multiple R-squared: 0.5806, Adjusted R-squared: 0.5671

F-statistic: 42.92 on 3 and 93 DF, p-value: < 2.2e-16

or, if you prefer,

```
tidy(summary(c11_modelA_cv), conf.int = TRUE) %>%
  kable(digits = 3)
```

term	estimate	std.error	statistic	p.value	conf.low	conf.high
(Intercept)	2.331	0.091	25.540	0.000	2.150	2.513
lcavol_c	0.586	0.082	7.145	0.000	0.423	0.749
svi	0.601	0.358	1.678	0.097	-0.110	1.313
`lcavol_c:svi`	0.065	0.266	0.243	0.808	-0.464	0.593

and

```
glance(summary(c11_modelA_cv)) %>%
  kable(digits = c(3,3,3,2,3,0))
```

r.squared	adj.r.squared	sigma	statistic	p.value	df	df.residual	nobs
0.581	0.567	0.759	42.92	0	3	93	97

13.7.3 Comparing Models with 5-fold Cross-Validation

To make this a bit more realistic, let's compare two models:

- our existing linear model $\text{lpsa} \sim \text{lcavol_c} * \text{svi}$, and
- a *robust* linear model fit with the `r1m` function in R, to predict `lpsa` using `lcavol_c` and `svi` but not the interaction between them.

The main purpose of *robust* linear models is to reduce the influence of specifically outlier or high leverage data points.

Here's that robust fit in the original `prost_train` data set. Note that fitting a robust linear model requires the choice of a ψ function, for which R provides three approaches, called the Huber, Hampel and Tukey bisquare approaches. In this fit, I'll just let R pick its default choice.

```
modelR <- prost_train %$% rlm(lpsa ~ lcavol_c + svi)

summary(modelR)
```

```
Call: rlm(formula = lpsa ~ lcavol_c + svi)
Residuals:
    Min      1Q  Median      3Q      Max 
-1.5904 -0.5869  0.1139  0.4902  1.7254 

Coefficients:
            Value   Std. Error t value
(Intercept) 2.3241  0.1356   17.1459
lcavol_c     0.6050  0.1139    5.3123
svi          0.7526  0.3296    2.2836

Residual standard error: 0.7819 on 67 degrees of freedom
```

Compare this with the standard ordinary least squares fit to the same data (again without the interaction term), and you'll see that in this case, the main differences are in the estimated standard errors, but the slope coefficients are also a bit smaller in the robust model.

```
model0 <- prost_train %$% lm(lpsa ~ lcavol_c + svi)

summary(model0)
```

```
Call:
lm(formula = lpsa ~ lcavol_c + svi)

Residuals:
    Min      1Q  Median      3Q      Max 
-1.5661 -0.5616  0.1128  0.5130  1.7535 

Coefficients:
            Estimate Std. Error t value Pr(>|t|)    
(Intercept) 2.30169    0.11291  20.386 < 2e-16 ***
lcavol_c    0.60820    0.09487   6.411 1.69e-08 ***
svi         0.79147    0.27451   2.883  0.00529 **  
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.7948 on 67 degrees of freedom
```

```
Multiple R-squared:  0.5924,    Adjusted R-squared:  0.5803
F-statistic:  48.7 on 2 and 67 DF,  p-value: 8.738e-14
```

So, how can we do 5-fold cross-validation on our model R, and also let the computer pick which of the three types of initial weights (Huber, Hampel or Tukey Bisquare) might be most appropriate? As follows...

```
c11_modelR_cv <- train(lpsa ~ lcavol_c + svi,
                         data = prost, method = "rlm",
                         trControl = train.control)
```

```
c11_modelR_cv
```

```
Robust Linear Model
```

```
97 samples
 2 predictor
```

```
No pre-processing
Resampling: Cross-Validated (5 fold)
Summary of sample sizes: 78, 77, 77, 78, 78
Resampling results across tuning parameters:
```

intercept	psi	RMSE	Rquared	MAE
FALSE	psi.huber	2.1261462	0.3643515	1.8841211
FALSE	psi.hampel	2.1253158	0.3652711	1.8837252
FALSE	psi.bisquare	2.1260854	0.3560150	1.8818567
TRUE	psi.huber	0.7683316	0.5741134	0.6358550
TRUE	psi.hampel	0.7634386	0.5773235	0.6352990
TRUE	psi.bisquare	0.7690475	0.5745092	0.6371126

```
RMSE was used to select the optimal model using the smallest value.
The final values used for the model were intercept = TRUE and psi = psi.hampel.
```

Compare these RMSE, Rquared and MAE values to those we observed in the interaction model with lm earlier...

```
c11_modelA_cv
```

```
Linear Regression
```

```
97 samples
 2 predictor
```

```
No pre-processing
Resampling: Cross-Validated (5 fold)
Summary of sample sizes: 77, 77, 77, 79, 78
Resampling results:
```

RMSE	Rsquared	MAE
0.7583897	0.5657784	0.6213951

Tuning parameter 'intercept' was held constant at a value of TRUE

The robust model shows a larger R-Squared, smaller RMSE and smaller MAE than the interaction model. Perhaps we'll focus further on model R going forward...

```
summary(c11_modelR_cv)
```

Call: rlm(formula = .outcome ~ ., data = dat, psi = psi)

Residuals:

Min	1Q	Median	3Q	Max
-1.6203	-0.5297	0.1191	0.4852	1.6939

Coefficients:

	Value	Std. Error	t value
(Intercept)	2.3340	0.0939	24.8529
lcavol_c	0.5930	0.0807	7.3521
svi	0.6685	0.2296	2.9111

Residual standard error: 0.7483 on 94 degrees of freedom

Let's stop there for now. Next, we'll consider the problem of considering adding more predictors to a linear model, and then making sensible selections as to which predictors actually should be incorporated.

Chapter 14

Multiple Imputation and Linear Regression

14.1 Developing a `smart_16` data set

In this chapter, we'll return to the `smart_ohio` file based on data from BRFSS 2017 that we built and cleaned back at the start of these Notes.

```
smart_ohio <- readRDS(here("data", "smart_ohio.Rds"))
```

We're going to look at a selection of variables from this tibble, among subjects who have been told they have diabetes, and who also provided a response to our `physhealth` (Number of Days Physical Health Not Good) variable, which asks "Now thinking about your physical health, which includes physical illness and injury, for how many days during the past 30 days was your physical health not good?" We'll build two models. In this chapter, we'll look at a linear model for `physhealth` and in the next chapter, we'll look at a logistic regression describing whether or not the subject's `physhealth` response was at least 1.

```
smart_16 <- smart_ohio %>%
  filter(dm_status == "Diabetes") %>%
  filter(complete.cases(physhealth)) %>%
  mutate(bad_phys = ifelse(physhealth > 0, 1, 0),
        comor = hx_mi + hx_chd + hx_stroke + hx_asthma +
        hx_skinc + hx_otherc + hx_copd + hx_arthr) %>%
  select(SEQNO, mmsa, physhealth, bad_phys, age_imp, smoke100,
         comor, hx_depress, bmi, activity)
```

The variables included in this `smart_16` tibble are:

Variable	Description
SEQNO	respondent identification number (all begin with 2016)
mmsa	
physhealth	Now thinking about your physical health, which includes physical illness and injury, for how many days during the past 30 days was your physical health not good?
bad_phys	Is physhealth 1 or more?
age_imp	Age in years (imputed from age categories)
smoke100	Have you smoked at least 100 cigarettes in your life? (1 = yes, 0 = no)
hx_depress	Has a doctor, nurse, or other health professional ever told you that you have a depressive disorder, including depression, major depression, dysthymia, or minor depression?
bmi	Body mass index, in kg/m ²
activity	Physical activity (Highly Active, Active, Insufficiently Active, Inactive)
comor	Sum of 8 potential groups of comorbidities (see below)

The **comor** variable is the sum of the following 8 variables, each of which is measured on a 1 = Yes, 0 = No scale, and begin with “Has a doctor, nurse, or other health professional ever told you that you had ...”

- **hx_mi**: a heart attack, also called a myocardial infarction?
- **hx_chd**: angina or coronary heart disease?
- **hx_stroke**: a stroke?
- **hx_asthma**: asthma?
- **hx_skinc**: skin cancer?
- **hx_otherc**: any other types of cancer?
- **hx_copd**: Chronic Obstructive Pulmonary Disease or COPD, emphysema or chronic bronchitis?
- **hx_arthr**: some form of arthritis, rheumatoid arthritis, gout, lupus, or fibromyalgia?

```
smart_16 %>% tabyl(comor)
```

comor	n	percent	valid_percent
0	224	0.211920530	0.221782178
1	315	0.298013245	0.311881188
2	228	0.215704825	0.225742574
3	130	0.122989593	0.128712871
4	72	0.068117313	0.071287129
5	29	0.027436140	0.028712871
6	9	0.008514664	0.008910891
7	3	0.002838221	0.002970297
NA	47	0.044465468	NA

14.1.1 Any missing values?

We have 1057 observations (rows) in the `smart_16` data set, of whom 860 have complete data on all variables.

```
dim(smart_16)

[1] 1057   10

n_case_complete(smart_16)

[1] 860
```

Which variables are missing?

```
miss_var_summary(smart_16)

# A tibble: 10 x 3
  variable n_miss pct_miss
  <chr>     <int>    <dbl>
1 activity      85     8.04
2 bmi          84     7.95
3 comor         47     4.45
4 smoke100     24     2.27
5 age_imp       12     1.14
6 hx_depress     3     0.284
7 SEQNO         0      0
8 mmsa          0      0
9 physhealth     0      0
10 bad_phys      0      0
```

Note that our outcomes (`physhealth` and the derived `bad_phys`) have no missing values here, by design. We will be performing multiple imputation to account appropriately for missingness in the predictors with missing values.

14.2 Obtaining a Simple Imputation with `mice`

The `mice` package provides several approaches we can use for imputation in building models of all kinds. Here, we'll use it just to obtain a single set of imputed results that we can apply to "complete" our data for the purposes of thinking about (a) transforming our outcome and (b) considering the addition of non-linear predictor terms.

```
# requires library(mice)

set.seed(432)

# create small data set including only variables to
```

```
# be used in building the imputation model

sm16 <- smart_16 %>%
  select(physhealth, activity, age_imp, bmi, comor,
         hx_depress, smoke100)

smart_16_mice1 <- mice(sm16, m = 1)

iter imp variable
1   1  activity  age_imp  bmi  comor  hx_depress  smoke100
2   1  activity  age_imp  bmi  comor  hx_depress  smoke100
3   1  activity  age_imp  bmi  comor  hx_depress  smoke100
4   1  activity  age_imp  bmi  comor  hx_depress  smoke100
5   1  activity  age_imp  bmi  comor  hx_depress  smoke100

smart_16_imp1 <- mice::complete(smart_16_mice1)

n_case_miss(smart_16_imp1)

[1] 0
```

And now we'll use this completed `smart_16_imp1` data set (the product of just a single imputation) to help us address the next two issues.

14.3 Linear Regression: Considering a Transformation of the Outcome

A plausible strategy here would be to try to identify an outcome transformation only after some accounting for missing predictor values, perhaps through a simple imputation approach. However, to keep things simple here, I'll just use the complete cases in this section.

Recall that our outcome here, `physhealth` can take the value 0, and is thus not strictly positive.

```
mosaic::favstats(~ physhealth, data = smart_16_imp1)
```

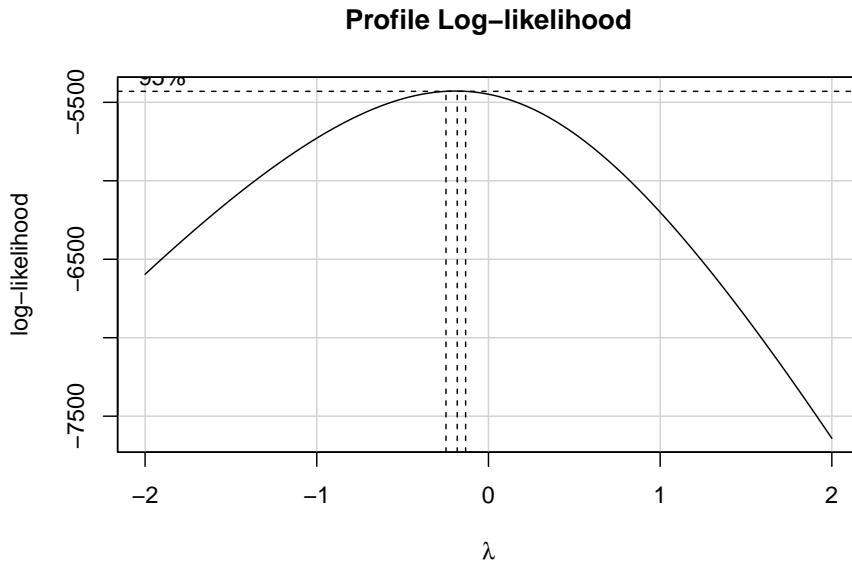
min	Q1	median	Q3	max	mean	sd	n	missing
0	0	2	20	30	9.227058	11.92676	1057	0

So, if we want to investigate a potential transformation with a Box-Cox plot, we'll have to add a small value to each `physhealth` value. We'll add 1, so that the range of potential values is now from 1-31.

```
# requires library(car)

smart_16_imp1 %$%
```

```
boxCox((physhealth + 1) ~ age_imp + comor + smoke100 +
       hx_depress + bmi + activity)
```



It looks like the logarithm is a reasonable transformation in this setting. So we'll create a new outcome, that is the natural logarithm of `(physhealth + 1)`, which we'll call `phys_tr` to remind us that a transformation is involved that we'll eventually need to back out of to make predictions. We'll build this new variable in both our original `smart_16` data set and in the simply imputed data set we're using for just these early stages.

```
smart_16_imp1 <- smart_16_imp1 %>%
  mutate(phys_tr = log(physhealth + 1))

smart_16 <- smart_16 %>%
  mutate(phys_tr = log(physhealth + 1))
```

So we have `phys_tr = log(physhealth + 1)`

- where we are referring above to the natural (base e logarithm).

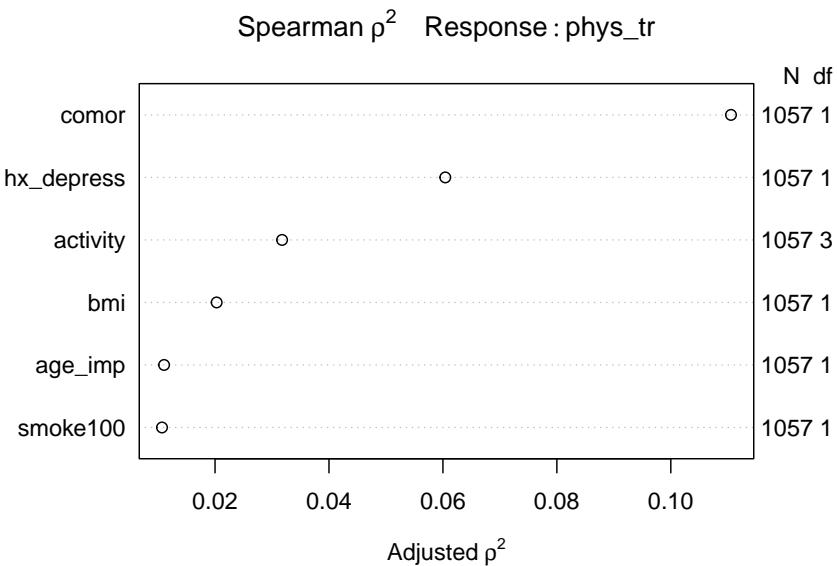
We can also specify our back-transformation to the original `physhealth` values from our new `phys_tr` as `physhealth = exp(phys_tr) - 1`.

14.4 Linear Regression: Considering Non-linearity in the Predictors

Consider the following Spearman ρ^2 plot.

```
# requires rms package
# (technically Hmisc, which is loaded by rms)

smart_16_imp1 %%%
  plot(spearman2(phys_tr ~ age_imp + comor + smoke100 +
    hx_depress + bmi + activity))
```



After our single imputation, we have the same N value in all rows of this plot, which is what we want to see. It appears that in considering potential non-linear terms, `comor` and `hx_depress` and perhaps `activity` are worthy of increased attention. I'll make a couple of arbitrary choices, to add a raw cubic polynomial to represent the `comor` information, and we'll add an interaction term between `hx_depress` and `activity`.

14.5 “Main Effects” Linear Regression with `lm` on the Complete Cases

Recall that we have 860 complete cases in our `smart_16` data, out of a total of 1057 observations in total. A model using only the complete cases should thus drop the remaining 197 subjects. Let’s see if a main effects only model for our newly transformed `phys_tr` outcome does in fact do this.

```
m_1cc <- smart_16 %%
  lm(phys_tr ~ age_imp + comor + smoke100 +
    hx_depress + bmi + activity)

summary(m_1cc)

Call:
lm(formula = phys_tr ~ age_imp + comor + smoke100 + hx_depress +
    bmi + activity)

Residuals:
    Min      1Q  Median      3Q      Max 
-3.0801 -1.0389 -0.2918  1.1029  2.8478 

Coefficients:
            Estimate Std. Error t value Pr(>|t|)    
(Intercept) 0.581959  0.370847   1.569  0.11696  
age_imp     -0.007043  0.003813  -1.847  0.06511 .  
comor       0.301773  0.033105   9.116 < 2e-16 *** 
smoke100    0.099038  0.090280   1.097  0.27295  
hx_depress  0.471949  0.104232   4.528 6.81e-06 *** 
bmi         0.016375  0.006295   2.601  0.00945 **  
activityActive -0.229927  0.154912  -1.484  0.13812  
activityInsufficiently_Active -0.116998  0.139440  -0.839  0.40168  
activityInactive  0.256118  0.115266   2.222  0.02655 *  
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 1.303 on 851 degrees of freedom
(197 observations deleted due to missingness)
Multiple R-squared:  0.1806,    Adjusted R-squared:  0.1729 
F-statistic: 23.45 on 8 and 851 DF,  p-value: < 2.2e-16
```

Note that the appropriate number of observations are listed as “deleted due to missingness.”

14.5.1 Quality of Fit Statistics

```
glance(m_1cc) %>%
  select(r.squared, adj.r.squared, sigma, AIC, BIC) %>%
  kable(digits = c(3, 3, 2, 1, 1))
```

r.squared	adj.r.squared	sigma	AIC	BIC
0.181	0.173	1.3	2906.3	2953.8

14.5.2 Interpreting Effect Sizes

```
tidy(m_1cc, conf.int = TRUE) %>%
  select(term, estimate, std.error, conf.low, conf.high) %>%
  kable(digits = 3)
```

term	estimate	std.error	conf.low	conf.high
(Intercept)	0.582	0.371	-0.146	1.310
age_imp	-0.007	0.004	-0.015	0.000
comor	0.302	0.033	0.237	0.367
smoke100	0.099	0.090	-0.078	0.276
hx_depress	0.472	0.104	0.267	0.677
bmi	0.016	0.006	0.004	0.029
activityActive	-0.230	0.155	-0.534	0.074
activityInsufficiently_Active	-0.117	0.139	-0.391	0.157
activityInactive	0.256	0.115	0.030	0.482

We'll interpret three of the predictors here to demonstrate ideas: `comor`, `hx_depress` and `activity`.

- If we have two subjects with the same values of `age_imp`, `smoke100`, `hx_depress`, `bmi`, and `activity`, but Harry has a `comor` score that is one point higher than Sally's, then the model predicts that Harry's transformed outcome (specifically the natural logarithm of (his `physhealth` days + 1)) will be 0.302 higher than Sally's, with a 95% confidence interval around that estimate ranging from (`round(a$conf.low,3)`, `round(a$conf.high,3)`).
- If we have two subjects with the same values of `age_imp`, `comor`, `smoke100`, `bmi`, and `activity`, but Harry has a history of depression (`hx_depress = 1`) while Sally does not have such a history (so Sally's `hx_depress = 0`), then the model predicts that Harry's transformed outcome (specifically the natural logarithm of (his `physhealth` days + 1)) will be 0.472 higher than Sally's, with a 95% confidence interval around that estimate ranging from (`round(a$conf.low,3)`, `round(a$conf.high,3)`).
- The `activity` variable has four categories as indicated in the table below. The model uses the "Highly_Active" category as the reference group.

14.5. “MAIN EFFECTS” LINEAR REGRESSION WITH LM ON THE COMPLETE CASES375

```
smart_16_imp1 %>% tabyl(activity)
```

	activity	n	percent
Highly_Active	Highly_Active	252	0.2384106
	Active	135	0.1277200
Insufficiently_Active	Insufficiently_Active	193	0.1825922
	Inactive	477	0.4512772

- From the tidied set of coefficients, we can describe the `activity` effects as follows.
 - If Sally is “Highly Active” and Harry is “Active” but they otherwise have the same values of all predictors, then our prediction is that Harry’s transformed outcome (specifically the natural logarithm of (his `physhealth` days + 1)) will be 0.23 lower than Sally’s, with a 95% confidence interval around that estimate ranging from `(round(a$conf.low,3), round(a$conf.high,3))`.
 - If instead Harry is “Insufficiently Active” but nothing else changes, then our prediction is that Harry’s transformed outcome will be 0.117 lower than Sally’s, with a 95% confidence interval around that estimate ranging from `(round(a2$conf.low,3), round(a2$conf.high,3))`.
 - If instead Harry is “Inactive” but nothing else changes, then our prediction is that Harry’s transformed outcome will be -0.117 higher than Sally’s, with a 95% confidence interval around that estimate ranging from `(round(a2$conf.low,3), round(a2$conf.high,3))`.

14.5.3 Making Predictions with the Model

Let’s describe two subjects, and use this model (and the ones that follow) to predict their `physhealth` values.

- Sheena is age 50, has 2 comorbidities, has smoked 100 cigarettes in her life, has no history of depression, a BMI of 25, and is Highly Active.
- Jacob is age 65, has 4 comorbidities, has never smoked, has a history of depression, a BMI of 32 and is Inactive.

We’ll first build predictions for Sheena and Jacob (with 95% prediction intervals) for `phys_tr`.

```
new2 <- tibble(  
  name = c("Sheena", "Jacob"),  
  age_imp = c(50, 65),  
  comor = c(2, 4),  
  smoke100 = c(1, 0),  
  hx_depress = c(0, 1),  
  bmi = c(25, 32),  
  activity = c("Highly_Active", "Inactive")
```

```
)
preds_m_1cc <- predict(m_1cc, newdata = new2,
                        interval = "prediction")

preds_m_1cc

      fit      lwr      upr
1 1.341778 -1.22937 3.912925
2 2.583336  0.01399 5.152681
```

The model makes predictions for our transformed outcome, `phys_tr`. Now, we need to back-transform the predictions and the confidence intervals to build predictions for `physhealth`.

```
preds_m_1cc <- preds_m_1cc %>%
 tbl_df() %>%
  mutate(names = c("Sheena", "Jacob"),
         pred_physhealth = exp(fit) - 1,
         conf_low = exp(lwr) - 1,
         conf_high = exp(upr) - 1) %>%
  select(names, pred_physhealth, conf_low, conf_high,
         everything())

preds_m_1cc %>% kable(digits = 3)
```

names	pred_physhealth	conf_low	conf_high	fit	lwr	upr
Sheena	2.826	-0.708	49.045	1.342	-1.229	3.913
Jacob	12.241	0.014	171.894	2.583	0.014	5.153

14.6 “Augmented” Linear Regression with `lm` on the Complete Cases

Now, we’ll add the non-linear terms we discussed earlier. We’ll add a (raw) cubic polynomial to represent the `comor` information, and we’ll add an interaction term between `hx_depress` and `activity`.

```
# requires rms package (and co-loading Hmisc)

m_2cc <- smart_16 %%%
  lm(phys_tr ~ age_imp + pol(comor, 3) + smoke100 +
     bmi + hx_depress*activity)

summary(m_2cc)
```

Call:

14.6. "AUGMENTED" LINEAR REGRESSION WITH LM ON THE COMPLETE CASES 377

```
lm(formula = phys_tr ~ age_imp + pol(comor, 3) + smoke100 + bmi +
hx_depress * activity)
```

Residuals:

Min	1Q	Median	3Q	Max
-2.907	-1.063	-0.267	1.143	2.924

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	0.514823	0.376203	1.368	0.17153
age_imp	-0.008100	0.003865	-2.096	0.03640
pol(comor, 3)comor	0.634274	0.160630	3.949	8.51e-05
pol(comor, 3)comor^2	-0.130626	0.073525	-1.777	0.07599
pol(comor, 3)comor^3	0.012508	0.008977	1.393	0.16386
smoke100	0.089345	0.090336	0.989	0.32294
bmi	0.015203	0.006315	2.408	0.01627
hx_depress	0.647054	0.229696	2.817	0.00496
activityActive	-0.202196	0.172300	-1.174	0.24092
activityInsufficiently_Active	-0.005815	0.166221	-0.035	0.97210
activityInactive	0.290380	0.132198	2.197	0.02832
hx_depress:activityActive	-0.124836	0.395415	-0.316	0.75230
hx_depress:activityInsufficiently_Active	-0.376355	0.310160	-1.213	0.22531
hx_depress:activityInactive	-0.172952	0.267427	-0.647	0.51798

(Intercept)	
age_imp	*
pol(comor, 3)comor	***
pol(comor, 3)comor^2	.
pol(comor, 3)comor^3	
smoke100	
bmi	*
hx_depress	**
activityActive	
activityInsufficiently_Active	
activityInactive	*
hx_depress:activityActive	
hx_depress:activityInsufficiently_Active	
hx_depress:activityInactive	

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 1.301 on 846 degrees of freedom

(197 observations deleted due to missingness)

Multiple R-squared: 0.187, Adjusted R-squared: 0.1745

F-statistic: 14.97 on 13 and 846 DF, p-value: < 2.2e-16

Note again that the appropriate number of observations are listed as “deleted due to missingness.”

14.6.1 Quality of Fit Statistics

```
glance(m_2cc) %>%
  select(r.squared, adj.r.squared, sigma, AIC, BIC) %>%
  kable(digits = c(3, 3, 2, 1, 1))
```

r.squared	adj.r.squared	sigma	AIC	BIC
0.187	0.175	1.3	2909.5	2980.9

14.6.2 ANOVA assessing the impact of the non-linear terms

```
anova(m_1cc, m_2cc)
```

Analysis of Variance Table

```
Model 1: phys_tr ~ age_imp + comor + smoke100 + hx_depress + bmi + activity
Model 2: phys_tr ~ age_imp + pol(comor, 3) + smoke100 + bmi + hx_depress *
          activity
Res.Df   RSS Df Sum of Sq    F Pr(>F)
1     851 1444.0
2     846 1432.8  5     11.265 1.3303 0.249
```

The difference between the models doesn't meet the standard for statistical detectability at our usual α levels.

14.6.3 Interpreting Effect Sizes

```
tidy(m_2cc, conf.int = TRUE) %>%
  select(term, estimate, std.error, conf.low, conf.high) %>%
  kable(digits = 3)
```

14.6. “AUGMENTED” LINEAR REGRESSION WITH LM ON THE COMPLETE CASES 379

term	estimate	std.error	conf.low	conf.high
(Intercept)	0.515	0.376	-0.224	1.253
age_imp	-0.008	0.004	-0.016	-0.001
pol(comor, 3)comor	0.634	0.161	0.319	0.950
pol(comor, 3)comor^2	-0.131	0.074	-0.275	0.014
pol(comor, 3)comor^3	0.013	0.009	-0.005	0.030
smoke100	0.089	0.090	-0.088	0.267
bmi	0.015	0.006	0.003	0.028
hx_depress	0.647	0.230	0.196	1.098
activityActive	-0.202	0.172	-0.540	0.136
activityInsufficiently_Active	-0.006	0.166	-0.332	0.320
activityInactive	0.290	0.132	0.031	0.550
hx_depress:activityActive	-0.125	0.395	-0.901	0.651
hx_depress:activityInsufficiently_Active	-0.376	0.310	-0.985	0.232
hx_depress:activityInactive	-0.173	0.267	-0.698	0.352

Let's focus first on interpreting the interaction terms between `hx_depress` and `activity`.

Assume first that we have a set of subjects with the same values of `age_imp`, `smoke100`, `bmi`, and `comor`.

- Arnold has `hx_depress` = 1 and is Inactive
- Betty has `hx_depress` = 1 and is Insufficiently Active
- Carlos has `hx_depress` = 1 and is Active
- Debbie has `hx_depress` = 1 and is Highly Active
- Eamon has `hx_depress` = 0 and is Inactive
- Florence has `hx_depress` = 0 and is Insufficiently Active
- Garry has `hx_depress` = 0 and is Active
- Harry has `hx_depress` = 0 and is Highly Active

So the model, essentially can be used to compare each of the first seven people on that list to Harry (who has the reference levels of both `hx_depress` and `activity`.) Let's compare Arnold to Harry.

For instance, as compared to Harry, Arnold is expected to have a transformed outcome (specifically the natural logarithm of (his `physhealth` days + 1)) that is:

- 0.647 higher because Arnold's `hx_depress` = 1, and
- 0.29 higher still because Arnold's `activity` is “Inactive,” and
- 0.173 lower because of the combination (see the ‘`hx_depress:activityInactive`’ row)

So, in total, we expect Arnold's transformed outcome to be $0.647 + 0.29 + (-0.173)$, or 0.764 higher than Harry's.

If we want to compare Arnold to, for instance, Betty, we first calculate Betty's difference from Harry, and then compare the two differences.

As compared to Harry, Betty is expected to have a transformed outcome (specifically the natural logarithm of (her `physhealth` days + 1)) that is:

- 0.647 higher because Betty's `hx_depress` = 1, and
- 0.006 lower still because Betty's `activity` is "Insufficiently Active," and
- 0.376 lower because of the combination (see the 'hx_depress:activityInsufficiently_Active' row)

So, in total, we expect Betty's transformed outcome to be $0.647 + (-0.006) + (-0.376)$, or 0.265 higher than Harry's.

And thus we can compare Betty and Arnold directly.

- Arnold is predicted to have an outcome that is 0.764 higher than Harry's.
- Betty is predicted to have an outcome that is 0.265 higher than Harry's.
- And so Arnold's predicted outcome (`phys_tr`) is 0.499 larger than Betty's.

Now, suppose we want to look at our cubic polynomial in `comor`.

- Suppose Harry and Sally have the same values for all other predictors in the model, but Harry has 1 comorbidity where Sally has none. Then the three terms in the model related to `comor` will be 1 for Harry and 0 for Sally, and the interpretation becomes pretty straightforward.
- But suppose instead that nothing has changed except Harry has 2 comorbidities and Sally has just 1. The size of the impact of this Harry - Sally difference is far larger in this situation, because the `comor` variable enters the model in a non-linear way. This is an area where fitting the model using `ols` can be helpful because of the ability to generate plots (of effects, nomograms, etc.) that can show this non-linearity in a clear way.

Suppose for instance, that Harry and Sally share the following values for the other predictors: each is age 40, has never smoked, has no history of depression, a BMI of 30 and is Highly Active.

- Now, if Harry has 1 comorbidity and Sally has none, the predicted `phys_tr` values for Harry and Sally are as indicated below.

```
hands1 <- tibble(
  name = c("Harry", "Sally"),
  age_imp = c(40, 40),
  comor = c(1, 0),
  smoke100 = c(0, 0),
  hx_depress = c(0, 0),
  bmi = c(30, 30),
  activity = c("Highly_Active", "Highly_Active")
)

predict(m_2cc, newdata = hands1)
```

```
1.1630840 0.6469282
```

But if Harry has 2 comorbidities and Sally 1, the predictions are:

```
hands2 <- tibble(
  name = c("Harry", "Sally"),
  age_imp = c(40, 40),
  comor = c(2, 1), # only thing that changes
  smoke100 = c(0, 0),
  hx_depress = c(0, 0),
  bmi = c(30, 30),
  activity = c("Highly_Active", "Highly_Active")
)

predict(m_2cc, newdata = hands2)
```

```
1          2
1.493035 1.163084
```

Note that the difference in predictions between Harry and Sally is much smaller now than it was previously.

14.6.4 Making Predictions with the Model

As before, we'll use the new model to predict `physhealth` values for Sheena and Jacob.

- Sheena is age 50, has 2 comorbidities, has smoked 100 cigarettes in her life, has no history of depression, a BMI of 25, and is Highly Active.
- Jacob is age 65, has 4 comorbidities, has never smoked, has a history of depression, a BMI of 32 and is Inactive.

We'll first build predictions for Sheena and Jacob (with 95% prediction intervals) for `phys_tr`.

```
new2 <- tibble(
  name = c("Sheena", "Jacob"),
  age_imp = c(50, 65),
  comor = c(2, 4),
  smoke100 = c(1, 0),
  hx_depress = c(0, 1),
  bmi = c(25, 32),
  activity = c("Highly_Active", "Inactive")
)

preds_m_2cc <- predict(m_2cc, newdata = new2,
                        interval = "prediction")
```

```
preds_m_2cc
```

	fit	lwr	upr
1	1.425362	-1.14707613	3.997801
2	2.486907	-0.08635658	5.060171

Now, we need to back-transform the predictions and the confidence intervals that describe `phys_tr` to build predictions for `physhealth`.

```
preds_m_2cc <- preds_m_2cc %>%
 tbl_df() %>%
  mutate(names = c("Sheena", "Jacob"),
        pred_physhealth = exp(fit) - 1,
        conf_low = exp(lwr) - 1,
        conf_high = exp(upr) - 1) %>%
  select(names, pred_physhealth, conf_low, conf_high,
         everything())

preds_m_2cc %>% kable(digits = 3)
```

names	pred_physhealth	conf_low	conf_high	fit	lwr	upr
Sheena	3.159	-0.682	53.478	1.425	-1.147	3.998
Jacob	11.024	-0.083	156.617	2.487	-0.086	5.060

14.7 Using `mice` to perform Multiple Imputation

Let's focus on the main effects model, and look at the impact of performing multiple imputation to account for the missing data. Recall that in our `smart_16` data, the most “missingness” is shown in the `activity` variable, which is still missing less than 10% of the time. So we'll try a set of 10 imputations, using the default settings in the `mice` package.

```
# requires library(mice)

set.seed(432)

# create small data set including only variables to
# be used in building the imputation model

sm16 <- smart_16 %>%
  select(physhealth, phys_tr, activity, age_imp, bmi, comor,
         hx_depress, smoke100)

smart_16_mice10 <- mice(sm16, m = 10)
```

```

iter imp variable
1 1 activity age_imp bmi comor hx_depress smoke100
1 2 activity age_imp bmi comor hx_depress smoke100
1 3 activity age_imp bmi comor hx_depress smoke100
1 4 activity age_imp bmi comor hx_depress smoke100
1 5 activity age_imp bmi comor hx_depress smoke100
1 6 activity age_imp bmi comor hx_depress smoke100
1 7 activity age_imp bmi comor hx_depress smoke100
1 8 activity age_imp bmi comor hx_depress smoke100
1 9 activity age_imp bmi comor hx_depress smoke100
1 10 activity age_imp bmi comor hx_depress smoke100
2 1 activity age_imp bmi comor hx_depress smoke100
2 2 activity age_imp bmi comor hx_depress smoke100
2 3 activity age_imp bmi comor hx_depress smoke100
2 4 activity age_imp bmi comor hx_depress smoke100
2 5 activity age_imp bmi comor hx_depress smoke100
2 6 activity age_imp bmi comor hx_depress smoke100
2 7 activity age_imp bmi comor hx_depress smoke100
2 8 activity age_imp bmi comor hx_depress smoke100
2 9 activity age_imp bmi comor hx_depress smoke100
2 10 activity age_imp bmi comor hx_depress smoke100
3 1 activity age_imp bmi comor hx_depress smoke100
3 2 activity age_imp bmi comor hx_depress smoke100
3 3 activity age_imp bmi comor hx_depress smoke100
3 4 activity age_imp bmi comor hx_depress smoke100
3 5 activity age_imp bmi comor hx_depress smoke100
3 6 activity age_imp bmi comor hx_depress smoke100
3 7 activity age_imp bmi comor hx_depress smoke100
3 8 activity age_imp bmi comor hx_depress smoke100
3 9 activity age_imp bmi comor hx_depress smoke100
3 10 activity age_imp bmi comor hx_depress smoke100
4 1 activity age_imp bmi comor hx_depress smoke100
4 2 activity age_imp bmi comor hx_depress smoke100
4 3 activity age_imp bmi comor hx_depress smoke100
4 4 activity age_imp bmi comor hx_depress smoke100
4 5 activity age_imp bmi comor hx_depress smoke100
4 6 activity age_imp bmi comor hx_depress smoke100
4 7 activity age_imp bmi comor hx_depress smoke100
4 8 activity age_imp bmi comor hx_depress smoke100
4 9 activity age_imp bmi comor hx_depress smoke100
4 10 activity age_imp bmi comor hx_depress smoke100
5 1 activity age_imp bmi comor hx_depress smoke100
5 2 activity age_imp bmi comor hx_depress smoke100
5 3 activity age_imp bmi comor hx_depress smoke100
5 4 activity age_imp bmi comor hx_depress smoke100
5 5 activity age_imp bmi comor hx_depress smoke100

```

```

5   6 activity age_imp bmi comor hx_depress smoke100
5   7 activity age_imp bmi comor hx_depress smoke100
5   8 activity age_imp bmi comor hx_depress smoke100
5   9 activity age_imp bmi comor hx_depress smoke100
5  10 activity age_imp bmi comor hx_depress smoke100

summary(smart_16_mice10)

Class: mids
Number of multiple imputations: 10
Imputation methods:
physhealth phys_tr activity age_imp      bmi      comor hx_depress
      ""        "" "polyreg"     "pmm"     "pmm"     "pmm"     "pmm"
smoke100
      "pmm"
PredictorMatrix:
          physhealth phys_tr activity age_imp bmi comor hx_depress smoke100
physhealth          0       1       1       1       1       1       1       1
phys_tr             1       0       1       1       1       1       1       1
activity            1       1       0       1       1       1       1       1
age_imp              1       1       1       0       1       1       1       1
bmi                  1       1       1       1       0       1       1       1
comor                1       1       1       1       1       0       1       1

```

14.8 Running the Linear Regression in lm with Multiple Imputation

Next, we'll run the linear model (main effects) on each of the 10 imputed data sets.

```

m10_mods <-
  with(smart_16_mice10, lm(phys_tr ~ age_imp + comor +
                            smoke100 + hx_depress +
                            bmi + activity))

summary(m10_mods)

# A tibble: 90 x 6
#>   term      estimate std.error statistic p.value    nobs
#>   <chr>     <dbl>     <dbl>     <dbl>    <dbl>    <int>
#> 1 (Intercept)  0.317     0.326     0.971 3.32e-1  1057
#> 2 age_imp     -0.00489   0.00334    -1.47  1.43e-1  1057
#> 3 comor        0.313     0.0295    10.6   4.72e-25 1057
#> 4 smoke100     0.135     0.0799    1.69   9.22e-2  1057
#> 5 hx_depress   0.500     0.0929    5.38   9.14e-8  1057

```

```

6 bmi                      0.0187   0.00564    3.31 9.64e- 4 1057
7 activityActive           -0.202    0.138     -1.46 1.44e- 1 1057
8 activityInsufficiently_Active -0.0695  0.124     -0.561 5.75e- 1 1057
9 activityInactive          0.262     0.103      2.54 1.11e- 2 1057
10 (Intercept)             0.363     0.332      1.10 2.74e- 1 1057
# ... with 80 more rows

```

Then, we'll pool results across the 10 imputations

```

m10_pool <- pool(m10_mods)
summary(m10_pool, conf.int = TRUE) %>%
  select(-statistic, -df) %>%
  kable(digits = 3)

```

term	estimate	std.error	p.value	2.5 %	97.5 %
(Intercept)	0.444	0.342	0.194	-0.227	1.114
age_imp	-0.005	0.003	0.128	-0.012	0.002
comor	0.309	0.031	0.000	0.249	0.369
smoke100	0.114	0.083	0.171	-0.049	0.278
hx_depress	0.512	0.094	0.000	0.327	0.696
bmi	0.016	0.006	0.009	0.004	0.027
activityActive	-0.204	0.140	0.146	-0.479	0.071
activityInsufficiently_Active	-0.044	0.129	0.735	-0.298	0.210
activityInactive	0.260	0.106	0.014	0.052	0.469

And we can compare these results to the complete case analysis we completed earlier.

```

tidy(m_1cc, conf.int = TRUE) %>%
  select(term, estimate, std.error, p.value, conf.low, conf.high) %>%
  kable(digits = 3)

```

term	estimate	std.error	p.value	conf.low	conf.high
(Intercept)	0.582	0.371	0.117	-0.146	1.310
age_imp	-0.007	0.004	0.065	-0.015	0.000
comor	0.302	0.033	0.000	0.237	0.367
smoke100	0.099	0.090	0.273	-0.078	0.276
hx_depress	0.472	0.104	0.000	0.267	0.677
bmi	0.016	0.006	0.009	0.004	0.029
activityActive	-0.230	0.155	0.138	-0.534	0.074
activityInsufficiently_Active	-0.117	0.139	0.402	-0.391	0.157
activityInactive	0.256	0.115	0.027	0.030	0.482

Note that there are some sizeable differences here, although nothing enormous.

If we want the pooled R^2 or pooled adjusted R^2 after imputation, R will provide it (and a 95% confidence interval around the estimate) with ...

```

pool.r.squared(m10_mods)

      est      lo 95      hi 95      fmi
R^2 0.1912561 0.1482819 0.2369623 0.08061427

pool.r.squared(m10_mods, adjusted = TRUE)

      est      lo 95      hi 95      fmi
adj R^2 0.1850807 0.1425132 0.2305277 0.08312639

```

We can see the fraction of missing information about each coefficient due to non-response (*fmi*) and other details with the following code...

```

m10_pool

Class: mipo    m = 10
      term   m   estimate      ubar      b
1 (Intercept) 10  0.44377168 1.078194e-01 8.002900e-03
2 age_imp     10 -0.00522810 1.123668e-05 4.824290e-07
3 comor      10  0.30871888 8.801039e-04 5.466082e-05
4 smoke100   10  0.11415718 6.474388e-03 4.130673e-04
5 hx_depress 10  0.51155722 8.669413e-03 1.582684e-04
6 bmi        10  0.01576150 3.182381e-05 3.425191e-06
7 activityActive 10 -0.20412627 1.914250e-02 4.851040e-04
8 activityInsufficiently_Active 10 -0.04383739 1.565925e-02 9.855183e-04
9 activityInactive 10  0.26046070 1.069627e-02 5.023887e-04
      t dfcom      df      riv      lambda      fmi
1 1.166226e-01 1048 599.8172 0.08164758 0.07548446 0.07855177
2 1.176735e-05 1048 814.9042 0.04722675 0.04509697 0.04743197
3 9.402308e-04 1048 677.6361 0.06831796 0.06394909 0.06669961
4 6.928762e-03 1048 666.2487 0.07018023 0.06557795 0.06837041
5 8.843508e-03 1048 982.0512 0.02008154 0.01968621 0.02167660
6 3.559152e-05 1048 432.0874 0.11839279 0.10585976 0.10996992
7 1.967611e-02 1048 939.5064 0.02787591 0.02711992 0.02918437
8 1.674332e-02 1048 672.0473 0.06922874 0.06474643 0.06751736
9 1.124890e-02 1048 785.1905 0.05166545 0.04912727 0.05154007

```

14.9 Fit the Multiple Imputation Model with `aregImpute`

Here, we'll use `aregImpute` to deal with missing values through multiple imputation, and use the `ols` function in the `rms` package to fit the model.

The first step is to fit the multiple imputation model. We'll use `n.impute = 10` imputations, with `B = 10` bootstrap samples for the predictive mean matching, and fit both linear models and models with restricted cubic splines with 3 knots

(nk = c(0, 3)) allowing the target variable to have a non-linear transformation when nk is 3, via tlinear = FALSE.

```
set.seed(43201602)
dd <- datadist(smart_16)
options(datadist = "dd")

fit16_imp <-
  aregImpute(~ phys_tr + age_imp + comor + smoke100 +
             hx_depress + bmi + activity,
             nk = c(0, 3), tlinear = FALSE,
             data = smart_16, B = 10, n.impute = 10)
```

Iteration 1 Iteration 2 Iteration 3 Iteration 4 Iteration 5 Iteration 6 Iteration 7 Iteration 8 1

Here are the results of that imputation model.

```
fit16_imp
```

Multiple Imputation using Bootstrap and PMM

```
aregImpute(formula = ~phys_tr + age_imp + comor + smoke100 +
            hx_depress + bmi + activity, data = smart_16, n.impute = 10,
            nk = c(0, 3), tlinear = FALSE, B = 10)
```

n: 1057 p: 7 Imputations: 10 nk: 0

Number of NAs:

phys_tr	age_imp	comor	smoke100	hx_depress	bmi	activity
0	12	47	24	3	84	85

type d.f.

phys_tr	s	1
age_imp	s	1
comor	s	1
smoke100	l	1
hx_depress	l	1
bmi	s	1
activity	c	3

R-squares for Predicting Non-Missing Values for Each Variable
Using Last Imputations of Predictors

age_imp	comor	smoke100	hx_depress	bmi	activity
0.224	0.206	0.059	0.167	0.169	0.057

Resampling results for determining the complexity of imputation models

Variable being imputed: age_imp

	nk=0	nk=3
Bootstrap bias-corrected R ²	0.186	0.215
10-fold cross-validated R ²	0.211	0.215
Bootstrap bias-corrected mean error	9.108	10.894
10-fold cross-validated mean error	65.169	10.919
Bootstrap bias-corrected median error	7.290	8.784
10-fold cross-validated median error	66.006	8.613
 Variable being imputed: comor		
	nk=0	nk=3
Bootstrap bias-corrected R ²	0.183	0.182
10-fold cross-validated R ²	0.184	0.193
Bootstrap bias-corrected mean error	0.987	1.184
10-fold cross-validated mean error	1.759	1.171
Bootstrap bias-corrected median error	0.828	0.910
10-fold cross-validated median error	1.574	0.892
 Variable being imputed: smoke100		
	nk=0	nk=3
Bootstrap bias-corrected R ²	0.0224	0.0187
10-fold cross-validated R ²	0.0358	0.0217
Bootstrap bias-corrected mean error	0.4853	0.4866
10-fold cross-validated mean error	0.9462	0.9561
Bootstrap bias-corrected median error	0.4788	0.4772
10-fold cross-validated median error	0.8479	0.8706
 Variable being imputed: hx_depress		
	nk=0	nk=3
Bootstrap bias-corrected R ²	0.157	0.138
10-fold cross-validated R ²	0.147	0.148
Bootstrap bias-corrected mean error	0.355	0.360
10-fold cross-validated mean error	0.801	0.783
Bootstrap bias-corrected median error	0.333	0.337
10-fold cross-validated median error	0.711	0.673
 Variable being imputed: bmi		
	nk=0	nk=3
Bootstrap bias-corrected R ²	0.125	0.122
10-fold cross-validated R ²	0.134	0.133
Bootstrap bias-corrected mean error	5.221	6.822
10-fold cross-validated mean error	32.458	6.884
Bootstrap bias-corrected median error	4.178	5.782
10-fold cross-validated median error	31.330	5.932
 Variable being imputed: activity		
	nk=0	nk=3

```

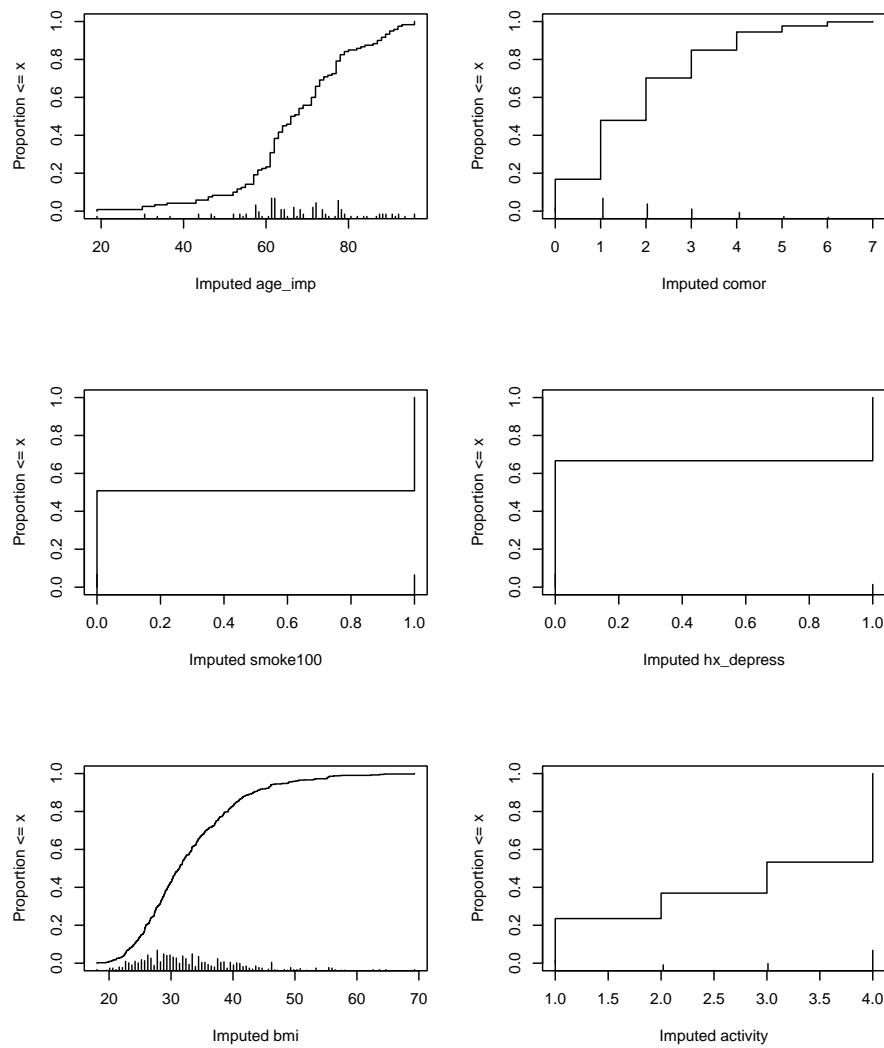
Bootstrap bias-corrected R^2          0.0312 0.0275
10-fold cross-validated   R^2          0.0450 0.0402
Bootstrap bias-corrected mean |error| 1.8774 1.8884
10-fold cross-validated  mean |error| 1.1121 1.1043
Bootstrap bias-corrected median |error| 2.0000 2.0000
10-fold cross-validated  median |error| 1.0000 1.0000

```

```

par(mfrow = c(3,2))
plot(fit16_imp)

```



```
par(mfrow = c(1,1))
```

The plot helps us see where the imputations are happening.

14.10 Fit Linear Regression using `ols` and `fit.mult.impute`

```
m16_imp <-
  fit.mult.impute(phys_tr ~ age_imp + comor + smoke100 +
    hx_depress + bmi + activity,
    fitter = ols, xtrans = fit16_imp,
    data = smart_16, x = TRUE, y = TRUE)
```

Variance Inflation Factors Due to Imputation:

Intercept	age_imp
1.03	1.01
comor	smoke100
1.03	1.06
hx_depress	bmi
1.02	1.06
activity=Active	activity=Insufficiently_Active
1.19	1.14
activity=Inactive	
1.23	

Rate of Missing Information:

Intercept	age_imp
0.03	0.01
comor	smoke100
0.03	0.06
hx_depress	bmi
0.02	0.06
activity=Active	activity=Insufficiently_Active
0.16	0.13
activity=Inactive	
0.19	

d.f. for t-distribution for Tests of Single Coefficients:

Intercept	age_imp
8176.67	45410.80

comor		smoke100	
13030.27		2670.64	
hx_depress		bmi	
28199.30		2935.89	
activity=Active	activity=Insufficiently_Active		
354.62		571.56	
activity=Inactive			
258.42			

The following fit components were averaged over the 10 model fits:

```
fitted.values stats linear.predictors
```

14.10.1 Summaries and Coefficients

Here are the results:

```
m16_imp
```

Linear Regression Model

```
fit.mult.impute(formula = phys_tr ~ age_imp + comor + smoke100 +
  hx_depress + bmi + activity, fitter = ols, xtrans = fit16_imp,
  data = smart_16, x = TRUE, y = TRUE)
```

	Model Likelihood	Discrimination	
	Ratio Test	Indexes	
Obs	1057	LR chi2 219.94	R2 0.188
sigma1.2881	d.f.	8	R2 adj 0.182
d.f.	1048	Pr(> chi2) 0.0000	g 0.687

Residuals

	Min	1Q	Median	3Q	Max
	-3.0621	-1.0327	-0.2878	1.1104	2.8018

	Coef	S.E.	t	Pr(> t)
Intercept	0.4052	0.3352	1.21	0.2271
age_imp	-0.0049	0.0034	-1.46	0.1437
comor	0.3078	0.0302	10.20	<0.0001
smoke100	0.1259	0.0830	1.52	0.1296
hx_depress	0.5120	0.0940	5.45	<0.0001
bmi	0.0164	0.0058	2.83	0.0048
activity=Active	-0.1773	0.1513	-1.17	0.2416
activity=Insufficiently_Active	-0.0396	0.1342	-0.30	0.7680

```
activity=Inactive          0.2401 0.1144  2.10 0.0360
```

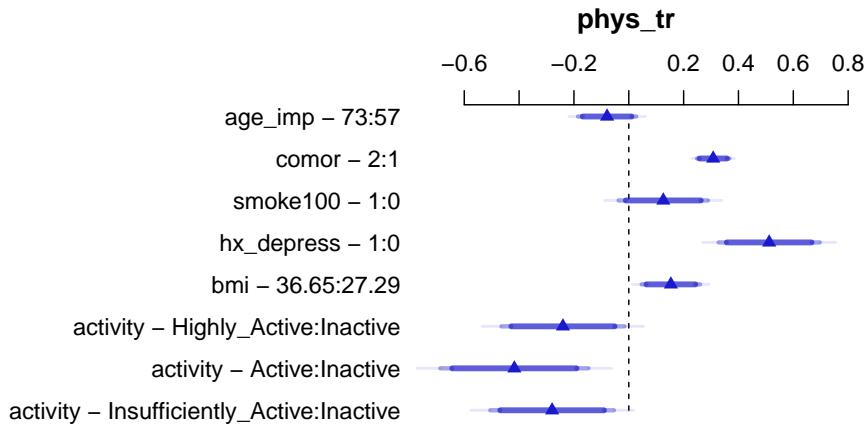
14.10.2 Effect Sizes

We can plot and summarize the effect sizes using the usual `ols` tools:

```
summary(m16_imp)
```

Effects	Response : phys_tr				
Factor	Low	High	Diff.	Effect	S.E.
age_imp	57.00	73.00	16.00	-0.079163	0.054107
comor	1.00	2.00	1.00	0.307790	0.030171
smoke100	0.00	1.00	1.00	0.125890	0.083000
hx_depress	0.00	1.00	1.00	0.511980	0.094007
bmi	27.29	36.65	9.36	0.153530	0.054322
activity - Highly_Active:Inactive	4.00	1.00	NA	-0.240070	0.114350
activity - Active:Inactive	4.00	2.00	NA	-0.417320	0.137640
activity - Insufficiently_Active:Inactive	4.00	3.00	NA	-0.279650	0.115000
Lower 0.95	Upper 0.95				
-0.185330	0.027008				
0.248590	0.366990				
-0.036973	0.288760				
0.327520	0.696450				
0.046932	0.260120				
-0.464450	-0.015686				
-0.687400	-0.147250				
-0.505310	-0.054002				

```
plot(summary(m16_imp))
```



14.10.3 Making Predictions with this Model

Once again, let's make predictions for our two subjects, and use this model (and the ones that follow) to predict their `physhealth` values.

- Sheena is age 50, has 2 comorbidities, has smoked 100 cigarettes in her life, has no history of depression, a BMI of 25, and is Highly Active.
- Jacob is age 65, has 4 comorbidities, has never smoked, has a history of depression, a BMI of 32 and is Inactive.

```
new2 <- tibble(
  name = c("Sheena", "Jacob"),
  age_imp = c(50, 65),
  comor = c(2, 4),
  smoke100 = c(1, 0),
  hx_depress = c(0, 1),
  bmi = c(25, 32),
  activity = c("Highly_Active", "Inactive")
)

preds_m_16imp <- predict(m16_imp,
                         newdata = data.frame(new2))

preds_m_16imp
```

```

1      2
1.309306 2.591649

preds_m_16imp <- preds_m_16imp %>%
 tbl_df() %>%
  mutate(names = c("Sheena", "Jacob"),
         pred_physhealth = exp(value) - 1) %>%
  select(names, pred_physhealth)

preds_m_16imp %>% kable(digits = 3)

```

names	pred_physhealth
Sheena	2.704
Jacob	12.352

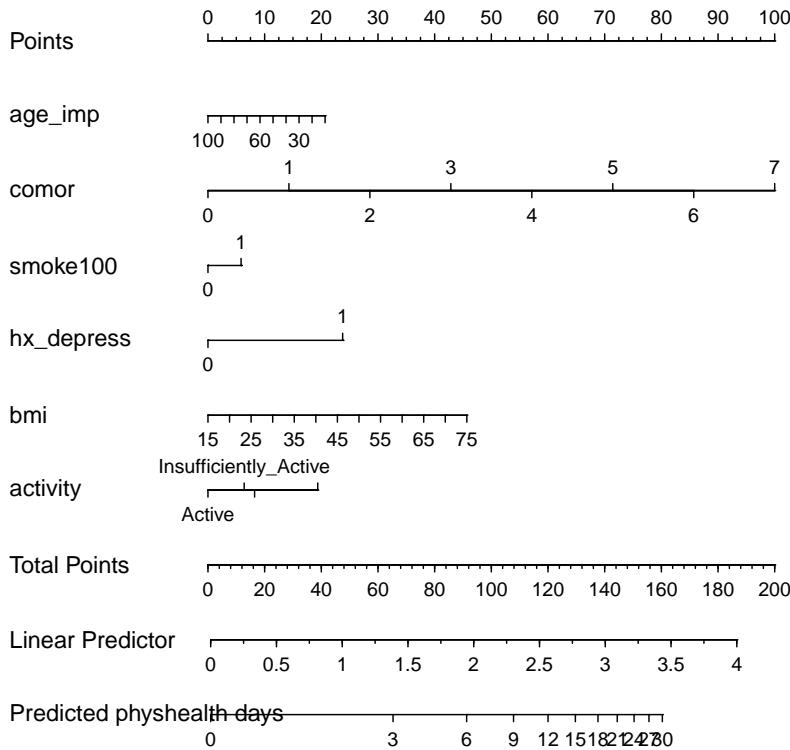
14.10.4 Nomogram

We can also develop a nomogram, if we like. As a special touch, we'll add a prediction at the bottom which back-transforms out of the predicted `phys_tr` back to the `physhealth` days.

```

plot(nomogram(m16_imp,
              fun = list(function(x) exp(x) - 1),
              funlabel = "Predicted physhealth days",
              fun.at = seq(0, 30, 3)))

```



We can see the big role of `comor` and `hx_depress` in this model.

14.10.5 Validating Summary Statistics

We can cross-validate summary measures, like R^2 ...

```
validate(m16_imp)
```

	index.orig	training	test	optimism	index.corrected	n
R-square	0.1867	0.1984	0.1793	0.0192	0.1676	40
MSE	1.6472	1.6168	1.6623	-0.0455	1.6927	40
g	0.6876	0.7031	0.6749	0.0282	0.6594	40
Intercept	0.0000	0.0000	0.0677	-0.0677	0.0677	40
Slope	1.0000	1.0000	0.9636	0.0364	0.9636	40

Chapter 15

Using `tidymodels` to fit linear regressions

Rather than create new chapters for `tidymodels` approaches, I will instead point you to the Class Slides, and to the relevant materials in *Tidy Modeling with R* by Max Kuhn and Julia Silge.

That text is available at <https://www.tmwr.org/>.

Chapter 16

Using `tidymodels` to fit logistic regressions

Rather than create new chapters for `tidymodels` approaches, I will instead point you to the Class Slides, and to the relevant materials in *Tidy Modeling with R* by Max Kuhn and Julia Silge.

That text is available at <https://www.tmwr.org/>.

Chapter 17

Using `tidymodels` approaches: Next Steps

Rather than create new chapters for `tidymodels` approaches, I will instead point you to the relevant materials in *Tidy Modeling with R* by Max Kuhn and Julia Silge.

That text is available at <https://www.tmwr.org/>.

Chapter 18

Colorectal Cancer Screening and Some Special Cases

In this Chapter, we discuss two issues not yet raised regarding regression on a binary outcome.

1. What do we do if our binary outcome is not available for each subject individually, but instead aggregated?
2. What is probit regression, and how can we use it as an alternative to logistic regression on a binary outcome?

18.1 Logistic Regression for Aggregated Data

18.1.1 Colorectal Cancer Screening Data

The `screening.csv` data (imported into the R tibble `cols`) are simulated. They mirror a subset of the actual results from the Better Health Partnership's pilot study of colorectal cancer screening in primary care clinics in Northeast Ohio, but the numbers have been fuzzed slightly, and the clinics have been de-identified and moved around from system to system.

Available to us are the following variables:

Variable	Description
<code>location</code>	clinic code
<code>subjects</code>	number of subjects reported by clinic

Variable	Description
screen_rate	proportion of subjects who were screened
screened	number of subjects who were screened
notscreened	number of subjects not screened
meanage	mean age of clinic's subjects, years
female	% of clinic's subjects who are female
pct_lowins	% of clinic's subjects who have Medicaid or are uninsured
system	system code

```
describe(colscr)
```

```
colscr
```

```
9 Variables      26 Observations
```

```
location
```

n	missing	distinct
26	0	26

```
lowest : A B C D E, highest: V W X Y Z
```

```
subjects
```

n	missing	distinct	Info	Mean	Gmd	.05	.10
26	0	26	1	3247	2134	1249	1475
.25	.50	.75	.90	.95			
1915	2766	3608	6523	7068			

```
lowest : 803 1179 1459 1491 1512, highest: 5451 6061 6985 7095 7677
```

```
screen_rate
```

n	missing	distinct	Info	Mean	Gmd	.05	.10
26	0	26	1	0.7659	0.08339	0.6682	0.6732
.25	.50	.75	.90	.95			
0.7179	0.7579	0.8088	0.8654	0.8899			

```
lowest : 0.6354901 0.6666667 0.6728914 0.6734521 0.6869919
```

```
highest: 0.8237129 0.8469110 0.8838745 0.8919112 0.9049108
```

```
screened
```

n	missing	distinct	Info	Mean	Gmd	.05	.10
26	0	26	1	2584	1895	843.5	993.0
.25	.50	.75	.90	.95			
1395.2	2169.5	2716.0	5293.5	6107.2			

```
lowest : 572 794 992 994 1088, highest: 4818 4848 5739 6230 6947
```

notscreened

	n	missing	distinct	Info	Mean	Gmd	.05	.10
26	0	26	1	663.2	303.5	336.0	352.5	
.25	.50	.75	.90	.95				
508.8	611.0	791.0	989.0	1172.5				

lowest : 231 335 339 366 371, highest: 881 927 1051 1213 1356

meanage

	n	missing	distinct	Info	Mean	Gmd	.05	.10
26	0	23	0.999	60.58	2.186	58.23	58.35	
.25	.50	.75	.90	.95				
58.82	60.50	61.98	62.50	62.90				

lowest : 58.0 58.2 58.3 58.4 58.5, highest: 62.2 62.4 62.6 63.0 65.9

female

	n	missing	distinct	Info	Mean	Gmd	.05	.10
26	0	23	0.999	58.72	7.118	46.93	48.45	
.25	.50	.75	.90	.95				
55.42	60.05	62.62	64.90	67.50				

lowest : 46.2 46.6 47.9 49.0 54.3, highest: 63.6 64.1 65.7 68.1 70.3

pct_lowins

	n	missing	distinct	Info	Mean	Gmd	.05	.10
26	0	24	0.999	24.47	22.12	0.675	1.800	
.25	.50	.75	.90	.95				
4.800	23.950	44.025	49.500	49.950				

lowest : 0.3 0.4 1.5 2.1 3.0, highest: 45.4 47.1 49.5 50.1 51.3

system

	n	missing	distinct
26	0	4	

Value	Sys_1	Sys_2	Sys_3	Sys_4
Frequency	7	7	6	6
Proportion	0.269	0.269	0.231	0.231

18.1.2 Fitting a Logistic Regression Model to Proportion Data

Here, we have a binary outcome (was the subject screened or not?) but we have aggregated results. We can use the counts of the numbers of subjects at each clinic (in `subjects`) and the proportion who were screened (in `screen_rate`) to fit a logistic regression model, as follows:

```
m_screen1 <- glm(screen_rate ~ meanage + female +
                     pct_lowins + system, family = binomial,
                     weights = subjects, data = colscr)

summary(m_screen1)
```

```
Call:
glm(formula = screen_rate ~ meanage + female + pct_lowins + system,
     family = binomial, data = colscr, weights = subjects)

Deviance Residuals:
    Min      1Q  Median      3Q      Max
-9.017 -3.705 -2.000   1.380  11.500

Coefficients:
            Estimate Std. Error z value Pr(>|z|)
(Intercept) -1.3270393  0.5530782 -2.399  0.0164 *
meanage      0.0679866  0.0089754  7.575 3.60e-14 ***
female       -0.0193142  0.0015831 -12.200 < 2e-16 ***
pct_lowins  -0.0134547  0.0008585 -15.672 < 2e-16 ***
systemSys_2 -0.1382189  0.0246591 -5.605 2.08e-08 ***
systemSys_3 -0.0400170  0.0254505 -1.572  0.1159
systemSys_4  0.0229273  0.0294207  0.779  0.4358
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

Null deviance: 2825.28  on 25  degrees of freedom
Residual deviance: 816.39  on 19  degrees of freedom
AIC: 1037.9

Number of Fisher Scoring iterations: 4
```

18.1.3 Fitting a Logistic Regression Model to Counts of Successes and Failures

```
m_screen2 <- glm(cbind(screened, notscreened) ~
  meanage + female + pct_lowins + system,
  family = binomial, data = colscr)
summary(m_screen2)

Call:
glm(formula = cbind(screened, notscreened) ~ meanage + female +
  pct_lowins + system, family = binomial, data = colscr)

Deviance Residuals:
    Min      1Q  Median      3Q     Max
-9.017 -3.705 -2.000   1.380  11.500

Coefficients:
            Estimate Std. Error z value Pr(>|z|)
(Intercept) -1.3270392  0.5530782 -2.399  0.0164 *
meanage      0.0679866  0.0089754  7.575 3.60e-14 ***
female       -0.0193142  0.0015831 -12.200 < 2e-16 ***
pct_lowins  -0.0134547  0.0008585 -15.672 < 2e-16 ***
systemSys_2 -0.1382189  0.0246591 -5.605 2.08e-08 ***
systemSys_3 -0.0400170  0.0254505 -1.572  0.1159
systemSys_4  0.0229273  0.0294207  0.779  0.4358
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

Null deviance: 2825.28 on 25 degrees of freedom
Residual deviance: 816.39 on 19 degrees of freedom
AIC: 1037.9

Number of Fisher Scoring iterations: 4
```

18.1.4 How does one address this problem in rms?

We can use `Glm`. As an example to mirror `m_screen1`, we have the following...

```
d <- datadist(colscr)
options(datadist = "d")

mod_screen_1 <- Glm(screen_rate ~ meanage + female +
  pct_lowins + system,
```

```

family = binomial, weights = subjects,
data = colscr, x = T, y = T)

mod_screen_1

General Linear Model

Glm(formula = screen_rate ~ meanage + female + pct_lowins + system,
    family = binomial, data = colscr, weights = subjects, x = T,
    y = T)

      Model Likelihood
      Ratio Test
Obs      26   LR chi2   2008.90
Residual d.f. 19   d.f.       6
g 0.4614539   Pr(> chi2) <0.0001

      Coef     S.E.   Wald Z Pr(>|Z| )
Intercept -1.3270 0.5531  -2.40 0.0164
meanage    0.0680 0.0090   7.57 <0.0001
female     -0.0193 0.0016 -12.20 <0.0001
pct_lowins -0.0135 0.0009 -15.67 <0.0001
system=Sys_2 -0.1382 0.0247  -5.61 <0.0001
system=Sys_3 -0.0400 0.0255  -1.57 0.1159
system=Sys_4  0.0229 0.0294   0.78 0.4358

```

18.2 Probit Regression

18.2.1 Colorectal Cancer Screening Data on Individuals

The data in the `colscr2` data frame describe (disguised) data on the status of 172 adults who were eligible for colon cancer screening, with the following information included:

Variable	Description
<code>subject</code>	subject ID code
<code>age</code>	subject's age (years)
<code>race</code>	subject's race (White/Black/Other)
<code>hispanic</code>	subject of Hispanic ethnicity (1 = yes / 0 = no)
<code>insurance</code>	Commercial, Medicaid, Medicare, Uninsured
<code>bmi</code>	body mass index at most recent visit
<code>sbp</code>	systolic blood pressure at most recent visit

Variable	Description
up_to_date	meets colon cancer screening standards

The goal is to use the other variables (besides subject ID) to predict whether or not a subject is up to date.

```
colscr2 %>% describe()
```

```
.
8 Variables      172 Observations
-----
subject
  n    missing   distinct      Info      Mean      Gmd      .05      .10
  172        0       172          1     186.5     57.67    109.6    118.1
  .25       .50       .75          .90       .95
  143.8    186.5    229.2        254.9    263.4

lowest : 101 102 103 104 105, highest: 268 269 270 271 272
-----
age
  n    missing   distinct      Info      Mean      Gmd      .05      .10
  172        0       19        0.995     57.8     5.536    51.00    52.00
  .25       .50       .75          .90       .95
  54.00    57.00    61.25        65.00    67.00

lowest : 51 52 53 54 55, highest: 65 66 67 68 69
-----
Value      51      52      53      54      55      56      57      58      59      60      61
Frequency  10      17      14       9      15      13      18      13       4      10       6
Proportion 0.058  0.099  0.081  0.052  0.087  0.076  0.105  0.076  0.023  0.058  0.035

Value      62      63      64      65      66      67      68      69
Frequency  11       4       5       7       6       3       4       3
Proportion 0.064  0.023  0.029  0.041  0.035  0.017  0.023  0.017
-----
race
  n    missing   distinct
  172        0       3

Value      Black   Other   White
Frequency  118      9      45
Proportion 0.686  0.052  0.262
-----
hispanic
```

	n	missing	distinct	Info	Sum	Mean	Gmd
	172	0	2	0.18	11	0.06395	0.1204

insurance			
	n	missing	distinct
	172	0	4

Value	Commercial	Medicaid	Medicare	Uninsured
Frequency	32	81	46	13
Proportion	0.186	0.471	0.267	0.076

	n	missing	distinct	Info	Mean	Gmd	.05	.10
	172	0	165	1	31.24	8.982	20.32	21.88
	.25	.50	.75	.90	.95			
	25.48	30.05	36.03	43.06	45.68			

lowest : 17.20 17.59 17.85 18.09 18.44, highest: 49.41 50.83 53.28 54.66 55.41

sbp								
	n	missing	distinct	Info	Mean	Gmd	.05	.10
	172	0	68	0.999	128.9	19.46	101.6	109.1
	.25	.50	.75	.90	.95			
	118.0	127.0	138.0	150.9	162.0			

lowest : 89 90 94 95 96, highest: 166 168 169 170 198

up_to_date							
	n	missing	distinct	Info	Sum	Mean	Gmd
	172	0	2	0.717	104	0.6047	0.4809

18.2.2 A logistic regression model

Here is a logistic regression model.

```
m_scr2_logistic <- glm(up_to_date ~ age + race + hispanic +
  insurance + bmi + sbp,
  family = binomial, data = colscr2)

summary(m_scr2_logistic)
```

Call:

```
glm(formula = up_to_date ~ age + race + hispanic + insurance +
bmi + sbp, family = binomial, data = colscr2)
```

Deviance Residuals:

Min	1Q	Median	3Q	Max
-1.8604	-1.1540	0.6933	0.9564	1.6108

Coefficients:

	Estimate	Std. Error	z value	Pr(> z)
(Intercept)	2.7040470	2.7418625	0.986	0.3240
age	0.0204901	0.0396920	0.516	0.6057
raceOther	-1.9722351	1.0023237	-1.968	0.0491 *
raceWhite	-0.3210458	0.4001744	-0.802	0.4224
hispanic	0.0005855	0.7953482	0.001	0.9994
insuranceMedicaid	-1.0151860	0.4945169	-2.053	0.0401 *
insuranceMedicare	-0.5216006	0.5629935	-0.926	0.3542
insuranceUninsured	0.1099966	0.7906196	0.139	0.8893
bmi	0.0155894	0.0213547	0.730	0.4654
sbp	-0.0241777	0.0099138	-2.439	0.0147 *

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

Null deviance: 230.85 on 171 degrees of freedom
Residual deviance: 210.55 on 162 degrees of freedom
AIC: 230.55

Number of Fisher Scoring iterations: 4

confint(m_scr2_logistic)

Waiting for profiling to be done...

	2.5 %	97.5 %
(Intercept)	-2.64274441	8.157738367
age	-0.05703934	0.099298904
raceOther	-4.24604744	-0.175953229
raceWhite	-1.10800168	0.469396096
hispanic	-1.53691431	1.688738936
insuranceMedicaid	-2.03640881	-0.079142467
insuranceMedicare	-1.66246932	0.564088260
insuranceUninsured	-1.40110091	1.759391281
bmi	-0.02568426	0.058563761
sbp	-0.04436143	-0.005282686

In this model, there appears to be some link between `sbp` and screening, as well as, perhaps, some statistically significant differences between some race groups and some insurance groups. We won't look at this much further for now, though. Instead, we'll simply describe predictions for two subjects, Harry and Sally.

18.2.3 Predicting status for Harry and Sally

- Harry is age 65, White, non-Hispanic, with Medicare insurance, a BMI of 28 and SBP of 135.
- Sally is age 60, Black, Hispanic, with Medicaid insurance, a BMI of 22 and SBP of 148.

```
newdat_s2 <- tibble(subject = c("Harry", "Sally"),
                      age = c(65, 60),
                      race = c("White", "Black"),
                      hispanic = c(0, 1),
                      insurance = c("Medicare", "Medicaid"),
                      bmi = c(28, 22),
                      sbp = c(135, 148))

predict(m_scr2_logistic, newdata = newdat_s2,
       type = "response")
```

```
1          2
0.5904364 0.4215335
```

The prediction for Harry is 0.59, and for Sally, 0.42, by this logistic regression model.

18.2.4 A probit regression model

Now, consider a probit regression, fit by changing the default link for the `binomial` family as follows:

```
m_scr2_probit <- glm(up_to_date ~ age + race + hispanic +
                       insurance + bmi + sbp,
                       family = binomial(link = "probit"),
                       data = colscr2)

summary(m_scr2_probit)
```

Call:
`glm(formula = up_to_date ~ age + race + hispanic + insurance +`
`bmi + sbp, family = binomial(link = "probit"), data = colscr2)`

Deviance Residuals:

```

      Min       1Q     Median      3Q      Max
-1.8608 -1.1561    0.6933    0.9607    1.6073

Coefficients:
            Estimate Std. Error z value Pr(>|z|)
(Intercept) 1.584604  1.658489  0.955   0.3394
age          0.013461  0.024107  0.558   0.5766
raceOther   -1.238445  0.587093 -2.109   0.0349 *
raceWhite   -0.199260  0.243505 -0.818   0.4132
hispanic     0.029483  0.484819  0.061   0.9515
insuranceMedicaid -0.619277  0.293205 -2.112   0.0347 *
insuranceMedicare -0.322881  0.333549 -0.968   0.3330
insuranceUninsured  0.052776  0.463798  0.114   0.9094
bmi          0.009652  0.012887  0.749   0.4539
sbp         -0.014696  0.005944 -2.472   0.0134 *
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

Null deviance: 230.85 on 171 degrees of freedom
Residual deviance: 210.49 on 162 degrees of freedom
AIC: 230.49

Number of Fisher Scoring iterations: 4
confint(m_scr2_probit)

Waiting for profiling to be done...

      2.5 %      97.5 %
(Intercept) -1.70739455  4.894713408
age          -0.03400934  0.061437784
raceOther   -2.53819772 -0.119812915
raceWhite   -0.67910311  0.282566349
hispanic    -0.92476109  1.011089937
insuranceMedicaid -1.21212866 -0.049749873
insuranceMedicare -0.99315267  0.329851215
insuranceUninsured -0.83335121  0.966504724
bmi          -0.01577917  0.035632196
sbp         -0.02659318 -0.003148488

```

18.2.5 Interpreting the Probit Model's Coefficients

It is possible to use any number of link functions to ensure that predicted values in a generalized linear model fall between 0 and 1. The probit regression model,

for instance, uses the inverse of the cumulative distribution function of the Normal model as its link function. Let's look more closely at the coefficients of the probit model we just fit.

```
m_scr2_probit$coef
```

	(Intercept)	age	raceOther	raceWhite
	1.584603569	0.013461338	-1.238445198	-0.199260184
hispanic		insuranceMedicaid	insuranceMedicare	insuranceUninsured
	0.029483051	-0.619276718	-0.322880519	0.052775722
bmi			sbp	
	0.009652339	-0.014695526		

The probit regression coefficients give the change in the z-score of the outcome of interest (here, `up_to_date`) for a one-unit change in the target predictor, holding all other predictors constant.

- So, for a one-year increase in age, holding all other predictors constant, the z-score for `up_to_date` increases by 0.013
- And for a Medicaid subject as compared to a Commercial subject of the same age, race, ethnicity, bmi and sbp, the z-score for the Medicaid subject is predicted to be -0.619 lower, according to this model.

18.2.6 What about Harry and Sally?

Do the predictions for Harry and Sally change much with this probit model, as compared to the logistic regression?

```
predict(m_scr2_probit, newdata = newdat_s2, type = "response")
```

1	2
0.5885511	0.4364027

Chapter 19

Modeling a Count Outcome in Ohio SMART

In this chapter, I use a count outcome (# of poor physical health days out of the last 30) to demonstrate regression models for count outcomes.

Methods discussed in the chapter include the following:

- Ordinary Least Squares
- Poisson Regression
- Overdispersed Quasi-Poisson Regression
- Negative Binomial Regression
- Zero-Inflated Poisson Regression
- Zero-Inflated Negative Binomial Regression
- Hurdle Models with Poisson counts
- Hurdle Models with Negative Binomial counts
- Tobit Regression

19.1 Preliminaries

```
library(boot)
library(lmtest)
library(sandwich)
library(countreg)
library(VGAM)

smart_oh <- readRDS(here("data", "smart_ohio.Rds"))
```

19.2 A subset of the Ohio SMART data

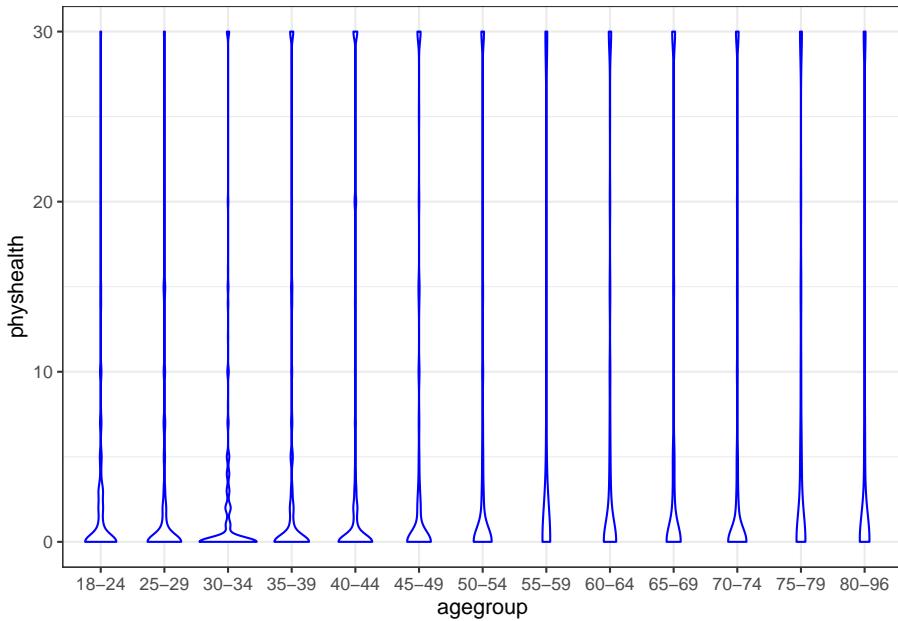
Let's consider the following data. I'll include the subset of all observations in `smart_oh` with complete data on these 14 variables.

Variable	Description
<code>SEQNO</code>	Subject identification code
<code>mmsa_name</code>	Name of statistical area
<code>genhealth</code>	Five categories (E, VG, G, F, P) on general health
<code>physhealth</code>	Now thinking about your physical health, which includes physical illness and injury, for how many days during the past 30 days was your physical health not good?
<code>menthlth</code>	Now thinking about your mental health, which includes stress, depression, and problems with emotions, for how many days during the past 30 days was your mental health not good?
<code>healthplan</code>	1 if the subject has any kind of health care coverage, 0 otherwise
<code>costprob</code>	1 indicates Yes to "Was there a time in the past 12 months when you needed to see a doctor but could not because of cost?"
<code>agegroup</code>	13 age groups from 18 through 80+
<code>female</code>	1 if subject is female
<code>incomegroup</code>	8 income groups from < 10,000 to 75,000 or more
<code>bmi</code>	body-mass index
<code>smoke100</code>	1 if Yes to "Have you smoked at least 100 cigarettes in your entire life?"
<code>alcdays</code>	# of days out of the past 30 on which the subject had at least one alcoholic drink

```
sm_oh_A <- smart_oh %>%
  select(SEQNO, mmsa_name, genhealth, physhealth,
         menthealth, healthplan, costprob,
         agegroup, female, incomegroup, bmi, smoke100,
         alcdays) %>%
  drop_na
```

19.2.1 Is age group associated with `physhealth`?

```
ggplot(sm_oh_A, aes(x = agegroup, y = physhealth)) +
  geom_violin(col = "blue")
```



It's hard to see much of anything here. The main conclusion seems to be that 0 is by far the most common response.

Here's a table by age group of:

- the number of respondents in that age group,
- the group's mean `physhealth` response (remember that these are the number of poor physical health days in the last 30),
- their median `physhealth` response (which turns out to be 0 in each group), and
- the percentage of group members who responded 0.

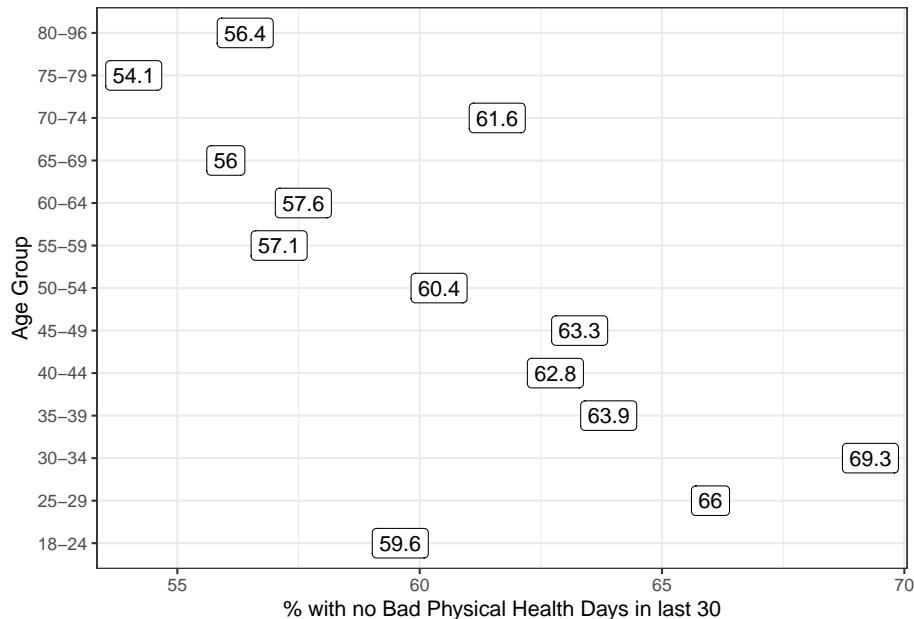
```
sm_oh_A %>% group_by(agegroup) %>%
  summarize(n = n(), mean = round(mean(physhealth), 2),
            median = median(physhealth),
            percent_0s = round(100*sum(physhealth == 0)/n, 1))
```

	agegroup	n	mean	median	percent_0s
	<fct>	<int>	<dbl>	<dbl>	<dbl>
1	18-24	297	2.31	0	59.6
2	25-29	259	2.53	0	66
3	30-34	296	2.14	0	69.3
4	35-39	366	3.67	0	63.9
5	40-44	347	4.18	0	62.8
6	45-49	409	4.46	0	63.3
7	50-54	472	4.76	0	60.4

8	55-59	608	6.71	0	57.1
9	60-64	648	5.9	0	57.6
10	65-69	604	6.09	0	56
11	70-74	490	4.89	0	61.6
12	75-79	338	6.38	0	54.1
13	80-96	374	6.1	0	56.4

We can see a real change between the 45-49 age group and the 50-54 age group. The mean difference is clear from the table above, and the plot below (of the percentage with a zero response) in each age group identifies the same story.

```
sm_oh_A %>% group_by(agegroup) %>%
  summarize(n = n(),
            percent_0s = round(100*sum(physhealth == 0)/n,1)) %>%
  ggplot(aes(y = agegroup, x = percent_0s)) +
  geom_label(aes(label = percent_0s)) +
  labs(x = "% with no Bad Physical Health Days in last 30",
       y = "Age Group")
```



It looks like we have a fairly consistent result in the younger age range (18-49) or the older range (50+). On the theory that most of the people reading this document are in that younger range, we'll focus on those respondents in what follows.

19.3 Exploratory Data Analysis (in the 18-49 group)

We want to predict the 0-30 `physhealth` count variable for the 18-49 year old respondents.

To start, we'll use two predictors:

- the respondent's body mass index, and
- whether the respondent has smoked 100 cigarettes in their lifetime.

We anticipate that each of these variables will have positive associations with the `physhealth` score. That is, heavier people, and those who have used tobacco will be less healthy, and thus have higher numbers of poor physical health days.

19.3.1 Build a subset of those ages 18-49

First, we'll identify the subset of respondents who are between 18 and 49 years of age.

```
sm_oh_A_young.raw <- sm_oh_A %>%
  filter(agegroup %in% c("18-24", "25-29", "30-34",
                        "35-39", "40-44", "45-49")) %>%
  droplevels()

sm_oh_A_young.raw %>%
  select(physhealth, bmi, smoke100, agegroup) %>%
  summary

  physhealth          bmi          smoke100      agegroup
Min.    : 0.000  Min.   :14.00  Min.   :0.0000  18-24:297
1st Qu.: 0.000  1st Qu.:23.74  1st Qu.:0.0000  25-29:259
Median  : 0.000  Median :27.32  Median :0.0000  30-34:296
Mean    : 3.337  Mean   :28.79  Mean   :0.4189  35-39:366
3rd Qu.: 2.000  3rd Qu.:32.43  3rd Qu.:1.0000  40-44:347
Max.    :30.000  Max.   :75.52  Max.   :1.0000  45-49:409
```

19.3.2 Centering `bmi`

I'm going to center the `bmi` variable to help me interpret the final models later.

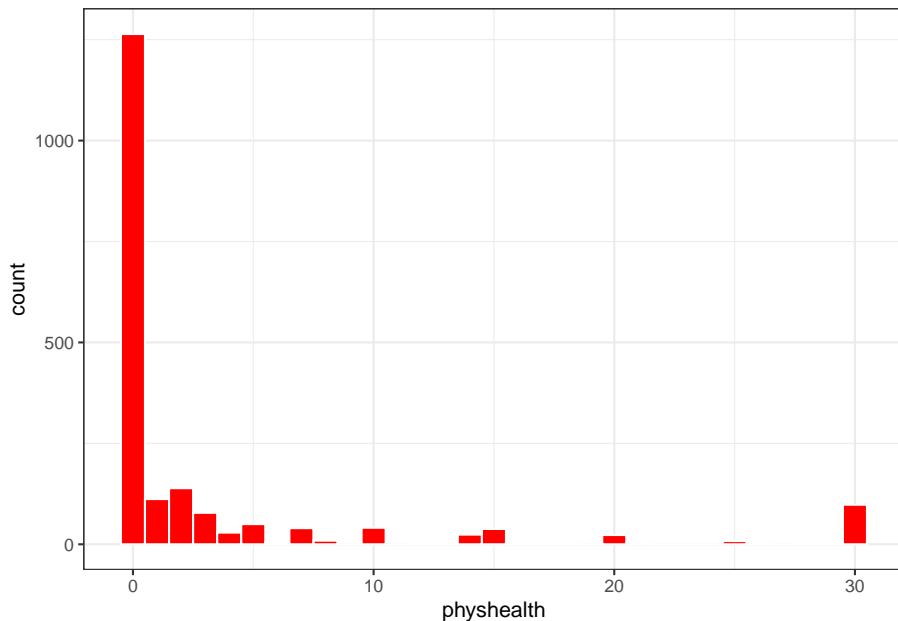
```
sm_oh_A_young <- sm_oh_A_young.raw %>%
  mutate(bmi_c = bmi - mean(bmi))
```

Now, let's look more closely at the distribution of these variables, starting with our outcome.

19.3.3 Distribution of the Outcome

What's the distribution of `physhealth`?

```
ggplot(sm_oh_A_young.raw, aes(x = physhealth)) +
  geom_histogram(binwidth = 1, fill = "red", col = "white")
```



```
sm_oh_A_young.raw %>%
  count(physhealth == 0, physhealth == 30)
```

	<code>physhealth == 0</code>	<code>physhealth == 30</code>	n
1	FALSE	FALSE	612
2	FALSE	TRUE	98
3	TRUE	FALSE	1264

Most of our respondents said zero, the minimum allowable value, although there is also a much smaller bump at 30, the maximum value we will allow.

Dealing with this distribution is going to be a bit of a challenge. We will develop a series of potential modeling approaches for this sort of data, but before we do that, let's look at the distribution of our other two variables, and the pairwise associations, in a scatterplot matrix.

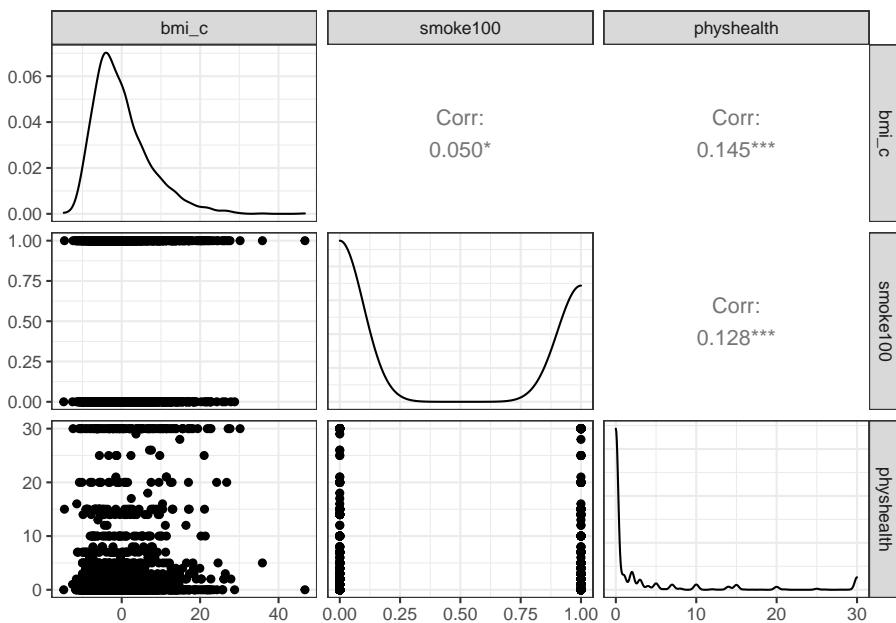
19.3.4 Scatterplot Matrix

Now, here's the scatterplot matrix for those 1974 subjects, using the centered `bmi` data captured in the `bmi_c` variable.

```
GGally::ggpairs(sm_oh_A_young %>%
                  select(bmi_c, smoke100, physhealth))
```

Registered S3 method overwritten by 'GGally':

```
method from
+.gg   ggplot2
```



So `bmi_c` and `smoke100` each have modest positive correlations with `physhealth` and only a very small correlation with each other. Here are some summary statistics for this final data.

19.3.5 Summary of the final subset of data

Remember that since the mean of `bmi` is 28.8, the `bmi_c` values are just `bmi` - 28.8 for each subject.

```
sm_oh_A_young %>%
  select(bmi, bmi_c, smoke100, physhealth) %>%
  summary()
```

bmi	bmi_c	smoke100	physhealth
-----	-------	----------	------------

Min.	:14.00	Min.	:-14.791
1st Qu.	:23.74	1st Qu.:	-5.051
Median	:27.32	Median :	-1.471
Mean	:28.79	Mean :	0.000
3rd Qu.	:32.43	3rd Qu.:	3.636
Max.	:75.52	Max. :	46.729
		Min. :	0.000
		1st Qu.:	0.000
		Median :	0.000
		Mean :	0.4189
		3rd Qu.:	1.0000
		Max. :	30.000

19.4 Modeling Strategies Explored Here

We are going to predict `physhealth` using `bmi_c` and `smoke100`.

- Remember that `physhealth` is a count of the number of poor physical health days in the past 30.
- As a result, `physhealth` is restricted to taking values between 0 and 30.

We will demonstrate the use of each of the following regression models, some of which are better choices than others.

1. Ordinary Least Squares (OLS) predicting `physhealth`
2. OLS predicting the logarithm of (`physhealth` + 1)
3. Poisson regression, which is appropriate for predicting counts
4. Poisson regression, adjusted to account for overdispersion
5. Negative binomial regression, also used for counts and which adjusts for overdispersion
6. Zero-inflated models, in both the Poisson and Negative Binomial varieties, which allow us to fit counts that have lots of zero values
7. A “hurdle” model, which allows us to separately fit a model to predict the incidence of “0” and then a separate model to predict the value of `physhealth` when we know it is not zero
8. Tobit regression, where a lower (and upper) bound may be set, but the underlying model describes a latent variable which can extend beyond these boundaries

19.4.1 What Will We Demonstrate?

With each approach, we will fit the model and specify procedures for doing so in R. Then we will:

1. Specify the fitted model equation
2. Interpret the model’s coefficient estimates and 95% confidence intervals around those estimates.
3. Perform a test of whether each variable adds value to the model, when the other one is already included.
4. Store the fitted values and appropriate residuals for each model.

5. Summarize the model's apparent R^2 value, the proportion of variation explained, and the model log likelihood.
6. Perform checks of model assumptions as appropriate.
7. Describe how predictions would be made for two new subjects.
 - Harry has a BMI that is 10 kg/m² higher than the average across all respondents and has smoked more than 100 cigarettes in his life.
 - Sally has a BMI that is 5 kg/m² less than the average across all respondents and has not smoked more than 100 cigarettes in her life.

In addition, for some of the new models, we provide a little of the mathematical background, and point to other resources you can use to learn more about the model.

19.4.2 Extra Data File for Harry and Sally

To make our lives a little easier, I'll create a little tibble containing the necessary data for Harry and Sally.

```
hs_data <- data_frame(subj = c("Harry", "Sally"),
                      bmi_c = c(10, -5),
                      smoke100 = c(1, 0))
hs_data

# A tibble: 2 x 3
  subj   bmi_c smoke100
  <chr> <dbl>    <dbl>
1 Harry     10        1
2 Sally     -5        0
```

19.5 The OLS Approach

```
mod_ols1 <- lm(physhealth ~ bmi_c + smoke100,
                data = sm_oh_A_young)

summary(mod_ols1)

Call:
lm(formula = physhealth ~ bmi_c + smoke100, data = sm_oh_A_young)

Residuals:
    Min      1Q  Median      3Q     Max 
-11.1472 -3.6639 -2.2426 -0.7807 28.8777 

Coefficients:
```

```

      Estimate Std. Error t value Pr(>|t|)
(Intercept) 2.57046   0.21648 11.874 < 2e-16 ***
bmi_c       0.14437   0.02305  6.263 4.61e-10 ***
smoke100    1.83061   0.33469  5.470 5.09e-08 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 7.328 on 1971 degrees of freedom
Multiple R-squared:  0.0356, Adjusted R-squared:  0.03462
F-statistic: 36.37 on 2 and 1971 DF, p-value: 3.072e-16
confint(mod_ols1)

      2.5 % 97.5 %
(Intercept) 2.1459143 2.995005
bmi_c       0.0991636 0.189573
smoke100    1.1742213 2.486995

```

19.5.1 Interpreting the Coefficients

- The intercept, 2.57, is the predicted `physhealth` (in days) for a subject with average BMI who has not smoked 100 cigarettes or more.
- The `bmi_c` coefficient, 0.144, indicates that for each additional kg/m² of BMI, while holding `smoke100` constant, the predicted `physhealth` value increases by 0.144 day.
- The `smoke100` coefficient, 1.83, indicates that a subject who has smoked 100 cigarettes or more has a predicted `physhealth` value 1.83 days larger than another subject with the same `bmi` but who has not smoked 100 cigarettes.

19.5.2 Store fitted values and residuals

We can use `broom` to do this. Here, for instance, is a table of the first six predictions and residuals.

```

sm_ols_1 <- augment(mod_ols1, sm_oh_A_young)

sm_ols_1 %>% select(physhealth, .fitted, .resid) %>% head()

# A tibble: 6 x 3
  physhealth .fitted .resid
     <dbl>     <dbl>   <dbl>
1         0     2.13  -2.13
2         0     2.25  -2.25
3         0     3.14  -3.14

```

```
4      30    5.72  24.3
5      0    3.20 -3.20
6      0    3.83 -3.83
```

It turns out that 0 of the 1974 predictions that we make are below 0, and the largest prediction made by this model is 11.15 days.

19.5.3 Specify the R^2 and log(likelihood) values

The `glance` function in the `broom` package gives us the raw and adjusted R^2 values, and the model log(likelihood), among other summaries.

```
glance(mod_ols1) %>% round(3)

# A tibble: 1 x 12
  r.squared adj.r.squared sigma statistic p.value    df logLik     AIC     BIC
        <dbl>         <dbl>   <dbl>     <dbl>    <dbl>  <dbl> <dbl>    <dbl>
1     0.036       0.035    7.33     36.4      0      2 -6731. 13470. 13492.
# ... with 3 more variables: deviance <dbl>, df.residual <dbl>, nobs <dbl>
```

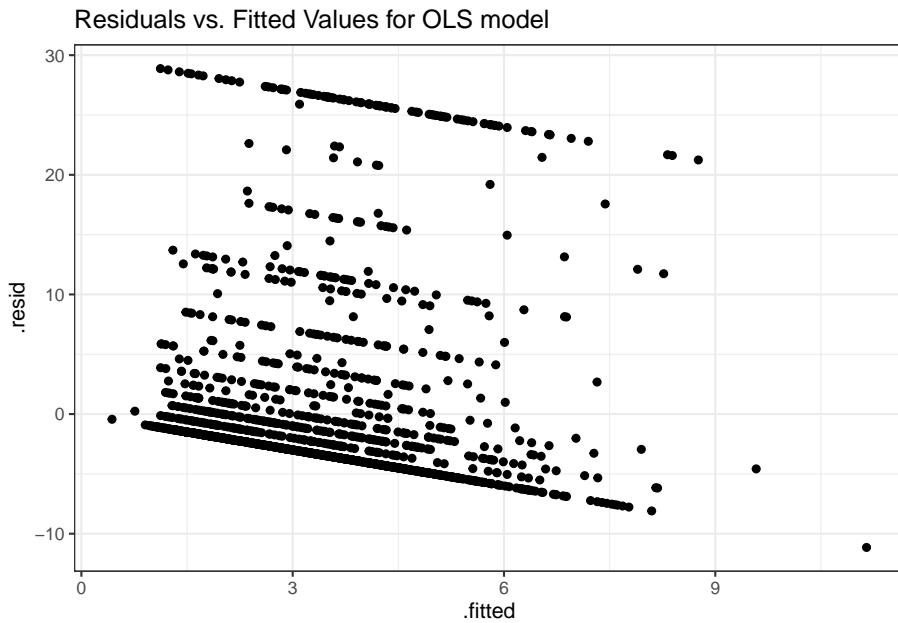
Here, we have

Model	R^2	log(likelihood)
OLS	0.036	-6730.98

19.5.4 Check model assumptions

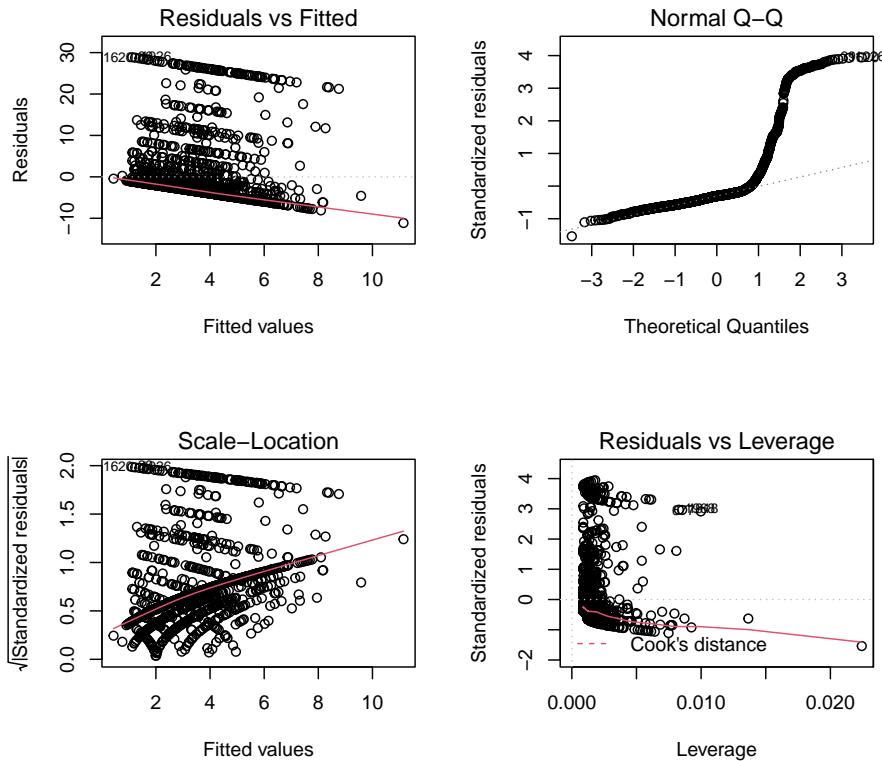
Here is a plot of the residuals vs. the fitted values for this OLS model.

```
ggplot(sm_ols_1, aes(x = .fitted, y = .resid)) +
  geom_point() +
  labs(title = "Residuals vs. Fitted Values for OLS model")
```



As usual, we can check OLS assumptions (linearity, homoscedasticity and normality) with R's complete set of residual plots.

```
par(mfrow = c(2,2))
plot(mod_ols1)
```



```
par(mfrow = c(1,1))
```

We see the problem with our residuals. They don't follow a Normal distribution.

19.5.5 Predictions for Harry and Sally

```
predict(mod_ols1, newdata = hs_data,
       interval = "prediction")
```

	fit	lwr	upr
1	5.844750	-8.541164	20.23067
2	1.848618	-12.529923	16.22716

The prediction for Harry is 5.8 days, and for Sally is 1.8 days. The prediction intervals for each include some values below 0, even though 0 is the smallest possible value.

19.5.6 Notes

- This model could have been estimated using the `ols` function in the `rms` package, as well.

```
dd <- datadist(sm_oh_A_young)
options(datadist = "dd")

(mod_ols1a <- ols(physhealth ~ bmi_c + smoke100,
                   data = sm_oh_A_young, x = TRUE, y = TRUE))
```

Linear Regression Model

```
ols(formula = physhealth ~ bmi_c + smoke100, data = sm_oh_A_young,
    x = TRUE, y = TRUE)
```

	Model	Likelihood	Discrimination
		Ratio Test	Indexes
Obs	1974	LR chi2	71.55
sigma	7.3276	d.f.	2
d.f.	1971	Pr(> chi2)	0.0000
		g	1.570

Residuals

	Min	1Q	Median	3Q	Max
	-11.1472	-3.6639	-2.2426	-0.7807	28.8777

	Coef	S.E.	t	Pr(> t)
Intercept	2.5705	0.2165	11.87	<0.0001
bmi_c	0.1444	0.0230	6.26	<0.0001
smoke100	1.8306	0.3347	5.47	<0.0001

19.6 OLS model on `log(physhealth + 1)` days

We could try to solve the problem of fitting some predictions below 0 by log-transforming the data, so as to force values to be at least 0. Since we have undefined values when we take the log of 0, we'll add one to each of the `physhealth` values before taking logs, and then transform back when we want to make predictions.

```
mod_ols_log1 <- lm(log(physhealth + 1) ~ bmi_c + smoke100,
                     data = sm_oh_A_young)
```

```
summary(mod_ols_log1)

Call:
lm(formula = log(physhealth + 1) ~ bmi_c + smoke100, data = sm_oh_A_young)

Residuals:
    Min      1Q  Median      3Q     Max 
-1.7079 -0.7058 -0.5051  0.5053  3.0484 

Coefficients:
            Estimate Std. Error t value Pr(>|t|)    
(Intercept) 0.57746   0.03099 18.634 < 2e-16 ***
bmi_c       0.01912   0.00330  5.796 7.91e-09 ***
smoke100    0.23679   0.04791  4.942 8.38e-07 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 1.049 on 1971 degrees of freedom
Multiple R-squared:  0.03003, Adjusted R-squared:  0.02905 
F-statistic: 30.51 on 2 and 1971 DF,  p-value: 8.897e-14

confint(mod_ols_log1)

          2.5 %    97.5 %    
(Intercept) 0.51668682 0.63823590  
bmi_c       0.01265192 0.02559421  
smoke100    0.14282518 0.33075147
```

19.6.1 Interpreting the Coefficients

- The intercept, 0.58, is the predicted logarithm of $(\text{physhealth} + 1)$ (in days) for a subject with average BMI who has not smoked 100 cigarettes or more.
 - We can exponentiate to see that the prediction for $(\text{physhealth} + 1)$ here is $\exp(0.58) = 1.79$ so the predicted physhealth for a subject with average BMI who has not smoked 100 cigarettes is 0.79 days.
- The `bmi_c` coefficient, 0.019, indicates that for each additional kg/m² of BMI, while holding `smoke100` constant, the predicted logarithm of $(\text{physhealth} + 1)$ increases by 0.019
- The `smoke100` coefficient, 0.24, indicates that a subject who has smoked 100 cigarettes or more has a predicted log of $(\text{physhealth} + 1)$ value that is 0.24 larger than another subject with the same `bmi` but who has not smoked 100 cigarettes.

19.6.2 Store fitted values and residuals

We can use `broom` to help us with this. Here, for instance, is a table of the first six predictions and residuals, on the scale of our transformed response, $\log(\text{physhealth} + 1)$.

```
sm_ols_log1 <- augment(mod_ols_log1, sm_oh_A_young)

sm_ols_log1 <- sm_ols_log1 %>%
  mutate(outcome = log(physhealth + 1))

sm_ols_log1 %>%
  select(physhealth, outcome, .fitted, .resid) %>%
  head()

# A tibble: 6 x 4
  physhealth outcome .fitted .resid
  <dbl>     <dbl>    <dbl>   <dbl>
1 0         0       0.520 -0.520
2 0         0       0.535 -0.535
3 0         0       0.647 -0.647
4 30        3.43   0.995  2.44 
5 0         0       0.656 -0.656
6 0         0       0.739 -0.739
```

Note that the `outcome` used in this model is $\log(\text{physhealth} + 1)$, so the `.fitted` and `.resid` values react to that outcome, and not to our original `physhealth`.

Another option would be to calculate the model-predicted `physhealth`, which I'll call `ph` for a moment, with the formula:

$$ph = e^{.fitted} - 1$$

```
sm_ols_log1 <- sm_ols_log1 %>%
  mutate(pred.physhealth = exp(.fitted) - 1,
        res.physhealth = physhealth - pred.physhealth)

sm_ols_log1 %>%
  select(physhealth, pred.physhealth, res.physhealth) %>%
  head()

# A tibble: 6 x 3
  physhealth pred.physhealth res.physhealth
  <dbl>          <dbl>        <dbl>
1 0            0.681        -0.681
2 0            0.708        -0.708
3 0            0.910        -0.910
```

4	30	1.70	28.3
5	0	0.926	-0.926
6	0	1.09	-1.09

It turns out that 0 of the 1974 predictions that we make are below 0, and the largest prediction made by this model is 4.52 days.

19.6.3 Specify the R^2 and log(likelihood) values

The `glance` function in the `broom` package gives us the raw and adjusted R^2 values, and the model log(likelihood), among other summaries.

```
glance(mod_ols_log1) %>% round(3)
```

```
# A tibble: 1 x 12
  r.squared adj.r.squared sigma statistic p.value    df logLik     AIC     BIC
  <dbl>        <dbl> <dbl>      <dbl>    <dbl> <dbl> <dbl> <dbl> <dbl>
1     0.03        0.029  1.05      30.5      0      2 -2894. 5796. 5818.
# ... with 3 more variables: deviance <dbl>, df.residual <dbl>, nobs <dbl>
```

Here, we have

Model	Scale	R^2	log(likelihood)
OLS on log $\log(\text{physhealth} + 1)$	0.03	-2893.83	

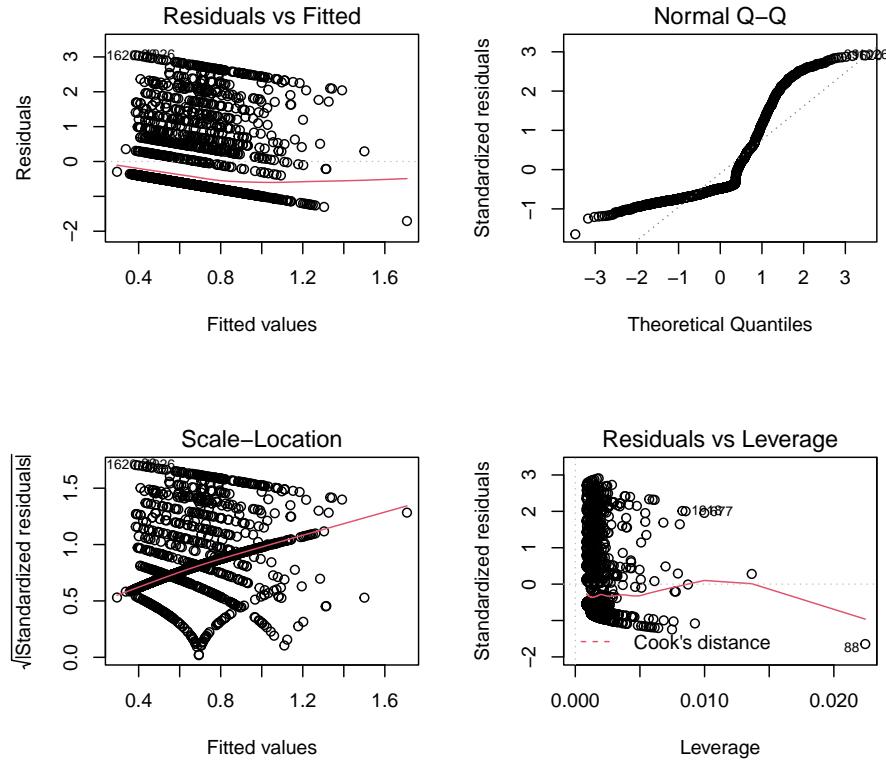
19.6.4 Getting R^2 on the scale of `physhealth`

We could find the correlation of our model-predicted `physhealth` values, after back-transformation, and our observed `physhealth` values, if we wanted to, and then square that to get a sort of R^2 value. But this model is not linear in `physhealth`, of course, so it's not completely comparable to our prior OLS model.

19.6.5 Check model assumptions

As usual, we can check OLS assumptions (linearity, homoscedasticity and normality) with R's complete set of residual plots. Of course, these residuals and fitted values are now on the $\log(\text{physhealth} + 1)$ scale.

```
par(mfrow = c(2,2))
plot(mod_ols_log1)
```



```
par(mfrow = c(1,1))
```

19.6.6 Predictions for Harry and Sally

```
predict(mod_ols_log1, newdata = hs_data,
       interval = "prediction", type = "response")
```

	fit	lwr	upr
1	1.005480	-1.053893	3.064854
2	0.481846	-1.576472	2.540164

Again, these predictions are on the $\log(\text{physhealth} + 1)$ scale, and so we have to exponentiate them, and then subtract 1, to see them on the original `physhealth` scale.

```
exp(predict(mod_ols_log1, newdata = hs_data,
            interval = "prediction", type = "response")) - 1
```

	fit	lwr	upr
1	1.7332198	-0.6514221	20.43133
2	0.6190605	-0.7932970	11.68175

The prediction for Harry is now 1.73 days, and for Sally is 0.62 days. The prediction intervals for each again include some values below 0, which is the smallest possible value.

19.7 A Poisson Regression Model

The `physhealth` data describe a count. Specifically a count of the number of days where the subject felt poorly in the last 30. Why wouldn't we model this count with linear regression?

- A count can only be positive. Linear regression would estimate some subjects as having negative counts.
- A count is unlikely to follow a Normal distribution. In fact, it's far more likely that the log of the counts will follow a Poisson distribution.

So, we'll try that. The Poisson distribution is used to model a *count* outcome - that is, an outcome with possible values (0, 1, 2, ...). The model takes a somewhat familiar form to the models we've used for linear and logistic regression¹. If our outcome is y and our linear predictors X , then the model is:

$$y_i \sim \text{Poisson}(\lambda_i)$$

The parameter λ must be positive, so it makes sense to fit a linear regression on the logarithm of this...

$$\lambda_i = \exp(\beta_0 + \beta_1 X_1 + \dots + \beta_k X_k)$$

The coefficients β can be exponentiated and treated as multiplicative effects.

We'll run a generalized linear model with a log link function, ensuring that all of the predicted values will be positive, and using a Poisson error distribution. This is called **Poisson regression**.

Poisson regression may be appropriate when the dependent variable is a count of events. The events must be independent - the occurrence of one event must not make any other more or less likely. That's hard to justify in our case, but we can take a look.

¹This discussion is motivated by Section 6.2 of Gelman and Hill.

```

mod_poiss1 <- glm(physhealth ~ bmi_c + smoke100,
                    family = poisson(),
                    data = sm_oh_A_young)

summary(mod_poiss1)

Call:
glm(formula = physhealth ~ bmi_c + smoke100, family = poisson(),
     data = sm_oh_A_young)

Deviance Residuals:
    Min      1Q  Median      3Q      Max
-6.588 -2.604 -2.087 -0.533 10.688

Coefficients:
            Estimate Std. Error z value Pr(>|z|)
(Intercept) 0.906991  0.018699 48.51   <2e-16 ***
bmi_c       0.035051  0.001421 24.66   <2e-16 ***
smoke100    0.532505  0.024903 21.38   <2e-16 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for poisson family taken to be 1)

Null deviance: 20222  on 1973  degrees of freedom
Residual deviance: 19151  on 1971  degrees of freedom
AIC: 21645

Number of Fisher Scoring iterations: 6

confint(mod_poiss1)

Waiting for profiling to be done...

          2.5 %    97.5 %
(Intercept) 0.87011622 0.94342323
bmi_c       0.03225125 0.03782255
smoke100    0.48373882 0.58136497

```

19.7.1 The Fitted Equation

The model equation is

```
log(physhealth) = 0.91 + 0.035 bmi_c + 0.53 smoke100
```

It looks like both `bmi` and `smoke_100` have confidence intervals excluding 0.

19.7.2 Interpreting the Coefficients

Our new model for y_i = counts of poor `physhealth` days in the last 30, follows the regression equation:

$$y_i \sim \text{Poisson}(\exp(0.91 + 0.035\text{bmi}_c + 0.53\text{smoke100}))$$

where `smoke100` is 1 if the subject has smoked 100 cigarettes (lifetime) and 0 otherwise, and `bmi_c` is just the centered body-mass index value in kg/m². We interpret the coefficients as follows:

- The constant term, 0.91, gives us the intercept of the regression - the prediction if `smoke100` = 0 and `bmi_c` = 0. In this case, because we've centered BMI, it implies that $\exp(0.91) = 2.48$ is the predicted days of poor `physhealth` for a non-smoker with average BMI.
- The coefficient of `bmi_c`, 0.035, is the expected difference in count of poor `physhealth` days (on the log scale) for each additional kg/m² of body mass index. The expected multiplicative *increase* is $e^{0.035} = 1.036$, corresponding to a 3.6% difference in the count.
- The coefficient of `smoke100`, 0.53, tells us that the predictive difference between those who have and who have not smoked 100 cigarettes can be found by multiplying the `physhealth` count by $\exp(0.53) = 1.7$, yielding a 70% increase of the `physhealth` count.

As with linear or logistic regression, each coefficient is interpreted as a comparison where one predictor changes by one unit, while the others remain constant.

19.7.3 Testing the Predictors

We can use the Wald tests (z tests) provided with the Poisson regression output, or we can fit the model and then run an ANOVA to obtain a test based on the deviance (a simple transformation of the log likelihood ratio.)

- By the Wald tests shown above, each predictor clearly adds significant predictive value to the model given the other predictor, and we note that the *p* values are as small as R will support.
- The ANOVA approach for this model lets us check the impact of adding `smoke100` to a model already containing `bmi_c`.

```
anova(mod_poiss1, test = "Chisq")
```

Analysis of Deviance Table

Model: poisson, link: log

Response: `physhealth`

```

Terms added sequentially (first to last)

          Df Deviance Resid. Df Resid. Dev Pr(>Chi)
NULL                  1973      20222
bmi_c     1    609.46      1972    19612 < 2.2e-16 ***
smoke100  1    461.46      1971    19151 < 2.2e-16 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

```

To obtain a p value for `smoke100`'s impact after `bmi_c` is accounted for, we compare the difference in deviance to a chi-square distribution with 1 degree of freedom. To check the effect of `bmi_c`, we could refit the model with `bmi_c` entering last, and again run an ANOVA.

We could also run a likelihood-ratio test for each predictor, by fitting the model with and without that predictor.

```

mod_poiss1_without_bmi <- glm(physhealth ~ smoke100,
                                family = poisson(),
                                data = sm_oh_A_young)

anova(mod_poiss1, mod_poiss1_without_bmi, test = "Chisq")

```

Analysis of Deviance Table

```

Model 1: physhealth ~ bmi_c + smoke100
Model 2: physhealth ~ smoke100
          Resid. Df Resid. Dev Df Deviance Pr(>Chi)
1           1971      19151
2           1972      19692 -1   -540.98 < 2.2e-16 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

```

19.7.4 Correcting for Overdispersion with `coeftest/coefci`

The main assumption we'll think about in a Poisson model is about **overdispersion**. We might deal with the overdispersion we see in this model by changing the nature of the tests we run within this model, using the `coeftest` or `coefci` approaches from the `lmtest` package, as I'll demonstrate next, or we might refit the model using a quasi-likelihood approach, as I'll show in the material to come.

Here, we'll use the `coeftest` and `coefci` approach from `lmtest` combined with robust sandwich estimation (via the `sandwich` package) to re-compute the Wald tests.

```

coeftest(mod_poiss1, vcov. = sandwich)

z test of coefficients:

      Estimate Std. Error z value Pr(>|z|)
(Intercept) 0.9069908 0.0717221 12.6459 < 2.2e-16 ***
bmi_c       0.0350508 0.0061178  5.7293 1.008e-08 ***
smoke100    0.5325053 0.1004300  5.3023 1.144e-07 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

coefci(mod_poiss1, vcov. = sandwich)

      2.5 %     97.5 %
(Intercept) 0.76641801 1.04756366
bmi_c       0.02306016 0.04704145
smoke100    0.33566606 0.72934460

```

Both predictors are still significant, but the standard errors are more appropriate. Later, we'll fit this approach by changing the estimation method to a quasi-likelihood approach.

19.7.5 Store fitted values and residuals

What happens if we try using the `broom` package in this case? We can, if we like, get our residuals and predicted values right on the scale of our `physhealth` response.

```

sm_poiss1 <- augment(mod_poiss1, sm_oh_A_young,
                      type.predict = "response")

sm_poiss1 %>%
  select(physhealth, .fitted) %>%
  head()

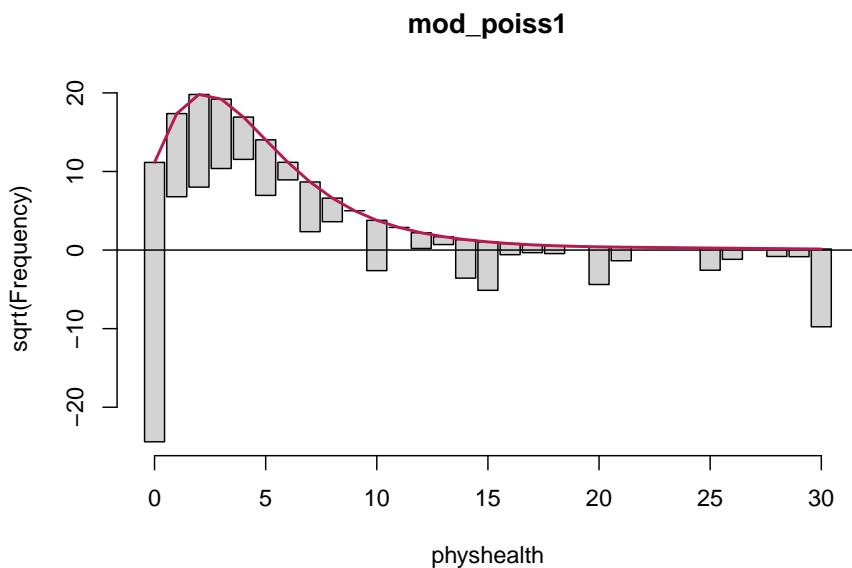
# A tibble: 6 x 2
  physhealth .fitted
  <dbl>     <dbl>
1 0         2.23
2 0         2.29
3 0         3.10
4 30        5.32
5 0         3.15
6 0         3.68

```

19.7.6 Rootogram: see the fit of a count regression model

A **rootogram** is a very useful way to visualize the fit of a count regression model². The `rootogram` function in the `countreg` package makes this pretty straightforward. By default, this fits a hanging rootogram on the square root of the frequencies.

```
rootogram(mod_poiss1, max = 30)
```



The red curved line is the theoretical Poisson fit. “Hanging” from each point on the red line is a bar, the height of which represents the difference between expected and observed counts. A bar hanging below 0 indicates underfitting. A bar hanging above 0 indicates overfitting. The counts have been transformed with a square root transformation to prevent smaller counts from getting obscured and overwhelmed by larger counts. We see a great deal of underfitting for counts of 0, and overfitting for most other counts, especially 1-6, with some underfitting again by `physhealth` above 14 days.

19.7.7 Specify the R^2 and log(likelihood) values

We can calculate the R^2 as the squared correlation of the fitted values and the observed values.

²See <http://data.library.virginia.edu/getting-started-with-negative-binomial-regression-modeling/>

```
# The correlation of observed and fitted values
(poiss_r <- with(sm_poiss1, cor(physhealth, .fitted)))

[1] 0.1846814

# R-square
poiss_r^2

[1] 0.03410723
```

The `glance` function in the `broom` package gives us model `log(likelihood)`, among other summaries.

```
glance(mod_poiss1) %>% round(3)

# A tibble: 1 x 8
  null.deviance df.null  logLik     AIC     BIC deviance df.residual nobs
        <dbl>    <dbl>    <dbl>    <dbl>    <dbl>    <dbl>    <dbl>    <dbl>
1      20222.     1973 -10820.  21645.  21662.   19151.     1971    1974
```

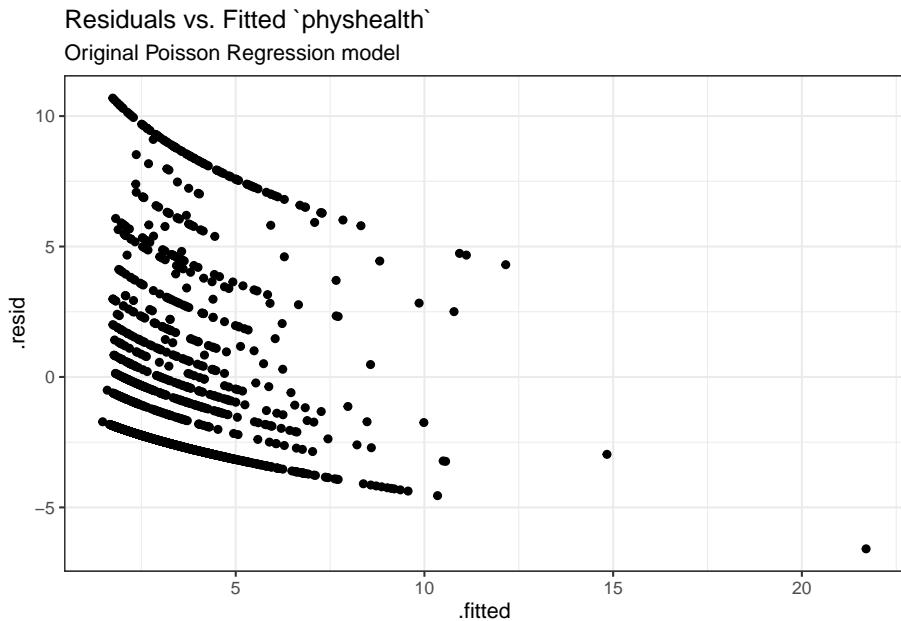
Here, we have

Model	Scale	R^2	<code>log(likelihood)</code>
Poisson	<code>log(physhealth)</code>	0.185	-10189.33

19.7.8 Check model assumptions

The Poisson model is a classical generalized linear model, estimated using the method of maximum likelihood. While the default `plot` option for a `glm` still shows the plots we would use to assess the assumptions of an OLS model, we don't actually get much from that, since our Poisson model has different assumptions. It can be useful to look at a plot of residuals vs. fitted values on the original `physhealth` scale.

```
ggplot(sm_poiss1, aes(x = .fitted, y = .resid)) +
  geom_point() +
  labs(title = "Residuals vs. Fitted `physhealth`",
       subtitle = "Original Poisson Regression model")
```

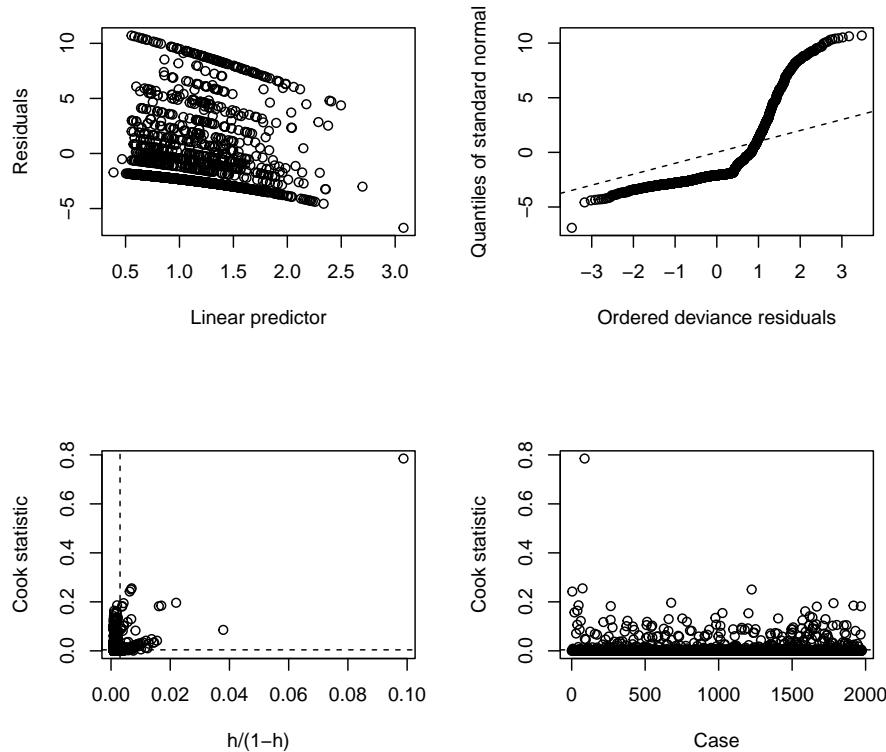


19.7.9 Using `glm.diag.plots` from the `boot` package

The `glm.diag.plots` function from the `boot` package makes a series of diagnostic plots for generalized linear models.

- (Top, Left) Jackknife deviance residuals against fitted values. This is essentially identical to what you obtain with `plot(mod_poiss1, which = 1)`. A *jackknife deviance* residual is also called a likelihood residual. It is the change in deviance when this observation is omitted from the data.
- (Top, Right) Normal Q-Q plot of standardized deviance residuals. (Dotted line shows expectation if those standardized residuals followed a Normal distribution, and these residuals generally should.) The result is similar to what you obtain with `plot(mod_poiss1, which = 2)`.
- (Bottom, Left) Cook statistic vs. standardized leverage
 - $n = \#$ of observations, $p = \#$ of parameters estimated
 - Horizontal dotted line is at $\frac{8}{n-2p}$. Points above the line have high influence on the model.
 - Vertical line is at $\frac{2p}{n-2p}$. Points to the right of the line have high leverage.
- (Bottom, Right) Index plot of Cook's statistic to help identify the observations with high influence. This is essentially the same plot as `plot(mod_poiss1, which = 4)`

```
glm.diag.plots(mod_poiss1)
```



When working with these plots, it is possible to use the `iden` command to perform some interactive identification of points in your R terminal. But that doesn't play out effectively in an HTML summary document like this, so we'll leave that out.

19.7.10 Predictions for Harry and Sally

The predictions from a `glm` fit like this don't include prediction intervals. But we can get predictions on the scale of our original response variable, `physhealth`, like this.

```
predict(mod_poiss1, newdata = hs_data, se.fit = TRUE,
       type = "response")
```

```
$fit
 1      2
```

```

5.989478 2.078688

$se.fit
      1          2
0.11544326 0.04314273

$residual.scale
[1] 1

```

By using `response` as the type, these predictions fall on the original `physhealth` scale. The prediction for Harry is now 5.99 days, and for Sally is 2.08 days.

19.8 Overdispersion in a Poisson Model

Poisson regressions do not supply an independent variance parameter σ , and as a result can be overdispersed, and usually are. Under the Poisson distribution, the variance equals the mean - so the standard deviation equals the square root of the mean. The notion of **overdispersion** arises here. When fitting generalized linear models with Poisson error distributions, the residual deviance and its degrees of freedom should be approximately equal if the model fits well.

If the residual deviance is far greater than the degrees of freedom, then overdispersion may well be a problem. In this case, the residual deviance is about 8.5 times the size of the residual degrees of freedom, so that's a clear indication of overdispersion. We saw earlier that the Poisson regression model requires that the outcome (here the `physhealth` counts) be independent. A possible reason for the overdispersion we see here is that `physhealth` on different days likely do not occur independently of one another but likely “cluster” together.

19.8.1 Testing for Overdispersion?

Gelman and Hill provide an overdispersion test in R for a Poisson model as follows...

```

yhat <- predict(mod_poiss1, type = "response")
n <- arm::display(mod_poiss1)$n

glm(formula = physhealth ~ bmi_c + smoke100, family = poisson(),
    data = sm_oh_A_young)
            coef.est  coef.se
(Intercept) 0.91      0.02
bmi_c        0.04      0.00
smoke100    0.53      0.02
---
n = 1974, k = 3

```

```

residual deviance = 19150.9, null deviance = 20221.8 (difference = 1070.9)
k <- arm::display(mod_poiss1)$k

glm(formula = physhealth ~ bmi_c + smoke100, family = poisson(),
    data = sm_oh_A_young)
    coef.est coef.se
(Intercept) 0.91      0.02
bmi_c        0.04      0.00
smoke100     0.53      0.02
---
n = 1974, k = 3
residual deviance = 19150.9, null deviance = 20221.8 (difference = 1070.9)

z <- (sm_oh_A_young$physhealth - yhat) / sqrt(yhat)
cat("overdispersion ratio is ", sum(z^2)/(n - k), "\n")

overdispersion ratio is 15.58261
cat("p value of overdispersion test is ",
    pchisq(sum(z^2), df = n-k, lower.tail = FALSE), "\n")

p value of overdispersion test is 0

```

The p value here is 0, indicating that the probability is essentially zero that a random variable from a χ^2 distribution with $(n - k) = 1971$ degrees of freedom would be as large as what we observed in this case. So there is significant overdispersion.

In summary, the `physhealth` counts are overdispersed by a factor of 15.581, which is enormous (even a factor of 2 would be considered large) and also highly statistically significant. The basic correction for overdispersion is to multiply all regression standard errors by $\sqrt{15.581} = 3.95$.

The `quasipoisson` model and the negative binomial model that we'll fit below are very similar. We write the overdispersed “quasiPoisson” model as:

$$y_i \sim \text{overdispersed Poisson}(\mu_i \exp(X_i \beta), \omega)$$

where ω is the overdispersion parameter, 15.581, in our case. The Poisson model we saw previously is then just the overdispersed Poisson model with $\omega = 1$.

19.9 Fitting the Quasi-Poisson Model

To deal with overdispersion, one useful approach is to apply a **quasi-likelihood estimation procedure**, as follows:

```

mod_poiss_od1 <- glm(physhealth ~ bmi_c + smoke100,
                      family = quasipoisson(),
                      data = sm_oh_A_young)

summary(mod_poiss_od1)

Call:
glm(formula = physhealth ~ bmi_c + smoke100, family = quasipoisson(),
     data = sm_oh_A_young)

Deviance Residuals:
    Min      1Q  Median      3Q      Max
-6.588 -2.604 -2.087 -0.533 10.688

Coefficients:
            Estimate Std. Error t value Pr(>|t|)
(Intercept) 0.906991  0.073818 12.287 < 2e-16 ***
bmi_c       0.035051  0.005611  6.247 5.11e-10 ***
smoke100   0.532505  0.098308  5.417 6.81e-08 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for quasipoisson family taken to be 15.58436)

Null deviance: 20222  on 1973  degrees of freedom
Residual deviance: 19151  on 1971  degrees of freedom
AIC: NA

Number of Fisher Scoring iterations: 6

confint(mod_poiss_od1)

Waiting for profiling to be done...

              2.5 %    97.5 %
(Intercept) 0.75877174 1.04831740
bmi_c       0.02383574 0.04583252
smoke100   0.34039368 0.72606948

```

This “quasi-Poisson regression” model uses the same mean function as Poisson regression, but now estimated by quasi-maximum likelihood estimation or, equivalently, through the method of generalized estimating equations, where the inference is adjusted by an estimated dispersion parameter. Sometimes, though I won’t demonstrate this here, people fit an “adjusted” Poisson regression model, where this estimation by quasi-ML is augmented to adjust the inference via

sandwich estimates of the covariances³.

19.9.1 The Fitted Equation

The model equation is still `log(physhealth) = 0.91 + 0.035 bmi_c + 0.53 smoke100`. The estimated coefficients are still statistically significant, but the standard errors for each coefficient are considerably larger when we account for overdispersion.

The dispersion parameter for the quasi-Poisson family is now taken to be a bit less than the square root of the ratio of the residual deviance and its degrees of freedom. This is a much more believable model, as a result.

19.9.2 Interpreting the Coefficients

No meaningful change from the Poisson model we saw previously.

19.9.3 Testing the Predictors

Again, we can use the Wald tests (z tests) provided with the Poisson regression output, or we can fit the model and then run an ANOVA to obtain a test based on the deviance (a simple transformation of the log likelihood ratio.)

- By the Wald tests shown above, each predictor clearly adds significant predictive value to the model given the other predictor, and we note that the *p* values are as small as R will support.
- The ANOVA approach for this model lets us check the impact of adding `smoke100` to a model already containing `bmi_c`.

```
anova(mod_poiss_od1, test = "Chisq")
```

Analysis of Deviance Table

Model: quasipoisson, link: log

Response: physhealth

Terms added sequentially (first to last)

	Df	Deviance	Resid. Df	Resid. Dev	Pr(>Chi)
NULL			1973	20222	
<code>bmi_c</code>	1	609.46	1972	19612	4.011e-10 ***

³See Zeileis A Kleiber C Jackman S *Regression Models for Count Data in R* Vignette at <https://cran.r-project.org/web/packages/pscl/vignettes/countreg.pdf>

```
smoke100 1 461.46 1971 19151 5.282e-08 ***
---
Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

The result is unchanged. To obtain a p value for `smoke100`'s impact after `bmi_c` is accounted for, we compare the difference in deviance to a chi-square distribution with 1 degree of freedom. The result is incredibly statistically significant.

To check the effect of `bmi_c`, we could refit the model with and without `bmi_c`, and again run an ANOVA. I'll skip that here.

19.9.4 Store fitted values and residuals

What happens if we try using the `broom` package in this case? We can, if we like, get our predicted values right on the scale of our `physhealth` response.

```
sm_poiss_od1 <- augment(mod_poiss_od1, sm_oh_A_young,
                         type.predict = "response")

sm_poiss_od1 %>%
  select(physhealth, .fitted) %>%
  head()

# A tibble: 6 x 2
  physhealth .fitted
  <dbl>     <dbl>
1 0         2.23
2 0         2.29
3 0         3.10
4 30        5.32
5 0         3.15
6 0         3.68
```

It turns out that 0 of the 1974 predictions that we make are below 0, and the largest prediction made by this model is 21.7 days.

The `rootogram` function we've shown doesn't support overdispersed Poisson models at the moment.

19.9.5 Specify the R^2 and log(likelihood) values

We can calculate the R^2 as the squared correlation of the fitted values and the observed values.

```
# The correlation of observed and fitted values
(poiss_od_r <- with(sm_poiss_od1, cor(physhealth, .fitted)))
```

```
[1] 0.1846814
```

```
# R-square
poiss_od_r^2
```

```
[1] 0.03410723
```

The `glance` function in the `broom` package gives us model log(likelihood), among other summaries.

```
glance(mod_poiss_od1) %>% round(3)
```

```
# A tibble: 1 x 8
  null.deviance df.null logLik   AIC   BIC deviance df.residual nobs
        <dbl>     <dbl>  <dbl> <dbl> <dbl>    <dbl>       <dbl> <dbl>
1      20222.     1973     NA     NA     NA    19151.     1971    1974
```

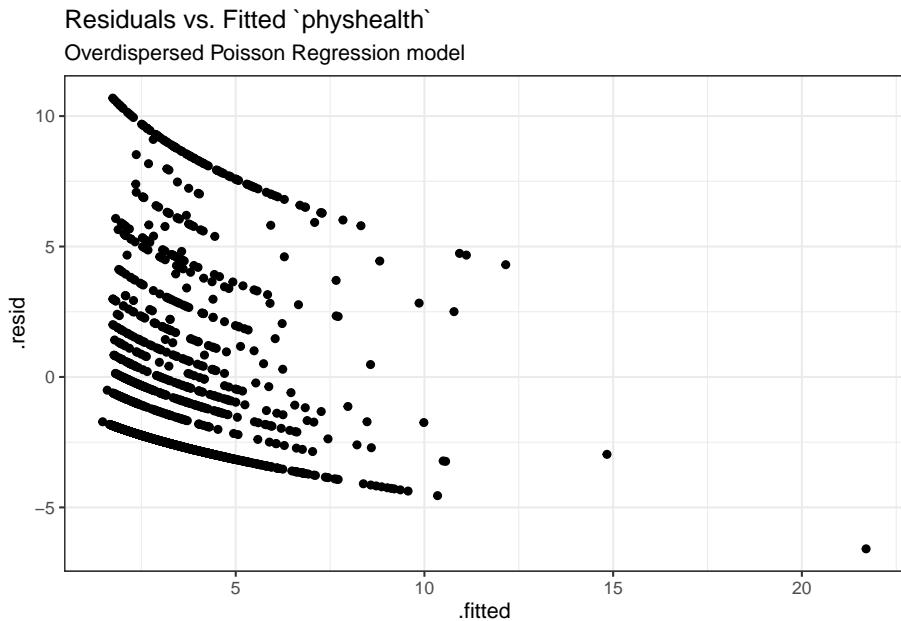
Here, we have

Model	Scale	R^2	log(likelihood)
Poisson	log(<code>physhealth</code>)	0.034	NA

19.9.6 Check model assumptions

Having dealt with the overdispersion, this should be a cleaner model in some ways, but the diagnostics (other than the dispersion) will be the same. Here is a plot of residuals vs. fitted values on the original `physhealth` scale.

```
ggplot(sm_poiss_od1, aes(x = .fitted, y = .resid)) +
  geom_point() +
  labs(title = "Residuals vs. Fitted `physhealth`",
       subtitle = "Overdispersed Poisson Regression model")
```



I'll skip the `glm.diag.plots` results, since you've already seen them.

19.9.7 Predictions for Harry and Sally

The predictions from this overdispersed Poisson regression will match those in the original Poisson regression, but the standard error will be larger.

```
predict(mod_poiss_od1, newdata = hs_data, se.fit = TRUE,
       type = "response")
```

```
$fit
 1          2
5.989478 2.078688

$se.fit
 1          2
0.4557357 0.1703147

$residual.scale
[1] 3.947703
```

By using `response` as the type, these predictions fall on the original `physhealth` scale. Again, the prediction for Harry is 5.99 days, and for Sally is 2.08 days.

19.10 Poisson and Quasi-Poisson models using `Glm` from the `rms` package

The `Glm` function in the `rms` package can be used to fit both the original Poisson regression and the quasi-Poisson model accounting for overdispersion.

19.10.1 Refitting the original Poisson regression with `Glm`

```
d <- datadist(sm_oh_A_young)
options(datadist = "d")

mod_poi_Glm_1 <- Glm(physhealth ~ bmi_c + smoke100,
                      family = poisson(),
                      data = sm_oh_A_young,
                      x = T, y = T)

mod_poi_Glm_1
```

General Linear Model

```
Glm(formula = physhealth ~ bmi_c + smoke100, family = poisson(),
     data = sm_oh_A_young, x = T, y = T)

              Model Likelihood
              Ratio Test
    Obs      1974   LR chi2    1070.93
    Residual d.f. 1971   d.f.        2
    g 0.4182128 Pr(> chi2) <0.0001

          Coef    S.E.   Wald Z Pr(>|Z|)
    Intercept 0.9070 0.0187 48.50 <0.0001
    bmi_c     0.0351 0.0014 24.66 <0.0001
    smoke100  0.5325 0.0249 21.38 <0.0001
```

19.10.2 Refitting the overdispersed Poisson regression with `Glm`

```
d <- datadist(sm_oh_A_young)
options(datadist = "d")

mod_poi_od_Glm_1 <- Glm(physhealth ~ bmi_c + smoke100,
```

```
family = quasipoisson(),
data = sm_oh_A_young,
x = T, y = T)
```

```
mod_poi_od_Glm_1
```

General Linear Model

```
Glm(formula = physhealth ~ bmi_c + smoke100, family = quasipoisson(),
     data = sm_oh_A_young, x = T, y = T)

              Model Likelihood
              Ratio Test
    Obs      1974    LR chi2    1070.93
    Residual d.f. 1971    d.f.          2
    g 0.4182128    Pr(> chi2) <0.0001

            Coef    S.E.   Wald Z Pr(>|Z|)
    Intercept 0.9070 0.0738 12.29 <0.0001
    bmi_c      0.0351 0.0056  6.25 <0.0001
    smoke100  0.5325 0.0983  5.42 <0.0001
```

The big advantage here is that we have access to the usual ANOVA, `summary`, and `nomogram` features that `rms` brings to fitting models.

19.10.3 ANOVA on a `Glm` fit

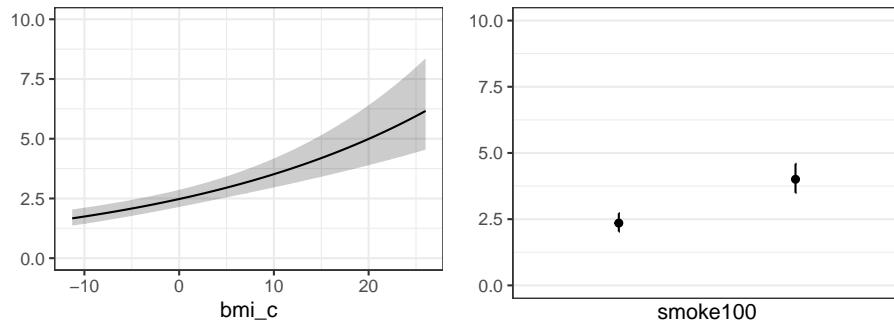
```
anova(mod_poi_od_Glm_1)
```

	Wald Statistics			Response: physhealth
Factor	Chi-Square	d.f.	P	
bmi_c	39.03	1	<.0001	
smoke100	29.34	1	<.0001	
TOTAL	74.39	2	<.0001	

This shows the individual Wald χ^2 tests without having to refit the model.

19.10.4 ggplots from `Glm` fit

```
ggplot(Predict(mod_poi_od_Glm_1, fun = exp))
```



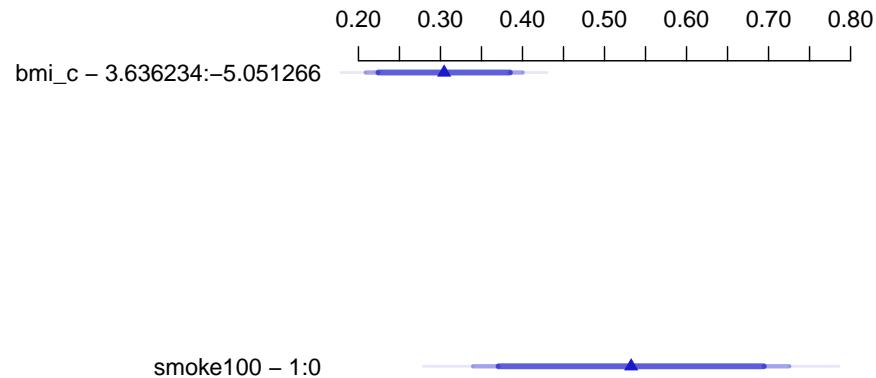
19.10.5 Summary of a `Glm` fit

```
summary(mod_poi_od_Glm_1)
```

	Effects			Response : physhealth					
Factor	Low	High	Diff.	Effect	S.E.	Lower	0.95	Upper	0.95
bmi_c	-5.0513	3.6362	8.6875	0.30450	0.048742	0.20891	0.4001		
smoke100	0.0000	1.0000	1.0000	0.53251	0.098308	0.33971	0.7253		

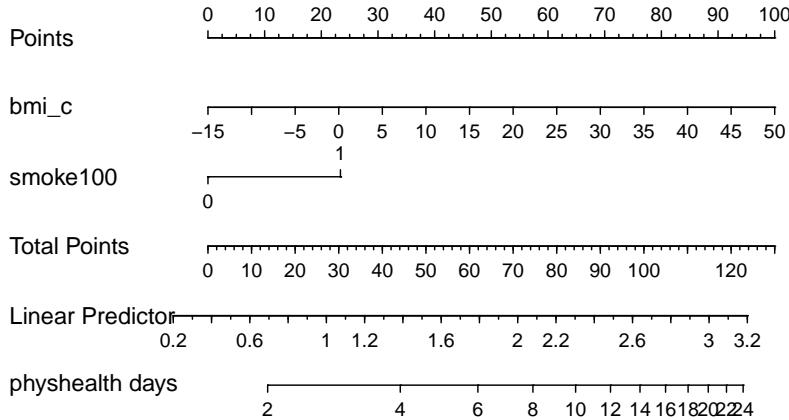
19.10.6 Plot of the Summary

```
plot(summary(mod_poi_od_Glm_1))
```



19.10.7 Nomogram of a `Glm` fit

```
plot(nomogram(mod_poi_od_Glm_1, fun = exp,  
              funlabel = "physhealth days"))
```



Note the use of `fun=exp` in both the ggplot of Predict and the nomogram. What's that doing?

19.11 Negative Binomial Model

Another approach to dealing with overdispersion is to fit a negative binomial model⁴ to predict the `log(physhealth)` counts. This involves the fitting of an additional parameter, θ . That's our dispersion parameter⁵

Sometimes, people will fit a model where θ is known, for instance a geometric model (where $\theta = 1$), and then this can be directly plugged into a `glm()` fit, but the more common scenario is that we are going to iteratively estimate the β coefficients and θ . To do this, I'll use the `glm.nb` function from the MASS package.

```
mod_nb1 <- MASS::glm.nb(physhealth ~ bmi_c + smoke100, link = log,
                         data = sm_oh_A_young)

summary(mod_nb1)
```

Call:

⁴See <https://cran.r-project.org/web/packages/pscl/vignettes/countreg.pdf> for more details.

⁵This θ is the inverse of the dispersion parameter estimated for these models by most other software packages, like SAS, Stata and SPSS. See <https://stats.idre.ucla.edu/r/dae/negative-binomial-regression/> for more details.

```

MASS::glm.nb(formula = physhealth ~ bmi_c + smoke100, data = sm_oh_A_young,
link = log, init.theta = 0.1487673114)

Deviance Residuals:
    Min      1Q   Median      3Q     Max
-1.2260 -0.9712 -0.8985 -0.1128  1.9929

Coefficients:
            Estimate Std. Error z value Pr(>|z|)
(Intercept) 0.874530  0.078994 11.071 < 2e-16 ***
bmi_c       0.035712  0.008317  4.294 1.76e-05 ***
smoke100    0.596396  0.121166  4.922 8.56e-07 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for Negative Binomial(0.1488) family taken to be 1)

Null deviance: 1468.1 on 1973 degrees of freedom
Residual deviance: 1422.6 on 1971 degrees of freedom
AIC: 6976.6

Number of Fisher Scoring iterations: 1

Theta:  0.14877
Std. Err.: 0.00705

2 x log-likelihood:  -6968.55300
confint(mod_nb1)

Waiting for profiling to be done...

          2.5 %    97.5 %
(Intercept) 0.72304817 1.03304660
bmi_c       0.02072601 0.05124925
smoke100    0.35977590 0.83584673

```

19.11.1 The Fitted Equation

The form of the model equation for a negative binomial regression is the same as that for Poisson regression.

```
log(physhealth) = 0.87 + 0.036 bmi_c + 0.60 smoke100
```

19.11.2 Comparison with the (raw) Poisson model

To compare the negative binomial model to the Poisson model (without the overdispersion) we can use the `logLik` function to make a comparison. Note that the Poisson model is a subset of the negative binomial.

```
logLik(mod_nb1)

'log Lik.' -3484.277 (df=4)
logLik(mod_poiss1)

'log Lik.' -10819.6 (df=3)
2 * (logLik(mod_nb1) - logLik(mod_poiss1))

'log Lik.' 14670.65 (df=4)
pchisq(2 * (logLik(mod_nb1) - logLik(mod_poiss1)), df = 1, lower.tail = FALSE)

'log Lik.' 0 (df=4)
```

Here, the difference in the log likelihoods is large enough that the resulting p value is very small. This strongly suggests that the negative binomial model, which adds the dispersion parameter, is more appropriate than the raw Poisson model.

However, both the regression coefficients and the standard errors are rather similar to the quasi-Poisson and the sandwich-adjusted Poisson results above. Thus, in terms of predicted means, all three models give very similar results; the associated Wald tests also lead to the same conclusions.

19.11.3 Interpreting the Coefficients

There's only a small change here from the Poisson models we saw previously.

- The constant term, 0.87, gives us the intercept of the regression - the prediction if `smoke100` = 0 and `bmi_c` = 0. In this case, because we've centered BMI, it implies that $\exp(0.87) = 2.39$ is the predicted days of poor `physhealth` for a non-smoker with average BMI.
- The coefficient of `bmi_c`, 0.036, is the expected difference in count of poor `physhealth` days (on the log scale) for each additional kg/m² of body mass index. The expected multiplicative *increase* is $e^{0.036} = 1.037$, corresponding to a 3.7% difference in the count.
- The coefficient of `smoke100`, 0.60, tells us that the predictive difference between those who have and who have not smoked 100 cigarettes can be found by multiplying the `physhealth` count by $\exp(0.6) = 1.82$, yielding essentially an 82% increase of the `physhealth` count.

19.11.4 Interpretation of Coefficients in terms of IRRs

We might be interested in looking at incident rate ratios rather than coefficients. The coefficients have an additive effect in the $\log(y)$ scale, and the IRR have a multiplicative effect in the y scale. To do this, we can exponentiate our model coefficients. This also applies to the confidence intervals.

```
exp(coef(mod_nb1))

(Intercept)      bmi_c      smoke100
2.397748     1.036357    1.815563

exp(confint(mod_nb1))

Waiting for profiling to be done...

          2.5 %   97.5 %
(Intercept) 2.060705 2.809613
bmi_c       1.020942 1.052585
smoke100    1.433008 2.306766
```

As an example, then, the incident rate for `smoke100 = 1` is 1.82 times the incident rate of `physhealth` days for the reference group (`smoke100 = 0`). The percent change in the incident rate of `physhealth` is a 3.6% increase for every kg/m^2 increase in centered `bmi`.

19.11.5 Testing the Predictors

Again, we can use the Wald tests (z tests) provided with the negative binomial regression output.

As an alternative, we probably should not use the standard `anova` process, because the models there don't re-estimate θ for each new model, as the warning message below indicates.

```
anova(mod_nb1)

Warning in anova.negbin(mod_nb1): tests made without re-estimating 'theta'
Analysis of Deviance Table

Model: Negative Binomial(0.1488), link: log

Response: physhealth

Terms added sequentially (first to last)

Df Deviance Resid. Df Resid. Dev Pr(>Chi)
NULL                      1973      1468.0
```

```
bmi_c      1   20.837      1972    1447.2 5.001e-06 ***
smoke100  1   24.584      1971    1422.6 7.115e-07 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

So, instead, if we want, for instance to assess the significance of `bmi_c`, after `smoke100` is already included in the model, we fit both models (with and without `bmi_c`) and then compare those models with a likelihood ratio test.

```
mod_nb1_without_bmi <- MASS::glm.nb(physhealth ~ smoke100,
                                      link = log,
                                      data = sm_oh_A_young)

anova(mod_nb1, mod_nb1_without_bmi)
```

Likelihood ratio tests of Negative Binomial Models

```
Response: physhealth
          Model     theta Resid. df   2 x log-lik.   Test   df LR stat.
1           smoke100 0.1452219    1972      -6991.195
2 bmi_c + smoke100 0.1487673    1971      -6968.553 1 vs 2      1  22.6421
Pr(Chi)
1
2 1.951615e-06
```

And we could compare the negative binomial models with and without `smoke100` in a similar way.

```
mod_nb1_without_smoke <- MASS::glm.nb(physhealth ~ bmi_c,
                                       link = log,
                                       data = sm_oh_A_young)

anova(mod_nb1, mod_nb1_without_smoke)
```

Likelihood ratio tests of Negative Binomial Models

```
Response: physhealth
          Model     theta Resid. df   2 x log-lik.   Test   df LR stat.
1           bmi_c 0.1449966    1972      -6992.839
2 bmi_c + smoke100 0.1487673    1971      -6968.553 1 vs 2      1  24.28569
Pr(Chi)
1
2 8.305388e-07
```

19.11.6 Store fitted values and residuals

The `broom` package works in this case, too. We'll look here at predicted (fitted) values on the scale of our `physhealth` response.

```
sm_nb1 <- augment(mod_nb1, sm_oh_A_young,
                    type.predict = "response")
```

Warning: Tidiers for objects of class `negbin` are not maintained by the `broom` team, and are only supported through the `glm` tidier method. Please be cautious in interpreting and reporting `broom` output.

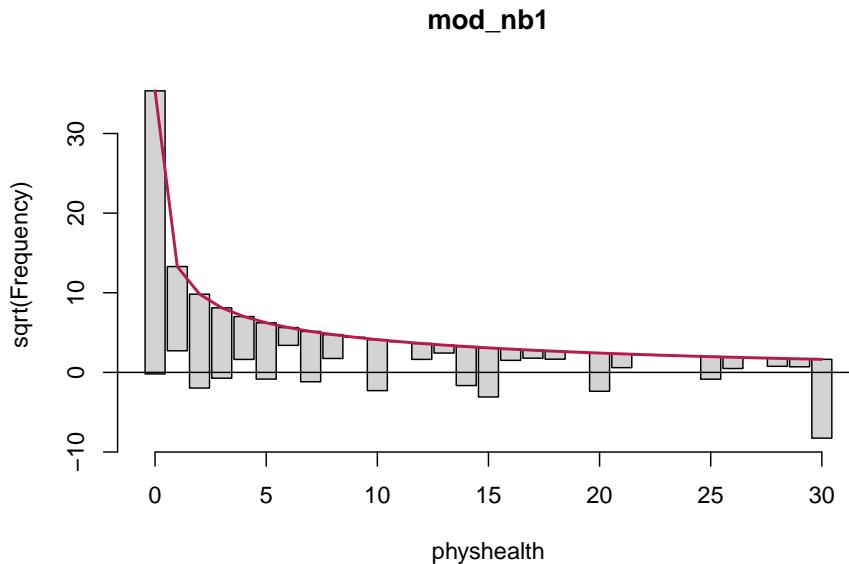
```
sm_nb1 %>%
  select(physhealth, .fitted) %>%
  head()
```

```
# A tibble: 6 x 2
  physhealth .fitted
  <dbl>     <dbl>
1 0         2.15
2 0         2.22
3 0         3.18
4 30        5.23
5 0         3.24
6 0         3.78
```

19.11.7 Rootogram for Negative Binomial model

Here's the rootogram for the negative binomial model.

```
rootogram(mod_nb1, max = 30)
```



Again, the red curved line is the theoretical (negative binomial) fit. “Hanging” from each point on the red line is a bar, the height of which represents the difference between expected and observed counts. A bar hanging below 0 indicates underfitting. A bar hanging above 0 indicates overfitting. The counts have been transformed with a square root transformation to prevent smaller counts from getting obscured and overwhelmed by larger counts.

The match looks much better than the Poisson model, which is a sign that accounting for overdispersion is very important. Even this model badly underfits the number of 30 values, however.

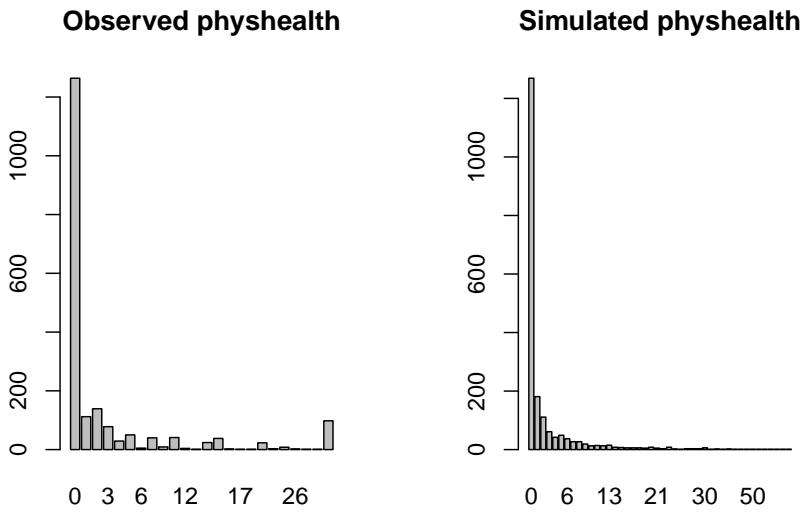
19.11.8 Simulating what the Negative Binomial model predicts

We can use the parameters of the negative binomial model to simulate data⁶ and compare the simulated results to our observed physhealth data.

```
par(mfrow=c(1,2))
sm_oh_A_young$physhealth %>%
  table() %>% barplot(main = "Observed physhealth")
set.seed(432122)
rnbinom(n = nrow(sm_oh_A_young),
        size = mod_nb1$theta,
```

⁶See <http://data.library.virginia.edu/getting-started-with-negative-binomial-regression-modeling/>

```
mu = exp(coef(mod_nb1)[1])) %>%
table() %>% barplot(main = "Simulated physhealth")
```



Again we see that the simulated data badly underfits the 30 values, and includes some predictions larger than 30.

19.11.9 Specify the R^2 and log(likelihood) values

We can calculate the R^2 as the squared correlation of the fitted values and the observed values.

```
# The correlation of observed and fitted values
(nb_r <- with(sm_nb1, cor(physhealth, .fitted)))
```

```
[1] 0.183675
```

```
# R-square
nb_r^2
```

```
[1] 0.03373649
```

The `glance` function in the `broom` package gives us model `log(likelihood)`, among other summaries.

```
glance(mod_nb1) %>% round(3)
```

```
# A tibble: 1 x 8
```

```
null.deviance df.null logLik      AIC    BIC deviance df.residual nobs
<dbl>     <dbl> <logLik> <dbl> <dbl>     <dbl>     <dbl> <dbl>
1        1468.    1973 -3484.277 6977. 6999.    1423.     1971 1974
```

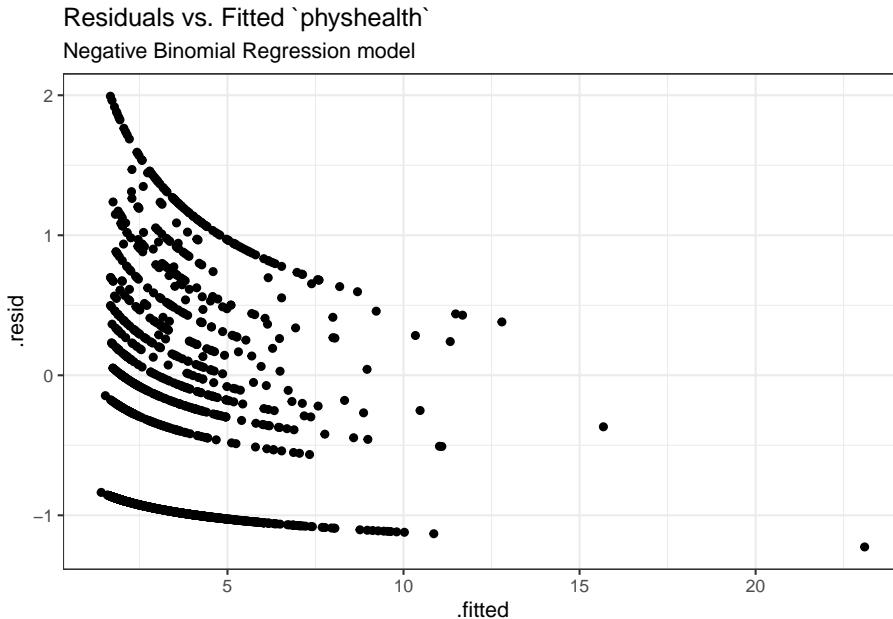
Here, we have

	Model	Scale	R^2	log(likelihood)
	Negative Binomial	log(physhealth)	.034	-3484.27

19.11.10 Check model assumptions

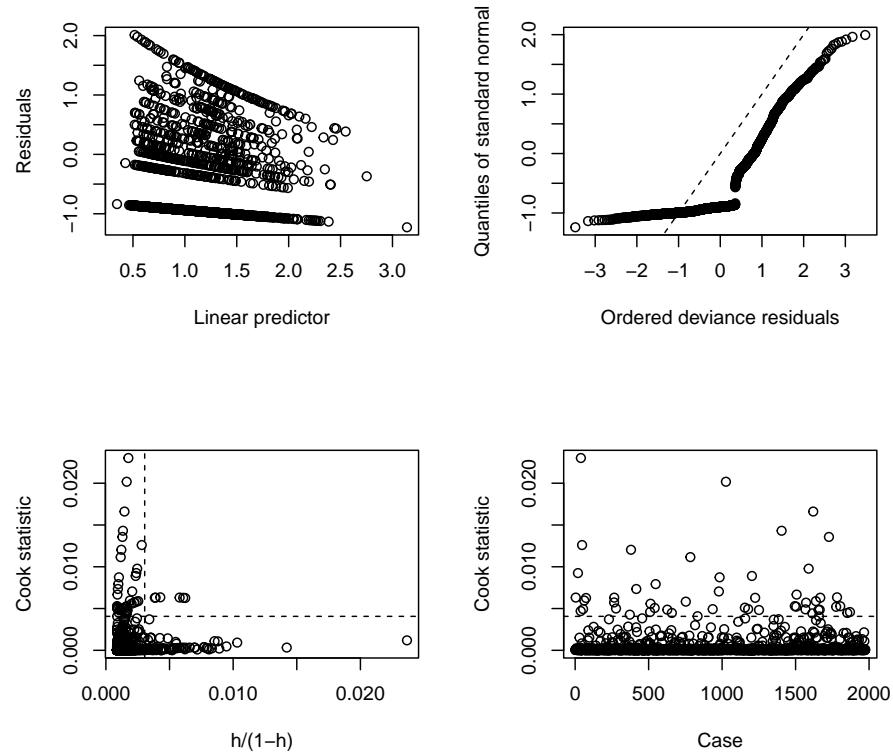
Here is a plot of residuals vs. fitted values on the original `physhealth` scale.

```
ggplot(sm_nb1, aes(x = .fitted, y = .resid)) +
  geom_point() +
  labs(title = "Residuals vs. Fitted `physhealth`",
       subtitle = "Negative Binomial Regression model")
```



Here are the `glm` diagnostic plots from the `boot` package.

```
glm.diag.plots(mod_nb1)
```



From the lower left plot, we see fewer points with large values of both Cook's distance and leverage, so that's a step in the right direction. The upper right plot still has some issues, but we're closer to a desirable result there, too.

19.11.11 Predictions for Harry and Sally

The predictions from this negative binomial regression model will be only a little different than those from the Poisson models.

```
predict(mod_nb1, newdata = hs_data, se.fit = TRUE,
       type = "response")
```

```
$fit
 1      2
6.221696 2.005657

$se.fit
 1      2
```

```
0.7537535 0.1773595
```

```
$residual.scale
[1] 1
```

As we've seen in the past, when we use `response` as the type, the predictions fall on the original `physhealth` scale. The prediction for Harry is 6.2 days, and for Sally is 2.0 days.

19.12 The Problem: Too Few Zeros

Remember that we observe more than 1000 zeros in our `physhealth` data.

```
sm_oh_A_young %>% count(physhealth == 0)
```

```
# A tibble: 2 x 2
  `physhealth == 0`     n
  <lg1>                 <int>
1 FALSE                  710
2 TRUE                   1264
```

Let's go back to our Poisson model (without overdispersion) for a moment, and concentrate on the zero values.

```
# predict expected mean physhealth for each subject
mu <- predict(mod_poiss1, type = "response")

# sum the probabilities of a zero count for each mean
exp <- sum(dpois(x = 0, lambda = mu))

# predicted number of zeros from Poisson model
round(exp)
```

```
[1] 124
```

As we've seen previously, we're severely underfitting zero counts. We can compare the observed number of zero `physhealth` results to the expected number of zero values from the likelihood-based models.

```
round(c("Obs" = sum(sm_oh_A_young$physhealth == 0),
      "Poisson" = sum(dpois(0, fitted(mod_poiss1))),
      "NB" = sum(dnbinom(0, mu = fitted(mod_nb1), size = mod_nb1$theta))),0)
```

Obs	Poisson	NB
1264	124	1250

There are at least two ways to tackle this problem.

- Fitting a model which deliberately inflates the number of zeros that are fitted
- Fitting a hurdle model

We'll look at those options, next.

19.13 The Zero-Inflated Poisson Regression Model

The zero-inflated Poisson or (ZIP) model is used to describe count data with an excess of zero counts⁷. The model posits that there are two processes involved:

- a logit model is used to predict excess zeros
- while a Poisson model is used to predict the counts, generally

The `pscl` package is used here, which can conflict with the `countreg` package we used to fit rootograms. That's why I'm loading it here.

```
library(pscl)
```

To run the zero-inflated Poisson model, we use the following:

```
mod_zip1 <- pscl::zeroinfl(physhealth ~ bmi_c + smoke100,
                           data = sm_oh_A_young)

summary(mod_zip1)
```

Call:

```
pscl::zeroinfl(formula = physhealth ~ bmi_c + smoke100, data = sm_oh_A_young)
```

Pearson residuals:

Min	1Q	Median	3Q	Max
-1.4389	-0.6987	-0.6138	-0.1947	9.3036

Count model coefficients (poisson with log link):

	Estimate	Std. Error	z value	Pr(> z)
(Intercept)	1.997594	0.018797	106.3	<2e-16 ***
bmi_c	0.018176	0.001398	13.0	<2e-16 ***
smoke100	0.393383	0.024889	15.8	<2e-16 ***

Zero-inflation model coefficients (binomial with logit link):

	Estimate	Std. Error	z value	Pr(> z)
(Intercept)	0.682427	0.062853	10.857	< 2e-16 ***
bmi_c	-0.027731	0.006509	-4.260	2.04e-05 ***
smoke100	-0.237224	0.095318	-2.489	0.0128 *

⁷See <https://stats.idre.ucla.edu/r/dae/zip/> for more on the zero-inflated poisson model.

```
---
Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Number of iterations in BFGS optimization: 8
Log-likelihood: -5680 on 6 Df
confint(mod_zip1)

              2.5 %    97.5 %
count_(Intercept) 1.96075280 2.03443546
count_bmi_c       0.01543576 0.02091718
count_smoke100    0.34460096 0.44216461
zero_(Intercept)  0.55923725 0.80561744
zero_bmi_c        -0.04048828 -0.01497359
zero_smoke100     -0.42404310 -0.05040515
```

The output describes two separate regression models. Below the model call, we see information on a Poisson regression model. Then we see another block describing the inflation model.

Each predictor (`bmi_c` and `smoke100`) appears to be statistically significant in each part of the model.

19.13.1 Comparison to a null model

To show that this model fits better than the null model (the model with intercept only), we can compare them directly with a chi-squared test. Since we have two predictors in the full model, the degrees of freedom for this test is 2.

```
mod_zipnull <- pscl::zeroinfl(physhealth ~ 1,
                               data = sm_oh_A_young)

summary(mod_zipnull)

Call:
pscl::zeroinfl(formula = physhealth ~ 1, data = sm_oh_A_young)

Pearson residuals:
      Min      1Q Median      3Q      Max
-0.6934 -0.6934 -0.6934 -0.2779  5.5400

Count model coefficients (poisson with log link):
Estimate Std. Error z value Pr(>|z|)
(Intercept) 2.22764   0.01233 180.7 <2e-16 ***

Zero-inflation model coefficients (binomial with logit link):
Estimate Std. Error z value Pr(>|z|)
```

```
(Intercept) 0.5767     0.0469    12.29    <2e-16 ***
---
Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Number of iterations in BFGS optimization: 2
Log-likelihood: -5908 on 2 Df
pchisq(2 * (logLik(mod_zip1) - logLik(mod_zipnull)), df = 2, lower.tail = FALSE)

'log Lik.' 8.59676e-100 (df=6)
```

19.13.2 Comparison to a Poisson Model with the Vuong test

```
vuong(mod_zip1, mod_poiss1)
```

```
Vuong Non-Nested Hypothesis Test-Statistic:
(test-statistic is asymptotically distributed N(0,1) under the
null that the models are indistinguishable)
-----
Vuong z-statistic      H_A      p-value
Raw                  19.90843 model1 > model2 < 2.22e-16
AIC-corrected       19.89681 model1 > model2 < 2.22e-16
BIC-corrected       19.86434 model1 > model2 < 2.22e-16
```

Certainly, the ZIP model is a significant improvement over the standard Poisson model, as the Vuong test reveals.

19.13.3 The Fitted Equation

The form of the model equation for a zero-inflated Poisson regression requires us to take two separate models into account. First we have a logistic regression model to predict the log odds of zero `physhealth` days. That takes care of the *extra* zeros. Then, to predict the number of `physhealth` days, we have a Poisson model, which may produce some additional zero count estimates.

19.13.4 Interpreting the Coefficients

We can exponentiate the logistic regression coefficients to obtain results in terms of odds ratios for that model, and that can be of some help in understanding the process behind excess zeros.

Also, exponentiating the coefficients of the count model help us describe those counts on the original scale of `physhealth`.

```
exp(coef(mod_zip1))

count_(Intercept)      count_bmi_c      count_smoke100  zero_(Intercept)
    7.3713003        1.0183427       1.4819856        1.9786748
zero_bmi_c            zero_smoke100
    0.9726500        0.7888145
```

For example,

- in the model for `physhealth = 0`, the odds of `physhealth = 0` are 79% as high for subjects with `smoke100 = 1` as for non-smokers with the same BMI.
- in the Poisson model for `physhealth`, the `physhealth` count is estimated to increase by 1.48 for smokers as compared to non-smokers with the same BMI.

19.13.5 Testing the Predictors

We can test the model with and without `bmi_c`, for example, by fitting the model both ways, and comparing the results with either a Wald or Likelihood Ratio test, each of which is available in the `lmtest` package.

```
mod_zip1_nobmi <- pscl::zeroinfl(physhealth ~ smoke100,
                                    data = sm_oh_A_young)

lmtest::waldtest(mod_zip1, mod_zip1_nobmi)

Wald test

Model 1: physhealth ~ bmi_c + smoke100
Model 2: physhealth ~ smoke100
  Res.Df Df Chisq Pr(>Chisq)
1     1968
2     1970 -2 187.2 < 2.2e-16 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

lmtest::lrtest(mod_zip1, mod_zip1_nobmi)

Likelihood ratio test

Model 1: physhealth ~ bmi_c + smoke100
Model 2: physhealth ~ smoke100
  #Df LogLik Df Chisq Pr(>Chisq)
1   6 -5679.8
2   4 -5769.5 -2 179.35 < 2.2e-16 ***
```

```
Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

19.13.6 Store fitted values and residuals

The `broom` package does not work with the `zeroinfl` tool. So we need to build up the fitted values and residuals ourselves.

```
sm_zip1 <- sm_oh_A_young %>%
  mutate(fitted = fitted(mod_zip1, type = "response"),
        resid = resid(mod_zip1, type = "response"))

sm_zip1 %>%
  dplyr::select(physhealth, fitted, resid) %>%
  head()

# A tibble: 6 x 3
  physhealth fitted resid
  <dbl>     <dbl> <dbl>
1 0         2.21 -2.21
2 0         2.28 -2.28
3 0         3.12 -3.12
4 30       5.27  24.7 
5 0         3.17 -3.17
6 0         3.71 -3.71
```

19.13.7 Modeled Number of Zero Counts

The zero-inflated model is designed to perfectly match the number of observed zeros. We can compare the observed number of zero `physhealth` results to the expected number of zero values from the likelihood-based models.

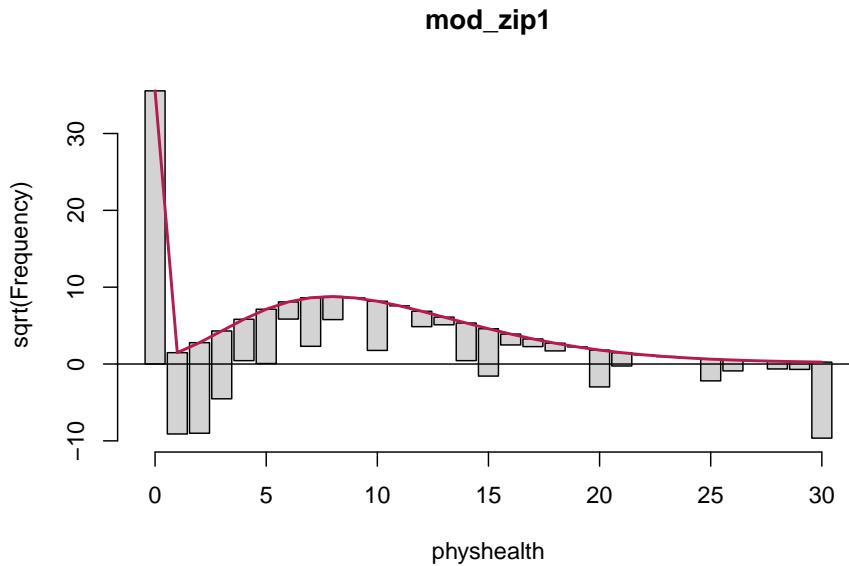
```
round(c("Obs" = sum(sm_oh_A_young$physhealth == 0),
      "Poisson" = sum(dpois(0, fitted(mod_poiss1))),
      "NB" = sum(dnbinom(0, mu = fitted(mod_nb1), size = mod_nb1$theta)),
      "ZIP" = sum(predict(mod_zip1, type = "prob")[,1])), 0)
```

Obs	Poisson	NB	ZIP
1264	124	1250	1264

19.13.8 Rootogram for ZIP model

Here's the rootogram for the zero-inflated Poisson model.

```
countreg::rootogram(mod_zip1, max = 30)
```



The zero frequencies are perfectly matched here, but we can see that counts of 1 and 2 are now substantially underfit, and values between 6 and 13 are overfit.

19.13.9 Specify the R^2 and log (likelihood) values

We can calculate a proxy for R^2 as the squared correlation of the fitted values and the observed values.

```
# The correlation of observed and fitted values
(zip_r <- with(sm_zip1, cor(physhealth, fitted)))
```

```
[1] 0.187328
```

```
# R-square
zip_r^2
```

```
[1] 0.03509178
```

```
logLik(mod_zip1)
```

```
'log Lik.' -5679.794 (df=6)
```

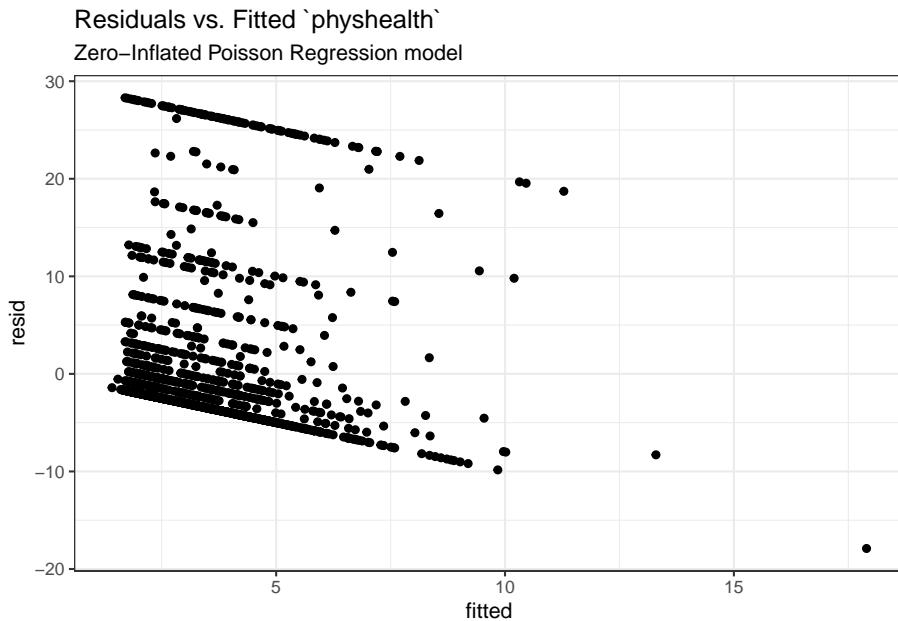
Here, we have

Model	Scale	R^2	log(likelihood)
Zero-Inflated Poisson	Complex: log(physhealth)	.035	-5679.83

19.13.10 Check model assumptions

Here is a plot of residuals vs. fitted values on the original physhealth scale.

```
ggplot(sm_zip1, aes(x = fitted, y = resid)) +
  geom_point() +
  labs(title = "Residuals vs. Fitted `physhealth`",
       subtitle = "Zero-Inflated Poisson Regression model")
```



19.13.11 Predictions for Harry and Sally

The predictions from this ZIP regression model are obtained as follows...

```
predict(mod_zip1, newdata = hs_data, type = "response")
```

1	2
6.002212	2.056525

As we've seen in the past, when we use `response` as the type, the predictions fall on the original physhealth scale. The prediction for Harry is 6.0 days, and for Sally is 2.1 days.

19.14 The Zero-Inflated Negative Binomial Regression Model

As an alternative to the ZIP model, we might consider a zero-inflated negative binomial regression⁸. This will involve a logistic regression to predict the probability of a 0, and then a negative binomial model to describe the counts of physhealth.

To run the zero-inflated negative binomial model, we use the following code:

```
mod_zinb1 <- pscl::zeroinfl(physhealth ~ bmi_c + smoke100,
                             dist = "negbin", data = sm_oh_A_young)

summary(mod_zinb1)

Call:
pscl::zeroinfl(formula = physhealth ~ bmi_c + smoke100, data = sm_oh_A_young,
dist = "negbin")

Pearson residuals:
    Min     1Q   Median     3Q     Max 
-0.5579 -0.4192 -0.3957 -0.1166  6.4358 

Count model coefficients (negbin with log link):
            Estimate Std. Error z value Pr(>|z|)    
(Intercept) 1.545599  0.099781 15.490 < 2e-16 ***
bmi_c       0.024614  0.006641  3.707  0.00021 ***
smoke100    0.517559  0.110772  4.672 2.98e-06 ***
Log(theta) -0.874034  0.143876 -6.075 1.24e-09 ***

Zero-inflation model coefficients (binomial with logit link):
            Estimate Std. Error z value Pr(>|z|)    
(Intercept) -0.071629  0.164039 -0.437  0.66236  
bmi_c        -0.027607  0.009313 -2.964  0.00303 ** 
smoke100     -0.127004  0.137171 -0.926  0.35451  
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Theta = 0.4173
Number of iterations in BFGS optimization: 14
Log-likelihood: -3469 on 7 Df
confint(mod_zinb1)
```

2.5 % 97.5 %

⁸See <https://stats.idre.ucla.edu/r/dae/zinb/>

```

count_(Intercept) 1.35003234 1.741165036
count_bmi_c 0.01159892 0.037629890
count_smoke100 0.30045008 0.734668671
zero_(Intercept) -0.39313901 0.249880357
zero_bmi_c -0.04585974 -0.009354261
zero_smoke100 -0.39585468 0.141845790

```

19.14.1 Comparison to a null model

To show that this model fits better than the null model (the model with intercept only), we can compare them directly with a chi-squared test. Since we have two predictors in the full model, the degrees of freedom for this test is 2.

```

mod_zinbnull <- pscl::zeroinfl(physhealth ~ 1, dist = "negbin",
                                data = sm_oh_A_young)

summary(mod_zinbnull)

Call:
pscl::zeroinfl(formula = physhealth ~ 1, data = sm_oh_A_young, dist = "negbin")

Pearson residuals:
    Min      1Q   Median      3Q      Max 
-0.4048 -0.4048 -0.4048 -0.1622  3.2340 

Count model coefficients (negbin with log link):
            Estimate Std. Error z value Pr(>|z|)    
(Intercept) 1.7766     0.0964 18.429 < 2e-16 ***
Log(theta) -1.0445     0.1612 -6.479 9.25e-11 ***

Zero-inflation model coefficients (binomial with logit link):
            Estimate Std. Error z value Pr(>|z|)    
(Intercept) -0.2605     0.1920 -1.357   0.175    
---
Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Theta = 0.3519
Number of iterations in BFGS optimization: 12
Log-likelihood: -3498 on 3 Df
pchisq(2 * (logLik(mod_nb1) - logLik(mod_zinbnull)), df = 2, lower.tail = FALSE)

'log Lik.' 8.538917e-07 (df=4)

```

19.14.2 Comparison to a Negative Binomial Model: Vuong test

```
vuong(mod_zinb1, mod_nb1)
```

Vuong Non-Nested Hypothesis Test-Statistic:
 (test-statistic is asymptotically distributed $N(0,1)$ under the
 null that the models are indistinguishable)

	Vuong z-statistic	H_A	p-value
Raw	3.0753499	model1 > model2	0.0010513
AIC-corrected	2.4597116	model1 > model2	0.0069524
BIC-corrected	0.7396743	model1 > model2	0.2297488

The zero-inflated negative binomial model is a significant improvement over the standard negative binomial model according to the raw or AIC-corrected Vuong tests, but not according to the BIC-corrected test.

19.14.3 The Fitted Equation

Like the ZIP, the zero-inflated negative binomial regression also requires us to take two separate models into account. First we have a logistic regression model to predict the log odds of zero `physhealth` days. That takes care of the *extra* zeros. Then, to predict the number of `physhealth` days, we have a negative binomial regression, with a θ term, and this negative binomial regression model may also produce some additional zero count estimates.

19.14.4 Interpreting the Coefficients

As with the zip, we can exponentiate the logistic regression coefficients to obtain results in terms of odds ratios for that model, and that can be of some help in understanding the process behind excess zeros.

```
exp(coef(mod_zinb1))
```

count_(Intercept)	count_bmi_c	count_smoke100	zero_(Intercept)
4.6907791	1.0249198	1.6779275	0.9308759
zero_bmi_c	zero_smoke100		
0.9727706	0.8807298		

For example,

- in the model for `physhealth` = 0, the odds of `physhealth` = 0 are 88.1% as high for subjects with `smoke100` = 1 as for non-smokers with the same BMI.

Interpreting the negative binomial piece works the same way as it did in the negative binomial regression.

19.14.5 Testing the Predictors

We can test the model with and without `bmi_c`, for example, by fitting the model both ways, and comparing the results with either a Wald or Likelihood Ratio test, each of which is available in the `lmtest` package.

```
mod_zinb1_nobmi <- pscl::zeroinfl(physhealth ~ smoke100,
                                    dist = "negbin",
                                    data = sm_oh_A_young)

lmtest::waldtest(mod_zinb1, mod_zinb1_nobmi)
```

Wald test

```
Model 1: physhealth ~ bmi_c + smoke100
Model 2: physhealth ~ smoke100
  Res.Df Df  Chisq Pr(>Chisq)
1     1967
2     1969 -2 29.589  3.757e-07 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
lmtest::lrtest(mod_zinb1, mod_zinb1_nobmi)
```

Likelihood ratio test

```
Model 1: physhealth ~ bmi_c + smoke100
Model 2: physhealth ~ smoke100
  #Df LogLik Df  Chisq Pr(>Chisq)
1    7 -3469.3
2    5 -3485.0 -2 31.418  1.506e-07 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

19.14.6 Store fitted values and residuals

Again, we need to build up the fitted values and residuals without the `broom` package.

```
sm_zinb1 <- sm_oh_A_young %>%
  mutate(fitted = fitted(mod_zinb1, type = "response"),
        resid = resid(mod_zinb1, type = "response"))
```

```
sm_zip1 %>%
  dplyr::select(physhealth, fitted, resid) %>%
  head()

# A tibble: 6 x 3
  physhealth fitted resid
  <dbl>     <dbl> <dbl>
1 0         2.21 -2.21
2 0         2.28 -2.28
3 0         3.12 -3.12
4 30        5.27 24.7 
5 0         3.17 -3.17
6 0         3.71 -3.71
```

19.14.7 Modeled Number of Zero Counts

Once again, we can compare the observed number of zero `physhealth` results to the expected number of zero values from the likelihood-based models.

```
round(c("Obs" = sum(sm_oh_A_young$physhealth == 0),
  "Poisson" = sum(dpois(0, fitted(mod_poiss1))),
  "NB" = sum(dnbinom(0, mu = fitted(mod_nb1), size = mod_nb1$theta)),
  "ZIP" = sum(predict(mod_zip1, type = "prob")[,1]),
  "ZINB" = sum(predict(mod_zinb1, type = "prob")[,1])),0)
```

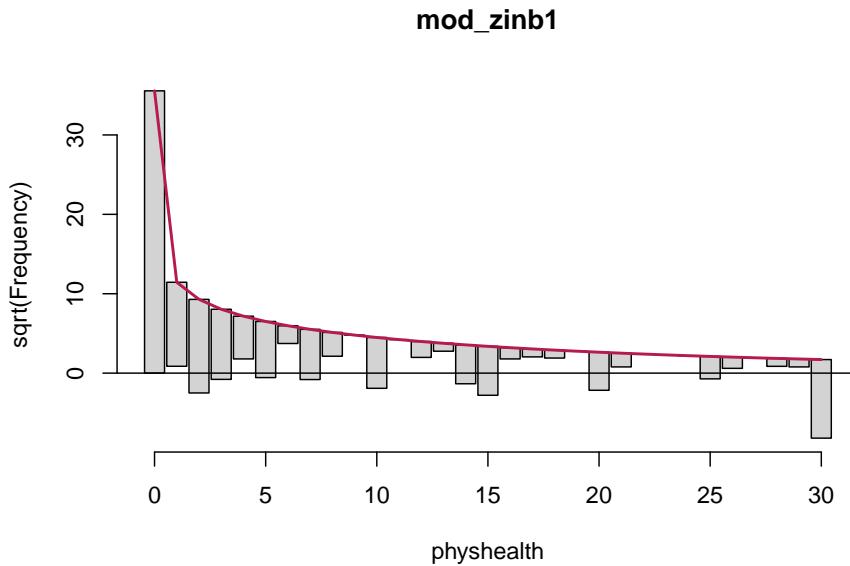
Obs	Poisson	NB	ZIP	ZINB
1264	124	1250	1264	1264

So, the Poisson model is clearly inappropriate, but the zero-inflated (Poisson and NB) and the negative binomial model all give reasonable fits in this regard.

19.14.8 Rootogram for Zero-Inflated Negative Binomial model

Here's the rootogram for the zero-inflated negative binomial model.

```
countreg::rootogram(mod_zinb1, max = 30)
```



As in the ZIP model, the zero frequencies are perfectly matched here, but we can see that counts of 1 and 2 are now closer to the data we observe than in the ZIP model. We are still substantially underfitting values of 30.

19.14.9 Specify the R^2 and log (likelihood) values

We can calculate a proxy for R^2 as the squared correlation of the fitted values and the observed values.

```
# The correlation of observed and fitted values
(zinb_r <- with(sm_zinb1, cor(physhealth, fitted)))
```

```
[1] 0.1858969
```

```
# R-square
zinb_r^2
```

```
[1] 0.03455767
```

```
logLik(mod_zinb1)
```

```
'log Lik.' -3469.29 (df=7)
```

Here, we have

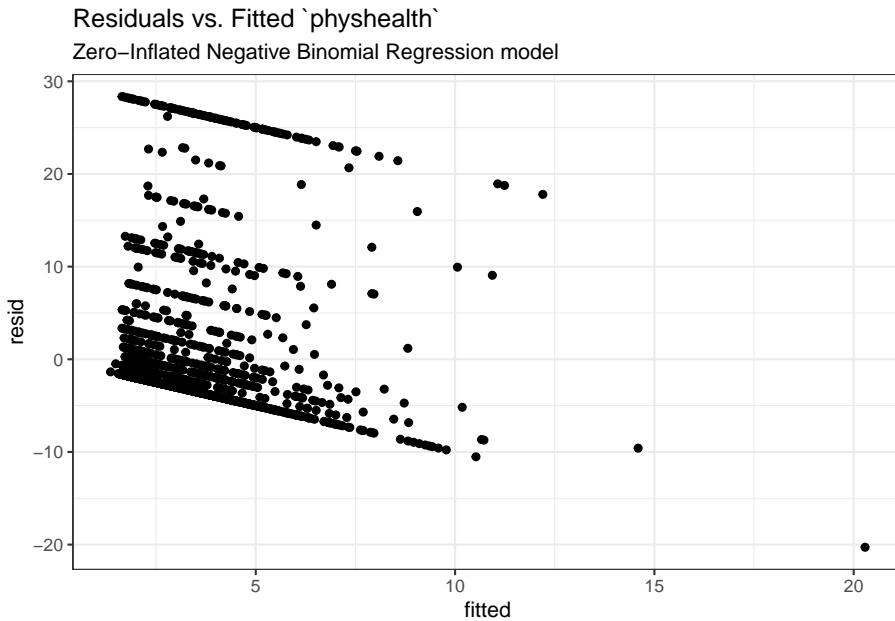
Model	Scale	R^2	log(likelihood)
Zero-Inflated Negative Binomial	Complex: <code>log(physhealth)</code>	.035	-3469.27

Model	Scale	R^2	log(likelihood)
-------	-------	-------	-----------------

19.14.10 Check model assumptions

Here is a plot of residuals vs. fitted values on the original physhealth scale.

```
ggplot(sm_zinb1, aes(x = fitted, y = resid)) +
  geom_point() +
  labs(title = "Residuals vs. Fitted `physhealth`",
       subtitle = "Zero-Inflated Negative Binomial Regression model")
```



19.14.11 Predictions for Harry and Sally

The predictions from this zero-inflated negative binomial regression model are obtained as follows...

```
predict(mod_zinb1, newdata = hs_data, type = "response")
```

```
1      2
6.206514 2.004963
```

As we've seen in the past, when we use `response` as the type, the predictions fall on the original `physhealth` scale. The prediction for Harry is 6.2 days, and for Sally is 2.0 days.

19.15 A “hurdle” model (with Poisson)

Much of the discussion here of hurdle models comes from Clay Ford at the University of Virginia⁹. Ford describes a hurdle model as follows:

The hurdle model is a two-part model that specifies one process for zero counts and another process for positive counts. The idea is that positive counts occur once a threshold is crossed, or put another way, a hurdle is cleared. If the hurdle is not cleared, then we have a count of 0.

The first part of the model is typically a binary logit model. This models whether an observation takes a positive count or not. The second part of the model is usually a truncated Poisson or Negative Binomial model. Truncated means we’re only fitting positive counts. If we were to fit a hurdle model to our [medicare] data, the interpretation would be that one process governs whether a patient visits a doctor or not, and another process governs how many visits are made.

To fit a hurdle model, we’ll use the `hurdle` function in the `pscl` package.

```
mod_hur1 <- pscl::hurdle(physhealth ~ bmi_c + smoke100,
                           dist = "poisson", zero.dist = "binomial",
                           data = sm_oh_A_young)

summary(mod_hur1)

Call:
pscl::hurdle(formula = physhealth ~ bmi_c + smoke100, data = sm_oh_A_young,
              dist = "poisson", zero.dist = "binomial")

Pearson residuals:
    Min      1Q  Median      3Q      Max
-1.4403 -0.6987 -0.6139 -0.1946  9.2997

Count model coefficients (truncated poisson with log link):
            Estimate Std. Error z value Pr(>|z|)
(Intercept) 1.997612   0.018797 106.3 <2e-16 ***
bmi_c       0.018178   0.001398   13.0 <2e-16 ***
```

⁹<http://data.library.virginia.edu/getting-started-with-hurdle-models/> is an excellent introduction, by Clay Ford, a Statistical Research Consultant at the University of Virginia Library. I can also recommend https://rpubs.com/kaz_yos/pscl-2 as a place to learn more about the `pscl` package, and the fitting and interpretation of both hurdle and zero-inflated regression models. That `rpubs` site has a link to this article by Hu, Pavlicova and Nunes from the Am J Drug Alcohol Abuse which provides a real set of examples from a trial of a behavioral health intervention meant to reduce the risk of unprotected sexual occasions as part of a strategy to reduce HIV risk.

```

smoke100    0.393348  0.024889   15.8   <2e-16 ***
Zero hurdle model coefficients (binomial with logit link):
  Estimate Std. Error z value Pr(>|z|)
(Intercept) -0.683509  0.062827 -10.879 < 2e-16 ***
bmi_c        0.027777  0.006508   4.268 1.97e-05 ***
smoke100     0.238300  0.095301   2.500   0.0124 *
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

```

Number of iterations in BFGS optimization: 14

Log-likelihood: -5680 on 6 Df

```
confint(mod_hur1)
```

	2.5 %	97.5 %
count_(Intercept)	1.96077105	2.03445261
count_bmi_c	0.01543722	0.02091901
count_smoke100	0.34456599	0.44212932
zero_(Intercept)	-0.80664805	-0.56037031
zero_bmi_c	0.01502124	0.04053230
zero_smoke100	0.05151314	0.42508607

We are using the default settings here, using the same predictors for both models:

- a **Binomial** model to predict the probability of `physhealth` = 0 given our predictors, as specified by the `zero.dist` argument in the `hurdle` function, and
- a (truncated) **Poisson** model to predict the positive-count of `physhealth` given those same predictors, as specified by the `dist` argument in the `hurdle` function.

19.15.1 Comparison to a null model

To show that this model fits better than the null model (the model with intercept only), we can compare them directly with a chi-squared test. Since we have two predictors in the full model, the degrees of freedom for this test is 2.

```

mod_hurnull <- pscl::hurdle(physhealth ~ 1, dist = "poisson",
                             zero.dist = "binomial",
                             data = sm_oh_A_young)

summary(mod_hurnull)

Call:
pscl::hurdle(formula = physhealth ~ 1, data = sm_oh_A_young, dist = "poisson",
              zero.dist = "binomial")

```

```

Pearson residuals:
    Min      1Q Median      3Q     Max
-0.6934 -0.6934 -0.6934 -0.2779  5.5399

Count model coefficients (truncated poisson with log link):
    Estimate Std. Error z value Pr(>|z|)
(Intercept) 2.22765   0.01233 180.7 <2e-16 ***
Zero hurdle model coefficients (binomial with logit link):
    Estimate Std. Error z value Pr(>|z|)
(Intercept) -0.5768    0.0469  -12.3 <2e-16 ***
---
Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Number of iterations in BFGS optimization: 8
Log-likelihood: -5908 on 2 Df
pchisq(2 * (logLik(mod_hur1) - logLik(mod_hurnull)), df = 2, lower.tail = FALSE)

'log Lik.' 8.577393e-100 (df=6)

```

19.15.2 Comparison to a Poisson Model: Vuong test

```
vuong(mod_hur1, mod_poiss1)
```

```

Vuong Non-Nested Hypothesis Test-Statistic:
(test-statistic is asymptotically distributed N(0,1) under the
null that the models are indistinguishable)
-----
          Vuong z-statistic      H_A      p-value
Raw            19.90847 model1 > model2 < 2.22e-16
AIC-corrected 19.89685 model1 > model2 < 2.22e-16
BIC-corrected 19.86438 model1 > model2 < 2.22e-16

```

The hurdle model shows a detectable improvement over the standard Poisson model according to this test.

19.15.3 Comparison to a Zero-Inflated Poisson Model: Vuong test

Is the hurdle model comparable to the zero-inflated Poisson?

```
vuong(mod_hur1, mod_zip1)
```

```

Vuong Non-Nested Hypothesis Test-Statistic:
(test-statistic is asymptotically distributed N(0,1) under the

```

```
null that the models are indistinguishable)
-----
```

	Vuong z-statistic	H_A	p-value
Raw	0.2046181	model1 > model2	0.41894
AIC-corrected	0.2046181	model1 > model2	0.41894
BIC-corrected	0.2046181	model1 > model2	0.41894

The hurdle model doesn’t show a detectable improvement over the zero-inflated Poisson model according to this test.

19.15.4 The Fitted Equation

The form of the model equation for this hurdle also requires us to take two separate models into account. First we have a logistic regression model to predict the log odds of zero `physhealth` days. That takes care of the zeros. Then, to predict the number of `physhealth` days, we use a truncated Poisson model, which is truncated to produce only estimates greater than zero.

19.15.5 Interpreting the Coefficients

We can exponentiate the logistic regression coefficients to obtain results in terms of odds ratios for that model, and that can be of some help in understanding the process behind excess zeros.

Also, exponentiating the coefficients of the count model help us describe those counts on the original scale of `physhealth`.

```
exp(coef(mod_hur1))
```

count_(Intercept)	count_bmi_c	count_smoke100	zero_(Intercept)
7.3714308	1.0183443	1.4819335	0.5048423
zero_bmi_c	zero_smoke100		
1.0281661	1.2690894		

For example,

- in the model for `physhealth` = 0, the odds of `physhealth` = 0 are 127% as high for subjects with `smoke100` = 1 as for non-smokers with the same BMI.
- in the Poisson model for `physhealth`, the `physhealth` count is estimated to increase by 1.48 for smokers as compared to non-smokers with the same BMI.

19.15.6 Testing the Predictors

We can test the model with and without `bmi_c`, for example, by fitting the model both ways, and comparing the results with either a Wald or Likelihood Ratio test, each of which is available in the `lmtest` package.

```
mod_hur1_nobmi <- pscl::hurdle(physhealth ~ smoke100,
                                 dist = "poisson",
                                 zero.dist = "binomial",
                                 data = sm_oh_A_young)

lmtest::waldtest(mod_hur1, mod_hur1_nobmi)

Wald test

Model 1: physhealth ~ bmi_c + smoke100
Model 2: physhealth ~ smoke100
  Res.Df Df  Chisq Pr(>Chisq)
1     1968
2     1970 -2 187.19 < 2.2e-16 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

lmtest::lrtest(mod_hur1, mod_hur1_nobmi)

Likelihood ratio test

Model 1: physhealth ~ bmi_c + smoke100
Model 2: physhealth ~ smoke100
  #Df  LogLik Df  Chisq Pr(>Chisq)
1    6 -5679.8
2    4 -5769.5 -2 179.35 < 2.2e-16 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

19.15.7 Store fitted values and residuals

The `broom` package does not work with the `hurdle` class of models. Again we need to build up the fitted values and residuals ourselves.

```
sm_hur1 <- sm_oh_A_young %>%
  mutate(fitted = fitted(mod_hur1, type = "response"),
        resid = resid(mod_hur1, type = "response"))

sm_hur1 %>%
  dplyr::select(physhealth, fitted, resid) %>%
  head()
```

```
# A tibble: 6 x 3
  physhealth fitted resid
  <dbl>    <dbl> <dbl>
1      0     2.21 -2.21
2      0     2.28 -2.28
3      0     3.12 -3.12
4     30     5.27 24.7
5      0     3.17 -3.17
6      0     3.71 -3.71
```

19.15.8 Modeled Number of Zero Counts

Once again, we can compare the observed number of zero `physhealth` results to the expected number of zero values from the likelihood-based models.

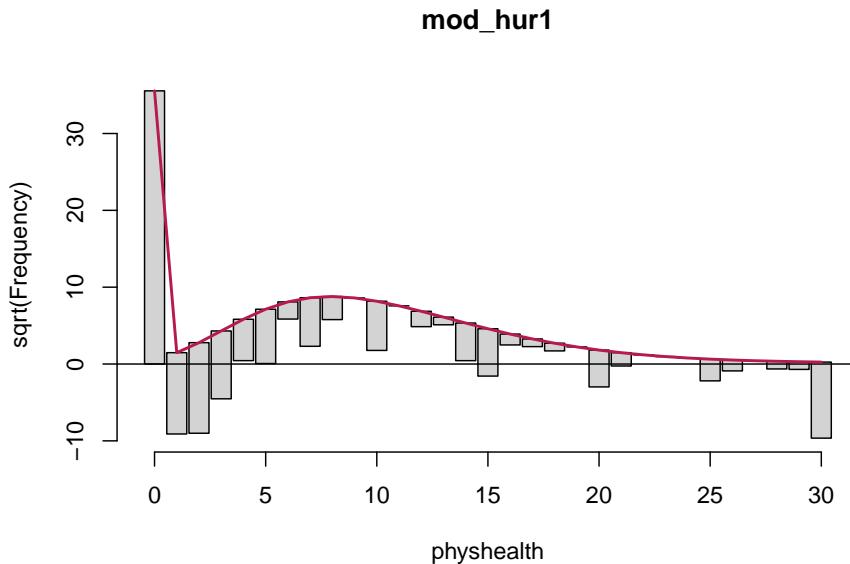
```
round(c("Obs" = sum(sm_oh_A_young$physhealth == 0),
       "Poisson" = sum(dpois(0, fitted(mod_poiss1))),
       "NB" = sum(dnbinom(0, mu = fitted(mod_nb1), size = mod_nb1$theta)),
       "ZIP" = sum(predict(mod_zip1, type = "prob")[,1]),
       "ZINB" = sum(predict(mod_zinb1, type = "prob")[,1]),
       "Hurdle" = sum(predict(mod_hur1, type = "prob")[,1])),0)
```

	Obs	Poisson	NB	ZIP	ZINB	Hurdle
1264	124	1250	1264	1264	1264	1264

The hurdle model does about as well as the negative binomial and zero-inflated models. All but the Poisson give reasonable fits in this regard.

19.15.9 Rootogram for Hurdle Model

```
countreg::rootogram(mod_hur1, max = 30)
```



The results are still not perfect, of course. In fitting the zeros exactly, we're underfitting counts of 1, 2, and 30, and overfitting many of the counts between 6 and 20. We still have a problem here with overdispersion. That's why we'll consider a hurdle model with a negative binomial regression for the counts in a moment.

19.15.10 Understanding the Modeled Counts in Detail

The expected mean count uses both parts of the hurdle model. Mathematically, we want...

$$E[y|\mathbf{x}] = \frac{1 - f_1(0|\mathbf{x})}{1 - f_2(0|\mathbf{x})} \mu_2(\mathbf{x})$$

where

- our count of `physhealth` is y
- our predictors are represented by \mathbf{x}
- and the expected count is the product of a ratio and a mean.

The ratio is the probability of a non-zero in the first process divided the probability of a non-zero in the second untruncated process. The f symbols represent distributions. Recall these are logistic and Poisson, respectively, by default but can be others. The mean is for the untruncated version of the positive-count process.

If we want to see the expected hurdle counts, we can get them using some clever applications of the `predict` function.

The first six expected mean counts ($E[y|\mathbf{x}]$ from the equation above) are:

```
head(predict(mod_hur1, type = "response"))
```

1	2	3	4	5	6
2.214179	2.281482	3.116124	5.267881	3.167890	3.712120

The ratio of non-zero probabilities, $\frac{1-f_1(0|\mathbf{x})}{1-f_2(0|\mathbf{x})}$, from the mathematical expression above can be extracted by:

```
head(predict(mod_hur1, type = "zero"))
```

1	2	3	4	5	6
0.3173312	0.3221978	0.3344433	0.4806542	0.3372292	0.3649208

The mean for the untruncated process, $\mu_2(\mathbf{x})$, can also be obtained by:

```
head(predict(mod_hur1, type = "count"))
```

1	2	3	4	5	6
6.977501	7.081000	9.317347	10.959814	9.393877	10.172401

and we can multiply these last two pieces together to verify that they match our expected hurdle counts.

```
head(predict(mod_hur1, type = "zero") * predict(mod_hur1, type = "count"), 5)
```

1	2	3	4	5
2.214179	2.281482	3.116124	5.267881	3.167890

19.15.11 Specify the R^2 and log (likelihood) values

We can calculate a proxy for R^2 as the squared correlation of the fitted values and the observed values.

```
# The correlation of observed and fitted values
(hur1_r <- with(sm_hur1, cor(physhealth, fitted)))
```

```
[1] 0.1873104
```

```
# R-square
```

```
hur1_r^2
```

```
[1] 0.03508517
```

```
logLik(mod_hur1)
```

```
'log Lik.' -5679.792 (df=6)
```

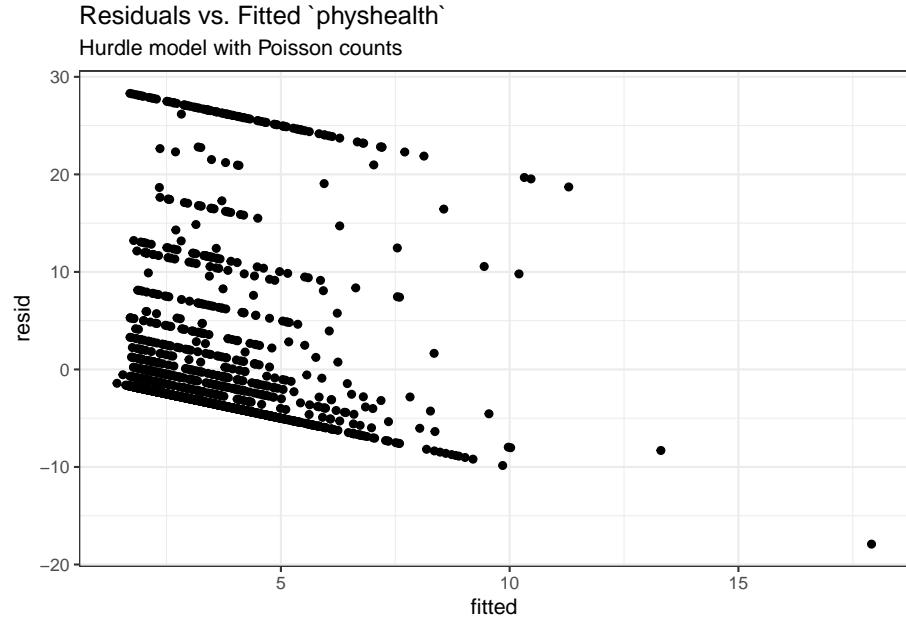
Here, we have

Model	Scale	R^2	log(likelihood)
Hurdle Model (Poisson)	Complex: log(physhealth)	.035	-5679.83

19.15.12 Check model assumptions

Here is a plot of residuals vs. fitted values on the original physhealth scale.

```
ggplot(sm_hur1, aes(x = fitted, y = resid)) +
  geom_point() +
  labs(title = "Residuals vs. Fitted `physhealth`",
       subtitle = "Hurdle model with Poisson counts")
```



19.15.13 Predictions for Harry and Sally

The predictions from this zero-inflated negative binomial regression model are obtained as follows...

```
predict(mod_hur1, newdata = hs_data, type = "response")
```

1	2
6.003689	2.057127

As we've seen in the past, when we use `response` as the type, the predictions fall on the original `physhealth` scale. The prediction for Harry is 6.0 days, and for Sally is 2.1 days.

19.16 A “hurdle” model (with negative binomial for overdispersion)

Let's account for overdispersion better with a negative binomial model for the counts in our hurdle model. We specify that the positive-count process be fit with this NB model using `dist = negbin`.

```
mod_hur_nb1 <- pscl::hurdle(physhealth ~ bmi_c + smoke100,
                             dist = "negbin", zero.dist = "binomial",
                             data = sm_oh_A_young)

summary(mod_hur_nb1)

Call:
pscl::hurdle(formula = physhealth ~ bmi_c + smoke100, data = sm_oh_A_young,
  dist = "negbin", zero.dist = "binomial")

Pearson residuals:
    Min      1Q   Median      3Q      Max
-0.5749 -0.4178 -0.3948 -0.1165  6.4023

Count model coefficients (truncated negbin with log link):
            Estimate Std. Error z value Pr(>|z|)
(Intercept) 1.54170   0.10101 15.263 < 2e-16 ***
bmi_c       0.02434   0.00677  3.595 0.000324 ***
smoke100    0.51792   0.11101  4.666 3.08e-06 ***
Log(theta)  -0.88245   0.14653 -6.023 1.72e-09 ***
Zero hurdle model coefficients (binomial with logit link):
            Estimate Std. Error z value Pr(>|z|)
(Intercept) -0.683509   0.062827 -10.879 < 2e-16 ***
bmi_c        0.027777   0.006508   4.268 1.97e-05 ***
smoke100     0.238300   0.095301   2.500  0.0124 *
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Theta: count = 0.4138
Number of iterations in BFGS optimization: 17
Log-likelihood: -3469 on 7 Df
confint(mod_hur_nb1)
```

	2.5 %	97.5 %
count_(Intercept)	1.34373051	1.73966938
count_bmi_c	0.01107002	0.03760871
count_smoke100	0.30035006	0.73549531
zero_(Intercept)	-0.80664805	-0.56037031
zero_bmi_c	0.01502124	0.04053230
zero_smoke100	0.05151314	0.42508607

19.16.1 Comparison to a null model

To show that this model fits better than the null model (the model with intercept only), we can compare them directly with a chi-squared test. Since we have two predictors in the full model, the degrees of freedom for this test is 2.

```
mod_hur_nb_null <- pscl::hurdle(physhealth ~ 1, dist = "negbin",
                                 zero.dist = "binomial",
                                 data = sm_oh_A_young)

summary(mod_hur_nb_null)

Call:
pscl::hurdle(formula = physhealth ~ 1, data = sm_oh_A_young, dist = "negbin",
              zero.dist = "binomial")

Pearson residuals:
    Min      1Q  Median      3Q     Max 
-0.4048 -0.4048 -0.4048 -0.1622  3.2340 

Count model coefficients (truncated negbin with log link):
            Estimate Std. Error z value Pr(>|z|)    
(Intercept) 1.77653   0.09641 18.427 < 2e-16 ***
Log(theta) -1.04455   0.16123 -6.479 9.25e-11 ***

Zero hurdle model coefficients (binomial with logit link):
            Estimate Std. Error z value Pr(>|z|)    
(Intercept) -0.5768    0.0469  -12.3  <2e-16 ***

---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Theta: count = 0.3518
Number of iterations in BFGS optimization: 12
Log-likelihood: -3498 on 3 Df
pchisq(2 * (logLik(mod_hur_nb1) - logLik(mod_hur_nb_null)), df = 2, lower.tail = FALSE)

'log Lik.' 2.17419e-13 (df=7)
```

19.16.2 Comparison to a Negative Binomial Model: Vuong test

```
vuong(mod_hur_nb1, mod_nb1)
```

Vuong Non-Nested Hypothesis Test-Statistic:
(test-statistic is asymptotically distributed $N(0,1)$ under the
null that the models are indistinguishable)

	Vuong z-statistic	H_A	p-value
Raw	3.0953118	model1 > model2	0.00098303
AIC-corrected	2.4837306	model1 > model2	0.00650071
BIC-corrected	0.7750288	model1 > model2	0.21916133

The hurdle model is a significant improvement over the standard negative binomial model according to the raw and AIC-corrected versions of this test, but not the BIC-corrected version.

19.16.3 Comparison to a Zero-Inflated NB Model: Vuong test

Is the hurdle model comparable to the zero-inflated Poisson?

```
vuong(mod_hur_nb1, mod_zinb1)
```

Vuong Non-Nested Hypothesis Test-Statistic:
(test-statistic is asymptotically distributed $N(0,1)$ under the
null that the models are indistinguishable)

	Vuong z-statistic	H_A	p-value
Raw	0.91178	model1 > model2	0.18094
AIC-corrected	0.91178	model1 > model2	0.18094
BIC-corrected	0.91178	model1 > model2	0.18094

The hurdle model is just barely a significant improvement over the zero-inflated Negative Binomial model.

19.16.4 Comparing the Hurdle Models with AIC and BIC

```
AIC(mod_hur1); BIC(mod_hur1)
```

```
[1] 11371.58
```

```
[1] 11405.11
```

AIC(mod_hur_nb1); BIC(mod_hur_nb1)

[1] 6952.186

[1] 6991.301

The negative binomial approach certainly looks better than the Poisson here.

19.16.5 The Fitted Equation

The form of the model equation for this hurdle also requires us to take two separate models into account. First we have a logistic regression model to predict the log odds of zero `physhealth` days. That takes care of the zeros. Then, to predict the number of `physhealth` days, we use a truncated negative binomial model, which is truncated to produce only estimates greater than zero, with θ estimated as `exp(-1.123)` or 0.325.

19.16.6 Interpreting the Coefficients

We can exponentiate the logistic regression coefficients to obtain results in terms of odds ratios for that model, and that can be of some help in understanding the process behind excess zeros.

```
exp(coef(mod_hur_nb1))
```

```

count_(Intercept)      count_bmi_c    count_smoke100 zero_(Intercept)
        4.6725266          1.0246380         1.6785372       0.5048423
zero_bmi_c            zero_smoke100
        1.0281661          1.2690894

```

For example,

- in the model for `physhealth` = 0, the odds of `physhealth` = 0 are 127% as high for subjects with `smoke100` = 1 as for non-smokers with the same BMI.

19.16.7 Testing the Predictors

We can test the model with and without `bmi_c`, for example, by fitting the model both ways, and comparing the results with either a Wald or Likelihood Ratio test, each of which is available in the `lmtest` package.

```

lmtest::waldtest(mod_hur_nb1, mod_hurnb1_nobmi)

Wald test

Model 1: physhealth ~ bmi_c + smoke100
Model 2: physhealth ~ smoke100
  Res.Df Df  Chisq Pr(>Chisq)
1     1967
2     1969 -2 31.141  1.729e-07 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

lmtest::lrtest(mod_hur_nb1, mod_hurnb1_nobmi)

Likelihood ratio test

Model 1: physhealth ~ bmi_c + smoke100
Model 2: physhealth ~ smoke100
  #Df  LogLik Df  Chisq Pr(>Chisq)
1    7 -3469.1
2    5 -3485.0 -2 31.812  1.236e-07 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

```

19.16.8 Store fitted values and residuals

Again we need to build up the fitted values and residuals, without `broom` to help.

```

sm_hur_nb1 <- sm_oh_A_young %>%
  mutate(fitted = fitted(mod_hur_nb1, type = "response"),
        resid = resid(mod_hur_nb1, type = "response"))

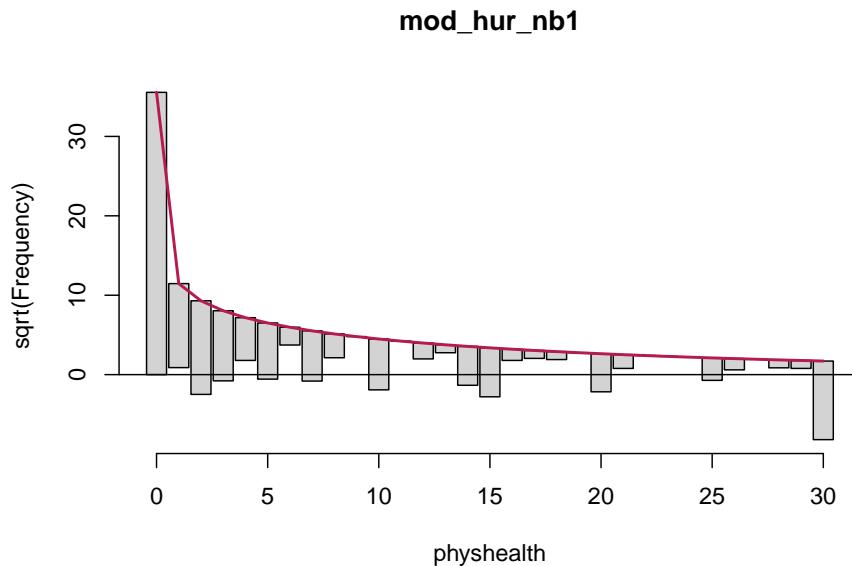
sm_hur_nb1 %>%
  dplyr::select(physhealth, fitted, resid) %>%
  head()

# A tibble: 6 x 3
  physhealth fitted resid
  <dbl>    <dbl> <dbl>
1 0        2.16 -2.16
2 0        2.23 -2.23
3 0        3.09 -3.09
4 30      5.37 24.6 
5 0        3.15 -3.15
6 0        3.72 -3.72

```

19.16.9 Rootogram for NB Hurdle Model

```
countreg::rootogram(mod_hur_nb1, max = 30)
```



This improves the situation, but we're still underfitting the 30s.

19.16.10 Specify the R^2 and log (likelihood) values

We can calculate a proxy for R^2 as the squared correlation of the fitted values and the observed values.

```
# The correlation of observed and fitted values
(hurnb1_r <- with(sm_hur_nb1, cor(physhealth, fitted)))
```

```
[1] 0.1856393
```

```
# R-square
hurnb1_r^2
```

```
[1] 0.03446194
```

```
logLik(mod_hur_nb1)
```

```
'log Lik.' -3469.093 (df=7)
```

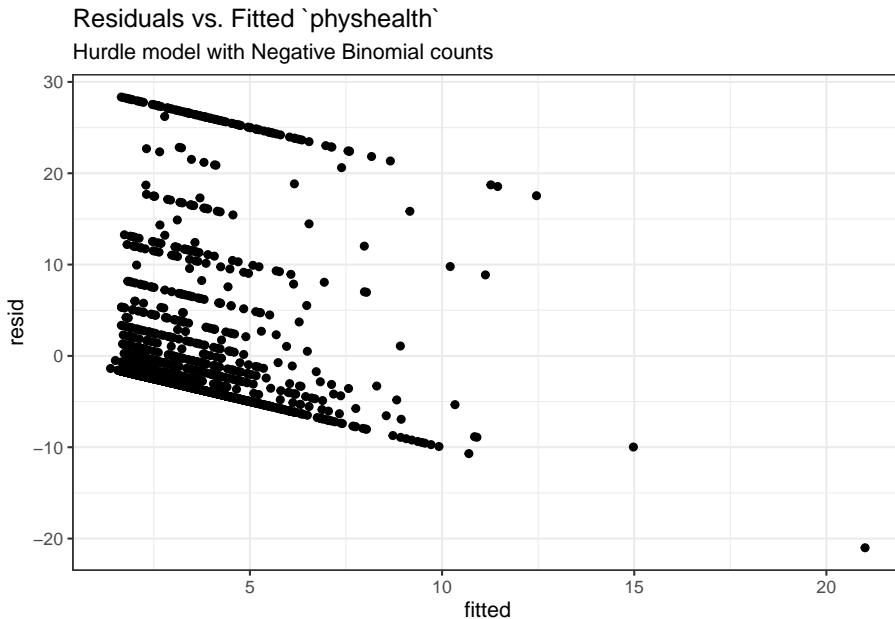
Here, we have

Model	Scale	R^2	log(likelihood)
Hurdle Model (Neg. Bin.)	Complex: log(physhealth)	.035	-3469.07

19.16.11 Check model assumptions

Here is a plot of residuals vs. fitted values on the original physhealth scale.

```
ggplot(sm_hur_nb1, aes(x = fitted, y = resid)) +
  geom_point() +
  labs(title = "Residuals vs. Fitted `physhealth`",
       subtitle = "Hurdle model with Negative Binomial counts")
```



19.16.12 Predictions for Harry and Sally

The predictions from this zero-inflated negative binomial regression model are obtained as follows...

```
predict(mod_hur_nb1, newdata = hs_data, type = "response")
```

```
1      2
6.222041 2.007095
```

The prediction for Harry is 6.22 days, and for Sally is 2.01 days.

19.16.13 Note: Fitting a Different Hurdle Model for Counts and Pr(zero)

Suppose we wanted to use only `bmi_c` to predict the probability of a zero count, but use both predictors in the model for the positive counts. We use the `|` command.

```
mod_hur_new1 <-
  pscl::hurdle(physhealth ~ bmi_c + smoke100 | bmi_c,
    dist = "negbin", zero.dist = "binomial",
    data = sm_oh_A_young)

summary(mod_hur_new1)

Call:
pscl::hurdle(formula = physhealth ~ bmi_c + smoke100 | bmi_c, data = sm_oh_A_young,
  dist = "negbin", zero.dist = "binomial")

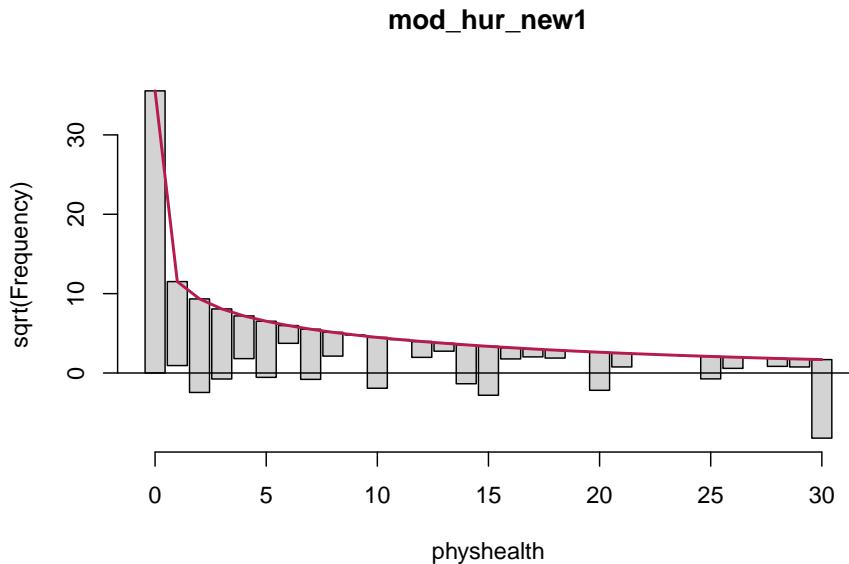
Pearson residuals:
    Min      1Q   Median      3Q      Max
-0.5630 -0.4179 -0.3971 -0.1150  6.1958

Count model coefficients (truncated negbin with log link):
            Estimate Std. Error z value Pr(>|z|)
(Intercept) 1.54170   0.10101 15.263 < 2e-16 ***
bmi_c       0.02434   0.00677  3.595 0.000324 ***
smoke100   0.51792   0.11101  4.666 3.08e-06 ***
Log(theta) -0.88245   0.14653 -6.023 1.72e-09 ***
Zero hurdle model coefficients (binomial with logit link):
            Estimate Std. Error z value Pr(>|z|)
(Intercept) -0.58186   0.04718 -12.332 < 2e-16 ***
bmi_c        0.02853   0.00649  4.396 1.1e-05 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Theta: count = 0.4138
Number of iterations in BFGS optimization: 17
Log-likelihood: -3472 on 6 Df
```

19.16.14 Hanging Rootogram for this new Hurdle Model

```
countreg::rootogram(mod_hur_new1, max = 30)
```



Not a meaningful improvement, certainly.

19.17 A Tobit (Censored) Regression Model

The idea of the **tobit** model (sometimes called a **censored regression** model) is to estimate associations for outcomes where we can see either left-censoring (censoring from below) or right-censoring (censoring from above.)

Consider the variable **physhealth**, which is restricted to fall between 0 and 30 by the way the measure was constructed. But suppose we think about a broader and latent (unobserved) variable describing physical health. Among the people with **physhealth** = 0, some would be incredible athletes and others would be in much poorer physical health but still healthy enough to truthfully answer 0. On the other end, some of the people responding 30 are in substantially worse physical health than others with that same response.

- Censoring from below takes place when values at or below a threshold (in this case 0) take that value.
- Censoring from above takes place when values at or above a threshold (here, 30) take that value.

Several examples of tobit analysis are available at <https://stats.idre.ucla.edu/r/dae/tobit-models/>, which is my primary source for the material here on those models.

The tobit model postulates that the value 0 in our model is just the lower limit of the underlying measure of poor physical health that we would actually observe in the population if we had a stronger measure. Similarly, we'll postulate that 30 is just the upper limit of "poor health" that we can see. The approach I'll take to run the tobit model comes from the `vglm` function in the `VGAM` package.

Here's the model, and its summary. Note that the default Lower value for a tobit model is 0, so we didn't technically have to list that here.

```
mod_tob1 <- vglm(physhealth ~ bmi_c + smoke100,
                  tobit(Lower = 0, Upper = 30),
                  type.fitted = "censored",
                  data = sm_oh_A_young)

summary(mod_tob1)
```

```
Call:
vglm(formula = physhealth ~ bmi_c + smoke100, family = tobit(Lower = 0,
Upper = 30), data = sm_oh_A_young, type.fitted = "censored")

Coefficients:
            Estimate Std. Error z value Pr(>|z|)
(Intercept):1 -9.29104   0.75509 -12.305 < 2e-16 ***
(Intercept):2  2.87081   0.03028  94.818 < 2e-16 ***
bmi_c         0.37434   0.06537   5.727 1.02e-08 ***
smoke100      4.35390   0.97223   4.478 7.52e-06 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Names of linear predictors: mu, loglink(sd)

Log-likelihood: -3420.605 on 3944 degrees of freedom

Number of Fisher scoring iterations: 6

No Hauck-Donner effect found in any of the estimates
confint(mod_tob1)
```

	2.5 %	97.5 %
(Intercept):1	-10.7709932	-7.8110877
(Intercept):2	2.8114721	2.9301560
bmi_c	0.2462244	0.5024551
smoke100	2.4483714	6.2594277

19.17.1 The Fitted Equation

Because we've used the censoring approach, our model will limit its predictions to the range of [0, 30], where any predictions outside that range are censored. Values below 0 are fitted as 0, and values above 30 are fitted as 30.

The model equation is

```
physhealth = -9.29 + 0.37 bmi_c + 4.35 smoke100
```

19.17.2 Interpreting the Coefficients

Tobit model regression coefficients are interpreted as we would a set of OLS coefficients, *except* that the linear effect is on the uncensored *latent variable*, rather than on the observed outcome.

In our case,

- a one-unit increase in `bmi_c` is associated with a 0.37 day increase in the predicted value of `physhealth`, holding `smoke100` constant
- a move from `smoke100 = 0` to 1 is associated with a 4.35 day increase in the predicted value of `physhealth`, holding `bmi_c` constant
- the coefficient labeled `(Intercept):1` is the intercept for the model and is the predicted value of `physhealth` when `smoke100 = 0` and `bmi_c = 0`. Note that this value is -9.29, which is outside the range of `physhealth` values we observed.
- the coefficient labeled `(Intercept):2` is a statistic we can use after we exponentiate it, as follows:
 - here `(Intercept):2 = 2.87`, and $\exp(2.87) = 17.6370182$, which is analogous to the square root of the residual variance in OLS regression, which is summarized for our OLS model as `Residual standard error: 17.64`.

19.17.3 Testing the Predictors

We can test the model with and without `bmi_c`, for example, by fitting the model both ways, and comparing the results with either a Wald or Likelihood Ratio test, each of which is available in the `lmtest` package.

```
mod_tob_nobmi <- vglm(physhealth ~ smoke100,
                        tobit(Lower = 0, Upper = 30),
                        type.fitted = "censored",
                        data = sm_oh_A_young)

lmtest::waldtest(mod_tob1, mod_tob_nobmi)
```

```
Wald test
```

```
Model 1: physhealth ~ bmi_c + smoke100
Model 2: physhealth ~ smoke100
  Res.Df Df  Chisq Pr(>Chisq)
1    3944
2    3945 -1 32.796 1.023e-08 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

The likelihood ratio test we have used in some other settings isn't available here.

19.17.4 Store fitted values and residuals

The residuals and fitted values from the tobit model can be stored and then summarized in several ways:

```
sm_tob1 <- sm_oh_A_young %>%
  mutate(fitted = fitted(mod_tob1,
                        type.fitted = "censored"),
         resid = physhealth - fitted)

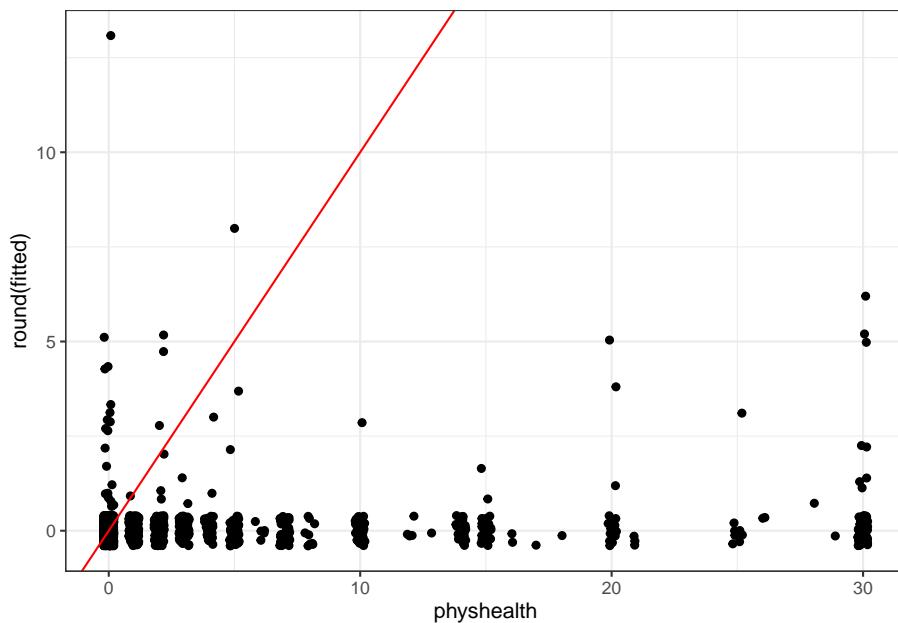
sm_tob1 %>%
  dplyr::select(physhealth, fitted, resid) %>%
  head()

# A tibble: 6 x 3
  physhealth fitted[,1] resid[,1]
  <dbl>      <dbl>     <dbl>
1 0          0          0
2 0          0          0
3 0          0          0
4 30         0          30
5 0          0          0
6 0          0          0
```

19.17.5 Building Something Like a Rootogram

Building a rootogram is tricky for a tobit model, to say the least, but we can approximate a piece of the issue by plotting the rounded fitted values against the observed `physhealth` data.

```
ggplot(sm_tob1, aes(x = physhealth, y = round(fitted))) +
  geom_jitter(width = 0.2) +
  geom_abline(intercept = 0, slope = 1, col = "red")
```



Note that the model never predicts a subject to have an underlying `physhealth` worse than about 13 (remember that larger numbers indicate worse health here.)

19.17.6 Tables of the Observed and Fitted `physhealth` from Tobit

```
addmargins(table(round(sm_tob1$physhealth)))
```

	0	1	2	3	4	5	6	7	8	10	12	13	14	15	16	17
1264	112	139	78	29	50	5	40	9	41	4	1	24	38	2	1	
18	20	21	25	26	28	29	30	Sum								
1	23	3	8	2	1	1	98	1974								

```
addmargins(table(round(sm_tob1$fitted)))
```

	0	1	2	3	4	5	6	8	13	Sum
1924	20	7	10	4	6	1	1	1	1	1974

19.17.7 Specify the R^2 and log (likelihood) values

We can calculate a proxy for R^2 as the squared correlation of the fitted values and the observed values.

```
# The correlation of observed and fitted values
(tob1_r <- with(sm_tob1, cor(physhealth, fitted)))

[,1]
[1,] 0.09329619

# R-square
tob1_r^2

[,1]
[1,] 0.008704178

logLik(mod_tob1)

[1] -3420.605
```

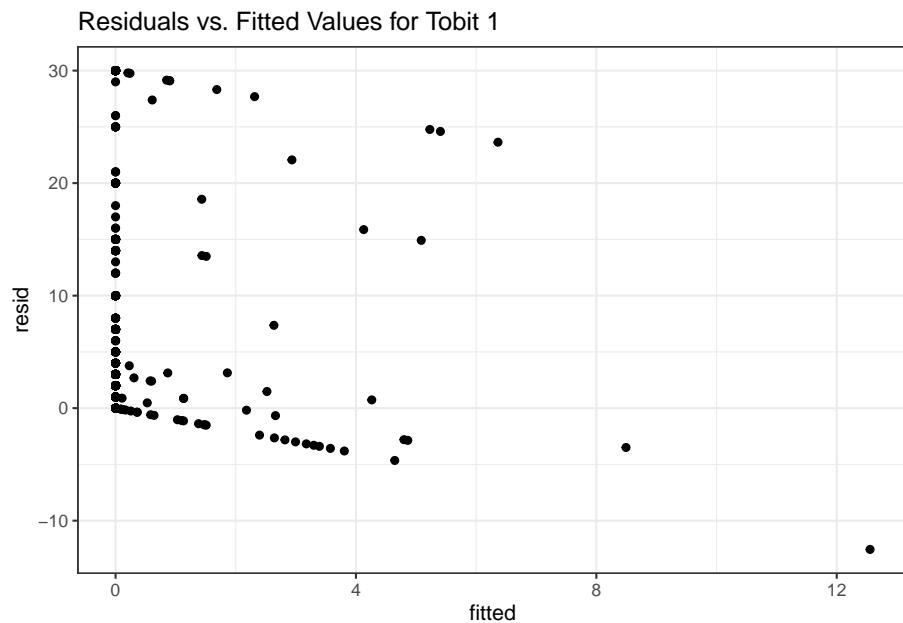
Here, we have

Model	Scale	R ²	log(likelihood)
Tobit	physhealth	.008	-3420.58

19.17.8 Check model assumptions

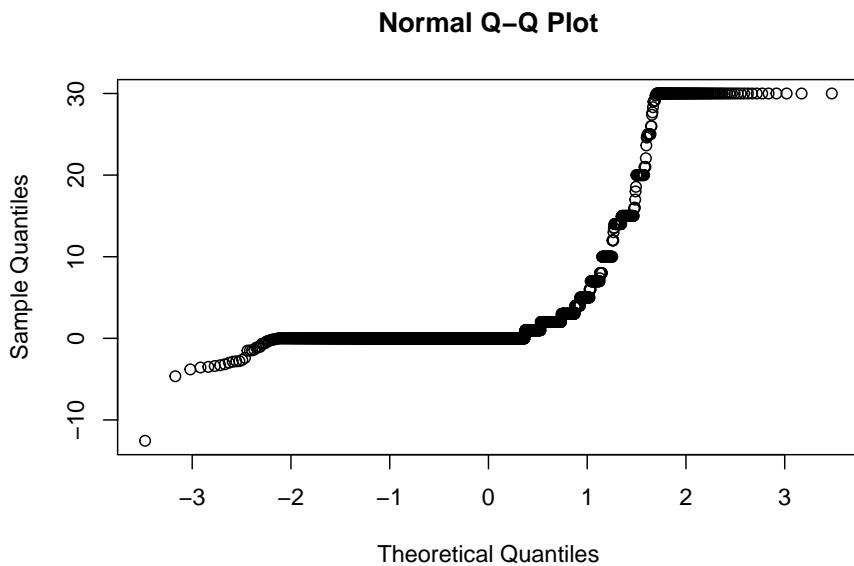
Here is a plot of residuals vs. fitted values.

```
ggplot(sm_tob1, aes(x = fitted, y = resid)) +
  geom_point() +
  labs(title = "Residuals vs. Fitted Values for Tobit 1")
```



Here is a normal Q-Q plot of the Tobit Model 1 residuals.

```
qqnorm(sm_tob1$resid)
```



19.17.9 Predictions for Harry and Sally

The predictions from this tobit model are obtained as follows...

```
predict(mod_tob1, newdata = hs_data, type = "response")
```

```
[,1]
1 -1.193743
2 -11.162739
```

The prediction for both Harry and Sally under the tobit model would be truncated to 0 days.

Chapter 20

Modeling an Ordinal Categorical Outcome in Ohio SMART

20.1 Preliminaries

```
library(gmodels)
library(nnet)

smart_oh <- readRDS(here("data", "smart_ohio.Rds"))
```

20.2 A subset of the Ohio SMART data

Let's consider the following data. The outcome we'll study now is `genhealth`, which has five ordered categories. I'll include the subset of all observations in `smart_oh` with complete data on these 7 variables.

Variable	Description
<code>SEQNO</code>	Subject identification code
<code>genhealth</code>	Five categories (1 = Excellent, 2 = Very Good, 3 = Good, 4 = Fair, 5 = Poor) on general health
<code>physhealth</code>	Now thinking about your physical health, which includes physical illness and injury, for how many days during the past 30 days was your physical health not good?
<code>veg_day</code>	mean number of vegetable servings consumed per day

Variable	Description
<code>costprob</code>	1 indicates Yes to “Was there a time in the past 12 months when you needed to see a doctor but could not because of cost?” and 0 otherwise.
<code>incomegroup</code>	8 income groups from < 10,000 to 75,000 or more
<code>bmi</code>	body-mass index

To make my life easier later, I’m going to drop any subjects with missing data on these variables. I’m also going to drop the subjects who have no missing data, but have a listed `bmi` above 60.

```
sm1 <- smart_oh %>%
  select(SEQNO, genhealth, physhealth, costprob, veg_day,
         incomegroup, bmi) %>%
  filter(bmi <= 60) %>%
  drop_na
```

In total, we have 5394 subjects in the `sm1` sample.

20.2.1 Several Ways of Storing Multi-Categorical data

We will store the information in our outcome, `genhealth` in both a numeric form (`gen_n`) and an ordered factor (`gen_h`) with some abbreviated labels) because we’ll have some use for each approach in this material.

```
sm1 <- sm1 %>%
  mutate(genh = fct_recode(genhealth,
                           "1-E" = "1_Excellent",
                           "2_VG" = "2_VeryGood",
                           "3_G" = "3_Good",
                           "4_F" = "4_Fair",
                           "5_P" = "5_Poor"),
        genh = factor(genh, ordered = TRUE),
        gen_n = as.numeric(genhealth))

sm1 %>% count(genh, gen_n, genhealth)

# A tibble: 5 x 4
  genh   gen_n genhealth     n
  <ord> <dbl> <fct>     <int>
1 1-E      1 1_Excellent  822
2 2_VG    2 2_VeryGood 1805
3 3_G      3 3_Good    1667
4 4_F      4 4_Fair     801
5 5_P      5 5_Poor     299
```

20.3 Building Cross-Tabulations

Is income group associated with general health?

20.3.1 Using base `table` functions

```
addmargins(table(sm1$incomegroup, sm1$genh))
```

	1-E	2_VG	3_G	4_F	5_P	Sum
0-9K	14	45	60	74	40	233
10-14K	12	41	66	93	46	258
15-19K	40	76	119	96	61	392
20-24K	51	129	175	100	50	505
25-34K	51	172	215	123	36	597
35-49K	97	270	303	118	24	812
50-74K	128	337	265	94	16	840
75K+	429	735	464	103	26	1757
Sum	822	1805	1667	801	299	5394

More people answer Very Good and Good than choose the other categories. It might be easier to look at percentages here.

20.3.1.1 Adding percentages within each row

Here are the percentages giving each `genhealth` response within each income group.

```
addmargins(
  round(100*prop.table(
    table(sm1$incomegroup, sm1$genh)
    ,1)
    ,1)
)
```

	1-E	2_VG	3_G	4_F	5_P	Sum
0-9K	6.0	19.3	25.8	31.8	17.2	100.1
10-14K	4.7	15.9	25.6	36.0	17.8	100.0
15-19K	10.2	19.4	30.4	24.5	15.6	100.1
20-24K	10.1	25.5	34.7	19.8	9.9	100.0
25-34K	8.5	28.8	36.0	20.6	6.0	99.9
35-49K	11.9	33.3	37.3	14.5	3.0	100.0
50-74K	15.2	40.1	31.5	11.2	1.9	99.9
75K+	24.4	41.8	26.4	5.9	1.5	100.0

```
Sum      91.0 224.1 247.7 164.3  72.9 800.0
```

So, for example, 11.3% of the `genhealth` responses in subjects with incomes between 25 and 34 thousand dollars were Excellent.

20.3.1.2 Adding percentages within each column

Here are the percentages in each `incomegroup` within each `genhealth` response.

```
addmargins(
  round(100*prop.table(
    table(sm1$incomegroup, sm1$genh)
    ,2)
    ,1)
)
```

	1-E	2_VG	3_G	4_F	5_P	Sum
0-9K	1.7	2.5	3.6	9.2	13.4	30.4
10-14K	1.5	2.3	4.0	11.6	15.4	34.8
15-19K	4.9	4.2	7.1	12.0	20.4	48.6
20-24K	6.2	7.1	10.5	12.5	16.7	53.0
25-34K	6.2	9.5	12.9	15.4	12.0	56.0
35-49K	11.8	15.0	18.2	14.7	8.0	67.7
50-74K	15.6	18.7	15.9	11.7	5.4	67.3
75K+	52.2	40.7	27.8	12.9	8.7	142.3
Sum	100.1	100.0	100.0	100.0	100.0	500.1

From this table, we see that 7.4% of the Excellent `genhealth` responses were given by people with incomes between 25 and 34 thousand dollars.

20.3.2 Using `xtabs`

The `xtabs` function provides a formula method for obtaining cross-tabulations.

```
xtabs(~ incomegroup + genh, data = sm1)
```

		genh					
		incomegroup	1-E	2_VG	3_G	4_F	5_P
		0-9K	14	45	60	74	40
		10-14K	12	41	66	93	46
		15-19K	40	76	119	96	61
		20-24K	51	129	175	100	50
		25-34K	51	172	215	123	36
		35-49K	97	270	303	118	24
		50-74K	128	337	265	94	16
		75K+	429	735	464	103	26

20.3.3 Storing a table in a tibble

We can store the elements of a cross-tabulation in a tibble, like this:

```
(sm1.tableA %>% count(incomegroup, genh))
```

```
# A tibble: 40 x 3
  incomegroup genh     n
  <fct>      <ord> <int>
  1 0-9K       1-E     14
  2 0-9K       2_VG    45
  3 0-9K       3_G     60
  4 0-9K       4_F     74
  5 0-9K       5_P     40
  6 10-14K     1-E    12
  7 10-14K     2_VG   41
  8 10-14K     3_G    66
  9 10-14K     4_F    93
  10 10-14K    5_P    46
# ... with 30 more rows
```

From such a tibble, we can visualize the data in many ways, but we can also return to `xtabs` and include the frequencies (`n`) in that setup.

```
xtabs(n ~ incomegroup + genh, data = sm1.tableA)
```

		genh					
		1-E	2_VG	3_G	4_F	5_P	
incomegroup		0-9K	14	45	60	74	40
		10-14K	12	41	66	93	46
		15-19K	40	76	119	96	61
		20-24K	51	129	175	100	50
		25-34K	51	172	215	123	36
		35-49K	97	270	303	118	24
		50-74K	128	337	265	94	16
		75K+	429	735	464	103	26

And, we can get the χ^2 test of independence, with:

```
summary(xtabs(n ~ incomegroup + genh, data = sm1.tableA))
```

```
Call: xtabs(formula = n ~ incomegroup + genh, data = sm1.tableA)
Number of cases in table: 5394
Number of factors: 2
Test for independence of all factors:
  Chisq = 894.2, df = 28, p-value = 3.216e-170
```

20.3.4 Using `CrossTable` from the `gmodels` package

The `CrossTable` function from the `gmodels` package produces a cross-tabulation with various counts and proportions like people often generate with SPSS and SAS.

```
CrossTable(sm1$incomegroup, sm1$genh, chisq = T)
```

Cell Contents	
	N
Chi-square contribution	
N / Row Total	
N / Col Total	
N / Table Total	

Total Observations in Table: 5394

		sm1\$genh					Row Total
		1-E	2_VG	3_G	4_F	5_P	
sm1\$incomegroup	0-9K	14	45	60	74	40	233
		13.027	13.941	2.002	44.865	56.796	
		0.060	0.193	0.258	0.318	0.172	0.043
		0.017	0.025	0.036	0.092	0.134	
		0.003	0.008	0.011	0.014	0.007	
	10-14K	12	41	66	93	46	256
		18.980	23.806	2.366	78.061	70.259	
		0.047	0.159	0.256	0.360	0.178	0.044
		0.015	0.023	0.040	0.116	0.154	
		0.002	0.008	0.012	0.017	0.009	
		40	76	119	96	61	392
		6.521	23.208	0.038	24.531	70.973	
		0.102	0.194	0.304	0.245	0.156	0.073
		0.049	0.042	0.071	0.120	0.204	
		0.007	0.014	0.022	0.018	0.011	
		51	129	175	100	50	508
		8.756	9.463	2.296	8.340	17.301	
		0.101	0.255	0.347	0.198	0.099	0.093

	0.062	0.071	0.105	0.125	0.167		
	0.009	0.024	0.032	0.019	0.009		
25-34K	51	172	215	123	36		597
	17.567	3.862	5.042	13.307	0.255		
	0.085	0.288	0.360	0.206	0.060	0.111	
	0.062	0.095	0.129	0.154	0.120		
	0.009	0.032	0.040	0.023	0.007		
35-49K	97	270	303	118	24		812
	5.779	0.011	10.798	0.055	9.808		
	0.119	0.333	0.373	0.145	0.030	0.151	
	0.118	0.150	0.182	0.147	0.080		
	0.018	0.050	0.056	0.022	0.004		
50-74K	128	337	265	94	16		840
	0.000	11.121	0.112	7.575	20.061		
	0.152	0.401	0.315	0.112	0.019	0.156	
	0.156	0.187	0.159	0.117	0.054		
	0.024	0.062	0.049	0.017	0.003		
75K+	429	735	464	103	26		1757
	97.108	36.780	11.492	95.573	52.335		
	0.244	0.418	0.264	0.059	0.015	0.326	
	0.522	0.407	0.278	0.129	0.087		
	0.080	0.136	0.086	0.019	0.005		
Column Total	822	1805	1667	801	299		5394
	0.152	0.335	0.309	0.148	0.055		

Statistics for All Table Factors

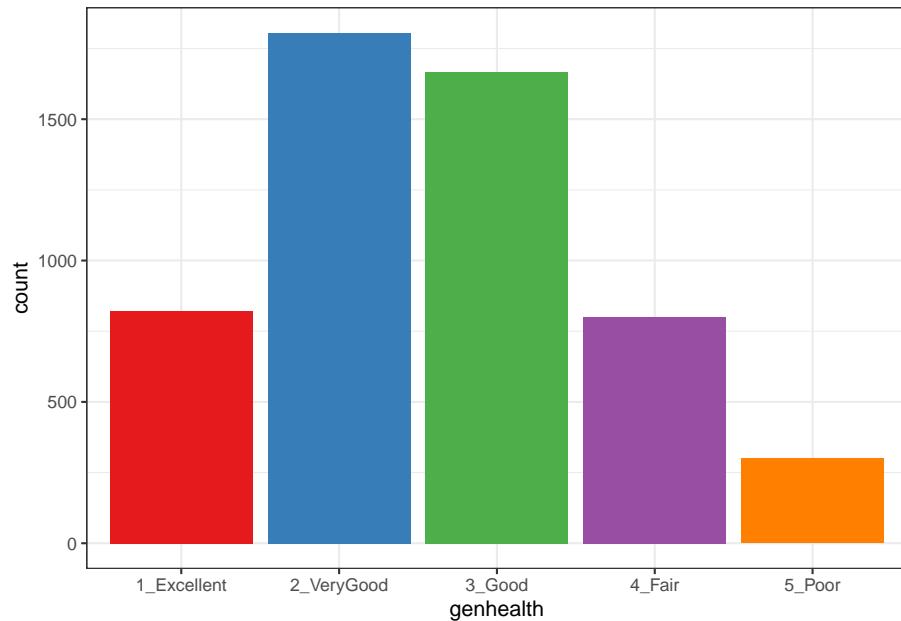
Pearson's Chi-squared test

Chi^2 = 894.1685 d.f. = 28 p = 3.216132e-170

20.4 Graphing Categorical Data

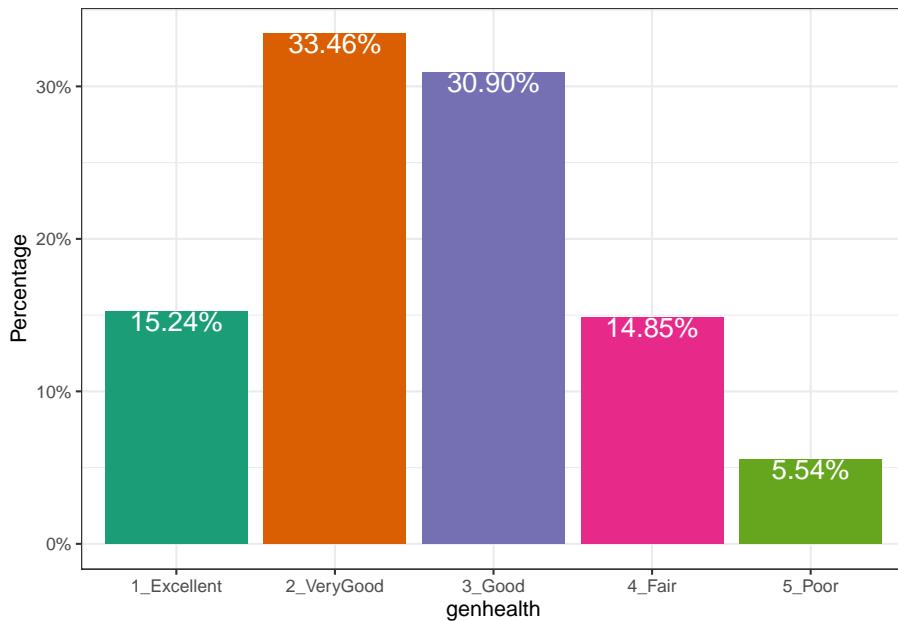
20.4.1 A Bar Chart for a Single Variable

```
ggplot(sml, aes(x = genhealth, fill = genhealth)) +
  geom_bar() +
  scale_fill_brewer(palette = "Set1") +
  guides(fill = "none")
```



or, you might prefer to plot percentages, perhaps like this:

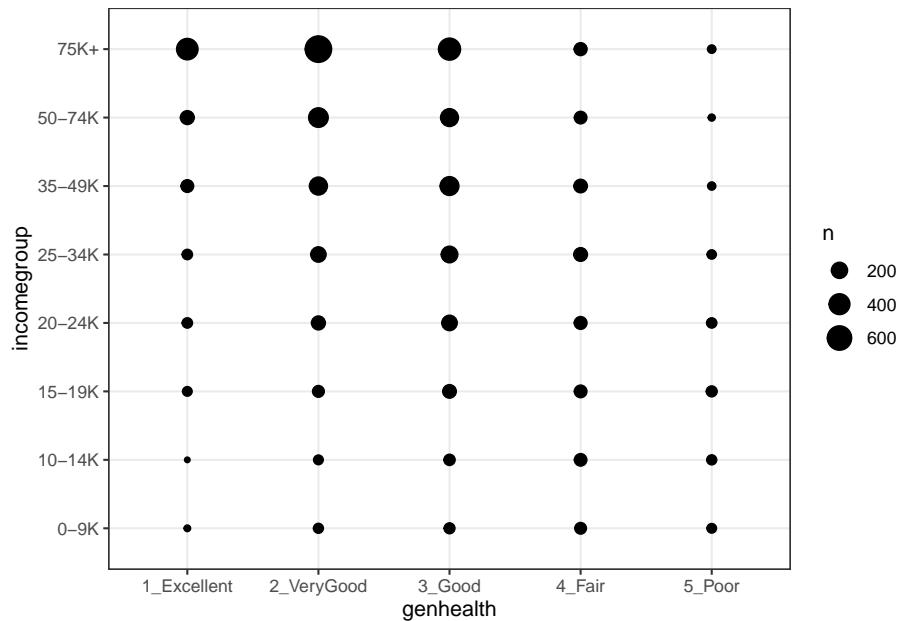
```
ggplot(sml, aes(x = genhealth, fill = genhealth)) +
  geom_bar(aes(y =(..count..)/sum(..count..))) +
  geom_text(aes(y =(..count..)/sum(..count..),
                label = scales::percent(..count..) /
                  sum(..count..))),
  stat = "count", vjust = 1,
  color = "white", size = 5) +
  scale_y_continuous(labels = scales::percent) +
  scale_fill_brewer(palette = "Dark2") +
  guides(fill = "none") +
  labs(y = "Percentage")
```



Use bar charts, rather than pie charts.

20.4.2 A Counts Chart for a 2-Way Cross-Tabulation

```
ggplot(sm1, aes(x = genhealth, y = incomegroup)) +  
  geom_count()
```



20.5 Building a Model for genh using veg_day

To begin, we'll predict each subject's genh response using just one predictor, `veg_day`.

20.5.1 A little EDA

Let's start with a quick table of summary statistics.

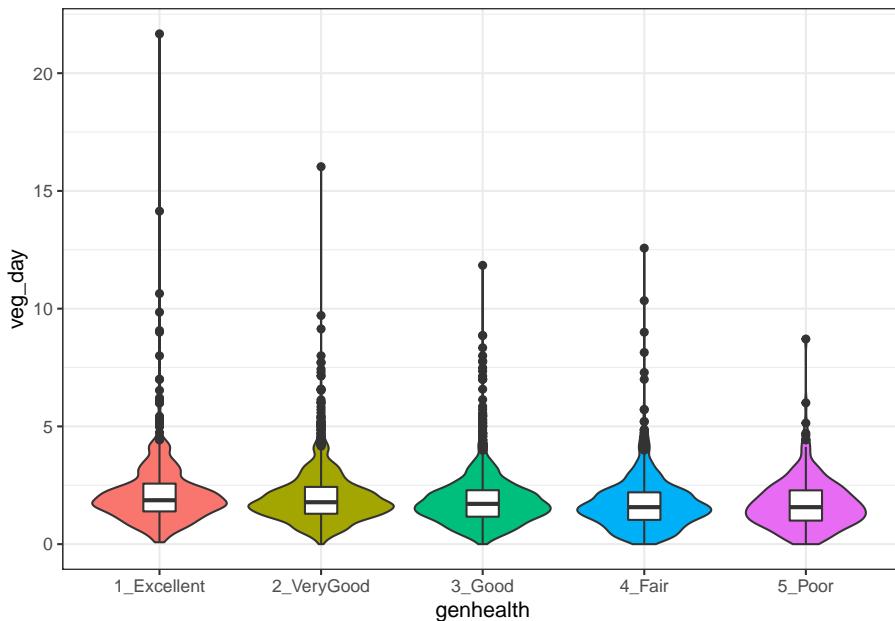
```
sm1 %>% group_by(genh) %>%
  summarize(n(), mean(veg_day), sd(veg_day), median(veg_day))

# A tibble: 5 x 5
  genh `n()`  `mean(veg_day)` `sd(veg_day)` `median(veg_day)`
  <ord> <int>        <dbl>          <dbl>          <dbl>
1 1_E   822         2.16          1.46          1.87
2 2_VG  1805        1.99          1.13          1.78
3 3_G   1667        1.86          1.11          1.71
4 4_F   801          1.74          1.18          1.57
5 5_P   299          1.71          1.06          1.57
```

To actually see what's going on, we might build a comparison boxplot, or violin plot. The plot below shows both, together, with the violin plot helping to indicate

the skewed nature of the `veg_day` data and the boxplot indicating quartiles and outlying values within each `genhealth` category.

```
ggplot(sm1, aes(x = genhealth, y = veg_day)) +
  geom_violin(aes(fill = genhealth), trim = TRUE) +
  geom_boxplot(width = 0.2) +
  guides(fill = "none", color = "none") +
  theme_bw()
```



20.5.2 Describing the Proportional-Odds Cumulative Logit Model

To fit the ordinal logistic regression model (specifically, a proportional-odds cumulative-logit model) in this situation, we'll use the `polr` function in the `MASS` library.

- Our outcome is `genh`, which has five ordered levels, with 1-E best and 5-P worst.
- Our model will include one quantitative predictor, `veg_day`.

The model will have four logit equations:

- one estimating the log odds that `genh` will be less than or equal to 1 (i.e. `genhealth` = 1_Excellent,)
- one estimating the log odds that `genh` ≤ 2 (i.e. `genhealth` = 1_Excellent or 2_VeryGood,)

- another estimating the log odds that `genh` ≤ 3 (i.e. `genhealth = 1_Excellent, 2_VeryGood or 3_Good,`) and, finally,
- one estimating the log odds that `genh` ≤ 4 (i.e. `genhealth = 1_Excellent, 2_VeryGood, 3_Good or 4_Fair`)

That's all we need to estimate the five categories, since $\Pr(\text{genh} \leq 5) = 1$, because (5_Poor) is the maximum category for `genhealth`.

We'll have a total of five free parameters when we add in the slope for `veg_day`, and I'll label these parameters as $\zeta_1, \zeta_2, \zeta_3, \zeta_4$ and β_1 . The ζ s are read as "zeta" values, and the people who built the `polr` function use that term.

The four logistic equations that will be fit differ only by their intercepts. They are:

$$\text{logit}[\Pr(\text{genh} \leq 1)] = \log \frac{\Pr(\text{genh} \leq 1)}{\Pr(\text{genh} > 1)} = \zeta_1 - \beta_1 \text{vegday}$$

which describes the log odds of a `genh` value of 1 (Excellent) as compared to a `genh` value greater than 1 (which includes Very Good, Good, Fair and Poor).

The second logit model is:

$$\text{logit}[\Pr(\text{genh} \leq 2)] = \log \frac{\Pr(\text{genh} \leq 2)}{\Pr(\text{genh} > 2)} = \zeta_2 - \beta_1 \text{vegday}$$

which describes the log odds of a `genh` value of 1 (Excellent) or 2 (Very Good) as compared to a `genh` value greater than 2 (which includes Good, Fair and Poor).

Next we have:

$$\text{logit}[\Pr(\text{genh} \leq 3)] = \log \frac{\Pr(\text{genh} \leq 3)}{\Pr(\text{genh} > 3)} = \zeta_3 - \beta_1 \text{vegday}$$

which describes the log odds of a `genh` value of 1 (Excellent) or 2 (Very Good) or 3 (Good) as compared to a `genh` value greater than 3 (which includes Fair and Poor).

Finally, we have

$$\text{logit}[\Pr(\text{genh} \leq 4)] = \log \frac{\Pr(\text{genh} \leq 4)}{\Pr(\text{genh} > 4)} = \zeta_4 - \beta_1 \text{vegday}$$

which describes the log odds of a `genh` value of 4 or less, which includes Excellent, Very Good, Good and Fair as compared to a `genh` value greater than 4 (which is Poor).

Again, the intercept term is the only piece that varies across the four equations.

In this case, a positive coefficient β_1 for `veg_day` means that increasing the value of `veg_day` would increase the `genh` category (describing a worse level of general health, since higher values of `genh` are associated with worse health.)

20.5.3 Fitting a Proportional Odds Logistic Regression with `polr`

Our model `m1` will use proportional odds logistic regression (sometimes called an *ordered logit* model) to predict `genh` on the basis of `veg_day`. The `polr` function can help us do this. Note that we include `Hess = TRUE` to retain what is called the *Hessian* matrix, which lets R calculate standard errors more effectively in `summary` and other follow-up descriptions of the model.

```
m1 <- polr(genh ~ veg_day,
            data = sm1, Hess = TRUE)
```

```
summary(m1)
```

Call:

```
polr(formula = genh ~ veg_day, data = sm1, Hess = TRUE)
```

Coefficients:

	Value	Std. Error	t value
<code>veg_day</code>	-0.1847	0.02178	-8.48

Intercepts:

	Value	Std. Error	t value
<code>1-E 2_VG</code>	-2.0866	0.0584	-35.7590
<code>2_VG 3_G</code>	-0.4065	0.0498	-8.1621
<code>3_G 4_F</code>	1.0202	0.0521	19.5771
<code>4_F 5_P</code>	2.5002	0.0710	35.2163

Residual Deviance: 15669.85

AIC: 15679.85

```
confint(m1)
```

Waiting for profiling to be done...

	2.5 %	97.5 %
<code>-0.2277073</code>	-0.1423088	

20.6 Interpreting Model `m1`

20.6.1 Looking at Predictions

Consider two individuals:

- Harry, who eats an average of 2.0 servings of vegetables per day, so Harry's `veg_day` = 2, and
- Sally, who eats an average of 1.0 serving of vegetables per day, so Sally's `veg_day` = 1.

We're going to start by using our model `m1` to predict the `genh` for Harry and Sally, so we can see the effect (on the predicted `genh` probabilities) of a change of one unit in `veg_day`.

For example, what are the log odds that Harry, with `veg_day` = 2, will describe his `genh` as Excellent (`genh` ≤ 1)?

$$\text{logit}[\Pr(\text{genh} \leq 1)] = \zeta_1 - \beta_1 \text{veg_day} \text{logit}[\Pr(\text{genh} \leq 1)] = -2.0866 - (-0.1847)\text{veg_day} \text{logit}[\Pr(\text{genh}$$

That's not much help. So we'll convert it to a probability by taking the inverse logit. The formula is

$$\Pr(\text{genh} \leq 1) = \frac{\exp(\zeta_1 + \beta_1 \text{vegday})}{1 + \exp(\zeta_1 + \beta_1 \text{vegday})} = \frac{\exp(-1.7172)}{1 + \exp(-1.7172)} = \frac{0.180}{1.180} = 0.15$$

So the model estimates a 15% probability that Harry will describe his `genh` as Excellent.

OK. Now, what are the log odds that Harry, who eats 2 servings per day, will describe his `genh` as either Excellent or Very Good (`genh` ≤ 2)?

$$\text{logit}[\Pr(\text{genh} \leq 2)] = \zeta_2 - \beta_1 \text{veg_day} \text{logit}[\Pr(\text{genh} \leq 2)] = -0.4065 - (-0.1847)\text{veg_day} \text{logit}[\Pr(\text{genh}$$

Again, we'll convert this to a probability by taking the inverse logit.

$$\Pr(\text{genh} \leq 2) = \frac{\exp(\zeta_2 + \beta_1 \text{vegday})}{1 + \exp(\zeta_2 + \beta_1 \text{vegday})} = \frac{\exp(-0.0371)}{1 + \exp(-0.0371)} = \frac{0.964}{1.964} = 0.49$$

So, the model estimates a probability of .49 that Harry will describe his `genh` as either Excellent or Very Good, so by subtraction, that's a probability of .34 that Harry describes his `genh` as Very Good.

Happily, that's the last time we'll calculate this by hand.

20.6.2 Making Predictions for Harry (and Sally) with predict

Suppose Harry eats 2 servings of vegetables per day on average, and Sally eats 1.

```
temp.dat <- data.frame(name = c("Harry", "Sally"),
                        veg_day = c(2,1))

predict(m1, temp.dat, type = "p")

  1-E      2_VG      3_G      4_F      5_P
1 0.1522351 0.3385119 0.3097906 0.1457864 0.05367596
2 0.1298931 0.3148971 0.3246105 0.1667285 0.06387071
```

The predicted probabilities of falling into each category of `genh` are:

Subject	veg_day	Pr(1_E)	Pr(2_VG)	Pr(3_G)	Pr(4_F)	Pr(5_P)
Harry	2	15.2	33.9	31.0	14.6	5.4
Sally	1	13.0	31.4	32.5	16.7	6.4

- Harry has a higher predicted probability of lower (healthier) values of `genh`. Specifically, Harry has a higher predicted probability than Sally of falling into the Excellent and Very Good categories, and a lower probability than Sally of falling into the Good, Fair and Poor categories.
- This means that Harry, with a higher `veg_day` is predicted to have, on average, a lower (that is to say, healthier) value of `genh`.
- As we'll see, this association will be indicated by a negative coefficient of `veg_day` in the proportional odds logistic regression model.

20.6.3 Predicting the actual classification of `genh`

The default prediction approach actually returns the predicted `genh` classification for Harry and Sally, which is just the classification with the largest predicted probability. Here, for Harry that is Very Good, and for Sally, that's Good.

```
predict(m1, temp.dat)
```

```
[1] 2_VG 3_G
Levels: 1-E 2_VG 3_G 4_F 5_P
```

20.6.4 A Cross-Tabulation of Predictions?

```
addmargins(table(predict(m1), sm1$genh))
```

	1-E	2_VG	3_G	4_F	5_P	Sum
1-E	6	3	3	3	0	15
2_VG	647	1398	1198	525	192	3960
3_G	169	404	466	273	107	1419
4_F	0	0	0	0	0	0
5_P	0	0	0	0	0	0
Sum	822	1805	1667	801	299	5394

The `m1` model classifies all subjects in the `sm1` sample as either Excellent, Very Good or Good, and most subjects as Very Good or Good.

20.6.5 The Fitted Model Equations

```
summary(m1)
```

Call:

```
polr(formula = genh ~ veg_day, data = sm1, Hess = TRUE)
```

Coefficients:

	Value	Std. Error	t value
veg_day	-0.1847	0.02178	-8.48

Intercepts:

	Value	Std. Error	t value
1-E 2_VG	-2.0866	0.0584	-35.7590
2_VG 3_G	-0.4065	0.0498	-8.1621
3_G 4_F	1.0202	0.0521	19.5771
4_F 5_P	2.5002	0.0710	35.2163

Residual Deviance: 15669.85

AIC: 15679.85

The first part of the output provides coefficient estimates for the `veg_day` predictor, and these are followed by the estimates for the various model intercepts. Plugging in the estimates, we have:

$$\text{logit}[Pr(\text{genh} \leq 1)] = -2.0866 - (-0.1847)\text{vegday} \quad \text{logit}[Pr(\text{genh} \leq 2)] = -0.4065 - (-0.1847)\text{vegday}$$

Note that we can obtain these pieces separately as follows:

```
m1$zeta
```

1-E 2_VG	2_VG 3_G	3_G 4_F	4_F 5_P
-2.0866313	-0.4064704	1.0202035	2.5001655

shows the boundary intercepts, and

```
m1$coefficients
```

```
    veg_day  
-0.1847272
```

shows the regression coefficient for `veg_day`.

20.6.6 Interpreting the `veg_day` coefficient

The first part of the output provides coefficient estimates for the `veg_day` predictor.

- The estimated slope for `veg_day` is -0.1847
 - Remember Harry and Sally, who have the same values of `bmi` and `costprob`, but Harry eats one more serving than Sally does. We noted that Harry is predicted by the model to have a smaller (i.e. healthier) `genh` response than Sally.
 - So a negative coefficient here means that higher values of `veg_day` are associated with more of the probability distribution falling in lower values of `genh`.
 - We usually don't interpret this slope (on the log odds scale) directly, but rather exponentiate it.

20.6.7 Exponentiating the Slope Coefficient to facilitate Interpretation

We can compute the odds ratio associated with `veg_day` and its confidence interval as follows...

```
exp(coef(m1))
```

```
    veg_day  
0.8313311
```

```
exp(confint(m1))
```

Waiting for profiling to be done...

```
    2.5 %     97.5 %  
0.7963573 0.8673534
```

- So, if Harry eats one more serving of vegetables than Sally, our model predicts that Harry will have 83.1% of the odds of Sally of having a larger `genh` score. That means that Harry is likelier to have a smaller `genh` score.
 - Since `genh` gets larger as a person's general health gets worse (moves from Excellent towards Poor), this means that since Harry is predicted

to have smaller odds of a larger `genh` score, he is also predicted to have smaller odds of worse general health.

- Our 95% confidence interval around that estimated odds ratio of 0.831 is (0.796, 0.867). Since that interval is entirely below 1, the odds of having the larger (worse) `genh` for Harry are detectably lower than the odds for Sally.
- So, an increase in `veg_day` is associated with smaller (better) `genh` scores.

20.6.8 Comparison to a Null Model

We can fit a model with intercepts only to test the significance of `veg_day` in our model `m1`, using the `anova` function.

```
m0 <- polr(genh ~ 1, data = sm1)
```

```
anova(m1, m0)
```

```
Likelihood ratio tests of ordinal regression models
```

```
Response: genh
```

	Model	Resid.	df	Resid.	Dev	Test	Df	LR stat.	Pr(Chi)
1			1	5390	15744.89				
2	<code>veg_day</code>		5389	15669.85	1 vs 2	1	75.04297	0	

We could also compare model `m1` to the null model `m0` with AIC or BIC.

```
AIC(m1, m0)
```

	df	AIC
<code>m1</code>	5	15679.85
<code>m0</code>	4	15752.89

```
BIC(m1,m0)
```

	df	BIC
<code>m1</code>	5	15712.81
<code>m0</code>	4	15779.26

20.7 The Assumption of Proportional Odds

Let us calculate the odds for all levels of `genh` if a person eats two servings of vegetables. First, we'll get the probabilities, in another way, to demonstrate how to do so...

```
(prob.2 <- exp(m1$zeta - 2*m1$coefficients)/(1 + exp(m1$zeta - 2*m1$coefficients)))
1-E|2_VG 2_VG|3_G 3_G|4_F 4_F|5_P
0.1522351 0.4907471 0.8005376 0.9463240
(prob.1 <- exp(m1$zeta - 1*m1$coefficients)/(1 + exp(m1$zeta - 1*m1$coefficients)))
1-E|2_VG 2_VG|3_G 3_G|4_F 4_F|5_P
0.1298931 0.4447902 0.7694008 0.9361293
```

Now, we'll calculate the odds, first for a subject eating two servings:

```
(odds.2 = prob.2/(1-prob.2))
```

```
1-E|2_VG 2_VG|3_G 3_G|4_F 4_F|5_P
0.1795724 0.9636607 4.0134766 17.6303153
```

And here are the odds, for a subject eating one serving per day:

```
(odds.1 = prob.1/(1-prob.1))
```

```
1-E|2_VG 2_VG|3_G 3_G|4_F 4_F|5_P
0.1492841 0.8011211 3.3365277 14.6566285
```

Now, let's take the ratio of the odds for someone who eats two servings over the odds for someone who eats one.

```
odds.2/odds.1
```

```
1-E|2_VG 2_VG|3_G 3_G|4_F 4_F|5_P
1.20289 1.20289 1.20289 1.20289
```

They are all the same. The odds ratios are equal, which means they are proportional. For any level of `genh`, the estimated odds that a person who eats 2 servings has better (lower) `genh` is about 1.2 times the odds for someone who eats one serving. Those who eat more vegetables have higher odds of better (lower) `genh`. Less than 1 means lower odds, and more than 1 means greater odds.

Now, let's take the log of the odds ratios:

```
log(odds.2/odds.1)
```

```
1-E|2_VG 2_VG|3_G 3_G|4_F 4_F|5_P
0.1847272 0.1847272 0.1847272 0.1847272
```

That should be familiar. It is the slope coefficient in the model summary, without the minus sign. R tacks on a minus sign so that higher levels of predictors correspond to the ordinal outcome falling in the higher end of its scale.

If we exponentiate the slope estimated by R (-0.1847), we get 0.83. If we have two people, and A eats one more serving of vegetables on average than B, then

the estimated odds of A having a higher ‘genh’ (i.e. worse general health) are 83% as high as B’s.

20.7.1 Testing the Proportional Odds Assumption

One way to test the proportional odds assumption is to compare the fit of the proportional odds logistic regression to a model that does not make that assumption. A natural candidate is a **multinomial logit** model, which is typically used to model unordered multi-categorical outcomes, and fits a slope to each level of the `genh` outcome in this case, as opposed to the proportional odds logit, which fits only one slope across all levels.

Since the proportional odds logistic regression model is nested in the multinomial logit, we can perform a likelihood ratio test. To do this, we first fit the multinomial logit model, with the `multinom` function from the `nnet` package.

```
(m1_multi <- multinom(genh ~ veg_day, data = sm1))
```

```
# weights:  15 (8 variable)
initial  value 8681.308100
iter  10 value 7890.985276
final  value 7835.248471
converged

Call:
multinom(formula = genh ~ veg_day, data = sm1)
```

```
Coefficients:
(Intercept)    veg_day
2_VG      0.9791063 -0.09296694
3_G       1.0911990 -0.19260067
4_F       0.5708594 -0.31080687
5_P      -0.3583310 -0.34340619
```

Residual Deviance: 15670.5

AIC: 15686.5

The multinomial logit fits four intercepts and four slopes, for a total of 8 estimated parameters. The proportional odds logit, as we’ve seen, fits four intercepts and one slope, for a total of 5. The difference is 3, and we use that number in the sequence below to build our test of the proportional odds assumption.

```
LL_1 <- logLik(m1)
LL_1m <- logLik(m1_multi)
(G <- -2 * (LL_1[1] - LL_1m[1]))
```

[1] -0.6488392

```
pchisq(G, 3, lower.tail = FALSE)
```

```
[1] 1
```

The p value is very large, so it indicates that the proportional odds model fits about as well as the more complex multinomial logit. A non-significant p value here isn't always the best way to assess the proportional odds assumption, but it does provide some evidence of model adequacy.

20.8 Can model m1 be fit using rms tools?

Yes.

```
d <- datadist(sm1)
options(datadist = "d")
m1_lrm <- lrm(genh ~ veg_day, data = sm1, x = T, y = T)

m1_lrm
```

Logistic Regression Model

```
lrm(formula = genh ~ veg_day, data = sm1, x = T, y = T)
```

Frequencies of Responses

	1-E	2_VG	3_G	4_F	5_P
822	1805	1667	801	299	

	Model Likelihood		Discrimination		Rank Discrim.		
	Ratio	Test	Indexes		Indexes		
Obs	5394	LR chi2	75.04	R2	0.015	C	0.555
max deriv	3e-13	d.f.	1	g	0.216	Dxy	0.111
		Pr(> chi2)	<0.0001	gr	1.241	gamma	0.111
				gp	0.053	tau-a	0.082
				Brier	0.247		

	Coef	S.E.	Wald Z	Pr(> Z)
y>=2_VG	2.0866	0.0584	35.76	<0.0001
y>=3_G	0.4065	0.0498	8.16	<0.0001
y>=4_F	-1.0202	0.0521	-19.58	<0.0001
y>=5_P	-2.5002	0.0710	-35.22	<0.0001
veg_day	-0.1847	0.0218	-8.48	<0.0001

The model is highly significant (remember the large sample size) but very weak, with a Nagelkerke R^2 of 0.015, and a C statistic of 0.555.

```
summary(m1_lrm)
```

	Effects				Response : genh		
Factor	Low	High	Diff.	Effect	S.E.	Lower	0.95
veg_day	1.21	2.36	1.15	-0.21243	0.025051	-0.26153	-0.16334
Odds Ratio	1.21	2.36	1.15	0.80861	NA	0.76987	0.84931

A change from 1.21 to 2.36 servings in `veg_day` is associated with an odds ratio of 0.81, with 95% confidence interval (0.77, 0.85). Since these values are all below 1, we have a clear indication of a statistically detectable effect of `veg_day` with higher `veg_day` associated with lower `genh`, which means, in this case, better health.

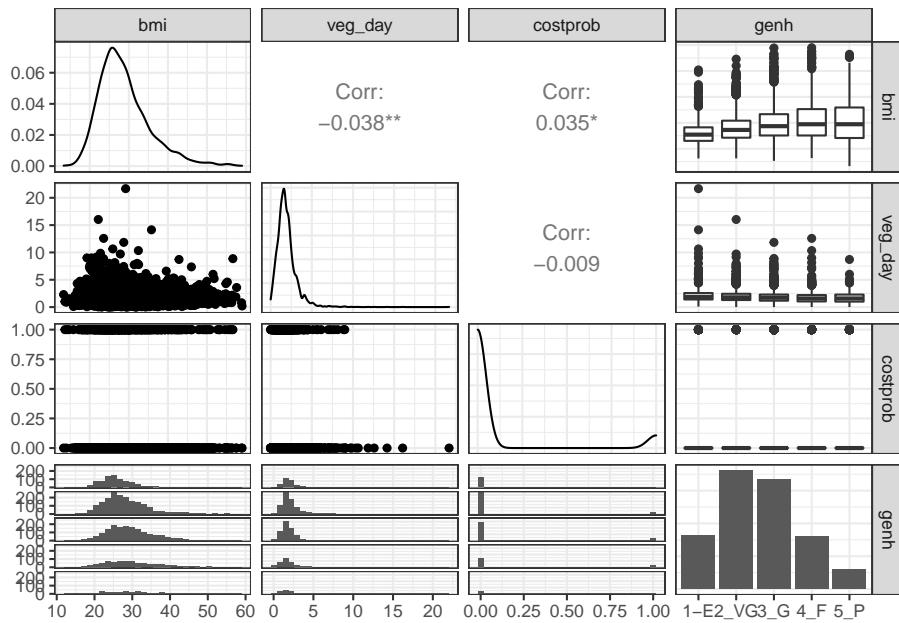
There is also a tool in `rms` called `orm` which may be used to fit a wide array of ordinal regression models. I suggest you read Frank Harrell's book on *Regression Modeling Strategies* if you want to learn more.

20.9 Building a Three-Predictor Model

Now, we'll model `genh` using `veg_day`, `bmi` and `costprob`.

20.9.1 Scatterplot Matrix

```
GGally::ggpairs(sml %>%
  select(bmi, veg_day, costprob, genh))
```



We might choose to plot the `costprob` data as a binary factor, rather than the raw 0-1 numbers included above, but not at this time.

20.9.2 Our Three-Predictor Model, m2

```
m2 <- polr(genh ~ veg_day + bmi + costprob, data = sm1)
```

```
summary(m2)
```

Re-fitting to get Hessian

Call:

```
polr(formula = genh ~ veg_day + bmi + costprob, data = sm1)
```

Coefficients:

	Value	Std. Error	t value
veg_day	-0.17130	0.021783	-7.864
bmi	0.06673	0.003855	17.311
costprob	0.96825	0.084871	11.409

Intercepts:

	Value	Std. Error	t value
1-E 2_VG	-0.1252	0.1229	-1.0183
2_VG 3_G	1.6358	0.1234	13.2572

```
3_G|4_F    3.1534  0.1294    24.3755
4_F|5_P    4.6881  0.1412    33.1928
```

Residual Deviance: 15229.24
AIC: 15243.24

This model contains four intercepts (to cover the five `genh` categories) and three slopes (one each for `veg_day`, `bmi` and `costprob`.)

20.9.3 Does the three-predictor model outperform `m1`?

```
anova(m1, m2)
```

Likelihood ratio tests of ordinal regression models

Response: `genh`

	Model	Resid. df	Resid. Dev	Test	Df	LR stat.	Pr(Chi)
1	<code>veg_day</code>	5389	15669.85				
2	<code>veg_day + bmi + costprob</code>	5387	15229.24	1 vs 2	2	440.6041	0

There is a statistically significant improvement in fit from model 1 to model 2. The AIC and BIC are also better for the three-predictor model than they were for the model with `veg_day` alone.

```
AIC(m1, m2)
```

df	AIC
<code>m1</code>	5 15679.85
<code>m2</code>	7 15243.24

```
BIC(m1, m2)
```

df	BIC
<code>m1</code>	5 15712.81
<code>m2</code>	7 15289.40

20.9.4 Wald tests for individual predictors

To obtain the appropriate Wald tests, we can use `lrm` to fit the model instead.

```
d <- datadist(sm1)
options(datadist = "d")
m2_lrm <- lrm(genh ~ veg_day + bmi + costprob,
               data = sm1, x = T, y = T)
m2_lrm
```

Logistic Regression Model

```
lrm(formula = genh ~ veg_day + bmi + costprob, data = sm1, x = T,
y = T)
```

Frequencies of Responses

	1-E	2_VG	3_G	4_F	5_P
822	1805	1667	801	299	

	Model Likelihood		Discrimination		Rank Discrim.		
	Ratio	Test		Indexes		Indexes	
Obs	5394	LR chi2	515.65	R2	0.096	C	0.629
max deriv	4e-09	d.f.	3	g	0.620	Dxy	0.258
		Pr(> chi2)	<0.0001	gr	1.859	gamma	0.258
				gp	0.143	tau-a	0.192
				Brier	0.231		
	Coef	S.E.	Wald Z	Pr(> Z)			
y>=2_VG	0.1252	0.1229	1.02	0.3085			
y>=3_G	-1.6358	0.1234	-13.26	<0.0001			
y>=4_F	-3.1534	0.1294	-24.38	<0.0001			
y>=5_P	-4.6881	0.1412	-33.19	<0.0001			
veg_day	-0.1713	0.0218	-7.86	<0.0001			
bmi	0.0667	0.0039	17.31	<0.0001			
costprob	0.9683	0.0849	11.41	<0.0001			

It appears that each of the added predictors (bmi and costprob) adds statistically detectable value to the model.

20.9.5 A Cross-Tabulation of Predictions?

```
addmargins(table(predict(m2), sm1$genh))
```

	1-E	2_VG	3_G	4_F	5_P	Sum
1-E	6	5	4	1	0	16
2_VG	686	1295	950	388	141	3460
3_G	128	495	672	374	135	1804
4_F	1	9	38	36	20	104
5_P	1	1	3	2	3	10
Sum	822	1805	1667	801	299	5394

At least the `m2` model predicted that a few of the cases will fall in the Fair and Poor categories, but still, this isn't impressive.

20.9.6 Interpreting the Effect Sizes

We can do this in two ways:

- By exponentiating the `polr` output, which shows the effect of increasing each predictor by a single unit
 - Increasing `veg_day` by 1 serving while holding the other predictors constant is associated with reducing the odds (by a factor of 0.84 with 95% CI 0.81, 0.88) of higher values of `genh`: hence increasing `veg_day` is associated with increasing the odds of a response indicating better health.
 - Increasing `bmi` by 1 kg/m² while holding the other predictors constant is associated with increasing the odds (by a factor of 1.07 with 95% CI 1.06, 1.08) of higher values of `genh`: hence increasing `bmi` is associated with reducing the odds of a response indicating better health.
 - Increasing `costprob` from 0 to 1 while holding the other predictors constant is associated with an increase (by a factor of 2.63 with 95% CI 2.23, 3.11) of a higher `genh` value. Since higher `genh` values indicate worse health, those with `costprob` = 1 are modeled to have generally worse health.

```
exp(coef(m2))
```

```
veg_day      bmi   costprob
0.8425722  1.0690045 2.6333356
```

```
exp(confint(m2))
```

Waiting for profiling to be done...

Re-fitting to get Hessian

2.5 %	97.5 %
veg_day	0.8071346 0.879096
bmi	1.0609722 1.077126
costprob	2.2301783 3.110633

- Or by looking at the summary provided by `lrm`, which like all such summaries produced by `rms` shows the impact of moving from the 25th to the 75th percentile on all continuous predictors.

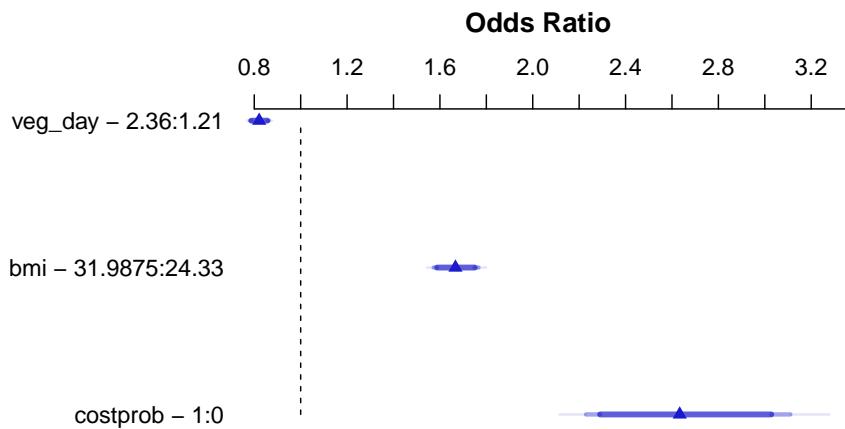
```
summary(m2_lrm)
```

Effects	Response : genh
---------	-----------------

```

Factor      Low   High  Diff.  Effect    S.E.      Lower 0.95 Upper 0.95
veg_day     1.21  2.360 1.1500 -0.19699 0.025051 -0.24609 -0.14789
Odds Ratio  1.21  2.360 1.1500  0.82120        NA  0.78185  0.86252
bmi        24.33 31.988 7.6575  0.51097 0.029515  0.45312  0.56882
Odds Ratio 24.33 31.988 7.6575  1.66690        NA  1.57320  1.76620
costprob    0.00  1.000 1.0000  0.96825 0.084871  0.80191  1.13460
Odds Ratio  0.00  1.000 1.0000  2.63330        NA  2.22980  3.10990
plot(summary(m2_lrm))

```



20.9.7 Quality of the Model Fit

Model `m2`, as we can see from the `m2_lrm` output, is still weak, with a Nagelkerke R^2 of 0.10, and a C statistic of 0.63.

20.9.8 Validating the Summary Statistics in `m2_lrm`

```

set.seed(43203); validate(m2_lrm)

      index.orig training      test optimism index.corrected n
Dxy       0.2583  0.2625  0.2579  0.0046       0.2537 40
R2        0.0964  0.0989  0.0960  0.0029       0.0935 40
Intercept 0.0000  0.0000 -0.0023  0.0023      -0.0023 40

```

Slope	1.0000	1.0000	0.9876	0.0124	0.9876	40
Emax	0.0000	0.0000	0.0031	0.0031	0.0031	40
D	0.0954	0.0980	0.0950	0.0030	0.0924	40
U	-0.0004	-0.0004	-1.5149	1.5146	-1.5149	40
Q	0.0958	0.0984	1.6099	-1.5115	1.6073	40
B	0.2314	0.2308	0.2315	-0.0007	0.2321	40
g	0.6199	0.6278	0.6182	0.0096	0.6103	40
gp	0.1434	0.1448	0.1430	0.0018	0.1417	40

As in our work with binary logistic regression, we can convert the index-corrected Dxy to an index-corrected C with $C = 0.5 + (Dxy/2)$. Both the R^2 and C statistics are pretty consistent with what we saw above.

20.9.9 Testing the Proportional Odds Assumption

Again, we'll fit the analogous multinomial logit model, with the `multinom` function from the `nnet` package.

```
(m2_multi <- multinom(genh ~ veg_day + bmi + costprob,
                        data = sm1))

# weights:  25 (16 variable)
initial  value 8681.308100
iter  10 value 8025.745934
iter  20 value 7605.878993
final  value 7595.767250
converged

Call:
multinom(formula = genh ~ veg_day + bmi + costprob, data = sm1)

Coefficients:
(Intercept)    veg_day        bmi   costprob
2_VG     -0.9126285 -0.0905958 0.06947231 0.3258568
3_G      -2.1886806 -0.1893454 0.11552563 1.0488262
4_F      -3.4095145 -0.3056028 0.13679908 1.4422074
5_P      -4.2629564 -0.3384199 0.13178846 1.8612088

Residual Deviance: 15191.53
AIC: 15223.53
```

The multinomial logit fits four intercepts and 12 slopes, for a total of 16 estimated parameters. The proportional odds logit in model `m2`, as we've seen, fits four intercepts and three slopes, for a total of 7. The difference is 9, and we use that number in the sequence below to build our test of the proportional odds assumption.

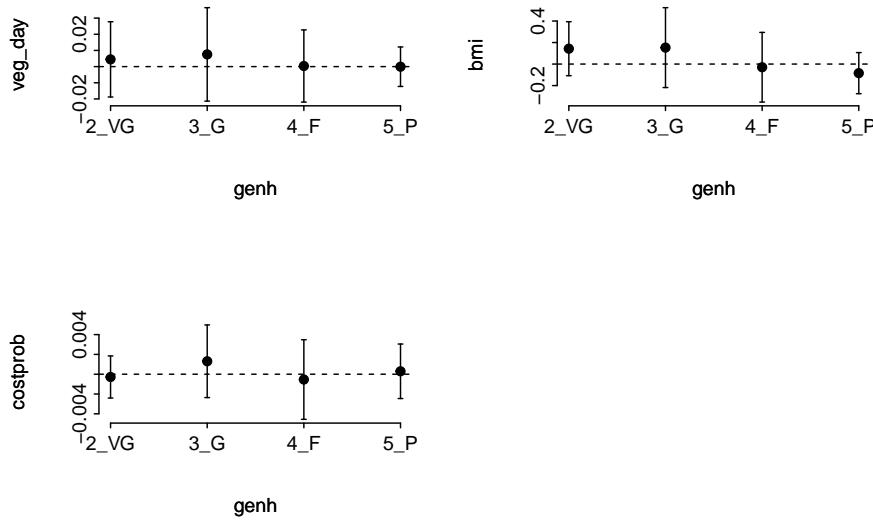
```
LL_2 <- logLik(m2)
LL_2m <- logLik(m2_multi)
(G <- -2 * (LL_2[1] - LL_2m[1]))
```

```
[1] 37.70952
pchisq(G, 9, lower.tail = FALSE)
```

```
[1] 1.965186e-05
```

The result is highly significant, suggesting that we have a problem somewhere with the proportional odds assumption. When this happens, I suggest you build the following plot of score residuals:

```
par(mfrow = c(2,2))
resid(m2_lrm, 'score.binary', pl=TRUE)
par(mfrow= c(1,1))
```



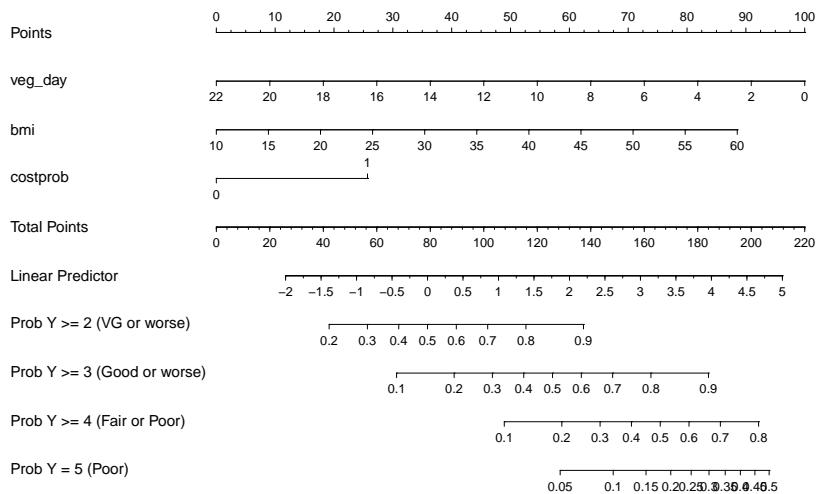
From this plot, **bmi** (especially) and **costprob** vary as we move from the Very Good toward the Poor cutpoints, relative to **veg_day**, which is more stable.

20.9.10 Plotting the Fitted Model

20.9.10.1 Nomogram

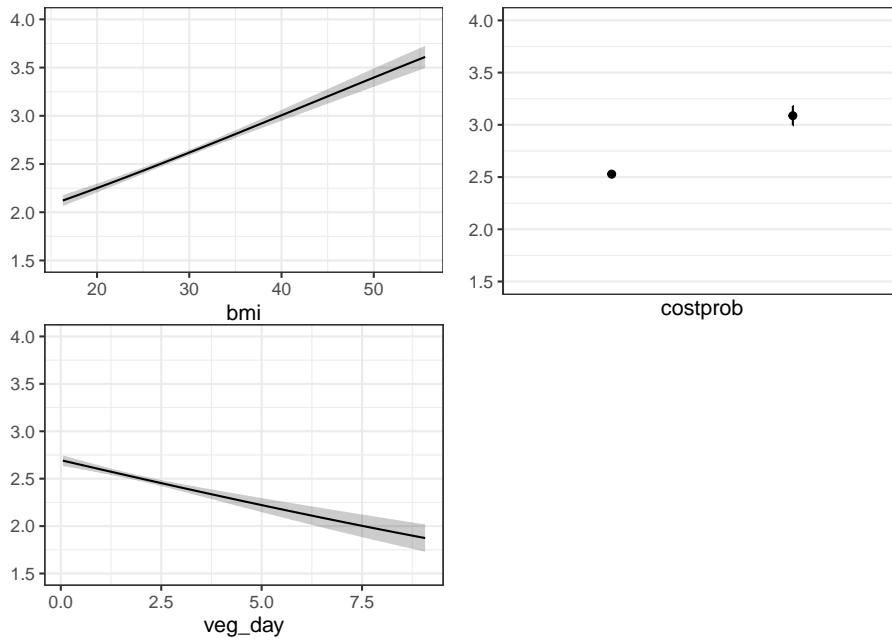
```
fun.ge3 <- function(x) plogis(x - m2_lrm$coef[1] + m2_lrm$coef[2])
fun.ge4 <- function(x) plogis(x - m2_lrm$coef[1] + m2_lrm$coef[3])
fun.ge5 <- function(x) plogis(x - m2_lrm$coef[1] + m2_lrm$coef[4])

plot(nomogram(m2_lrm, fun=list('Prob Y >= 2 (VG or worse)' = plogis,
                                'Prob Y >= 3 (Good or worse)' = fun.ge3,
                                'Prob Y >= 4 (Fair or Poor)' = fun.ge4,
                                'Prob Y = 5 (Poor)' = fun.ge5)))
```



20.9.10.2 Using Predict and showing mean prediction on 1-5 scale

```
ggplot(Predict(m2_lrm, fun = Mean(m2_lrm, code = TRUE)))
```



The nomogram and Predict results would be more interesting, of course, if we included a spline or interaction term. Let's do that in model `m3_lrm`, and also add the `incomegroup` information.

20.10 A Larger Model, including income group

```
m3_lrm <- lrm(gen_n ~ rcs(veg_day, 3) + rcs(bmi, 4) +
  incomegroup + catg(costprob) +
  bmi %ia% costprob,
  data = sm1, x = T, y = T)

m3_lrm
```

Logistic Regression Model

```
lrm(formula = gen_n ~ rcs(veg_day, 3) + rcs(bmi, 4) + incomegroup +
  catg(costprob) + bmi %ia% costprob, data = sm1, x = T, y = T)
```

Frequencies of Responses

1	2	3	4	5
822	1805	1667	801	299

	Model	Likelihood	Discrimination		Rank	Discrim.
		Ratio Test	Indexes			
Obs	5394	LR chi2	1190.35	R2	0.209	C 0.696
max deriv	7e-12	d.f.	14	g	1.036	Dxy 0.391
		Pr(> chi2)	<0.0001	gr	2.819	gamma 0.392
				gp	0.226	tau-a 0.291
				Brier	0.214	
	Coef	S.E.	Wald Z	Pr(> Z)		
y>=2	3.7535	0.4852	7.74	<0.0001		
y>=3	1.8717	0.4838	3.87	0.0001		
y>=4	0.2035	0.4831	0.42	0.6737		
y>=5	-1.4386	0.4846	-2.97	0.0030		
veg_day	-0.2602	0.0633	-4.11	<0.0001		
veg_day'	0.1756	0.0693	2.53	0.0113		
bmi	-0.0325	0.0203	-1.60	0.1086		
bmi'	0.5422	0.0989	5.48	<0.0001		
bmi ''	-1.4579	0.2663	-5.47	<0.0001		
incomegroup=10-14K	0.2445	0.1705	1.43	0.1516		
incomegroup=15-19K	-0.2626	0.1582	-1.66	0.0969		
incomegroup=20-24K	-0.6434	0.1501	-4.29	<0.0001		
incomegroup=25-34K	-0.7427	0.1459	-5.09	<0.0001		
incomegroup=35-49K	-1.1621	0.1415	-8.21	<0.0001		
incomegroup=50-74K	-1.4579	0.1418	-10.28	<0.0001		
incomegroup=75K+	-1.8592	0.1361	-13.66	<0.0001		
costprob=1	1.4576	0.3528	4.13	<0.0001		
bmi * costprob	-0.0259	0.0116	-2.24	0.0250		

Another option here would have been to consider building `incomegroup` as a `scored` variable, with an order on its own, but I won't force that here. Here's the `polr` version...

```
m3 <- polr(genh ~ rcs(veg_day, 3) + rcs(bmi, 4) +
  incomegroup + costprob +
  bmi %ia% costprob, data = sm1)
```

20.10.1 Cross-Tabulation of Predicted/Observed Classifications

```
addmargins(table(predict(m3), sm1$genh))
```

```
1-E 2_VG 3_G 4_F 5_P Sum
```

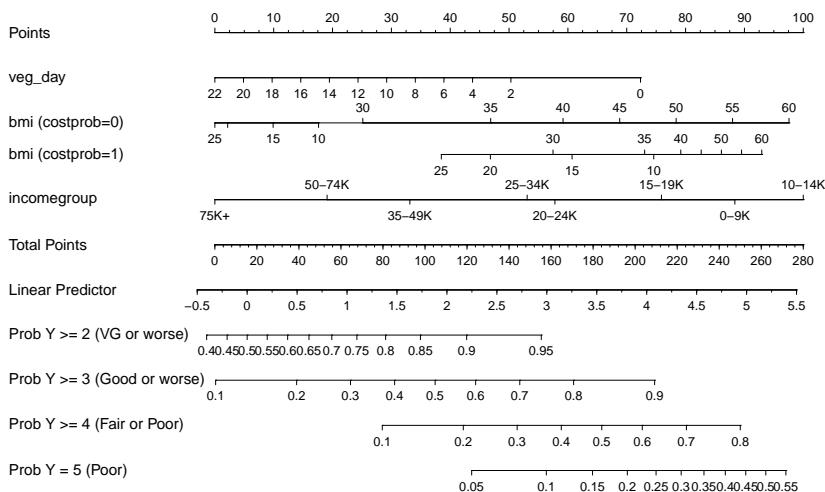
1-E	3	2	0	0	5
2_VG	642	1200	815	221	49 2927
3_G	170	565	754	468	182 2139
4_F	7	37	96	108	65 313
5_P	0	1	2	4	3 10
Sum	822	1805	1667	801	299 5394

This model predicts more Fair results, but still far too many Very Good with no Excellent at all.

20.10.2 Nomogram

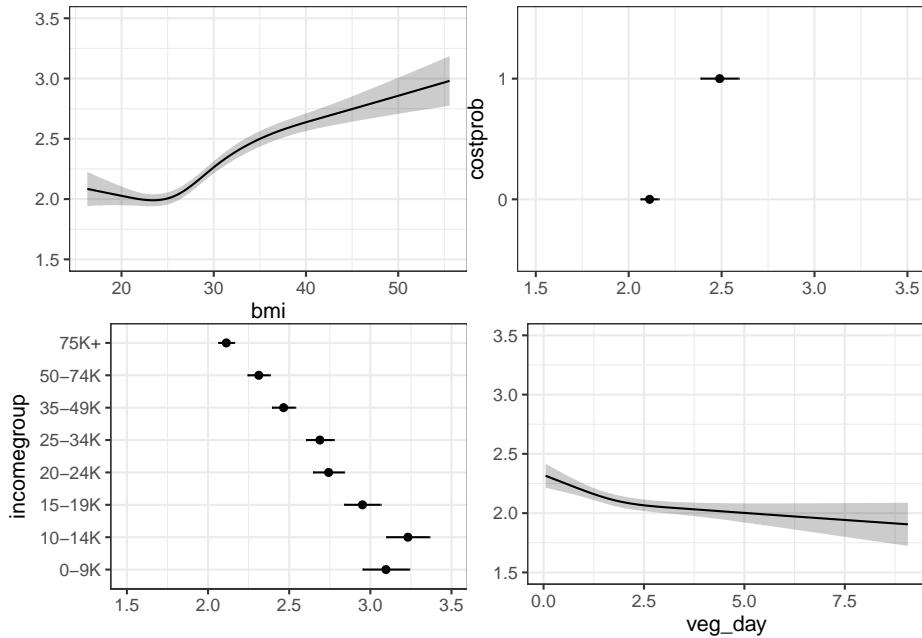
```
fun.ge3 <- function(x) plogis(x - m3_lrm$coef[1] + m3_lrm$coef[2])
fun.ge4 <- function(x) plogis(x - m3_lrm$coef[1] + m3_lrm$coef[3])
fun.ge5 <- function(x) plogis(x - m3_lrm$coef[1] + m3_lrm$coef[4])

plot(nomogram(m3_lrm, fun=list('Prob Y >= 2 (VG or worse)' = plogis,
                                'Prob Y >= 3 (Good or worse)' = fun.ge3,
                                'Prob Y >= 4 (Fair or Poor)' = fun.ge4,
                                'Prob Y = 5 (Poor)' = fun.ge5)))
```



20.10.3 Using Predict and showing mean prediction on 1-5 scale

```
ggplot(Predict(m3_lrm, fun = Mean(m3_lrm, code = TRUE)))
```



Here, we're plotting the mean score on the 1-5 gen_n scale.

20.10.4 Validating the Summary Statistics in m3_lrm

```
set.seed(43221); validate(m3_lrm)
```

	index.orig	training	test	optimism	index.corrected	n
Dxy	0.3915	0.3928	0.3899	0.0029	0.3886	40
R2	0.2093	0.2112	0.2071	0.0041	0.2053	40
Intercept	0.0000	0.0000	0.0019	-0.0019	0.0019	40
Slope	1.0000	1.0000	0.9871	0.0129	0.9871	40
Emax	0.0000	0.0000	0.0032	0.0032	0.0032	40
D	0.2205	0.2227	0.2179	0.0049	0.2156	40
U	-0.0004	-0.0004	-1.4667	1.4663	-1.4667	40
Q	0.2209	0.2231	1.6845	-1.4614	1.6823	40
B	0.2137	0.2134	0.2142	-0.0007	0.2145	40
g	1.0363	1.0410	1.0276	0.0133	1.0229	40
gp	0.2262	0.2265	0.2244	0.0022	0.2240	40

Still not very impressive, but much better than where we started. It's not crazy to suggest that in new data, we might expect a Nagelkerke R^2 of 0.205 and a C statistic of $0.5 + (0.3886/2) = 0.6943$.

20.11 References for this Chapter

1. Some of the material here is adapted from <http://stats.idre.ucla.edu/r/dae/ordinal-logistic-regression/>.
2. I also found great guidance at <http://data.library.virginia.edu/fitting-and-interpreting-a-proportional-odds-model/>
3. Other parts are based on the work of Jeffrey S. Simonoff (2003) *Analyzing Categorical Data* in Chapter 10. Related data and R code are available at <http://people.stern.nyu.edu/jsimonof/AnalCatData/Splus/>.
4. Another good source for a simple example is <https://onlinecourses.science.psu.edu/stat504/node/177>.
5. Also helpful is <https://onlinecourses.science.psu.edu/stat504/node/178> which shows a more complex example nicely.

Chapter 21

Analyzing Literary Styles with Multinomial Logistic Regression

21.1 The Authorship Example

This example is based on the work of Jeffrey S. Simonoff (2003) *Analyzing Categorical Data* in Chapter 10. Related data and R code are available at <http://people.stern.nyu.edu/jsimonof/AnalCatData/Splus/>. Meanwhile, the data set and analysis are based on the work of Peng RD and Hengartner NW (2002) Quantitative analysis of literary styles, *The American Statistician*, 56, 175-185.

The `authorship.csv` data file contains 841 rows. Each row describes a block of text that contains 1700 total words from one of several books by four authors: Jane Austen (samples from 7 books), Jack London (6 books), John Milton (2 books), or William Shakespeare (12 books). The data include counts within the blocks of text of 69 function words, such as “a,” “by,” “no,” “that” and “with.” The goal of our analysis, mirroring that of Simonoff, will be to use the incidence of these function words to build a model that distinguishes the authors.

```
authorship$Author <- factor(authorship$Author,  
  levels = c("Shakespeare", "Austen", "London", "Milton"))  
  
authorship  
  
# A tibble: 841 x 71  
BookID Author    a    all    also    an    and    any    are    as    at    be  
<dbl> <fct> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl>
```

```

1      1 Austen    46    12    0    3    66    9    4    16    13    13
2      1 Austen    35    10    0    7    44    4    3    18    16    9
3      1 Austen    46     2    0    3    40    1   13    11    9   23
4      1 Austen    40     7    0    4    64    3    3    20    13   20
5      1 Austen    29     5    0    6    52    5   14    17    6   16
6      1 Austen    27     8    0    3    42    2   15    11   14   12
7      1 Austen    34     8    0   15    44    2    6    16   14   11
8      1 Austen    38     6    1    2    67    3    6    17    4   21
9      1 Austen    34    12    0    5    50    2    8    7   13    7
10     1 Austen   54     8    0    6    44    4    8   13   17   14
# ... with 831 more rows, and 59 more variables: been <dbl>, but <dbl>,
# by <dbl>, can <dbl>, do <dbl>, down <dbl>, even <dbl>, every <dbl>,
# `for` <dbl>, from <dbl>, had <dbl>, has <dbl>, have <dbl>, her <dbl>,
# his <dbl>, `if` <dbl>, `in` <dbl>, into <dbl>, is <dbl>, it <dbl>,
# its <dbl>, may <dbl>, more <dbl>, must <dbl>, my <dbl>, no <dbl>,
# not <dbl>, now <dbl>, of <dbl>, on <dbl>, one <dbl>, only <dbl>, or <dbl>,
# our <dbl>, should <dbl>, so <dbl>, some <dbl>, such <dbl>, than <dbl>, ...

```

To-morrow, and to-morrow, and to-morrow, Creeps in this petty
 pace from day to day, To the last syllable of recorded time; And all
 our yesterdays have lighted fools The way to dusty death. Out, out,
 brief candle! Life's but a walking shadow, a poor player, That struts
 and frets his hour upon the stage, And then is heard no more. It is
 a tale Told by an idiot, full of sound and fury, Signifying nothing.

21.2 Focus on 11 key words

Again, following Simonoff, we will focus on 11 words from the set of 69 potential predictors in the data, specifically...

- “be,” “been,” “had,” “it,” “may,” “not,” “on,” “the,” “upon,” “was” and
“which”

```

auth2 <- authorship %>%
  select(Author, BookID, be, been, had, it, may, not,
         on, the, upon, was, which)

auth2.long <- auth2 %>%
  gather("word", "n", 3:13)

auth2.long

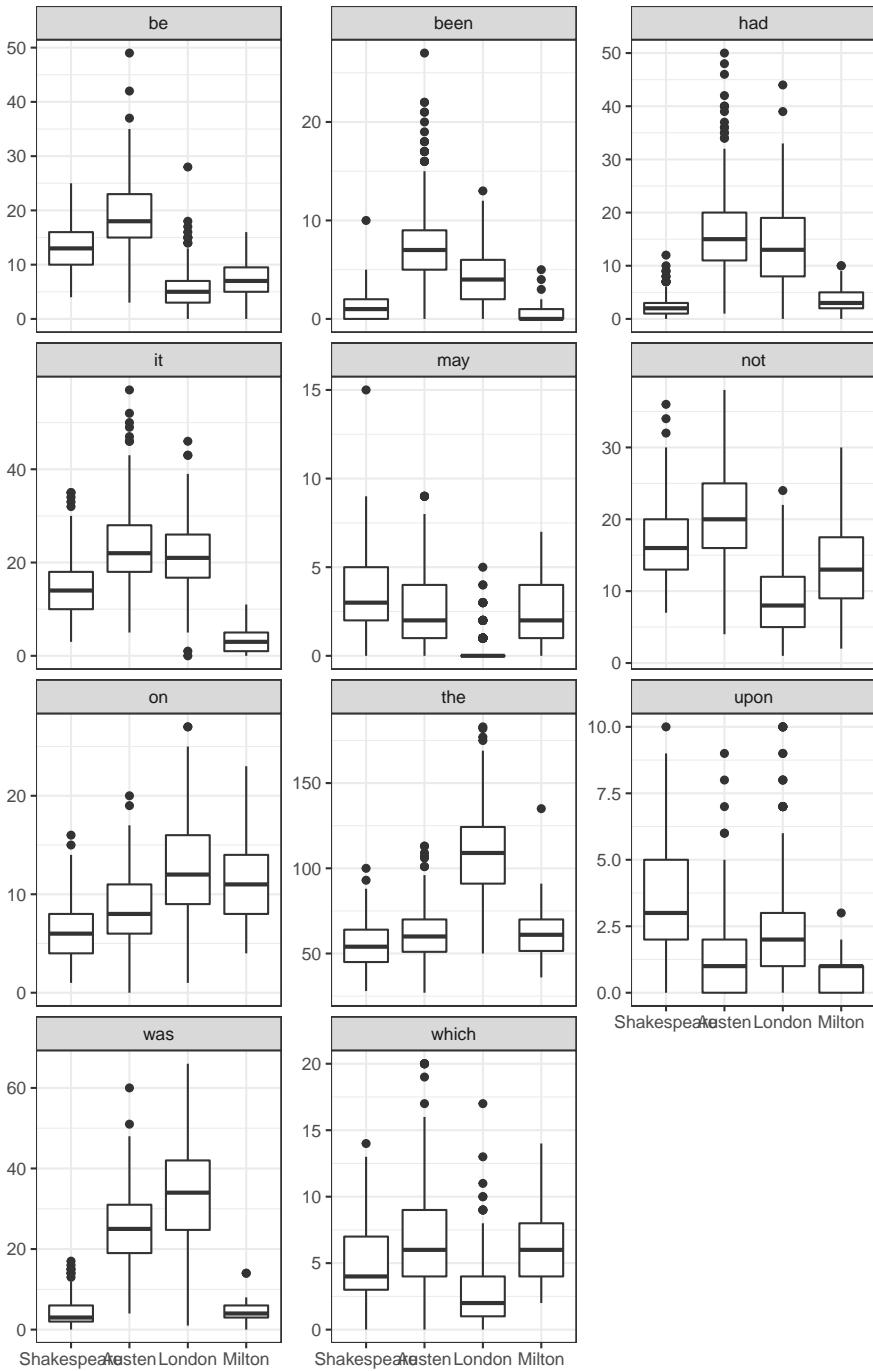
# A tibble: 9,251 x 4
  Author BookID word      n
  <fct>   <dbl> <chr> <dbl>
1 Austen     1 be      13

```

```
2 Austen      1 be      9
3 Austen      1 be     23
4 Austen      1 be     20
5 Austen      1 be     16
6 Austen      1 be     12
7 Austen      1 be     11
8 Austen      1 be     21
9 Austen      1 be      7
10 Austen     1 be    14
# ... with 9,241 more rows
```

21.2.1 Side by Side Boxplots

```
ggplot(auth2.long, aes(x = Author, y = n)) +
  geom_boxplot() +
  facet_wrap(~ word, ncol = 3, scales = "free_y") +
  theme_bw() +
  labs(x = "", y = "")
```



Oh! do not attack me with your watch. A watch is always too fast

or too slow. I cannot be dictated to by a watch.

21.3 A Multinomial Logistic Regression Model

Let's start with a multinomial model to predict `Author` on the basis of these 11 key predictors, using the `multinom` function from the `nnet` package.

```
authnom1 <- multinom(Author ~ be + been + had + it + may + not + on +
                      the + upon + was + which, data=authorship,
                      maxit=200)

# weights:  52 (36 variable)
initial  value 1165.873558
iter  10 value 293.806160
iter  20 value 273.554538
iter  30 value 192.309644
iter  40 value 71.091334
iter  50 value 48.419335
iter  60 value 46.808141
iter  70 value 46.184752
iter  80 value 46.026162
iter  90 value 45.932823
iter 100 value 45.897793
iter 110 value 45.868017
iter 120 value 45.863256
final  value 45.863228
converged

summary(authnom1)

Call:
multinom(formula = Author ~ be + been + had + it + may + not +
on + the + upon + was + which, data = authorship, maxit = 200)

Coefficients:
(Intercept)      be      been      had      it      may
Austen    -15.504834  0.48974946  0.5380318  0.4620513  0.00388835 -0.15025084
London    -14.671720 -0.07497073  0.1733116  0.4842272  0.08674782 -0.01590702
Milton    -1.776866 -0.10891178 -0.9127155  0.5319573 -0.82046587 -0.06760436
not          on      the      upon      was      which
Austen   -0.08861462  0.5967404 -0.02361614 -2.119001  0.7021371  0.10370827
London   -0.32567063  0.5749969  0.12039782 -1.914428  0.6767581 -0.59121054
Milton    0.05575887  0.5198173  0.08739368 -2.042475  0.3048202 -0.05939104

Std. Errors:
(Intercept)      be      been      had      it      may      not
```

```

Austen    4.892258 0.1643694 0.3117357 0.2695081 0.09554376 0.3008164 0.1078329
London    5.372898 0.1916618 0.3308759 0.2812555 0.11355697 0.4946804 0.1440760
Milton    5.417300 0.1613282 0.5910561 0.3187304 0.23015421 0.3500753 0.1309306
          on       the      upon      was      which
Austen 0.2213827 0.03457288 0.6426484 0.1808681 0.1646472
London 0.2251642 0.03088881 0.7072129 0.1768901 0.2542886
Milton 0.2223575 0.04783646 0.6436399 0.1820885 0.2111105

Residual Deviance: 91.72646
AIC: 163.7265

```

21.3.1 Testing Model 1

```

z1 <- summary(authnom1)$coefficients/summary(authnom1)$standard.errors
round(z1,2)

(Intercept)    be   been   had   it    may   not   on    the   upon   was
Austen        -3.17 2.98  1.73  1.71  0.04 -0.50 -0.82 2.70 -0.68 -3.30 3.88
London        -2.73 -0.39  0.52  1.72  0.76 -0.03 -2.26 2.55  3.90 -2.71 3.83
Milton        -0.33 -0.68 -1.54  1.67 -3.56 -0.19  0.43  2.34  1.83 -3.17 1.67
          which
Austen        0.63
London       -2.32
Milton       -0.28

p1 <- (1 - pnorm(abs(z1), 0, 1)) * 2
pander(round(p1,3))

```

Table 21.1: Table continues below

	(Intercept)	be	been	had	it	may	not
Austen	0.002	0.003	0.084	0.086	0.968	0.617	0.411
London	0.006	0.696	0.6	0.085	0.445	0.974	0.024
Milton	0.743	0.5	0.123	0.095	0	0.847	0.67

	on	the	upon	was	which
Austen	0.007	0.495	0.001	0	0.529
London	0.011	0	0.007	0	0.02
Milton	0.019	0.068	0.002	0.094	0.778

Simonoff suggests that “been” and “may” can be dropped. What do we think?

The proper function of man is to live, not to exist. I shall not waste my days in trying to prolong them. I shall use my time.

21.4 Model 2

```

authnom2 <- multinom(Author ~ be + had + it + not + on +
                      the + upon + was + which, data=authorship,
                      maxit=200)

# weights: 44 (30 variable)
initial value 1165.873558
iter 10 value 304.985478
iter 20 value 285.428679
iter 30 value 143.301103
iter 40 value 54.589791
iter 50 value 52.140470
iter 60 value 51.421454
iter 70 value 51.012790
iter 80 value 50.888718
iter 90 value 50.834262
iter 100 value 50.743136
final value 50.743111
converged

summary(authnom2)

Call:
multinom(formula = Author ~ be + had + it + not + on + the +
upon + was + which, data = authorship, maxit = 200)

Coefficients:
              (Intercept)      be      had       it      not      on
Austen     -16.55647  0.45995950  0.6698612  0.02621612 -0.03684654  0.4676716
London     -16.06419 -0.13378141  0.6052164  0.10517792 -0.27934022  0.4958923
Milton     -2.22344 -0.07031256  0.1737526 -0.81984885  0.05444678  0.5363108
                  the      upon      was      which
Austen   -0.001852454 -1.950761  0.6543956  0.06363998
London    0.128565811 -1.643829  0.6418607 -0.54690144
Milton    0.074236636 -1.762533  0.2932065 -0.08748272

Std. Errors:
              (Intercept)      be      had       it      not      on
Austen      4.723001  0.1293729  0.2201823  0.08657746  0.08771157  0.1949021
London      5.202732  0.1587639  0.2306803  0.10117217  0.11608348  0.2072383
Milton      4.593806  0.1499103  0.2057258  0.21551377  0.12103678  0.1895226

```

```

          the      upon      was      which
Austen 0.02945139 0.5620273 0.1524982 0.1466250
London 0.02739965 0.6219927 0.1512911 0.2087120
Milton 0.04463721 0.6246766 0.1601393 0.1928361

Residual Deviance: 101.4862
AIC: 161.4862

```

21.4.1 Comparing Model 2 to Model 1

```

anova(authnom1, authnom2)

Likelihood ratio tests of Multinomial Models

Response: Author
              Model Resid. df
1           be + had + it + not + on + the + upon + was + which      2493
2 be + been + had + it + may + not + on + the + upon + was + which      2487
  Resid. Dev   Test   Df LR stat.  Pr(Chi)
1 101.48622
2 91.72646 1 vs 2       6 9.759767 0.1351402

```

21.4.2 Testing Model 2

```

z2 <- summary(authnom2)$coefficients/summary(authnom2)$standard.errors
round(z2,2)

          (Intercept)    be    had     it     not     on     the     upon     was     which
Austen      -3.51  3.56  3.04  0.30 -0.42  2.40 -0.06 -3.47  4.29  0.43
London      -3.09 -0.84  2.62  1.04 -2.41  2.39  4.69 -2.64  4.24 -2.62
Milton      -0.48 -0.47  0.84 -3.80  0.45  2.83  1.66 -2.82  1.83 -0.45

p2 <- (1 - pnorm(abs(z2), 0, 1)) * 2
round(p2,3)

          (Intercept)    be    had     it     not     on     the     upon     was     which
Austen      0.000  0.000  0.002  0.762  0.674  0.016  0.950  0.001  0.000  0.664
London      0.002  0.399  0.009  0.299  0.016  0.017  0.000  0.008  0.000  0.009
Milton      0.628  0.639  0.398  0.000  0.653  0.005  0.096  0.005  0.067  0.650

```

21.4.3 A little history

Simonoff has an interesting note: Consider the lifetimes of these four authors:

- William Shakespeare was born in 1564 and died in 1616
- John Milton was born in 1608 (44 years after Shakespeare) and died in 1674
- Jane Austen was born in 1775 (211 years after Shakespeare) and died in 1817
- Jack London was born in 1876 (312 years after Shakespeare) and died in 1916

How many significant coefficients does each author display relative to Shakespeare?

21.5 Classification Table

How well does this model (model 2) distinguish these authors based on blocks of 1700 words of text?

```
table(authorship$Author, predict(authnom2))
```

	Shakespeare	Austen	London	Milton
Shakespeare	168	3	1	1
Austen	4	308	5	0
London	0	1	294	1
Milton	2	0	1	52

Based on this classification table, I'd say it does a nice job. Almost 98% of the blocks of text are correctly classified.

Fly, envious Time, till thou run out thy race; Call on the lazy leaden-stepping hours, Whose speed is but the heavy plummet's pace; And glut thyself with what thy womb devours, Which is no more then what is false and vain, And merely mortal dross; So little is our loss, So little is thy gain. For when, as each thing bad thou hast entomb'd And last of all thy greedy self consumed, Then long Eternity shall greet our bliss, With an individual kiss; And Joy shall overtake us, as a flood, When every thing that is sincerely good, And perfectly divine, With truth, and peace, and love, shall ever shine, About the supreme throne Of Him, to whose happy-making sight, alone, When once our heavenly-guided soul shall climb, Then all this earthly grossness quit, Attired with stars, we shall for ever sit, Triumphing over Death, and Chance, and thee, O Time!

21.6 Probability Curves based on a Single Predictor

In situations where only one predictor is used, we can develop nice plots of estimated probabilities for each group as a function of the predictor. Suppose we look at the single word “been” (note that this was left out of Model 2.)

Note that the possible values for counts of “been” in the data range from 0 to 27...

```
summary(authorship$been)
```

Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
0.000	2.000	4.000	4.614	7.000	27.000

Now, we’ll build a model to predict the author based solely on the counts of the word “been.”

```
authnom3 <- multinom(Author ~ been,
                      data=authorship, maxit=200)

# weights: 12 (6 variable)
initial value 1165.873558
iter 10 value 757.915093
iter 20 value 755.454631
final value 755.454551
converged
```

Next, we’ll build a grid of the predicted log odds for each author (as compared to Shakespeare) using the fitted coefficients. The grid will cover every possible value from 0 to 27, increasing by 0.1, using the following trick in R.

```
beengrid <- cbind(1,c(0:270)/10)
austenlogit <- beengrid %*% coef(authnom3)[1,]
londonlogit <- beengrid %*% coef(authnom3)[2,]
miltonlogit <- beengrid %*% coef(authnom3)[3,]
```

Next, we’ll use that grid of logit values to estimate the fitted probabilities for each value of “been” between 0 and 27.

```
austenprob <- exp(austenlogit) /
  (exp(austenlogit) + exp(londonlogit) +
   exp(miltonlogit) + 1)
londonprob <- exp(londonlogit) /
  (exp(austenlogit) + exp(londonlogit) +
   exp(miltonlogit) + 1)
miltonprob <- exp(miltonlogit) /
  (exp(austenlogit) + exp(londonlogit) +
   exp(miltonlogit) + 1)
```

```

shakesprob <- 1 - austenprob - londonprob - miltonprob

been_dat <- data_frame(been_count = beengrid[,2],
                        austen = austenprob[,1],
                        london = londonprob[,1],
                        milton = miltonprob[,1],
                        shakespeare = shakesprob[,1])
been_dat

# A tibble: 271 x 5
  been_count austen london milton shakespeare
  <dbl>     <dbl>   <dbl>   <dbl>      <dbl>
1 0         0.0258  0.136  0.285      0.553
2 0.1       0.0288  0.147  0.272      0.553
3 0.2       0.0321  0.158  0.258      0.551
4 0.3       0.0357  0.171  0.245      0.548
5 0.4       0.0396  0.184  0.232      0.545
6 0.5       0.0438  0.197  0.219      0.540
7 0.6       0.0484  0.211  0.207      0.534
8 0.7       0.0534  0.225  0.195      0.527
9 0.8       0.0587  0.240  0.183      0.518
10 0.9      0.0644  0.256  0.171      0.509
# ... with 261 more rows

```

Now, we gather the data by author name and probability

```

been_dat_long <- been_dat %>%
  gather("name", "prob", 2:5)
been_dat_long

```

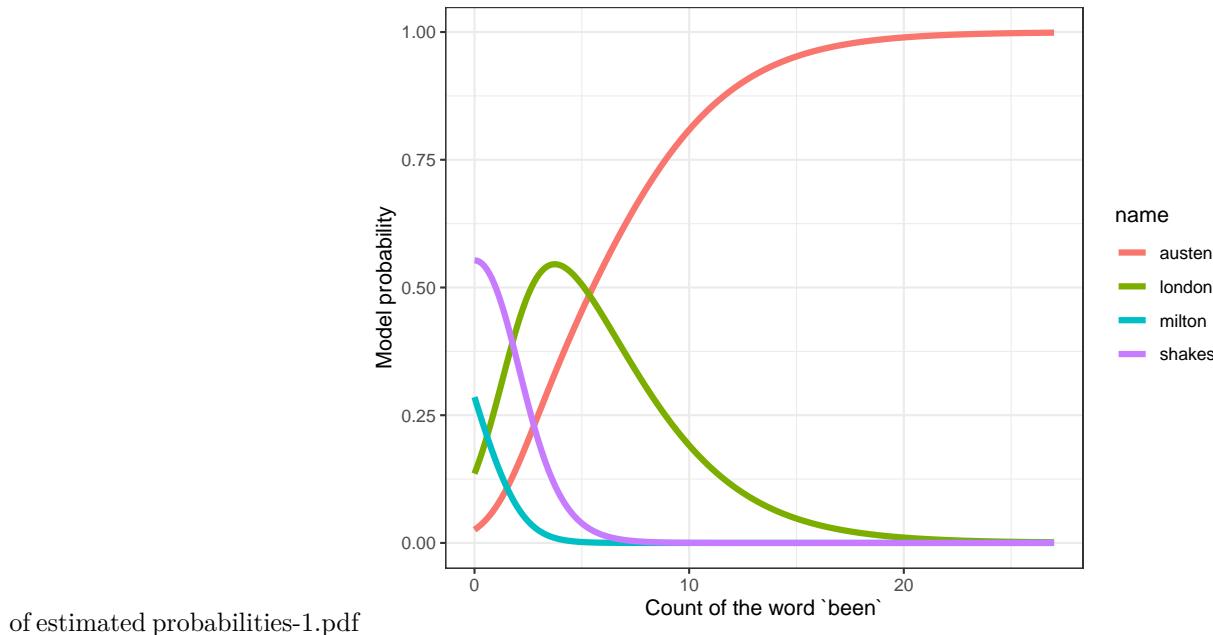
```

# A tibble: 1,084 x 3
  been_count name      prob
  <dbl>     <chr>    <dbl>
1 0         austen  0.0258
2 0.1       austen  0.0288
3 0.2       austen  0.0321
4 0.3       austen  0.0357
5 0.4       austen  0.0396
6 0.5       austen  0.0438
7 0.6       austen  0.0484
8 0.7       austen  0.0534
9 0.8       austen  0.0587
10 0.9      austen  0.0644
# ... with 1,074 more rows

```

21.6.1 Produce the Plot of Estimated Probabilities based on “been” counts

```
ggplot(bean_dat_long, aes(x = bean_count, y = prob,
                           col = name)) +
  geom_line(size = 1.5) +
  theme_bw() +
  labs(x = "Count of the word `been`",
       y = "Model probability")
```

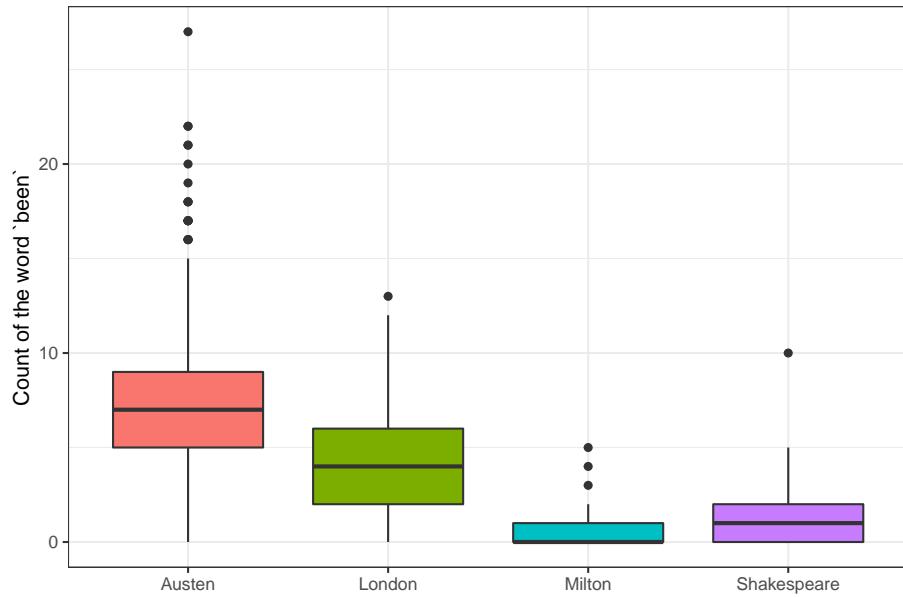


21.6.2 Boxplot of “been” counts

Compare this to what we see in the raw counts of the word “been.”

```
bean.long <- filter(auth2.long, word == "been")
bean.long$Auth <- fct_relevel(bean.long$Author,
                               "Austen", "London", "Milton", "Shakespeare")
# releveling to make the colors match the model plot

ggplot(bean.long, aes(x = Auth, y = n, fill = Auth)) +
  geom_boxplot() +
  guides(fill = "none") +
  theme_bw() +
  labs(x = "", y = "Count of the word `been`")
```



by side boxplot of been by author-1.pdf

21.6.3 Quote Sources

1. To-morrow, and to-morrow, and to-morrow . . . Shakespeare *Macbeth* Act 5.
2. Oh! do not attack me with your watch. . . . Jane Austen *Mansfield Park*
3. The proper function of man is to live, not to exist. . . . Jack London *The Bulletin* San Francisco 1916-12-02.
4. Fly, envious Time, till thou run out thy race . . . John Milton *On Time*

Chapter 22

Exploring Time To Event / Survival Data

In many medical studies, the main outcome variable is the time to the occurrence of a particular event.

- In a randomized controlled trial of cancer, for instance, surgery, radiation, and chemotherapy might be compared with respect to time from randomization and the start of therapy until death.
 - In this case, the event of interest is the death of a patient, but in other situations it might be remission from a disease, relief from symptoms or the recurrence of a particular condition.
 - Such observations are generally referred to by the generic term **survival data** even when the endpoint or event being considered is not death but something else.

These notes on survival analysis are just an introduction to the key ideas of the field. The PQHS department offers an entire course on survival analysis (PQHS 435) and I recommend that to those of you interested in deeper learning about the approaches we'll discuss, or in learning more about other approaches to survival analysis.

The OpenIntro Statistics extra material on Survival Analysis in R, written by David Diez is a very useful guide to survival analysis in R, using the **survival** package and supplemented by the **KMsurv** and **OIsurv** packages. A PDF version of that material is available, along with a full set of the code used in that guide.

22.1 An Outline of Key Topics Discussed in these Notes

In this chapter, we tackle the building blocks of survival analysis, and use R to work with survival objects.

- The Survival Function, $S(t)$
 - The Kaplan-Meier Estimate/Plot
 - Comparing Survival Functions with log rank test
- The Hazard Function, $H(t) = -\log(S(t))$
- Using `survival` and related packages in R

In the next chapter, we introduce the Cox Proportional Hazards Regression Model, one of several available models for fitting regressions to time-to-event (survival) outcomes.

22.2 Foundations of Survival Analysis

Survival analysis is concerned with *prospective* studies, where we start with a cohort of subjects and follow them forwards in time to determine some clinical outcome. Follow-up continues until either some event of interest occurs, the study ends, or further observation becomes impossible.

The outcomes in a survival analysis consist of the subject's **fate** and **length of follow-up** at the end of the study.

- For some patients, the outcome of interest may not occur during follow-up.
- For such patients, whose follow-up time is *censored*, we know only that this event did not occur while the patient was being followed. We do not know whether or not it will occur at some later time.

The primary problems with survival data are *non-normality* and *censoring*...

1. Survival data are quantitative, but not symmetrically distributed. They will often appear positively skewed, with a few people surviving a very long time compared with the majority; so assuming a normal distribution will not be reasonable.
2. At the completion of the study, some patients may not have reached the endpoint of interest (death, relapse, etc.). Consequently, the exact survival times are not known.
 - All that is known is that the survival times are greater than the amount of time the individual has been in the study.
 - The survival times of these individuals are said to be **censored** (precisely, they are right-censored).

22.2.1 The Survival Function, $S(t)$

The **survival function**, $S(t)$ (sometimes called the survivor function) is the probability that the survival time, T , is greater than or equal to a particular time, t .

- $S(t)$ = proportion of people surviving to time t or beyond

If there's no censoring, the survival function is easy to estimate.

$$\hat{S}(t) = \frac{\# \text{ of subjects with survival times } \geq t}{n}$$

but this won't work if there is censoring.

22.2.2 Kaplan-Meier Estimator of the Survival Function

The survival function $S(t)$ is the probability of surviving until at least time t . It is essentially estimated by the number of patients alive at time t divided by the total number of study subjects remaining at that time.

The Kaplan-Meier estimator first orders the (unique) survival times from smallest to largest, then estimates the survival function at each unique survival time.

- The survival function at the second death time, $t_{(2)}$ is equal to the estimated probability of not dying at time $t_{(2)}$ conditional on the individual being still at risk at time $t_{(2)}$.

In the presence of censoring, the survival function is estimated as follows.

1. Order the survival times from smallest to largest, where $t_{\{(j)\}}$ is the j th largest unique survival time, so we have...

$$t_{(1)} \leq t_{(2)} \leq t_{(3)} \leq \dots t_{(n)}$$

2. The Kaplan-Meier estimate of the survival function is

$$\hat{S}(t) = \prod_{j:t_{(j)} \leq t} \left(1 - \frac{d_j}{r_j}\right)$$

where r_j is the number of people at risk just before $t_{(j)}$, including those censored at time $t_{(j)}$, and d_j is the number of people who experience the event at time $t_{(j)}$.

When we want to compare survival functions (or their Kaplan-Meier estimates, at least) we'll use a **log rank test** or one of several extensions of that test.

22.2.3 Creating a Survival Object in R

To do survival analysis in R, we're going to start with three main functions, all in the `survival` package:

- `Surv` creates a survival object
- `survfit` builds a Kaplan-Meier test, and the results may be plotted, as we've seen.
- `survdiff` builds a log rank test, that will let us compare two survival functions, as well as running several alternatives.

Plus, we'll build out some estimates of the hazard function.

The `Surv` function, part of the `survival` package in R, will create a **survival object** from two arguments:

1. `time` = follow-up time
2. `event` = a status indicator, where
 - `event = 1` or `TRUE` means the event was observed (for instance, the patient died)
 - `event = 0` or `FALSE` means the follow-up time was censored

22.3 A First Example: Recurrent Lobar Intracerebral Hemorrhage

O'Donnell et al. (2000) studied the effect of the apolipoprotein E gene on the risk of recurrent lobar intracerebral hemorrhage in 70 patients who survived such a hemorrhage¹. Patients in the study are classified by:

- `time` = follow-up time, in months
- `recur` = indicator of whether or not they had a recurrent event (`1` = subject had a recurrence, `0` subject did not have a recurrence), and
- `genotype` = the subject's apolipoprotein E genotype (`0` = Homozygous $\epsilon 3/\epsilon 3$ and `1` = At least one $\epsilon 2$ or $\epsilon 4$ allele)

```
hem %>% head(., 4)
```

```
# A tibble: 4 x 4
  id genotype  time  recur
  <dbl>    <dbl> <dbl> <dbl>
1     1        0  0.230     1
2     2        0  1.05      0
3     3        1  1.38      0
4     4        1  1.41      1
```

¹These data come from Dupont WD (2002) Statistical Modeling for Biomedical Researchers. New York: Cambridge U. Press, section 6.4.

```
table(hem$recur)
```

```
0 1
52 18
```

We have 70 patients at the start, and observe 18 events (rest are censored.)

```
mosaic::favstats(time ~ recur, data = hem)
```

recur	min	Q1	median	Q3	max	mean	sd	n	
1	0	1.0513350	14.414783	23.01437	38.64477	53.88091	25.51129	14.68242	52
2	1	0.2299795	3.367556	10.92403	24.84600	42.87474	14.98517	13.88893	18
		missing							
1	0								
2	0								

The median survival time looks like 23 weeks in the patients who do not exhibit a recurrence, but only 11 weeks in those who do.

22.3.1 Building a Survival Object

```
hemsurv <- Surv(time = hem$time, event = hem$recur)
```

```
head(hemsurv, 4)
```

```
[1] 0.2299795 1.0513350+ 1.3798770+ 1.4127311
```

This object both displays the survival time for each subject, and indicates whether or not the subject's follow-up was *censored* before a recurrent event occurred. Survival times with a + sign indicate censoring.

- Subject 1 lived for 0.23 months hemorrhage-free and then had a recurrence.
- Subject 2 lived for 1.05 months hemorrhage-free, at which point they were censored (perhaps because the study ended, or perhaps because the subject was no longer available for follow-up)

Remember that 18 of the subjects experienced a recurrent hemorrhage, and the other 52 are therefore censored.

22.3.2 Kaplan-Meier Estimate of the Survival Function

To build a Kaplan-Meier estimate of the survival function (to account properly for censoring), we take the survival object we have created, and use the `survfit` function from the `survival` package.

```
hemfit1 <- survfit(hemsurv ~ 1)
```

We can look at the `hemfit1` object directly, although the K-M estimate is usually plotted.

```
print(hemfit1, print.rmean=TRUE)
```

```
Call: survfit(formula = hemsurv ~ 1)
```

	n	events	rmean*	se(rmean)	median	0.95LCL	0.95UCL
[1,]	70	18	40.4	2.63	NA	42.9	NA
					*	restricted mean with upper limit =	53.9

We see that 18 events occurred out of a total of 70 subjects. The median survival time is listed as `NA` (missing) which implies it cannot be estimated by this simple model.

- This is because only 18 of our 70 subjects have a known recurrence-free survival time (the rest are censored), so we don't actually know what the median survival time will be across our 70 subjects. Apparently, R can produce a lower bound on a 95% confidence interval for the median survival time, but not the actual point estimate.

We also observe a **restricted mean survival time** estimate. The restricted mean uses as its upper limit the largest observed or censored survival time, which here is a censored value: 53.9 months. So it is the mean survival time, assuming all censored subjects lived hemorrhage-free for 53.9 months.

```
summary(hemfit1)
```

```
Call: survfit(formula = hemsurv ~ 1)
```

time	n.risk	n.event	survival	std.err	lower	95% CI	upper	95% CI
0.23	70	1	0.986	0.0142		0.958		1.000
1.41	67	1	0.971	0.0202		0.932		1.000
1.58	65	1	0.956	0.0248		0.909		1.000
3.06	63	1	0.941	0.0287		0.886		0.999
3.32	62	1	0.926	0.0320		0.865		0.991
3.52	61	1	0.911	0.0349		0.845		0.982
3.55	60	1	0.895	0.0375		0.825		0.972
4.76	57	1	0.880	0.0400		0.805		0.962
9.53	54	1	0.863	0.0424		0.784		0.951
12.32	50	1	0.846	0.0449		0.762		0.939
15.57	46	1	0.828	0.0476		0.740		0.926
19.15	38	1	0.806	0.0511		0.712		0.912
24.77	32	1	0.781	0.0553		0.679		0.897
24.87	31	1	0.756	0.0590		0.648		0.881
28.09	26	1	0.726	0.0635		0.612		0.862
33.61	22	1	0.693	0.0687		0.571		0.842
37.52	17	1	0.653	0.0758		0.520		0.819

22.3. A FIRST EXAMPLE: RECURRENT LOBAR INTRACEREBRAL HEMORRHAGE 559

```
42.87      8      1      0.571  0.1011      0.404      0.808
```

This written summary provides us with lots of detail on the Kaplan-Meier estimate. In particular, the first two lines of this summary can be read to indicate the following.

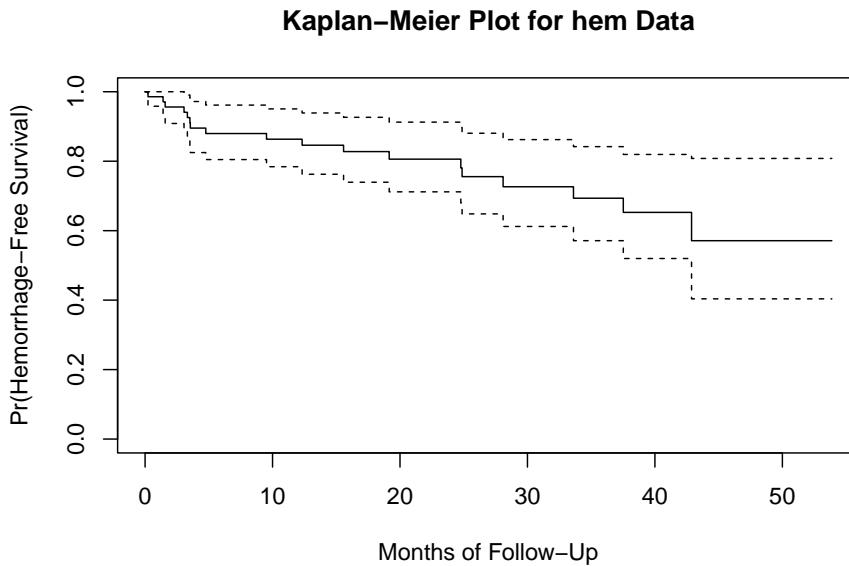
- Up to time 0.23 months, no patients had a recurrence. Then, an event occurred, and the estimated survival (i.e. non-recurrence) probability is reduced from 1 to 0.986.
- By time 1.41 months, when the next event occurred, only 67 patients remained at risk. This is because one of them had a recurrent hemorrhage already (at 0.23 months) and two others had been right-censored. The estimated hemorrhage-free survival probability estimate starting at time 1.41 months is now 0.971.

A Kaplan-Meier plot graphically represents this summary.

22.3.3 The Kaplan-Meier Plot, using Base R

Now, let's plot the Kaplan-Meier estimate, so we can see what is going on.

```
plot(hemfit1, ylab="Pr(Hemorrhage-Free Survival)",  
     xlab="Months of Follow-Up",  
     main="Kaplan-Meier Plot for hem Data")
```



The solid line indicates estimated hemorrhage-free survival probability. The

dotted lines identify pointwise confidence intervals (default 95%).

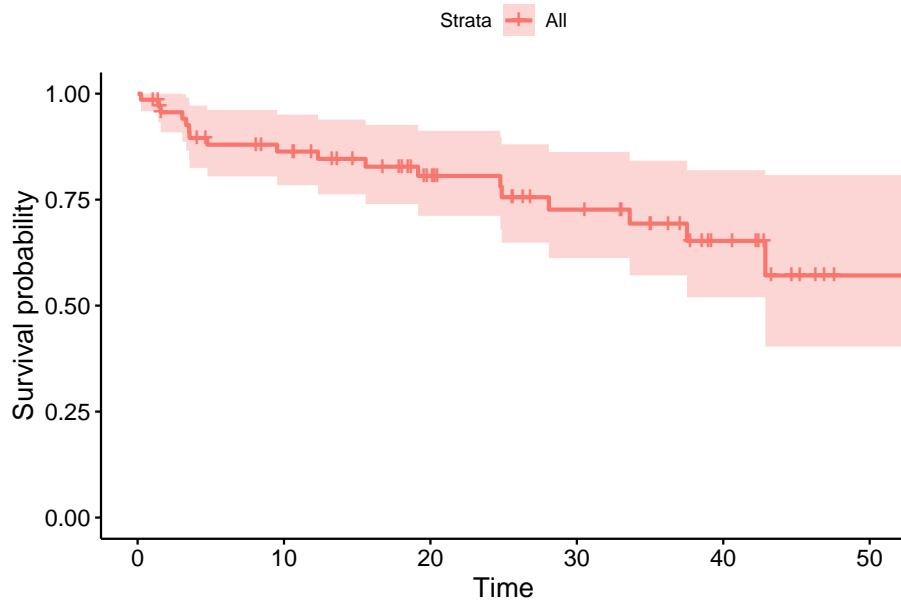
- For example, we see that the estimated probability of hemorrhage-free survival to 20 months is estimated to be about 0.8
- The estimated probability of hemorrhage-free survival to 50 months is estimated to be about 0.6

The steps down indicate events (recurrences.) The estimated probability of survival to 0 months starts at 1, and drops down at each time point where an event (or more than one event) is observed.

22.3.4 Using survminer to draw survival curves

Another approach to plotting the Kaplan-Meier estimate comes from `ggsurvplot`, from the `survminer` package.

```
ggsurvplot(hemfit1, data = hem)
```



Again, the solid line indicates estimated hemorrhage-free survival probability. The crosses indicate censoring. The steps down indicate events (recurrences,) and the shading indicates (default 95%) pointwise confidence intervals. By **pointwise** confidence intervals, I mean that these bounds apply only to individual points in the time scale.

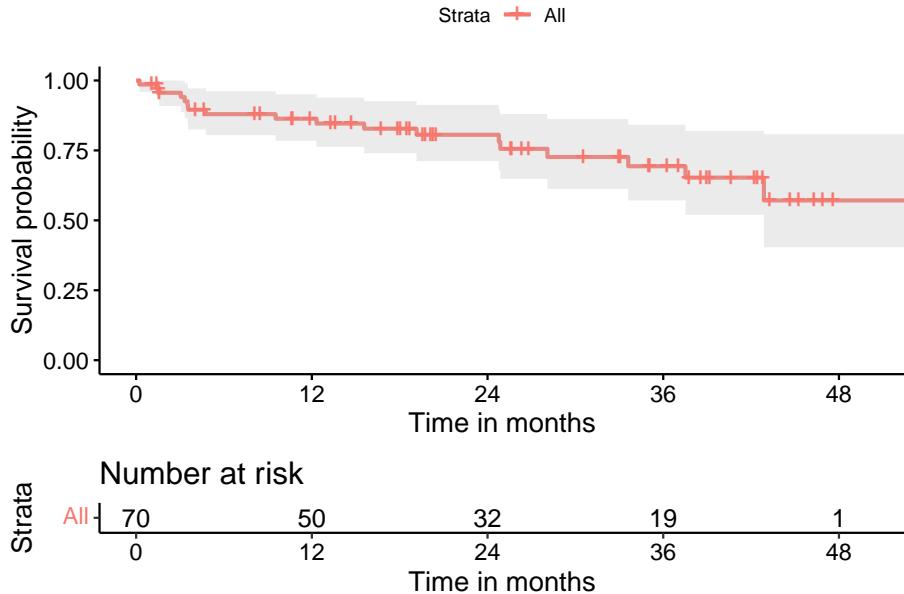
For more on an alternative approach, using *simultaneous confidence bands*, visit the OpenIntro Statistics Survival Analysis in R materials, written by David Diez,

which are also posted on our web site.

22.3.5 A “Fancy” K-M Plot with a number at risk table

We can do a lot more with these plots. Following the suggestions at <https://github.com/kassambara/survminer/> we can create the following...

```
ggsurvplot(hemfit1, data = hem,
            conf.int = TRUE, # Add confidence interval
            risk.table = TRUE, # Add risk table
            xlab = "Time in months", # adjust X axis label
            break.time.by = 12 # add tick every 12 months
)
```



This sort of plot is really designed to work best when we compare multiple groups in terms of their survival. So let's do that.

22.4 Comparing Survival Across the Two Genotypes

Now, suppose we want to compare the hemorrhage-free survival functions for subjects classified by their apolipoprotein E genotype. Working with the same

survival object `hemsurv` we now run the `survfit` function to compare across the two genotype groups.

```
hemfit2 <- survfit(hemsurv ~ hem$genotype)
print(hemfit2, print.rmean=TRUE)
```

```
Call: survfit(formula = hemsurv ~ hem$genotype)

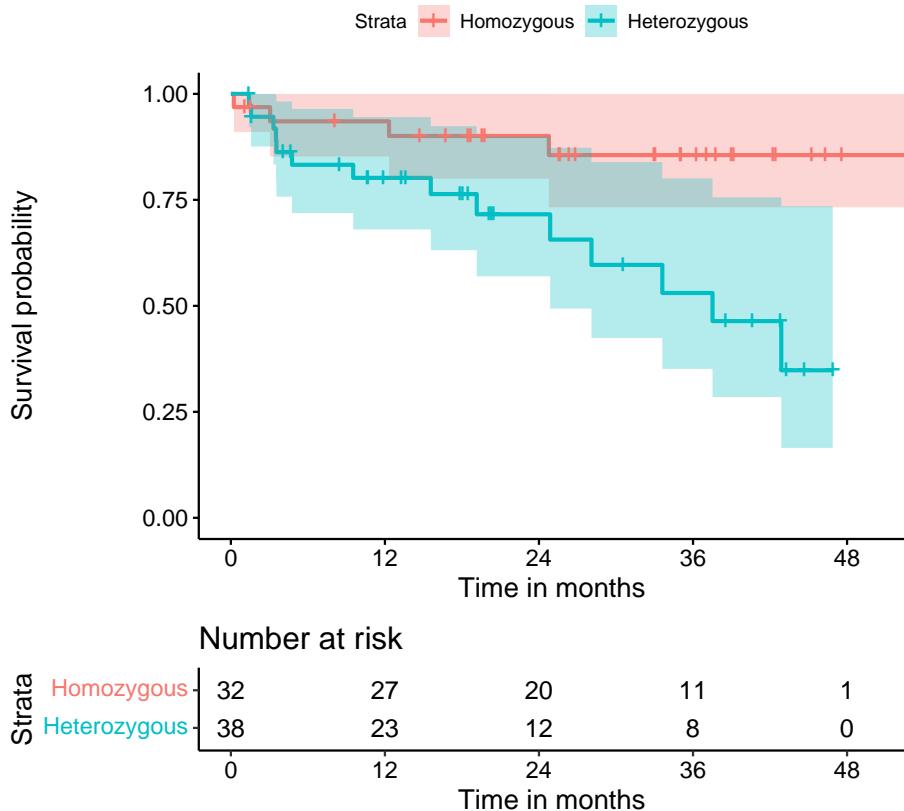
      n  events rmean* se(rmean) median 0.95LCL 0.95UCL
hem$genotype=0 32      4    47.8      2.87     NA     NA     NA
hem$genotype=1 38     14    33.9      3.77    37.5    24.9     NA
* restricted mean with upper limit = 53.9
```

- In `genotype = 0` (the subjects who are Homozygous $\epsilon 3/\epsilon 3$), we had 32 subjects, and observed 4 recurrent hemorrhages. Our estimated restricted mean survival time in those subjects is 44.8 months and we cannot estimate a median survival time because only a small fraction of our subjects were not censored.
- In `genotype = 1` (subjects who have at least one $\epsilon 2$ or $\epsilon 4$ allele,) we had 38 subjects and observed 14 recurrences. The estimated restricted mean survival time is 32.7 months in these subjects, and we can (it seems) estimate a median survival time in this group of 37.5 months. Note that we don't actually need to observe the event in half of the subjects to estimate a median survival time.

22.4.1 Kaplan-Meier Survival Function Estimates, by Genotype

I find I have to crank the `fig.height` in Markdown up to at least 6 to get the risk table to show up nicely in this setting.

```
ggsurvplot(hemfit2, data = hem,
            conf.int = TRUE,
            xlab = "Time in months",
            break.time.by = 12,
            legend.labs = c("Homozygous", "Heterozygous"),
            risk.table = TRUE,
            risk.table.height = 0.25
          )
```



It appears that patients who were homozygous for the $\epsilon 3$ allele of this gene (i.e. genotype = 0 in the hemorrhage data) had a much better prognosis than others (genotype = 1.)

22.5 Testing the difference between two survival curves

To obtain a significance test comparing these two survival curves, we turn to a *log rank* test, which tests the null hypothesis $H_0 : S_1(t) = S_2(t)$ for all t where the two exposures have survival functions $S_1(t)$ and $S_2(t)$. We use the `survdiff` function to explore this test, which uses a χ^2 statistic to do the testing.

```
survdiff(hem$surv ~ hem$genotype)
```

```
Call:  
survdiff(formula = hem$surv ~ hem$genotype)
```

```

N Observed Expected (O-E)^2/E (O-E)^2/V
hem$genotype=0 32      4     9.28     3.00     6.28
hem$genotype=1 38      14    8.72     3.19     6.28

Chisq= 6.3 on 1 degrees of freedom, p= 0.01

```

Based on the log rank test, we conclude that there is a statistically significant difference ($p = .0122$) between the hemorrhage-free survival curves for the two genotypes, as shown in the Kaplan-Meier plot.

- The log rank test generalizes to permit survival comparisons across more than two groups, with the test statistic having an asymptotic chi-squared distribution with one degree of freedom less than the number of patient groups being compared.

22.5.1 Alternative log rank tests

An alternative approach to testing is the *Peto and Peto modification of the Gehan-Wilcoxon test*, which results from adding `rho=1` to the `survdiff` function (`rho=0`, the default, yields the log rank test.)

```
survdiff(hemsurv ~ hem$genotype, rho=1)
```

Call:

```
survdiff(formula = hemsurv ~ hem$genotype, rho = 1)
```

```

N Observed Expected (O-E)^2/E (O-E)^2/V
hem$genotype=0 32      3.63     7.87     2.29     5.46
hem$genotype=1 38      11.79    7.54     2.39     5.46

```

Chisq= 5.5 on 1 degrees of freedom, p= 0.02

As compared to the log rank test, this Peto-Peto modification (and others using $\text{rho} > 0$) give greater weight to the left hand (earlier) side of the survival curves.

- To obtain chi-square tests that give greater weight to the right hand (later) side of the survival curves than the log rank test, use a `rho` value which is less than 0.

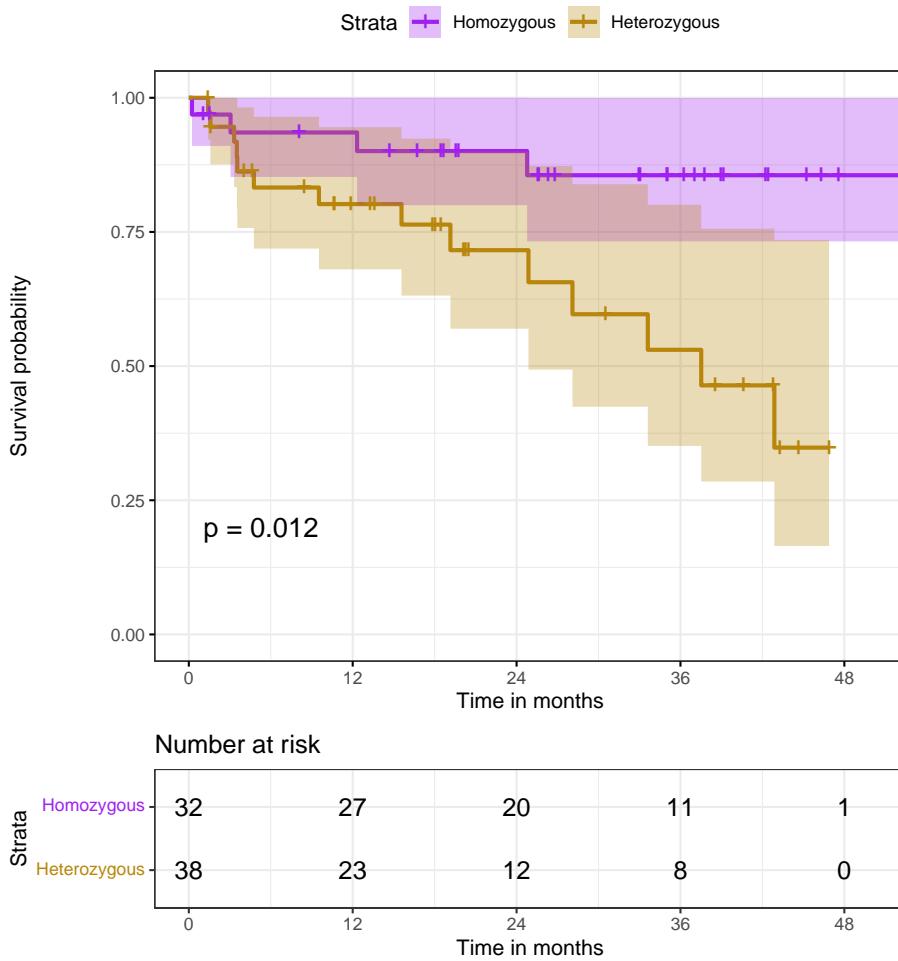
22.6 A “Fancy” K-M Plot with a number at risk table

We can add the log rank test result to our “fancy” K-M plot. Visit <https://github.com/kassambara/survminer/> for more options.

```

ggsurvplot(hemfit2, data = hem, size = 1,
  palette = c("purple", "darkgoldenrod"), # custom colors
  conf.int = TRUE, # Add confidence interval
  pval = TRUE, # Add p-value
  risk.table = TRUE, # Add risk table
  risk.table.height = 0.25, # change if you have >2 groups
  risk.table.y.text.col = T, # show colors in table listing
  xlab = "Time in months", # adjust X axis label
  break.time.by = 12, # break X axis in time intervals
  legend.labs = c("Homozygous", "Heterozygous"), # labels
  ggtheme = theme_bw() # Change ggplot2 theme
)

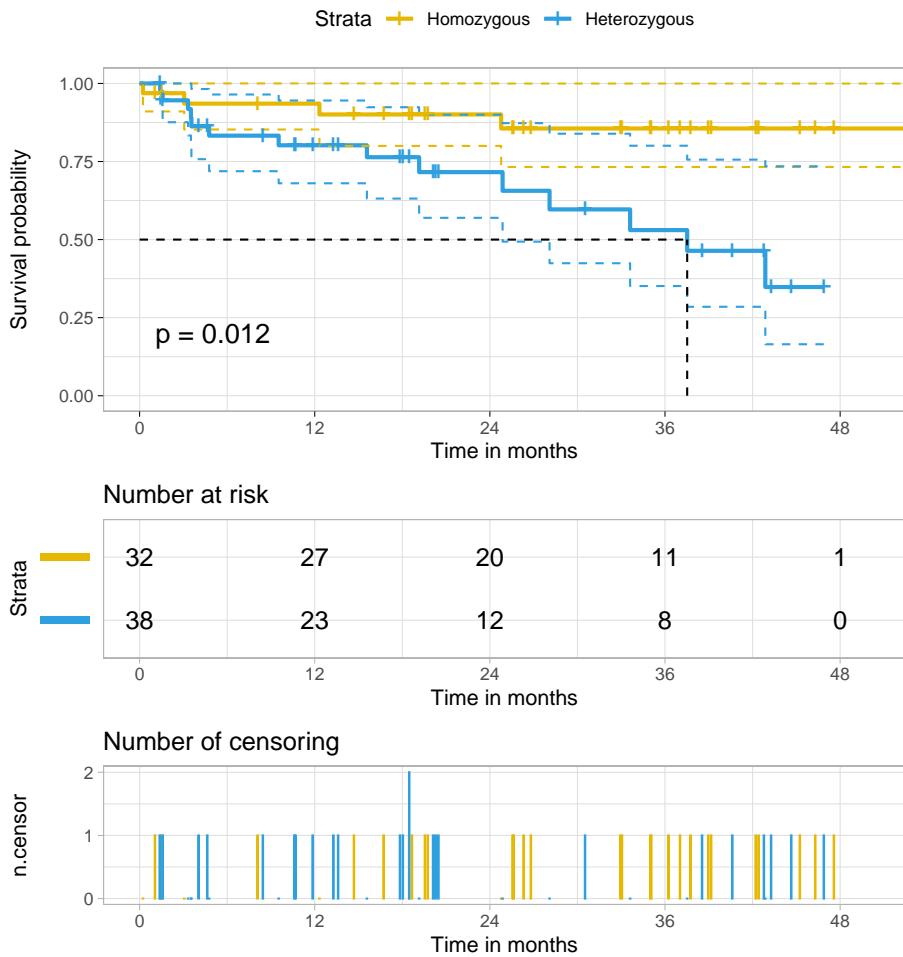
```



22.6.1 Customizing the Kaplan-Meier Plot Presentation Further

We can even add a plot of the number of censored subjects at each time point, as well as a median survival pointer (which, of course, we've seen that we can't estimate in one of the groups), and customize the style of the confidence intervals. Again, see <https://github.com/kassambara/survminer/> for even more customized results.

```
ggsurvplot(hemfit2,
            data = hem,
            palette = c("#E7B800", "#2E9FDF"),
            risk.table = TRUE,
            pval = TRUE,
            conf.int = TRUE,
            xlab = "Time in months",
            break.time.by = 12,
            ggtheme = theme_light(),
            risk.table.y.text.col = T,
            risk.table.height = 0.25,
            risk.table.y.text = FALSE,
            ncensor.plot = TRUE,
            ncensor.plot.height = 0.25,
            conf.int.style = "step",
            surv.median.line = "hv",
            legend.labs = c("Homozygous", "Heterozygous")
)
```



22.7 The Hazard Function

To build regression models for time-to-event data, we will need to introduce the **hazard function**. Consider a subject in the hemorrhage study who has a hemorrhage-free survival time of 9 months.

- For this subject to have had a recurrent hemorrhage at 9 months, they had to be hemorrhage-free for the first 8 months.
- The subject's hazard at 9 months is the failure rate "per month" conditional on the subject being hemorrhage-free for the first 8 months.

If $S(t)$ is the survival function, and time t is taken to be continuous, then $S(t) = e^{H(t)}$ defines the hazard function $H(t)$.

- Note that $H(t) = -\ln(S(t))$.
- The function $H(t)$ is an important analytic tool.
 - It's used to describe the concept of the risk of “failure” in an interval after time t , conditioned on the subject having survived to time t .
 - It's often called the *cumulative hazard function*, to emphasize the fact that its value is the “sum” of the hazard up to time t .

There are several different methods to estimate $H(t)$, but we'll focus on two...

1. The inverse Kaplan-Meier estimator
2. The Nelson-Aalen estimator

22.7.1 The Inverse Kaplan-Meier Estimator of $H(t)$

Our first estimator of the hazard function, $H(t)$ will be the inverse Kaplan-Meier estimate, which I'll place in an R object called `H_est1`.

- To start, we will take the negative of the log of the Kaplan-Meier survival estimate. That takes care of the first $t-1$ levels of the eventual estimate.
- To complete the process, we will repeat the final one of those time-specific estimates at the end.

```
# hemsurv <- Surv(hem$time, hem$recur)
# hemfit1 <- survfit(hemsurv ~ 1)
H_est1 <- -log(hemfit1$surv)
H_est1 <- c(H_est1, tail(H_est1, 1))
```

Here are the first five, and last five values of the hazard function estimate.

```
head(H_est1, 5) # first 5 values
[1] 0.01438874 0.01438874 0.01438874 0.02942661 0.02942661
tail(H_est1, 5) # last 5
[1] 0.5602049 0.5602049 0.5602049 0.5602049 0.5602049
```

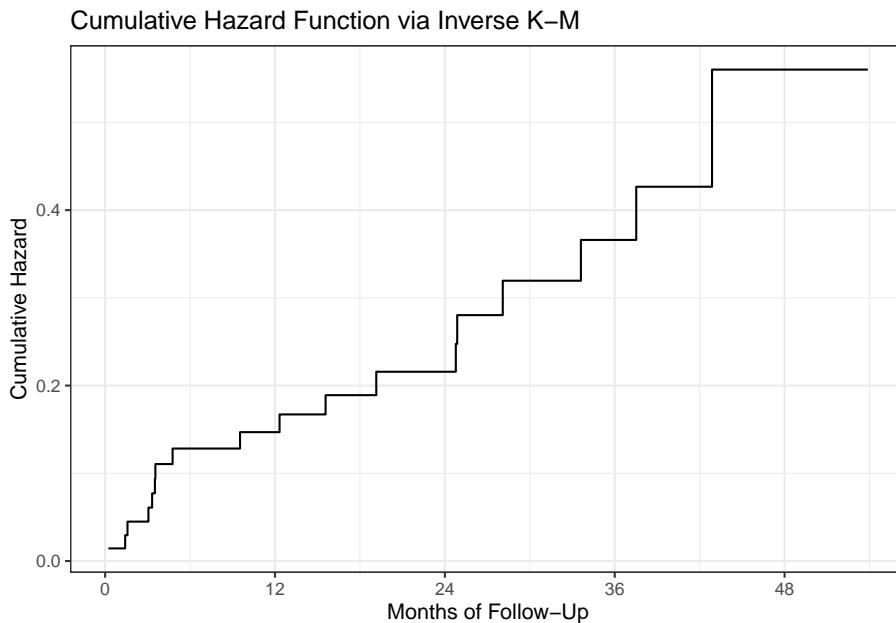
We can create a little data frame containing the times and hazard estimates, like this:

```
haz_hem <- data_frame(
  time = c(hemfit1$time, tail(hemfit1$time, 1)),
  inverse_KM = H_est1
)
```

22.7.2 Cumulative Hazard Function from Inverse K-M

Since we've built the data set of times and hazard values, we can use the `geom_step` function in `ggplot2`.

```
ggplot(haz_hem, aes(x = time, y = inverse_KM)) +
  geom_step() +
  scale_x_continuous(breaks = c(0, 12, 24, 36, 48)) +
  labs(x = "Months of Follow-Up",
       y = "Cumulative Hazard",
       title = "Cumulative Hazard Function via Inverse K-M")
```



22.7.3 The Nelson-Aalen Estimator of $H(t)$

An alternative estimate of the cumulative hazard is called the Nelson-Aalen estimate, captured here in `H_est2`.

```
h_st <- hemfit1$n.event / hemfit1$n.risk
H_est2 <- cumsum(h_st)
H_est2 <- c(H_est2, tail(H_est2, 1))
haz_hem$Nelson_Aalen <- H_est2

head(haz_hem)

# A tibble: 6 x 3
  time    inverse_KM Nelson_Aalen
  <dbl>      <dbl>        <dbl>
1 0.230     0.0144     0.0143
2 1.05      0.0144     0.0143
```

```

3 1.38      0.0144      0.0143
4 1.41      0.0294      0.0292
5 1.51      0.0294      0.0292
6 1.58      0.0449      0.0446

```

22.7.4 Convert Wide Data to Long

In order to easily plot the two hazard function estimates in the same graph, we'll want to convert these data from wide format to long format, with the `gather` function.

```

haz_hem_comp <- gather(haz_hem, key = "method",
                        value = "hazardest",
                        inverse_KM:Nelson_Aalen)

head(haz_hem_comp)

```

```

# A tibble: 6 x 3
  time method    hazardest
  <dbl> <chr>      <dbl>
1 0.230 inverse_KM 0.0144
2 1.05  inverse_KM 0.0144
3 1.38  inverse_KM 0.0144
4 1.41  inverse_KM 0.0294
5 1.51  inverse_KM 0.0294
6 1.58  inverse_KM 0.0449

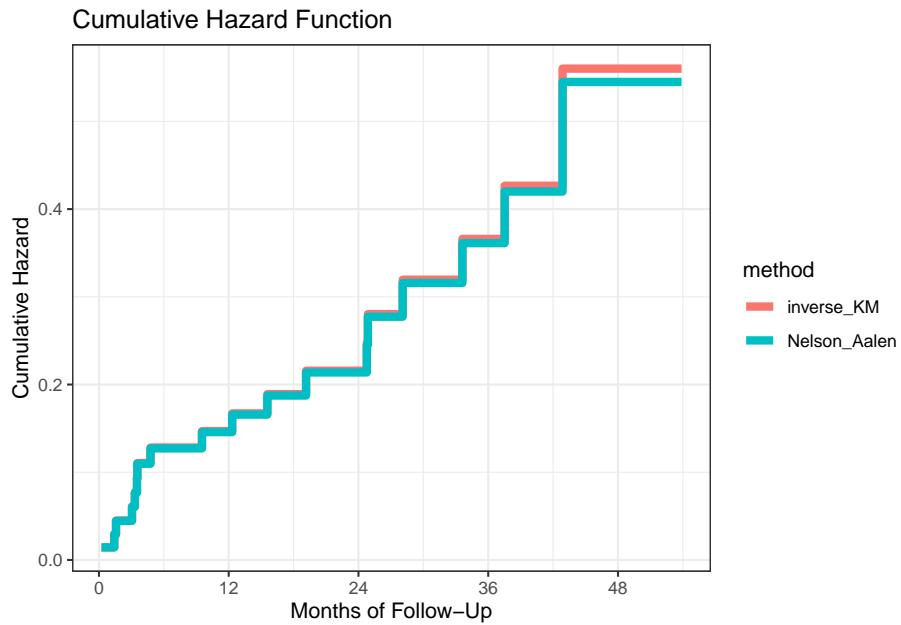
```

22.7.5 Plot Comparison of Hazard Estimates

```

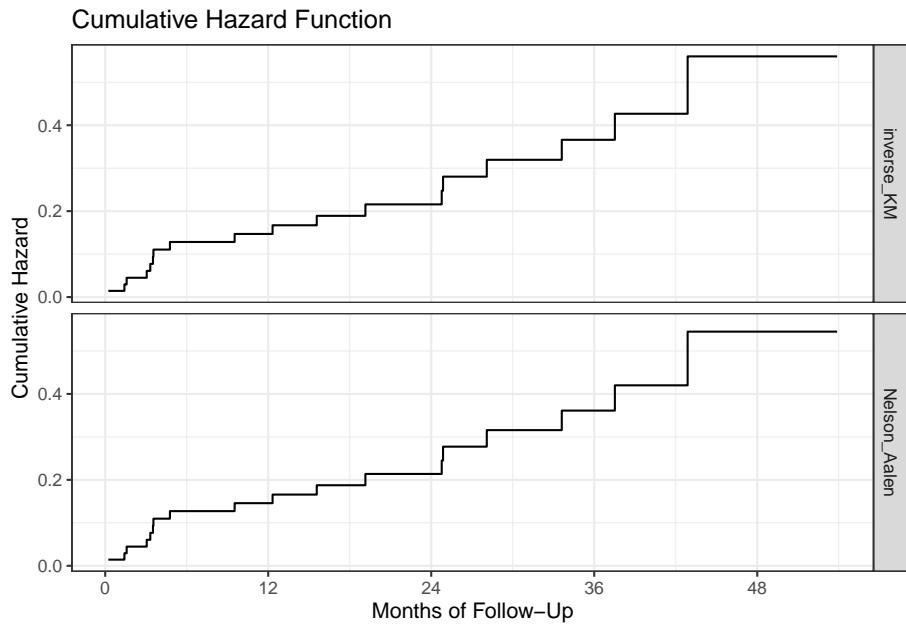
ggplot(haz_hem_comp, aes(x = time, y = hazardest,
                          col = method)) +
  geom_step(size = 2) +
  scale_x_continuous(breaks = c(0, 12, 24, 36, 48)) +
  labs(x = "Months of Follow-Up",
       y = "Cumulative Hazard",
       title = "Cumulative Hazard Function") +
  theme_bw()

```



We can see that the two cumulative hazard function estimates are nearly identical in this case. We could instead compare the two functions in faceted plots, if that would be helpful.

```
ggplot(haz_hem_comp, aes(x = time, y = hazardest)) +
  geom_step() +
  scale_x_continuous(breaks = c(0, 12, 24, 36, 48)) +
  labs(x = "Months of Follow-Up",
       y = "Cumulative Hazard",
       title = "Cumulative Hazard Function") +
  facet_grid(method ~ .) + theme_bw()
```



Next, we will consider the issue of modeling a survival outcome using Cox proportional hazards regression.

Chapter 23

Cox Regression Models for Survival Data: Example 1

The Cox proportional hazards (Cox regression) model fits survival data with a constant (i.e. not varying over time) covariate x to a hazard function of the form:

$$h(t|x) = h_0(t)\exp[\beta_1 x]$$

where we will estimate the unknown value of β_1 and where $h_0(t)$ is the baseline hazard, which is a non-parametric and unspecified value which depends on t but not on x .

- For particular x values, we will be able to estimate the survival function if we have an estimate of the baseline survival function, $\hat{S}_0(t)$.

The estimated survival function for an individual with covariate value x_k turns out to be

$$\hat{S}(t|x_k) = [\hat{S}_0(t)]^{\exp(\beta_1 x_k)}$$

From Wikipedia (yes, really) ...

Survival models can be viewed as consisting of two parts: the underlying hazard function, describing how the risk of event per time unit changes over time at baseline levels of covariates; and the effect parameters, describing how the hazard varies in response to explanatory covariates.

The key assumption in a Cox model is that the hazards are **proportional** - other types of survival models need not have this restriction. Quoting the always reliable (well, it's better than you think) Wikipedia ...

In a proportional hazards model, the unique effect of a unit increase in a covariate is multiplicative with respect to the hazard rate. For example, taking a drug may halve one's hazard rate for a stroke occurring, or, changing the material from which a manufactured component is constructed may double its hazard rate for failure.

There are two main approaches to fitting Cox models in R.

- the `coxph` function in the `survival` package, and
- the `cph` function in the `rms` package.

23.1 Sources used in building this material

- David Diez's excellent supplement for the OpenIntro Statistics project, on Survival Analysis in R. I've posted it on our web site, as well.
- Some tools in R to do some fancier work can be viewed at <https://cran.r-project.org/web/views/Survival.html>
- You might also look at these two blog posts, originally from the Just Another Data blog.
 - <https://www.r-bloggers.com/survival-analysis-1/>
 - <https://www.r-bloggers.com/survival-analysis-2/>
- <https://rpubs.com/daspringate/survival> has some great slides, and I've stolen from them quite a bit here.

23.2 Fitting a Cox Model in R with `coxph`

As a first example, I'll fit a model to predict time to recurrence in the `hem` data, on the basis of a single predictor: `genotype`.

```
cfit <- with(hem, coxph(Surv(time, recur) ~ genotype))
cfit
```

```
Call:
coxph(formula = Surv(time, recur) ~ genotype)
```

	coef	exp(coef)	se(coef)	z	p
genotype	1.3317	3.7874	0.5699	2.337	0.0195

```
Likelihood ratio test=6.61 on 1 df, p=0.01015
n= 70, number of events= 18
```

This summary provides an overall comparison of the two genotypes, using a proportional hazards model.

- The default approach in R is to use the “efron” method of breaking ties; other options include “breslow” and “exact.”

23.2.1 Summarizing the Fit

```
summary(cfit)

Call:
coxph(formula = Surv(time, recur) ~ genotype)

n= 70, number of events= 18

      coef exp(coef) se(coef)   z Pr(>|z|)
genotype 1.3317    3.7874   0.5699 2.337   0.0195 *
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

      exp(coef) exp(-coef) lower .95 upper .95
genotype     3.787      0.264     1.239     11.57

Concordance= 0.622  (se = 0.061 )
Likelihood ratio test= 6.61  on 1 df,   p=0.01
Wald test            = 5.46  on 1 df,   p=0.02
Score (logrank) test = 6.28  on 1 df,   p=0.01
```

This provides estimates of the β value for `genotype`, including standard errors and p values for a Wald test. Also included is an estimate of the hazard ratio and its confidence interval.

- Here we have a hazard ratio estimate of `exp(coef)` = 3.787, with 95% CI (1.24, 11.57).
- The hazard ratio is the multiplicative effect of the covariate (here, having at least one of the $\epsilon 2$ or $\epsilon 4$ allele) on the hazard function for recurrent hemorrhage
 - A hazard ratio of 1 indicates no effect
 - A hazard ratio > 1 indicates an increase in the hazard as the covariate rises
 - A hazard ratio < 1 indicates a decrease in the hazard as the covariate rises

We can also `tidy` the hazard ratio estimate with the `broom` package.

```
tidy(cfit, exponentiate = TRUE)
```

```
# A tibble: 1 x 5
  term      estimate std.error statistic p.value
  <chr>     <dbl>     <dbl>     <dbl>    <dbl>
1 genotype   3.79     0.570     2.34    0.0195
```

In addition, we have several other summaries:

- The *concordance* measure is only appropriate when we have at least one continuous predictor in our Cox model.
- The Cox & Snell pseudo- R^2 reflects the improvement of the model we've fit over the model with an intercept alone, but isn't a proportion of anything (hence the listing of the maximum possible value).
- The Likelihood ratio, Wald and Score (logrank) tests provide insight into the overall significance of the model.

We can obtain a more detailed description of the likelihood-ratio test of the model with `anova`.

```
anova(cfit)
```

```
Analysis of Deviance Table
Cox model: response is Surv(time, recur)
Terms added sequentially (first to last)

      loglik  Chisq Df Pr(>|Chi|)
NULL      -66.675
genotype -63.371 6.6078  1    0.01015 *
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

23.2.2 Glancing at the model?

```
glance(cfit)
```

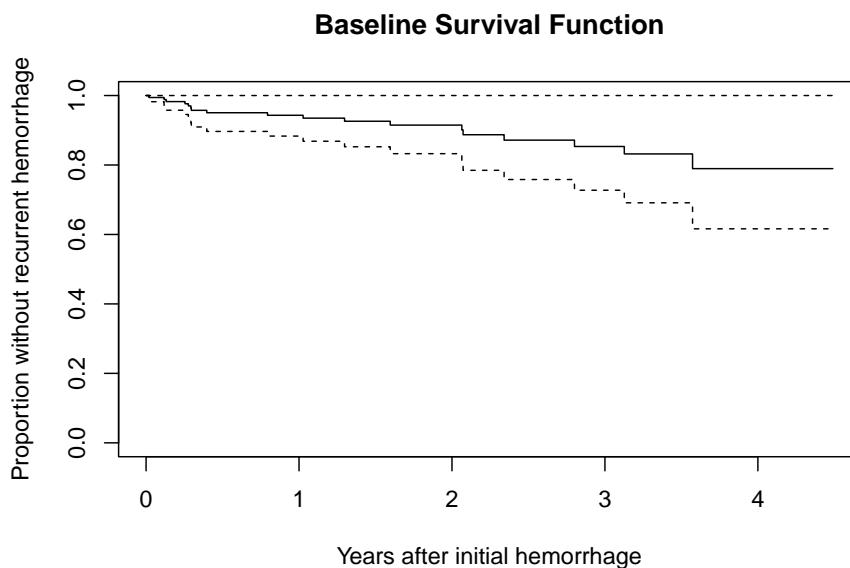
```
# A tibble: 1 x 18
  n nevent statistic.log p.value.log statistic.sc p.value.sc statistic.wald
  <int> <dbl>        <dbl>        <dbl>        <dbl>        <dbl>        <dbl>
1    70    18         6.61       0.0102       6.28       0.0122       5.46
# ... with 11 more variables: p.value.wald <dbl>, statistic.robust <dbl>,
#   p.value.robust <dbl>, r.squared <dbl>, r.squared.max <dbl>,
#   concordance <dbl>, std.error.concordance <dbl>, logLik <dbl>, AIC <dbl>,
#   BIC <dbl>, nobs <int>
```

Here, we obtain several additional summaries of the model, including most of the important information from a summary of `cfit`.

23.2.3 Plot the baseline survival function

Here, we'll plot the time in terms of months, but scaled to 12 month (one year) groups.

```
plot(survfit(cfit), xscale = 12,
      xlab = "Years after initial hemorrhage",
      ylab = "Proportion without recurrent hemorrhage",
      main = "Baseline Survival Function")
```



23.2.4 Plot the genotype effect

There are several ways to build these plots. One approach follows. Another uses a `cph` fit and the `survplot` function from the `rms` package.

```
newdat <- with(hem,
                 data.frame(
                   genotype = c(1, 0)
                 )
               )

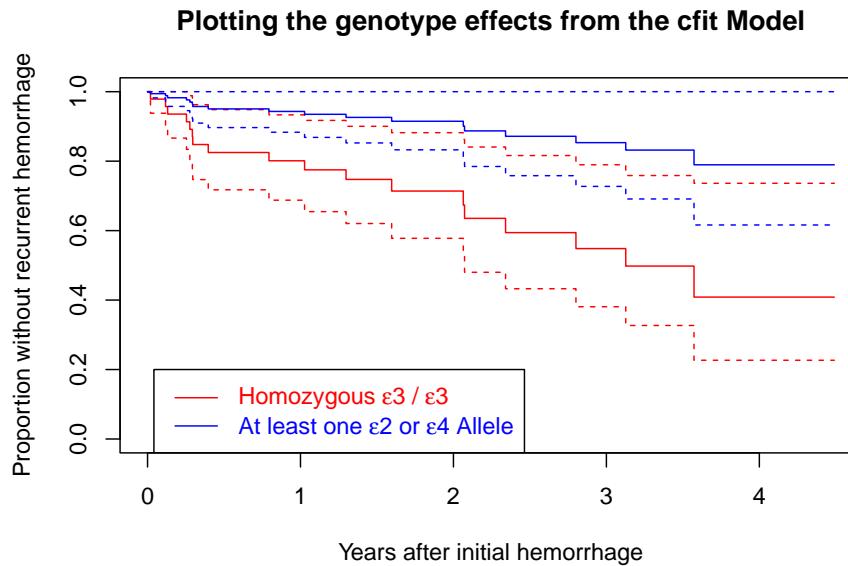
newdat

genotype
1      1
```

```

2      0
plot(survfit(cfit, newdata = newdat), xscale = 12,
      conf.int = TRUE,
      col = c("red", "blue"),
      xlab = "Years after initial hemorrhage",
      ylab = "Proportion without recurrent hemorrhage",
      main = "Plotting the genotype effects from the cfit Model")
legend(0.5, 0.2,
      legend=c(
        expression(paste("Homozygous ", epsilon, "3 / ",
                       epsilon, "3")),
        expression(paste("At least one ", epsilon,"2 or ",
                       epsilon,"4 Allele")))
      ),
      lty = 1,
      col = c("red", "blue"),
      text.col = c("red", "blue"))

```



23.2.5 Testing the Key Assumption: Proportional Hazards

The `cox.zph` function in the `survival` package will test the proportionality of all of the predictors included in your model by creating interactions with time.

- A small p value would indicate a violation of the proportionality assumption.

```
cox.zph(cfit, transform="km", global=TRUE)
```

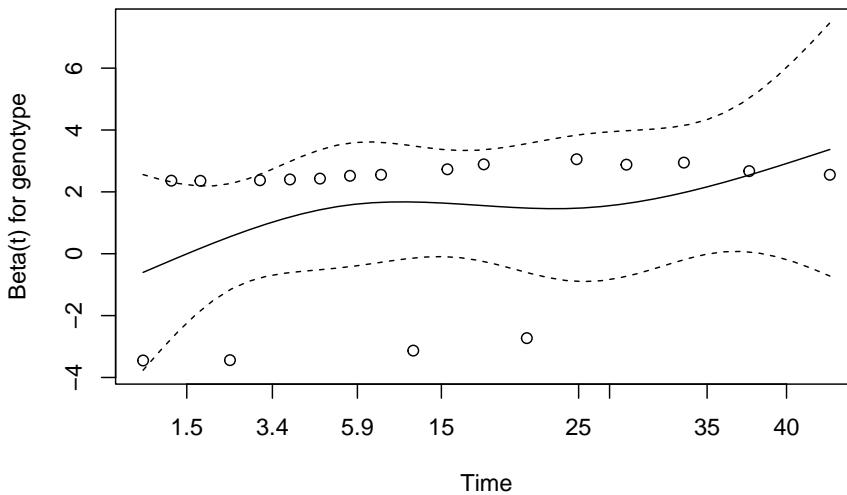
```
      chisq df      p
genotype  2.09  1 0.15
GLOBAL     2.09  1 0.15
```

Since the p value here is not significant, we may be all right. But it's sensible to focus further on plots derived from the model fit, rather than relying solely on this test.

23.2.6 Plotting the `cox.zph` results for the `cfit` model

`cox.zph` function can be used to generate a plot for each of the individual predictors in the model. Of course, in this case, we have just one predictor: `genotype`. If the proportional hazards assumption is appropriate, then we should see a slope of essentially zero in each such plot. A slope that is seriously different from zero suggests a violation of the proportional hazards assumption.

```
plot(cox.zph(cfit, transform="km", global=TRUE))
```



The plot suggests only a slight rise in the plotted values over time, suggesting no serious problem with the proportional hazards assumption. This combined testing and plotting approach is a reasonable starting place for assessing the

proportional hazards assumption, but it's likely insufficient for good practical work.

Should the proportional hazards assumption fit less well, we have two main options: (1) fit a non-linear term in the covariate in question, and (2) fit a different type of regression model that doesn't require the proportional hazards assumption.

23.3 Fitting a Cox Model using `cph` from the `rms` package

To set up a `cph` fit for our comparison of genotypes in the `hem` data, we'll follow these steps.

```
units(hem$time) <- "month"
d <- datadist(hem)
options(datadist = "d")

hemsurv <- Surv(time = hem$time, event = hem$recur)

model_hem <- cph(hemsurv ~ genotype, data = hem,
                   x = TRUE, y = TRUE, surv = TRUE)
```

Note that the `surv = TRUE` piece is required to get some of the follow-up analyses to work smoothly.

23.3.1 The Main `cph` results

```
model_hem

Cox Proportional Hazards Model

cph(formula = hemsurv ~ genotype, data = hem, x = TRUE, y = TRUE,
     surv = TRUE)

      Model Tests      Discrimination
                           Indexes
      Obs      70      LR chi2      6.61      R2      0.106
      Events    18      d.f.           1      Dxy      0.244
      Center 0.7229      Pr(> chi2) 0.0102      g      0.671
                           Score chi2     6.28      gr      1.955
                           Pr(> chi2) 0.0122

      Coef    S.E.      Wald Z Pr(>|Z|)
```

23.3. FITTING A COX MODEL USING CPH FROM THE RMS PACKAGE 581

```
genotype 1.3317 0.5699 2.34 0.0195
```

Included here are likelihood ratio and score tests for the model as a whole (as compared to the intercept-only model), as well as the usual discrimination indexes.

- These include both an R^2 analog due to Nagelkerke (which can go all the way up to 1), and
- Somers' D_{xy} , which can also produce an estimate of the C statistic (area under the curve) via the formula $C = 0.5 + D_{xy} / 2$, so here $C = 0.5 + (.244/2) = 0.622$
- For lots more on survival analysis C statistics, look at the **survAUC** package in R.

These results are followed by a table of Wald tests for each of the coefficients in the model.

23.3.2 Using **anova** with **cph**

As in other **rms** fits, we can use **anova** to obtain more detailed (in terms of combining nonlinear terms and, if available, interactions) tests.

```
anova(model_hem)
```

Wald Statistics			Response: hemsurv		
Factor	Chi-Square	d.f.	P		
genotype	5.46	1	0.0195		
TOTAL	5.46	1	0.0195		

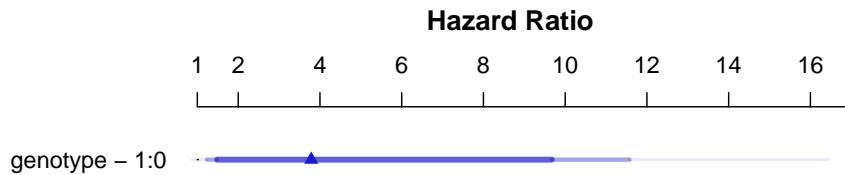
23.3.3 Effect Sizes after **cph** fit

We can use **summary** on a **cph** object to get and plot effect size estimates (here, these are hazard ratios.)

```
summary(model_hem)
```

Effects			Response : hemsurv						
Factor	Low	High	Diff.	Effect	S.E.	Lower	0.95	Upper	0.95
genotype	0	1	1	1.3317	0.5699	0.21468	2.4486		
Hazard Ratio	0	1	1	3.7873	NA	1.23950	11.5730		

```
plot(summary(model_hem))
```



23.3.4 Validating cph summaries

For details on these last few indices (D, U, Q, etc.), visit `?validate.cph` in R.

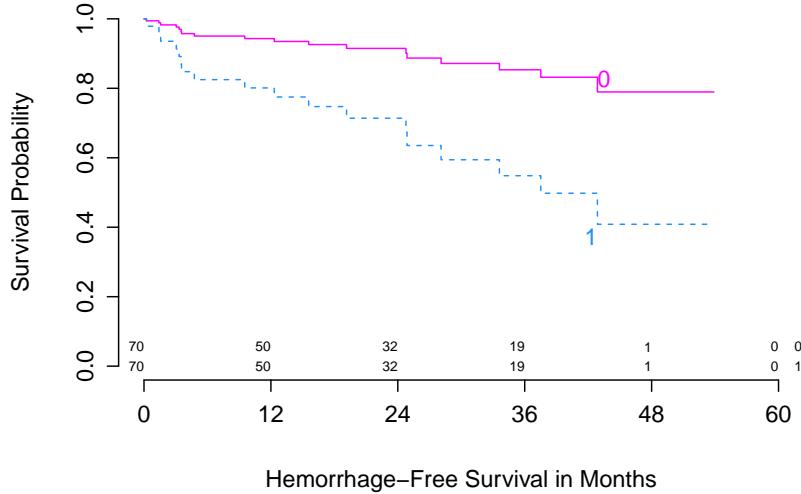
```
set.seed(43201); validate(model_hem)
```

	index.orig	training	test	optimism	index.corrected	n
Dxy	0.2441	0.2508	0.2197	0.0311	0.2130	40
R2	0.1058	0.1203	0.1058	0.0144	0.0914	40
Slope	1.0000	1.0000	-6.4747	7.4747	-6.4747	40
D	0.0421	0.0532	0.0421	0.0111	0.0310	40
U	-0.0150	-0.0168	0.0198	-0.0366	0.0216	40
Q	0.0571	0.0699	0.0222	0.0477	0.0094	40
g	0.6705	0.7908	0.6705	0.1203	0.5502	40

23.3.5 Plotting Survival Functions for each Genotype

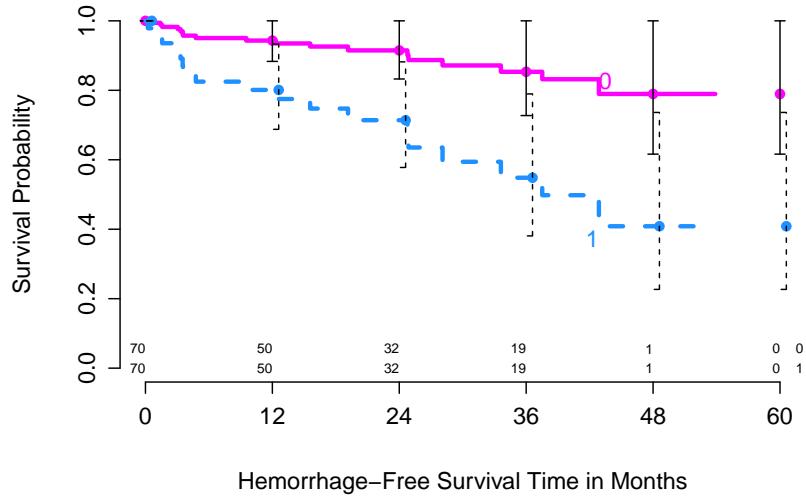
Here is the `survplot` approach I mentioned earlier.

```
survplot(model_hem, genotype,
          lty = c(1,2), n.risk=TRUE, time.inc=12,
          col=c("magenta", "dodgerblue"),
          xlab="Hemorrhage-Free Survival in Months")
```



We can add, for instance, confidence interval bars, with:

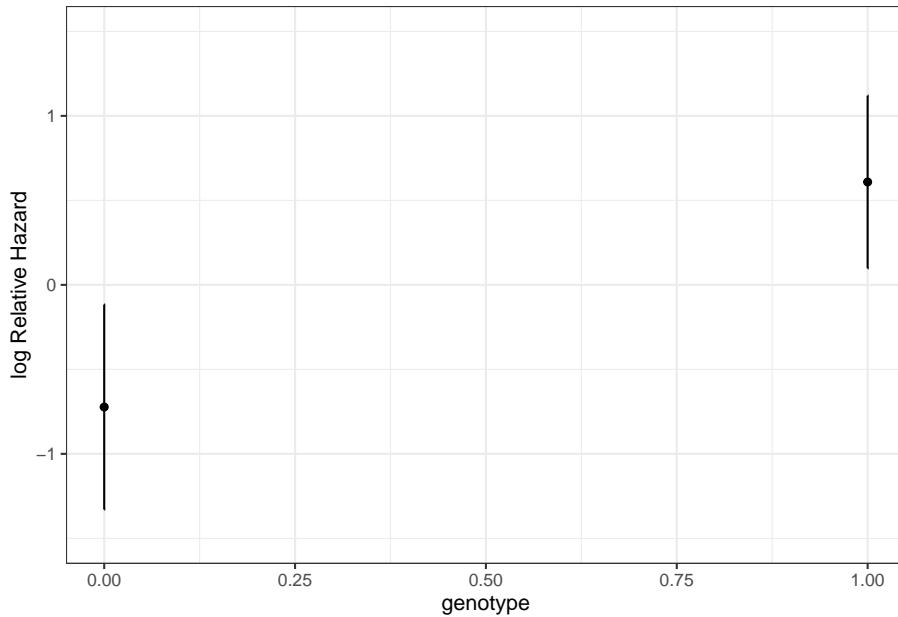
```
survplot(model_hem, genotype,
         lwd=3, lty = c(1,2), conf.int=.95,
         n.risk=TRUE, time.inc = 12, conf='bars',
         col=c("magenta", "dodgerblue"),
         xlab="Hemorrhage-Free Survival Time in Months")
```



For more details, check out R's help file on `survplot`.

23.3.6 Genotype's effect on log relative hazard

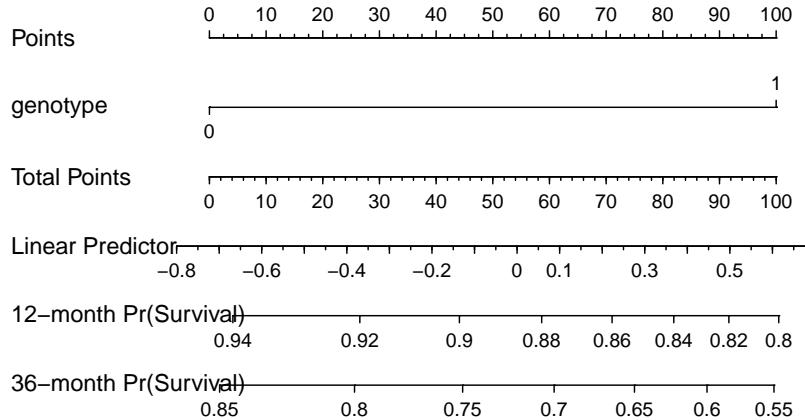
```
ggplot(Predict(model_hem, genotype))
```



23.3.7 Nomogram of our simple hem model

We can estimate 1-year and 3-year hemorrhage-free survival probabilities, for example, with this model, and incorporate these results into our nomogram.

```
survx <- Survival(model_hem)
plot(nomogram(model_hem, fun=list(function(x) survx(12, x),
                                    function(x) survx(36, x)),
             funlabel=c("12-month Pr(Survival)",
                       "36-month Pr(Survival)")))
```



Again, this is just a very simple model, with one binary predictor.

23.3.8 Assessing the Proportional Hazards Assumption

```
cox.zph(model_hem, transform="km")
```

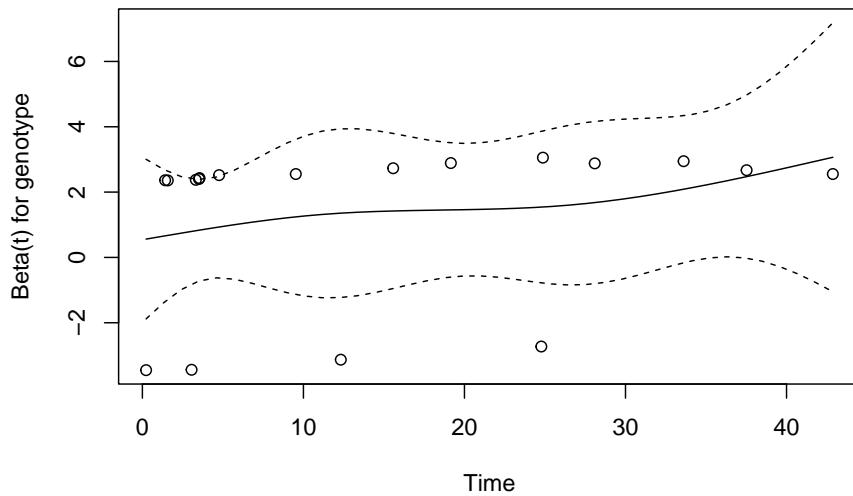
	chisq	df	p
genotype	2.09	1	0.15
GLOBAL	2.09	1	0.15

- Consider using `transform="rank"` to transform the survival times by their ranks prior to performing the test.
- Or use `transform="identity"` as we'll do in the plot below.

23.3.9 Plot to Check PH Assumption

```
plot(cox.zph(model_hem, "identity"))
```

23.3. FITTING A COX MODEL USING CPH FROM THE RMS PACKAGE 587



Chapter 24

Cox Regression Models for Survival Data: Example 2

24.1 A Second Example: The `leukem` data

```
leukem

# A tibble: 51 x 8
  id    age  pblasts  pinf   plab maxtemp months alive
  <dbl> <dbl>    <dbl> <dbl> <dbl>    <dbl>    <dbl> <dbl>
1     1    20      78    39     7    99     18      0
2     2    25      64    61    16   103     31      1
3     3    26      61    55    12   98.2    31      0
4     4    26      64    64    16   100     31      0
5     5    27      95    95     6    98     36      0
6     6    27      80    64     8   101      1      0
7     7    28      88    88    20   98.6     9      0
8     8    28      70    70    14   101     39      1
9     9    31      72    72     5   98.8    20      1
10   10    33      58    58     7   98.6     4      0
# ... with 41 more rows
```

The data describe 51 leukemia patients. The variables are:

- `id`, a patient identification code
- `age`, age at diagnosis
- `pblasts`, the Smear differential percentage of blasts
- `pinf`, the Percentage of absolute marrow leukemia infiltrate
- `plab`, the Percentage labeling index of the bone marrow leukemia cells
- `maxtemp`, Highest temperature prior to treatment (in $^{\circ}F$)

- `months`, which is Survival time from diagnosis (in months)
- `alive`, which indicates Status as of the end of the study (1 = alive and thus censored, 0 = dead)

```
glimpse(leukem)
```

```
Rows: 51
Columns: 8
$ id      <dbl> 1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13, 14, 15, 16, 17, 18, ~
$ age     <dbl> 20, 25, 26, 26, 27, 27, 28, 28, 31, 33, 33, 33, 34, 36, 37, 40, ~
$ pblasts <dbl> 78, 64, 61, 64, 95, 80, 88, 70, 72, 58, 92, 42, 26, 55, 71, 91, ~
$ pinf    <dbl> 39, 61, 55, 64, 95, 64, 88, 70, 72, 58, 92, 38, 26, 55, 71, 91, ~
$ plab    <dbl> 7, 16, 12, 16, 6, 8, 20, 14, 5, 7, 5, 12, 7, 14, 15, 9, 12, 4, ~
$ maxtemp <dbl> 99.0, 103.0, 98.2, 100.0, 98.0, 101.0, 98.6, 101.0, 98.8, 98.6, ~
$ months  <dbl> 18, 31, 31, 31, 36, 1, 9, 39, 20, 4, 45, 36, 12, 8, 1, 15, 24, ~
$ alive   <dbl> 0, 1, 0, 0, 0, 0, 1, 1, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, ~
```

24.1.1 Creating our response: A survival time object

Regardless of how we're going to fit a survival model, we start by creating a *survival time* object that combines the information in `months` (the survival times, possibly censored) and `alive` (the censoring indicator) into a single variable we'll call `stime` in this example.

The function below correctly registers the survival time, and censors subjects who are alive at the end of the study (we need to indicate those whose times are known, and they are identified by `alive == 0`). All other subjects are alive for at least as long as we observe them, but their exact survival times are *right-censored*.

```
stime <- Surv(leukem$months, leukem$alive == 0)
stime
```

```
[1] 18 31+ 31 31 36 1 9 39+ 20+ 4 45+ 36 12 8 1 15 24 2 33
[20] 29+ 7 0 1 2 12 9 1 1 9 5 27+ 1 13 1 5 1 3 4
[39] 1 18 1 2 1 8 3 4 14 3 13 13 1
```

24.1.2 Models We'll Fit

We'll fit several models here, including:

- Model A: A model for survival time using `age` at diagnosis alone.
- Model B: A model for survival time using the main effects of 5 predictors, specifically, `age`, `pblasts`, `pinf`, `plab`, and `maxtemp`.
- Model B2: The model we get after applying stepwise variable selection to Model B, which will include `age`, `pinf` and `plab`.

- Model C: A model using `age` (with a restricted cubic spline), `p1ab` and `maxtemp`

24.2 Model A: coxph Model for Survival Time using `age` at diagnosis

We'll start by using `age` at diagnosis to predict our survival object (survival time, accounting for censoring).

```
modA <- coxph(Surv(months, alive==0) ~ age,
                 data=leukem, model=TRUE)

summary(modA)

Call:
coxph(formula = Surv(months, alive == 0) ~ age, data = leukem,
      model = TRUE)

n= 51, number of events= 45

      coef exp(coef)  se(coef)    z Pr(>|z|)
age  0.032397  1.032927 0.009521 3.403 0.000667 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

      exp(coef) exp(-coef) lower .95 upper .95
age     1.033    0.9681    1.014    1.052

Concordance= 0.65  (se = 0.047 )
Likelihood ratio test= 11.85  on 1 df,   p=6e-04
Wald test       = 11.58  on 1 df,   p=7e-04
Score (logrank) test = 12.29  on 1 df,   p=5e-04

glance(modA) %>%
  select(r.squared, r.squared.max,
         concordance, std.error.concordance)

# A tibble: 1 x 4
  r.squared r.squared.max concordance std.error.concordance
  <dbl>        <dbl>        <dbl>                <dbl>
1 0.207       0.996       0.650               0.0465
```

Across these 51 subjects, we observe 45 events (deaths) and 6 subjects are censored. The hazard ratio (shown under `exp(coef)`) is 1.0329272, and this means each additional year of age at diagnosis is associated with a 1.03-fold increase in the hazard of death.

592CHAPTER 24. COX REGRESSION MODELS FOR SURVIVAL DATA: EXAMPLE 2

For this simple Cox regression model, we will focus on interpreting

1. the **hazard ratio** (specified by the `exp(coef)` result and associated confidence interval) as a measure of effect size,
 - Here, the hazard ratio associated with a 1-year increase in `age` is 1.033, and its 95% confidence interval is: (1.014, 1.052).
 - Since this confidence interval ratio does not include 1 (barely), we can conclude that there is a significant association between `age` and `stime` at the 5% significance level.
2. the **concordance** and **Rsquare** as measures of fit quality, and
 - **concordance** is only appropriate when we have at least one continuous predictor in our Cox model, in which case it assesses the probability of agreement between the survival time and the risk score generated by the predictor (or set of predictors.) A value of 1 indicates perfect agreement, but values of 0.6 to 0.7 are more common in survival data. 0.5 is an agreement that is no better than chance. Here, our concordance is 0.65, which is a fairly typical value.
 - **Rsquare** in this setting is Cox and Snell's pseudo- R^2 , which reflects the improvement of the model we have fit over the model with the intercept alone - a comparison that is tested by the likelihood ratio test. The maximum value of this statistic is often less than one, in which case R will tell you that. Here, our observed pseudo- R^2 is 0.207 and that is out of a possible maximum of 0.996.
3. the significance tests, particularly the **Wald** test (shown next to the coefficient estimates in the position of a t test in linear regression), and the **Likelihood ratio** test at the bottom of the output, which compares this model to a null model which predicts the mean survival time for all subjects.
 - The Wald test for an individual predictor compares the coefficient to its standard error, just like a t test in linear regression.
 - The likelihood ratio test compares the entire model to the null model (intercept-only). Again, run an ANOVA (technically an analysis of deviance) to get more details on the likelihood-ratio test.

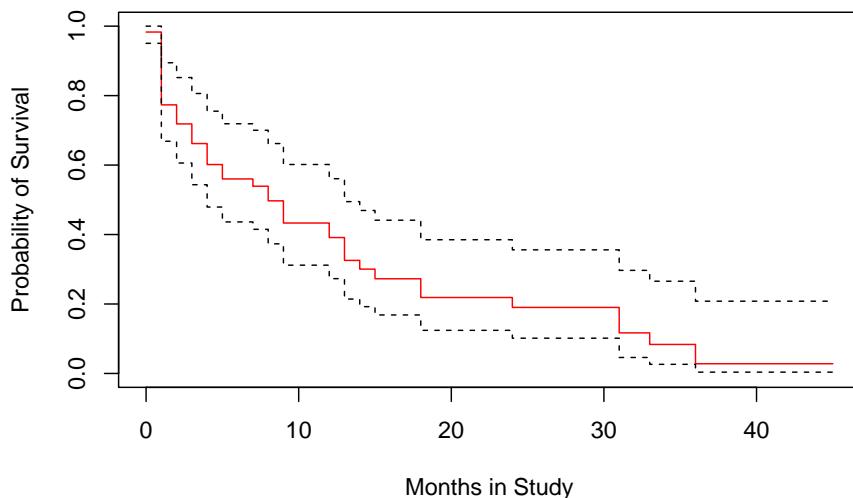
```
anova(modA)
```

```
Analysis of Deviance Table
Cox model: response is Surv(months, alive == 0)
Terms added sequentially (first to last)

loglik  Chisq Df Pr(>|Chi|)
NULL   -142.94
age    -137.02 11.849  1   0.000577 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

24.2.1 Plotting the Survival Curve implied by Model A

```
plot(survfit(modA), ylab="Probability of Survival",
      xlab="Months in Study", col=c("red", "black", "black"))
```



24.2.2 Testing the Proportional Hazards Assumption

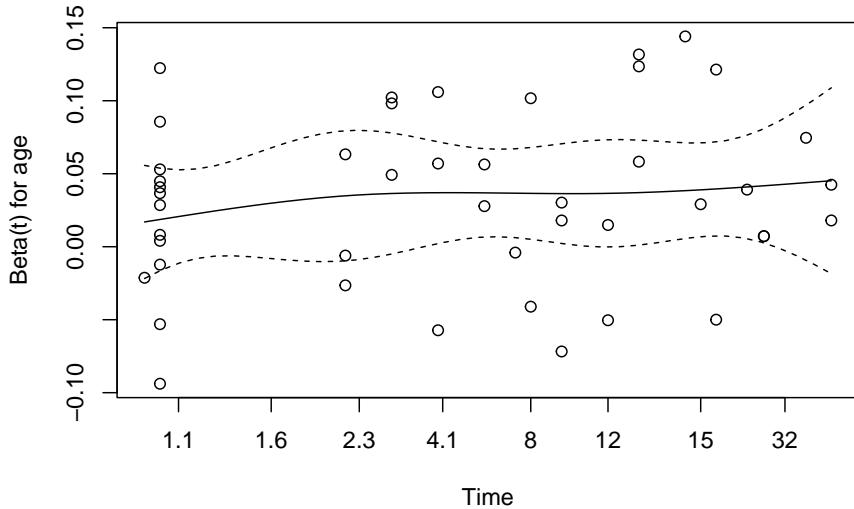
As we've noted, the key assumption in a Cox model is that the hazards are **proportional**.

```
cox.zph(modA)
```

	chisq	df	p
age	1.05	1	0.31
GLOBAL	1.05	1	0.31

A significant result here would indicate a problem with the proportional hazards assumption - again, not the case here. We can also plot the results:

```
plot(cox.zph(modA))
```



We're looking for the smooth curve to be fairly level across the time horizon here, as opposed to substantially increasing or decreasing in level as time passes.

24.3 Building Model A with cph for the leukem data

```
units(leukem$months) <- "month"
d <- datadist(leukem)
options(datadist="d")
modA_cph <- cph(Surv(months, alive==0) ~ age, data=leukem,
                  x=TRUE, y=TRUE, surv=TRUE, time.inc=12)

modA_cph
```

Cox Proportional Hazards Model

```
cph(formula = Surv(months, alive == 0) ~ age, data = leukem,
      x = TRUE, y = TRUE, surv = TRUE, time.inc = 12)
```

Model Tests	Discrimination
	Indexes

```

Obs      51    LR chi2     11.85    R2      0.208
Events    45    d.f.          1    Dxy     0.301
Center 1.6152    Pr(> chi2) 0.0006    g      0.621
                           Score chi2 12.29    gr     1.861
                           Pr(> chi2) 0.0005

```

```

Coef    S.E.    Wald Z Pr(>|Z|)
age 0.0324 0.0095 3.40  0.0007

```

```
exp(coef(modA_cph)) # hazard ratio estimate
```

```

age
1.032923

```

```
exp(confint(modA_cph)) # hazard ratio 95% CI
```

```

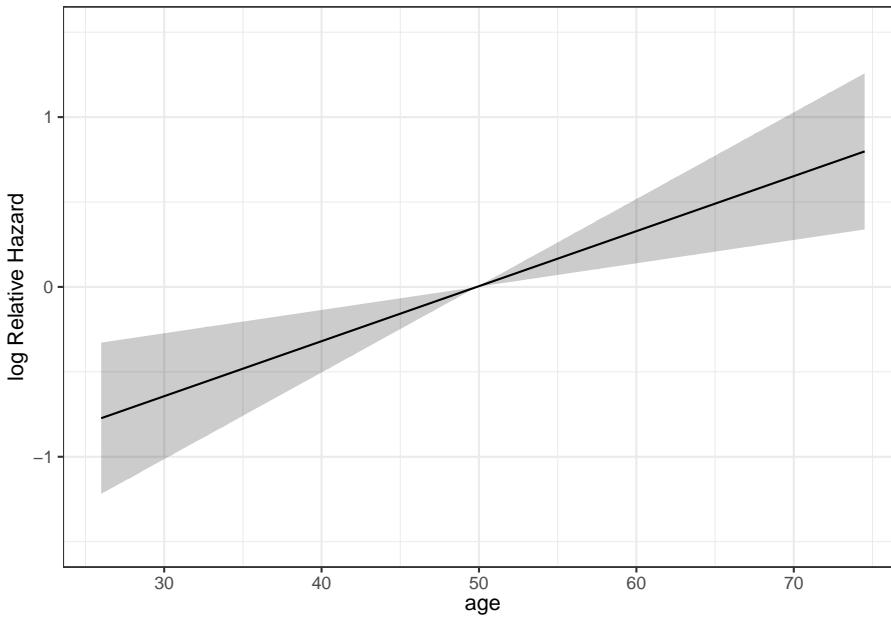
2.5 %   97.5 %
age 1.013826 1.052379

```

24.3.1 Plotting the age effect implied by our model.

We can plot the age effect implied by the model, using ggplot2, as follows...

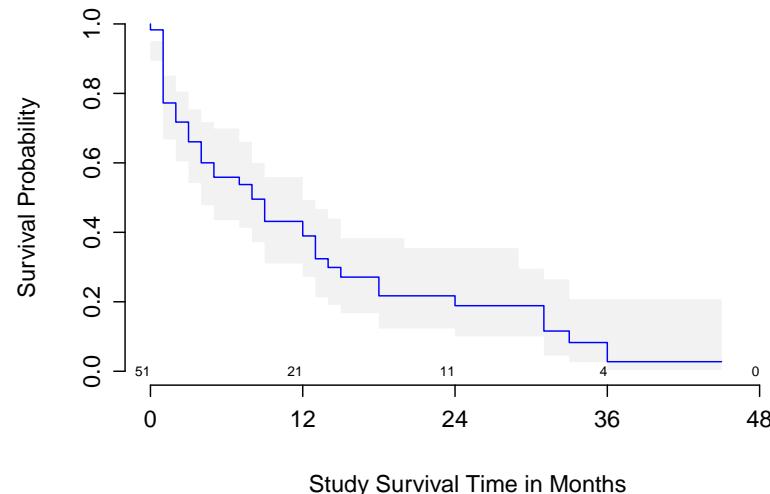
```
ggplot(Predict(modA_cph, age))
```



24.3.2 Survival Plots (Kaplan-Meier) of the age effect

The first survival plot I'll show displays 95% confidence intervals for the probability of survival at the median age at diagnosis in the sample, which turns out to be 50 years, with numbers of patients still at risk indicated every 12 months of time in the study. We can substitute in `conf = bars` to get a different flavor for this sort of plot.

```
survplot(modA_cph, age=median(leukem$age), conf.int=.95,
         col='blue', time.inc=12, n.risk=TRUE,
         conf='bands', type="kaplan-meier",
         xlab="Study Survival Time in Months")
```



Or we can generate a survival plot that shows survival probabilities over time across a range of values for `age` at diagnosis, as follows...

```
survplot(modA_cph, levels.only=TRUE, time.inc=12,
         type="kaplan-meier",
         xlab="Study Survival Time in Months")
```

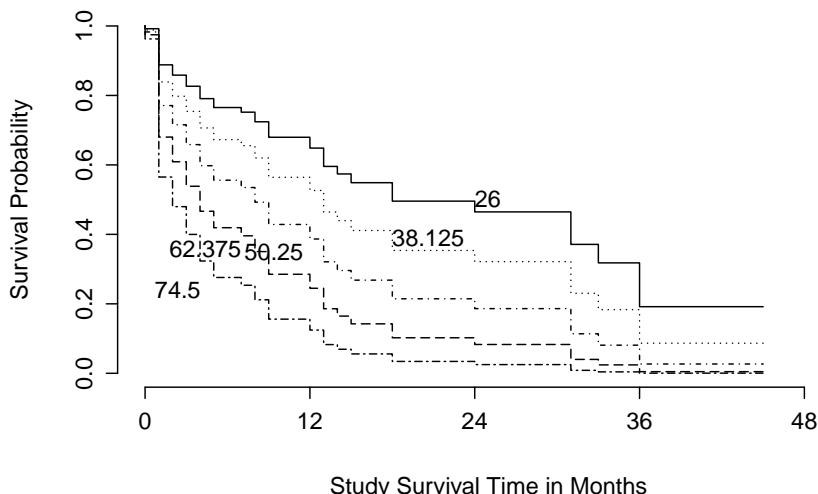
Warning in `regularize.values(x, y, ties, missing(ties), na.rm = na.rm)`:
collapsing to unique 'x' values

Warning in `regularize.values(x, y, ties, missing(ties), na.rm = na.rm)`:
collapsing to unique 'x' values

```
Warning in regularize.values(x, y, ties, missing(ties), na.rm = na.rm):
collapsing to unique 'x' values
```

```
Warning in regularize.values(x, y, ties, missing(ties), na.rm = na.rm):
collapsing to unique 'x' values
```

```
Warning in regularize.values(x, y, ties, missing(ties), na.rm = na.rm):
collapsing to unique 'x' values
```



This plot shows a series of modeled survival probabilities, for five different diagnosis `age` levels, as identified by the labels. Generally we see that the younger the subject is at diagnosis, the longer their survival time in the study.

24.3.3 ANOVA test for the cph-built model for `leukem`

We can run a likelihood-ratio (drop in deviance) test of the significance of the `age` effect...

```
anova(modA_cph)
```

Wald Statistics			Response: Surv(months, alive == 0)		
Factor	Chi-Square	d.f.	P		
age	11.57	1	7e-04		
TOTAL	11.57	1	7e-04		

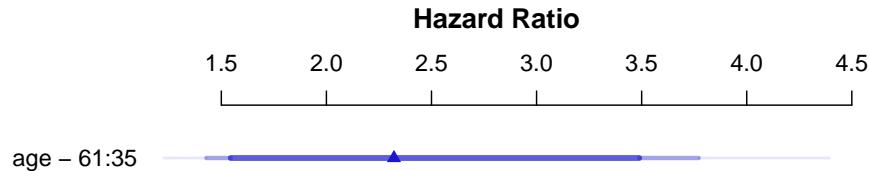
24.3.4 Summarizing the Effect Sizes from `modA_cph`

We can generate the usual summaries of effect size in this context, too.

```
summary(modA_cph)
```

Effects				Response : Surv(months, alive == 0)			
Factor	Low	High	Diff.	Effect	S.E.	Lower	0.95 Upper
age	35	61	26	0.8422	0.24755	0.35701	0.95 1.3274
Hazard Ratio	35	61	26	2.3215	NA	1.42910	3.7712

```
plot(summary(modA_cph))
```



As with all `rms` package effect estimates, this quantitative predictor (`age`) yields an effect comparing `age` at the 25th percentile of the sample (`age = 35`) to `age` at the 75th percentile (`age = 61`). So the hazard ratio is 2.32, with 95% CI (1.43, 3.77) for the effect of moving 26 years. Our `coxph` version of this same model showed a hazard ratio for the effect of moving just a single year.

24.3.5 Validating the Cox Model Summary Statistics

```
set.seed(432410); validate(modA_cph)
```

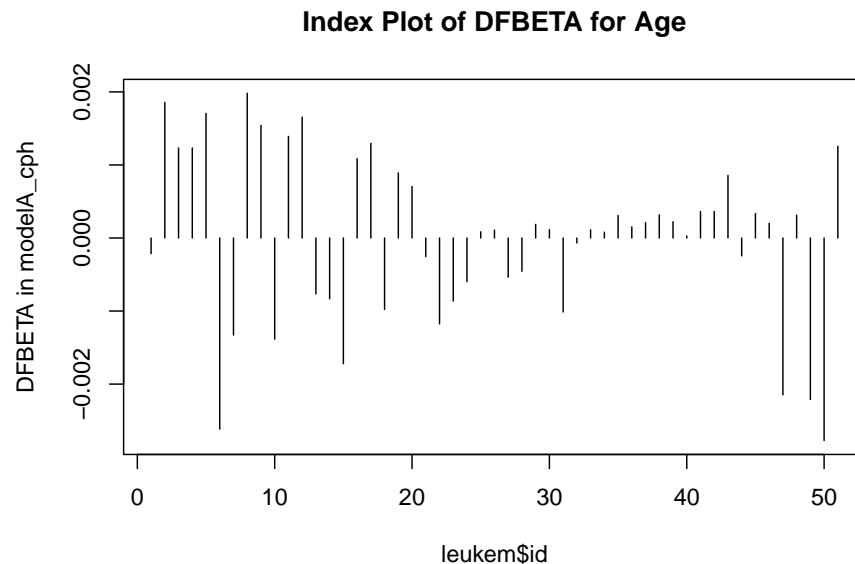
	index.orig	training	test	optimism	index.corrected	n
Dxy	0.3007	0.2950	0.3007	-0.0057	0.3064	40
R2	0.2081	0.2226	0.2081	0.0145	0.1936	40
Slope	1.0000	1.0000	1.0253	-0.0253	1.0253	40
D	0.0379	0.0425	0.0379	0.0045	0.0334	40
U	-0.0070	-0.0070	0.0037	-0.0107	0.0037	40
Q	0.0449	0.0495	0.0343	0.0152	0.0297	40
g	0.6210	0.6492	0.6210	0.0282	0.5928	40

The R^2 statistic barely moves, and neither does the Somers' d estimate, so at least in this simple model, the nominal summary statistics are likely to hold up pretty well in new data.

24.3.6 Looking for Influential Points

This plot shows the influence of each point, in terms of DFBETA - the impact on the coefficient of `age` in the model were that specific point to be removed from the data set. We can also identify the row numbers of the largest (positive and negative) DFBETAs.

```
plot(residuals(modA_cph, type="dfbeta",
               collapse = leukem$id) ~
      leukem$id, main="Index Plot of DFBETA for Age",
      type="h", ylab="DFBETA in modelA_cph")
```



```
which.max(residuals(modA_cph, type="dfbeta"))
```

```
8  
8
```

```
which.min(residuals(modA_cph, type="dfbeta"))
```

```
50  
50
```

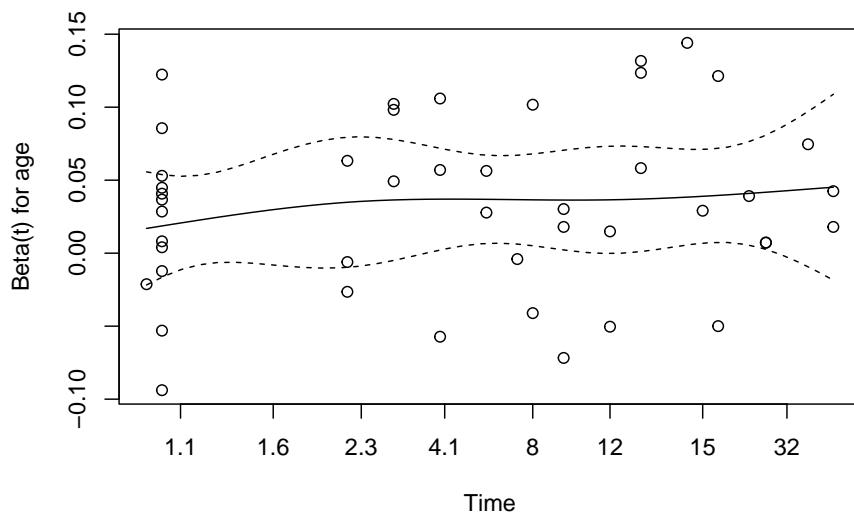
The DFBETAs look very small here. Changes in the β estimates as large as 0.002 don't have a meaningful impact in this case, so I don't see anything particularly influential.

24.3.7 Checking the Proportional Hazards Assumption

As before, we can check the proportional hazards assumption with a test, or plot.

```
cox.zph(modA_cph)
```

```
      chisq df      p
age     1.05  1 0.31
GLOBAL  1.05  1 0.31
plot(cox.zph(modA_cph))
```



Still no serious signs of trouble, of course. We'll see what happens when we fit a bigger model.

24.4 Model B: Fitting a 5-Predictor Model with `coxph`

Next, we use the `coxph` function from the `survival` package to apply a Cox regression model to predict the survival time using the main effects of the five predictors: `age`, `pblasts`, `pinf`, `plab` and `maxtemp`.

```
modB <- coxph(Surv(months, alive==0) ~
                 age + pblasts + pinf + plab + maxtemp, data=leukem)
modB
```

Call:

```
coxph(formula = Surv(months, alive == 0) ~ age + pblasts + pinf +
      plab + maxtemp, data = leukem)
```

	coef	exp(coef)	se(coef)	z	p
age	0.033080	1.033633	0.010163	3.255	0.00113
pblasts	0.009452	1.009497	0.013959	0.677	0.49831
pinf	-0.017102	0.983043	0.012244	-1.397	0.16248
plab	-0.066000	0.936131	0.038651	-1.708	0.08771
maxtemp	0.155448	1.168182	0.111978	1.388	0.16507

Likelihood ratio test=18.48 on 5 df, p=0.002405
n= 51, number of events= 45

The Wald tests suggest that `age` and perhaps `plab` are the only terms which appear to be significant after accounting for the other predictors in the model.

24.4.1 Interpreting the Results from Model B

```
summary(modB)
```

Call:

```
coxph(formula = Surv(months, alive == 0) ~ age + pblasts + pinf +
      plab + maxtemp, data = leukem)
```

n= 51, number of events= 45

	coef	exp(coef)	se(coef)	z	Pr(> z)
age	0.033080	1.033633	0.010163	3.255	0.00113 **
pblasts	0.009452	1.009497	0.013959	0.677	0.49831
pinf	-0.017102	0.983043	0.012244	-1.397	0.16248
plab	-0.066000	0.936131	0.038651	-1.708	0.08771 .
maxtemp	0.155448	1.168182	0.111978	1.388	0.16507

602CHAPTER 24. COX REGRESSION MODELS FOR SURVIVAL DATA: EXAMPLE 2

```

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

      exp(coef) exp(-coef) lower .95 upper .95
age       1.0336     0.9675    1.0132    1.054
pblasts   1.0095     0.9906    0.9823    1.037
pinf       0.9830     1.0172    0.9597    1.007
plab       0.9361     1.0682    0.8678    1.010
maxtemp   1.1682     0.8560    0.9380    1.455

Concordance= 0.705 (se = 0.039 )
Likelihood ratio test= 18.48 on 5 df,  p=0.002
Wald test            = 17.62 on 5 df,  p=0.003
Score (logrank) test = 18.96 on 5 df,  p=0.002

```

Again, it appears that only `age` has a statistically significant effect, as last predictor in.

```
anova(modB)
```

```

Analysis of Deviance Table
Cox model: response is Surv(months, alive == 0)
Terms added sequentially (first to last)

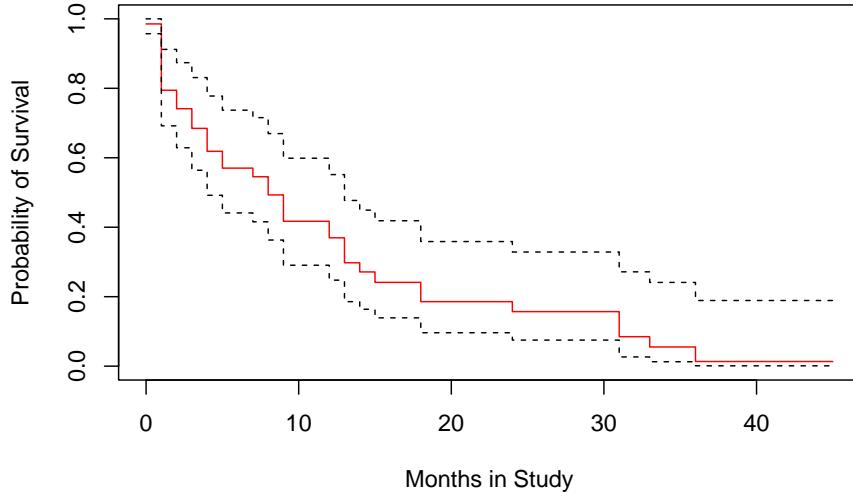
      loglik  Chisq Df Pr(>|Chi|)
NULL    -142.94
age     -137.02 11.8488  1  0.000577 ***
pblasts -136.64  0.7537  1  0.385319
pinf     -135.77  1.7308  1  0.188314
plab     -134.60  2.3370  1  0.126334
maxtemp -133.70  1.8057  1  0.179022
---
Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

```

At least taken in this order, none of the variables appear to add significant predictive value, given that we have already accounted for the preceding variables.

24.4.2 Plotting the Survival Curve implied by Model B

```
plot(survfit(modB), ylab="Probability of Survival",
      xlab="Months in Study", col=c("red", "black", "black"))
```



The crosses in the plot indicate censoring points, while the drops indicate people who have died, and are thus no longer at risk.

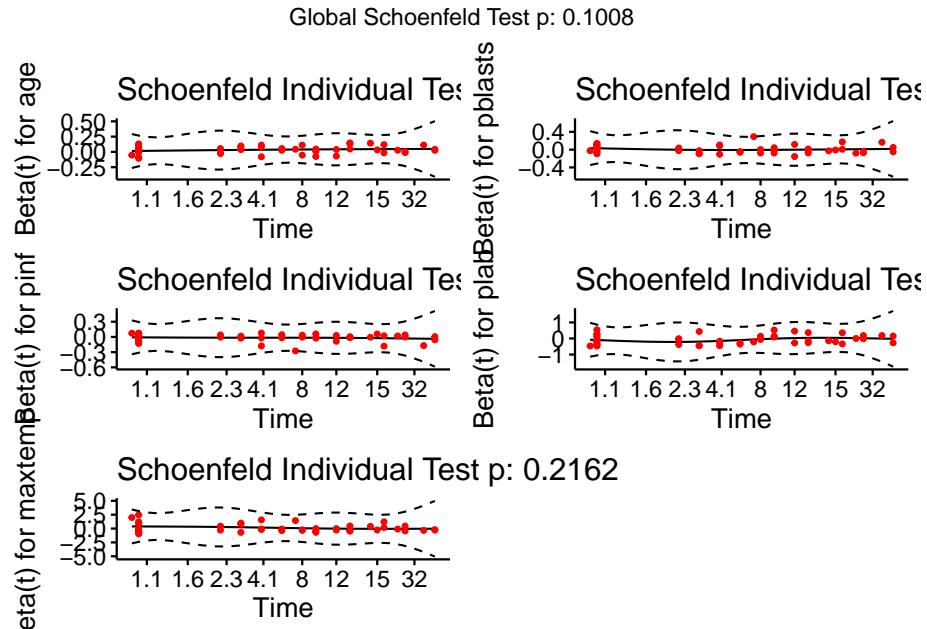
24.4.3 Testing the Proportional Hazards Assumption

```
cox.zph(modB, transform="km", global=TRUE)
```

	chisq	df	p
age	1.87	1	0.171
pblasts	4.37	1	0.037
pinf	3.51	1	0.061
plab	1.19	1	0.275
maxtemp	1.53	1	0.216
GLOBAL	9.22	5	0.101

Note that we get a global test, and a separate test for each predictor. None show significant problems. We can plot the scaled Schoenfeld residuals directly with `ggcoxph` from the `survminer` package.

```
ggcoxzph(cox.zph(modB))
```



24.4.4 Assessing Collinearity

Perhaps we have some collinearity here, which might imply that we could sensibly fit a smaller model, which would be appealing anyway, with only 45 actual events - we should probably be sticking to a model with no more than 2 or perhaps as many as 3 coefficients to be estimated.

```
rms::vif(modB)
```

	age	pblasts	pinf	plab	maxtemp
1.081775	3.029862	3.000944	1.035400	1.045249	

The variance inflation factors don't look enormous - it may be that removing one of these variables will help make the others look more significant. Let's consider a stepwise variable selection algorithm to see what results...

- Note that the `leaps` library, which generates best subsets output, is designed for linear regression, as is the `lars` library, which generates the lasso. Either could be used here for some guidance, but not with the survival object `stime = Surv(months, age)` as the response, but instead only with `months` as the outcome, which ignores the censoring. The `step` procedure can be used on the survival object, though.

24.5 Model B2: A Stepwise Reduction of Model B

```
stats::step(modB)
```

Start: AIC=277.4
 Surv(months, alive == 0) ~ age + pblast + pinf + plab + maxtemp

	Df	AIC
- pblast	1	275.85
- pinf	1	277.17
- maxtemp	1	277.21
<none>		277.40
- plab	1	278.42
- age	1	286.47

Step: AIC=275.85
 Surv(months, alive == 0) ~ age + pinf + plab + maxtemp

	Df	AIC
- maxtemp	1	275.63
<none>		275.85
- pinf	1	275.89
- plab	1	276.86
- age	1	284.47

Step: AIC=275.63
 Surv(months, alive == 0) ~ age + pinf + plab

	Df	AIC
<none>		275.63
- pinf	1	275.69
- plab	1	275.95
- age	1	285.52

Call:
 coxph(formula = Surv(months, alive == 0) ~ age + pinf + plab,
 data = leukem)

	coef	exp(coef)	se(coef)	z	p
age	0.033171	1.033727	0.009733	3.408	0.000655
pinf	-0.010147	0.989905	0.007088	-1.432	0.152246
plab	-0.057558	0.944067	0.038476	-1.496	0.134662

Likelihood ratio test=16.25 on 3 df, p=0.001008

606CHAPTER 24. COX REGRESSION MODELS FOR SURVIVAL DATA: EXAMPLE 2

```
n= 51, number of events= 45
```

The stepwise procedure lands on a model with three predictors. How does this result look, practically?

```
modB2 <- coxph(Surv(months, alive==0) ~ age + pinf + plab, data=leukem)
summary(modB2)
```

Call:

```
coxph(formula = Surv(months, alive == 0) ~ age + pinf + plab,
      data = leukem)
```

```
n= 51, number of events= 45
```

	coef	exp(coef)	se(coef)	z	Pr(> z)
age	0.033171	1.033727	0.009733	3.408	0.000655 ***
pinf	-0.010147	0.989905	0.007088	-1.432	0.152246
plab	-0.057558	0.944067	0.038476	-1.496	0.134662

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

	exp(coef)	exp(-coef)	lower	upper
age	1.0337	0.9674	1.0142	1.054
pinf	0.9899	1.0102	0.9762	1.004
plab	0.9441	1.0592	0.8755	1.018

Concordance= 0.676 (se = 0.046)

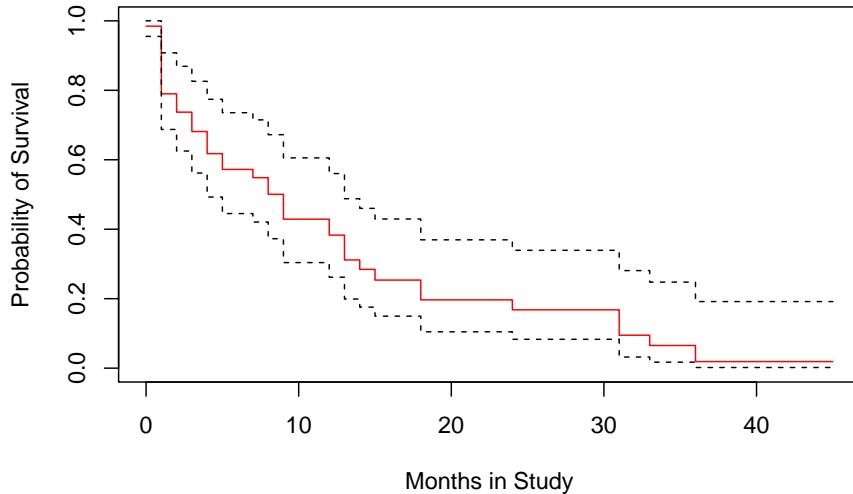
Likelihood ratio test= 16.25 on 3 df, p=0.001

Wald test = 15.28 on 3 df, p=0.002

Score (logrank) test = 16.21 on 3 df, p=0.001

24.5.1 The Survival Curve implied by Model B2

```
plot(survfit(modB2), ylab="Probability of Survival", xlab="Months in Study",
      col=c("red", "black", "black"))
```

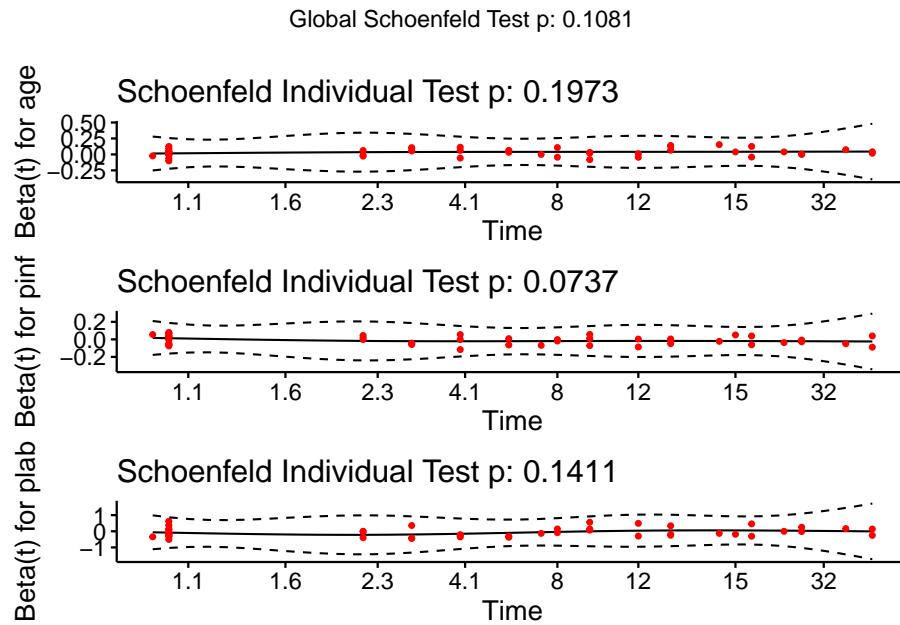


24.5.2 Checking Proportional Hazards for Model B2

```
cox.zph(modB2, transform="km", global=TRUE)
```

```
      chisq df      p
age     1.66  1 0.197
pinf    3.20  1 0.074
plab    2.17  1 0.141
GLOBAL  6.07  3 0.108
```

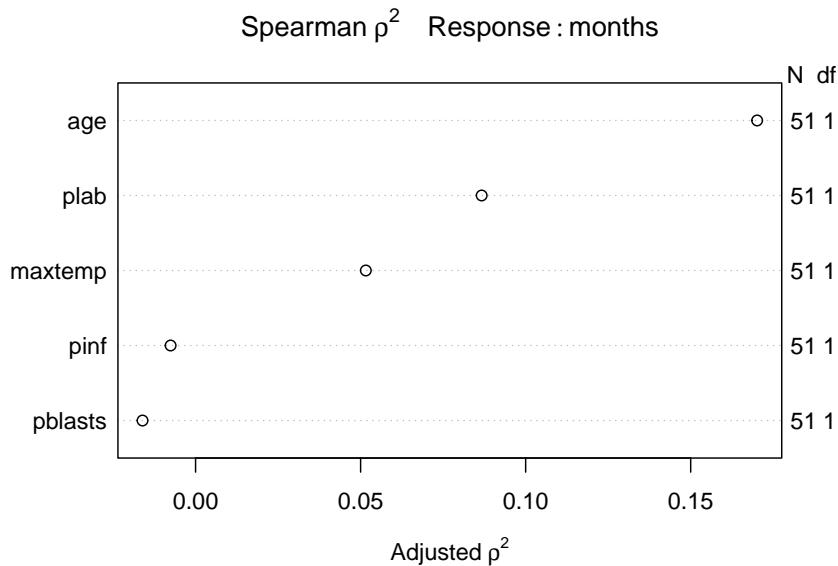
```
ggcoxzph(cox.zph(modB2))
```



24.6 Model C: Using a Spearman Plot to pick a model

If we want to use the **Spearman** ρ^2 plot to consider how we might perhaps incorporate non-linear terms describing any or all of the five potential predictors (`age`, `pblasts`, `pinf`, `plab` and `maxtemp`) for survival time, we need to do so on the raw `months` variable, rather than the survival object (`stime = Surv(months, alive==0)`) which accounts for censoring...

```
plot(spearman2(months ~ age + pblasts + pinf + plab + maxtemp, data=leukem))
```



Recognizing that we can probably only fit a small model safely (since we observe only 45 actual [uncensored] survival times) I will consider a non-linear term in `age` (specifically a restricted cubic spline with 3 knots), along with linear terms for `plab` and `maxtemp`. I'm mostly just looking for a new model to study for this example.

24.6.1 Fitting Model C

```
# still have datadist set up for leukem
modC <- cph(Surv(months, alive==0) ~ rcs(age, 3) + plab +
              maxtemp, data=leukem, x=TRUE, y=TRUE, surv=TRUE, time.inc=12)
modC

Cox Proportional Hazards Model

cph(formula = Surv(months, alive == 0) ~ rcs(age, 3) + plab +
      maxtemp, data = leukem, x = TRUE, y = TRUE, surv = TRUE,
      time.inc = 12)

      Model Tests      Discrimination
                           Indexes
Obs       51    LR chi2     19.75      R2      0.322
Events     45    d.f.          4      Dxy      0.438
Center 20.3856  Pr(> chi2) 0.0006      g      0.947
```

610CHAPTER 24. COX REGRESSION MODELS FOR SURVIVAL DATA: EXAMPLE 2

```

Score chi2 18.66    gr      2.577
Pr(> chi2) 0.0009

Coef     S.E.   Wald Z Pr(>|Z|)
age      0.0804 0.0283  2.84  0.0045
age'     -0.0629 0.0332 -1.90  0.0580
plab     -0.0736 0.0381 -1.93  0.0536
maxtemp  0.1788 0.1145  1.56  0.1184

```

24.6.2 ANOVA for Model C

```
anova(modC)
```

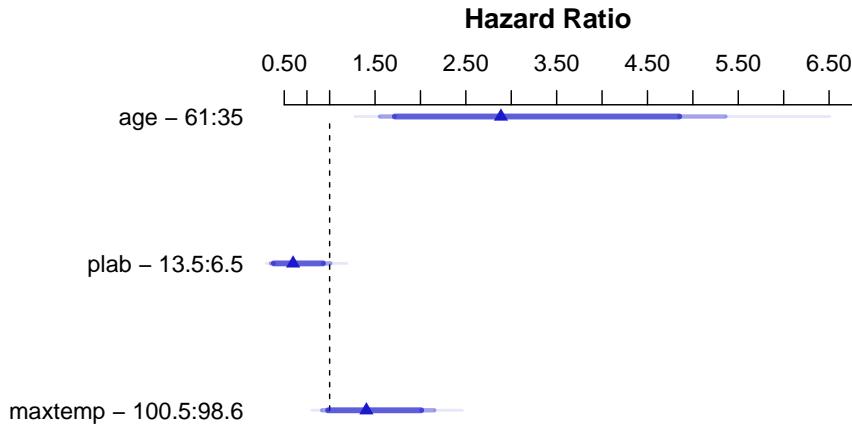
	Wald Statistics			Response: Surv(months, alive == 0)	
Factor	Chi-Square	d.f.	P		
age	11.55	2	0.0031		
Nonlinear	3.59	1	0.0580		
plab	3.73	1	0.0536		
maxtemp	2.44	1	0.1184		
TOTAL	17.13	4	0.0018		

24.6.3 Summarizing Model C Effect Sizes

```
summary(modC)
```

	Effects						Response : Surv(months, alive == 0)		
Factor	Low	High	Diff.	Effect	S.E.	Lower	0.95	Upper	0.95
age	35.0	61.0	26.0	1.06010	0.31559	0.44152	1.6786000		
Hazard Ratio	35.0	61.0	26.0	2.88650		NA	1.55510	5.3580000	
plab	6.5	13.5	7.0	-0.51511	0.26686	-1.03820	0.0079317		
Hazard Ratio	6.5	13.5	7.0	0.59744		NA	0.35411	1.0080000	
maxtemp	98.6	100.5	1.9	0.33979	0.21758	-0.08666	0.7662400		
Hazard Ratio	98.6	100.5	1.9	1.40470		NA	0.91699	2.1517000	

```
plot(summary(modC))
```



24.6.4 Plotting the diagnosis `age` effect in Model C

Of course, we're no longer assuming that the log relative hazard is linear in `age`, once we include a restricted cubic spline for `age` in our Model C. So our hazard ratio and confidence intervals for `age` are a bit trickier to understand.

```
exp(coef(modC))
```

```
age      age'      plab    maxtemp
1.0837479 0.9390008 0.9290553 1.1958265
```

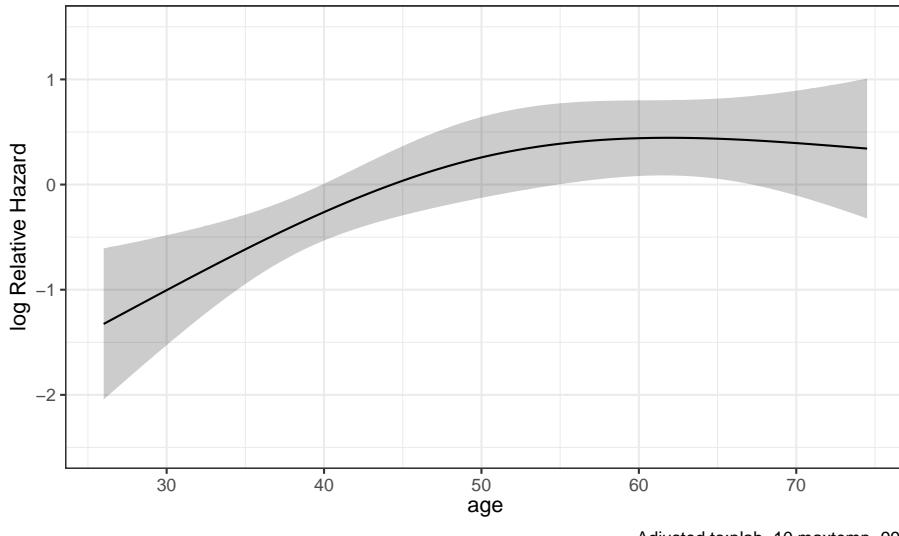
```
exp(confint(modC))
```

	2.5 %	97.5 %
<code>age</code>	1.0252687	1.145563
<code>age'</code>	0.8798328	1.002148
<code>plab</code>	0.8621662	1.001134
<code>maxtemp</code>	0.9554141	1.496734

We can use `ggplot` and the `Predict` function to produce plots of the log Relative Hazard associated with any of our predictors, while holding the others constant at their medians. The effects of `maxtemp` and `plab` in our Model C are linear in the log Relative Hazard, but `age`, thanks to our use of a restricted cubic spline with 3 knots, shows a single bend.

612CHAPTER 24. COX REGRESSION MODELS FOR SURVIVAL DATA: EXAMPLE 2

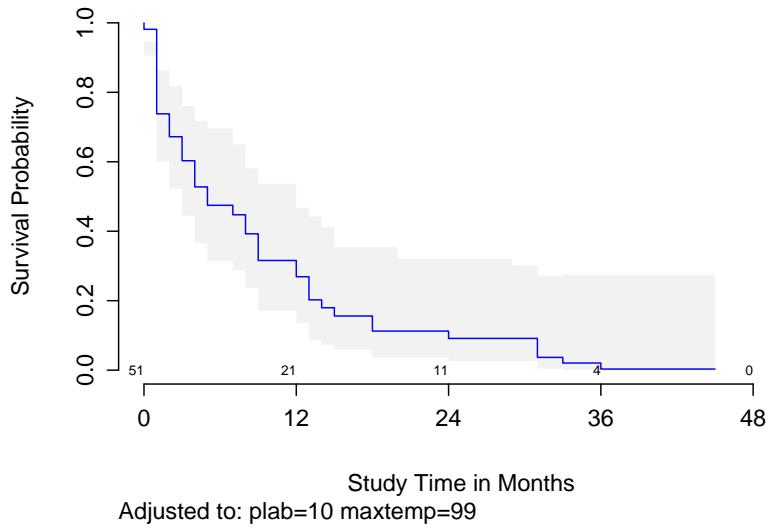
```
ggsplot(Predict(modC, age))
```



24.6.5 Survival Plot associated with Model C

Let's look at a survival plot associated with Model C for a subject with median values of our three predictors.

```
survplot(modC, age=median(leukem$age), conf.int=0.95, col="blue",
         time.inc=12, n.risk=TRUE, conf='bands',
         xlab="Study Time in Months")
```



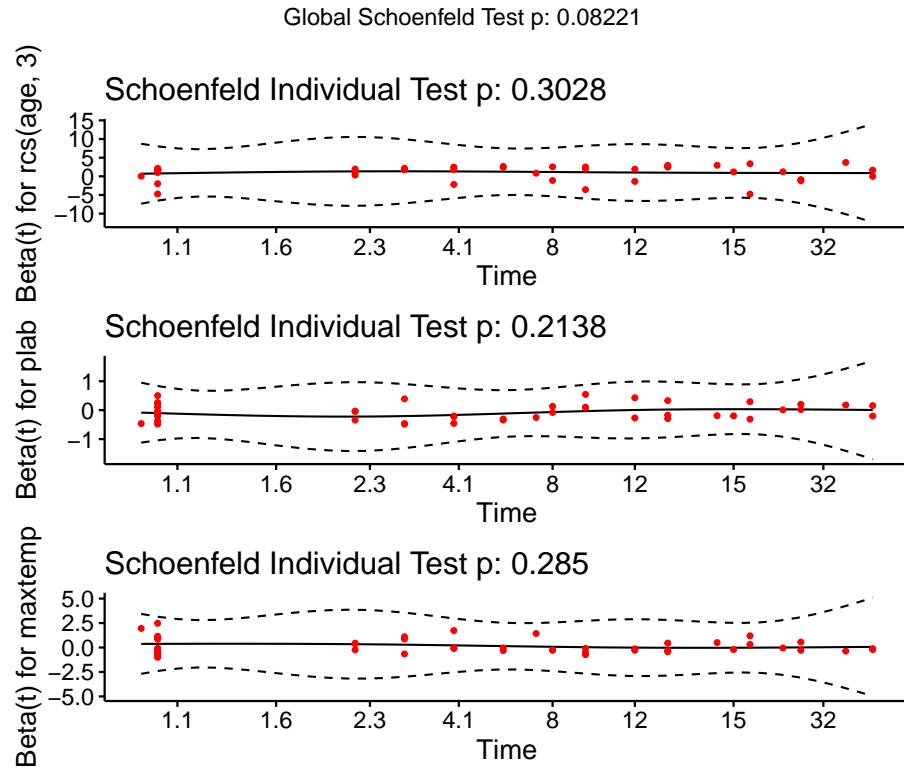
As before, we could fit such a plot to compare results across multiple `age` values, if desired.

24.6.6 Checking the Proportional Hazards Assumption

```
cox.zph(modC, transform="km", global=TRUE)
```

	chisq	df	p
<code>rcs(age, 3)</code>	2.39	2	0.303
<code>plab</code>	1.55	1	0.214
<code>maxtemp</code>	1.14	1	0.285
<code>GLOBAL</code>	8.27	4	0.082

```
ggcooxzph(cox.zph(modC))
```



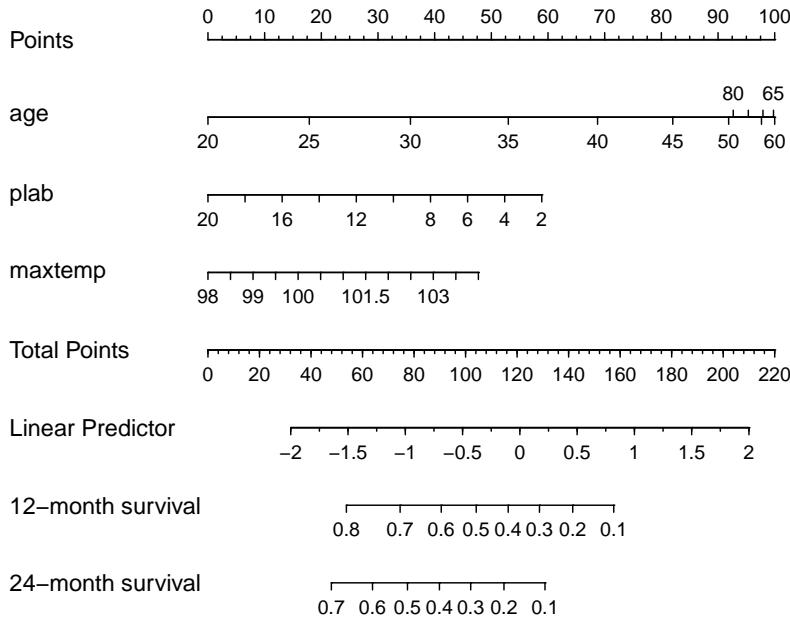
24.6.7 Model C Nomogram

```

sv <- Survival(modC)
surv12 <- function(x) sv(12, lp=x)
surv24 <- function(x) sv(24, lp=x)

plot(nomogram(modC, fun=list(surv12, surv24),
               funlabel=c('12-month survival', '24-month survival'))))

```



24.6.8 Validating Model C's Summary Statistics

We can validate the model for Somers' D_{xy} , which is the rank correlation between the predicted log hazard and observed survival time, and for slope shrinkage.

```
set.seed(43234); validate(modC, B=100)
```

	index.orig	training	test	optimism	index.corrected	n
Dxy	0.4385	0.4749	0.4144	0.0605	0.3780	100
R2	0.3222	0.3716	0.2785	0.0931	0.2292	100
Slope	1.0000	1.0000	0.8101	0.1899	0.8101	100
D	0.0656	0.0814	0.0547	0.0266	0.0390	100
U	-0.0070	-0.0070	0.0101	-0.0171	0.0101	100
Q	0.0726	0.0884	0.0446	0.0437	0.0288	100
g	0.9466	1.0751	0.8364	0.2387	0.7078	100

24.6.9 Calibration of Model C (12-month survival estimates)

Finally, we validate the model for calibration accuracy in predicting the probability of surviving one year.

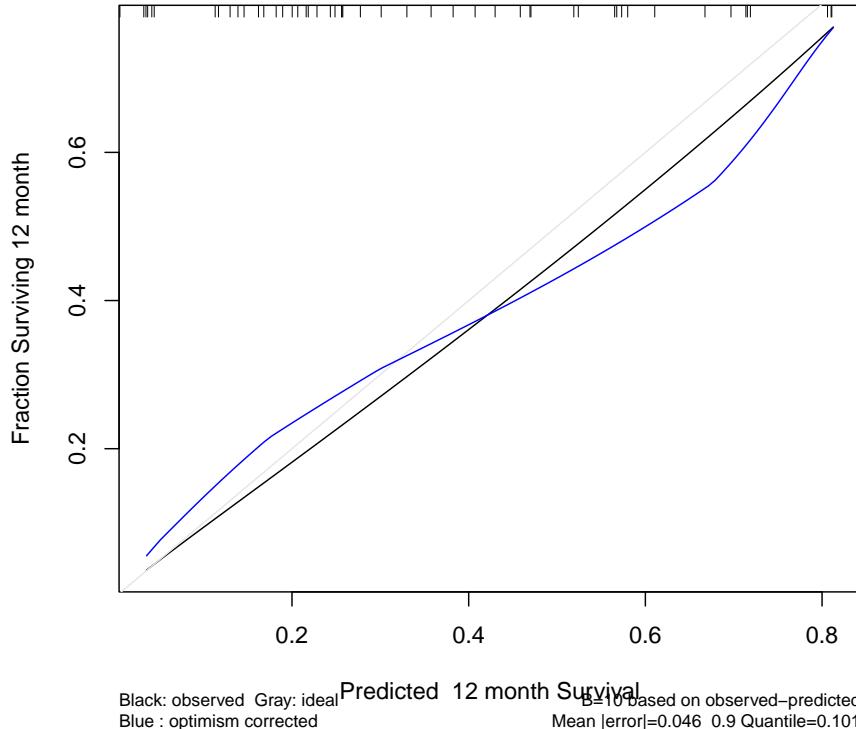
Quoting Harrell (page 529, RMS second edition):

The bootstrap is used to estimate the optimism in how well predicted [one]-year survival from the final Cox model tracks flexible smooth estimates, without any binning of predicted survival probabilities or assuming proportional hazards.

The `u` variable specifies the length of time at which we look at the calibration. I've specified the `units` earlier to be months.

```
set.seed(43233); plot(rms::calibrate(modC, B = 10, u = 12))
```

Using Cox survival estimates at 12 months



The model seems neither especially well calibrated nor especially poorly so -

looking at the comparison of the blue curve to the gray, our predictions basically aren't aggressive enough - more people are surviving to a year in our low predicted probability of 12 month survival group, and fewer people are surviving on the high end of the x-axis than should be the case.

- Barnett, Peggy A., Simon Roman-Golstein, Fred Ramsey, et al. 1995. "Differential Permeability and Quantitative MR Imaging of a Human Lung Carcinoma Brain Xenograft in the Nude Rat." *American Journal of Pathology* 146(2): 436–49. <https://www.ncbi.nlm.nih.gov/pmc/articles/PMC1869863/>.
- Berkhemer, Olvert A., Puck S. S. Fransen, Debbie Buemer, et al. 2015. "A Randomized Trial of Intraarterial Treatment for Acute Ischemic Stroke." *New England Journal of Medicine* 372: 11–20. <http://www.nejm.org/doi/full/10.1056/NEJMoa1411587>.
- Faraway, Julian J. 2006. *Extending the Linear Model with r*. Boca Raton, FL: CRC Press.
- Gelman, Andrew, and Jennifer Hill. 2007. *Data Analysis Using Regression and Multilevel/Hierarchical Models*. New York: Cambridge University Press.
- Harrell, Frank E. 2001. *Regression Modeling Strategies*. New York: Springer.
- . 2018. *Regression Modeling Strategies Course Notes*. <http://www.fharrell.com/#links>.
- Kim, Hae-Young. 2014. "Statistical Notes for Clinical Researchers: Two-Way Analysis of Variance (ANOVA) - Exploring Possible Interaction Between Factors." *Restorative Dentistry & Endodontics* 39(2): 143–47. <https://www.ncbi.nlm.nih.gov/pmc/articles/PMC3978106/>.
- Long, J. Scott. 1997. *Regression Models for Categorical and Limited Dependent Variables*. Thousand Oaks, CA: Sage Publications.
- McDonald, Gary C., and Richard C. Schwing. 1973. "Instabilities of Regression Estimates Relating Air Pollution to Mortality." *Technometrics* 15 (3): 463–81.
- Peduzzi, Peter, John Concato, Elizabeth Kemper, Theodore R. Holford, and Alvan R. Feinstein. 1996. "A Simulation Study of the Number of Events Per Variable in Logistic Regression Analysis." *Journal of Clinical Epidemiology* 49 (12): 1373–79.
- Ramsey, Fred L., and Daniel W. Schafer. 2002. *The Statistical Sleuth: A Course in Methods of Data Analysis*. Second Edition. Pacific Grove, CA: Duxbury.
- Riffenburgh, Robert H. 2006. *Statistics in Medicine*. Second Edition. Burlington, MA: Elsevier Academic Press.
- Rosenbaum, Paul R. 2017. *Observation and Experiment: An Introduction to Causal Inference*. Cambridge, MA: Harvard University Press.
- Roy, Denis, Mario Talajic, Stanley Nattel, et al. 2008. "Rhythm Control Versus Rate Control for Atrial Fibrillation and Heart Failure." *New England Journal of Medicine* 358: 2667–77. <http://www.nejm.org/doi/full/10.1056/NEJMoa0708789>.
- Stamey, J. N. Kabalin, T. A. et al. 1989. "Prostate Specific Antigen in the Diagnosis and Treatment of Adenocarcinoma of the Prostate: II. Radical Prostatectomy Treated Patients." *Journal of Urology* 141(5): 1076–83. <https://www.ncbi.nlm.nih.gov/pubmed/2468795>.
- Tolaney, Sara M, William T. Barry, T. Dang Chau, et al. 2015. "Adjuvant

- Paclitaxel and Trastuzumab for Node-Negative, Her2-Positive Breast Cancer." *New England Journal of Medicine* 372: 134–41. <http://www.nejm.org/doi/full/10.1056/NEJMoa1406281>.
- Vittinghoff, Eric, David V. Glidden, Stephen C. Shiboski, and Charles E. McCulloch. 2012. *Regression Methods in Biostatistics: Linear, Logistic, Survival, and Repeated Measures Models*. Second Edition. Springer-Verlag, Inc. <http://www.biostat.ucsf.edu/vgsm/>.